HELIOS USER GUIDE · LEADERSHIP

Platform admin

Procedures, usage, guidelines, and your roles and access

The person who owns the HELIOS platform for the MSSP: accounts and roles, platform settings, AI configuration, the price book and every client's set-up. There should be few of them, and each should also hold a day-to-day role account for ordinary work.

Platform	HELIOS managed detection and response
Guide	01 of 11 · Platform admin
Date	1 November 2026
Version	Written against HELIOS aaea8f3
Screenshots	Taken in the HELIOS test environment with test data, signed in as a platform admin. Your screens show your own clients and data.

Contents

1. Your role at a glance
2. Roles and access
3. Getting started
4. Your screens
5. Procedures
6. Guidelines
7. Reference

> You have **54** of 54 permissions and **30** pages in your sidebar: Overview, Cases & Alerts, Approvals, Endpoints, Reports, Compliance, AI Agents, SOC Operations, Explore, Security testing, Investigations, Email & File Forensics, Response Playbooks, Runbooks, Detection Engineering, Agent Governance, Threat Intelligence, Attack Graph, MITRE ATT&CK Coverage, Tools & Integrations, Rulebooks, Rule changes, Clients, Team & Roles, Data Retention, Audit Log, API Tokens, Settings, Build & Price, Contracts & Billing.

1. Your role at a glance

The person who owns the HELIOS platform for the MSSP: accounts and roles, platform settings, AI configuration, the price book and every client's set-up. There should be few of them, and each should also hold a day-to-day role account for ordinary work.

What you are responsible for

- Set up and maintain the platform: environment, company profile, provider keys, alert channels and sign-in security.
- Manage accounts: invite people, give them the right role and queues, deactivate leavers at once, reset two-step sign-in.
- Own the commercial reference data: service plans, the price book, exchange rates.
- Configure AI: providers, agent profiles, autonomy, the kill switch, and evaluations before a model acts on its own.
- Look after client data: retention periods, legal holds, data keys (rotation, destruction at offboarding).
- Keep the audit chain intact and exportable for auditors and regulators.

A typical day or shift

1. Check the Overview checklist (Get HELIOS ready) until every item is done.
2. Review new and changed accounts on Team & Roles; deactivate anyone who has left.
3. Look at Settings → Alert channels for deliveries that failed or were skipped.
4. Verify the audit chain (Audit Log → Verify chain integrity) now and then, and before any export.

Who to turn to

Security concerns about the platform itself go to your security lead; client-facing incidents go to the SOC manager on duty.

2. Roles and access

What you can see

You work across the clients the MSSP serves. The client switcher at the top focuses every page on one client, or All clients. You see every case queue regardless of security domains.

Your access comes from your role, and HELIOS checks it on the server for every page, button and API call; hiding a button is only a convenience. A change to your role applies on your next click. Everything you do is written to the hash-chained audit log under your name.

Your pages

Overview, Cases & Alerts, Approvals, Endpoints, Reports, Compliance, AI Agents, SOC Operations, Explore, Security testing, Investigations, Email & File Forensics, Response Playbooks, Runbooks, Detection Engineering, Agent Governance, Threat Intelligence, Attack Graph, MITRE ATT&CK Coverage, Tools & Integrations, Rulebooks, Rule changes, Clients, Team & Roles, Data Retention, Audit Log, API Tokens, Settings, Build & Price, Contracts & Billing, plus My alerts and your two-step sign-in settings in the user menu.

Your permissions

What the permission lets you do	You
Cases	
See the overview dashboard (dashboard.view)	Yes
See cases in their queues (or their own tenant, for clients) (cases.view)	Yes
See every queue, regardless of assigned security domains (cases.view_all)	Yes
Acknowledge cases, change status and severity (cases.triage)	Yes
Assign cases to other analysts (cases.assign)	Yes
Escalate or hand off a case to the next tier (cases.escalate)	Yes
Comment on cases and @mention teammates (cases.comment)	Yes
Share cases with the client and external channels (cases.share)	Yes
Response	
Run containment playbooks (subject to approval) (containment.request)	Yes
Approve or reject pending containment (containment.approve)	Yes
Isolate and release endpoints directly (endpoints.contain)	Yes
Analysis	
Run log investigations (investigations.run)	Yes
Submit emails and files for forensic analysis (forensics.run)	Yes
Upload raw logs into the detection pipeline (telemetry.ingest)	Yes

Govern AI agents: identities, tokens, revocation (governance.manage)	Yes
Intelligence and detection	
Look up threat indicators (intel.view)	Yes
Edit the blocklist and indicator store (intel.manage)	Yes
See ATT&CK coverage (mitre.view)	Yes
Switch SIEM detectors on and off (detections.tune)	Yes
See response playbooks (playbooks.view)	Yes
Author and enable response playbooks (playbooks.manage)	Yes
AI agents	
See AI agent runs and configuration (ai.view)	Yes
Dispatch AI agents onto a case (ai.dispatch)	Yes
Rate AI agent output (ai.feedback)	Yes
Configure AI providers, agent profiles and autonomy (ai.configure)	Yes
Tools and endpoints	
See onboarded security tools (integrations.view)	Yes
Onboard and configure security tools (integrations.manage)	Yes
See the endpoint inventory and protection coverage (endpoints.view)	Yes
Clients and rulebooks	
See client tenants (tenants.view)	Yes
Create and edit client tenants (tenants.manage)	Yes
Read a client's rulebook: what HELIOS may do on its estate (rulebook.view)	Yes
Draft changes to a client's rulebook and its allowlist (rulebook.edit)	Yes
Put a client's rulebook in force (client admins sign; staff only on the client's recorded instruction) (rulebook.sign)	Yes
Log search and security testing	
Search the logs HELIOS holds in Explore (client users: their own) (events.search)	Yes
See security tests of a client's internet-facing assets and their reports (clients: their own) (testing.view)	Yes
Ask for a security test of a client's own assets, and import scanner results into it (testing.request)	Yes
Authorise a security test: client admins for their organisation, HELIOS approvers for the SOC (testing.authorise)	Yes
Compliance and reports	
See regulatory clocks and notices (CERT-In, DPDP) and the ISO 27001 evidence map (compliance.view)	Yes
Draft regulatory notices, record the client's approval, and file them (compliance.prepare)	Yes

Approve a regulatory notice before it is filed (the report is the client's; staff record a client's approval given elsewhere) (compliance.approve)	Yes
Read client-facing incident reports (reports.view)	Yes
SOC operations	
See the on-call rota, pages, runbooks, QA results and digests (operations.view)	Yes
Edit the on-call rota, runbooks and operations settings (operations.manage)	Yes
Review the weekly sample of what the AI did on its own (qa.review)	Yes
Commercial	
Use the Build & Price configurator (pricing.view)	Yes
See internal cost and margin (pricing.margin)	Yes
Edit the price book (pricing.manage)	Yes
See contracts, pilots and invoices (clients: their own) (billing.view)	Yes
Draft and record contracts, issue invoices and record payments (billing.manage)	Yes
Platform	
Invite users and change roles (users.manage)	Yes
Edit platform settings: company profile, provider keys, forensics thresholds (settings.manage)	Yes
Read the audit chain (audit.view)	Yes
Export the audit chain (audit.export)	Yes
Create personal API tokens for the REST API (acting with their own permissions) (api.tokens)	Yes

How your role fits with the others

Role	Group	What it is for
Platform admin	Leadership	Owns the platform: users, AI configuration, price book, everything else.
SOC manager	Leadership	Runs the SOC: every queue, approvals, AI autonomy, tool onboarding.
Tier-1 triage analyst	Security operations	Works the first-line queue: acknowledges, validates AI triage, escalates.
Tier-2 security analyst	Security operations	Primary analysis: investigations, forensics, containment requests.
Tier-3 incident responder	Security operations	Owns major incidents: approves containment, isolates endpoints, client comms.
Threat hunter	Security operations	Hunts across all tenants with investigations, forensics and intel.
Detection engineer	Engineering	Maintains playbooks, detections, tool integrations and ATT&CK coverage.

Auditor / compliance	Governance	Read-only oversight: cases, AI decisions and the audit chain.
Account manager	Commercial	Client onboarding, Build & Price quotes and margins.
Client admin	Client	Client-side: their own tenant's cases and tools; approves containment on their estate.
Client stakeholder	Client	Client-side read-only: dashboards, incident reports and endpoint coverage.

3. Getting started

Signing in

1. Open the HELIOS address you were given and sign in with your email and password. A new account gets a temporary password from its administrator.
2. If two-step sign-in is required, set it up the first time: scan the code with an authenticator app, enter a code, and keep the ten recovery codes somewhere safe. Each recovery code works once.
3. If your organisation uses single sign-on, use the sign-in button for your identity provider instead; it satisfies two-step sign-in.
4. Five wrong passwords or codes in a row lock the account for 15 minutes. An administrator can reset your password or your two-step sign-in if you lose your phone.

Finding your way

- **Sidebar**: your pages, grouped (Operate, Analyze, Intelligence, Platform, Commercial). The arrow at the bottom collapses it.
- **Search or jump to** (Ctrl K or Cmd K): find a case, page or client by name.
- **Client switcher**: All clients, or focus every page on one client.
- **The bell**: your notifications (mentions, hand-offs, approvals, pages). Choose which also reach you by email, SMS or WhatsApp under My alerts.
- **Light, dark or system theme**, and your **user menu**: My alerts, two-step sign-in, Sign out.
- A yellow **Test environment** banner means everything on that installation is test data.

Install it as an app

HELIOS can be installed from the browser (Install app in the address bar) as a desktop or mobile app with its own window and shortcuts.

4. Your screens

Each page you can open, what it is for, and how to use it with your access. Where a page is long, further screenshots show the rest.

Overview

Your starting page. It shows what needs you now, and it changes with your role: the SOC sees live posture, queues, SLAs and AI activity; account managers see the commercial picture; client users see their own organisation.

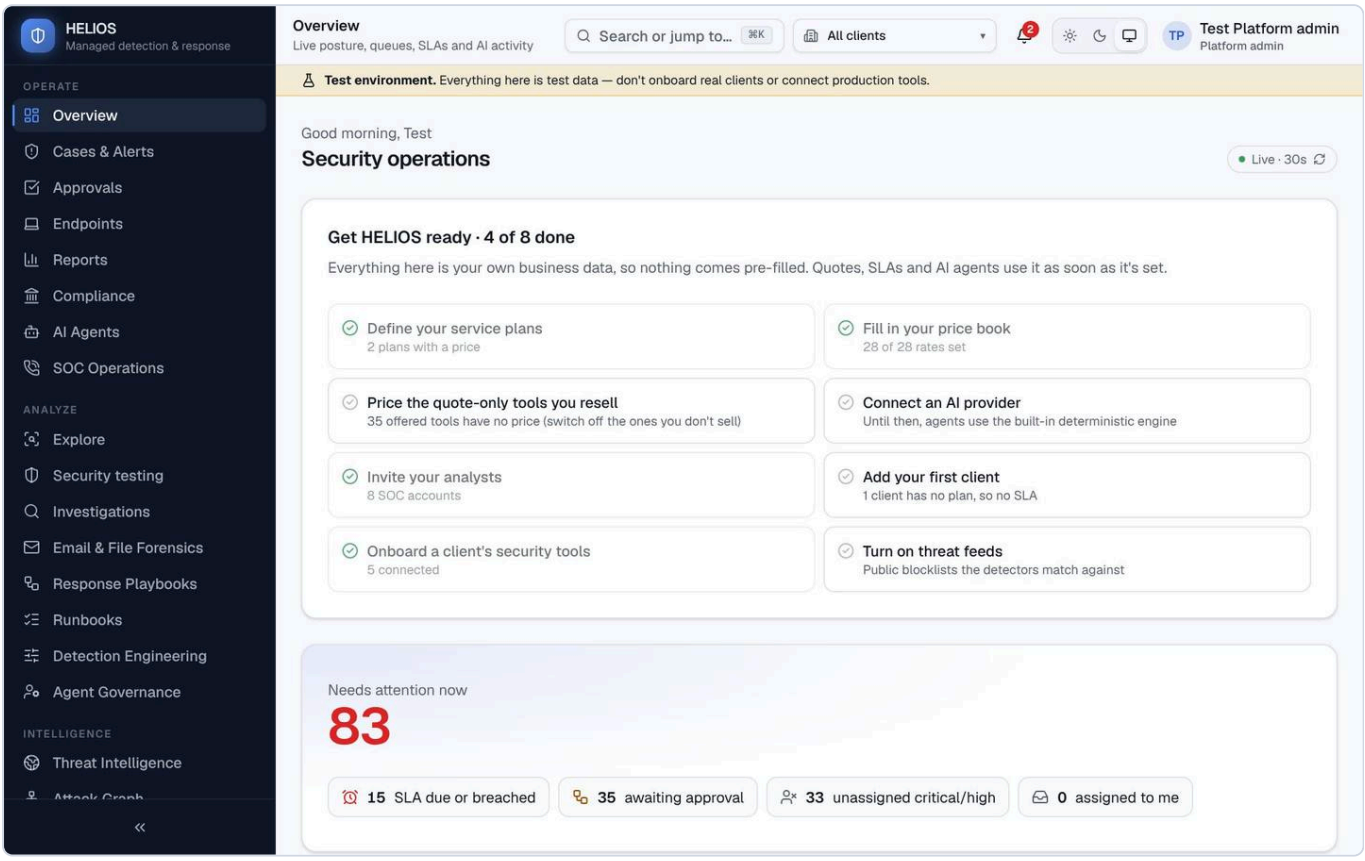


Figure 1. Overview, as seen by a platform admin (test environment).

How to use it

- Needs attention now adds up what is due: cases whose SLA is due or breached, containment awaiting approval, unassigned critical and high cases, and cases assigned to you. Click a figure to open the matching list.
- The Live badge means the figures refresh on their own every 30 seconds; the arrow refreshes now.
- Get HELIOS ready is the first-run checklist: service plans, price book, prices for quote-only tools, an AI provider, analysts, the first client, its security tools and threat feeds. Each item links to where it is done and ticks itself once it is.

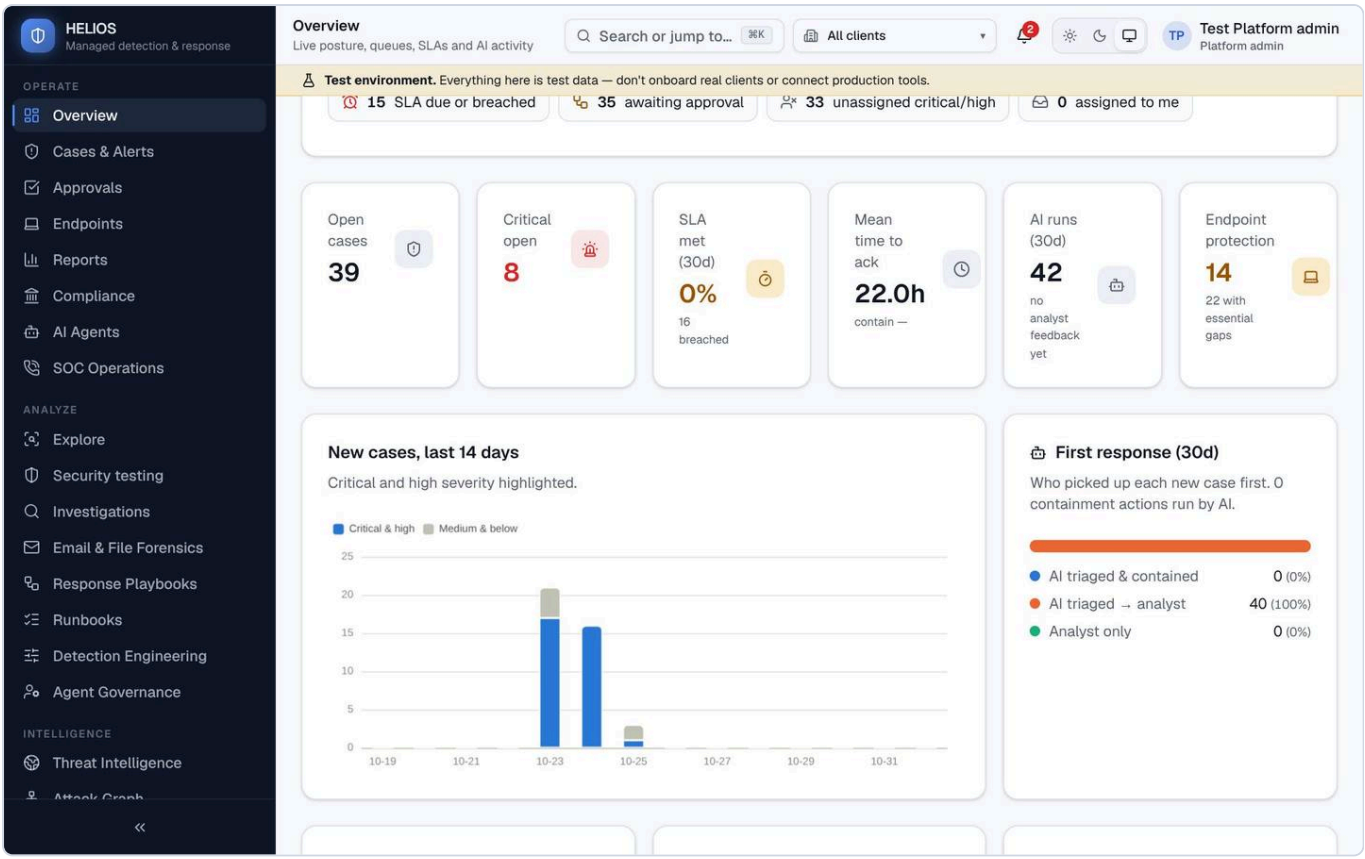


Figure 2. Overview, further down the page.

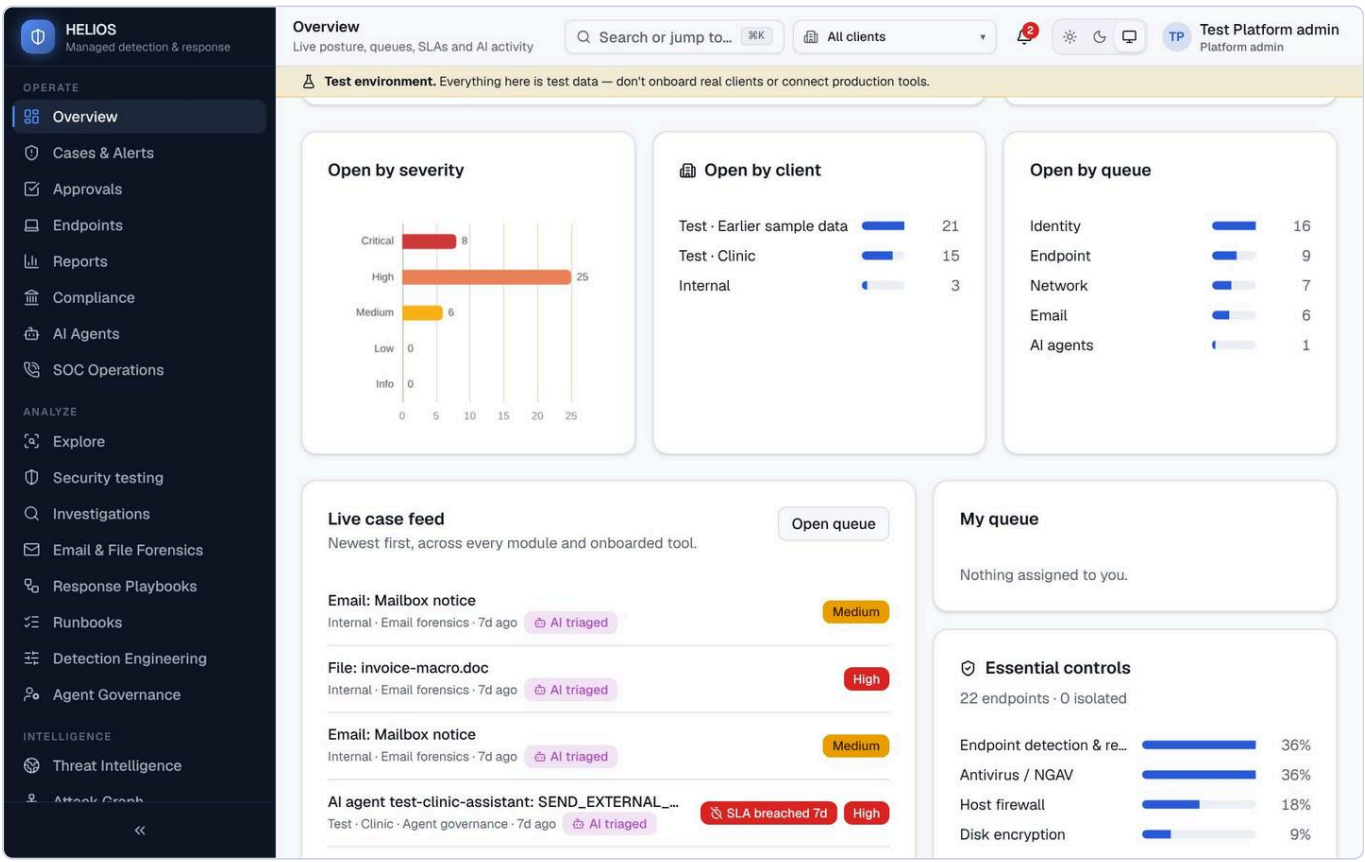


Figure 3. Overview, further down the page.

Cases & Alerts

Every case HELIOS has opened: detections from your clients' telemetry, vendor alerts, uploaded investigations, email and file analysis, and governed AI agents. Correlated detections for the same user, host or IP join one case instead of opening many.

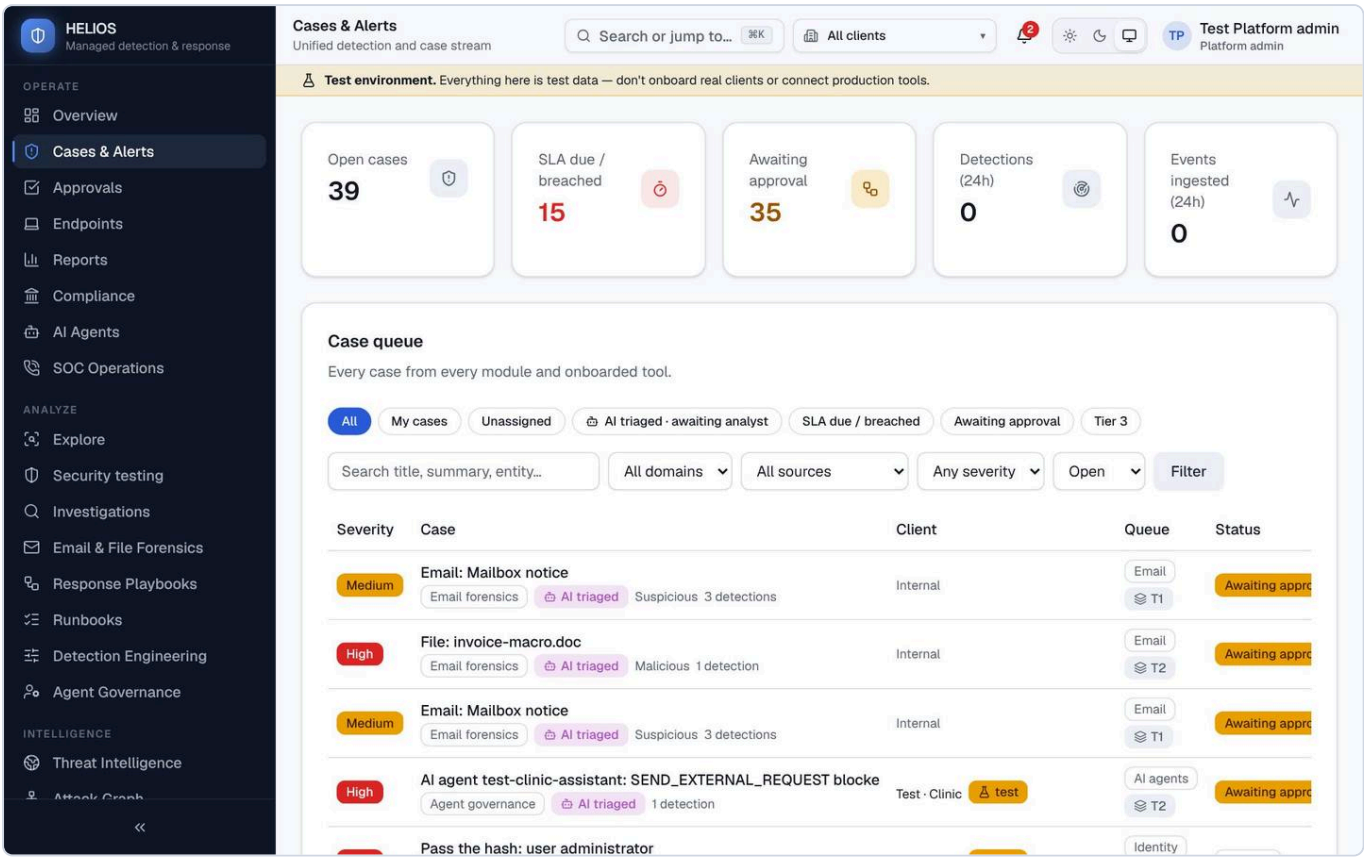


Figure 4. Cases & Alerts, as seen by a platform admin (test environment).

How to use it

- Filter by status, severity, source, client and assignee, or search by title. Open a case by clicking its row.
- You see every queue regardless of security domains.
- The SLA badge counts down to the acknowledgement deadline from the client's plan; a red badge means it has been breached.

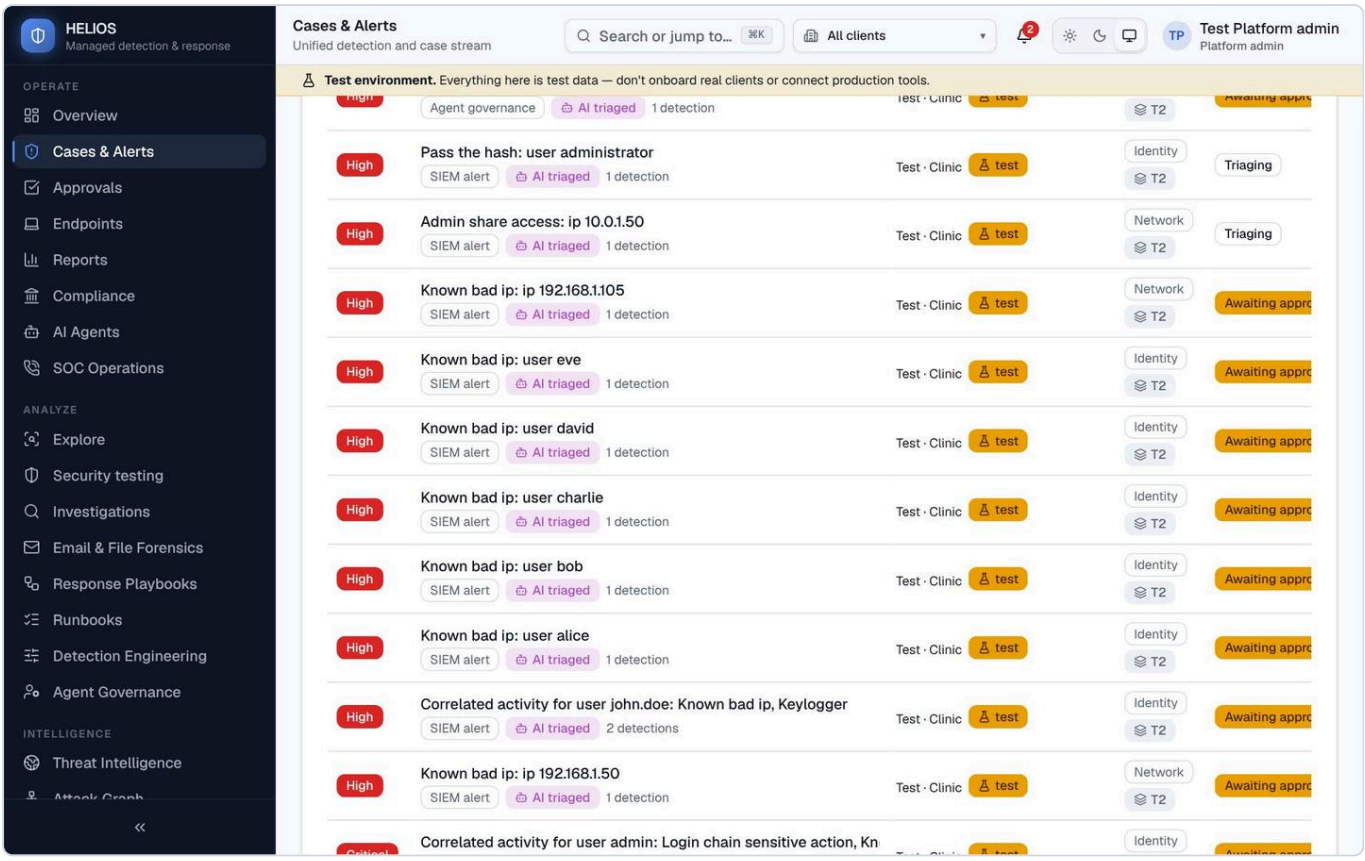


Figure 5. Cases & Alerts, further down the page.

A case

Everything about one case on one page: what was detected, what the AI made of it, what has been done

and what is waiting, the evidence, and the discussion.

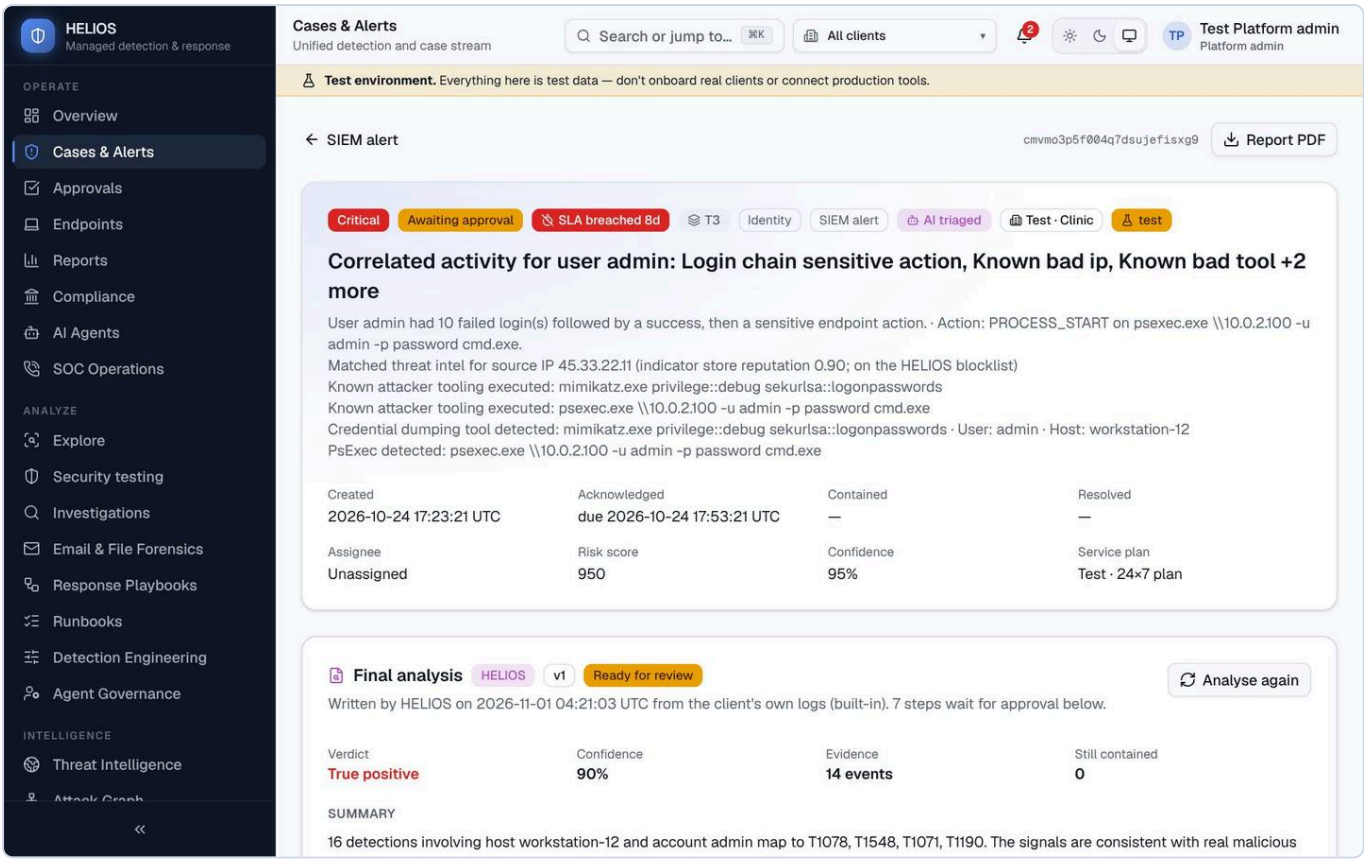


Figure 6. A case, as seen by a platform admin (test environment).

How to use it

- The header shows severity, status, the SLA, the tier (T1 triage, T2 analysis, T3 incident response), the security domain, the source, the client and whether AI triaged it. Below are the dates (created, acknowledged, contained, resolved), the assignee, the risk score, the confidence and the client's service plan. Report PDF downloads the incident report.
- Final analysis (HELIOS) appears once HELIOS has worked the case on its own: it collects every log the client's own tools hold about it, investigates, and writes the verdict and confidence, the summary and full report, the root cause, the scope, what happened, the findings and techniques, what is contained, the logs collected by type (and the types the client doesn't send), and open questions.
- Lifting containment: HELIOS never lifts containment on its own. Each release is proposed with its conditions (for a laptop: a clean EDR scan, persistence removed, passwords reset, 24 quiet hours); approve it when they hold.
- Remediation plan: the steps HELIOS can carry out (reset a password, purge a phishing mail, block a C2 address) wait as actions to approve or reject; the rest are tasks with an owner.
- Mark reviewed, with an optional note, records the SOC's review of the final analysis; Analyse again asks HELIOS for a fresh one (after new evidence, say).
- AI first response shows the agent's verdict, confidence and engine, its summary, its key findings (each citing evidence: D detections, E events, I indicators, T timeline, C context) and what it asks the analyst to check. Initial containment (isolate a host, block an IP, quarantine a file, block a sender, collect logs) runs on its own when the client's approved rules allow it; anything further is proposed for a person.
- Agree or Disagree records whether the analysis was right; a disagreement asks what it got wrong. These ratings drive the AI's autonomy ladder.

- Dispatch agent runs another agent on the case: the triage first responder, the phishing analyst, the threat-intel enricher, the investigation agent (works the case with read-only tools) or the client report writer.
- In Triage: Acknowledge stops the SLA clock; Status and Severity change the case; Assignee gives it to someone.
- Watch case sends you its updates. Escalate hands the case to another tier or person with a note; Share makes it visible to the client or sends it to an external channel.
- Page on-call (in the On call card) pages whoever is on call about this case, with what you need from them.
- Response actions lists containment the AI proposed and people requested, each with its status. Approve or Reject what is awaiting approval; the client's rulebook and the response policy still apply.
- Request containment asks for an action (isolate host, block IP, disable account, kill process, quarantine file, block sender, collect logs, purge email, reset password, collect an investigation package…) against a target, with a reason shown to the approver.
- Indicators shows the case's IPs, domains, URLs, hashes and addresses with their reputation; the arrow re-checks one. Related cases lists earlier or concurrent cases sharing hosts, accounts, indicators or the same agent.
- Regulatory reporting shows whether CERT-In's 6-hour report or the DPDP breach notices apply, the incident types HELIOS proposes, and the notices' state.
- Runbooks attach themselves to matching cases; tick each step or skip it with a reason. A case can't be closed as a true positive while a required step is open.
- The discussion takes comments and @mentions. Internal comments stay with the SOC; client-visible ones reach the client.
- Closing: set the status to Resolved or False positive, then give a disposition (true positive, benign true positive, false positive) and a reason.

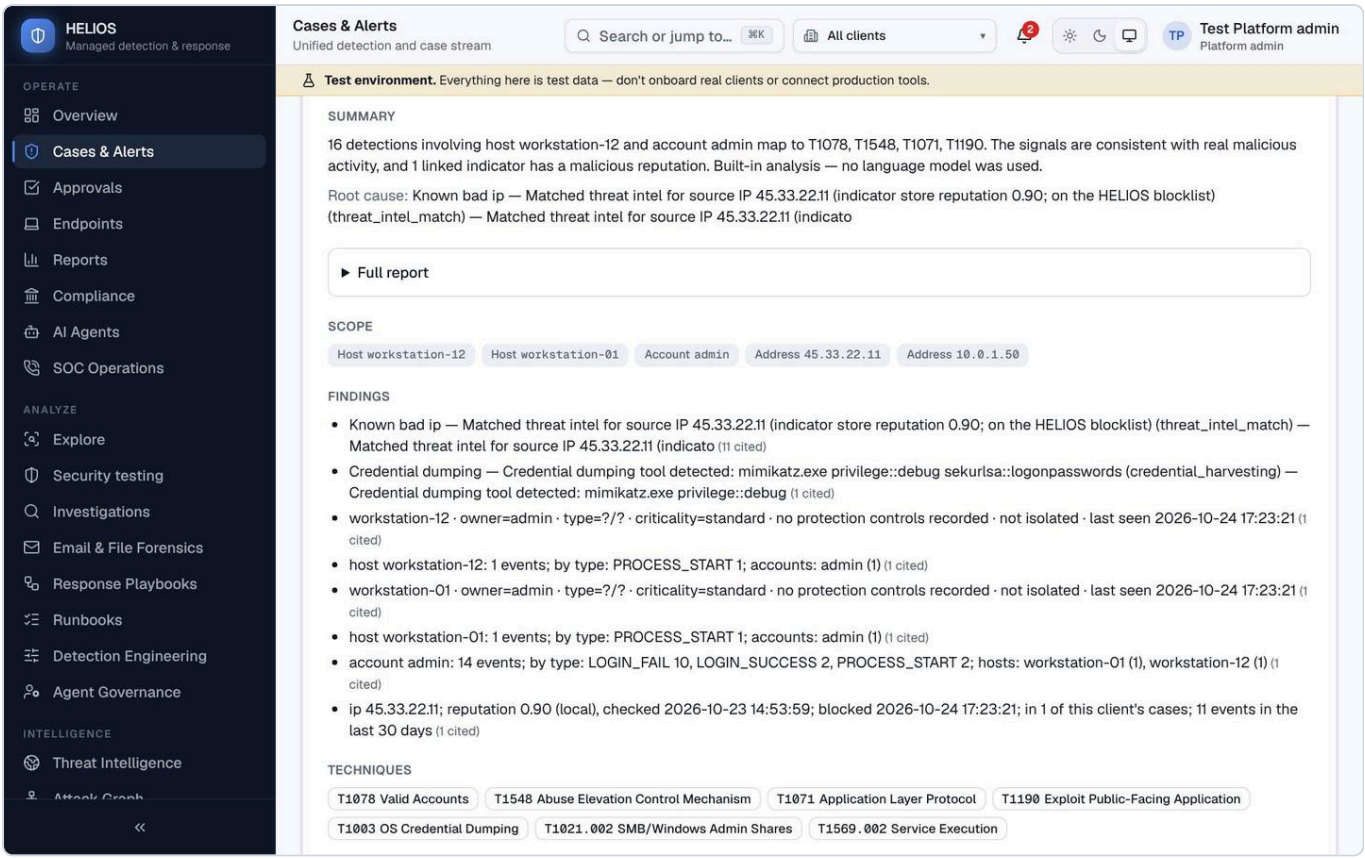


Figure 7. A case, further down the page.

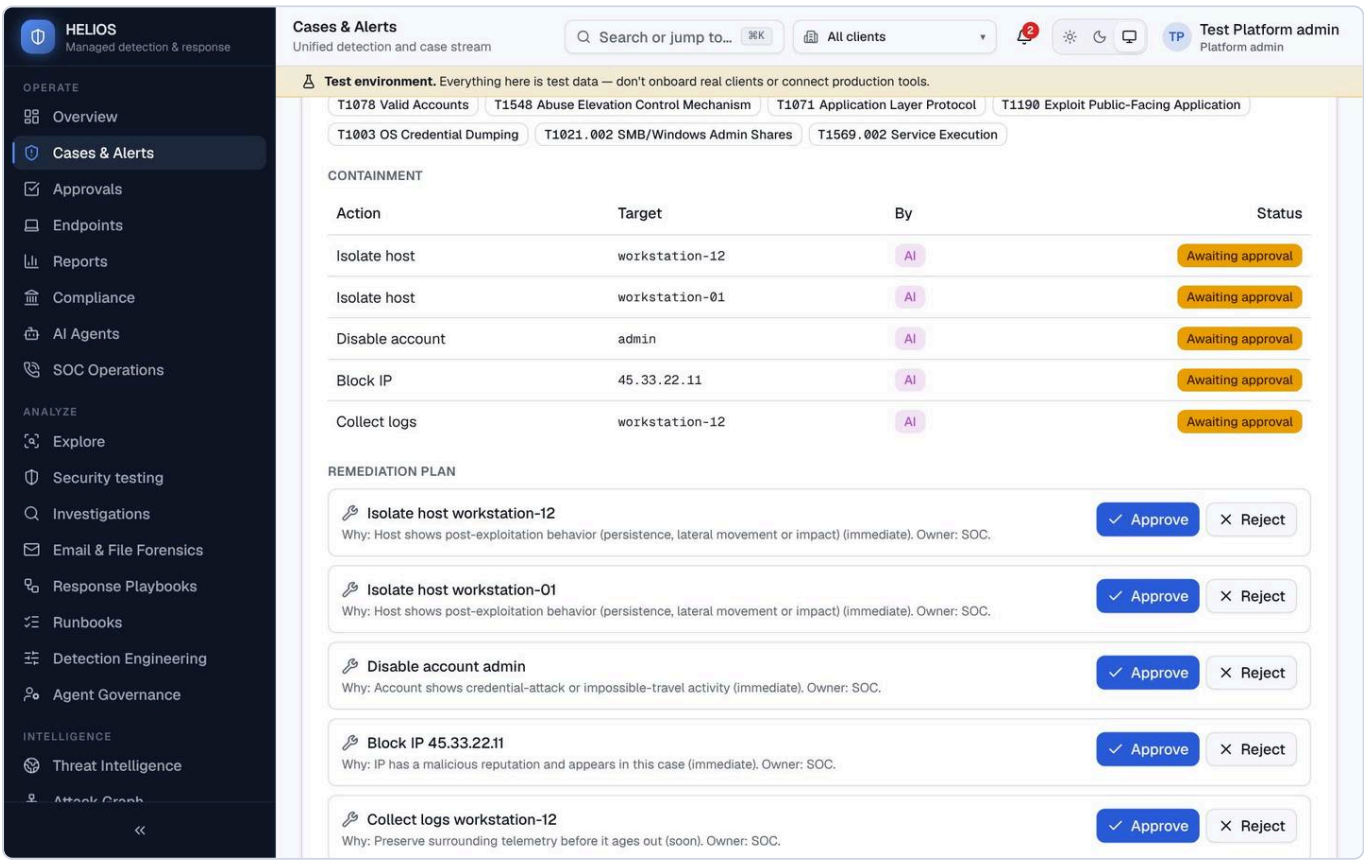


Figure 8. A case, further down the page.

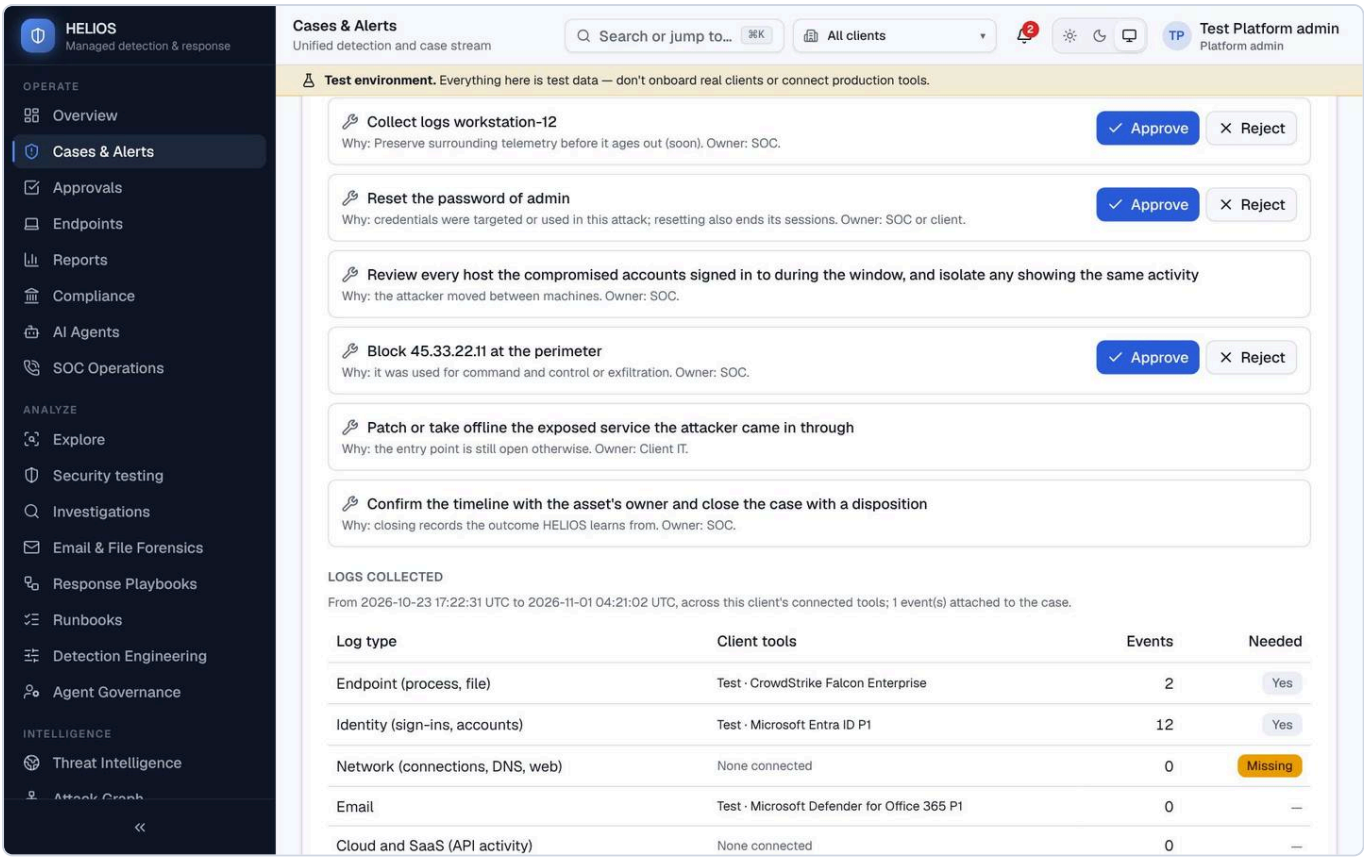


Figure 9. A case, further down the page.

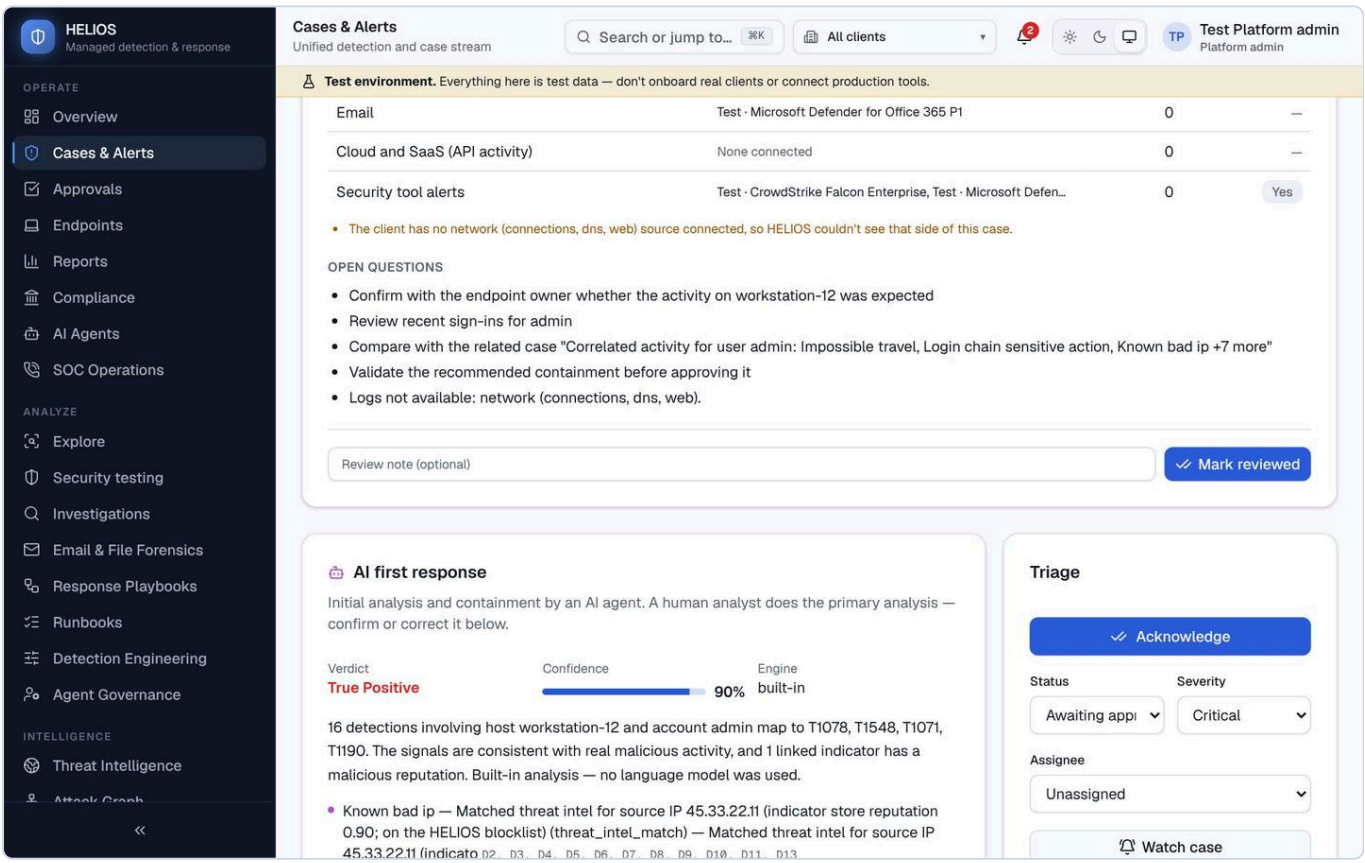


Figure 10. A case, further down the page.

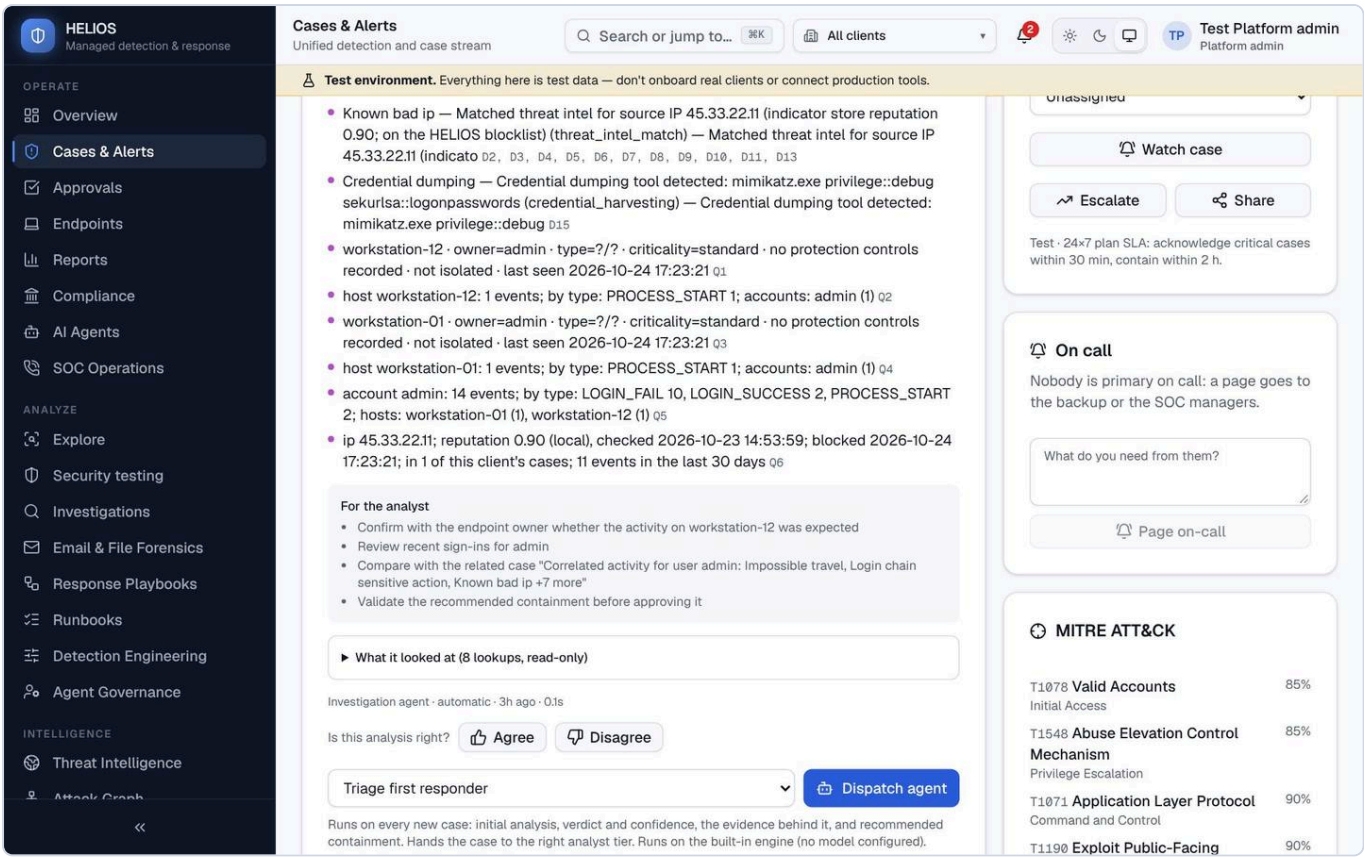


Figure 11. A case, further down the page.

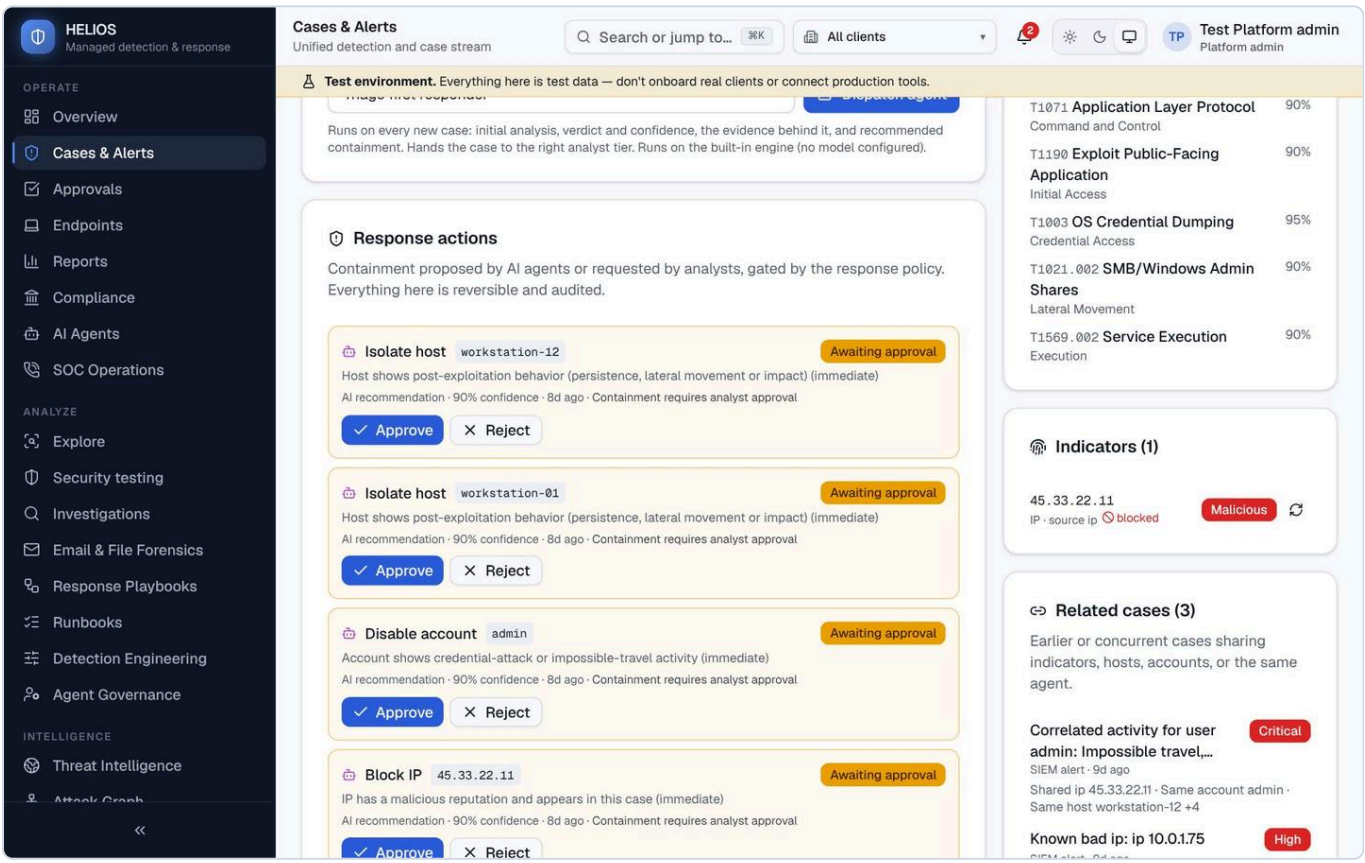


Figure 12. A case, further down the page.

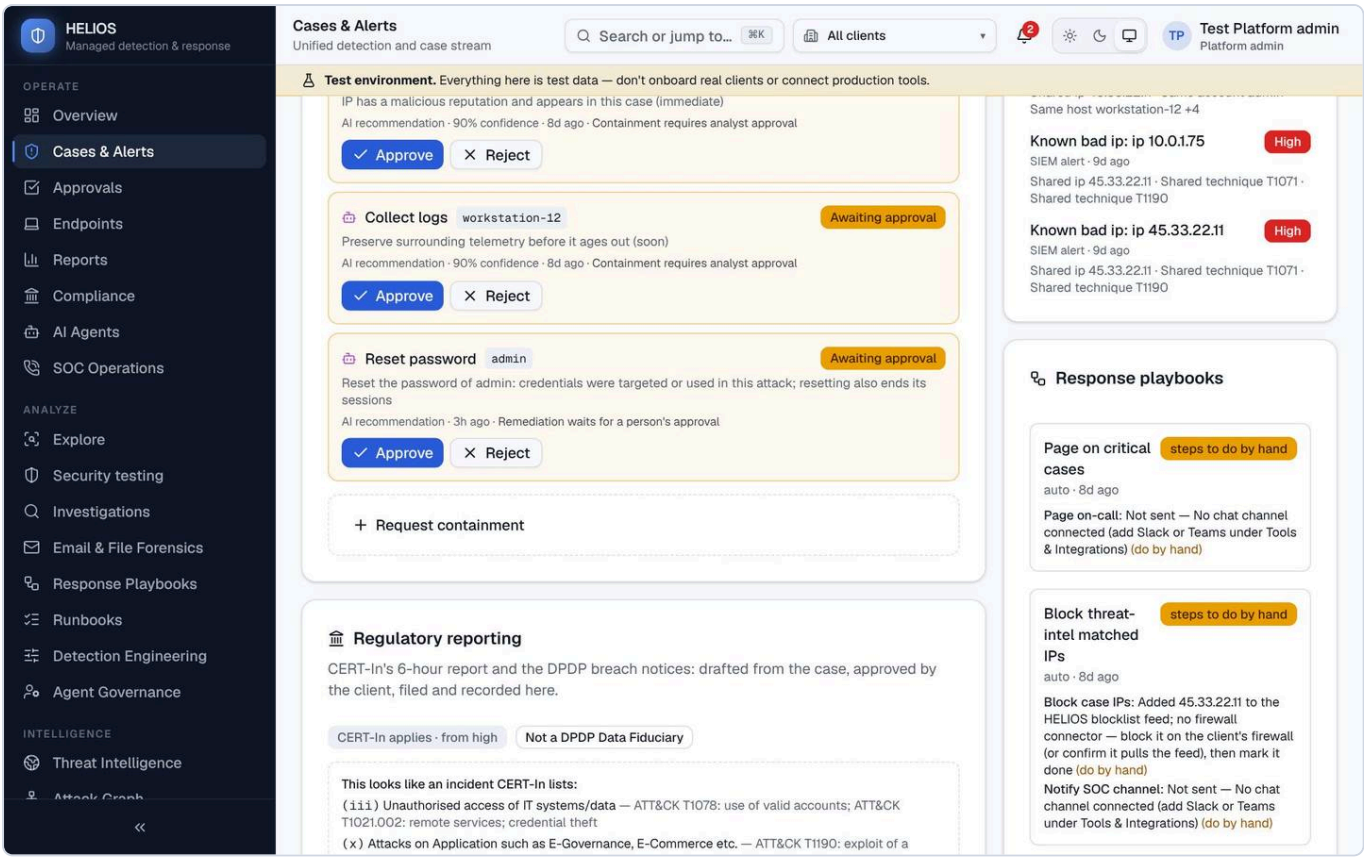


Figure 13. A case, further down the page.

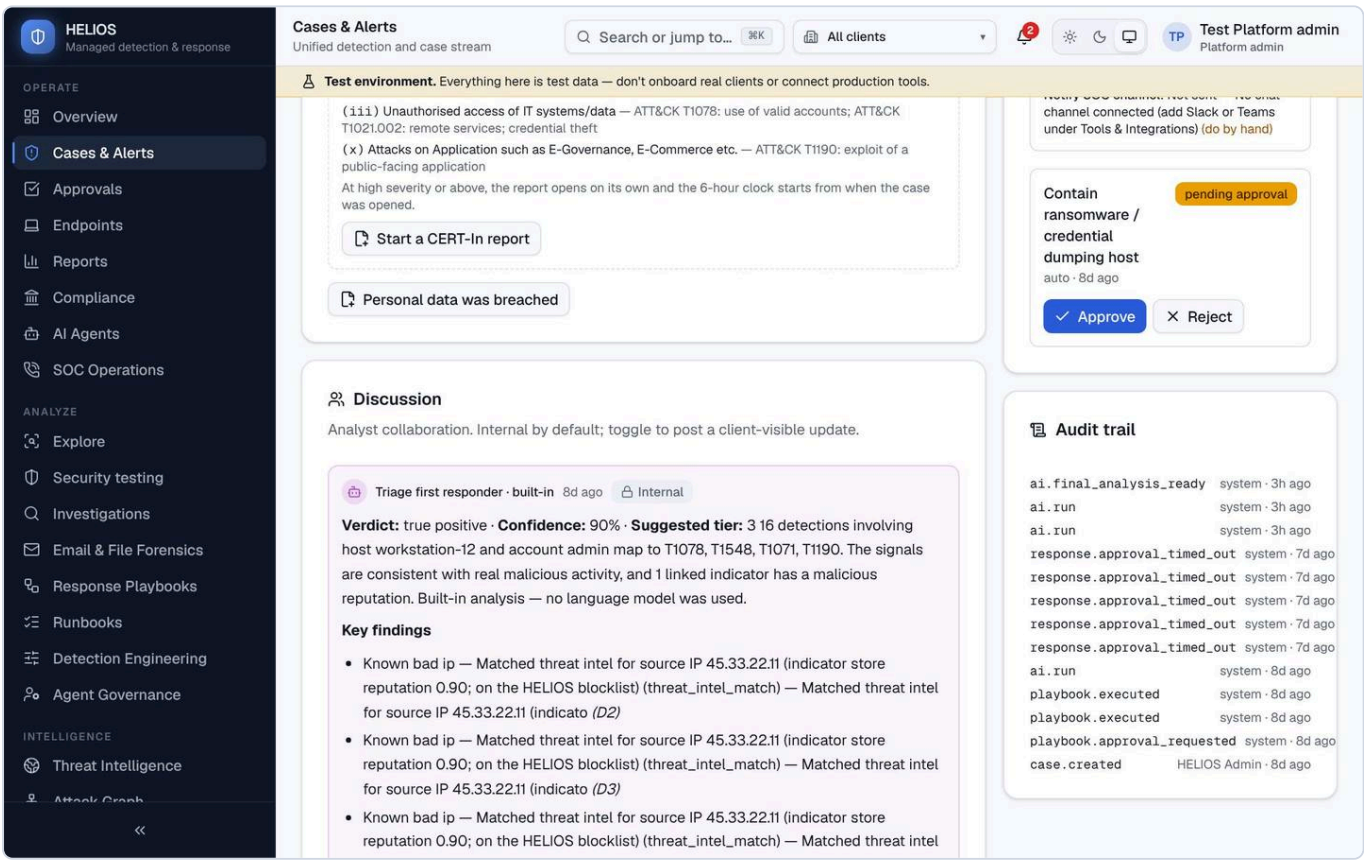


Figure 14. A case, further down the page.

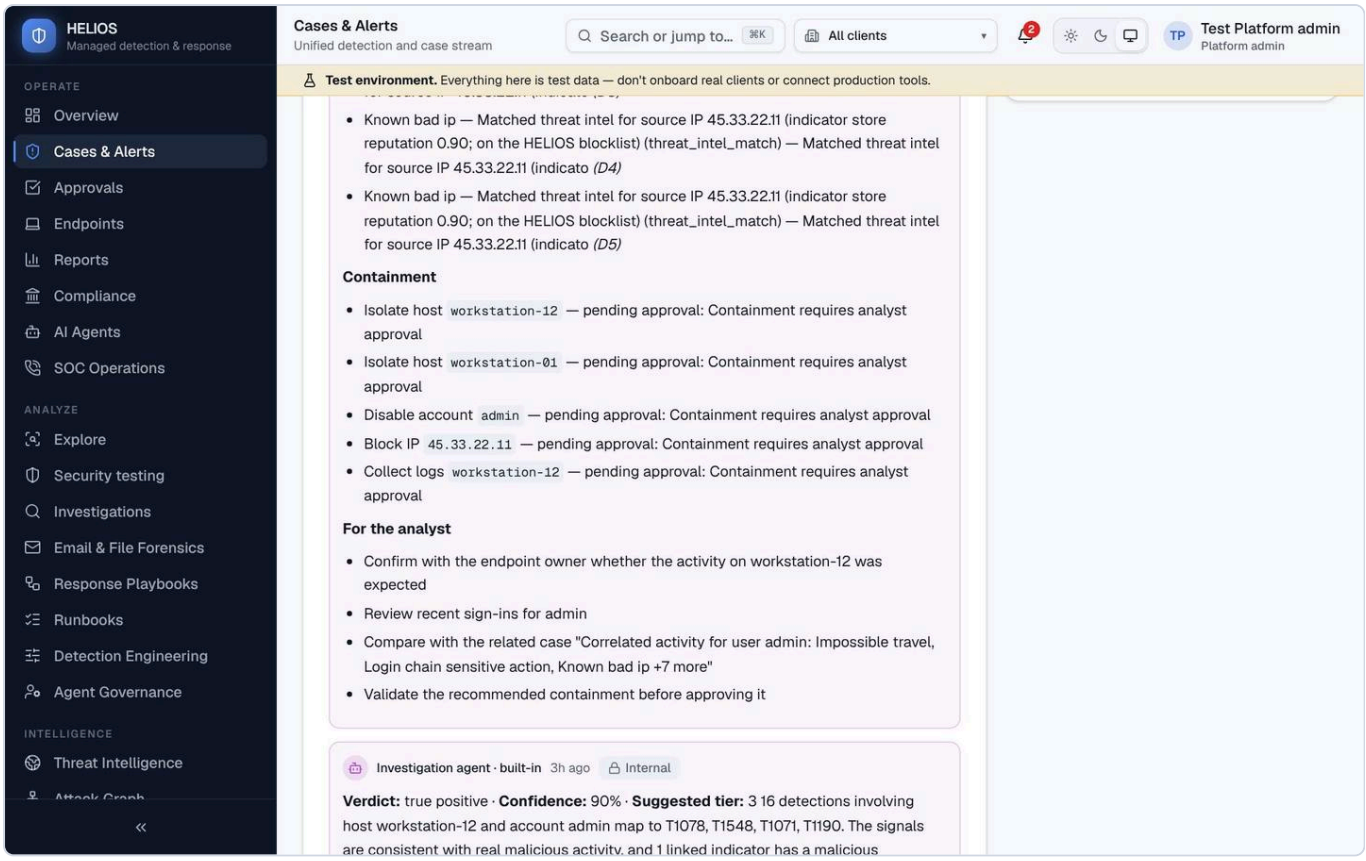


Figure 15. A case, further down the page.

Approvals

Everything waiting on someone's yes, in your scope: containment the response policy held, HELIOS's final analyses to review, playbook runs that need approval, regulatory notices waiting on the client, and

rule changes to approve.

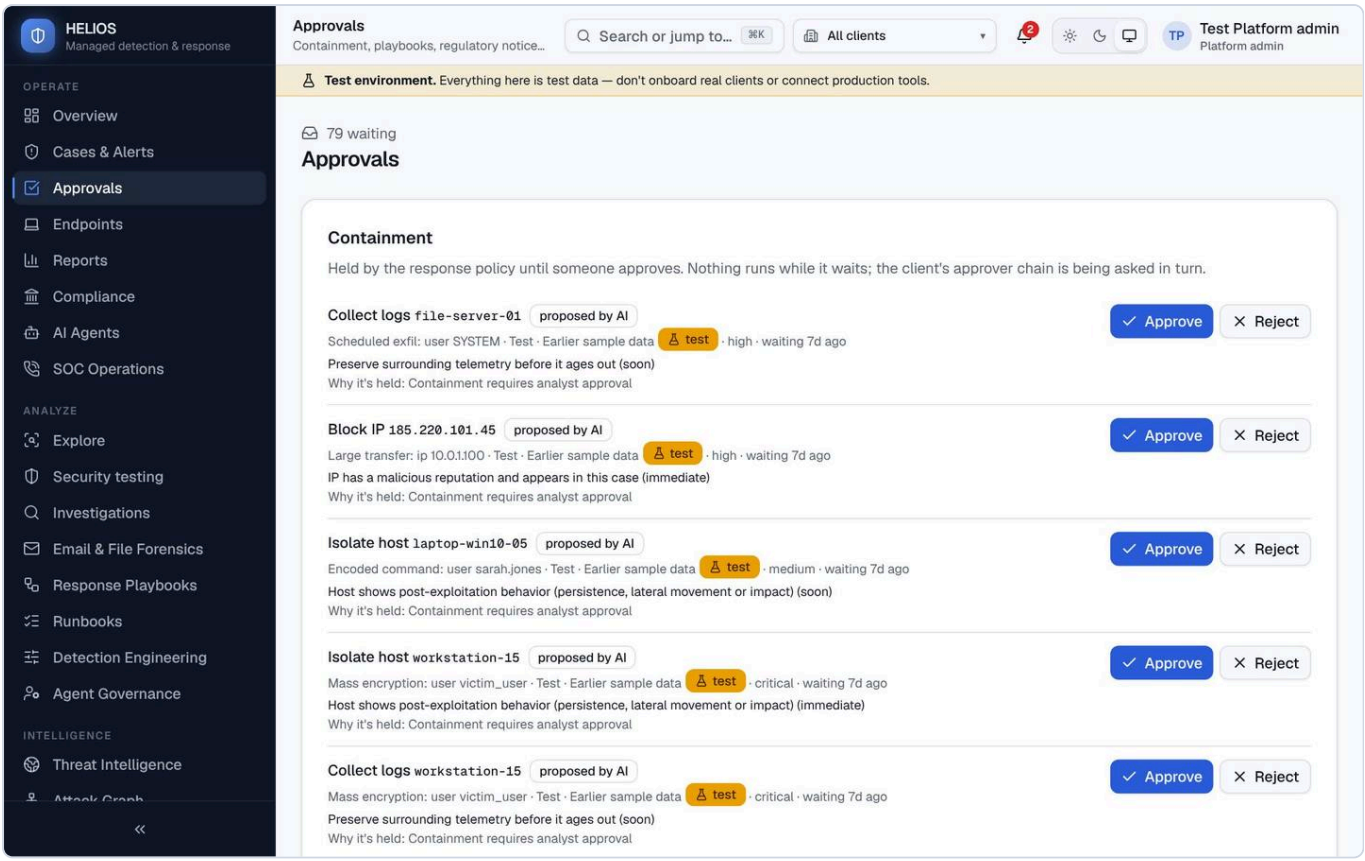


Figure 16. Approvals, as seen by a platform admin (test environment).

How to use it

- Each item says what would happen, against what, why it was held, who proposed it, the client, the severity and how long it has waited. Approve or Reject right there.
- Nothing held runs while it waits. The client's approver chain is asked in turn (in the app, by Slack or Teams, email, SMS or WhatsApp); when the chain runs out, the action holds or becomes a suggestion, as the client's rulebook says.
- A two-person rule in the rulebook needs two different approvers; one approval leaves it waiting for the second.
- Final analyses to review lists the analyses HELIOS wrote on its own; Open goes to the case, where the lifting and remediation steps are approved and the analysis is marked reviewed.
- Rule changes to approve lists changes to the rules containment runs under; Review opens Rule changes.

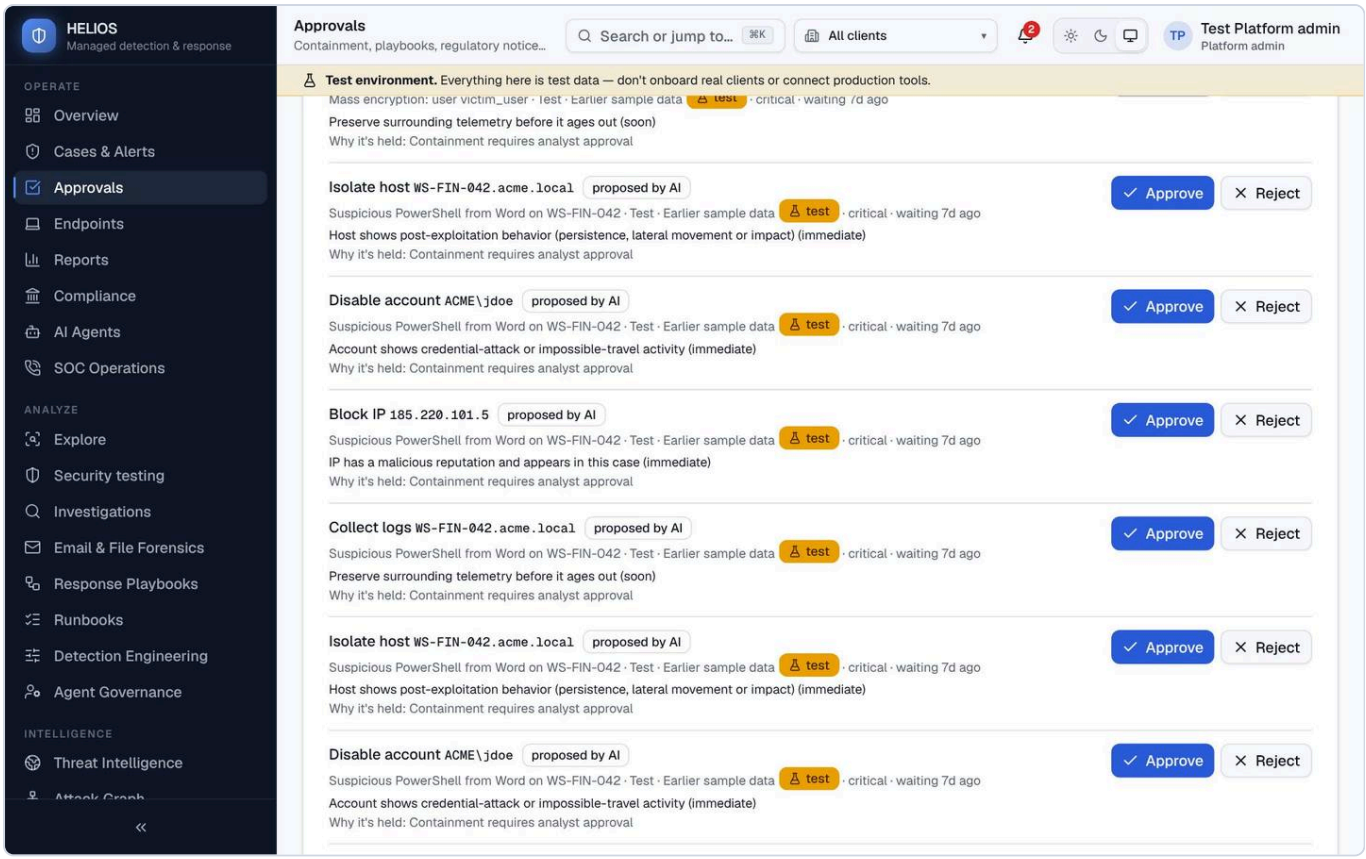


Figure 17. Approvals, further down the page.

Endpoints

The endpoint inventory and the whole attack surface: devices, servers, cloud, network, identity, IoT/OT

and applications, with how well each is protected.

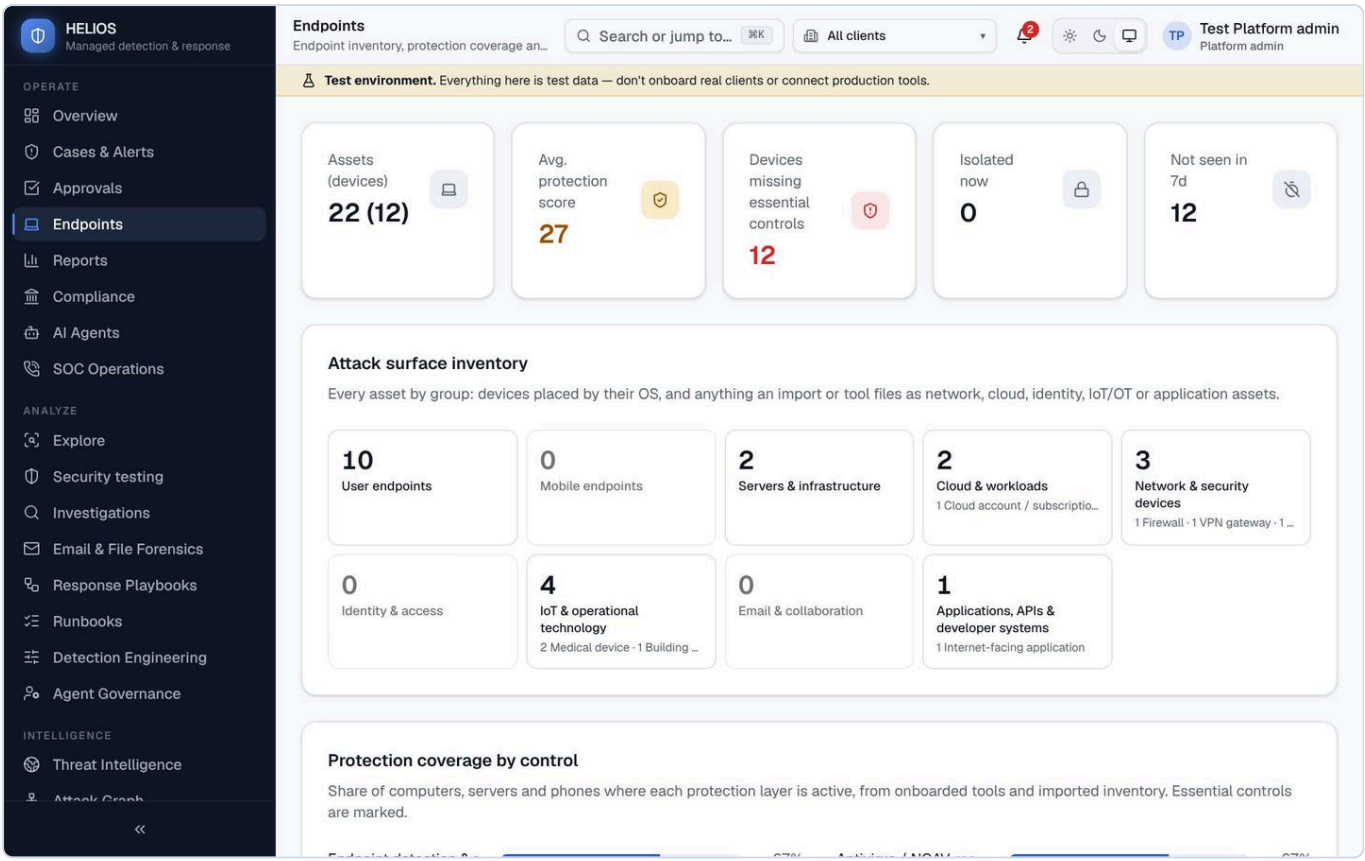


Figure 18. Endpoints, as seen by a platform admin (test environment).

How to use it

- The tiles count assets, the average protection score, devices missing essential controls, devices isolated now and devices not seen for seven days.
- Attack surface inventory groups every asset; Protection coverage by control shows the share of computers, servers and phones where each protection layer (EDR, antivirus, disk encryption, MFA...) is active, with the essential ones marked.
- Isolate or Release a device from its row. With an EDR connector HELIOS asks the vendor and confirms the device's state; without one it records the change and asks you to make it in the EDR console. The client's rulebook still decides: a never-isolate asset can't be isolated, a crown jewel needs approval.

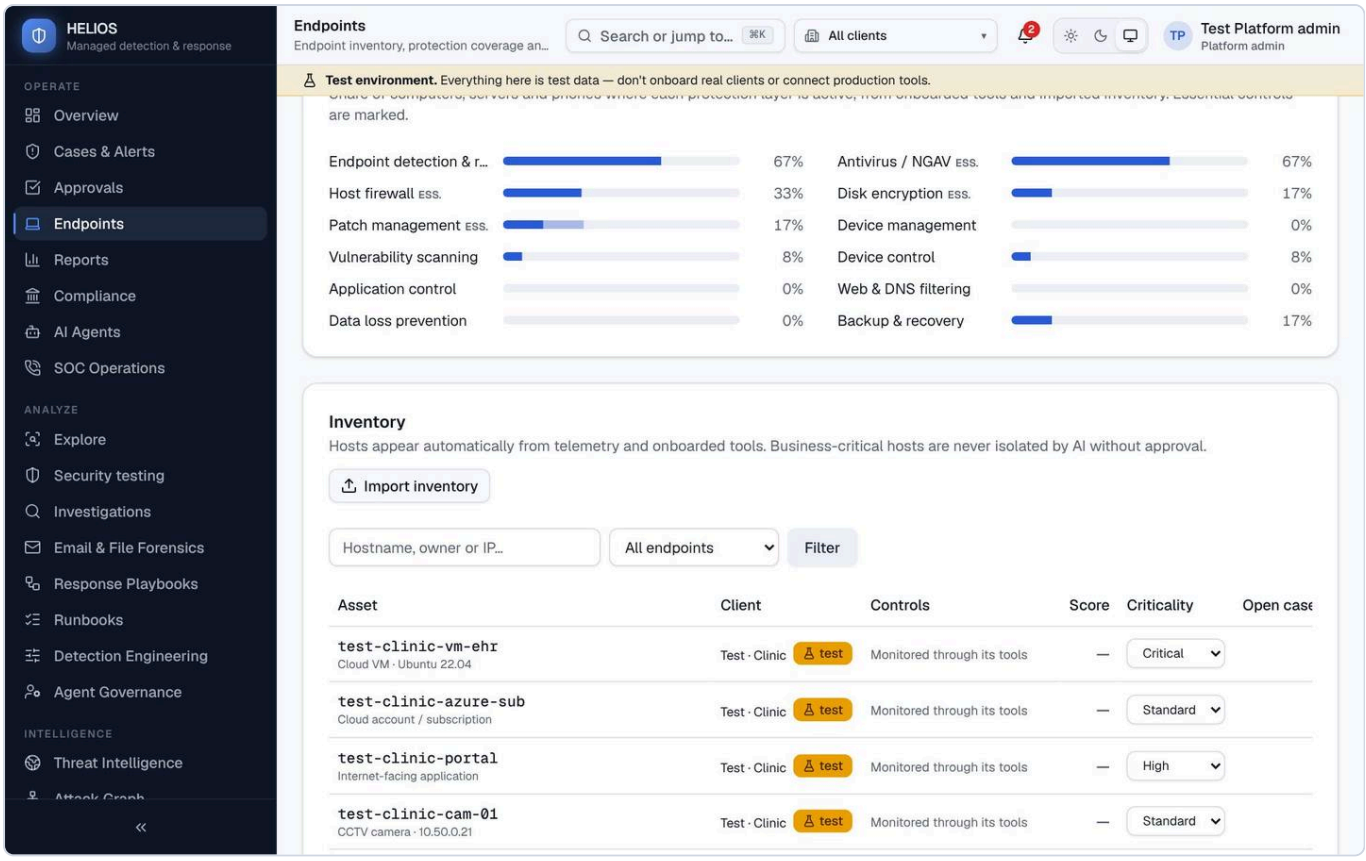


Figure 19. Endpoints, further down the page.

Reports

How the service performed over a period: volume, response times, SLA compliance, AI performance and

analyst workload, each compared with the period before.

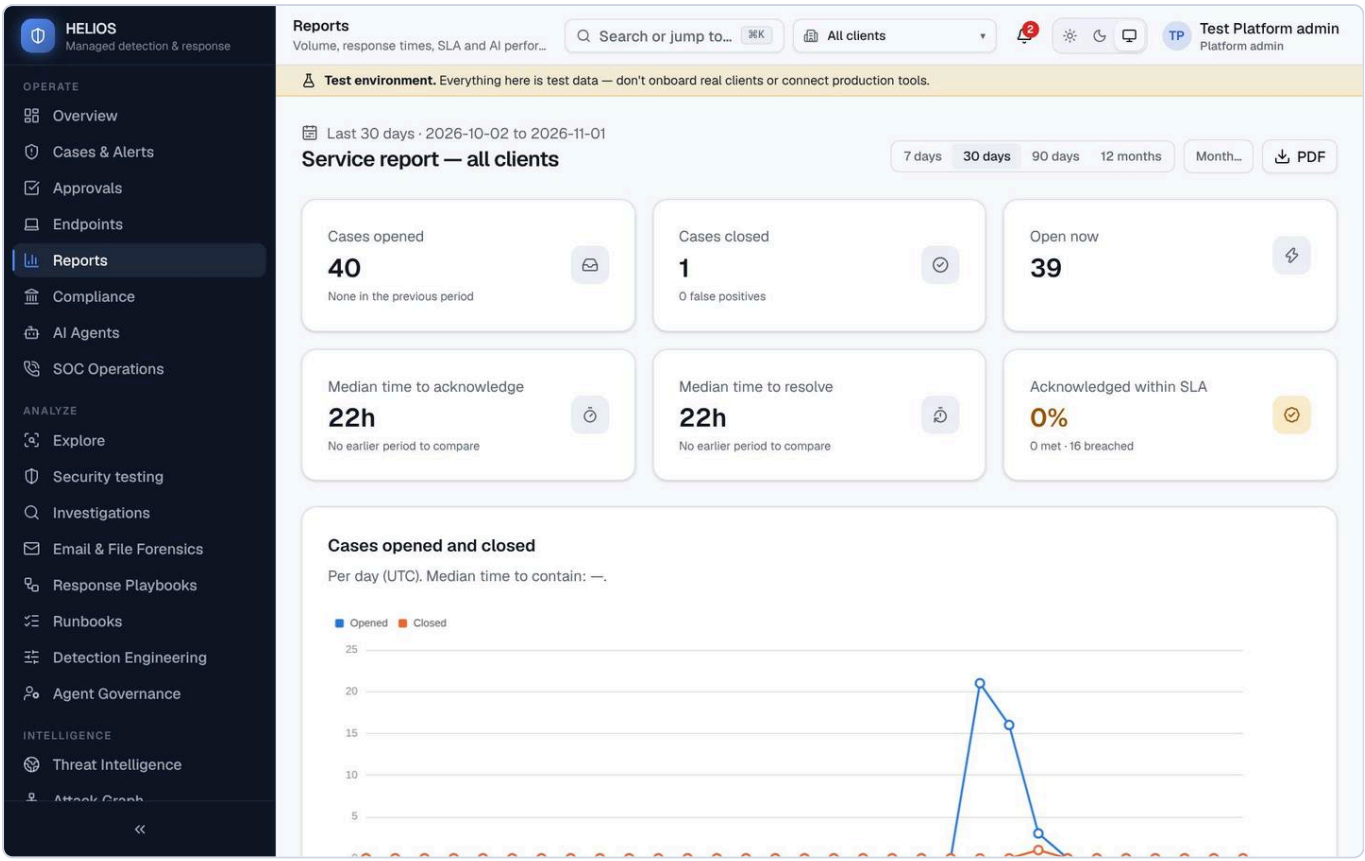


Figure 20. Reports, as seen by a platform admin (test environment).

How to use it

- Pick a period or a month. The figures include median time to acknowledge, contain and resolve, acknowledgement SLA compliance, the mix by severity, source, domain and client, top ATT&CK techniques and indicators, AI runs and analyst agreement, and containment outcomes.
- Download the service report as a PDF for the client.

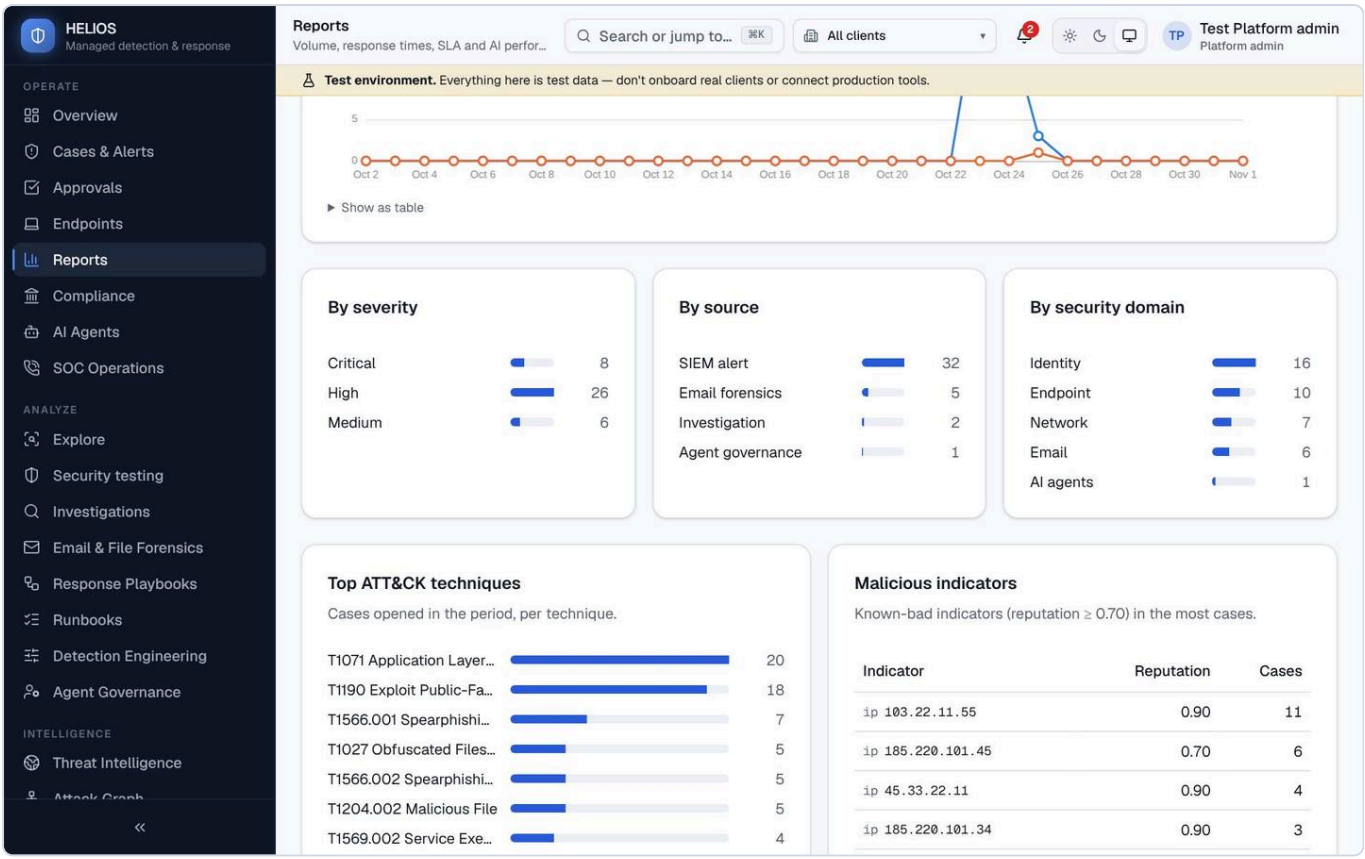


Figure 21. Reports, further down the page.

Compliance

Regulatory clocks and evidence: CERT-In's 6-hour incident reports, the DPDP breach notices, and the

ISO/IEC 27001 evidence map.

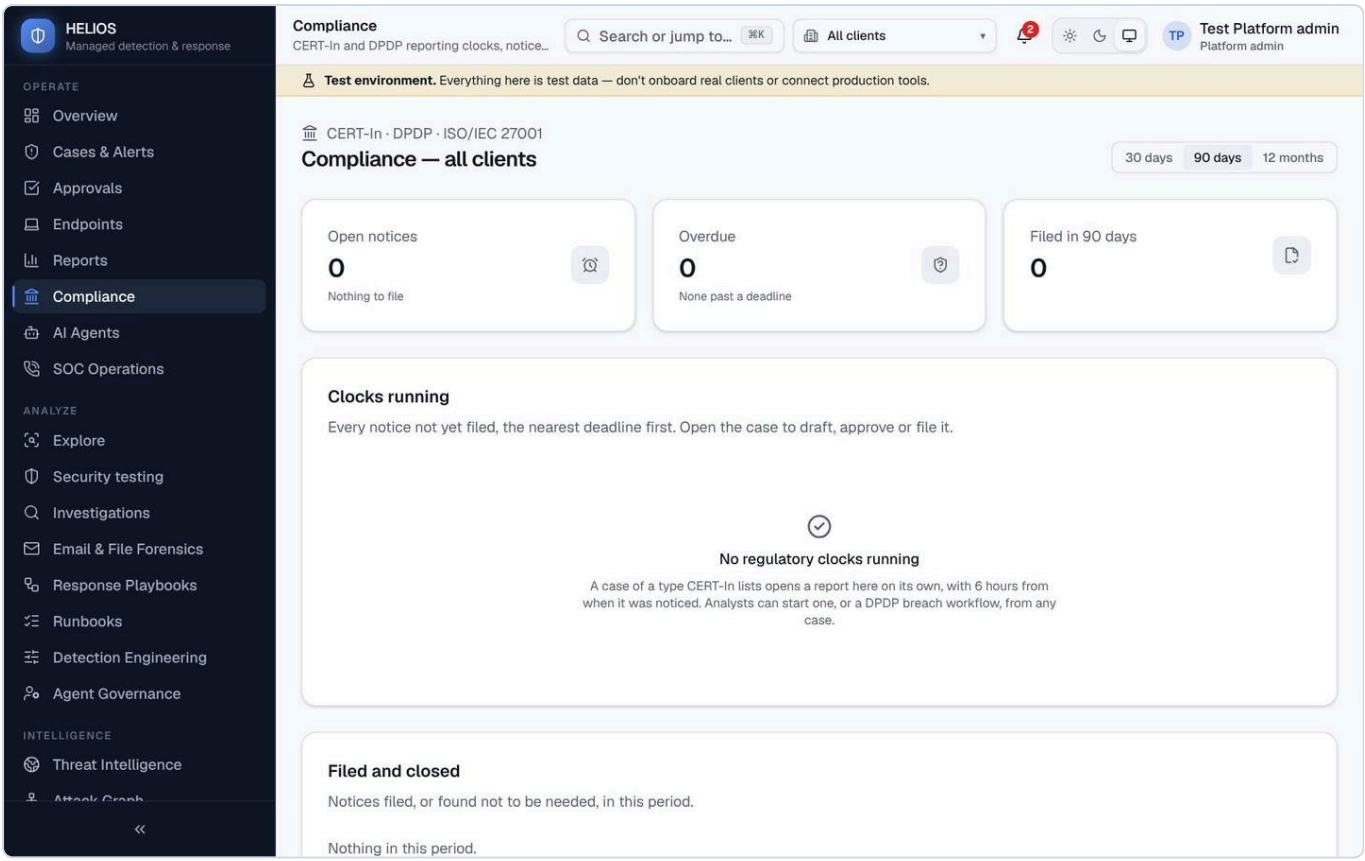


Figure 22. Compliance, as seen by a platform admin (test environment).

How to use it

- Clocks running lists every notice not yet filed, nearest deadline first. CERT-In: incidents of the 20 listed types are reported within 6 hours of being noticed. DPDP: the Board and each affected person are told without delay, and the Board gets a detailed report within 72 hours.
- Open the case to draft, approve or file a notice. Each notice is drafted from the case, approved by the client, then filed (emailed to CERT-In with its PDF, or recorded with the regulator's reference).
- A notice goes to the client for approval before it is filed: approve it here or from Approvals. An edit after approval sends it back.
- The ISO 27001 view maps 15 Annex A controls to what HELIOS's records show over the period, and shows gaps (controls with nothing behind them) plainly.
- Switch the period (30 days, 90 days, 12 months) at the top.

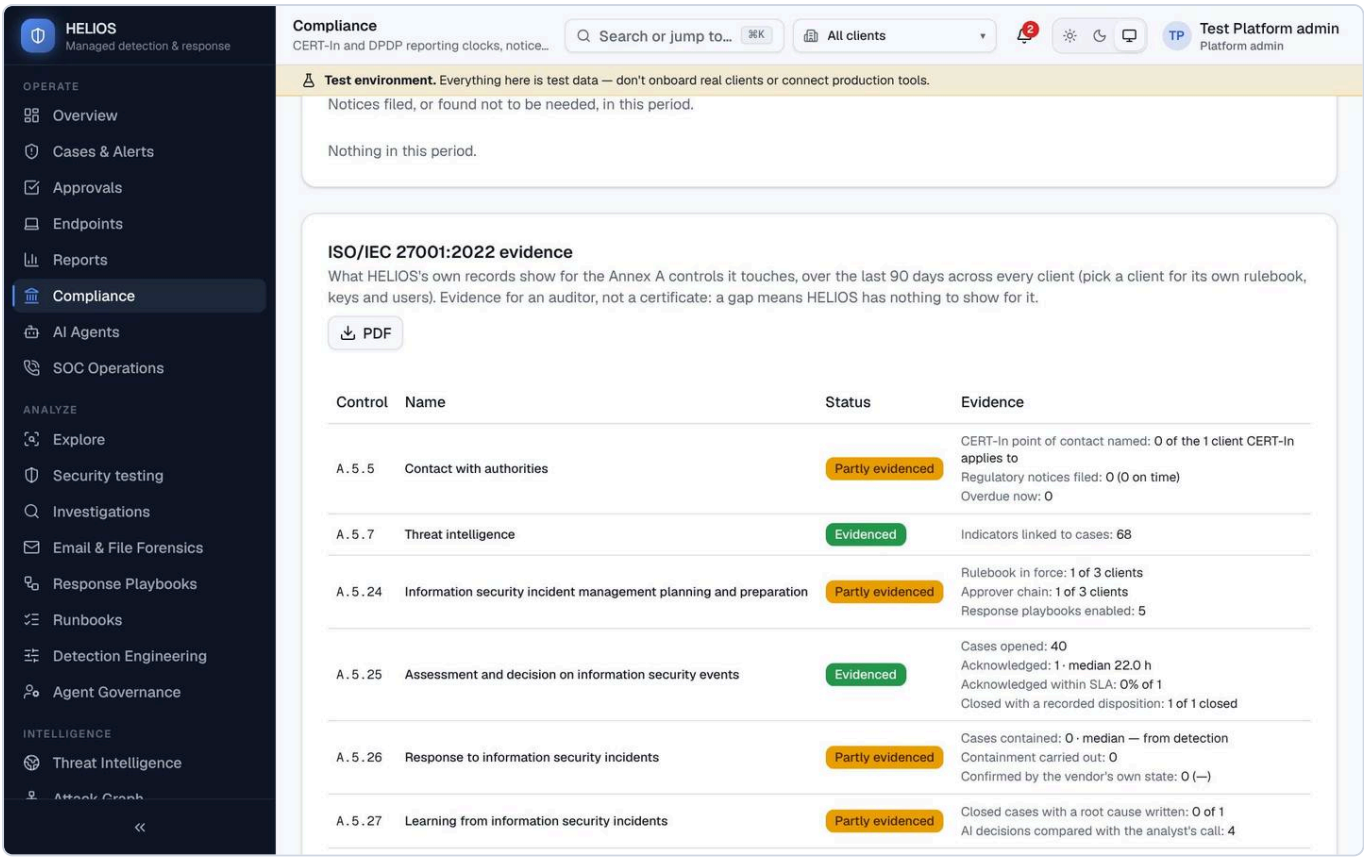


Figure 23. Compliance, further down the page.

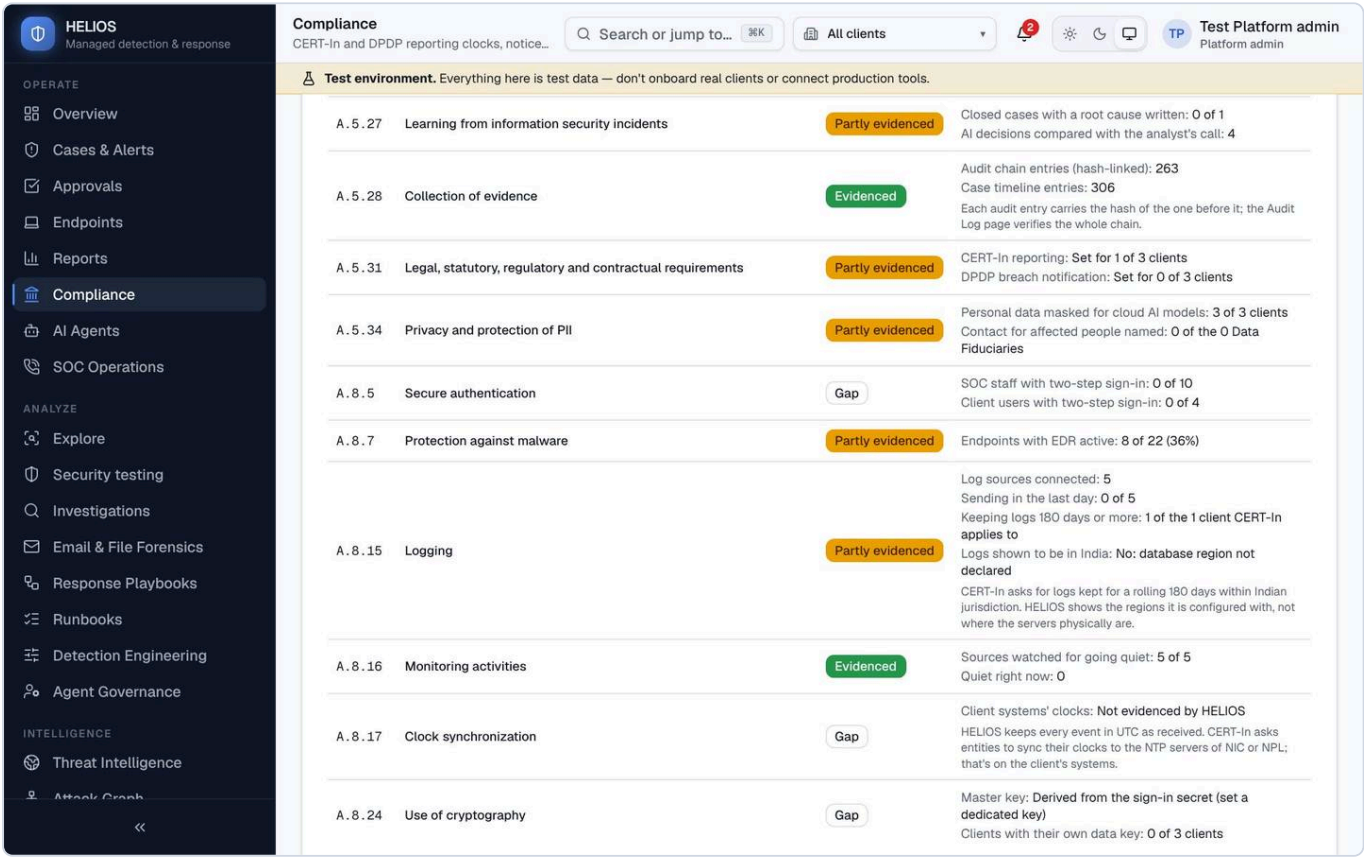


Figure 24. Compliance, further down the page.

AI Agents

The AI agents that work cases: what they did, how far they may act on their own, which models run them, and how those models scored.

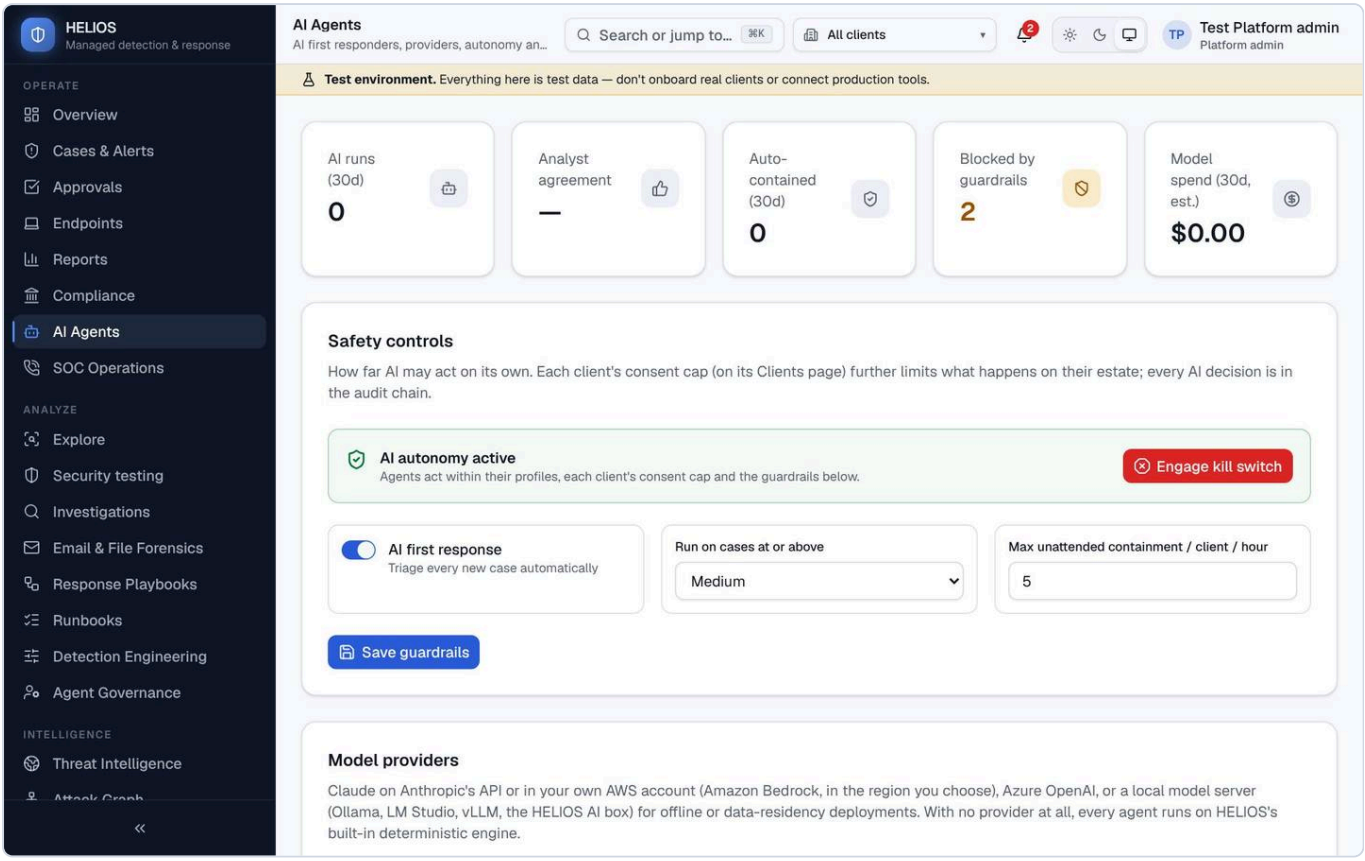


Figure 25. AI Agents, as seen by a platform admin (test environment).

How to use it

- The tiles show AI runs in the last 30 days, analyst agreement, actions contained automatically, actions the guardrails blocked and the estimated model spend (with the share of input read from the prompt cache).
- Safety controls: Engage kill switch stops all AI action at once; AI first response triages every new case at or above a severity; Max unattended containment caps automatic actions per client per hour. Anything that gives the AI more room (the kill switch off, a higher cap) waits for a second HELIOS approver under Rule changes; less room applies at once.
- Model providers: add Claude (Anthropic's API, or in your own AWS account through Amazon Bedrock), Azure OpenAI, or a local model server (Ollama, LM Studio, vLLM, the HELIOS AI box). Test checks the connection; Make default sets the model every agent uses unless told otherwise. Keys are encrypted and never shown again.
- Agent types: for each agent choose the provider, a model override, the reasoning effort, the autonomy (off, suggest only, contain with approval, auto-contain pre-approved actions) and extra instructions.
- Evaluations score a model on a labelled set (attacks, benign activity, look-alikes and planted instructions); the investigation agent is scored on worlds its tools search. A model a client or user picks only suggests until it passes for that agent.
- The run log lists recent runs with their verdict, model, tokens and outcome.

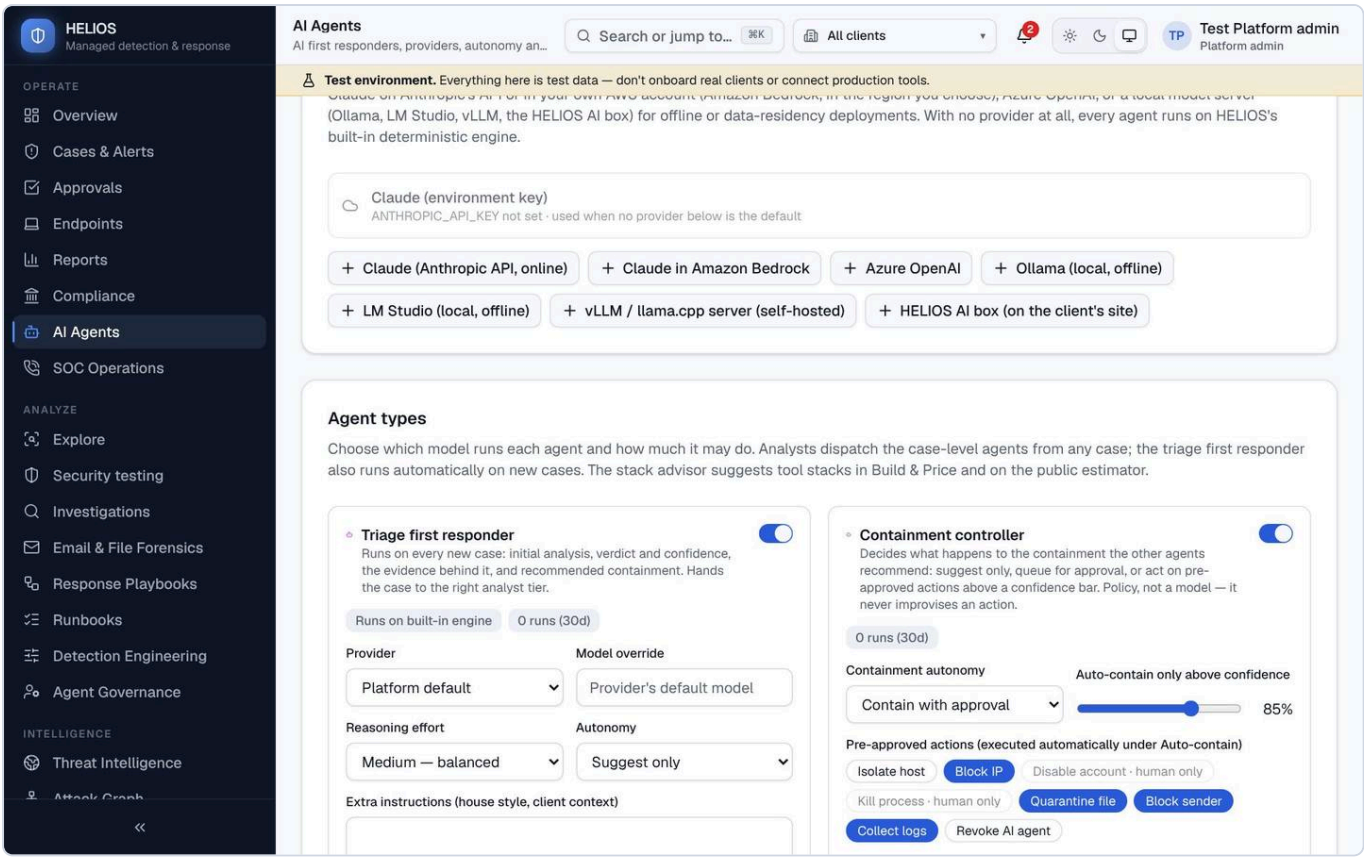


Figure 26. AI Agents, further down the page.

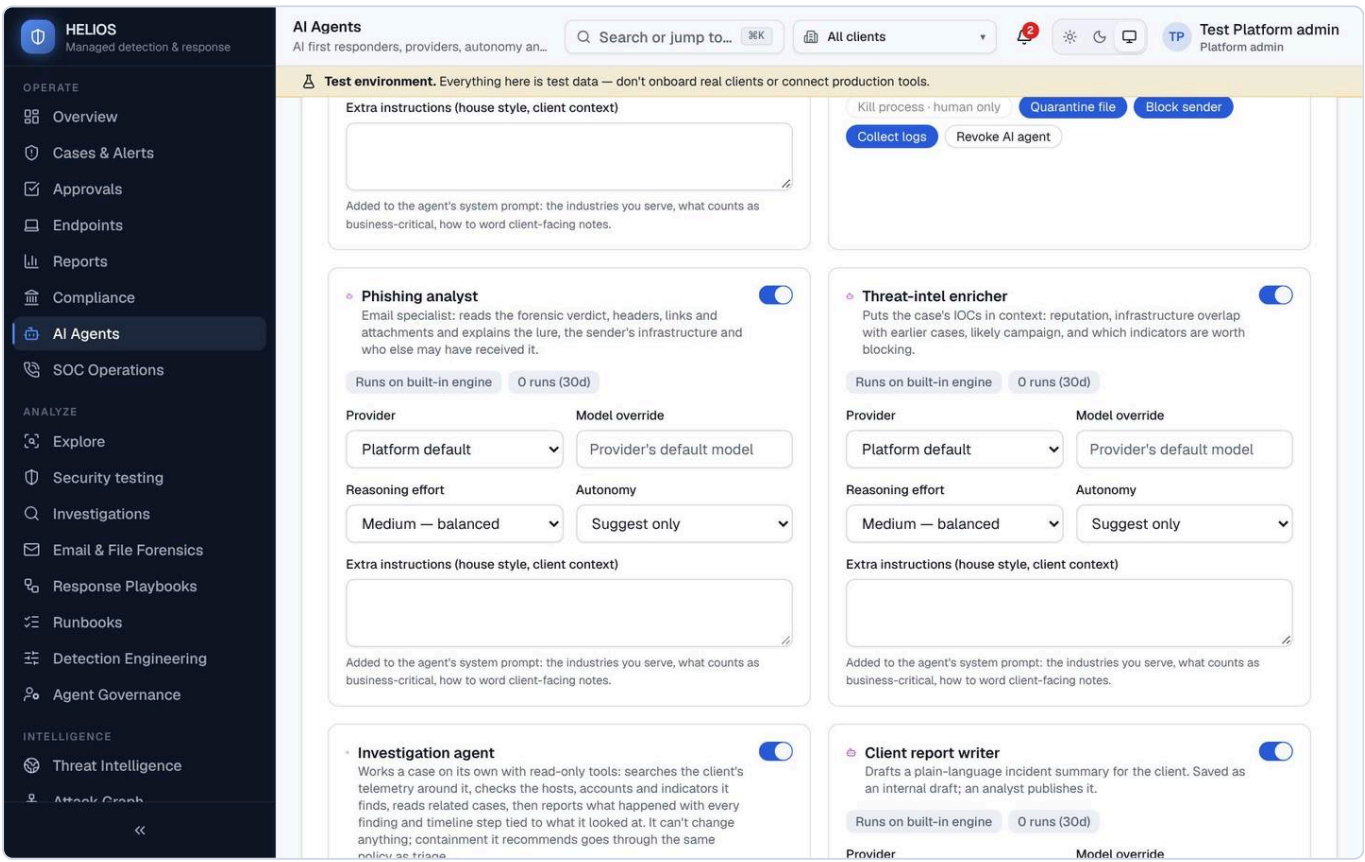


Figure 27. AI Agents, further down the page.

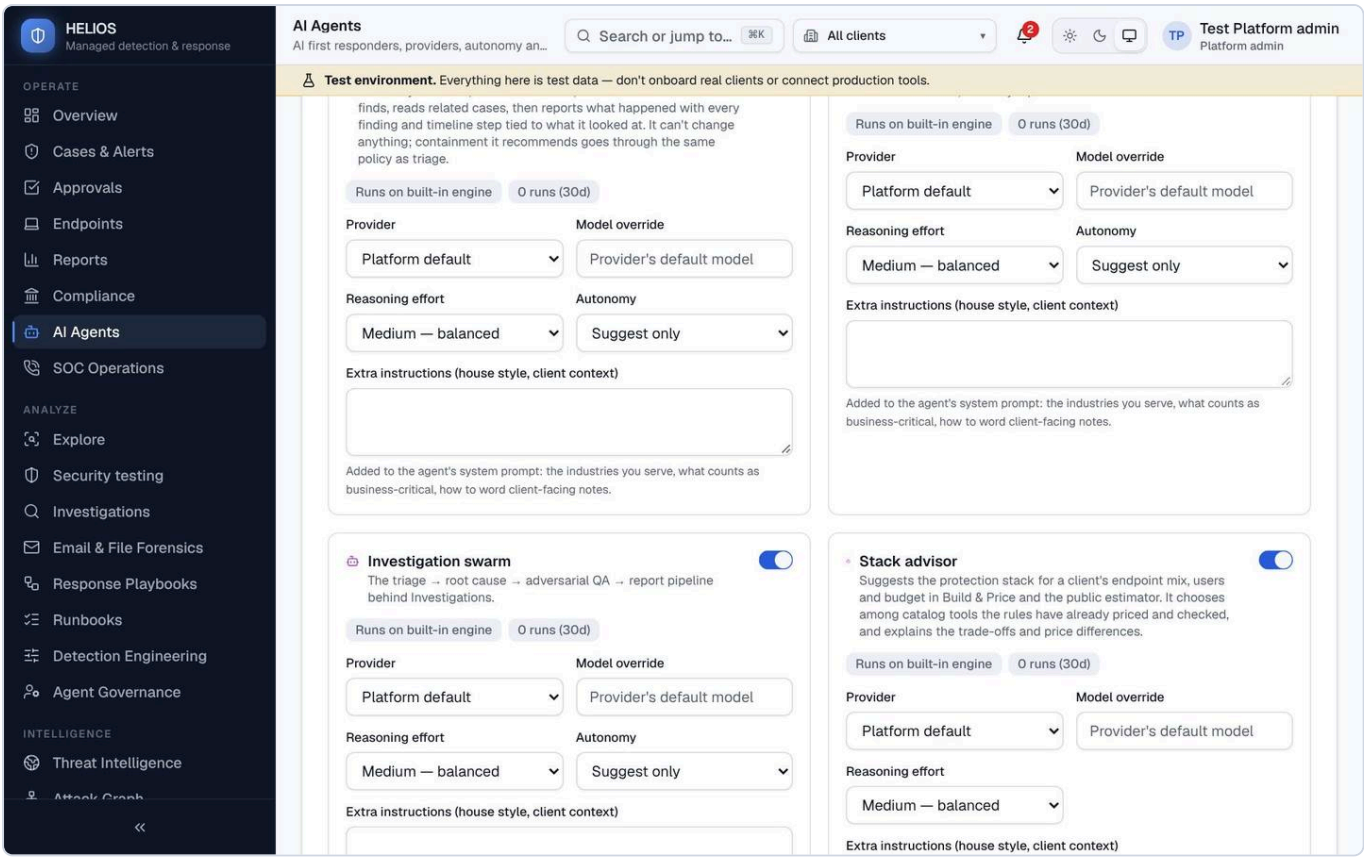


Figure 28. AI Agents, further down the page.

SOC Operations

Who is on call and the pages they get, the weekly QA of what the AI did on its own, and the daily digests.

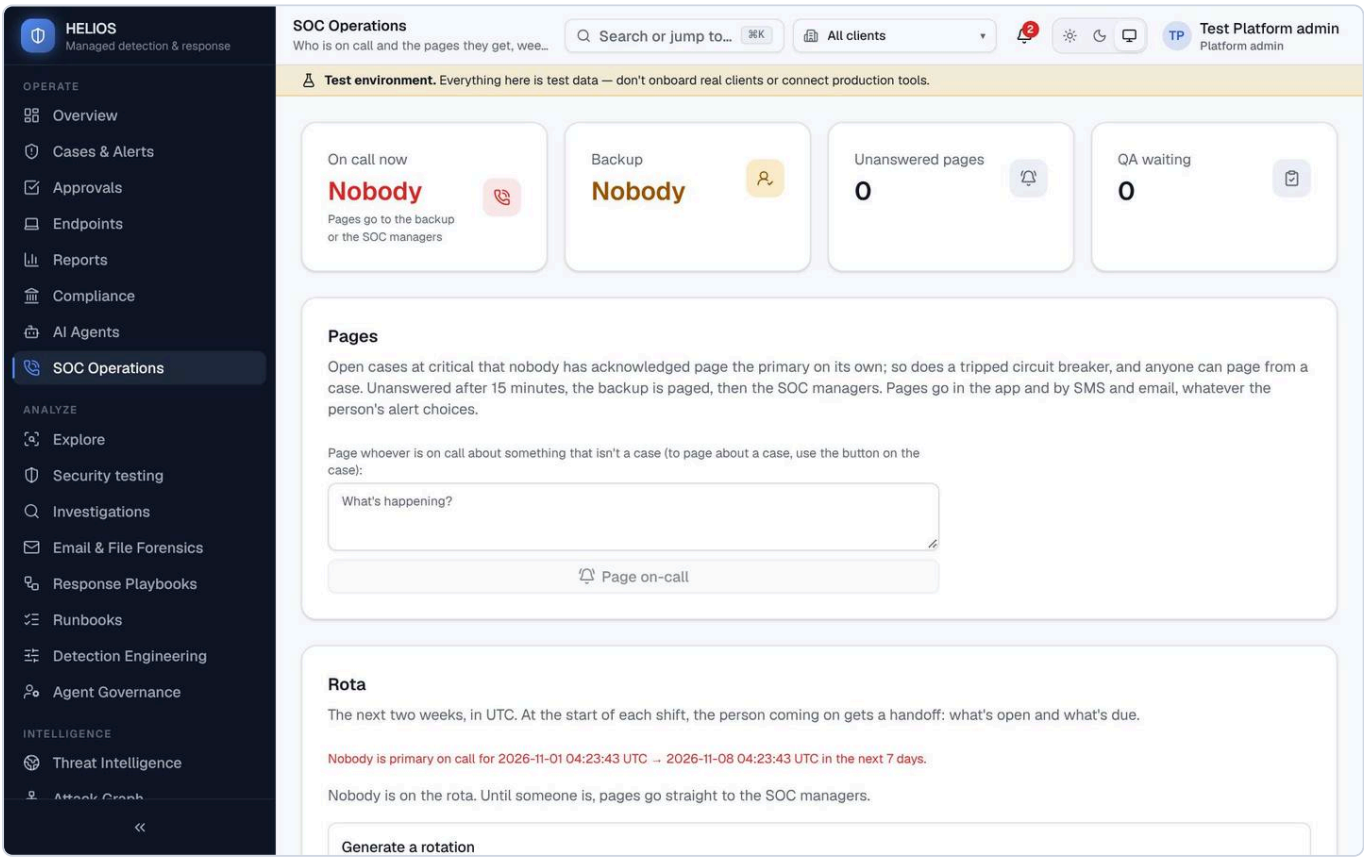


Figure 29. SOC Operations, as seen by a platform admin (test environment).

How to use it

- The tiles show who is primary and backup on call now, unanswered pages and QA reviews waiting.
- Pages: open critical cases nobody acknowledged page the primary on their own; unanswered after the acknowledgement window (15 minutes by default) the backup is paged, then the SOC managers. Pages go in the app and by SMS and email. Page on-call sends a page about something that isn't a case.
- Rota: Generate a rotation from people taking turns (8, 12 or 24-hour shifts, or a week); add cover for a swap or leave. Whoever starts a shift gets a handoff: open high and critical cases, unanswered pages, approvals waiting, notices due, quiet sources.
- QA: each Monday HELIOS samples last week's cases where the AI acted without anyone approving or called a case benign (weighted towards critical cases, new rules and critical assets). Review each one on the case: should the action have run, was the verdict right. A 'shouldn't have run' counts against the AI's autonomy at once.
- Settings: paging severity, acknowledgement window, QA share and floor, and the digest time and time zone.

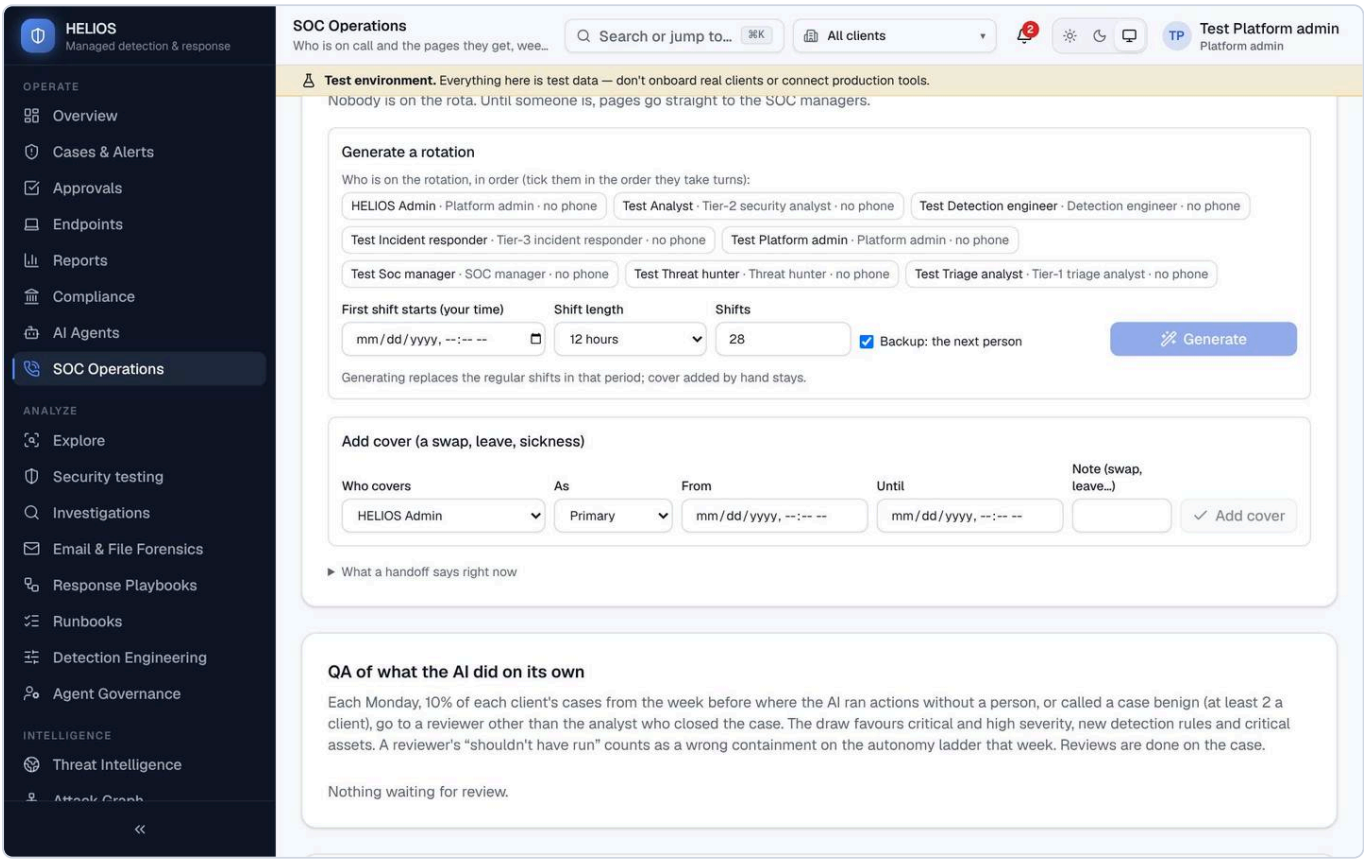


Figure 30. SOC Operations, further down the page.

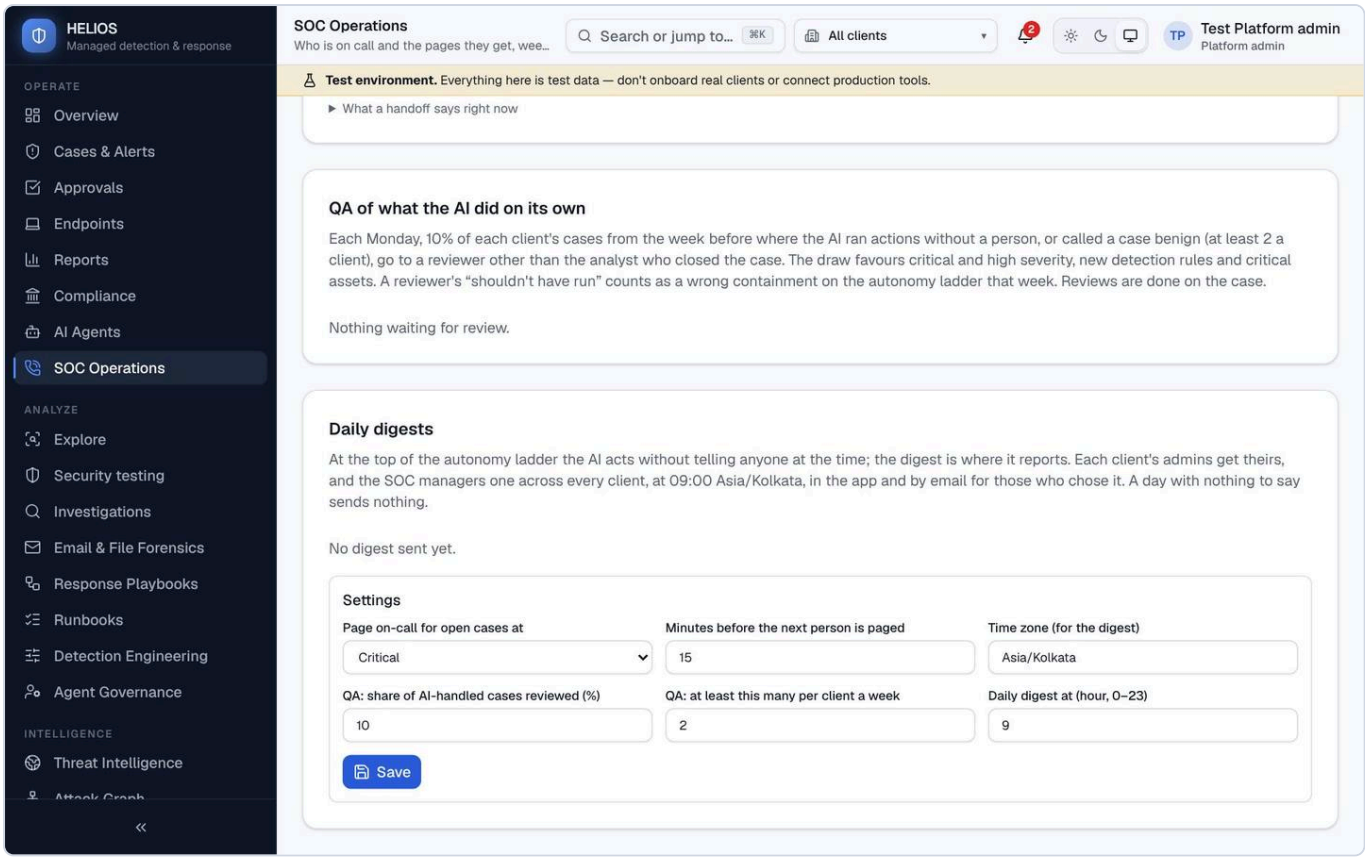


Figure 31. SOC Operations, further down the page.

Explore

One search over every log HELIOS holds, from all of a client's security tools, the way analysts search Splunk, Chronicle, QRadar or Wazuh. The SOC searches across clients or one at a time; client users

search their own organisation's logs only.

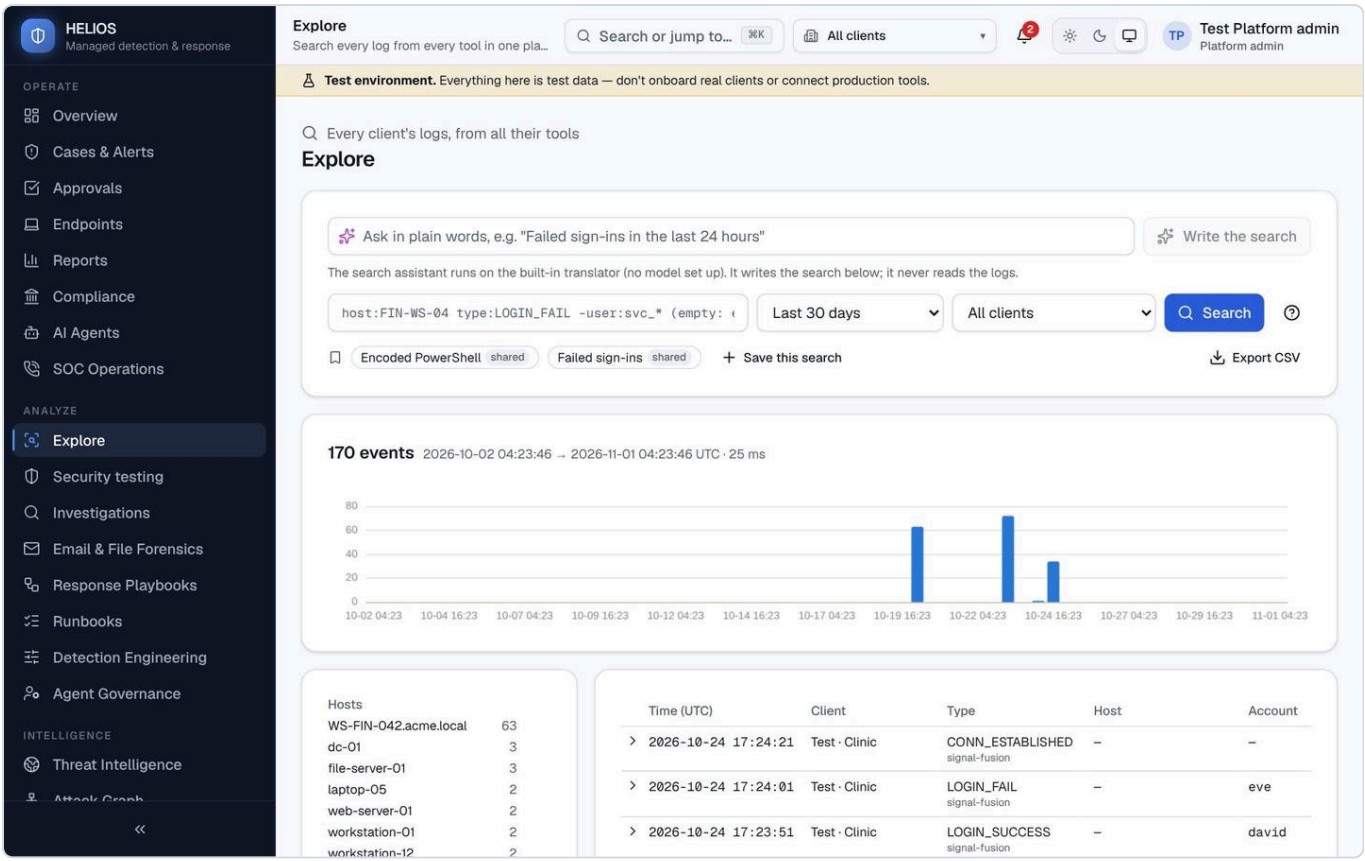


Figure 32. Explore, as seen by a platform admin (test environment).

How to use it

- Ask in plain words ("failed sign-ins for alice this week") and Write the search: the search assistant writes it on the model chosen for it, under the client's AI rules. It sees only your question and the field list, never the logs, and what it writes is checked like anything typed. Without a model a built-in translator does it.
- Or type the search: host:FIN-WS-04 type:LOGIN_FAIL -user:svc_*. Terms side by side all match, OR joins two, - or NOT excludes, * is a wildcard, port:>1024 compares, ip:a,b is any of a list, hash:* means it has a value, and a bare word or "phrase" searches everywhere in the event, raw fields included. The ? button lists the fields.
- Pick the time range (15 minutes to 90 days) and, for the SOC, the client or All clients, then Search.
- The histogram shows events over the window: click a bar to zoom into it. The side panel lists the top hosts, accounts, event types, IPs, ports, processes, log sources, products and client tools: click a value to keep only it, or - to exclude it.
- Events are listed newest first; click one for every field, the raw record the tool sent and its OCSF form, with + and - beside each value to filter. Load older fetches the next page.
- Save this search keeps it for later, private or shared with your side (your organisation's users, or the SOC). Export CSV downloads up to 10,000 matching events.
- Scripts and other tools can run the same searches through the REST API (GET /api/v1/events) with an API token.

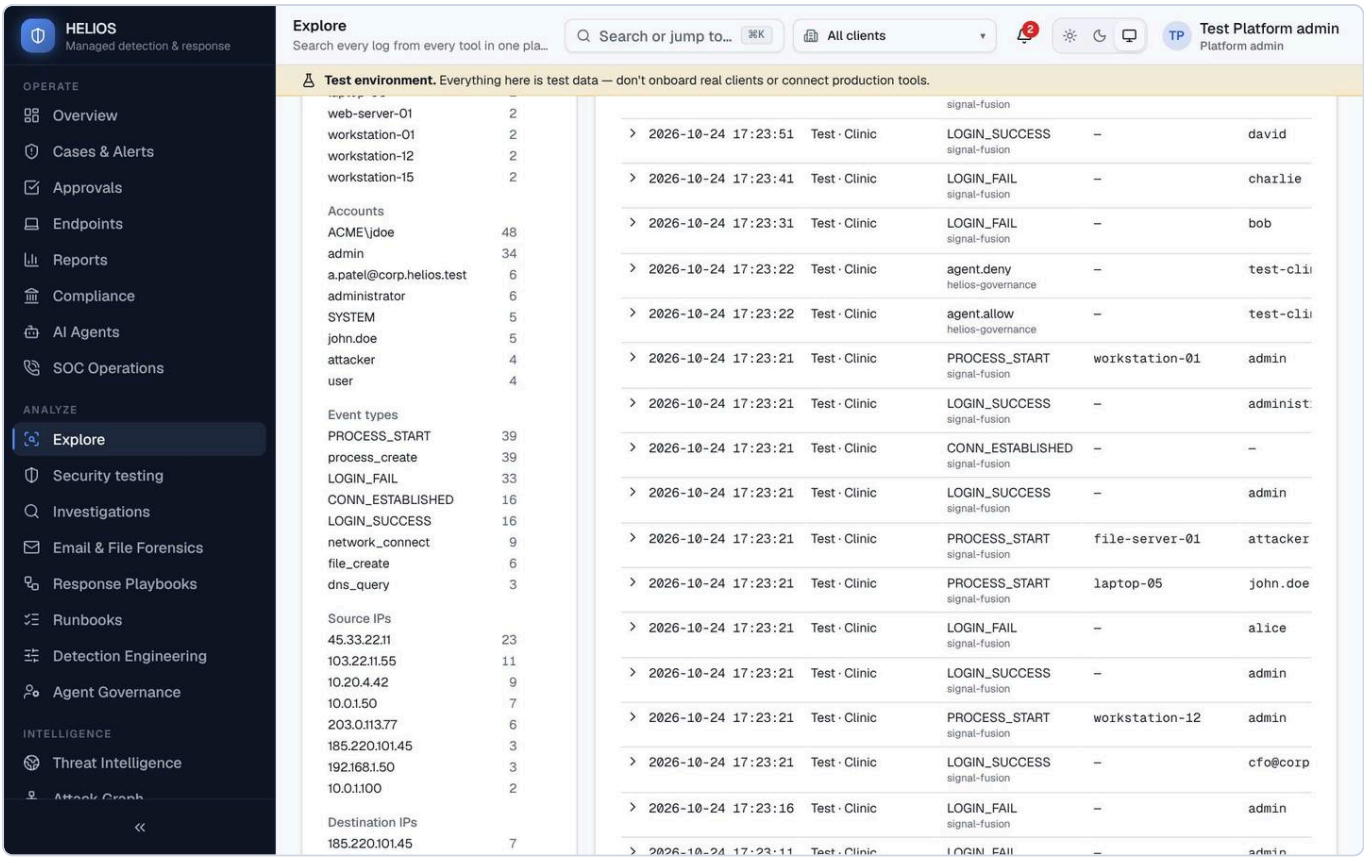


Figure 33. Explore, further down the page.

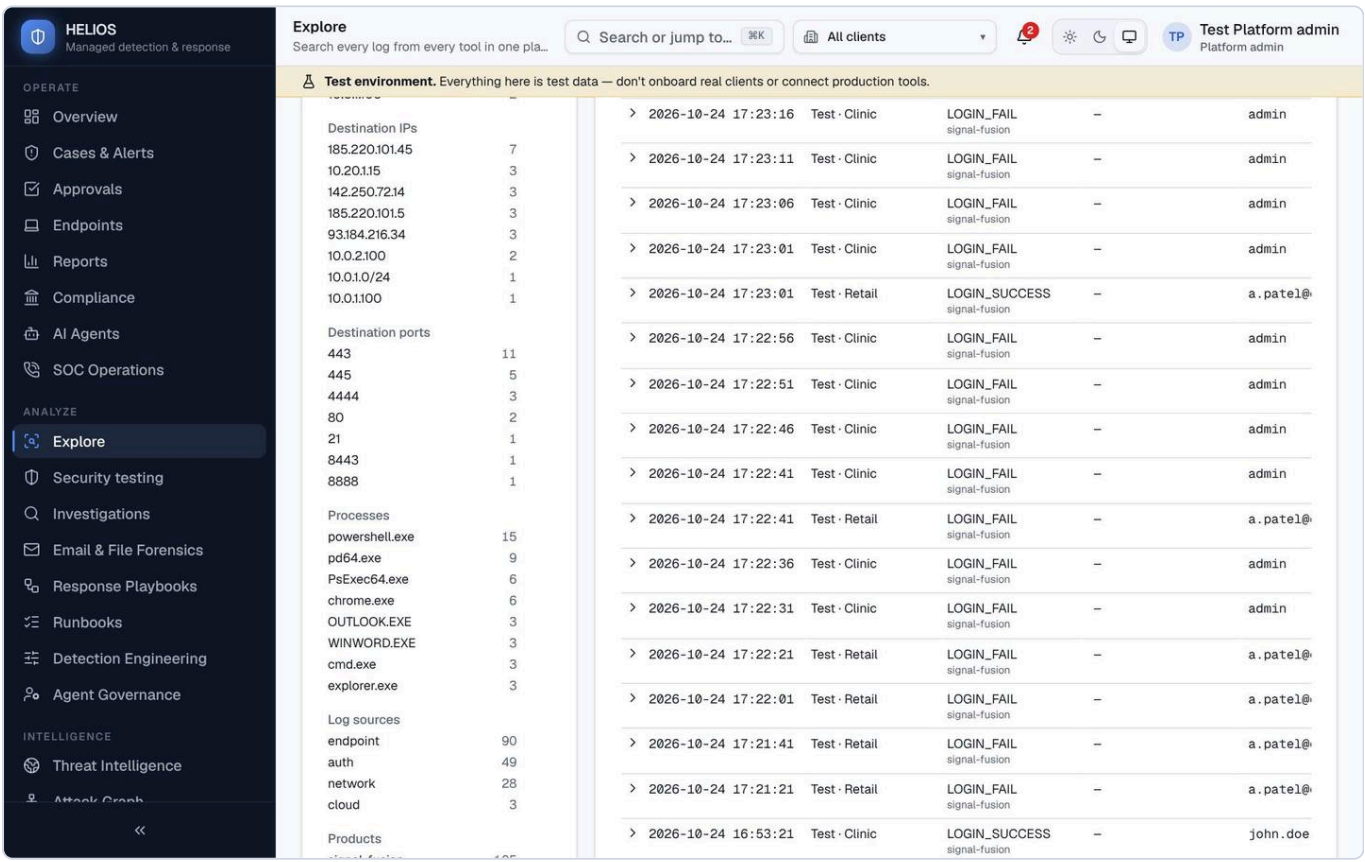


Figure 34. Explore, further down the page.

Security testing

Tests of a client's own internet-facing assets, on request and under written authorisation from both the client and HELIOS: what an attacker on the internet sees (exposed services, TLS, web headers, email spoofing protection), with results from established scanners in the same report.

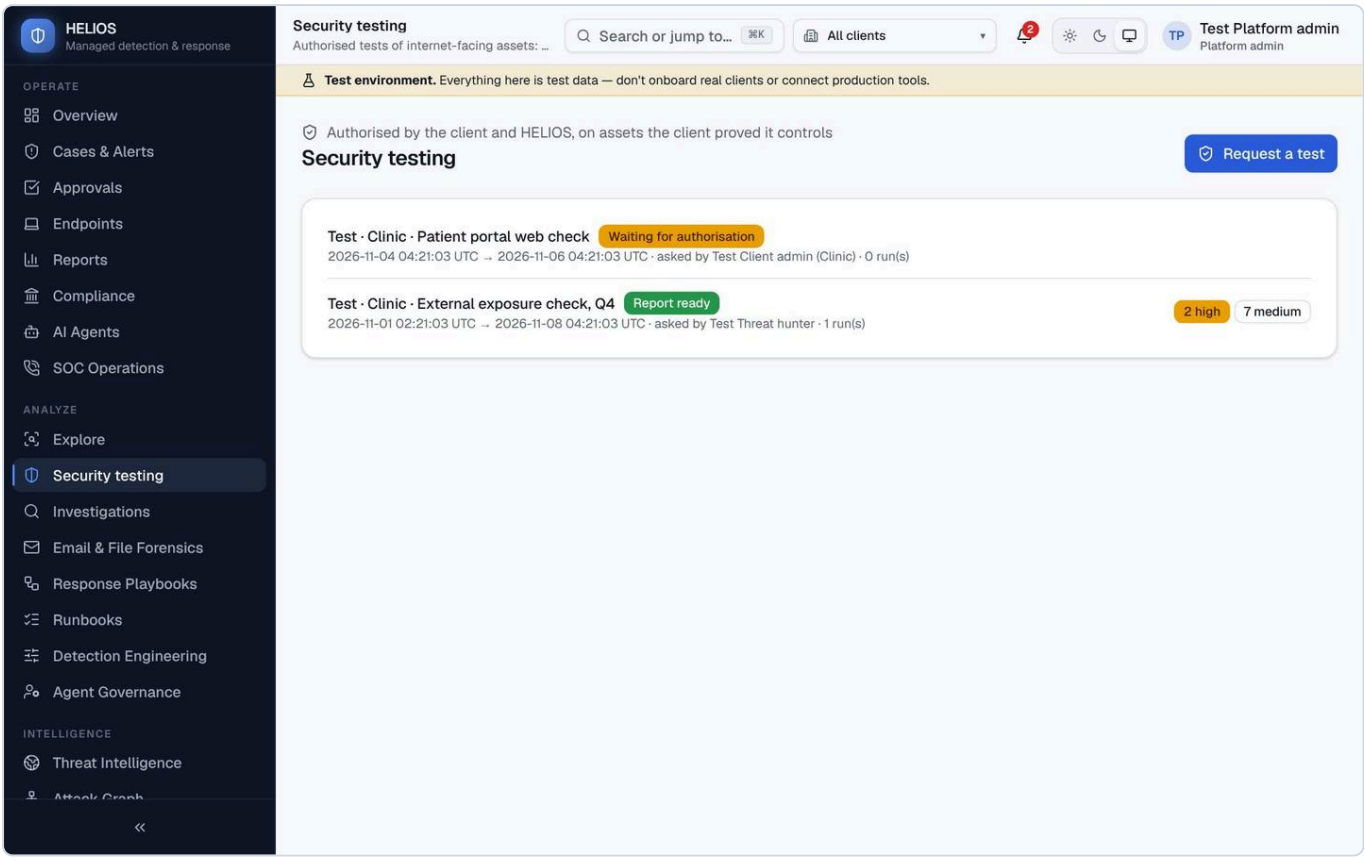


Figure 35. Security testing, as seen by a platform admin (test environment).

How to use it

- Each test shows its status (waiting for authorisation, authorised, running, report ready, rejected or stopped), the window, who asked, and its open findings by severity. Open one for the detail.
- Request a test: the targets (domains, host names, public IP addresses or URLs; up to 25, no ranges, no private addresses), the checks, the window (up to 30 days) and the rules of engagement (who to call, what's out of bounds, times to avoid).

Investigations

Upload logs from an incident and get a root-cause analysis: an agent pipeline (triage, root cause,

adversarial QA, report) works through them and opens a case with the findings.

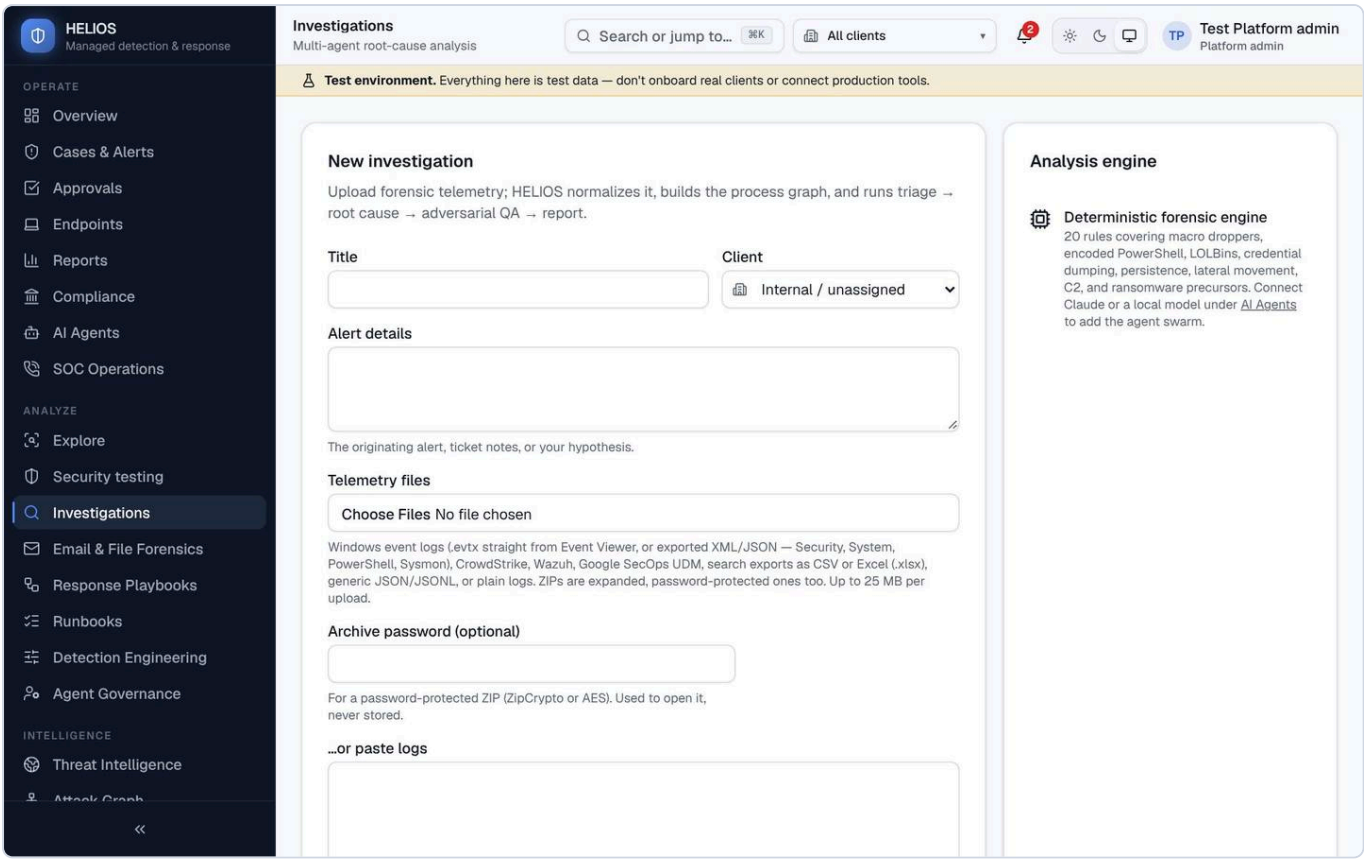


Figure 36. Investigations, as seen by a platform admin (test environment).

How to use it

- New investigation: a title, the client, the alert details (the originating alert, ticket notes or your hypothesis), and the telemetry files: Windows event logs (.evtx, or exported XML/JSON), EDR and SIEM exports, CSV or Excel, JSON/JSONL, plain logs, or a .zip (password-protected ones too, with the archive password, which isn't stored). Up to 25 MB per upload; or paste the logs.
- HELIOS normalizes the telemetry, builds the process graph and runs triage, root cause, adversarial QA and the report. Without a model it uses its deterministic forensic engine (20 rules for droppers, encoded PowerShell, living-off-the-land binaries, credential dumping, persistence, lateral movement, C2 and ransomware precursors).
- Open a past investigation to read its timeline, root cause, the adversarial QA's challenges and the report.

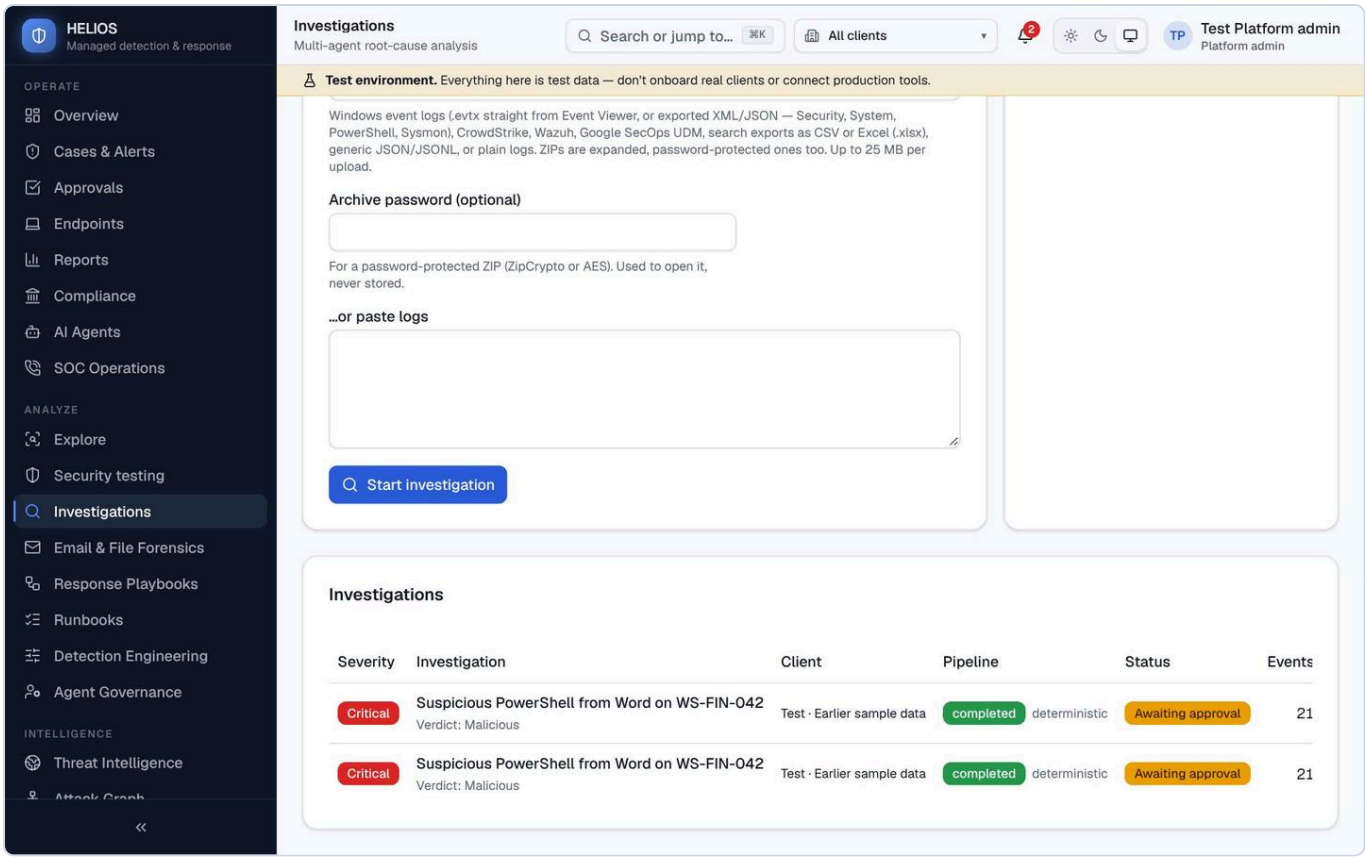


Figure 37. Investigations, further down the page.

Email & File Forensics

Submit a suspicious email or file for analysis: header authentication, links, attachments, macros, archives

and images are examined and scored to a verdict.

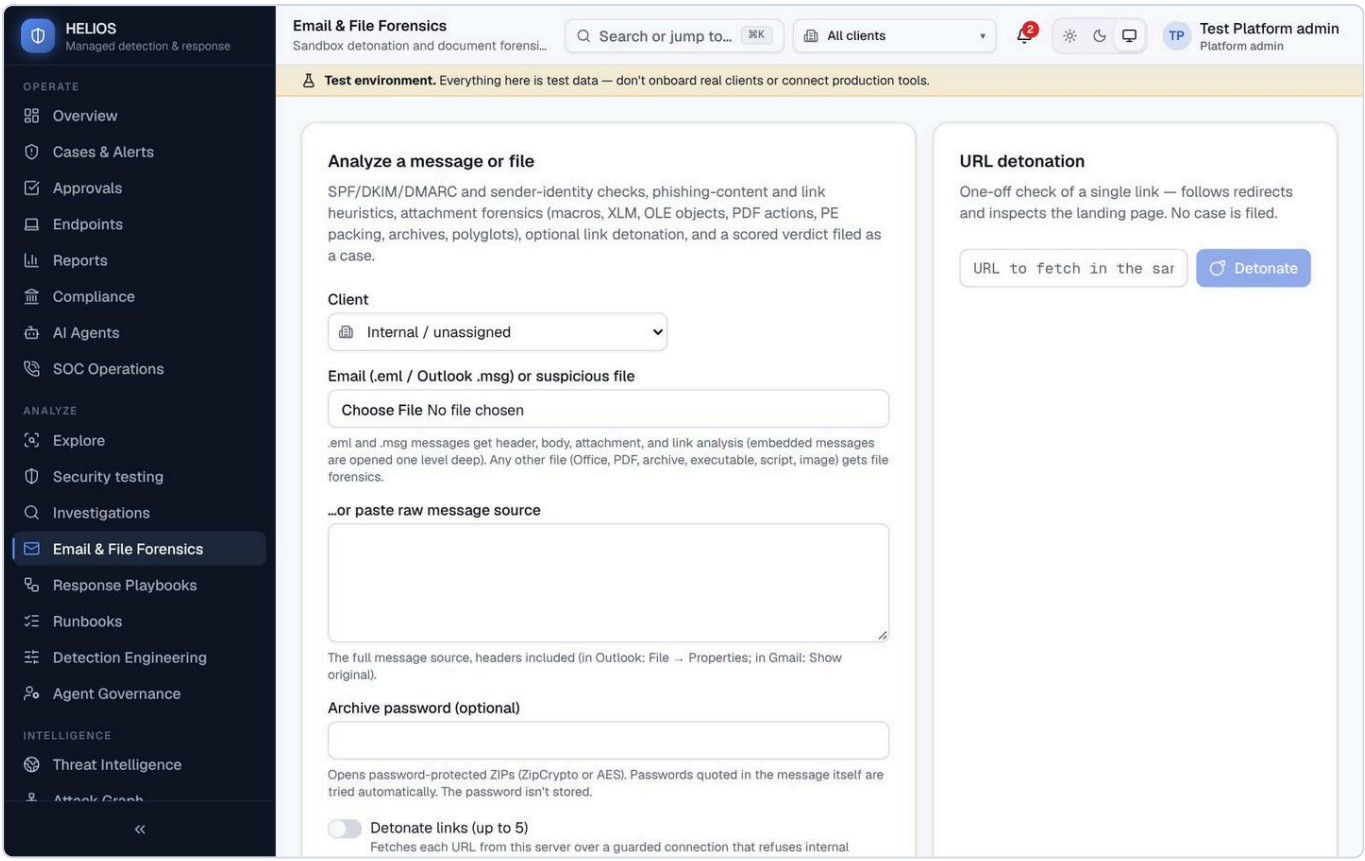


Figure 38. Email & File Forensics, as seen by a platform admin (test environment).

How to use it

- Analyze a message or file: choose the client, then the .eml or Outlook .msg (or any suspicious file), or paste the raw message source. Give an archive password if it's a protected ZIP (it isn't stored), and switch on Detonate links (up to 5) if you want links fetched through the guarded connection.
- HELIOS checks the sender's authentication (SPF, DKIM, DMARC) and identity, phishing content and links, attachments (macros, OLE objects, PDF actions, packing, archives, polyglots) and text and QR codes in images. URL detonation on the right checks a single link without filing a case.
- The verdict comes from the total score across the header, content, attachment, sandbox, intel and sender checks, against the suspicious and malicious thresholds set in Settings. A malicious result opens a case, where the message can be purged from mailboxes and the sender blocked.

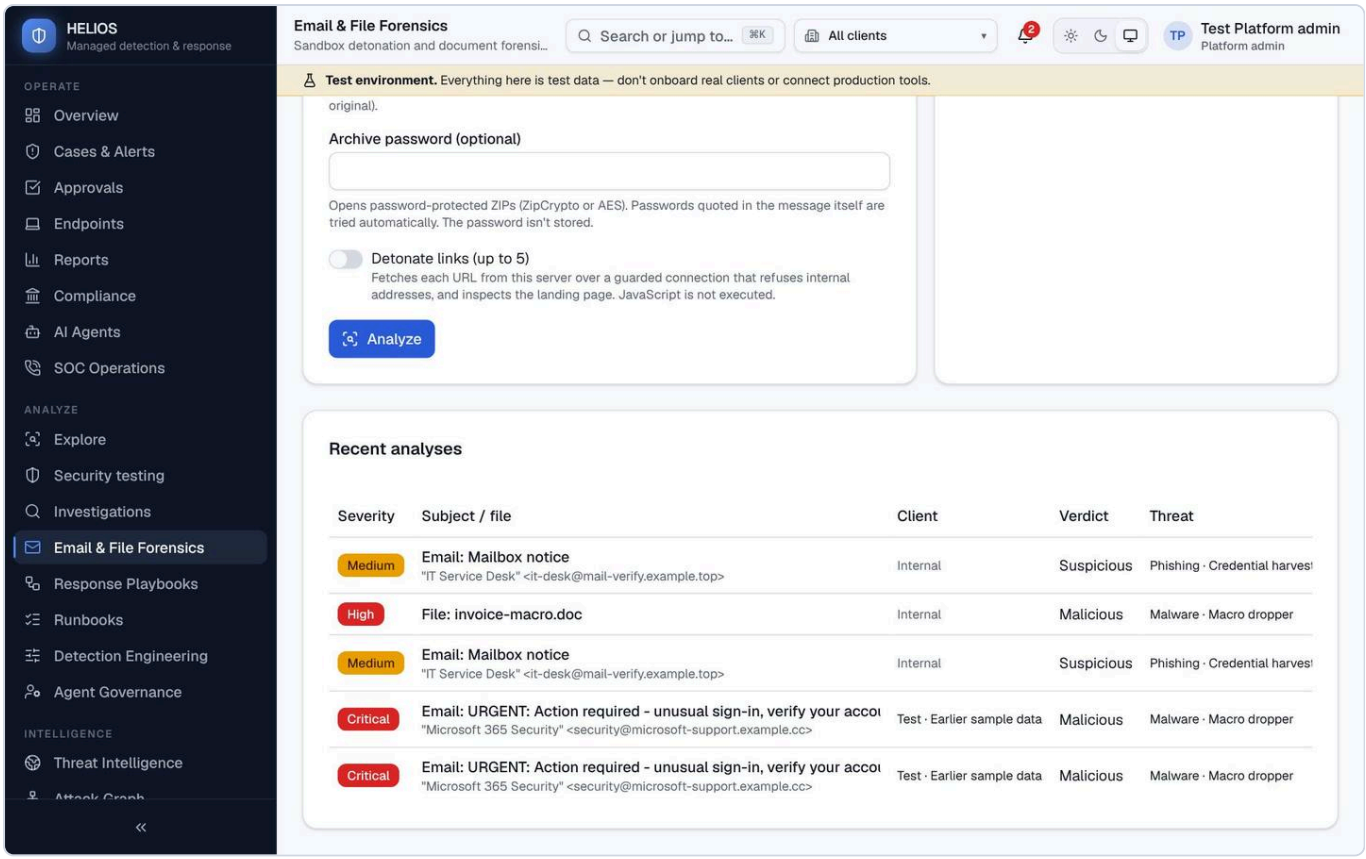


Figure 39. Email & File Forensics, further down the page.

Response Playbooks

Automated response workflows: what triggers them and the steps they run (containment, collecting logs,

notifying, opening tickets).

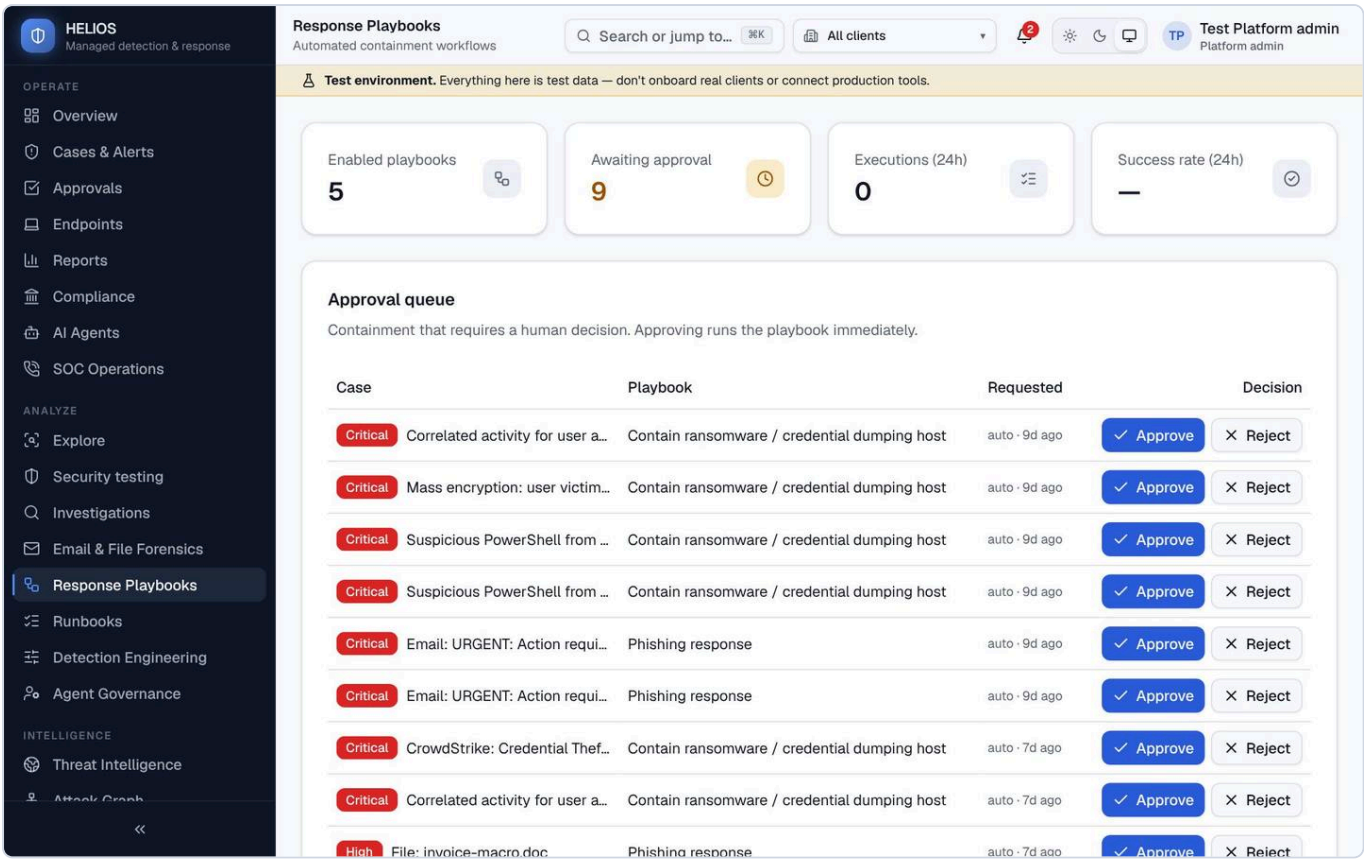


Figure 40. Response Playbooks, as seen by a platform admin (test environment).

How to use it

- Each playbook shows its trigger, its steps and whether it is enabled. Its runs appear on the cases they ran on; steps that need approval wait in Approvals, and steps without a connector are left to do by hand.
- Create or edit a playbook: name, trigger (severity, source, technique), and an ordered list of steps with their parameters. Enable it when it's ready.

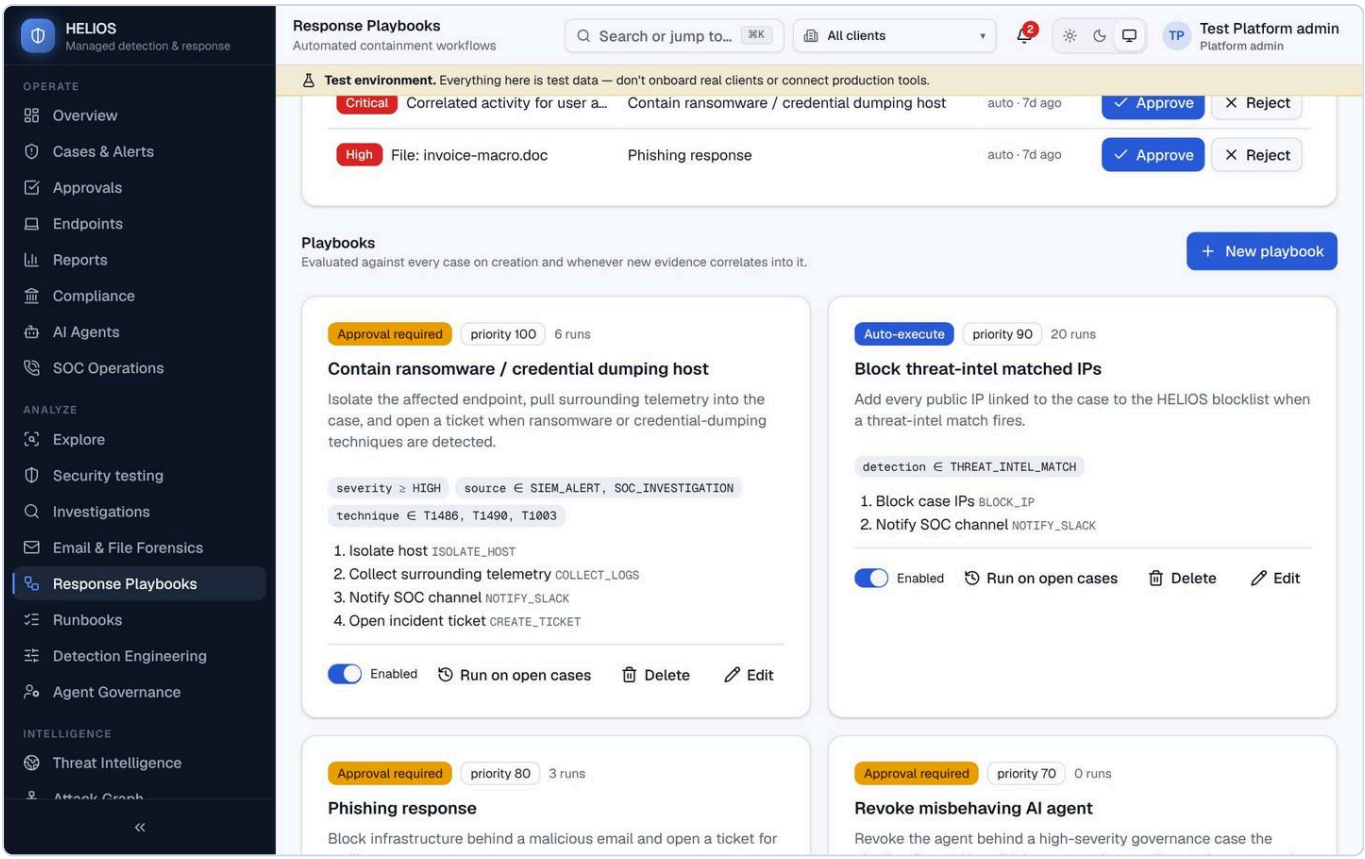


Figure 41. Response Playbooks, further down the page.

Runbooks

Procedures for analysts, attached to the cases they match and worked as checklists on the case.

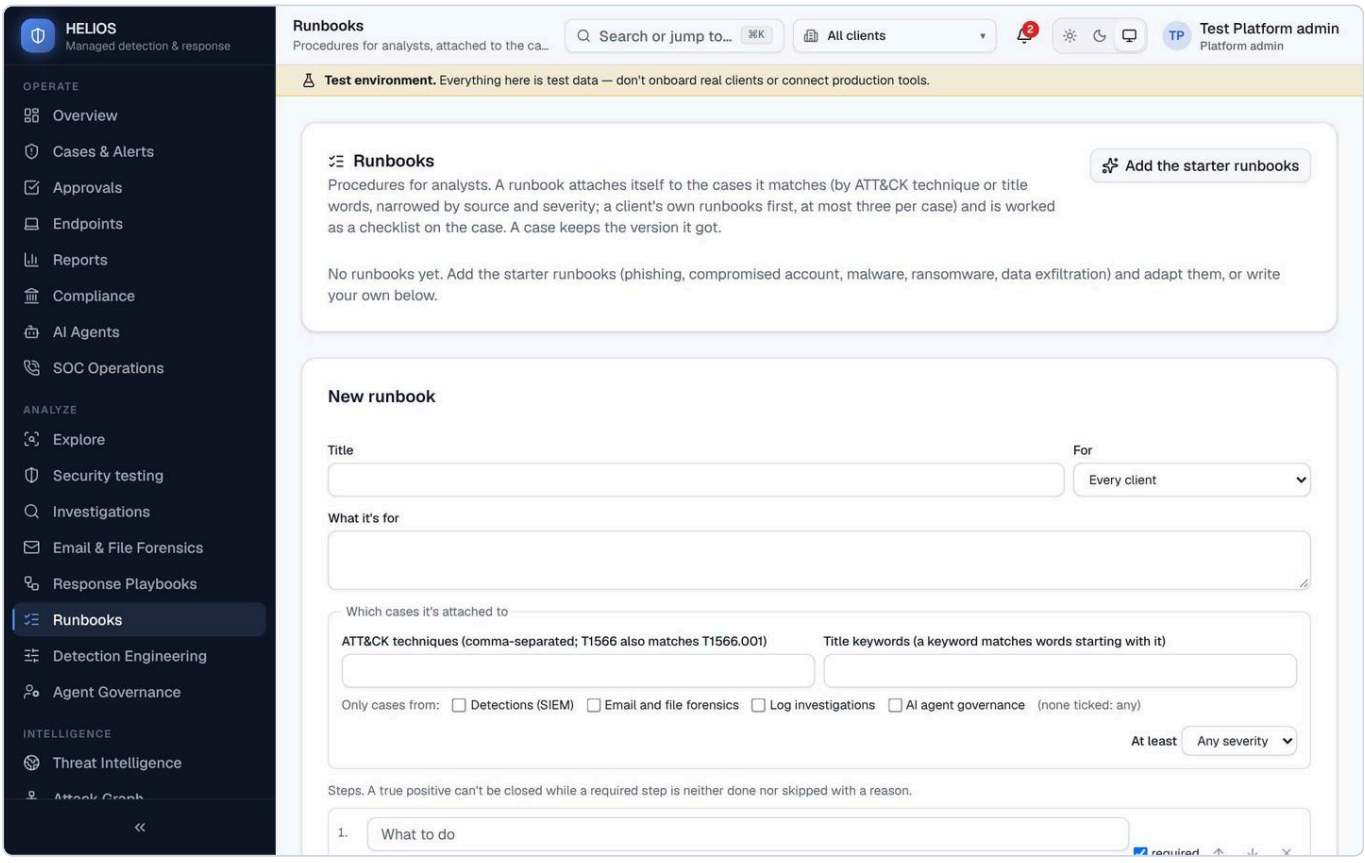


Figure 42. Runbooks, as seen by a platform admin (test environment).

How to use it

- A runbook matches cases by ATT&CK technique or title words, narrowed by source and minimum severity. Up to three matching runbooks attach to a case when it is opened, the client's own first.
- New runbook: a title, which client it's for (or every client), what it's for, the ATT&CK techniques and title keywords it matches, the sources and minimum severity, and its steps (tick required for steps a true positive can't close without). Add the starter runbooks to begin from phishing, compromised account, malware, ransomware and data exfiltration. Editing makes a new version; cases already attached keep the steps they had.

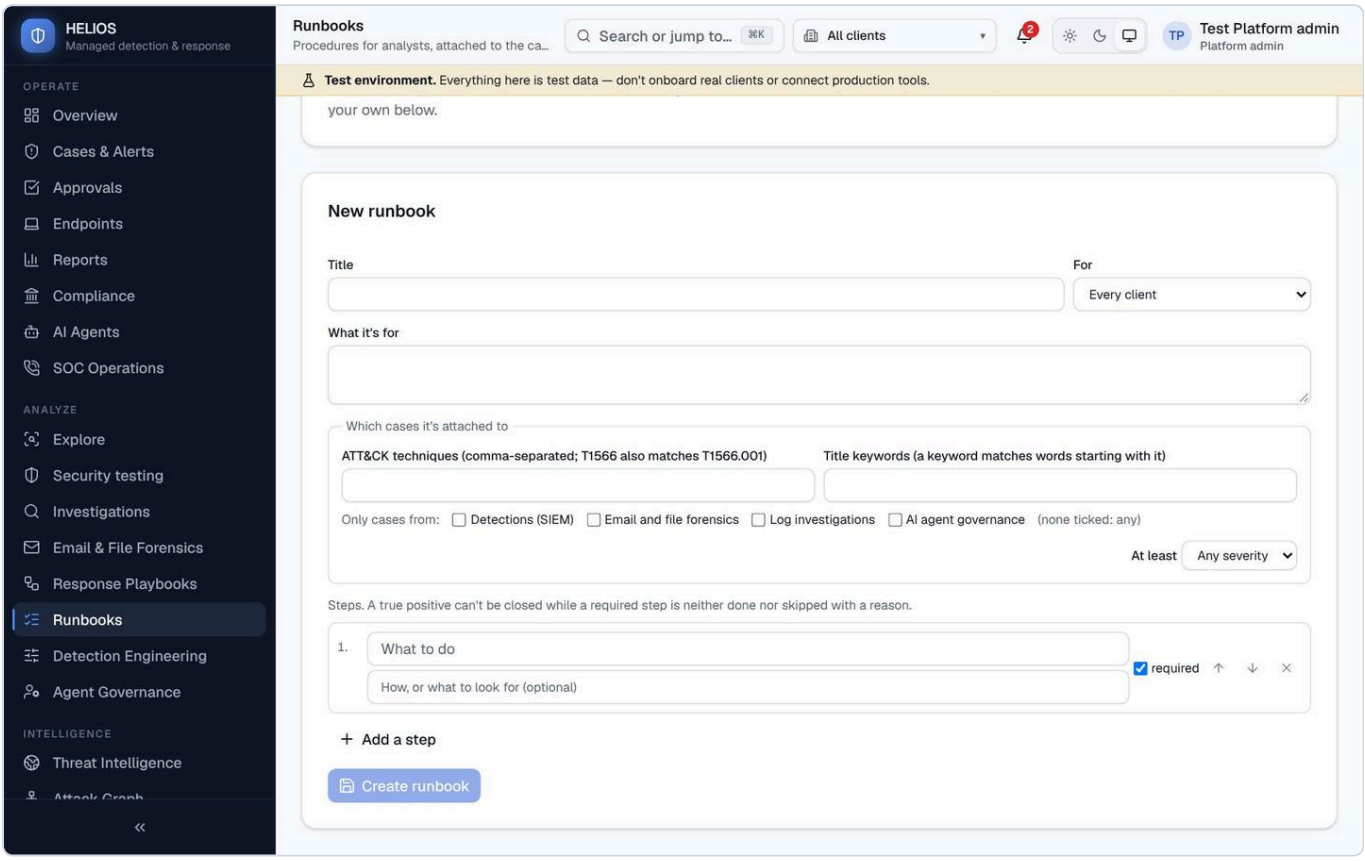


Figure 43. Runbooks, further down the page.

Detection Engineering

The detectors that turn telemetry into cases: the built-in detectors, imported Sigma rules tested before

they go live, and per-client tuning.

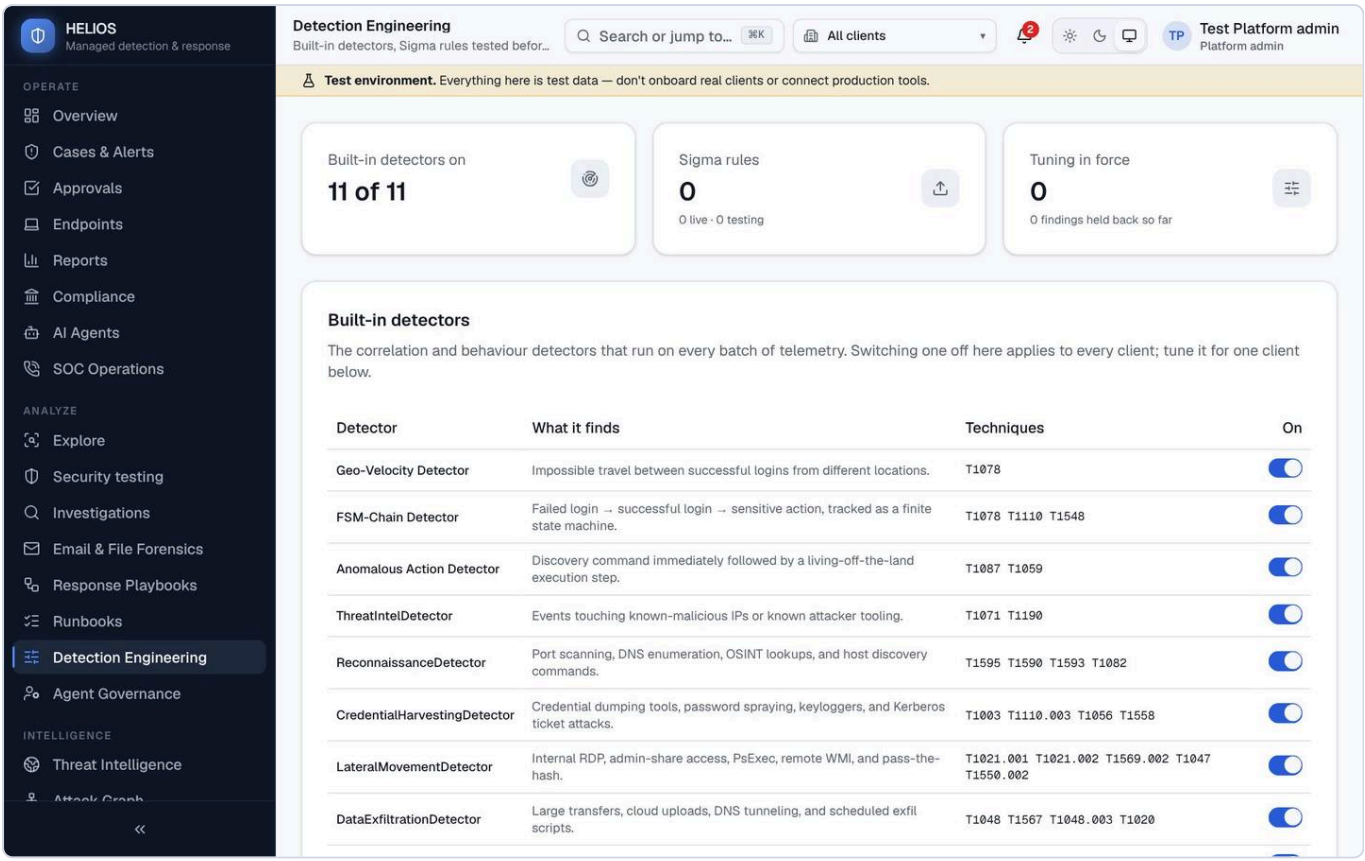


Figure 44. Detection Engineering, as seen by a platform admin (test environment).

How to use it

- Built-in detectors: switch one off for every client here (tune for one client below instead where you can).
- Import a Sigma rule: it is compiled into a matcher (anything it can't evaluate faithfully is refused with the reason), backtested on recent history, run in testing (matches counted, nothing opened) and then made live.
- Tune per client: switch a detector or rule off there, leave out known-good hosts, accounts or processes, or change the severity, each with a reason and an optional expiry. The count shows what tuning held back.

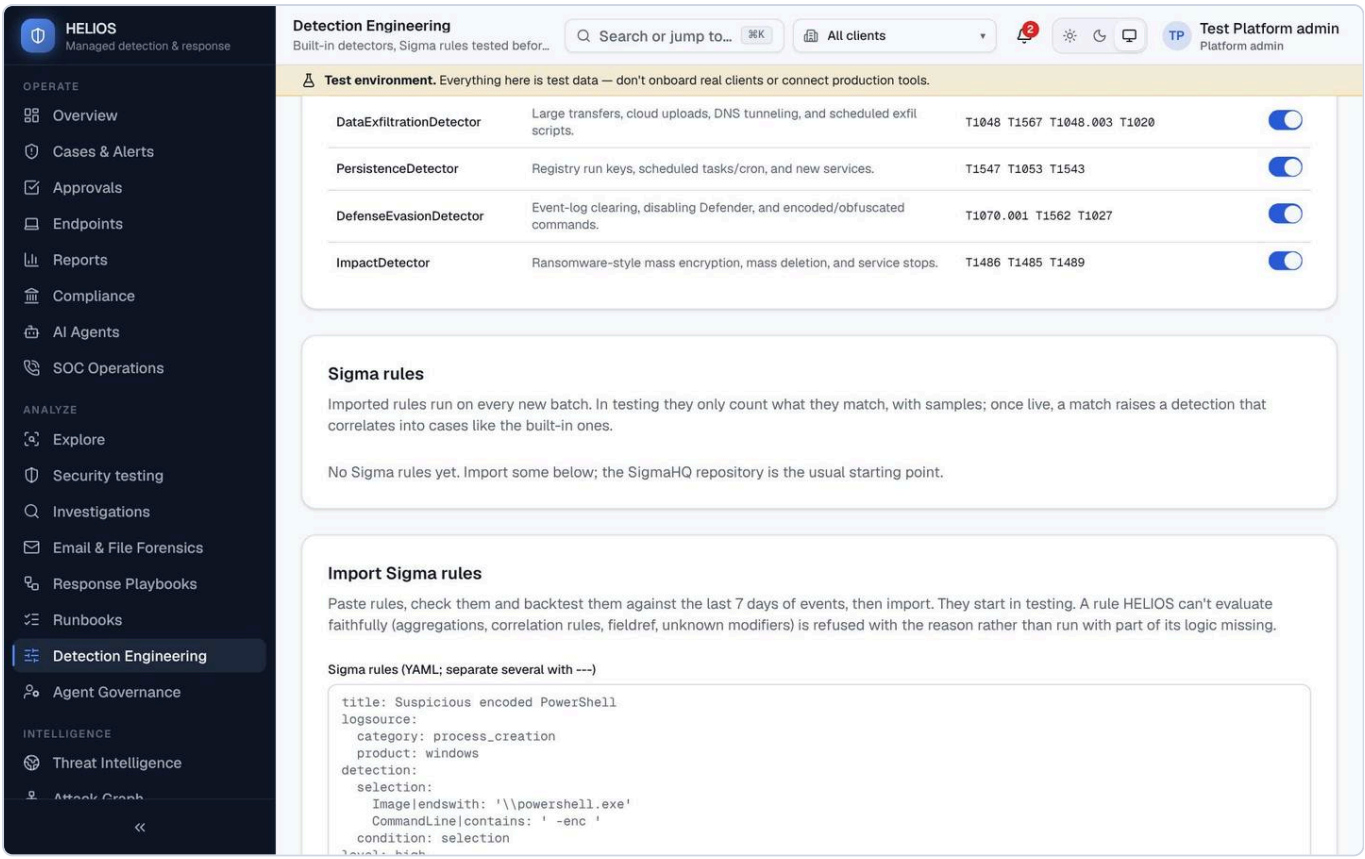


Figure 45. Detection Engineering, further down the page.

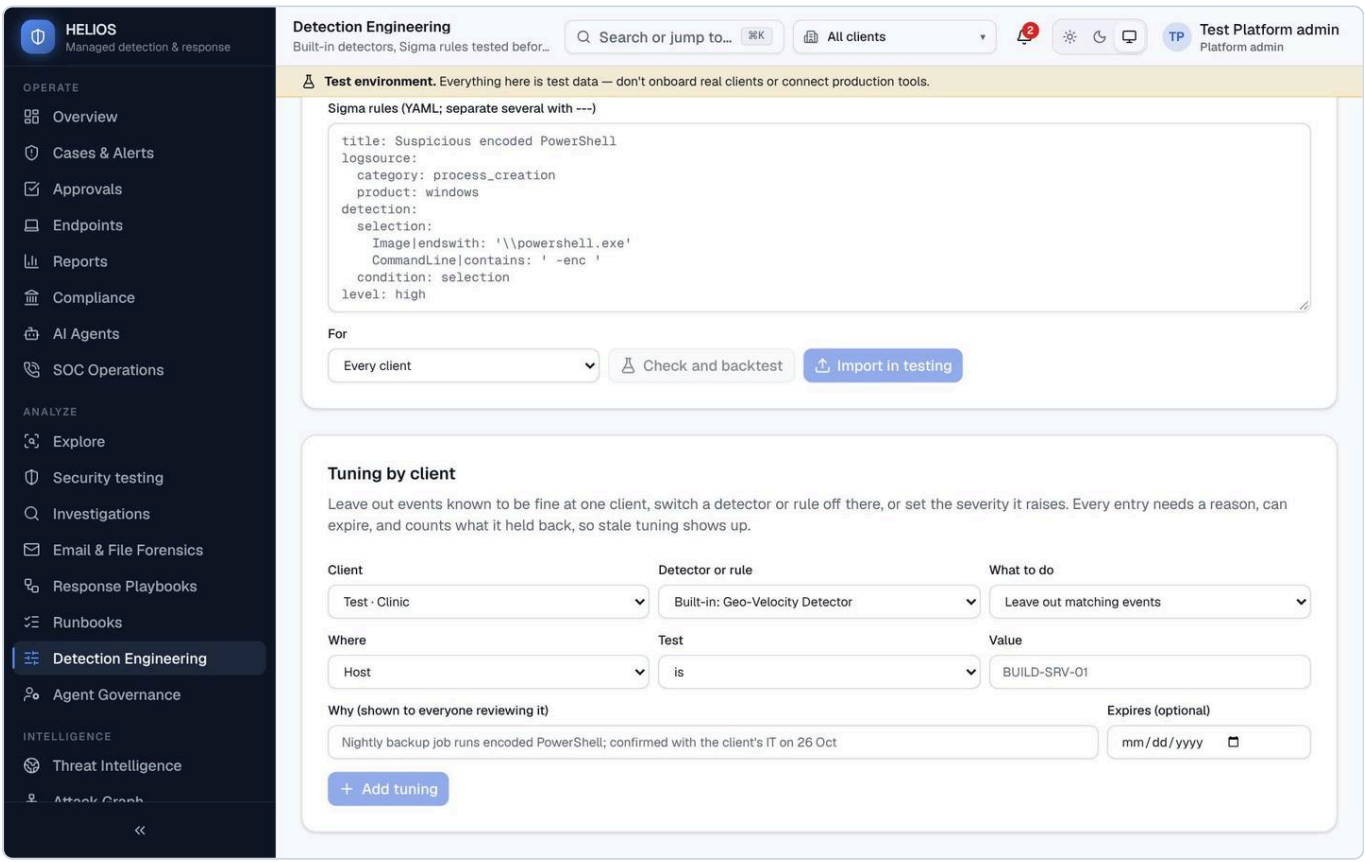


Figure 46. Detection Engineering, further down the page.

Agent Governance

Govern the AI agents your clients run: identities and tokens, the policies they act under, their trust, and revocation.

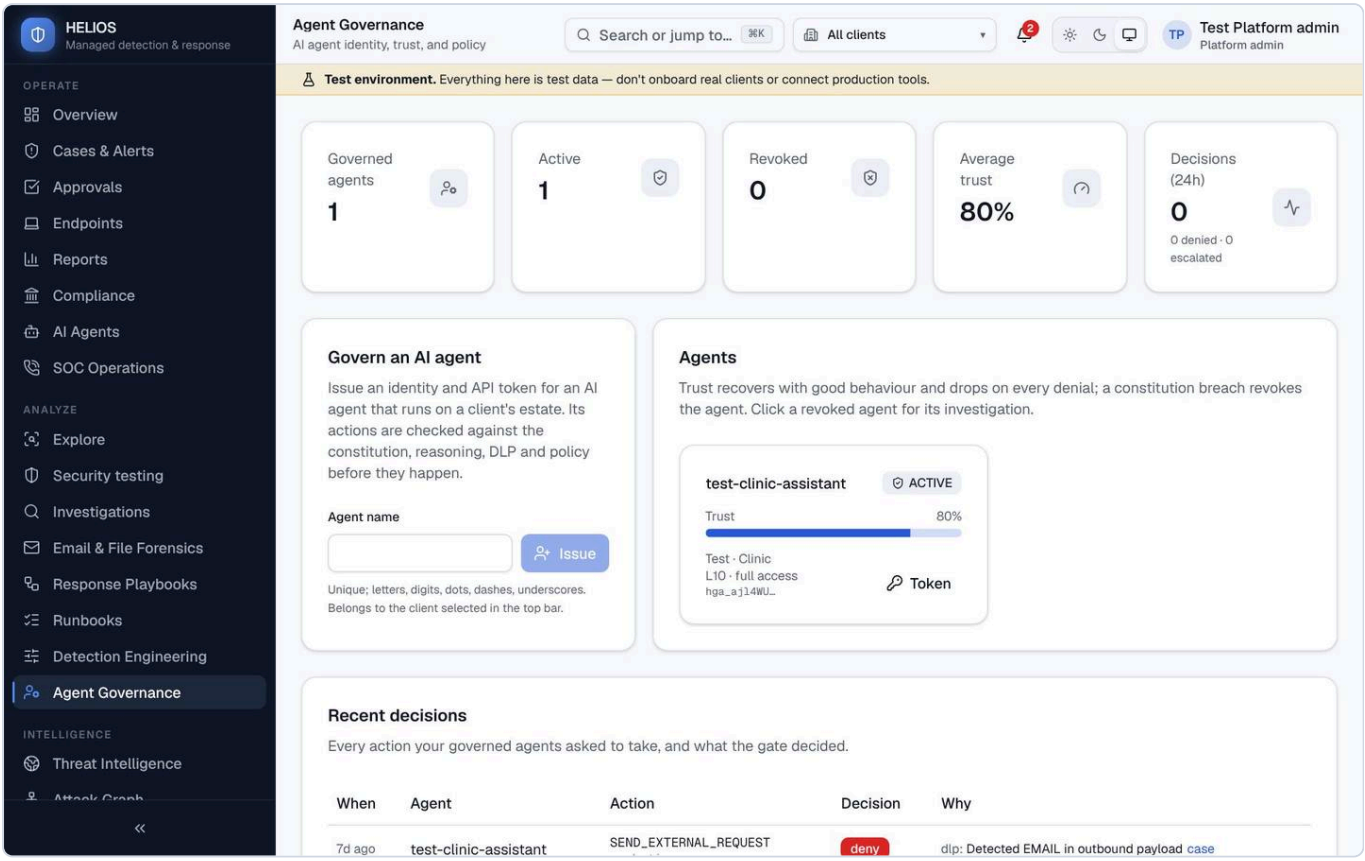


Figure 47. Agent Governance, as seen by a platform admin (test environment).

How to use it

- Govern an AI agent: enter the agent's name and Issue, for the client selected in the top bar. Its token is shown once. The agent asks HELIOS before each action; constitution, reasoning, data-loss and policy checks decide, and Recent decisions shows each request and why it was allowed or denied.
- Trust recovers with good behaviour and drops on every denial.
- A breach revokes the agent and opens a case. Revoke or restore an agent by hand here or from its case.

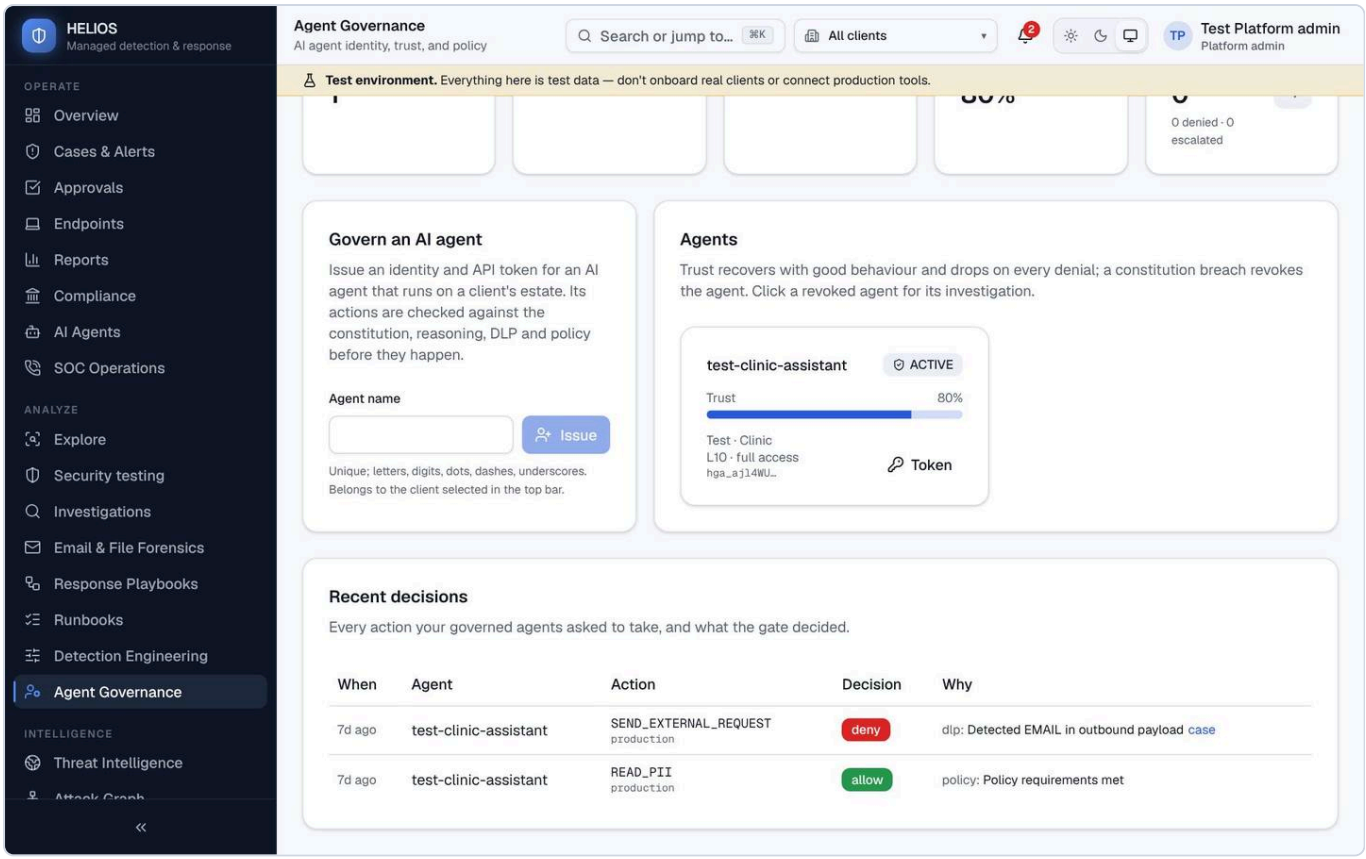


Figure 48. Agent Governance, further down the page.

Threat Intelligence

Look up indicators (IP, domain, URL, file hash, email address) against the SOC's indicator store, the

configured reputation providers and the threat feeds.

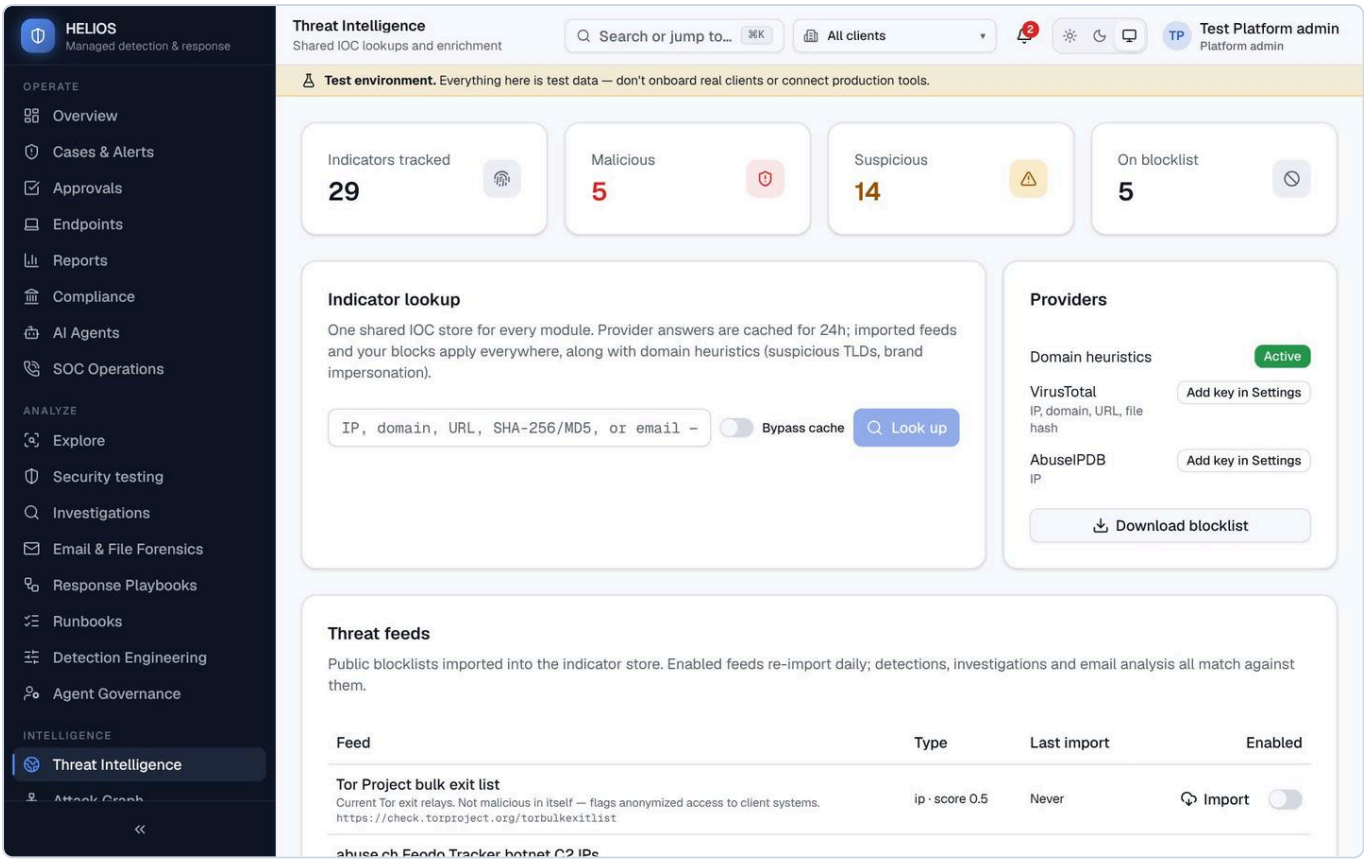


Figure 49. Threat Intelligence, as seen by a platform admin (test environment).

How to use it

- Paste an indicator (defanged forms like hxxp and [.] are fine) and Look up. The result shows the verdict, score, sources, why, and for an IP where it is (country, city and network, from the server's GeoIP database). Results are cached for 24 hours; Bypass cache forces a fresh lookup.
- Block or unblock an indicator: the blocklist feeds detections and firewalls that pull it. Enable and import threat feeds.

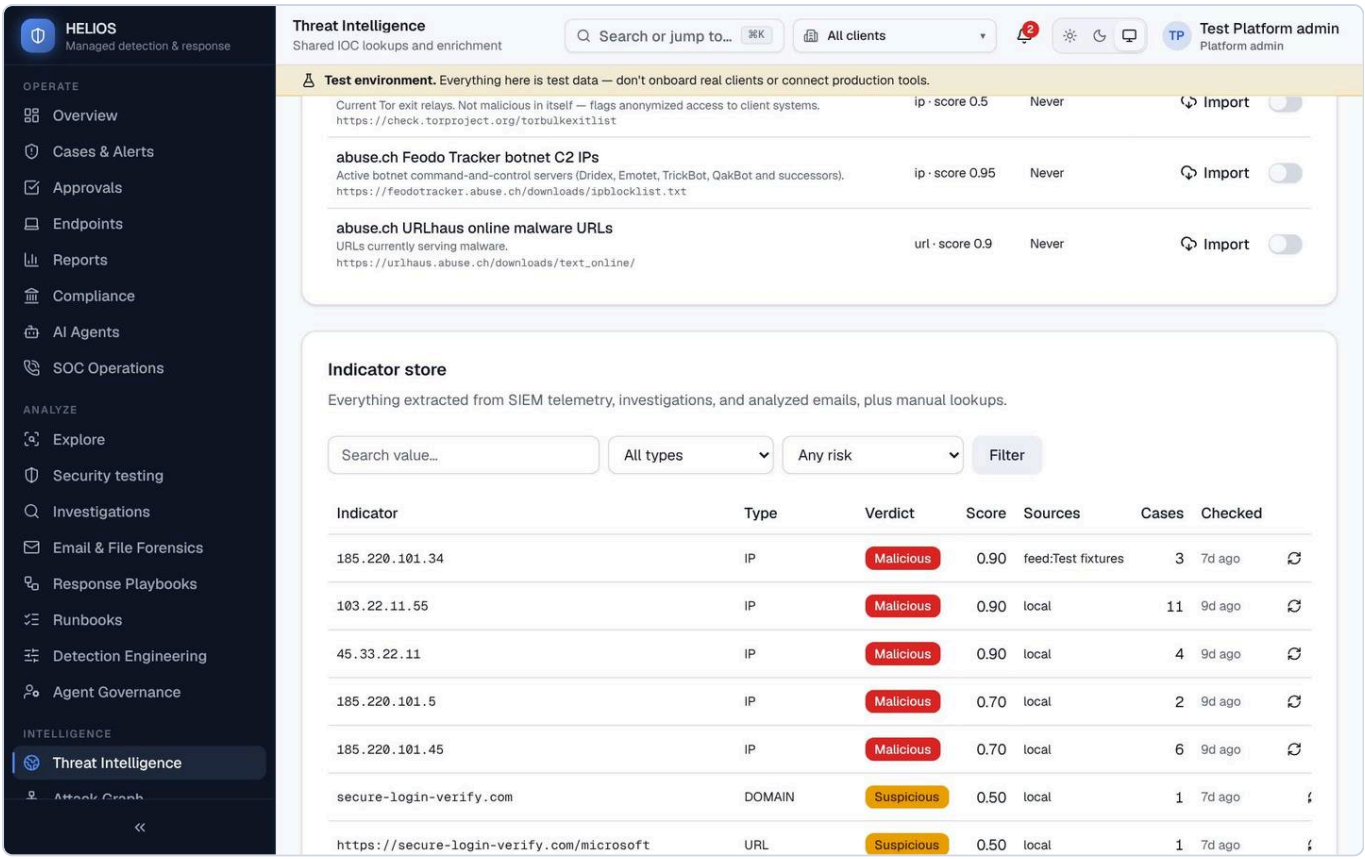


Figure 50. Threat Intelligence, further down the page.

Attack Graph

Cases joined by the hosts, accounts, addresses, indicators and agents they share. An entity is an attack

chain when its cases together reach more ATT&CK tactics than any one alone.

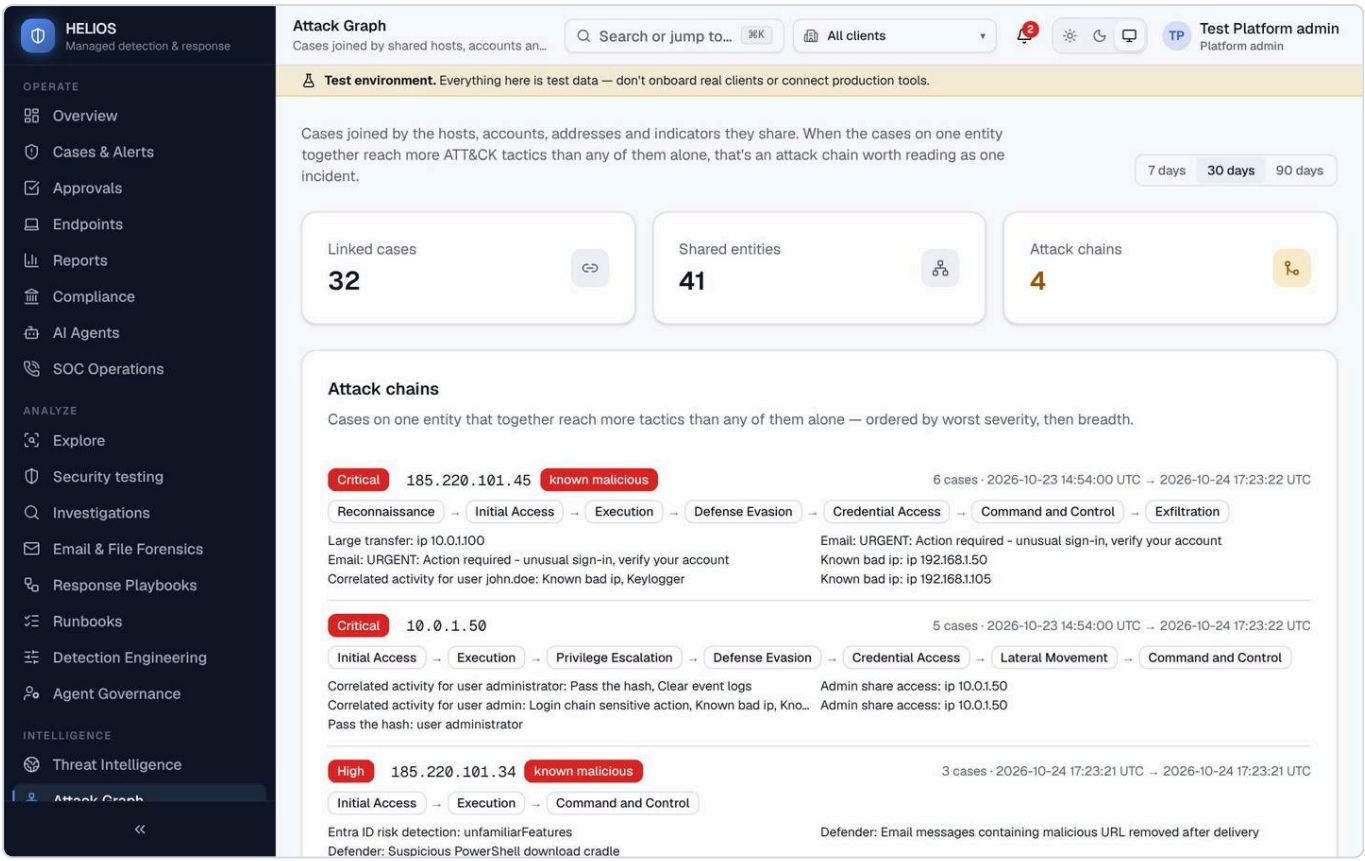


Figure 51. Attack Graph, as seen by a platform admin (test environment).

How to use it

- Hover or click a node to see the cases behind it; follow chains across tactics to see an intrusion that several smaller cases make up.

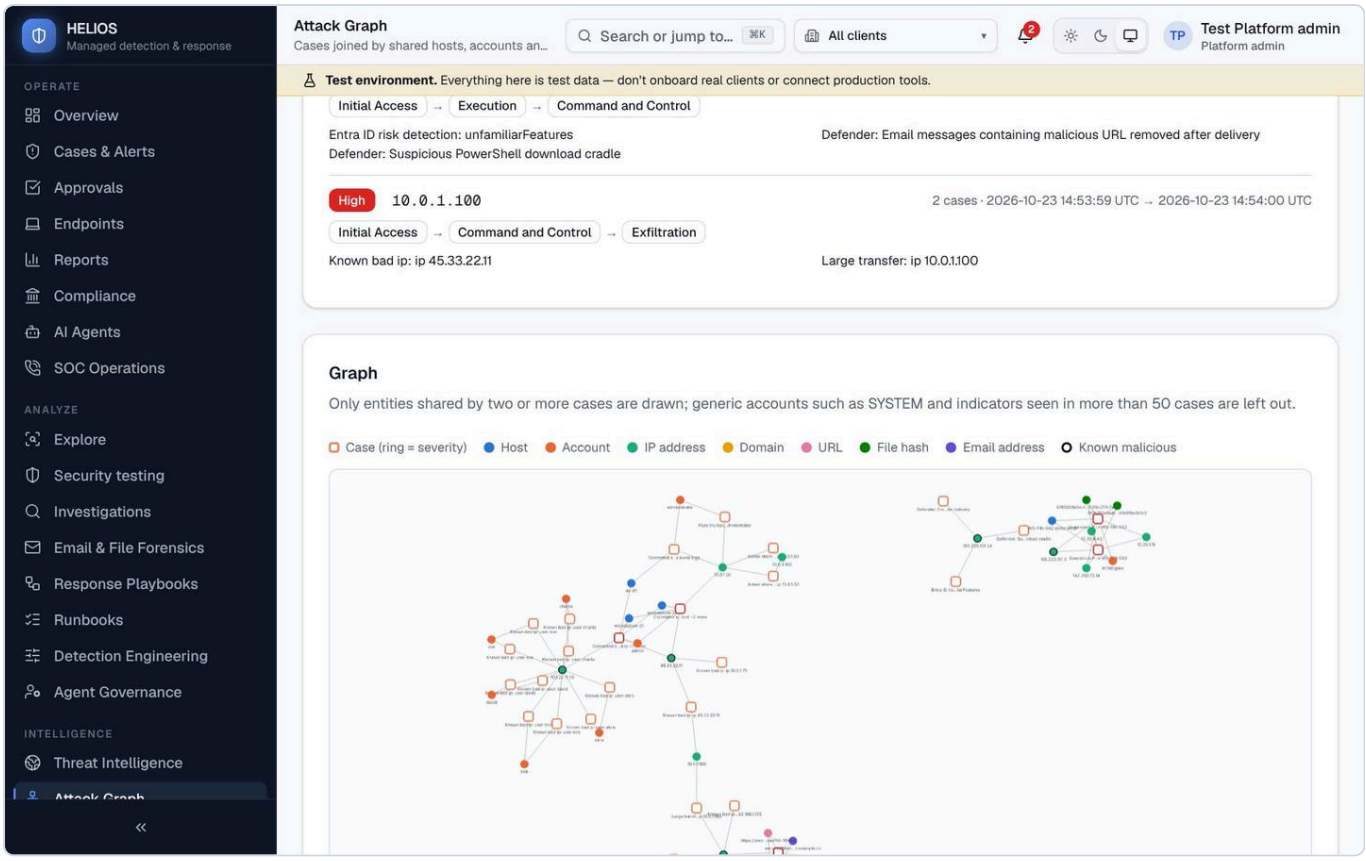


Figure 52. Attack Graph, further down the page.

MITRE ATT&CK Coverage

Which ATT&CK techniques HELIOS can detect, across built-in detectors, Sigma rules and onboarded

tools, and which it has seen.

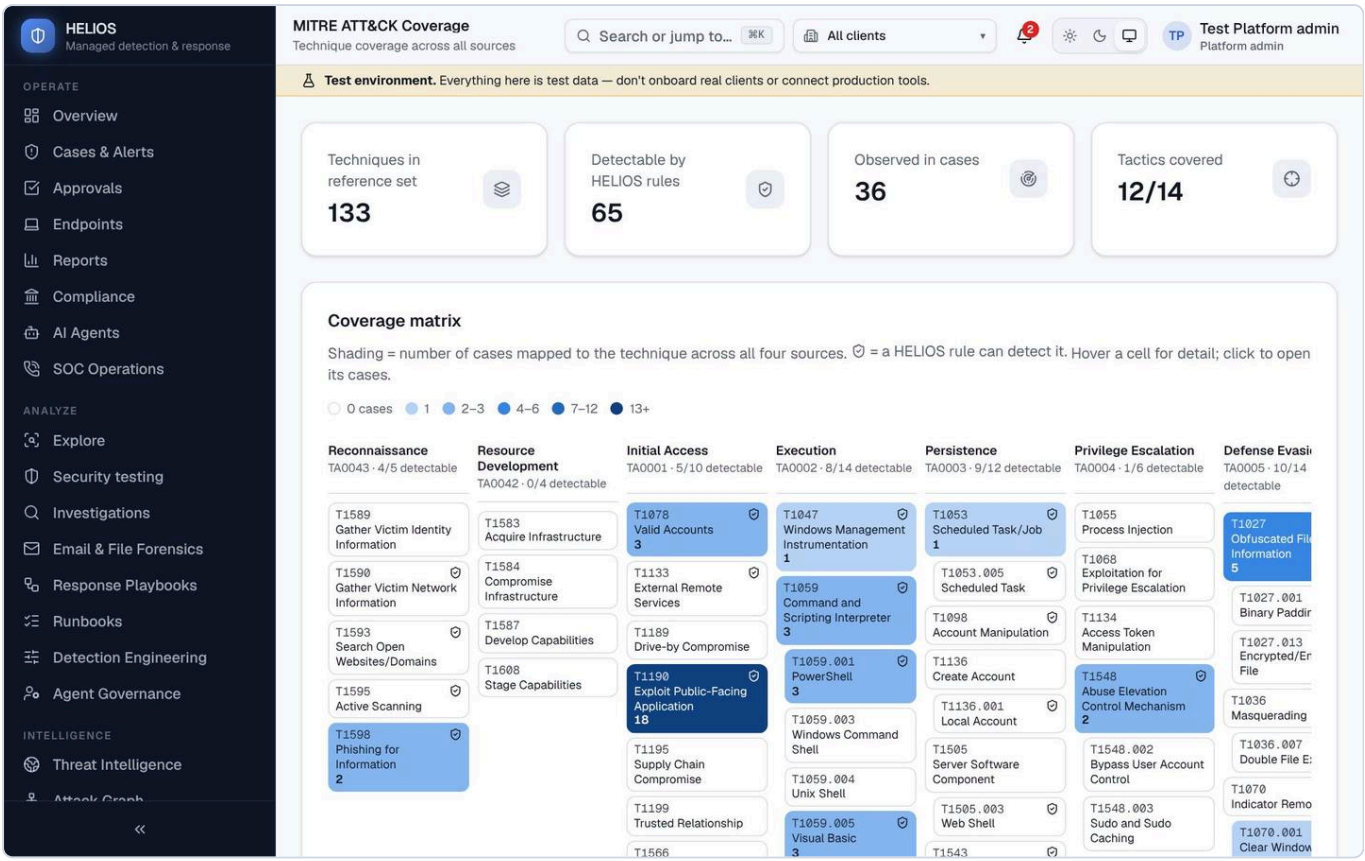


Figure 53. MITRE ATT&CK Coverage, as seen by a platform admin (test environment).

How to use it

- Read coverage by tactic and technique; gaps show where detection is missing.
- Switch built-in detectors on and off from here too.

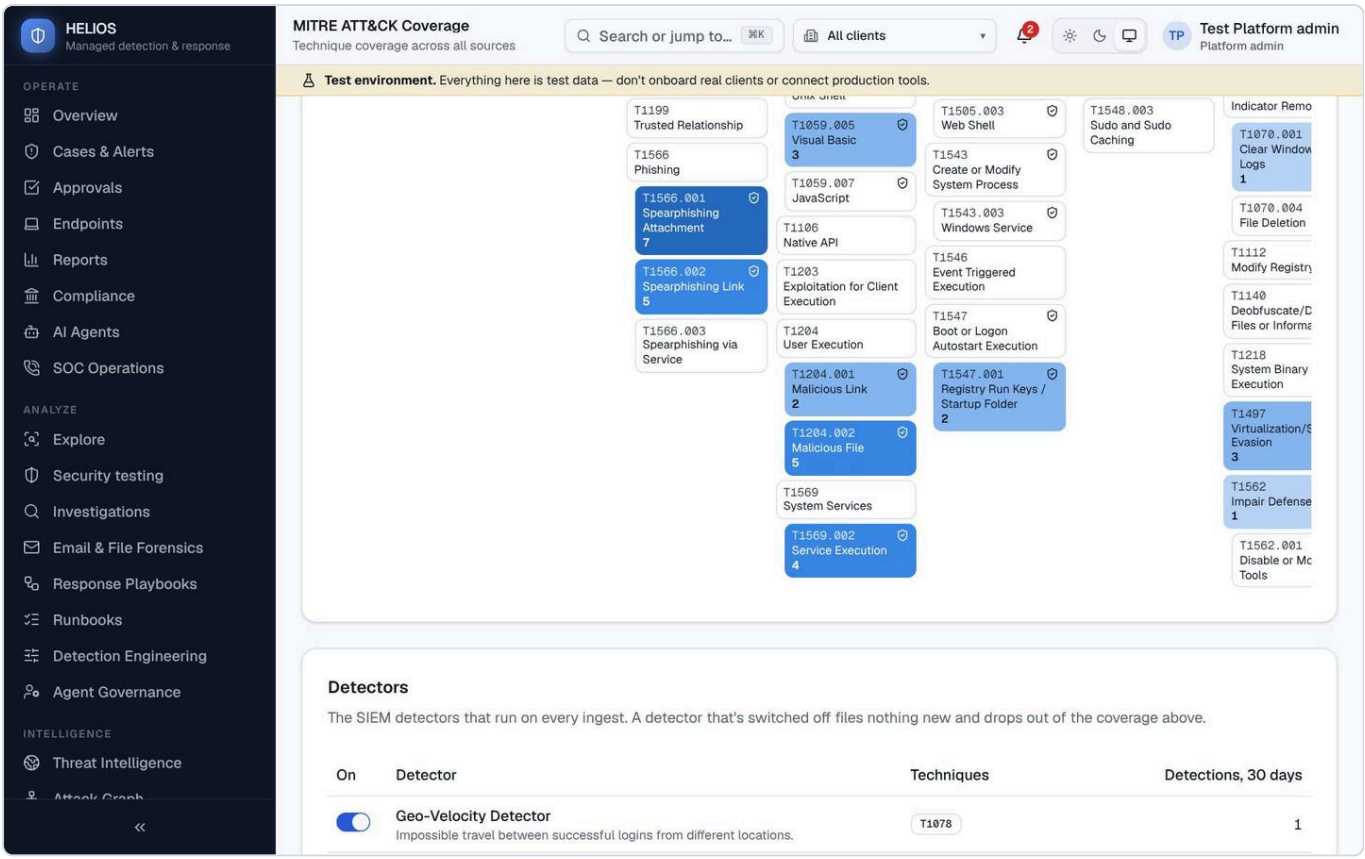


Figure 54. MITRE ATT&CK Coverage, further down the page.

Tools & Integrations

The security tools onboarded for each client: how their logs reach HELIOS and how HELIOS acts through

them.

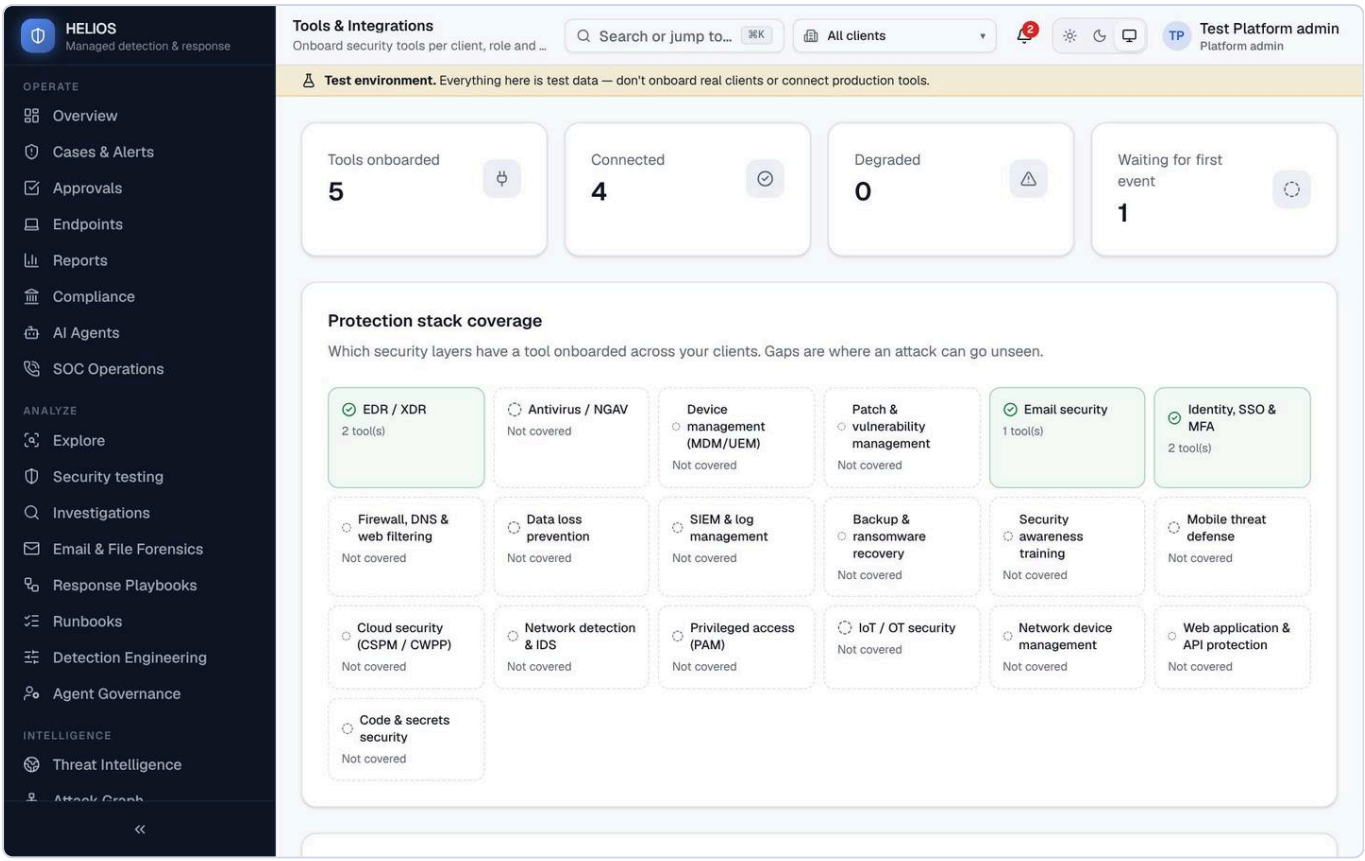


Figure 55. Tools & Integrations, as seen by a platform admin (test environment).

How to use it

- The tiles count tools onboarded, connected, degraded and waiting for their first event. Protection stack coverage shows which security layers have a tool across your clients; gaps are where an attack can go unseen.
- Each tool shows how it sends logs (push to its ingest URL, HELIOS pulling from the vendor's API, or the HELIOS forwarder), its source health, and whether a response connector can act through it.
- Onboard a tool for a client: pick it from the catalog, enter its read-only credentials (encrypted with the client's own key), test the connection, set the collectors running, and turn on the response connector (dry-run first). Generate the forwarder configuration for syslog, Windows event logs and files. Replay records waiting in the dead-letter queue.
- A source silent longer than its limit opens a case: attackers often stop logging first. A client can't go live while a source is dead.

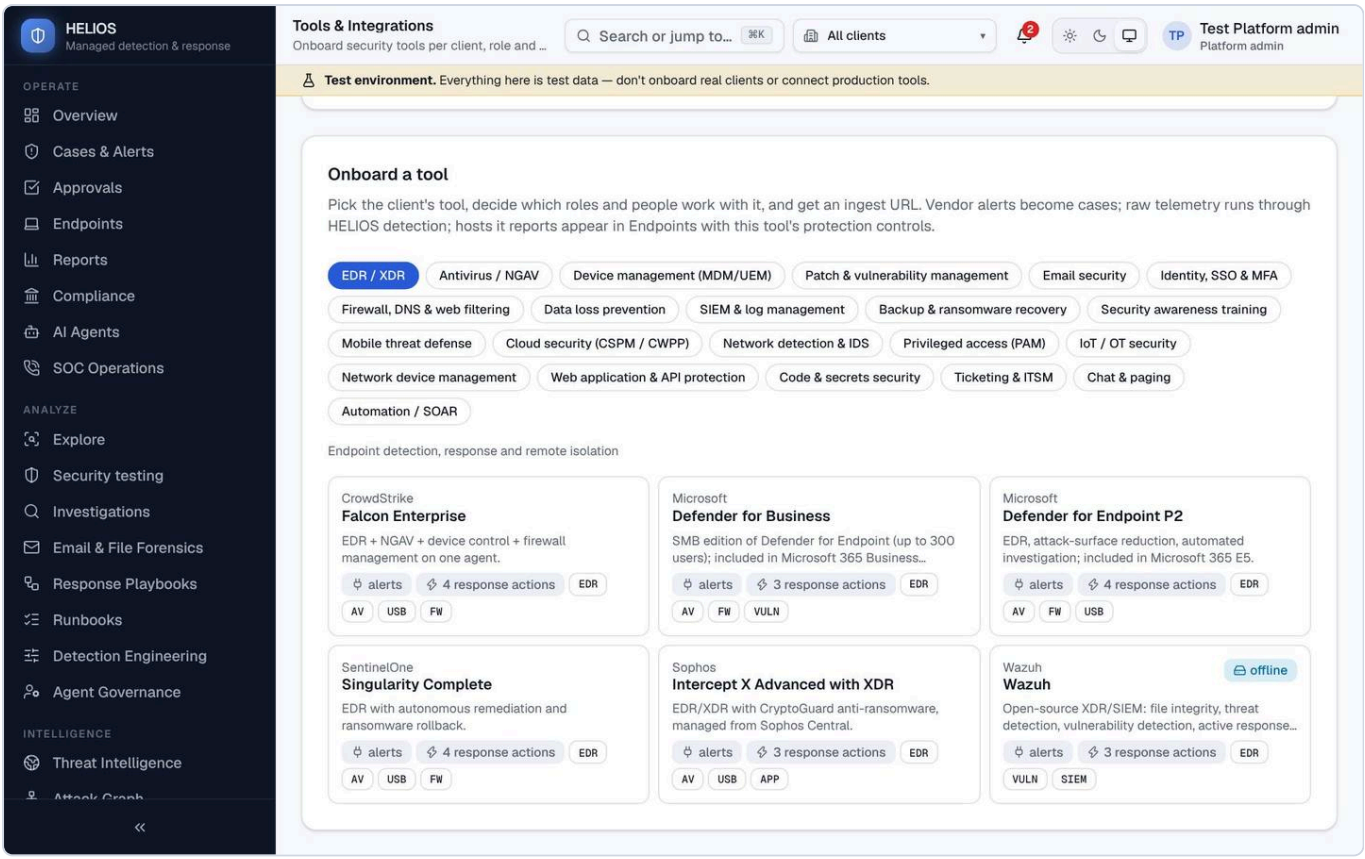


Figure 56. Tools & Integrations, further down the page.

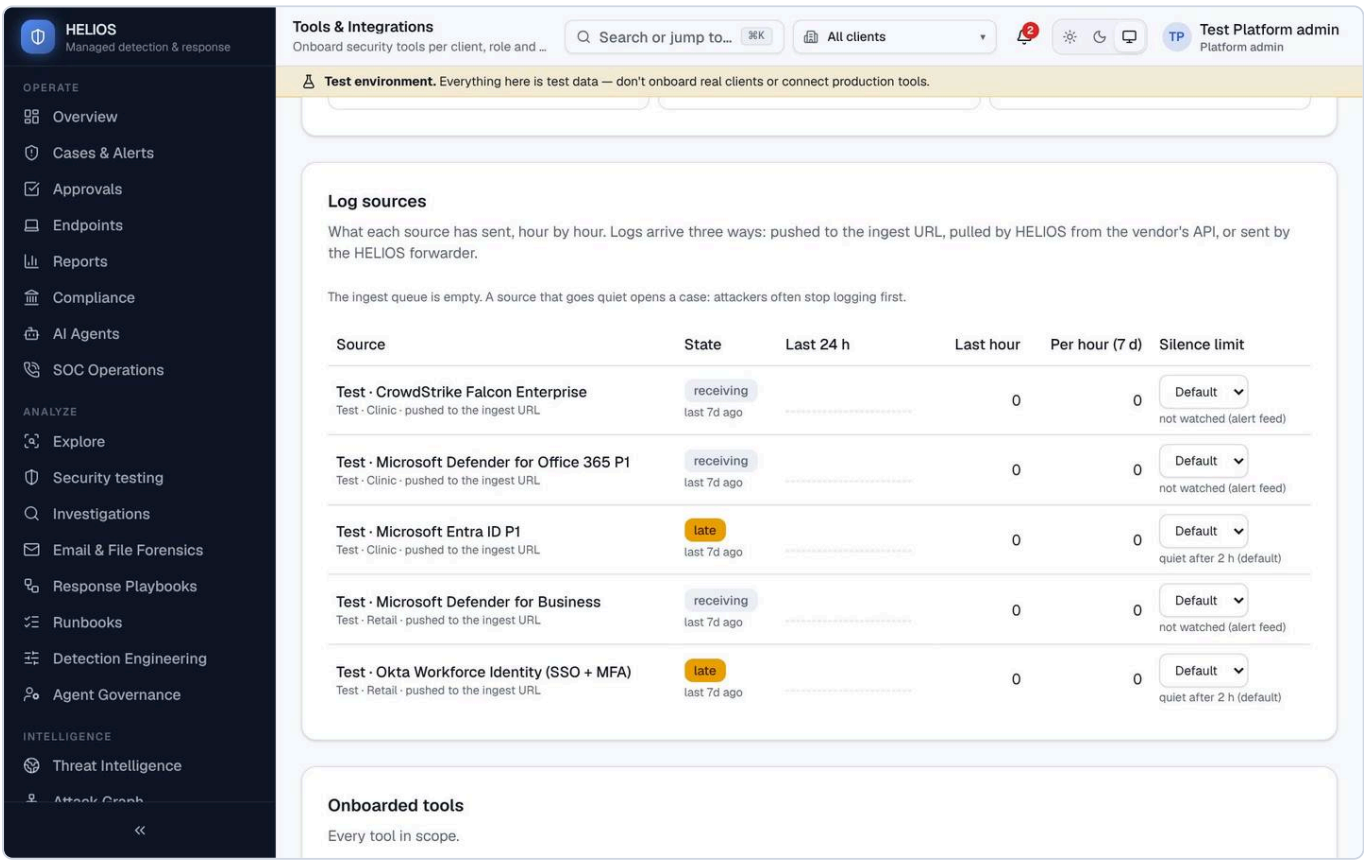


Figure 57. Tools & Integrations, further down the page.

Rulebooks

Each client's approved rules for what HELIOS may do on its estate. The rules only ever restrict: protected assets and accounts, which actions run automatically, need approval or never happen, blast-radius limits, business hours and change freezes, the approver chain, AI data handling and CERT-In contacts.

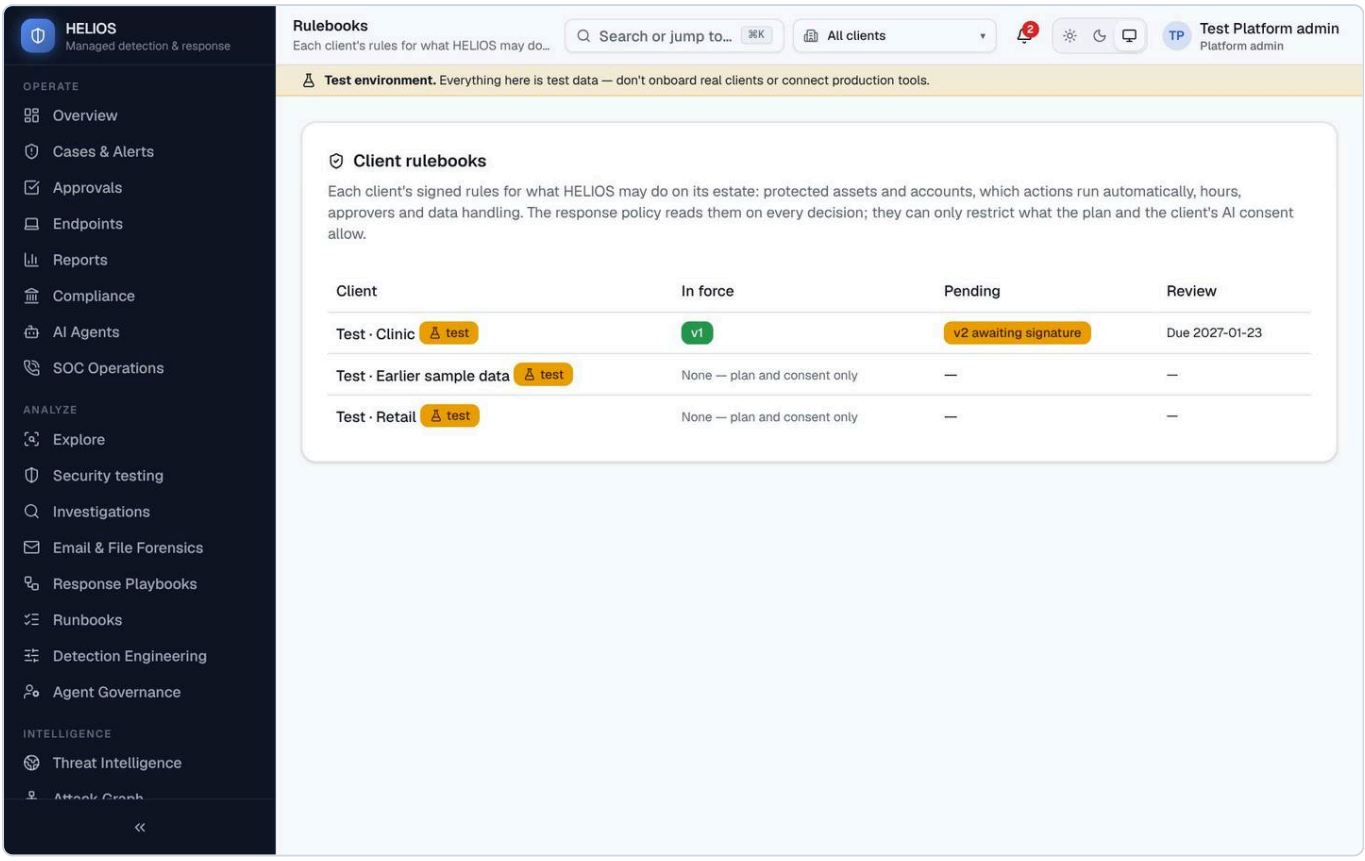


Figure 58. Rulebooks, as seen by a platform admin (test environment).

How to use it

- Pick the client at the top. The page shows the version in force, its hash and when it is due for review.
- The tabs: In force (the rules now applying), Edit, Dry run, Allowlist, Versions, AI models, Autonomy and Policy as code.
- Edit drafts a change, from an industry template or the version in force. Dry run replays the last 30 days of AI proposals under the draft and probes each action. Versions compares any two versions side by side.
- Send the draft for approval. It takes effect only once the client's admin and a HELIOS approver have both approved it, each after reviewing what changes and the dry run on the last 30 days; until then the version in force keeps deciding.
- Approve: client admins approve for their organisation; HELIOS's approvers (SOC managers and admins) approve for HELIOS, and can record a client's approval given elsewhere, naming it. Someone other than the person who asked approves for HELIOS.
- Allowlist: known-good hosts, IPs and accounts, each with an owner, a reason and an expiry. They open no cases, and AI containment against them is only suggested. A new entry applies once the client and HELIOS have approved it.
- Policy as code: export or import the rulebook as YAML, or keep it in Git (GitHub, GitLab or an https URL): a changed file becomes a draft that still needs both approvals here.

- AI models: which models may read the client's evidence (any, region-locked or on-premises only), the client's own providers (bring your own key), a provider per agent, and the monthly AI budget.
- Autonomy: the AI's autonomy ladder per action type for this client. New clients contain on their own from day one (initial containment); a promotion beyond that takes effect once the client and HELIOS approve it.

Rule changes

Changes to the rules containment runs under, waiting for approval: new rulebook versions, allowlist entries, more AI autonomy or consent, and platform AI settings. Nothing here is in effect yet: the rules in force keep deciding until every approval a change needs is in.

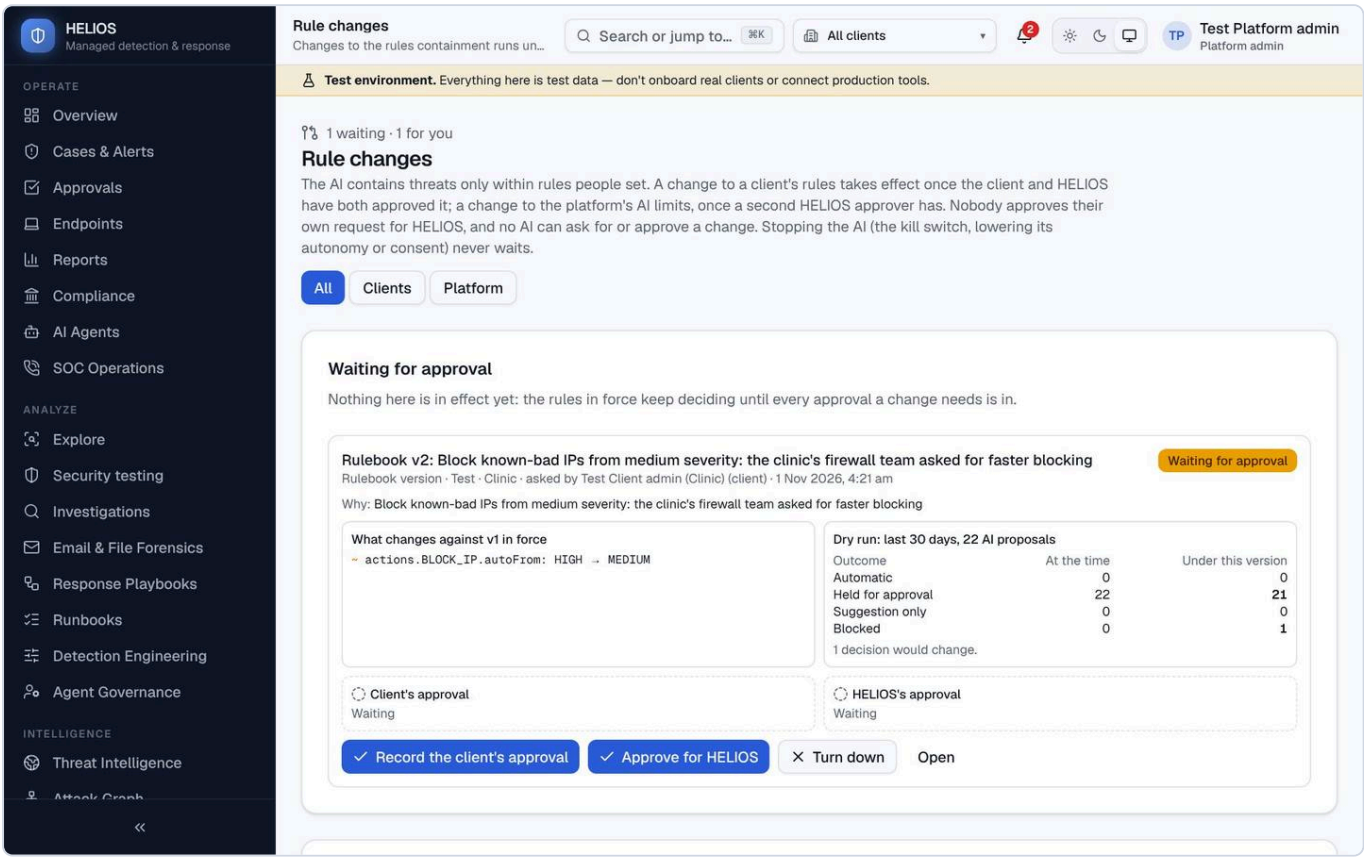


Figure 59. Rule changes, as seen by a platform admin (test environment).

How to use it

- Each change says what it is, the client, who asked and why, and which approvals it still needs: the client's and HELIOS's, or for a platform change, a second HELIOS approver.
- Impact shows exactly what changes (each rule, before and after) and, for a rulebook, the dry run: how the last 30 days of AI containment proposals would have gone under the new version.
- Approve or Reject for your side. Someone other than the person who asked approves for HELIOS. If the rules in force change while it waits, the change goes stale and must be asked again.
- Lowering the AI's autonomy or consent, and the kill switch, apply at once without waiting here.
- Decided lists past changes and who approved or rejected them.

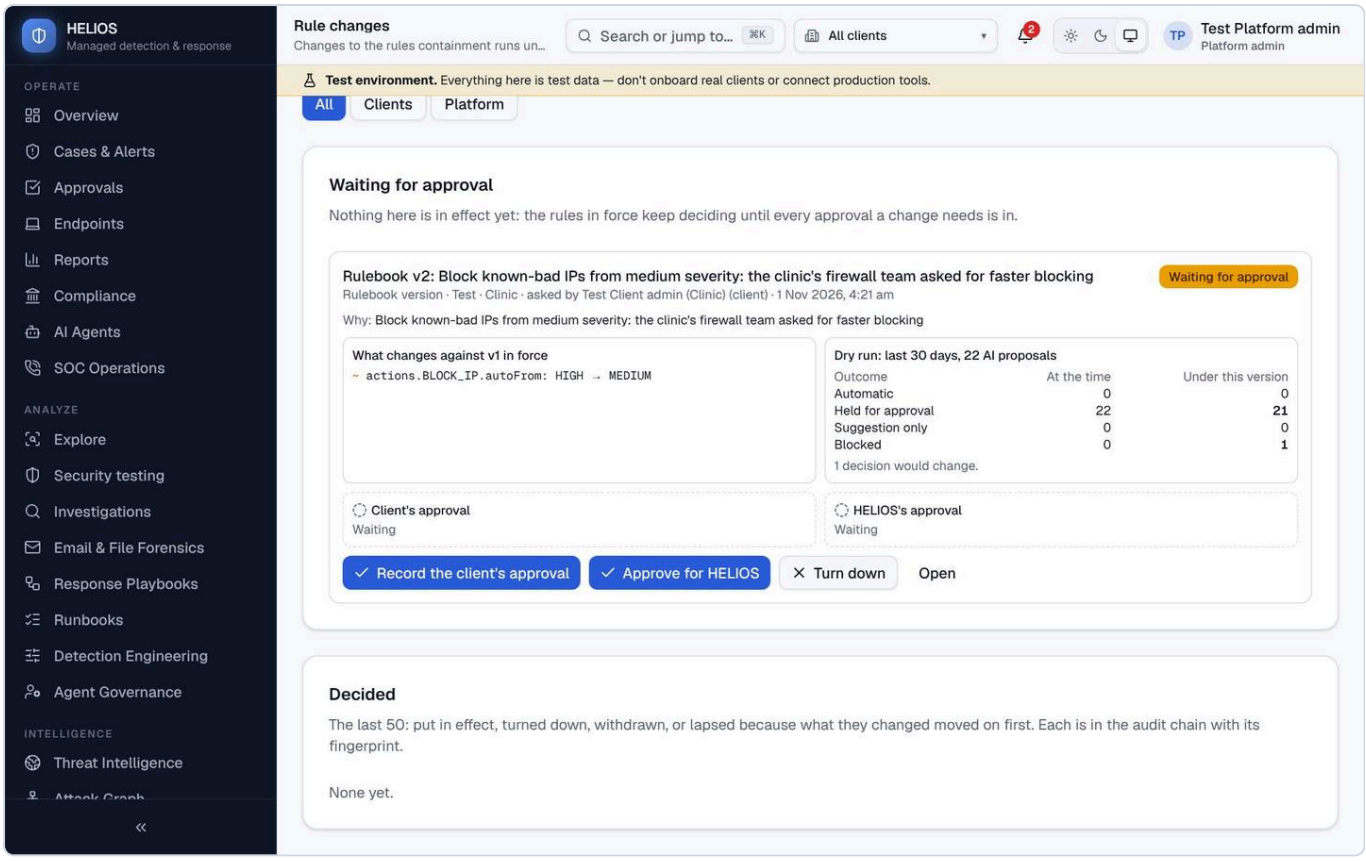


Figure 60. Rule changes, further down the page.

Clients

The client organisations (tenants): their plan and SLA, status, contacts, and how far they consent to AI

acting on their estate.

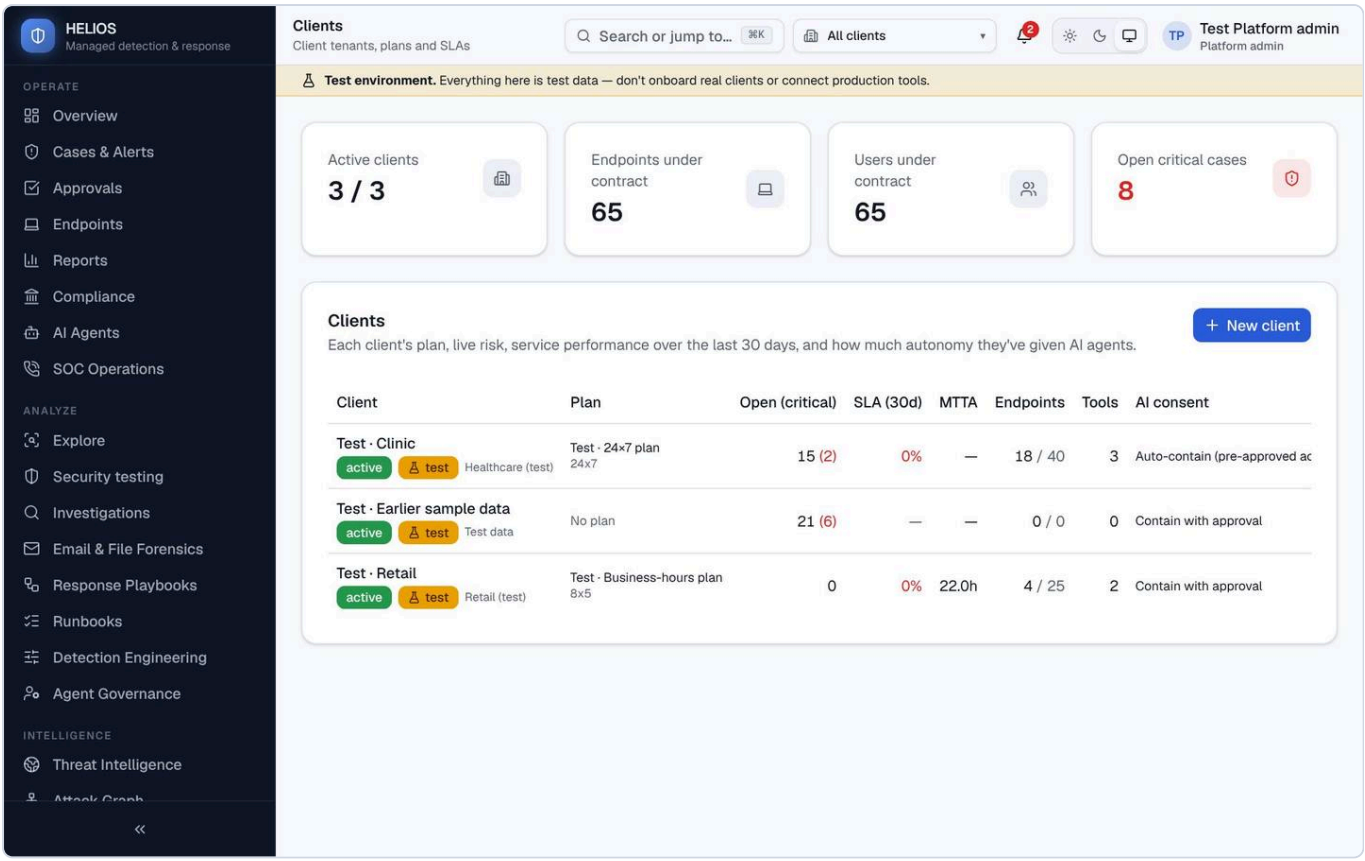


Figure 61. Clients, as seen by a platform admin (test environment).

How to use it

- Each client shows its plan, status (onboarding, active, suspended), licensed endpoints and users, and test clients are badged.
- Add or edit a client: name, industry, contacts, plan, licensed counts, and the AI autonomy cap the client has consented to (it caps every agent on their estate). By default HELIOS contains on its own from day one within the client's approved rulebook; Start in shadow mode has the AI only observe and earn autonomy first. A client moves from onboarding to active only once every data source has sent something in the last day.

Team & Roles

The people who use HELIOS, their roles, the security domains (queues) analysts work, and what each role

can see and do.

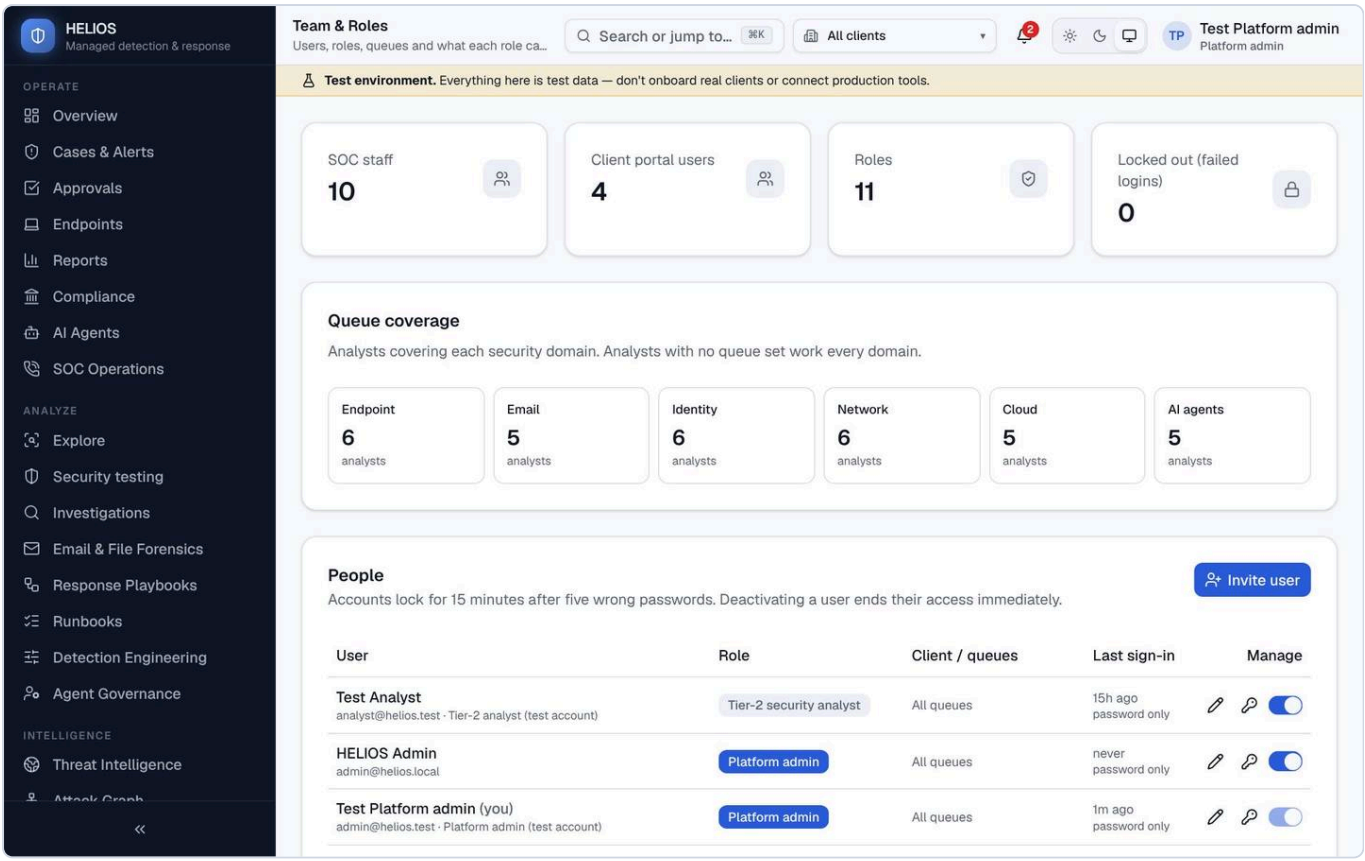


Figure 62. Team & Roles, as seen by a platform admin (test environment).

How to use it

- Invite user: name, email, role, and for client users their organisation. HELIOS shows a temporary password once; pass it to the person securely. Limit an analyst to security domains to give them those queues only.
- On each person's row: the pencil edits the role, organisation and queues; the key resets the password; the switch deactivates the account. Changes apply on the person's next request. Reset two-step sign-in for someone who lost their phone.
- Queue coverage shows how many analysts cover each security domain. Accounts lock for 15 minutes after five wrong passwords; the Locked out tile counts them.
- The role matrix shows every permission each role has.

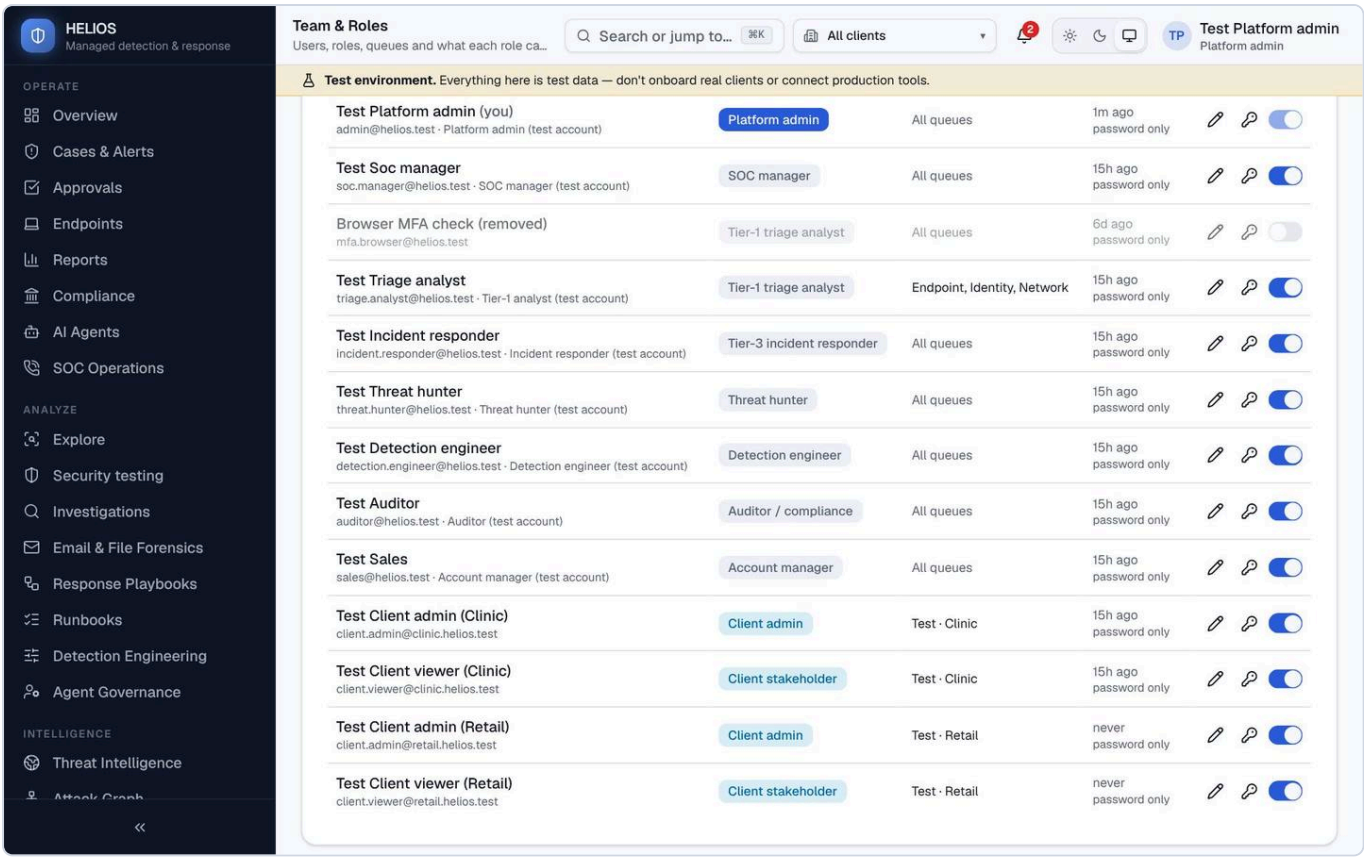


Figure 63. Team & Roles, further down the page.

Data Retention

Where each client's logs are kept and for how long, the encrypted archive, and legal holds.

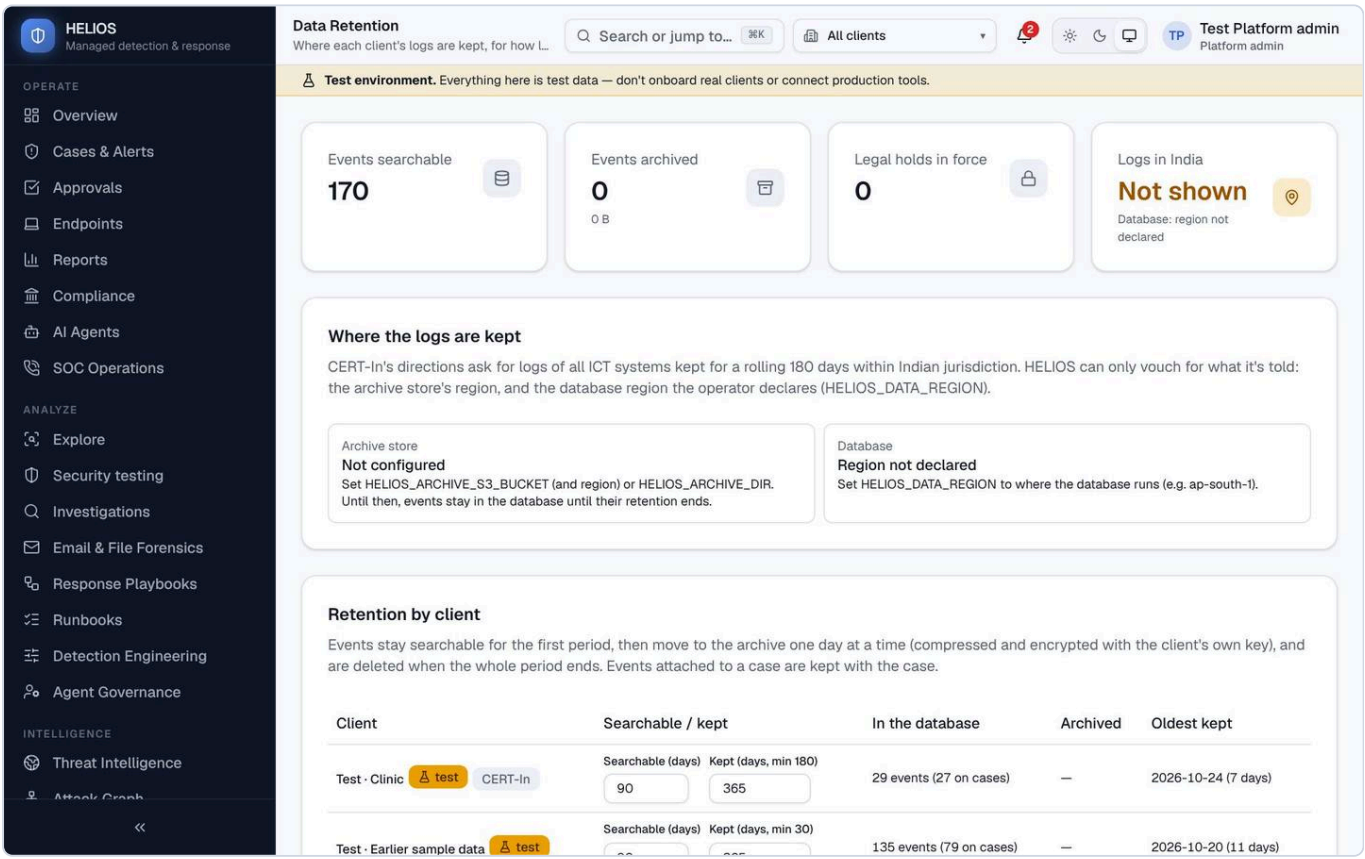


Figure 64. Data Retention, as seen by a platform admin (test environment).

How to use it

- Retention by client: Searchable (days) and Kept (days). Where CERT-In's directions apply, Kept can't go below 180. Events attached to a case are kept with the case.
- The page says whether the logs are shown to be in India (the database region and the archive store's region), from what the server is configured with.
- Legal holds stop deletion over a date range (optionally tied to a case); placing and releasing need a reason. Archives download decrypted, and each download is audited.

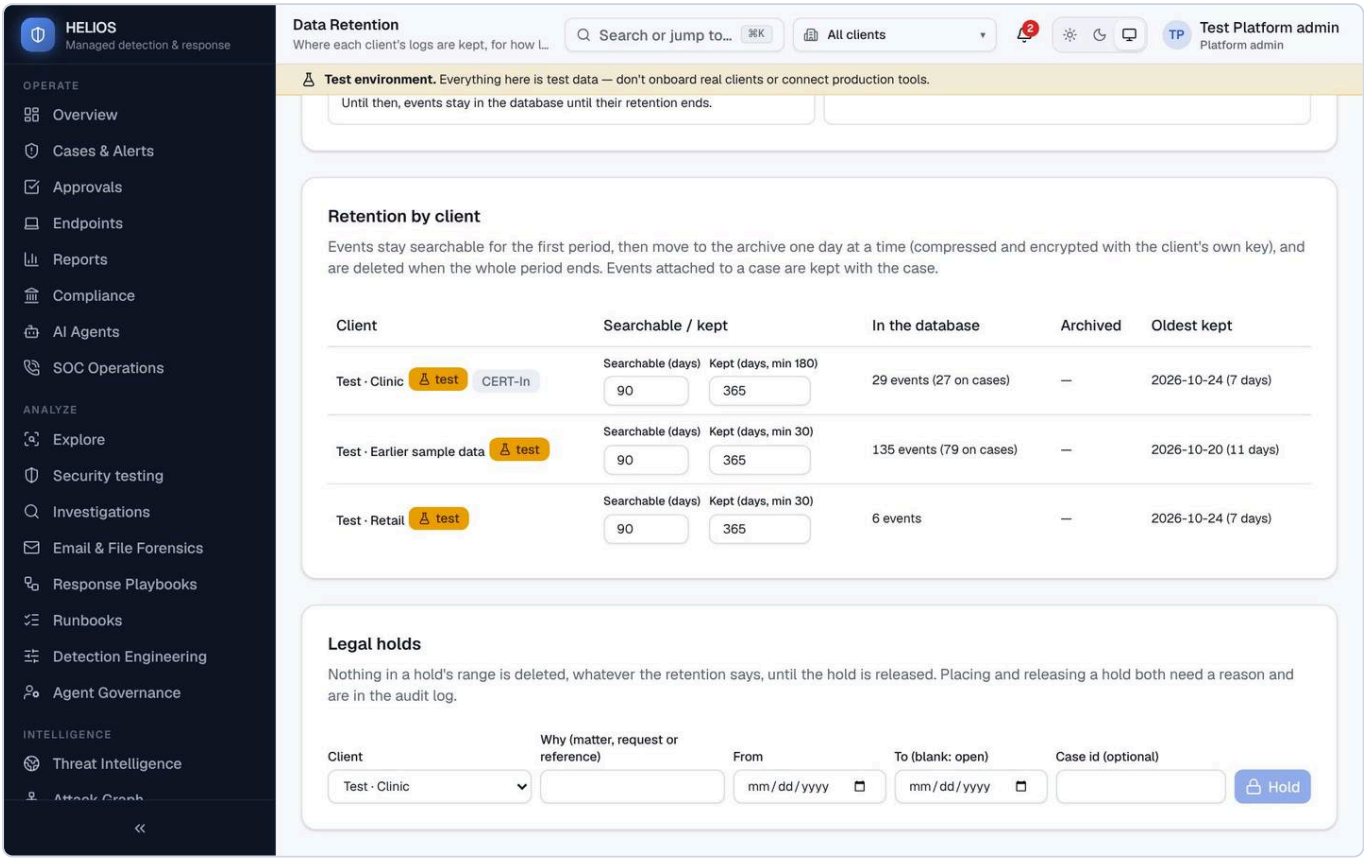


Figure 65. Data Retention, further down the page.

Audit Log

The hash-chained record of every action in HELIOS: who did what, when, to which case or object.

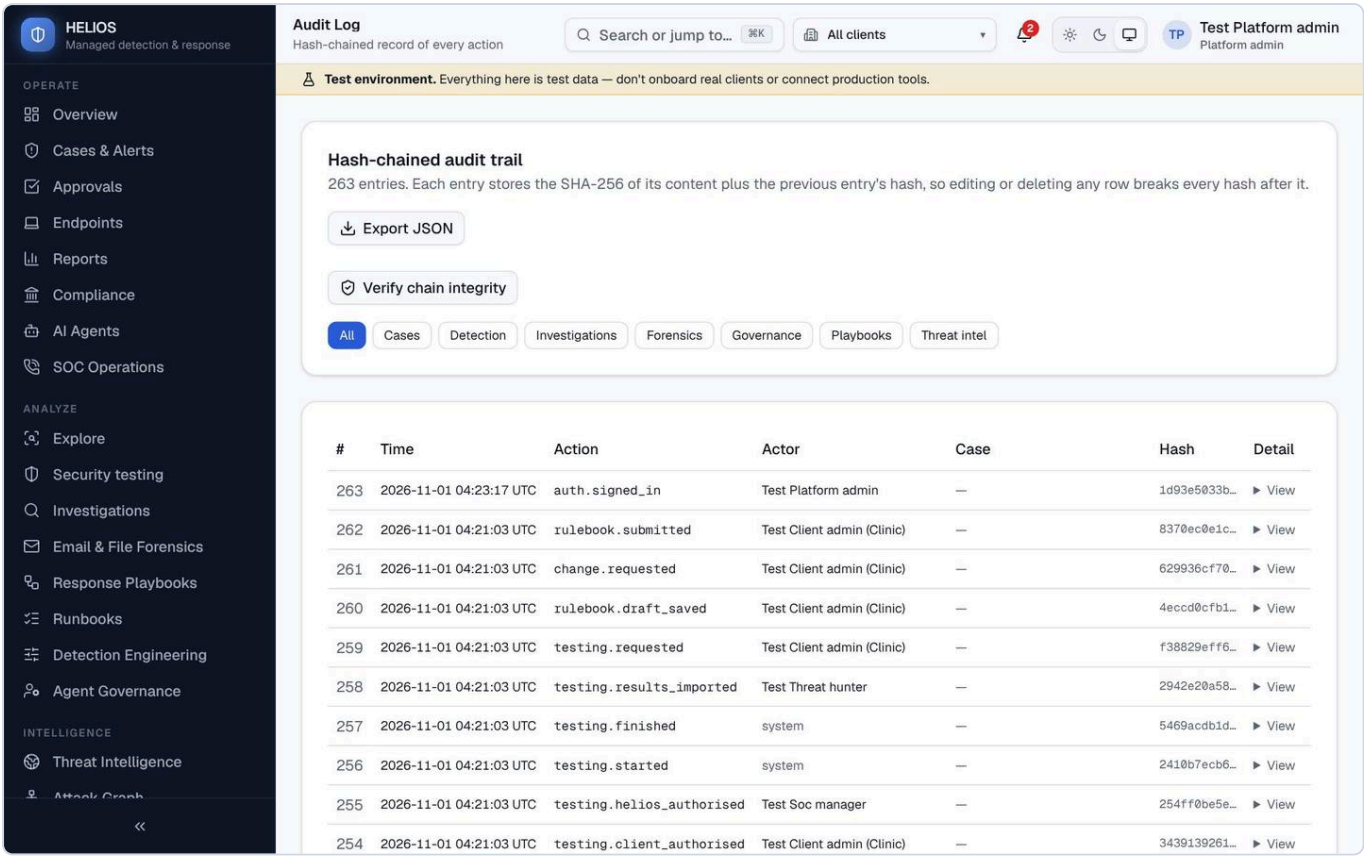


Figure 66. Audit Log, as seen by a platform admin (test environment).

How to use it

- Each entry stores the SHA-256 of its content plus the previous entry's hash, so editing or deleting any row breaks every hash after it. Verify chain integrity recomputes them all and names the first altered entry if there is one.
- Filter by area: cases, detection, investigations, forensics, governance, playbooks, threat intel.
- Export JSON exports the chain for an auditor or a regulator.

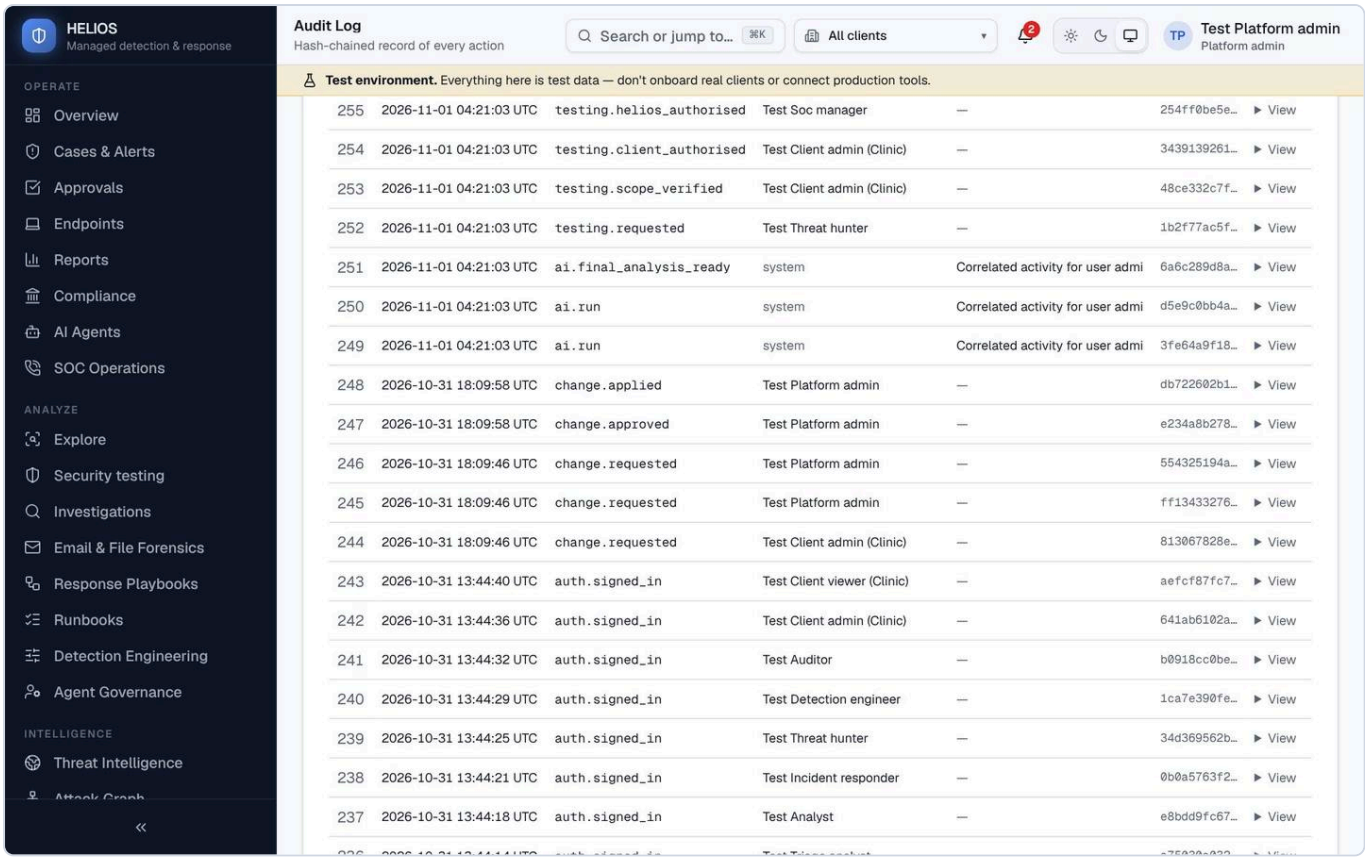


Figure 67. Audit Log, further down the page.

API Tokens

Personal tokens for the REST API, for ticketing systems, SOAR and client portals. A token acts as you:

same role, client and queues.

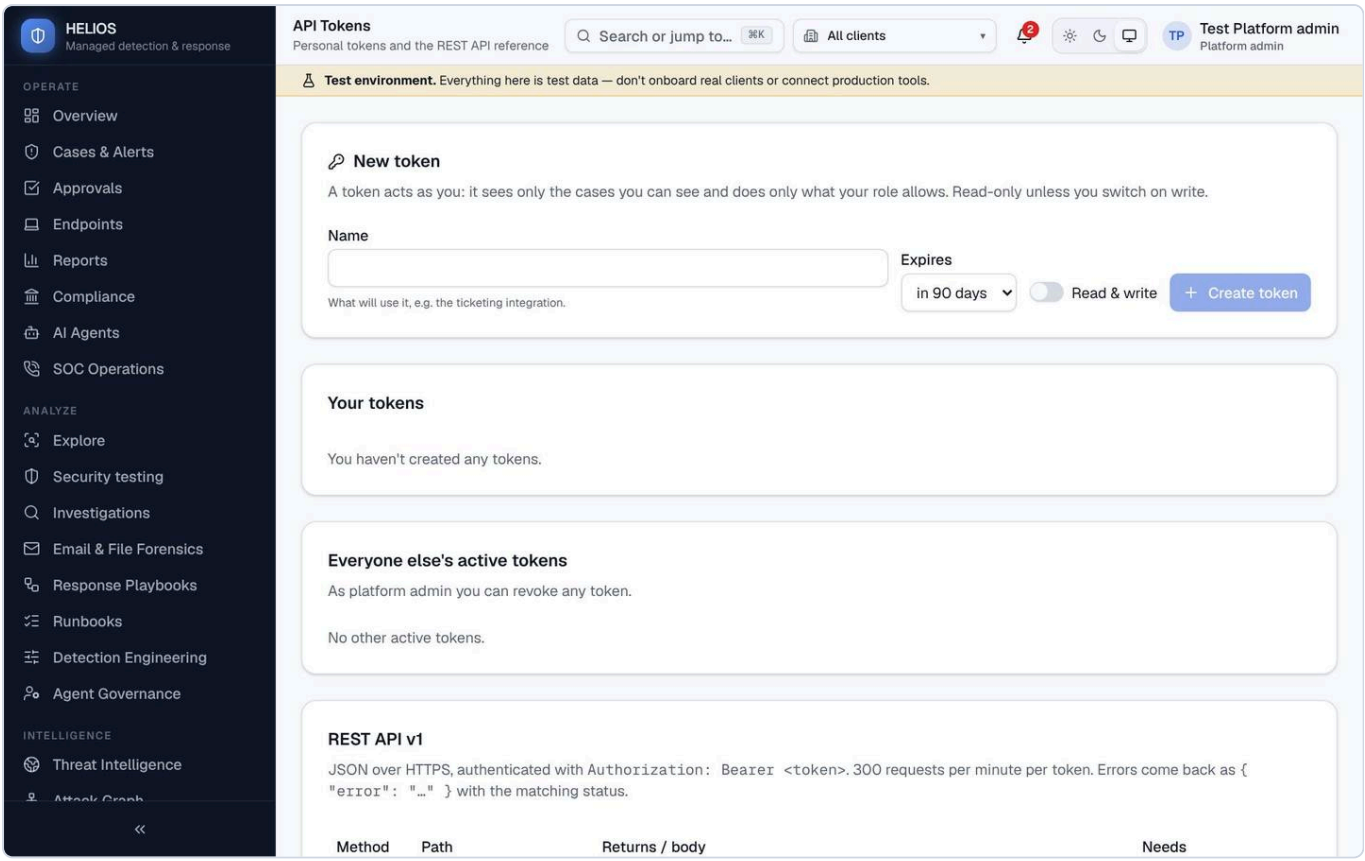


Figure 68. API Tokens, as seen by a platform admin (test environment).

How to use it

- New token: give it a name (what will use it), an expiry (90 days by default), and switch on Read & write only if it must change things. Create token shows it once; copy it then. Your tokens lists yours to revoke.
- As platform admin you also see everyone else's active tokens and can revoke any of them.
- The page includes the API reference: endpoints, parameters and examples. Calls are limited to 300 a minute per token, and every change through the API is audited.

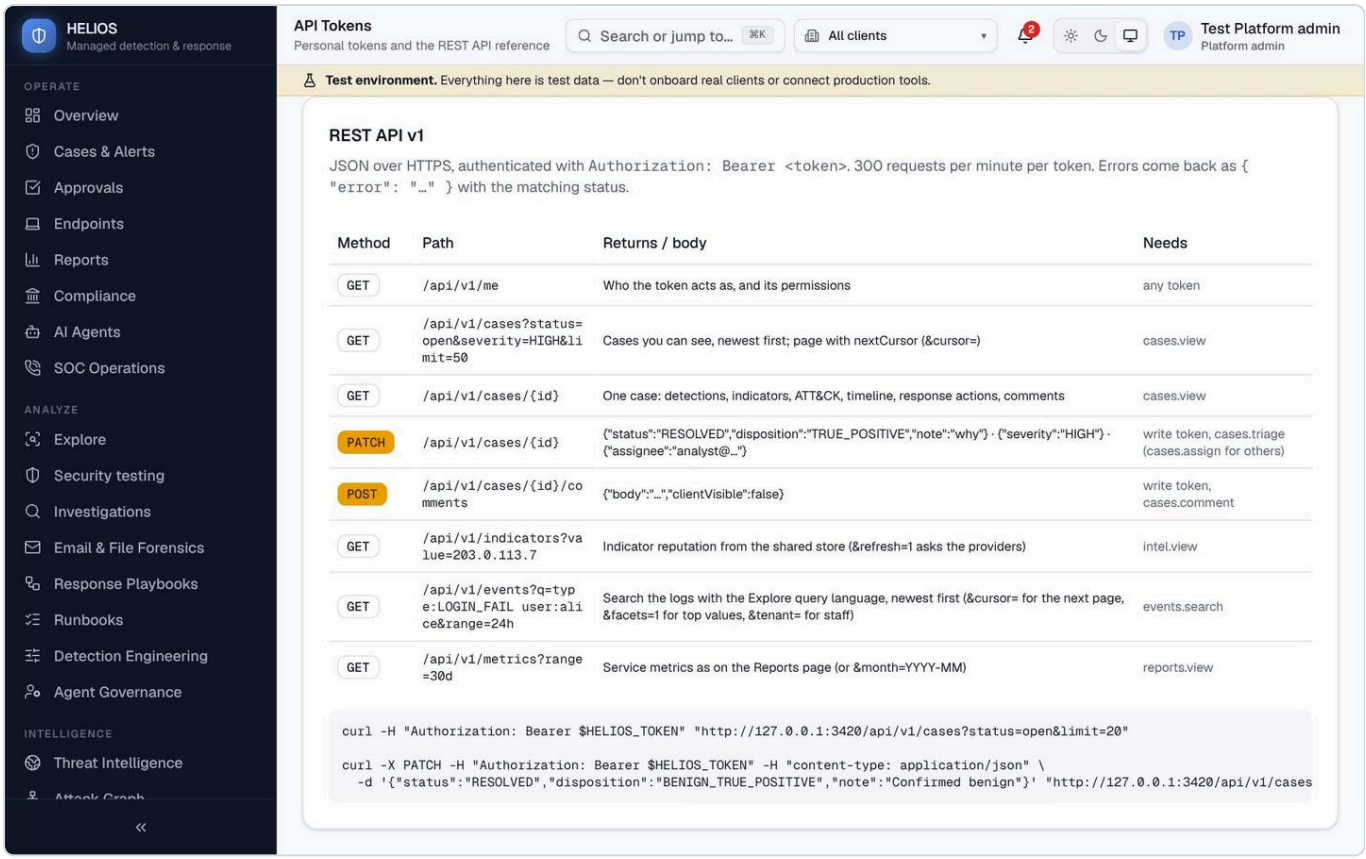


Figure 69. API Tokens, further down the page.

Settings

Platform-wide settings: environment, sign-in and encryption, company profile, public website, alert

channels, provider keys, forensics thresholds and IP geolocation.

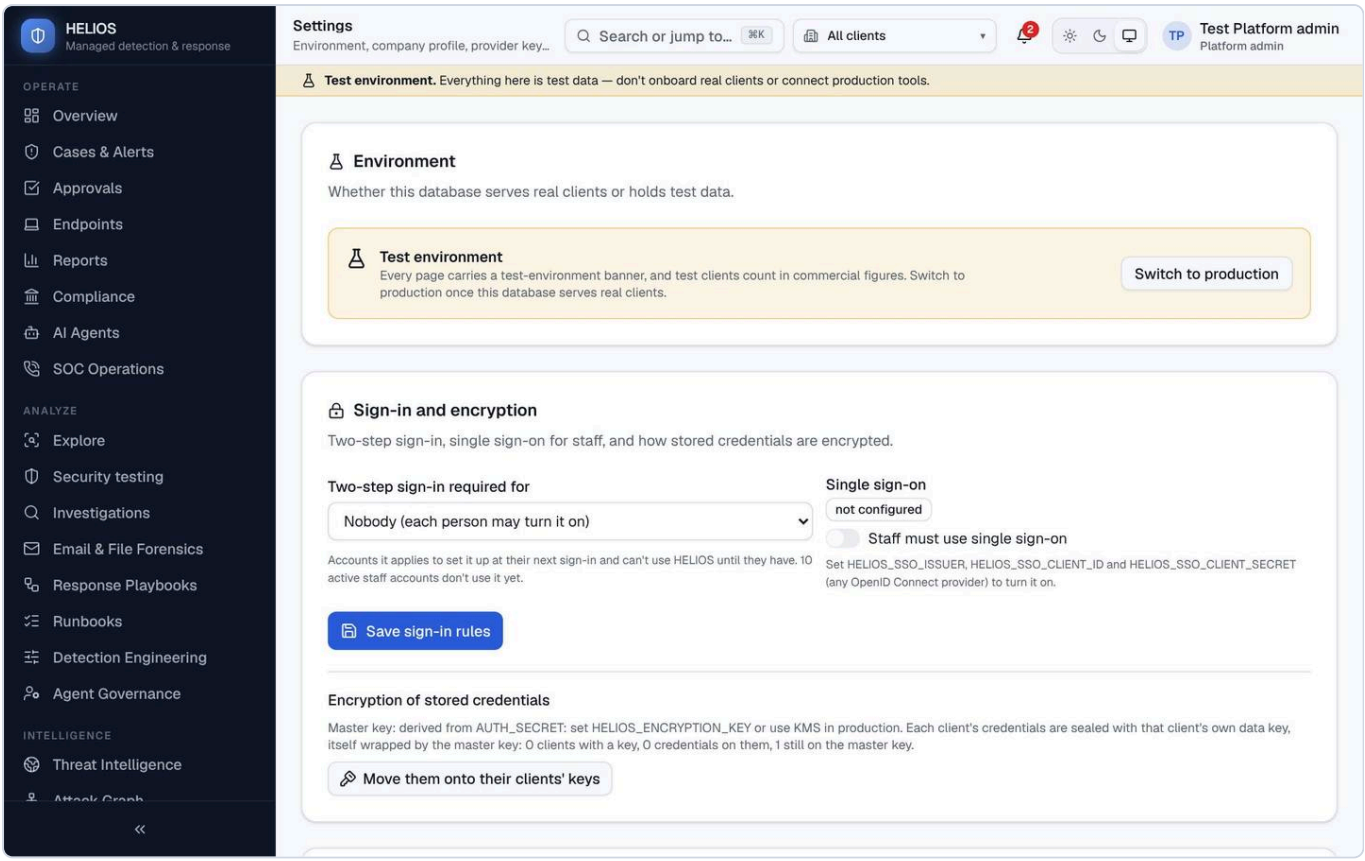


Figure 70. Settings, as seen by a platform admin (test environment).

How to use it

- Environment: production or test (a test environment shows a banner on every page). Switch to production once the database serves real clients.
- Sign-in and encryption: Two-step sign-in required for (nobody, staff or everyone) and Staff must use single sign-on, then Save sign-in rules. Under encryption, Move them onto their clients' keys re-seals older credentials with each client's own key.
- Company profile (printed on quotes, reports and invoices), the public website, alert channels (email, SMS, WhatsApp) and provider keys (VirusTotal, AbuseIPDB, MxToolbox, Slack). Keys are encrypted; only their last four characters are ever shown.
- Forensics thresholds and trusted link domains, and the IP geolocation databases in use.

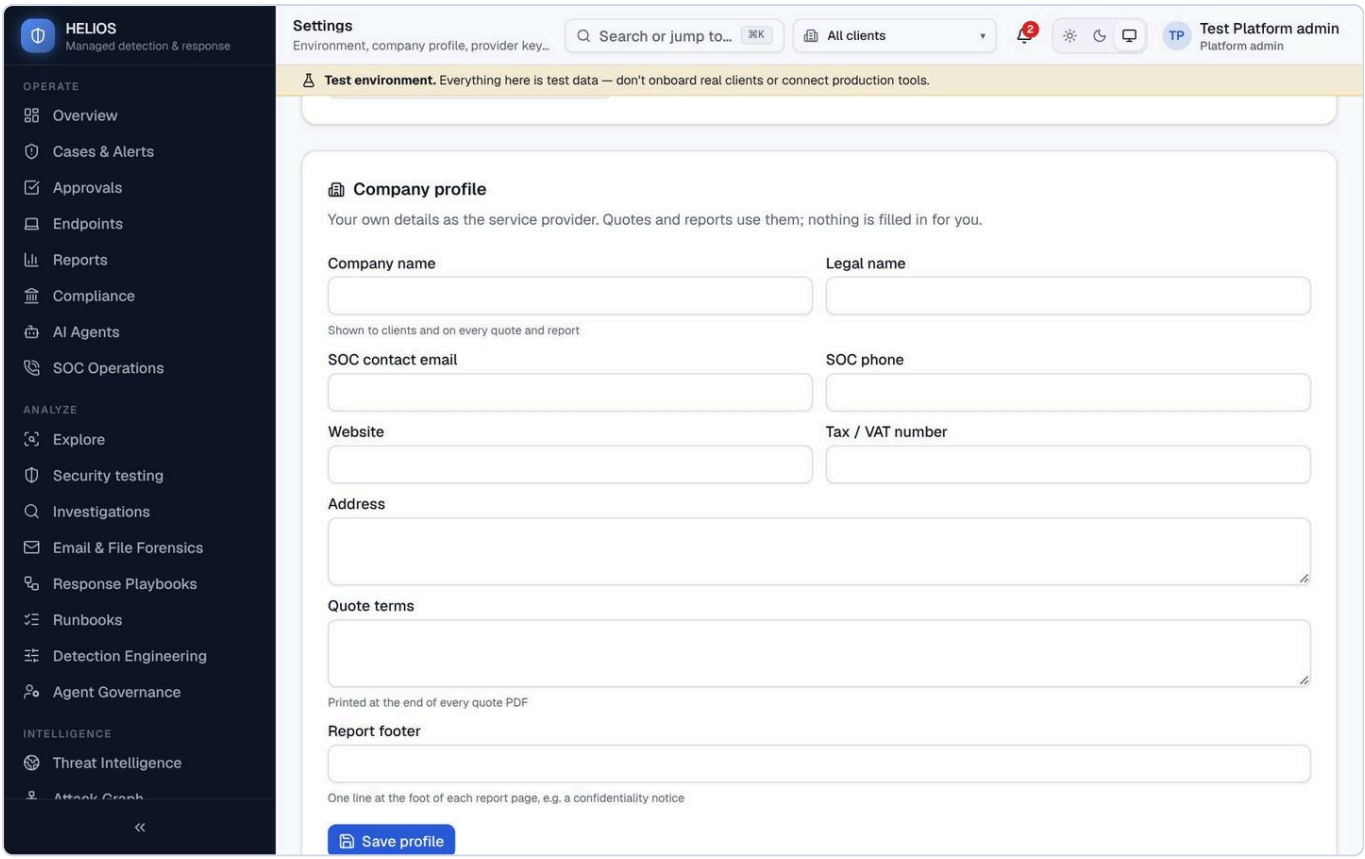


Figure 71. Settings, further down the page.

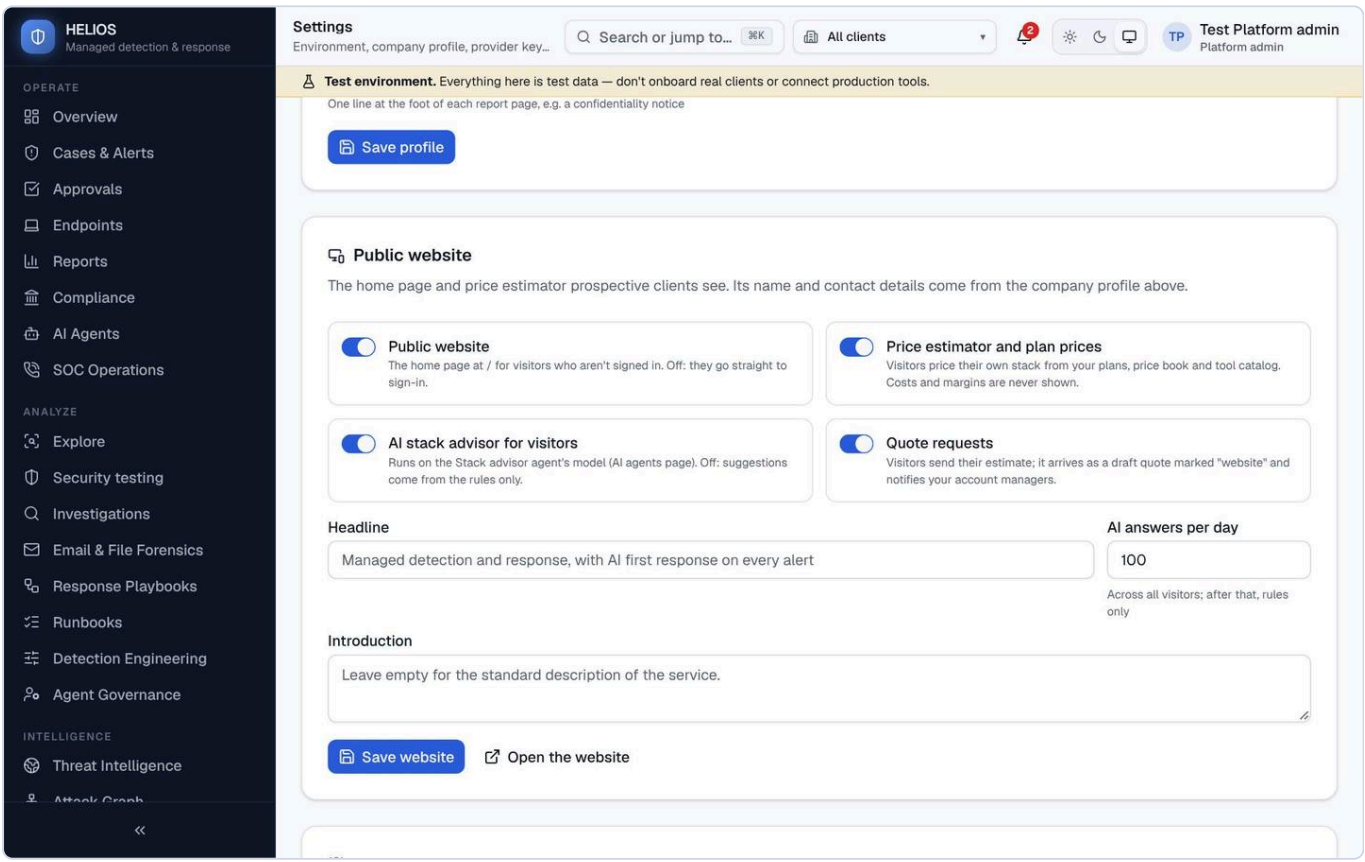


Figure 72. Settings, further down the page.

Build & Price

Configure a client's protection stack, compare vendors and plans, and estimate what it costs: monthly, annual, one-time, maintenance and over three years.

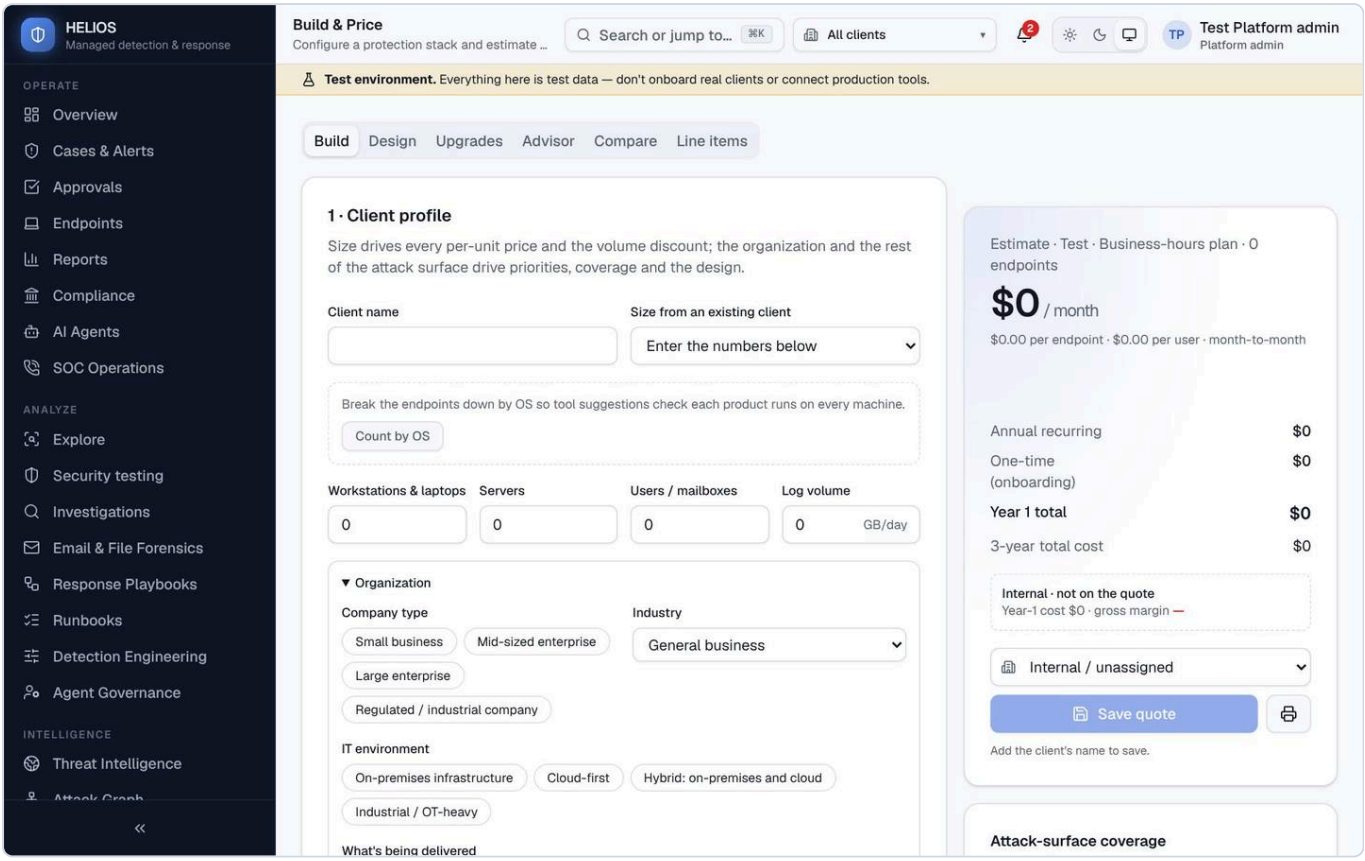


Figure 73. Build & Price, as seen by a platform admin (test environment).

How to use it

- The tabs: Build, Design, Upgrades, Advisor, Compare and Line items. The estimate on the right (monthly, annual recurring, one-time, year 1 and three-year totals) updates as you go.
- Build: describe the client (size from an existing client, or the numbers: workstations, servers, users, log volume, endpoints counted by OS; company type, industry, IT environment) and choose a tool per protection layer. Advisor suggests one per layer for a goal (best value, lowest cost, strongest protection) with its reasons.
- The upgrade review compares each tool with every alternative: what it adds or removes, whether it runs on the client's machines, and the price difference. The solution design lays out priorities, a phased rollout and coverage gaps.
- See internal cost and margin next to the price.
- Save quote (it is versioned), and print or download it as a PDF; a quote never shows cost or margin.
- Edit the price book, plans and exchange rates (Pricing setup).

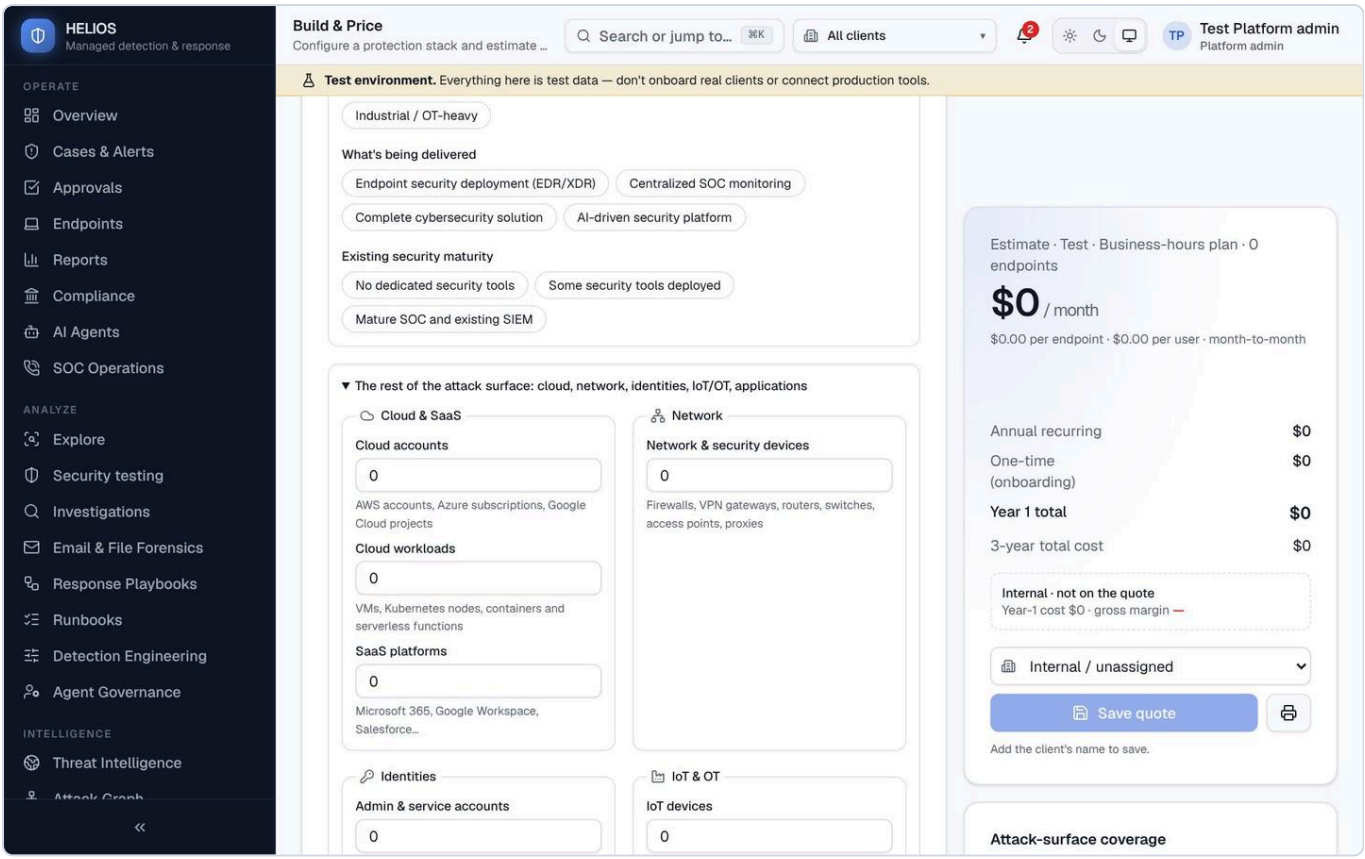


Figure 74. Build & Price, further down the page.

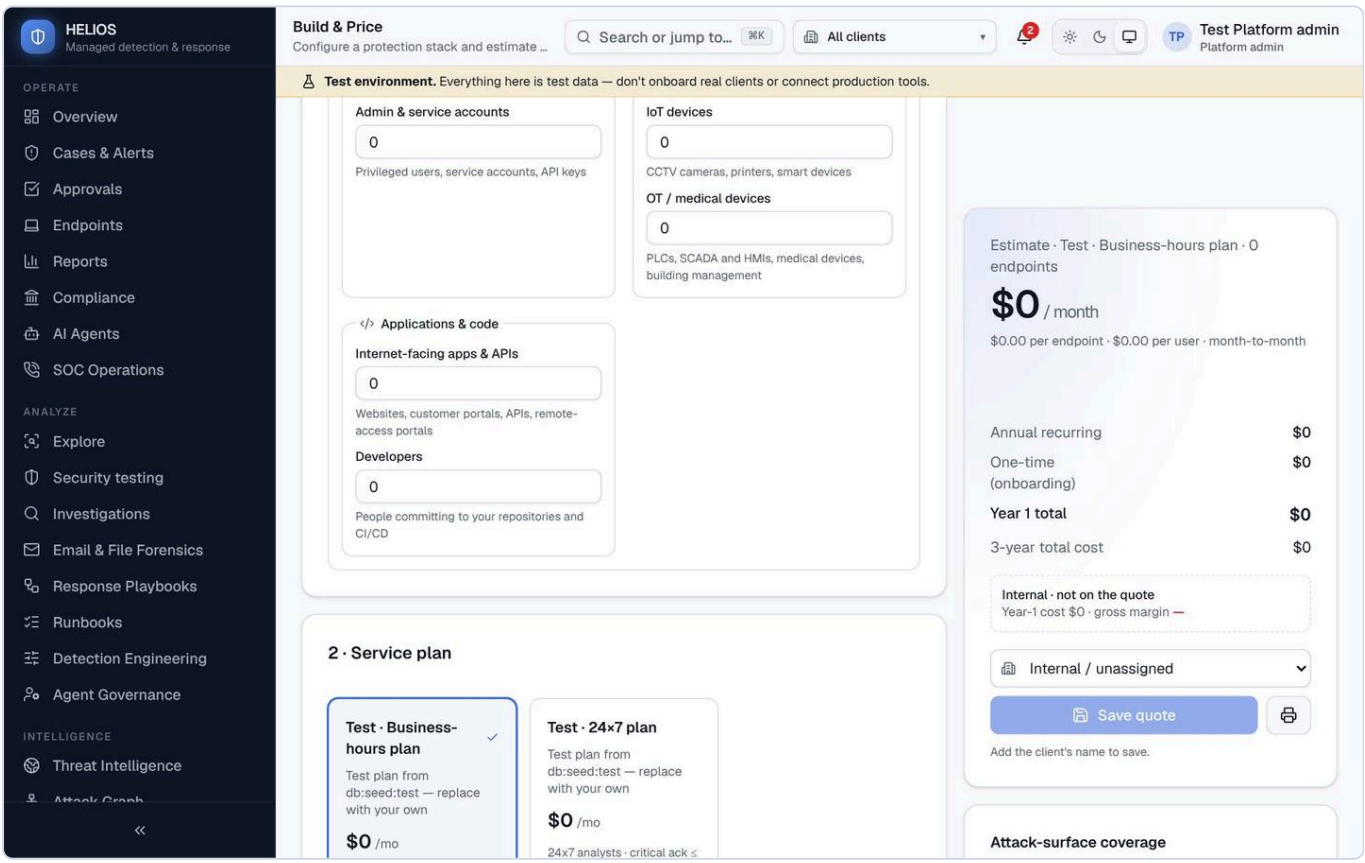


Figure 75. Build & Price, further down the page.

Contracts & Billing

Pilots with measured success criteria, contracts and GST invoices.

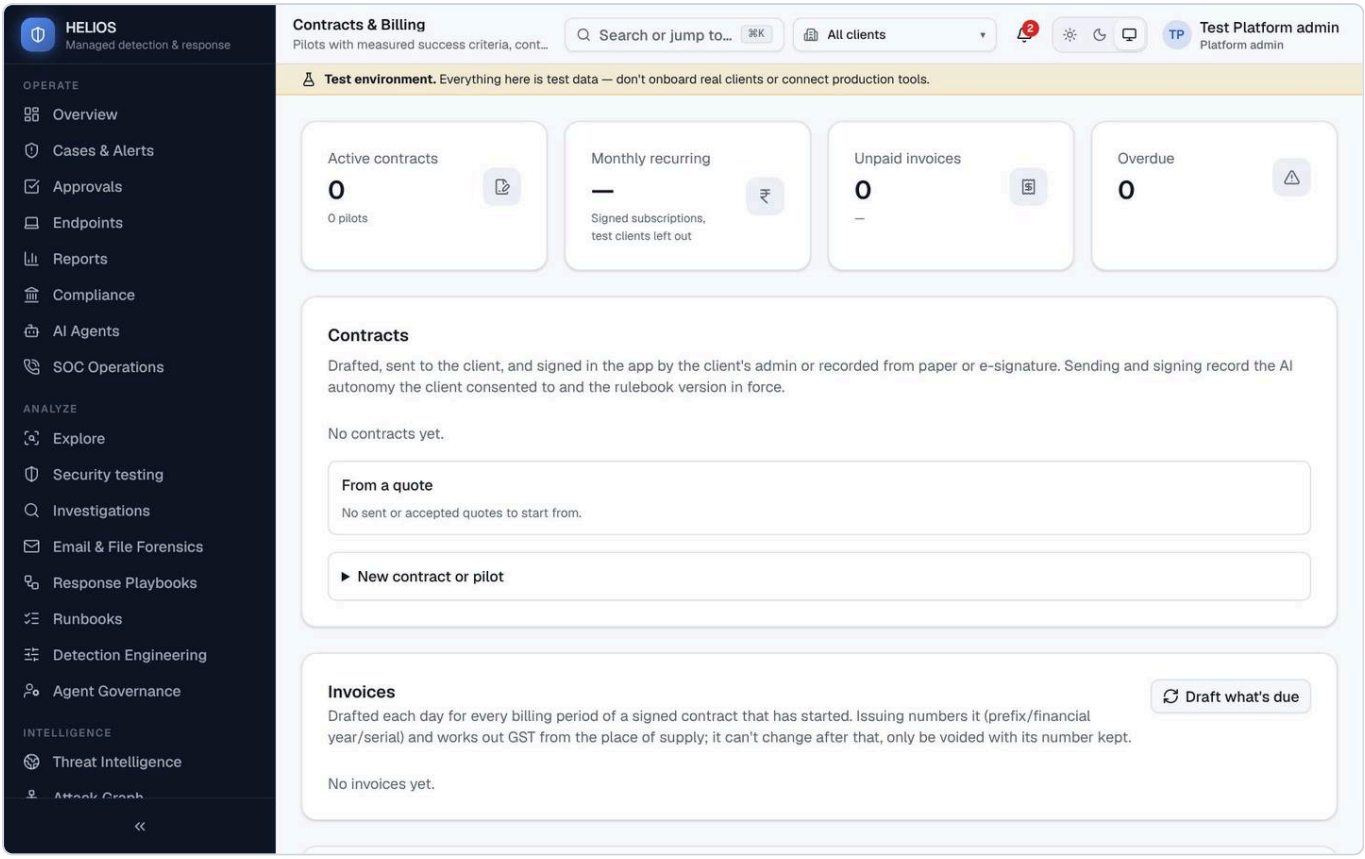


Figure 76. Contracts & Billing, as seen by a platform admin (test environment).

How to use it

- Pilots carry success criteria measured from HELIOS's own records (time to acknowledge and contain, SLA, AI agreement, wrong containments, sources live). A week before the end you get their status; convert the pilot to a subscription or record that it isn't continuing.
- Contracts: start From a quote (sent or accepted) or with New contract or pilot; send it with your terms of service; the client's admin signs it in the app (or you record a signature on paper or an e-signature). Sending and signing record the AI autonomy the client consented to and the rulebook version in force; signing is refused if either changed since sending.
- Invoices are drafted each day for every billing period of a signed contract that has started (Draft what's due does it now). Issuing numbers the invoice (prefix, financial year, serial) and works out GST from the place of supply; it can't change after that, only be voided with its number kept. Mark it paid when the money arrives.

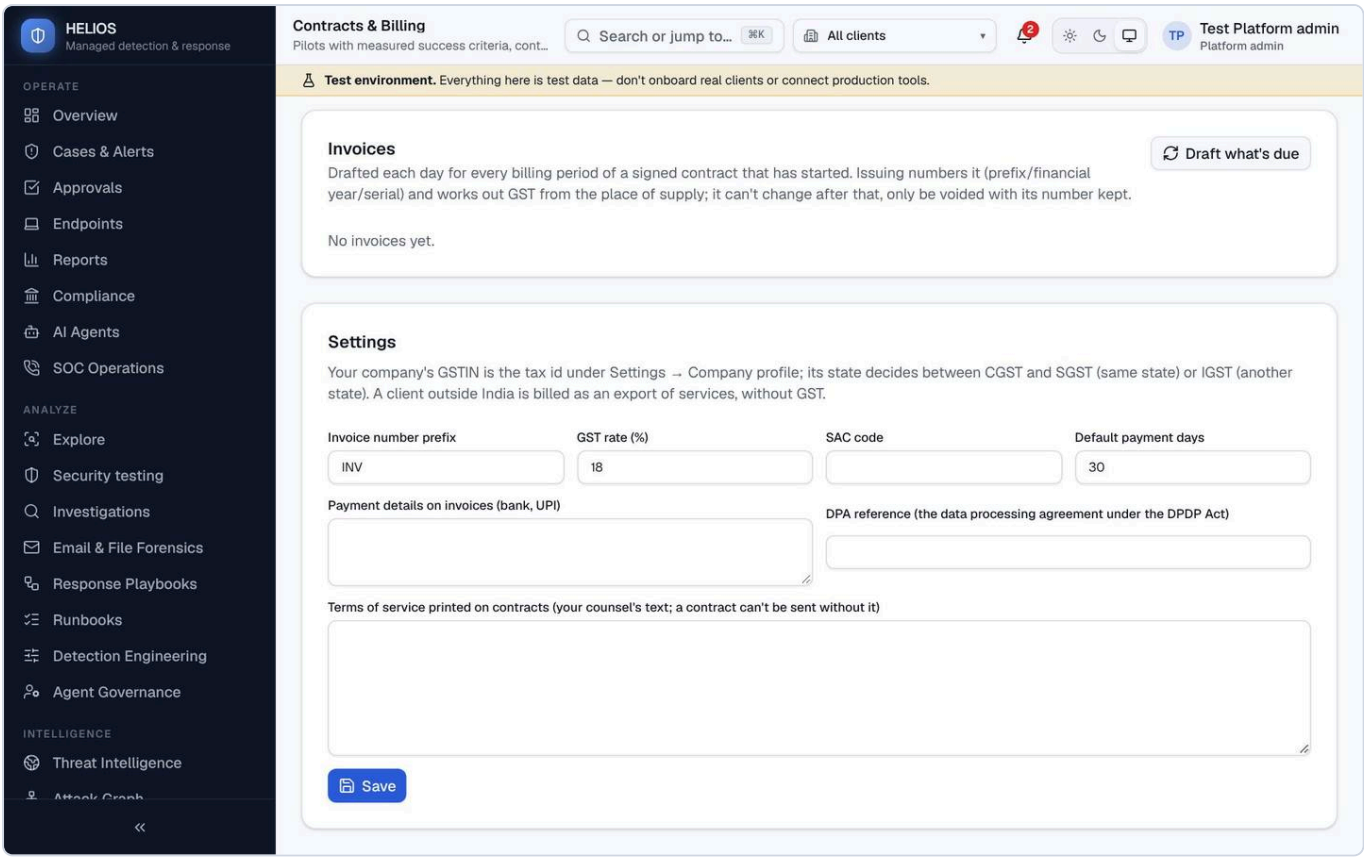


Figure 77. Contracts & Billing, further down the page.

A client's rulebook

The rulebook page for one client (here the test client Clinic).

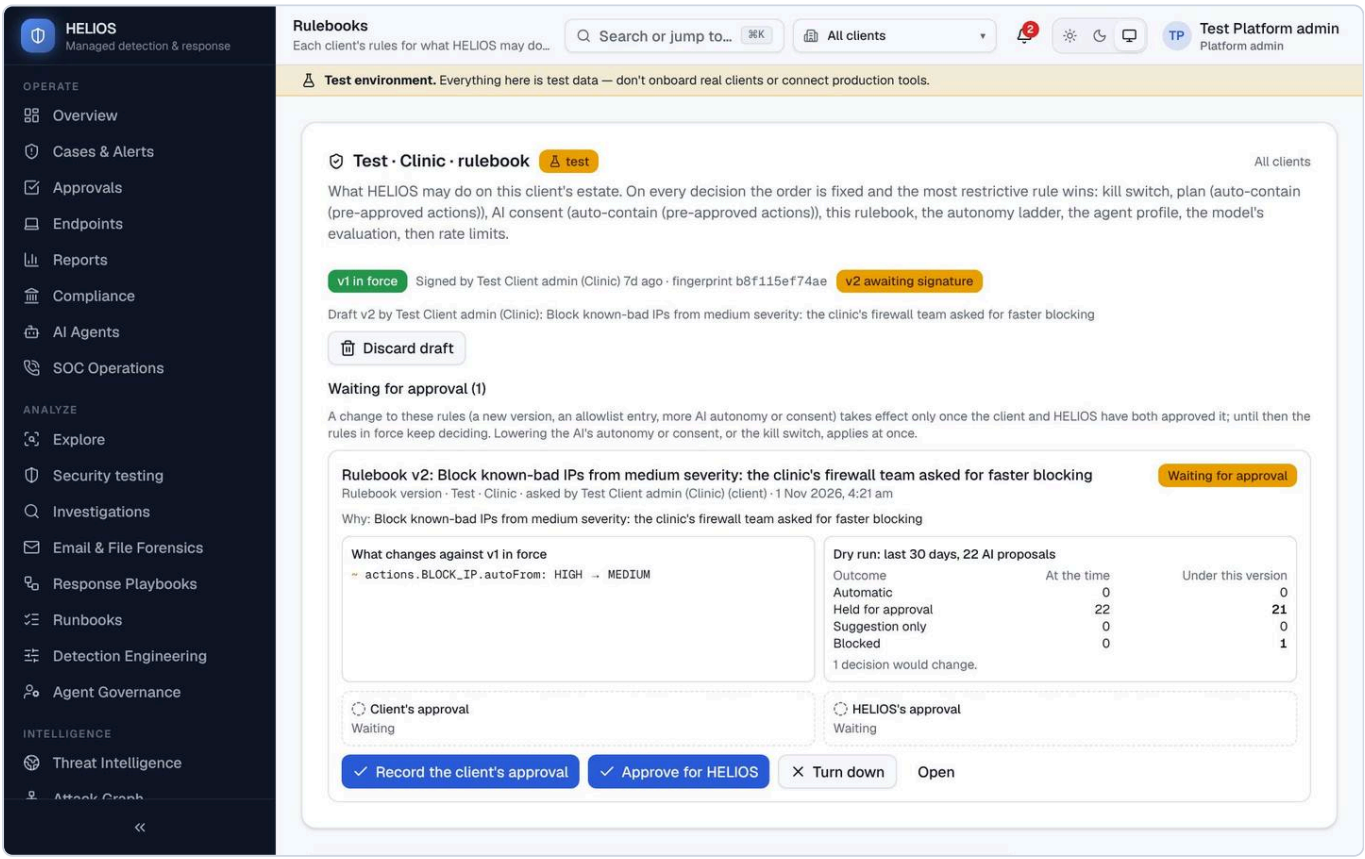


Figure 78. A client's rulebook, as seen by a platform admin (test environment).

How to use it

- In force shows the assets (never isolated, crown jewels), the people (never disabled, VIPs), each action's mode (automatic from a severity, needs approval, two people), the schedule, the approver chain, and data handling and obligations. The other tabs are as described above.

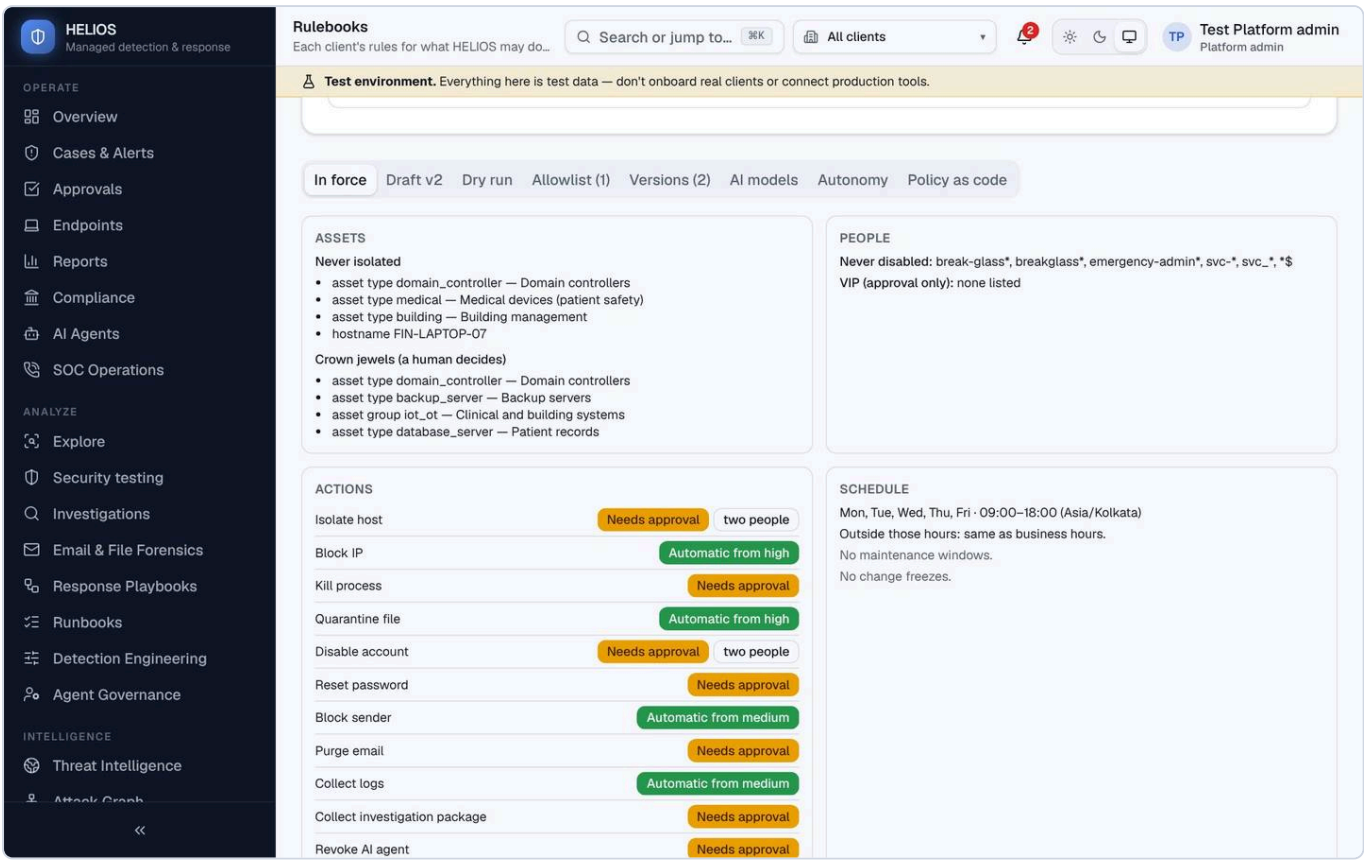


Figure 79. A client's rulebook, further down the page.

A client

One client's page: plan and SLA, contacts, consent, onboarding status, tools, users and data keys.

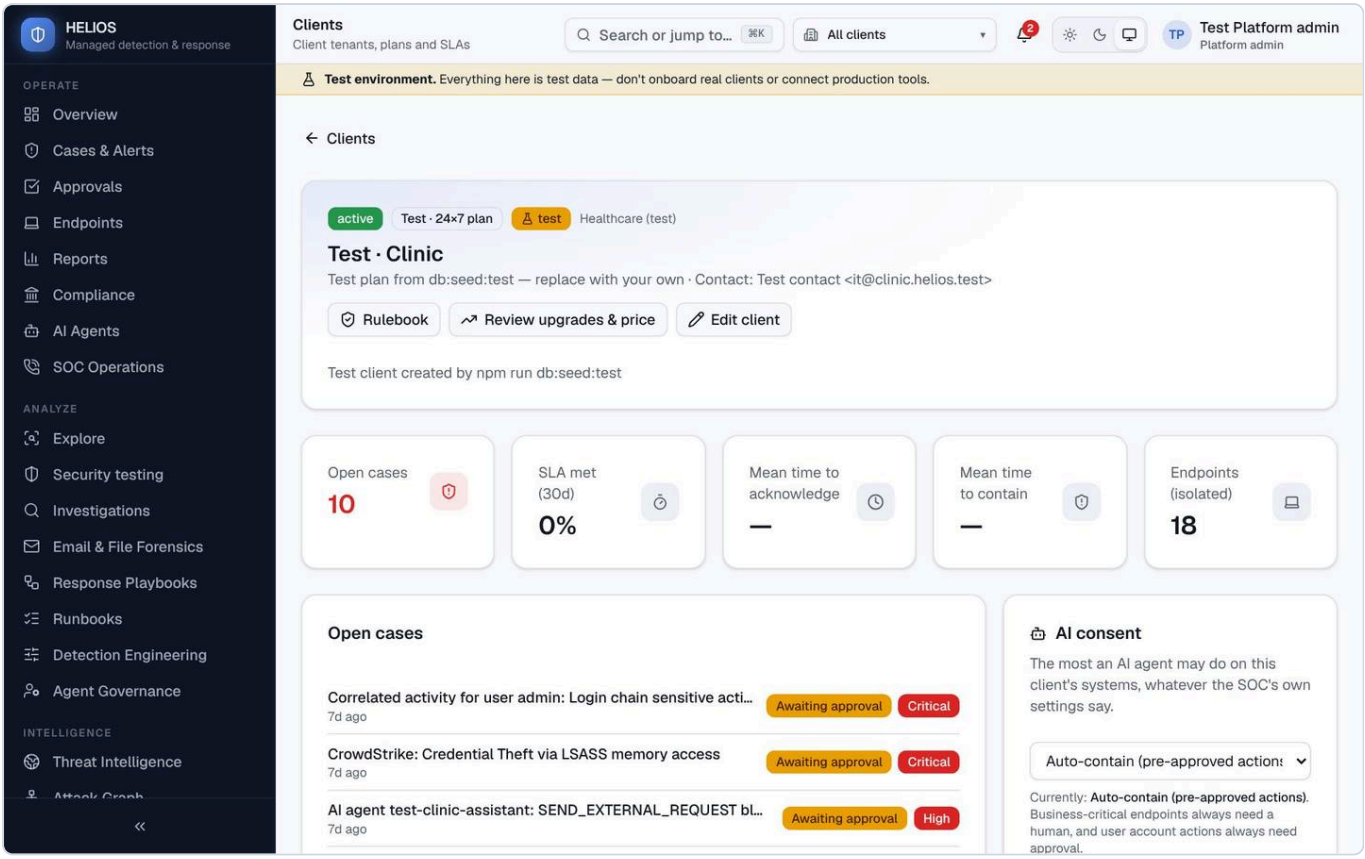


Figure 80. A client, as seen by a platform admin (test environment).

How to use it

- The page shows the client's open cases, SLA met, mean time to acknowledge and contain, endpoints, and buttons for its Rulebook, Review upgrades & price, and Edit client.
- AI consent: the most an AI agent may do on the client's systems, whatever the SOC's own settings say. Set it as the client agreed.
- Encryption key: Rotate key re-seals the client's credentials under a new data key; Destroy keys (for an offboarded, suspended client) makes its stored credentials unreadable for good, and asks you to type the client's name. Only platform admins can.

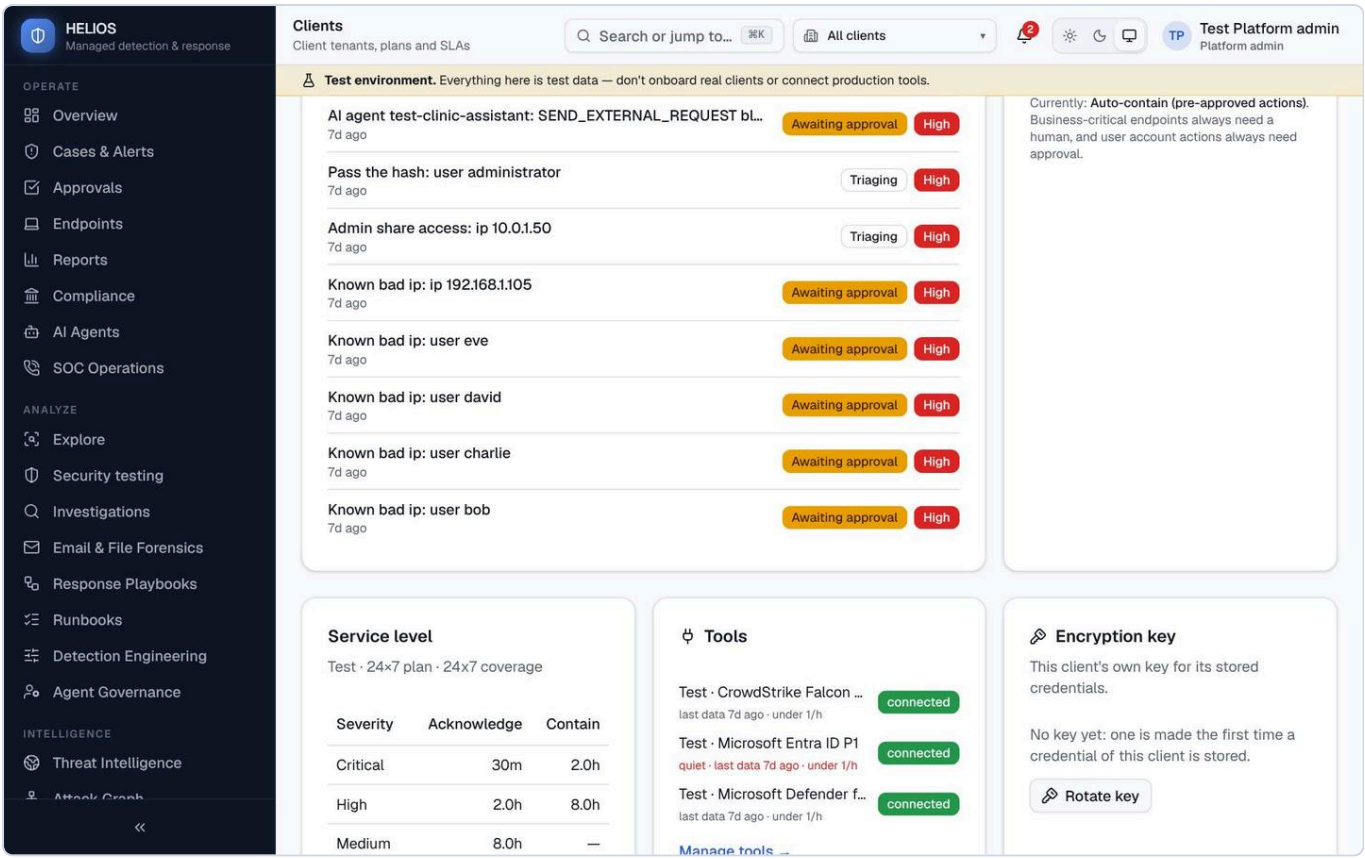


Figure 81. A client, further down the page.

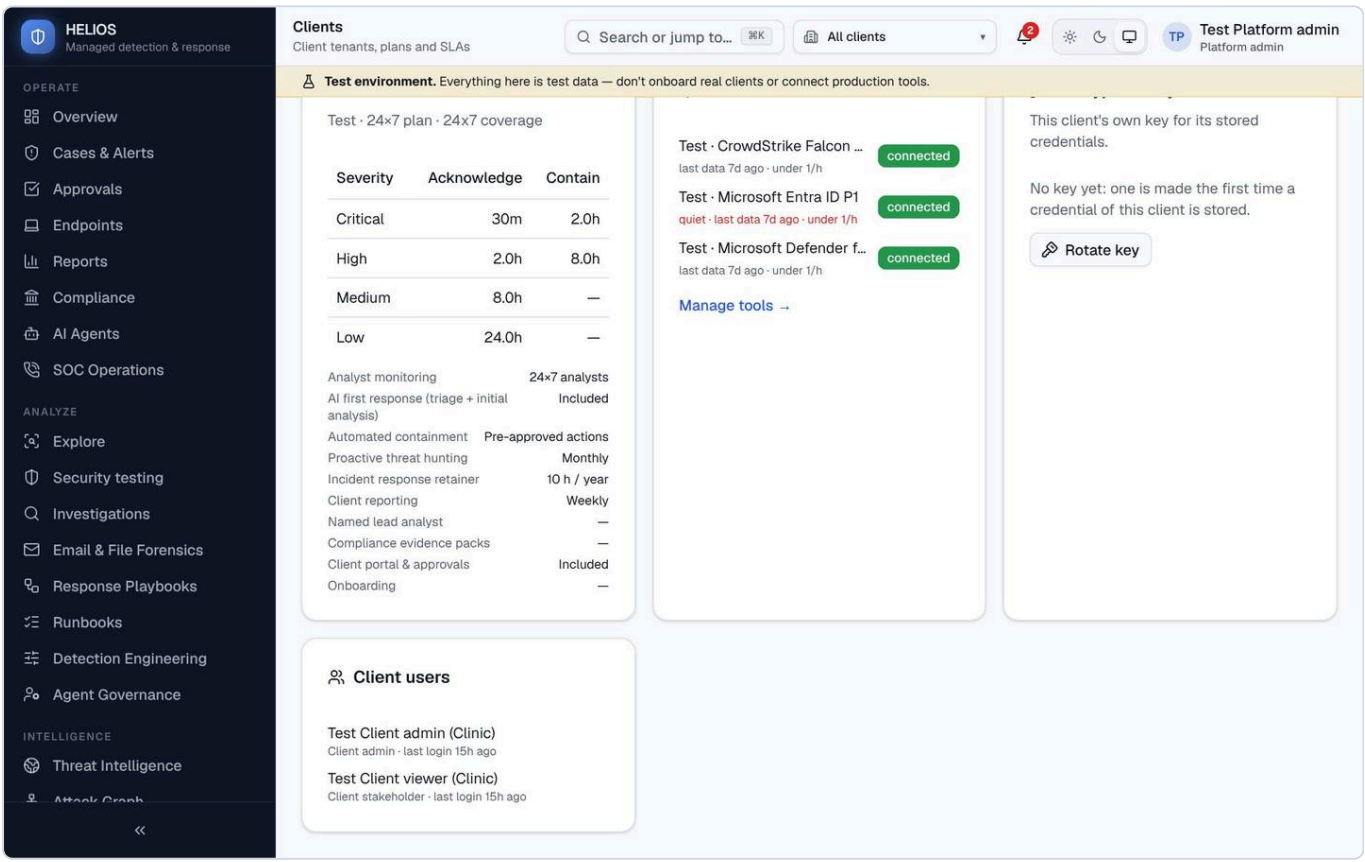


Figure 82. A client, further down the page.

A saved quote

One quote: its configuration, the price breakdown and its versions.

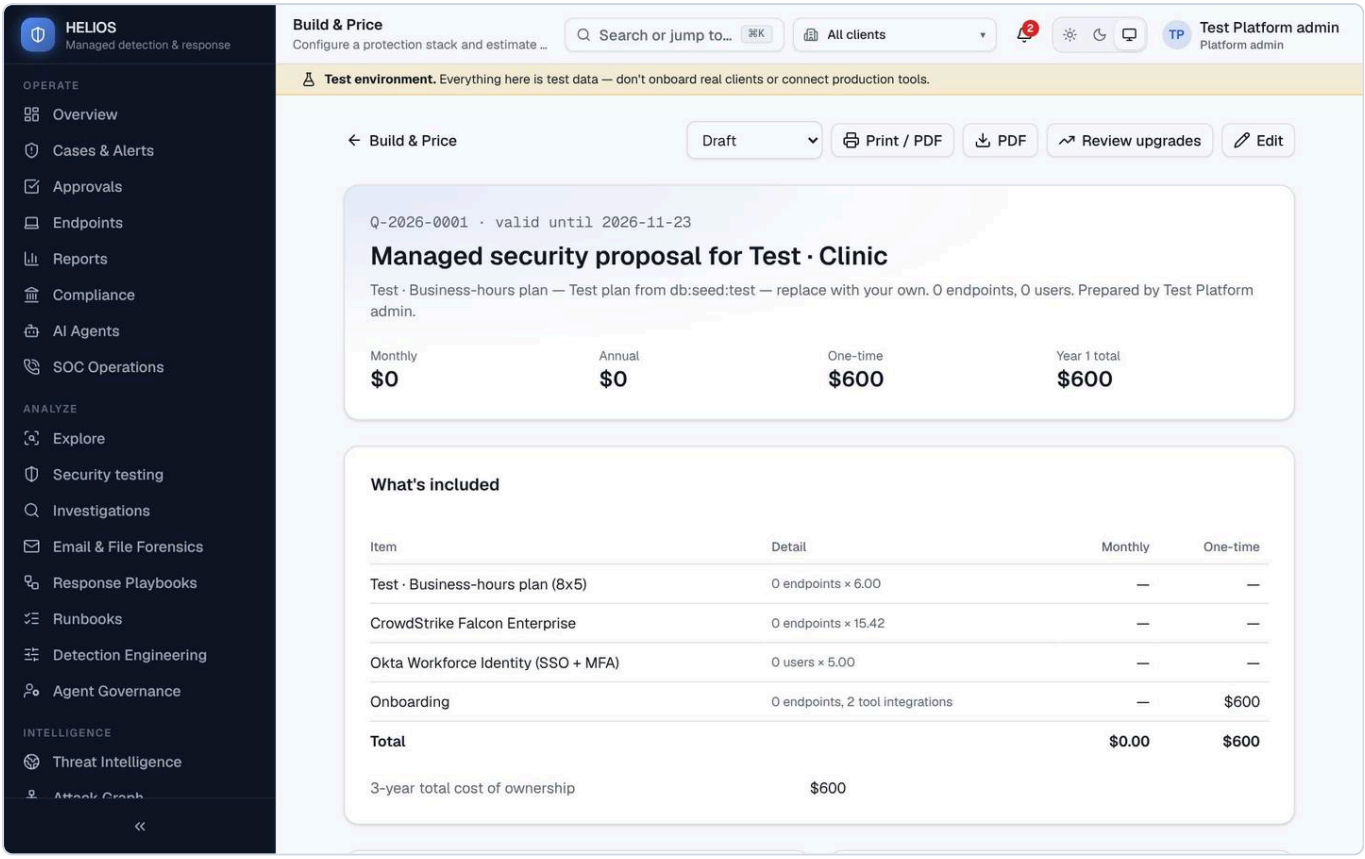


Figure 83. A saved quote, as seen by a platform admin (test environment).

How to use it

- Reopen it in the configurator, download the PDF, or start a contract draft from it on Contracts & Billing.

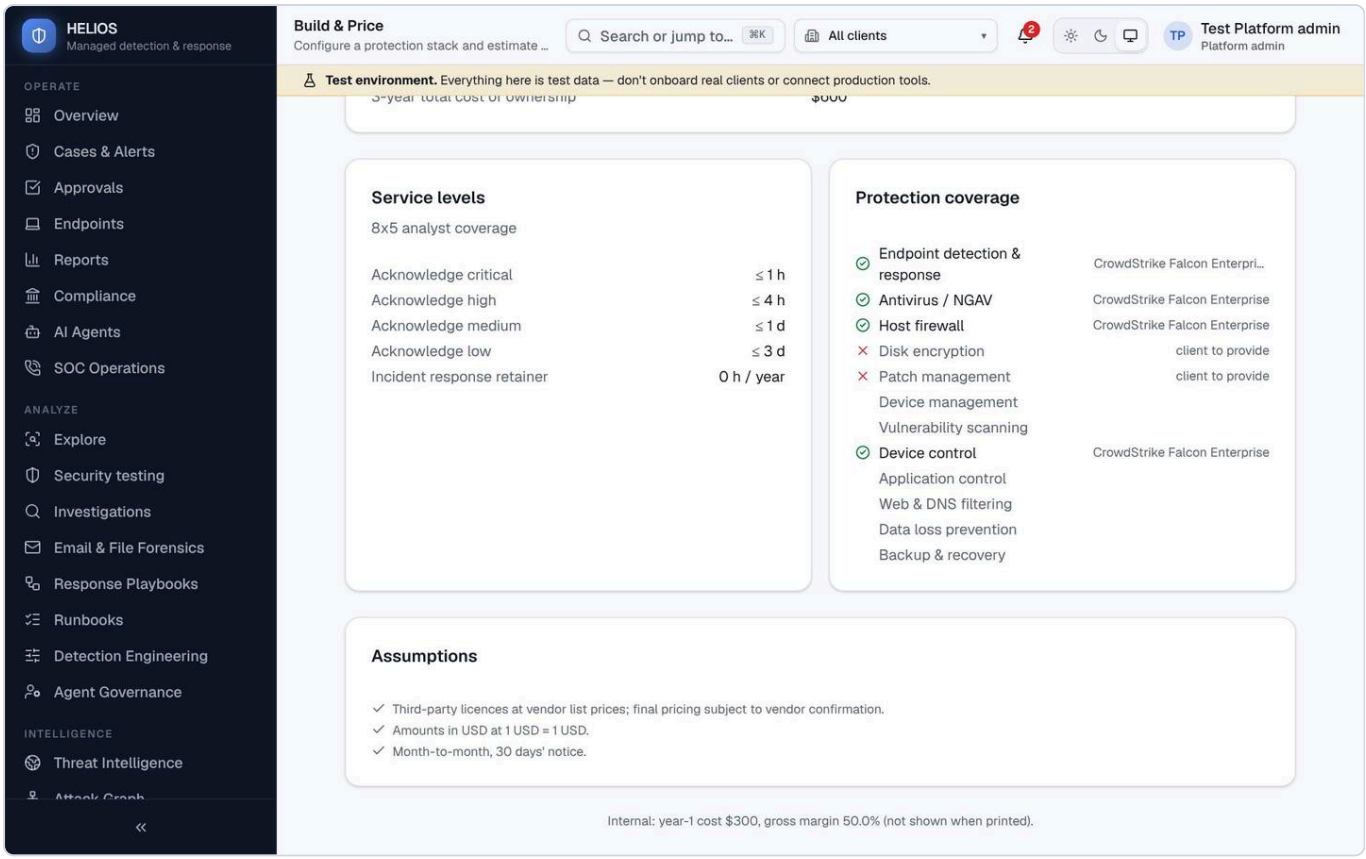


Figure 84. A saved quote, further down the page.

Daily digests

Once a day, what HELIOS did on its own, what was waiting, and the deadlines coming up.

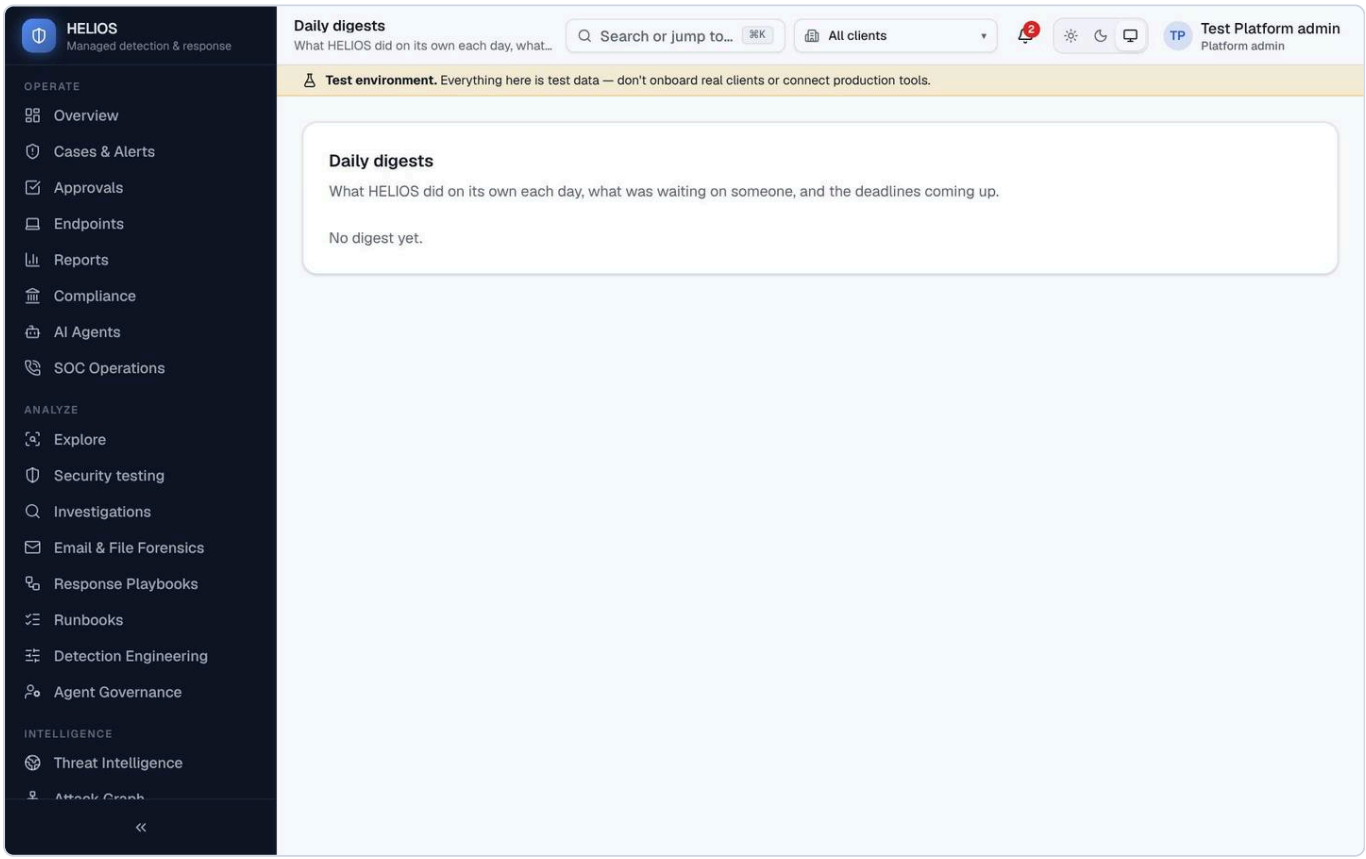


Figure 85. Daily digests, as seen by a platform admin (test environment).

How to use it

- Each digest lists actions the AI took without asking (with its autonomy level, rollbacks and QA calls), cases opened and closed, containment waiting for approval, regulatory deadlines within a day, quiet sources and the week's QA results.
- Choose under My alerts whether it also reaches you by email.

A security test

One security test: its scope and the proof of control, both authorisations, the findings and how to fix them, the scanner reports added, and what HELIOS probed.

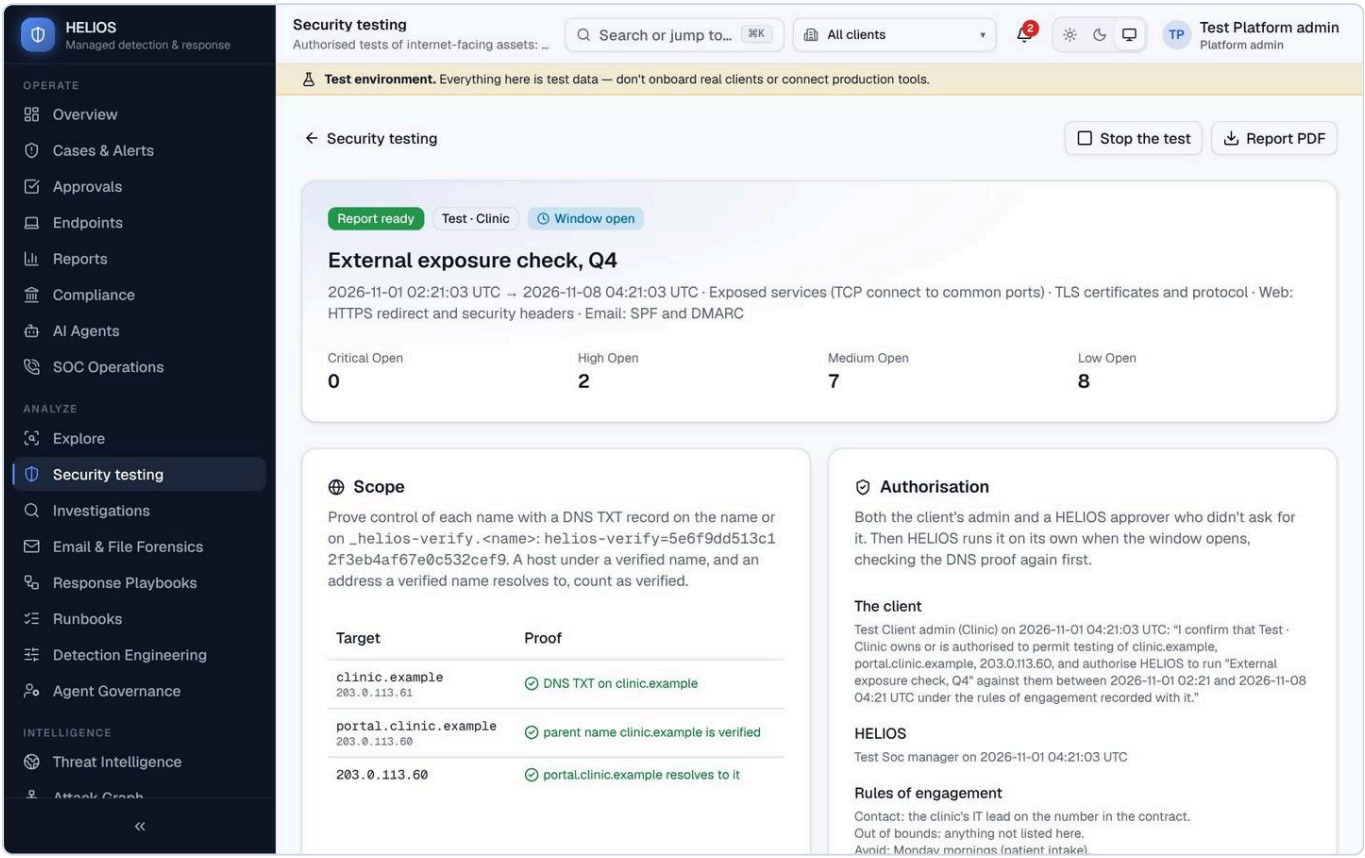


Figure 86. A security test, as seen by a platform admin (test environment).

How to use it

- Scope: the client proves it controls each name with a DNS TXT record (shown on the page) on the name or on _helios-verify.<name>. A host under a verified name, and an address a verified name resolves to, count as verified. Check DNS now re-checks.
- Authorisation: the client's admin authorises, confirming they may permit testing of exactly these assets in this window; a HELIOS approver (SOC manager or admin) who didn't ask for it authorises for HELIOS. Then HELIOS runs it on its own when the window opens, checking the DNS proof again first.
- HELIOS looks and doesn't exploit anything: a TCP connect to well-known ports and the banner a service volunteers, the TLS certificate and protocol, the HTTPS redirect and security headers, SPF and DMARC. Every probe is listed at the bottom.
- Findings are listed worst first, with the evidence and how to fix each. Mark one fixed, or accepted as a risk with the reason (client admins and SOC managers). Run again now re-tests inside the window: what it no longer finds is marked fixed.

- Scanner results: add a report from Tenable Nessus, Greenbone/OpenVAS, OWASP ZAP or Nuclei run under the same authorisation; findings on hosts outside the verified scope are left out.
- Reject (before it runs) or Stop the test (at any point): either side can. Report PDF downloads the report for the client and its auditors.

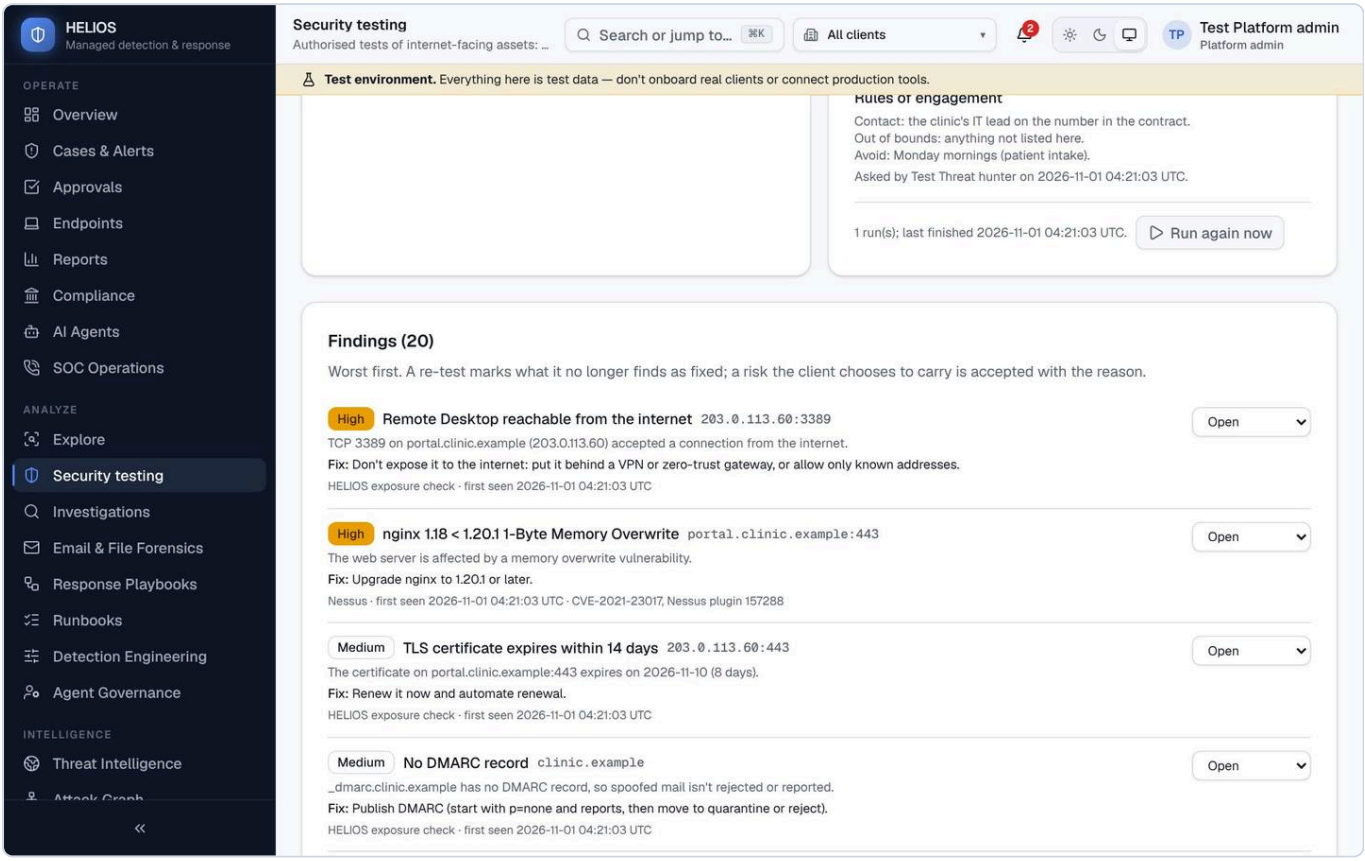


Figure 87. A security test, further down the page.

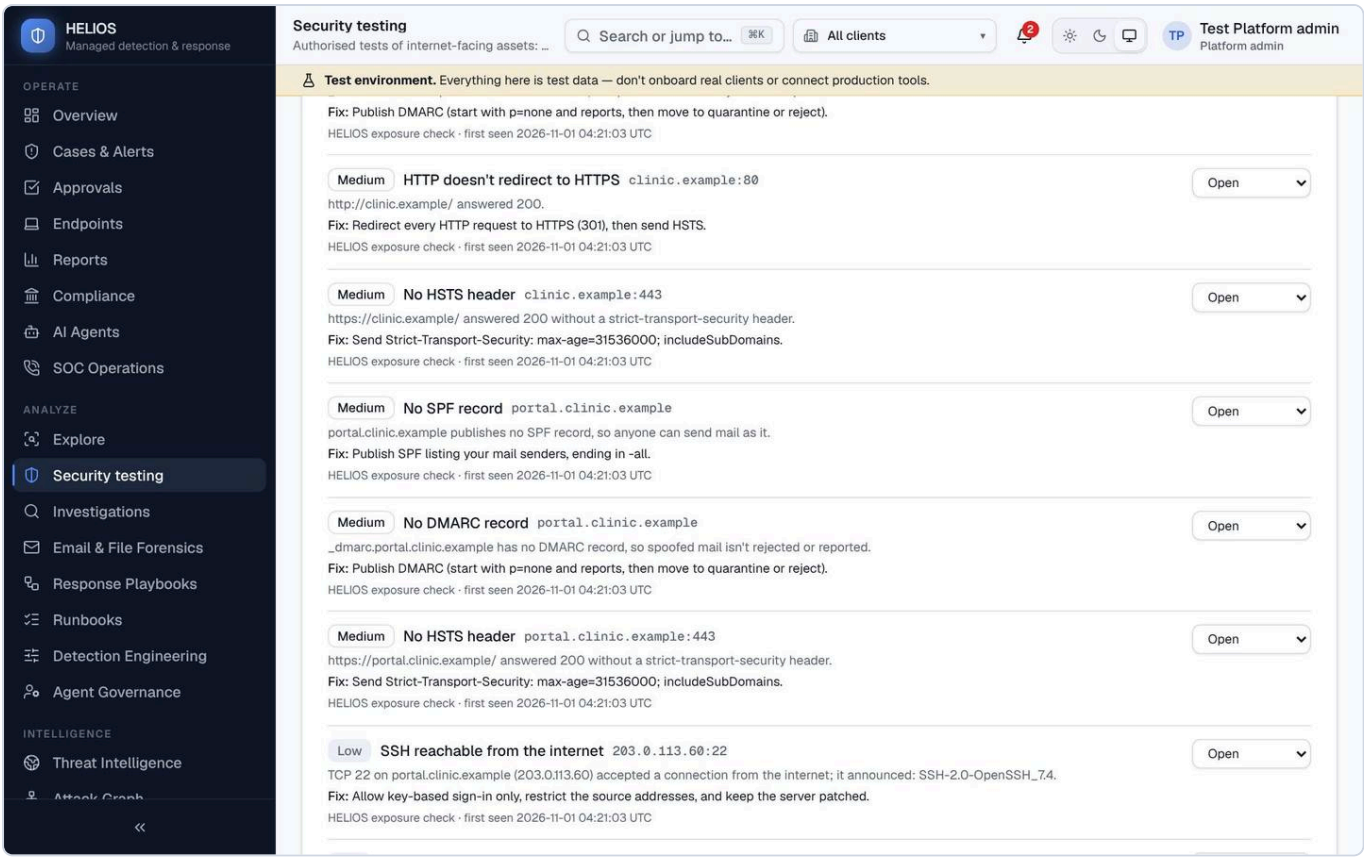


Figure 88. A security test, further down the page.

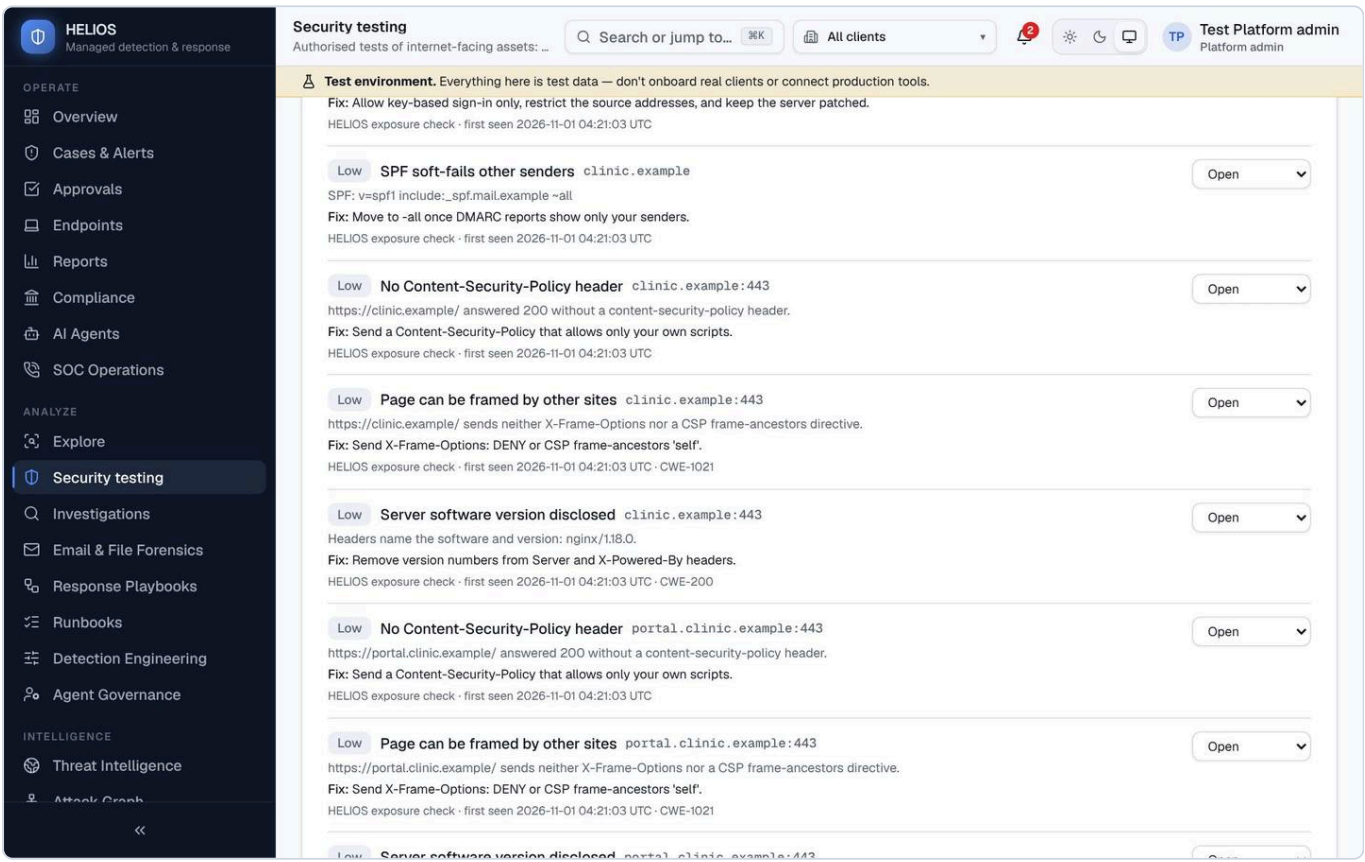


Figure 89. A security test, further down the page.

My alerts

Which notifications also reach you outside HELIOS, by email, SMS or WhatsApp.

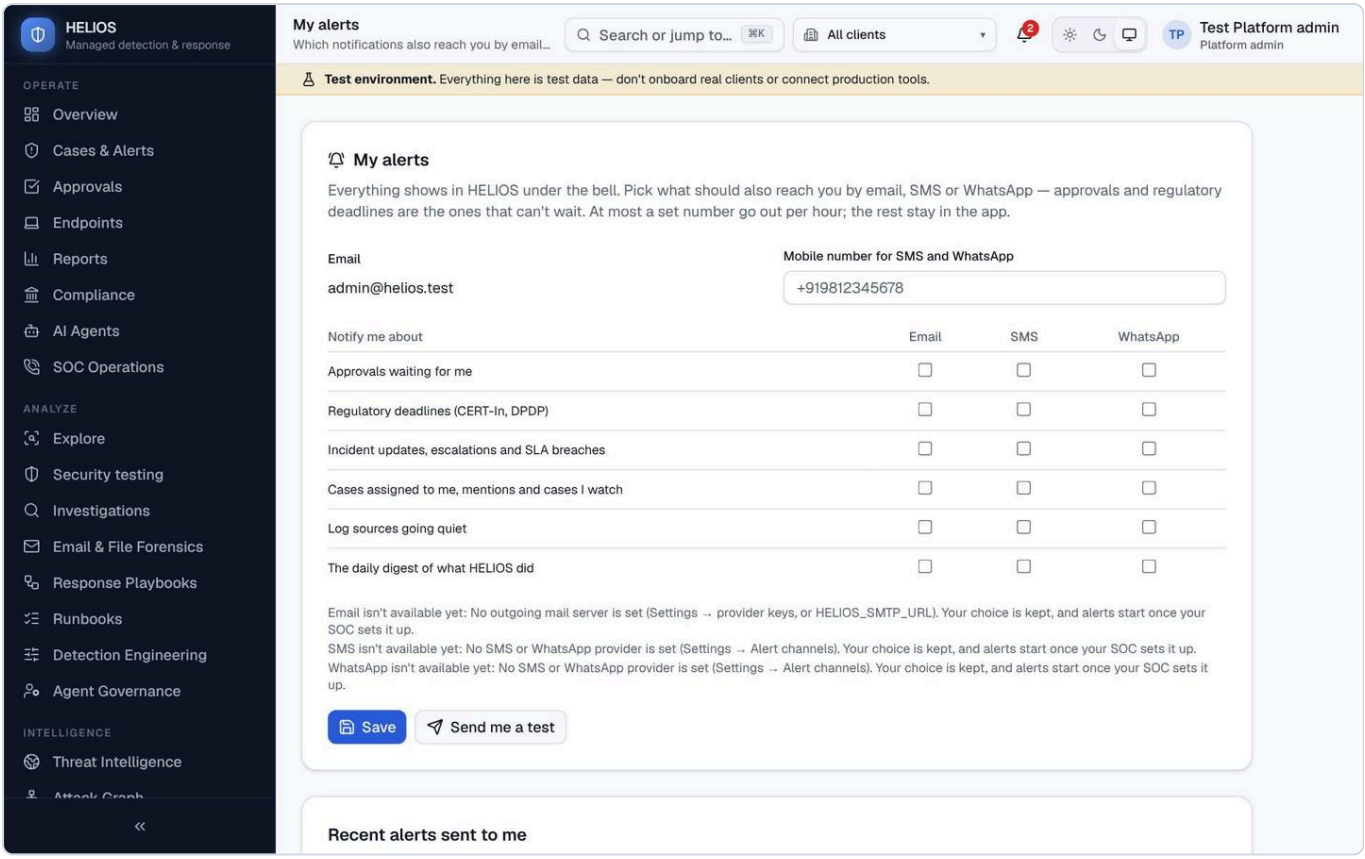


Figure 90. My alerts, as seen by a platform admin (test environment).

How to use it

- Every notification lands in the app (the bell at the top). Here choose, per kind of notification, the channels it also comes by, and give your phone number in international format (+91…).
- Deliveries that couldn't go (a channel not set up, no number, over the hourly cap) are listed with the reason.

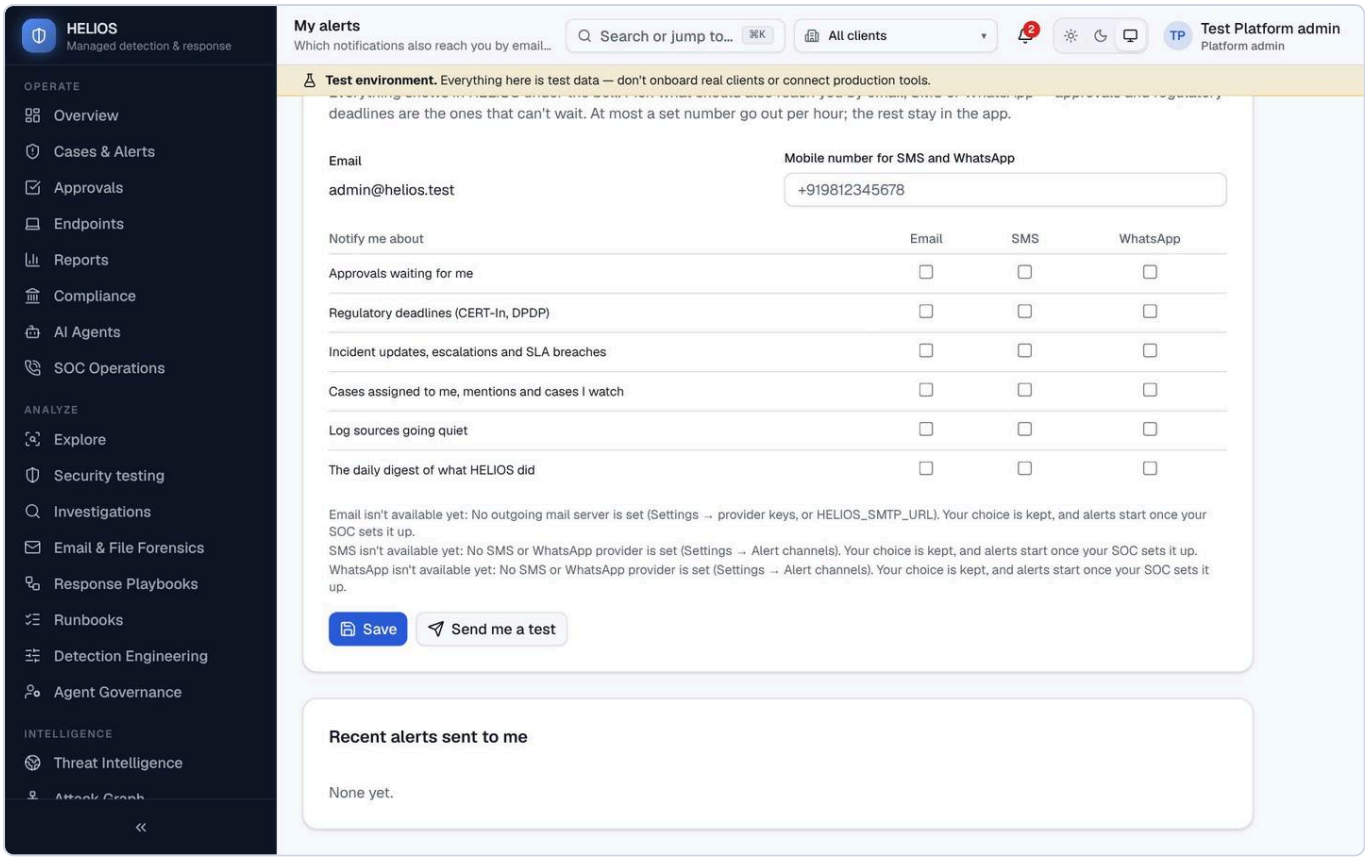


Figure 91. My alerts, further down the page.

5. Procedures

Step by step, for the tasks your role does most.

5.1 First-run set-up

1. Install HELIOS: on a Mac, ./helios.sh installs Docker if needed, generates the secrets, builds, starts and prints the first admin's sign-in; in the cloud, follow the deployment guide (docs/deployment.md).
2. Sign in as the first admin, change the password and set up two-step sign-in. Open Overview and work through Get HELIOS ready from top to bottom.
3. Define your service plans (SLA targets per severity, the most autonomy each plan allows) and fill in the price book under Build & Price → Pricing setup.
4. Price the quote-only tools you resell, or switch off the ones you don't sell.
5. Connect an AI provider on AI Agents (or leave agents on the built-in engine for now).
6. Invite your analysts on Team & Roles.
7. Add your first client on Clients, then onboard its security tools on Tools & Integrations.
8. Turn on the threat feeds you want on Threat Intelligence.

5.2 Invite a person and give them the right access

1. Open Team & Roles and choose Invite user.
2. Enter name and email, and pick the role that matches the job (see the role matrix on the page). Give the least access that does the job.
3. For a client user, choose their organisation; they will only ever see it.
4. For an analyst who works some queues only, tick their security domains (endpoint, identity, email, network, cloud, AI agents).
5. Save. HELIOS shows a temporary password once: pass it to the person by a secure channel. They sign in with it, and set up two-step sign-in if your policy requires it.

5.3 Someone leaves or changes job

1. Open Team & Roles, find the person, and deactivate the account (leaver) or change the role (mover). It applies on their next request; there is no session to wait out.
2. Revoke their personal API tokens if they had any (they stop working with the account in any case).
3. Check the on-call rota on SOC Operations and the approver chains in clients' rulebooks for their name.

5.4 Require two-step sign-in

1. Open Settings → Sign-in and encryption.
2. Choose whether two-step sign-in is required for staff or for everyone. People it applies to are taken to set it up (an authenticator app, plus ten recovery codes) before anything else.
3. Optionally turn on single sign-on only for staff once your identity provider is connected.
4. For someone who lost their phone, open Team & Roles → their account → Reset two-step.

5.5 Configure AI safely

1. On AI Agents → Model providers, add a provider and press Test. Keys are encrypted and never shown again.
2. Run an evaluation for each agent you will use (Evaluations). A model a client picks only suggests until it passes.
3. Under Agent types, set each agent's provider, reasoning effort and autonomy. The containment agent starts on auto-contain for the five initial actions (isolate a host, block an IP, quarantine a file, block a sender, collect logs) above 85% confidence; the client's approved rulebook still decides each one.
4. Under Safety controls, set the minimum severity for AI first response and the cap on unattended containment per client per hour.
5. In an emergency, Engage kill switch stops all AI action at once, for every client.

5.6 Offboard a client

1. Make sure contracts and invoices are settled (Contracts & Billing) and no legal hold is needed (Data Retention).
2. On Clients, set the client to suspended.
3. Once the retention period your contract promises is honoured or the data is returned, open the client's page → Encryption key → Destroy keys and type the client's name to confirm. Their stored credentials become unreadable, backups included.

5.7 Give an auditor or regulator the audit trail

1. Open Audit Log and press Verify chain integrity: it should report the chain intact.
2. Filter to the area asked for, and Export JSON.
3. Record in the case or a comment who received the export and why.

6. Guidelines

Do

- Keep a separate day-to-day account for ordinary SOC work; use the admin account only for administration.
- Require two-step sign-in for every staff account.
- Change roles rather than sharing accounts; every action is attributed to a person in the audit chain.
- Test AI models on Evaluations before letting them act. New clients contain on their own from day one within their approved rulebook; start one in shadow mode if it asks for that.
- Keep the test environment for testing: it shows a banner on every page; never onboard real clients or production tools in it.

Don't

- Don't paste keys, passwords or client secrets into comments or case notes; use the encrypted provider-key and credential fields.
- Don't destroy a client's keys until offboarding is final: it can't be undone.
- Don't edit the database directly: changes outside HELIOS bypass the audit chain and the row-level isolation.

For everyone

- Never share your account or password; every action is recorded under your name.
- Keep client information inside HELIOS and to the people who need it.
- Treat text inside evidence (command lines, emails, file names) as possibly written by an attacker: it may address 'the AI' or 'the analyst' to mislead. HELIOS flags such text; don't follow it.
- Every action in HELIOS is reversible or recorded, and audited; write down why you did it.

> The client's approved rulebook binds everyone: the AI, playbooks and people. If HELIOS refuses an action, the reason names the rule. Don't work around it in a vendor's console; ask the client to change the rule.

7. Reference

Terms you will see

Case statuses

Open	New; nobody has acknowledged it yet.
Triaging	Someone is looking at it.
Investigating	Under analysis.
Awaiting approval	Containment is waiting on someone's yes.
Contained	The threat is stopped; clean-up may continue.
Resolved	Closed with a disposition.
False positive	Closed: it wasn't a threat.

Severities

Critical, High, Medium, Low, Info	Set by the detections and adjustable by analysts. The client's plan sets the acknowledgement SLA for each.

Dispositions (when closing)

True positive	A real threat.
Benign true positive	Real activity, but expected or authorised (an admin's tool, a pen test).
False positive	Not what the detection thought.

Tiers

T1	Triage.
T2	Analysis.
T3	Incident response.

Response action statuses

Suggested	A recommendation only; autonomy doesn't allow acting.
Pending approval	Waiting on an approver.
In progress	The vendor was asked; waiting for it to confirm. Doesn't count as containment yet.
Executed	Done and confirmed by the vendor (or by an analyst for a manual action).
Manual required	Approved, but no connector can do it: someone does it in the tool's console and records it.
Blocked	A guardrail or the client's rulebook refused it; the reason is shown.
Rejected	An approver said no.
Rolled back	Undone.

Failed	The vendor or HELIOS couldn't do it.
Shadow	Shadow mode: what the AI would have done, recorded silently and shown once the case closes.

AI autonomy (per agent)

Off	The agent doesn't run.
Suggest only	It proposes; people decide.
Contain with approval	Its containment waits for approval.
Auto-contain	It may act on pre-approved actions above its confidence bar.

Autonomy ladder (per client and action type)

L0 observe	Shadow mode, if the client chose to start there.
L1 suggest	Suggestions only.
L2 approve	Acts once approved.
L3 act and notify	Acts, and tells people at once. New clients start here for initial containment (isolate a host, block an IP, quarantine a file, block a sender, collect logs) and at L2 for everything else.
L4 act	Acts, and reports in the daily digest.

Final analysis

Ready for review	HELIOS has written it; the SOC reviews it and the client acknowledges it.
Reviewed	A person in the SOC has read and accepted it, with an optional note.
Lift proposal	A release of containment HELIOS proposes with its conditions; it runs only once someone approves it.

Rule changes

Waiting for approval	Not in effect; the rules in force keep deciding.
Approved	Every approval it needed is in, and it took effect.
Rejected or withdrawn	Turned down by an approver, or withdrawn by the person who asked.
Stale	The rules in force changed while it waited; ask again.

Security tests

Waiting for authorisation	Asked for; needs every target verified and both authorisations.
Authorised	Runs on its own when the window opens.
Running	HELIOS is probing now.
Report ready	Finished; can run again inside the window.
Rejected or stopped	Turned down before it ran, or stopped by either side.

Evidence references

D, E, I, T, C	Detections, events, indicators, timeline entries and context in the case; Q for what the investigation agent's tools returned. Every AI finding cites them.

If something doesn't work

- **A page says you don't have access**: your role doesn't include it. Ask your administrator if your job needs it.
- **An action was refused**: the message names the rule (the client's rulebook, a guardrail, the plan or the kill switch).
- **An action is 'Manual required'**: no connector can do it. Do it in the tool's own console, then record what you did on the case.
- **You don't get alerts by SMS or email**: check My alerts for your channel choices, your phone number, and deliveries skipped with their reason.
- **You can't sign in**: after five failures the account locks for 15 minutes; your administrator can reset your password or two-step sign-in.