HELIOS USER GUIDE · LEADERSHIP

Platform admin

Procedures, usage, guidelines, and your roles and access

The person who owns the HELIOS platform for the MSSP: accounts and roles, platform settings, AI configuration, the price book and every client's set-up. There should be few of them, and each should also hold a day-to-day role account for ordinary work.

Platform	HELIOS managed detection and response
Guide	01 of 11 · Platform admin
Date	31 October 2026
Version	Written against HELIOS 809e500
Screenshots	Taken in the HELIOS test environment with test data, signed in as a platform admin. Your screens show your own clients and data.

Contents

1. Your role at a glance
2. Roles and access
3. Getting started
4. Your screens
5. Procedures
6. Guidelines
7. Reference

You have **50** of 50 permissions and **27** pages in your sidebar: Overview, Cases & Alerts, Approvals, Endpoints, Reports, Compliance, AI Agents, SOC Operations, Investigations, Email & File Forensics, Response Playbooks, Runbooks, Detection Engineering, Agent Governance, Threat Intelligence, Attack Graph, MITRE ATT&CK Coverage, Tools & Integrations, Rulebooks, Clients, Team & Roles, Data Retention, Audit Log, API Tokens, Settings, Build & Price, Contracts & Billing.

1. Your role at a glance

The person who owns the HELIOS platform for the MSSP: accounts and roles, platform settings, AI configuration, the price book and every client's set-up. There should be few of them, and each should also hold a day-to-day role account for ordinary work.

What you are responsible for

- Set up and maintain the platform: environment, company profile, provider keys, alert channels and sign-in security.
- Manage accounts: invite people, give them the right role and queues, deactivate leavers at once, reset two-step sign-in.
- Own the commercial reference data: service plans, the price book, exchange rates.
- Configure AI: providers, agent profiles, autonomy, the kill switch, and evaluations before a model acts on its own.
- Look after client data: retention periods, legal holds, data keys (rotation, destruction at offboarding).
- Keep the audit chain intact and exportable for auditors and regulators.

A typical day or shift

1. Check the Overview checklist (Get HELIOS ready) until every item is done.
2. Review new and changed accounts on Team & Roles; deactivate anyone who has left.
3. Look at Settings → Alert channels for deliveries that failed or were skipped.
4. Verify the audit chain (Audit Log → Verify chain integrity) now and then, and before any export.

Who to turn to

Security concerns about the platform itself go to your security lead; client-facing incidents go to the SOC manager on duty.

2. Roles and access

What you can see

You work across the clients the MSSP serves. The client switcher at the top focuses every page on one client, or All clients. You see every case queue regardless of security domains.

Your access comes from your role, and HELIOS checks it on the server for every page, button and API call; hiding a button is only a convenience. A change to your role applies on your next click. Everything you do is written to the hash-chained audit log under your name.

Your pages

Overview, Cases & Alerts, Approvals, Endpoints, Reports, Compliance, AI Agents, SOC Operations, Investigations, Email & File Forensics, Response Playbooks, Runbooks, Detection Engineering, Agent Governance, Threat Intelligence, Attack Graph, MITRE ATT&CK Coverage, Tools & Integrations, Rulebooks, Clients, Team & Roles, Data Retention, Audit Log, API Tokens, Settings, Build & Price, Contracts & Billing, plus My alerts and your two-step sign-in settings in the user menu.

Your permissions

What the permission lets you do	You
Cases	
See the overview dashboard (dashboard.view)	Yes
See cases in their queues (or their own tenant, for clients) (cases.view)	Yes
See every queue, regardless of assigned security domains (cases.view_all)	Yes
Acknowledge cases, change status and severity (cases.triage)	Yes
Assign cases to other analysts (cases.assign)	Yes
Escalate or hand off a case to the next tier (cases.escalate)	Yes
Comment on cases and @mention teammates (cases.comment)	Yes
Share cases with the client and external channels (cases.share)	Yes
Response	
Run containment playbooks (subject to approval) (containment.request)	Yes
Approve or reject pending containment (containment.approve)	Yes
Isolate and release endpoints directly (endpoints.contain)	Yes
Analysis	
Run log investigations (investigations.run)	Yes
Submit emails and files for forensic analysis (forensics.run)	Yes
Upload raw logs into the detection pipeline (telemetry.ingest)	Yes
Govern AI agents: identities, tokens, revocation (governance.manage)	Yes

Intelligence and detection	
Look up threat indicators (intel.view)	Yes
Edit the blocklist and indicator store (intel.manage)	Yes
See ATT&CK coverage (mitre.view)	Yes
Switch SIEM detectors on and off (detections.tune)	Yes
See response playbooks (playbooks.view)	Yes
Author and enable response playbooks (playbooks.manage)	Yes
AI agents	
See AI agent runs and configuration (ai.view)	Yes
Dispatch AI agents onto a case (ai.dispatch)	Yes
Rate AI agent output (ai.feedback)	Yes
Configure AI providers, agent profiles and autonomy (ai.configure)	Yes
Tools and endpoints	
See onboarded security tools (integrations.view)	Yes
Onboard and configure security tools (integrations.manage)	Yes
See the endpoint inventory and protection coverage (endpoints.view)	Yes
Clients and rulebooks	
See client tenants (tenants.view)	Yes
Create and edit client tenants (tenants.manage)	Yes
Read a client's rulebook: what HELIOS may do on its estate (rulebook.view)	Yes
Draft changes to a client's rulebook and its allowlist (rulebook.edit)	Yes
Put a client's rulebook in force (client admins sign; staff only on the client's recorded instruction) (rulebook.sign)	Yes
Compliance and reports	
See regulatory clocks and notices (CERT-In, DPDP) and the ISO 27001 evidence map (compliance.view)	Yes
Draft regulatory notices, record the client's approval, and file them (compliance.prepare)	Yes
Approve a regulatory notice before it is filed (the report is the client's; staff record a client's approval given elsewhere) (compliance.approve)	Yes
Read client-facing incident reports (reports.view)	Yes
SOC operations	
See the on-call rota, pages, runbooks, QA results and digests (operations.view)	Yes
Edit the on-call rota, runbooks and operations settings (operations.manage)	Yes
Review the weekly sample of what the AI did on its own (qa.review)	Yes

Commercial	
Use the Build & Price configurator (pricing.view)	Yes
See internal cost and margin (pricing.margin)	Yes
Edit the price book (pricing.manage)	Yes
See contracts, pilots and invoices (clients: their own) (billing.view)	Yes
Draft and record contracts, issue invoices and record payments (billing.manage)	Yes
Platform	
Invite users and change roles (users.manage)	Yes
Edit platform settings: company profile, provider keys, forensics thresholds (settings.manage)	Yes
Read the audit chain (audit.view)	Yes
Export the audit chain (audit.export)	Yes
Create personal API tokens for the REST API (acting with their own permissions) (api.tokens)	Yes

How your role fits with the others

Role	Group	What it is for
Platform admin	Leadership	Owns the platform: users, AI configuration, price book, everything else.
SOC manager	Leadership	Runs the SOC: every queue, approvals, AI autonomy, tool onboarding.
Tier-1 triage analyst	Security operations	Works the first-line queue: acknowledges, validates AI triage, escalates.
Tier-2 security analyst	Security operations	Primary analysis: investigations, forensics, containment requests.
Tier-3 incident responder	Security operations	Owns major incidents: approves containment, isolates endpoints, client comms.
Threat hunter	Security operations	Hunts across all tenants with investigations, forensics and intel.
Detection engineer	Engineering	Maintains playbooks, detections, tool integrations and ATT&CK coverage.
Auditor / compliance	Governance	Read-only oversight: cases, AI decisions and the audit chain.
Account manager	Commercial	Client onboarding, Build & Price quotes and margins.
Client admin	Client	Client-side: their own tenant's cases and tools; approves containment on their estate.
Client stakeholder	Client	Client-side read-only: dashboards, incident reports and endpoint coverage.

3. Getting started

Signing in

1. Open the HELIOS address you were given and sign in with your email and password. A new account gets a temporary password from its administrator.
2. If two-step sign-in is required, set it up the first time: scan the code with an authenticator app, enter a code, and keep the ten recovery codes somewhere safe. Each recovery code works once.
3. If your organisation uses single sign-on, use the sign-in button for your identity provider instead; it satisfies two-step sign-in.
4. Five wrong passwords or codes in a row lock the account for 15 minutes. An administrator can reset your password or your two-step sign-in if you lose your phone.

Finding your way

- **Sidebar**: your pages, grouped (Operate, Analyze, Intelligence, Platform, Commercial). The arrow at the bottom collapses it.
- **Search or jump to** (Ctrl K or Cmd K): find a case, page or client by name.
- **Client switcher**: All clients, or focus every page on one client.
- **The bell**: your notifications (mentions, hand-offs, approvals, pages). Choose which also reach you by email, SMS or WhatsApp under My alerts.
- **Light, dark or system theme**, and your **user menu**: My alerts, two-step sign-in, Sign out.
- A yellow **Test environment** banner means everything on that installation is test data.

Install it as an app

HELIOS can be installed from the browser (Install app in the address bar) as a desktop or mobile app with its own window and shortcuts.

4. Your screens

Each page you can open, what it is for, and how to use it with your access. Where a page is long, further screenshots show the rest.

Overview

Your starting page. It shows what needs you now, and it changes with your role: the SOC sees live posture, queues, SLAs and AI activity; account managers see the commercial picture; client users see their own organisation.

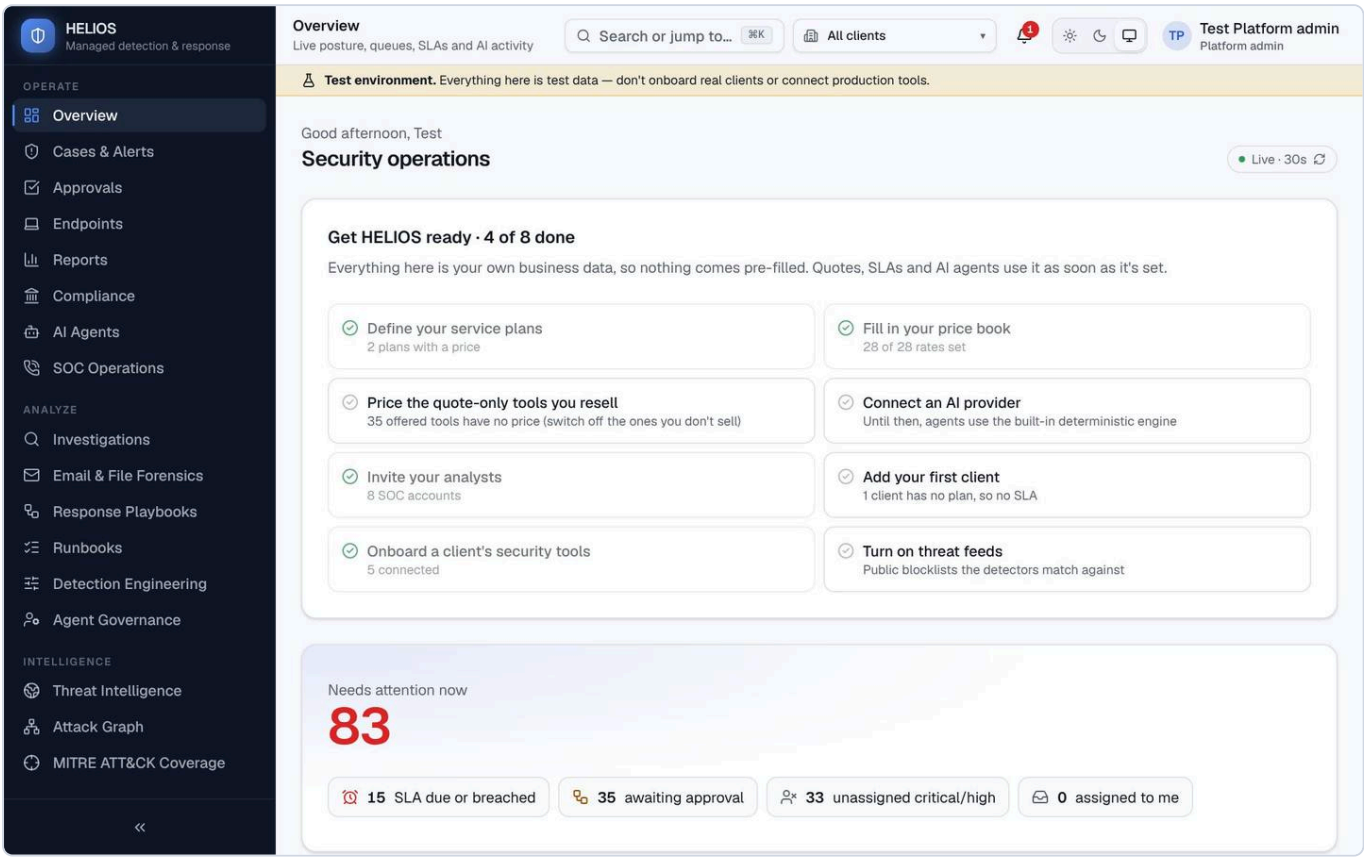


Figure 1. Overview, as seen by a platform admin (test environment).

How to use it

- Needs attention now adds up what is due: cases whose SLA is due or breached, containment awaiting approval, unassigned critical and high cases, and cases assigned to you. Click a figure to open the matching list.
- The Live badge means the figures refresh on their own every 30 seconds; the arrow refreshes now.
- Get HELIOS ready is the first-run checklist: service plans, price book, prices for quote-only tools, an AI provider, analysts, the first client, its security tools and threat feeds. Each item links to where it is done and ticks itself once it is.

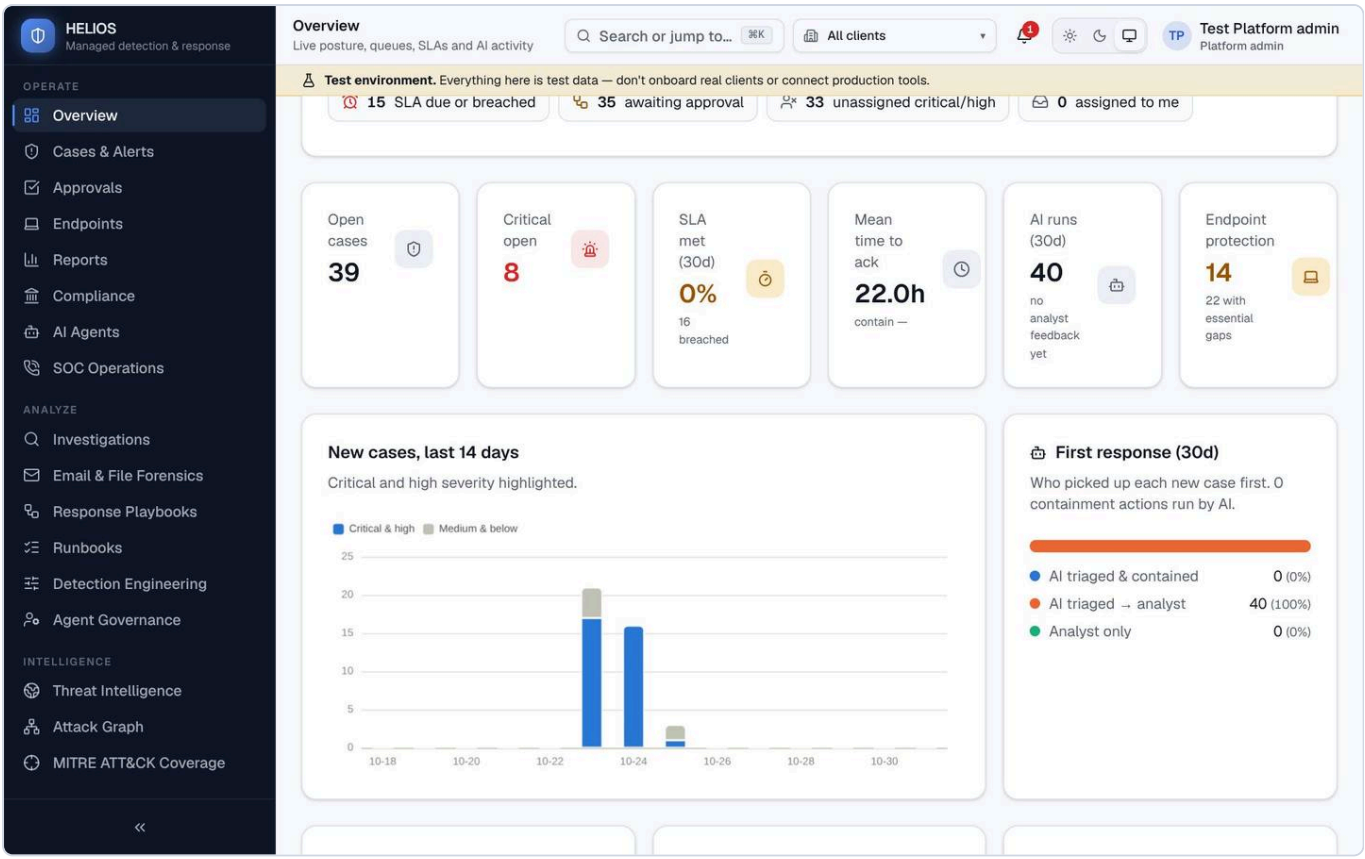


Figure 2. Overview, further down the page.

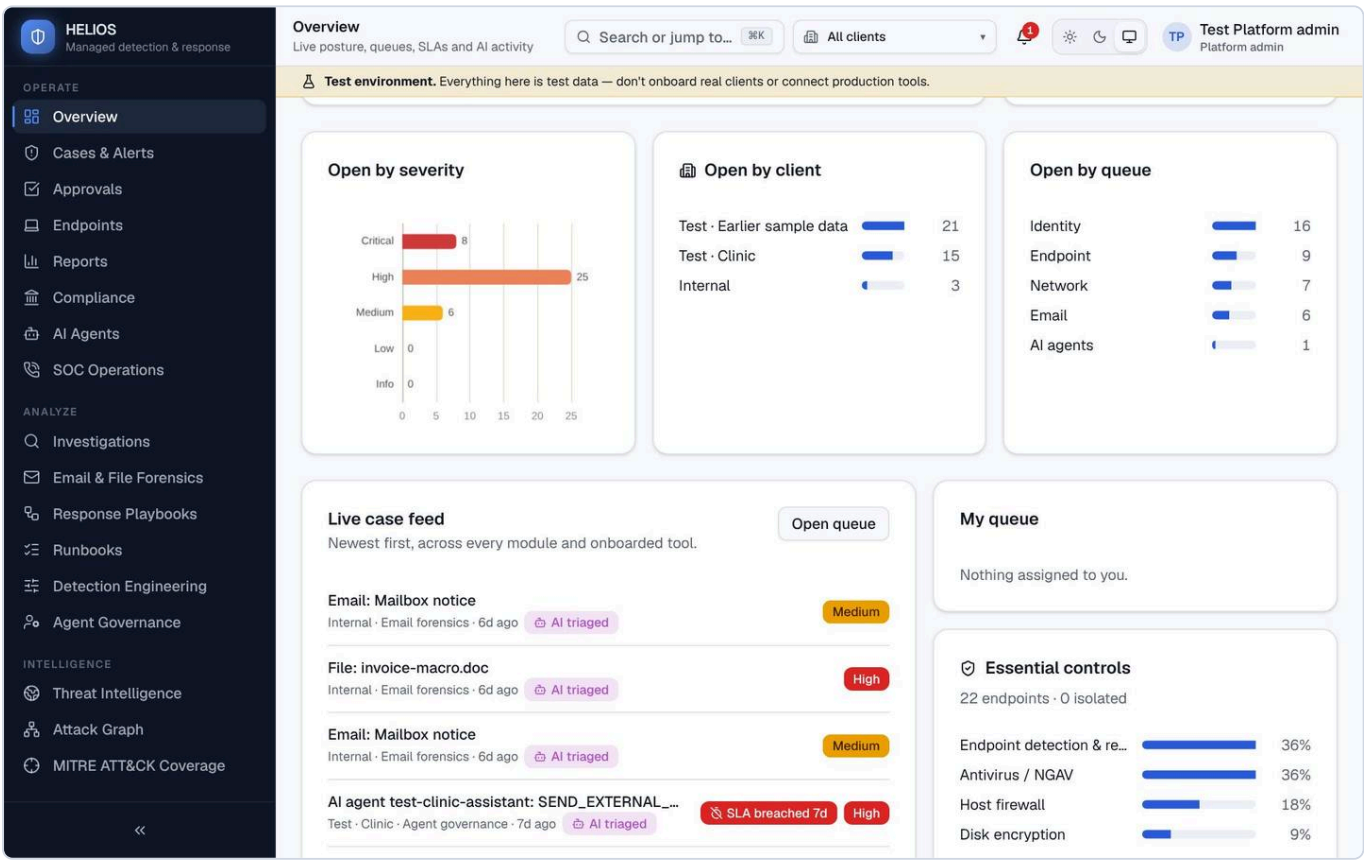


Figure 3. Overview, further down the page.

Cases & Alerts

Every case HELIOS has opened: detections from your clients' telemetry, vendor alerts, uploaded investigations, email and file analysis, and governed AI agents. Correlated detections for the same user, host or IP join one case instead of opening many.

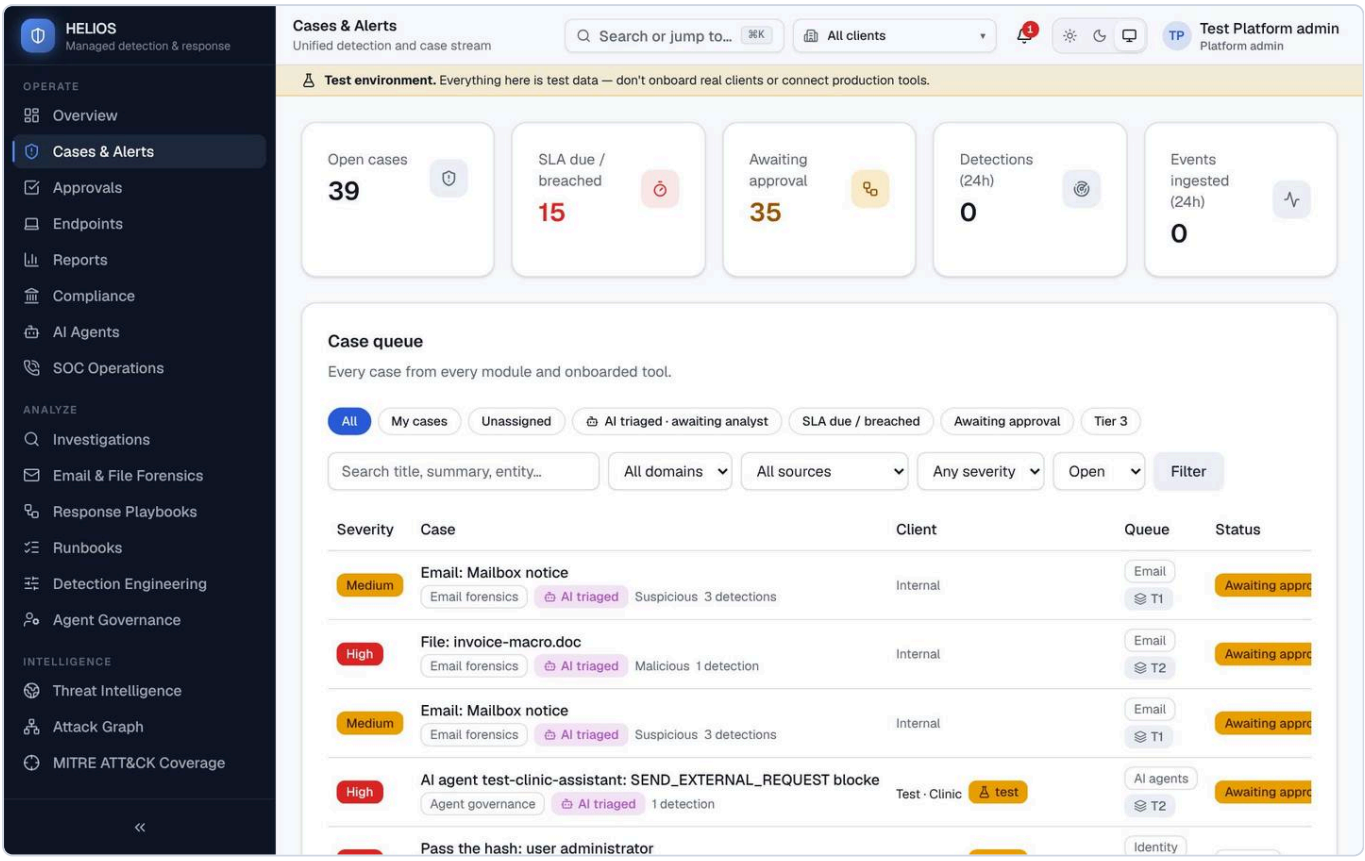


Figure 4. Cases & Alerts, as seen by a platform admin (test environment).

How to use it

- Filter by status, severity, source, client and assignee, or search by title. Open a case by clicking its row.
- You see every queue regardless of security domains.
- The SLA badge counts down to the acknowledgement deadline from the client's plan; a red badge means it has been breached.

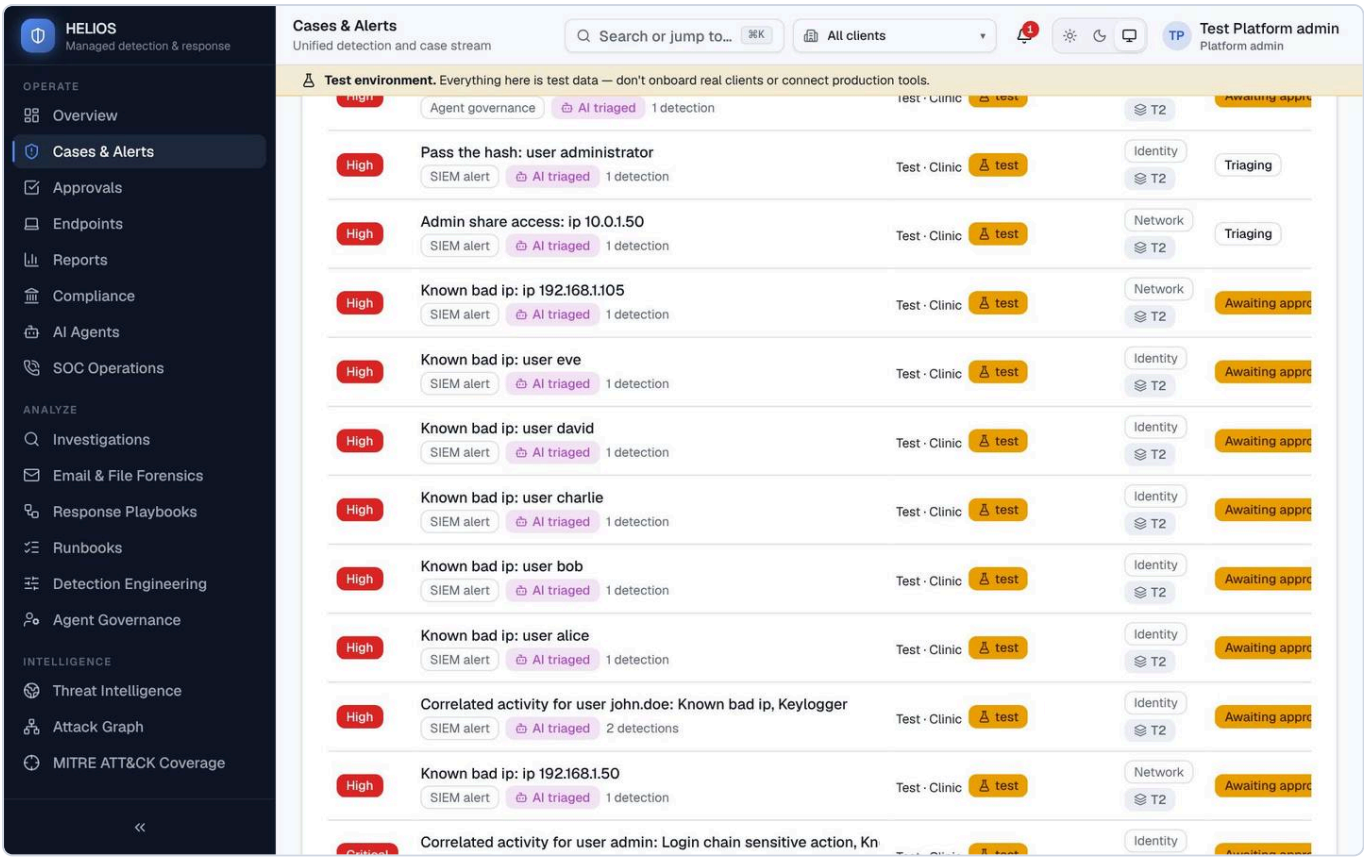


Figure 5. Cases & Alerts, further down the page.

A case

Everything about one case on one page: what was detected, what the AI made of it, what has been done

and what is waiting, the evidence, and the discussion.

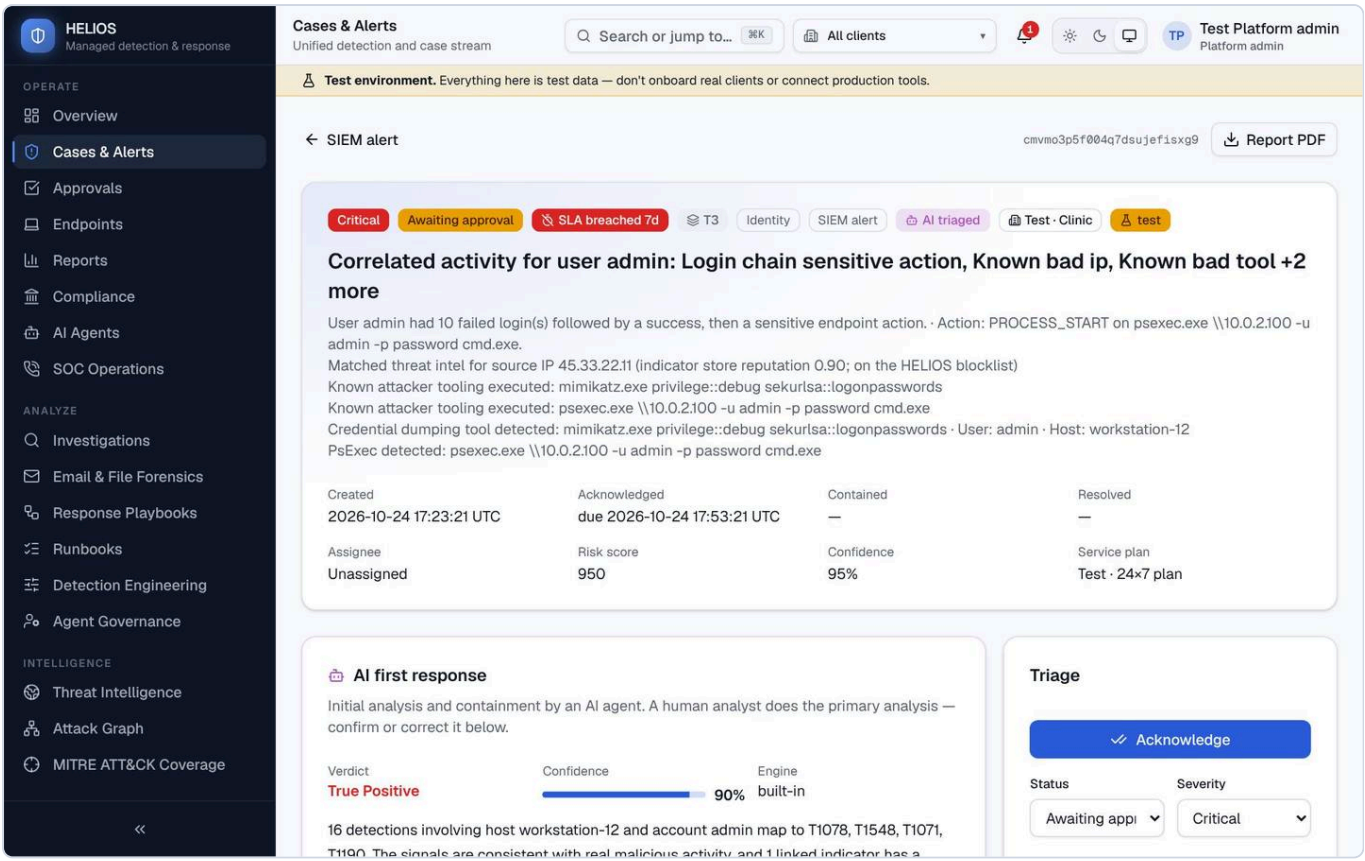


Figure 6. A case, as seen by a platform admin (test environment).

How to use it

- The header shows severity, status, the SLA, the tier (T1 triage, T2 analysis, T3 incident response), the security domain, the source, the client and whether AI triaged it. Below are the dates (created, acknowledged, contained, resolved), the assignee, the risk score, the confidence and the client's service plan. Report PDF downloads the incident report.
- AI first response shows the agent's verdict, confidence and engine, its summary, its key findings (each citing evidence: D detections, E events, I indicators, T timeline, C context) and what it asks the analyst to check. It is a first pass: a person confirms or corrects it.
- Agree or Disagree records whether the analysis was right; a disagreement asks what it got wrong. These ratings drive the AI's autonomy ladder.
- Dispatch agent runs another agent on the case: the triage first responder, the phishing analyst, the threat-intel enricher, the investigation agent (works the case with read-only tools) or the client report writer.
- In Triage: Acknowledge stops the SLA clock; Status and Severity change the case; Assignee gives it to someone.
- Watch case sends you its updates. Escalate hands the case to another tier or person with a note; Share makes it visible to the client or sends it to an external channel.
- Page on-call (in the On call card) pages whoever is on call about this case, with what you need from them.
- Response actions lists containment the AI proposed and people requested, each with its status. Approve or Reject what is awaiting approval; the client's rulebook and the response policy still apply.

- Request containment asks for an action (isolate host, block IP, disable account, kill process, quarantine file, block sender, collect logs, purge email, reset password, collect an investigation package…) against a target, with a reason shown to the approver.
- Indicators shows the case's IPs, domains, URLs, hashes and addresses with their reputation; the arrow re-checks one. Related cases lists earlier or concurrent cases sharing hosts, accounts, indicators or the same agent.
- Regulatory reporting shows whether CERT-In's 6-hour report or the DPDP breach notices apply, the incident types HELIOS proposes, and the notices' state.
- Runbooks attach themselves to matching cases; tick each step or skip it with a reason. A case can't be closed as a true positive while a required step is open.
- The discussion takes comments and @mentions. Internal comments stay with the SOC; client-visible ones reach the client.
- Closing: set the status to Resolved or False positive, then give a disposition (true positive, benign true positive, false positive) and a reason.

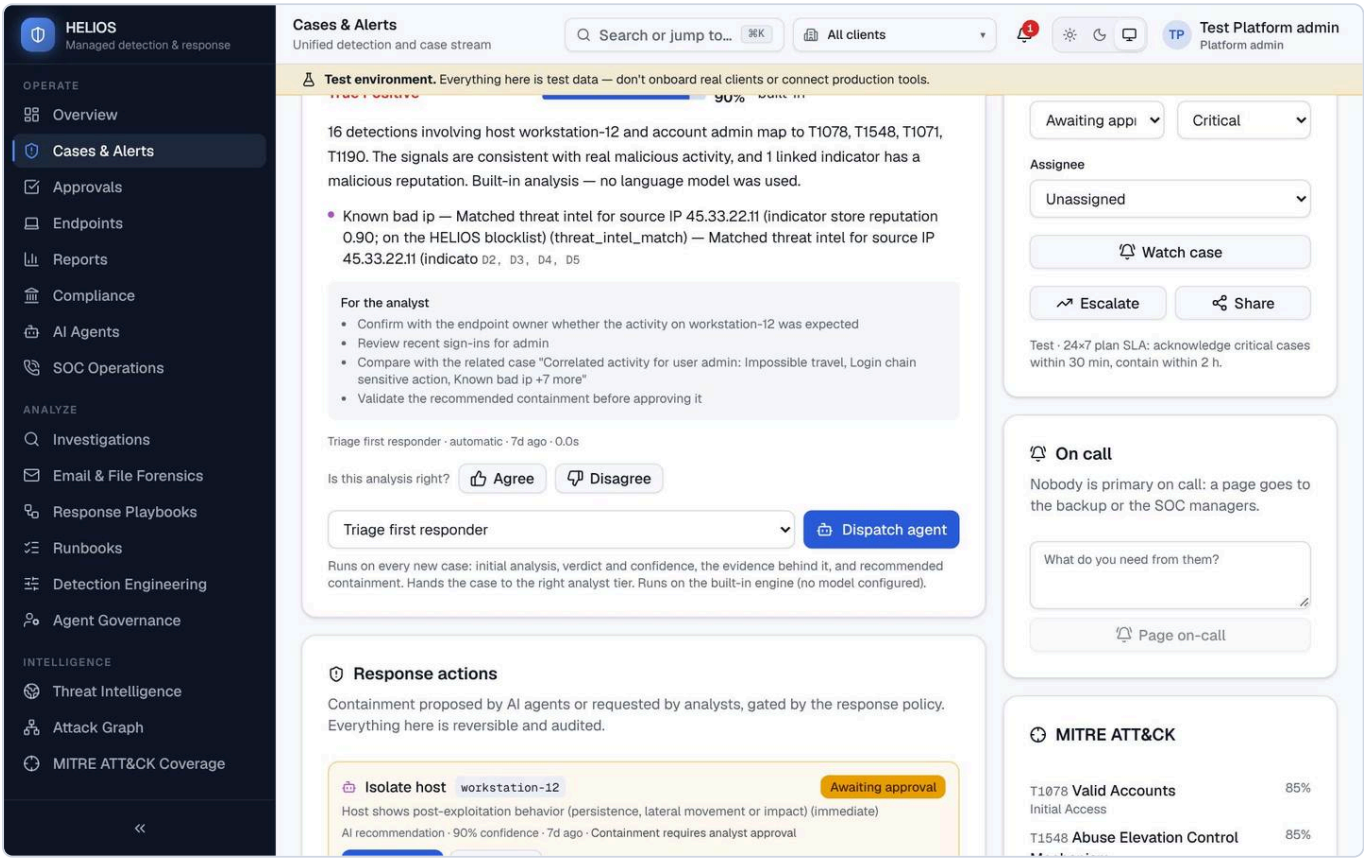


Figure 7. A case, further down the page.

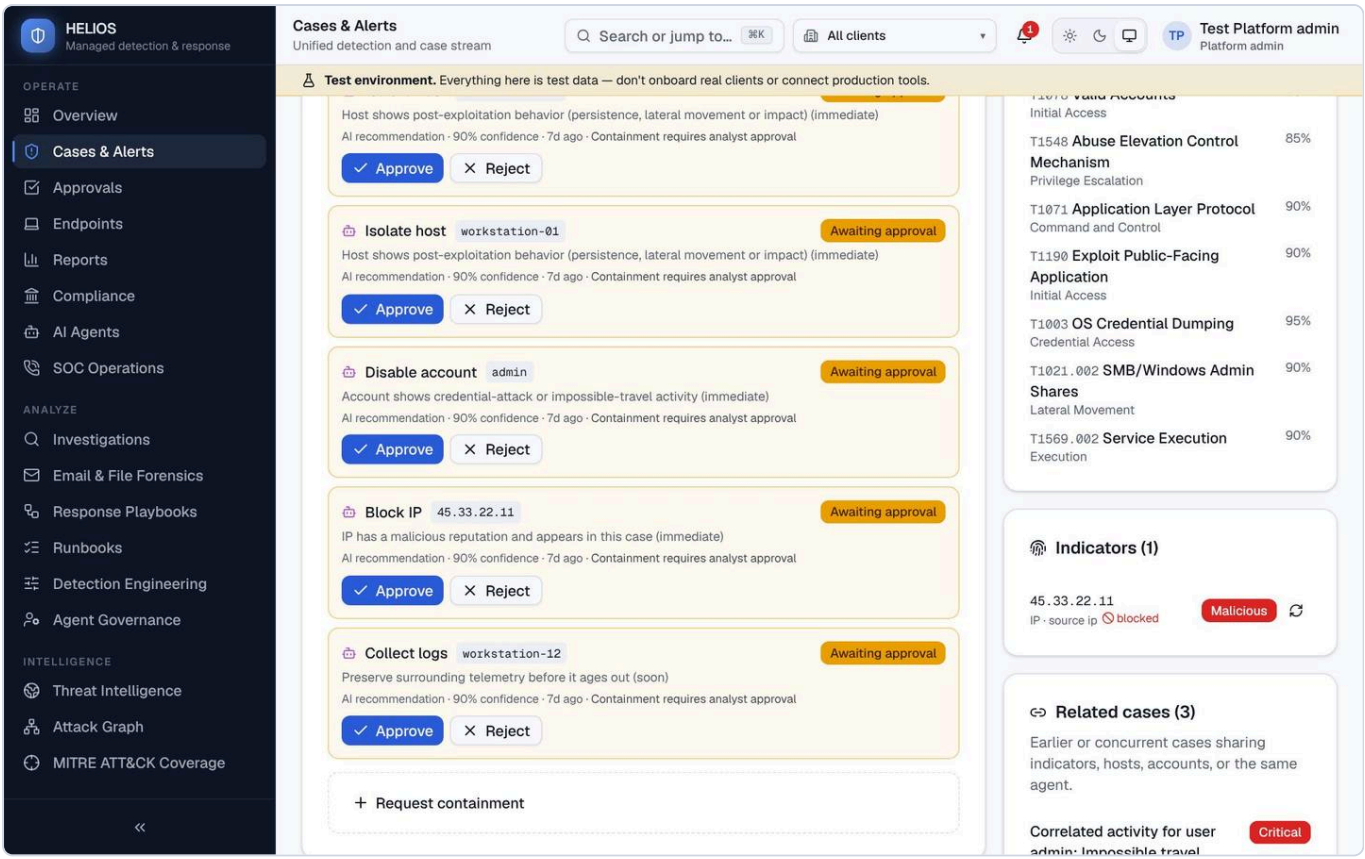


Figure 8. A case, further down the page.

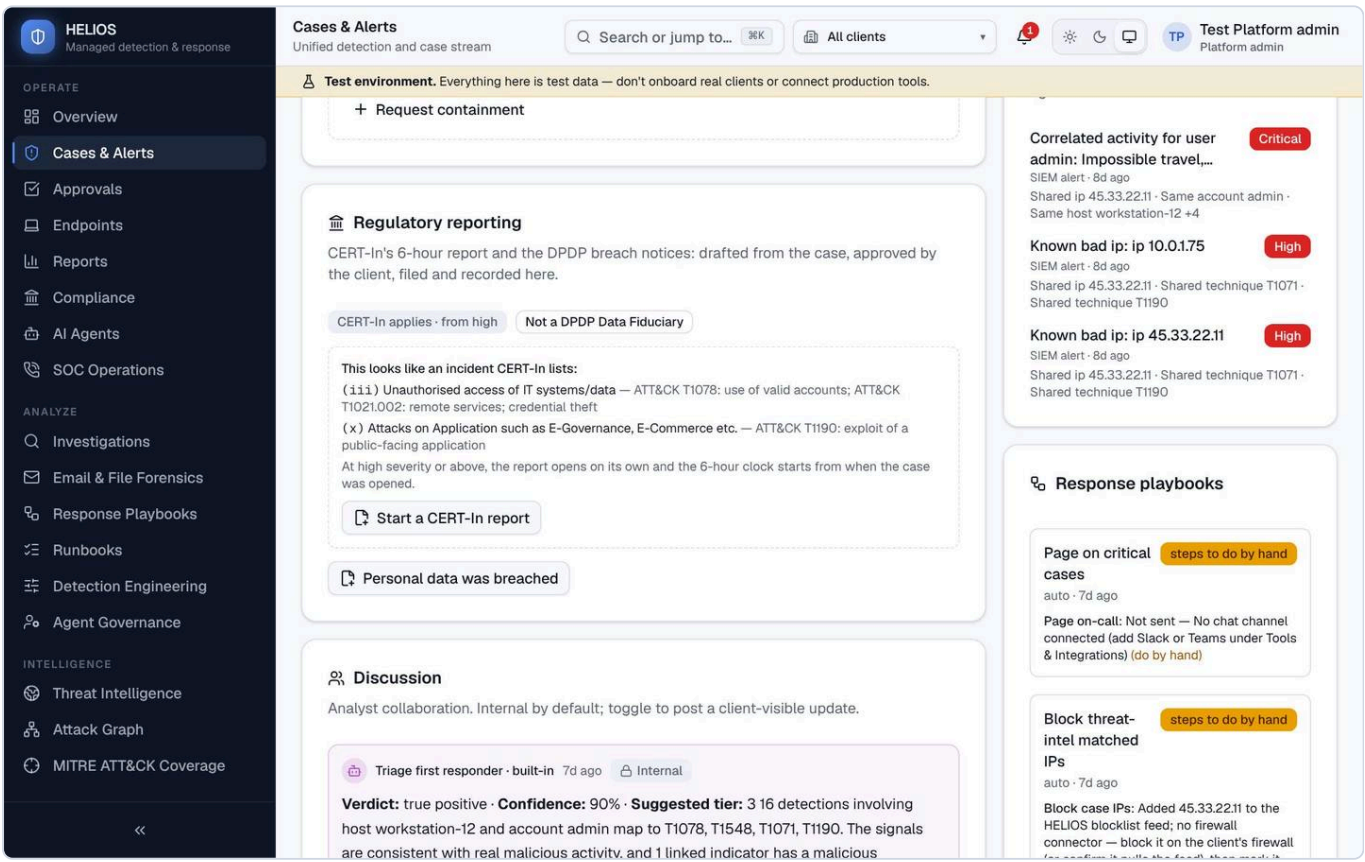


Figure 9. A case, further down the page.

Approvals

Everything waiting on someone's yes, in your scope: containment the response policy held, playbook runs that need approval, regulatory notices waiting on the client, and rulebooks to sign.

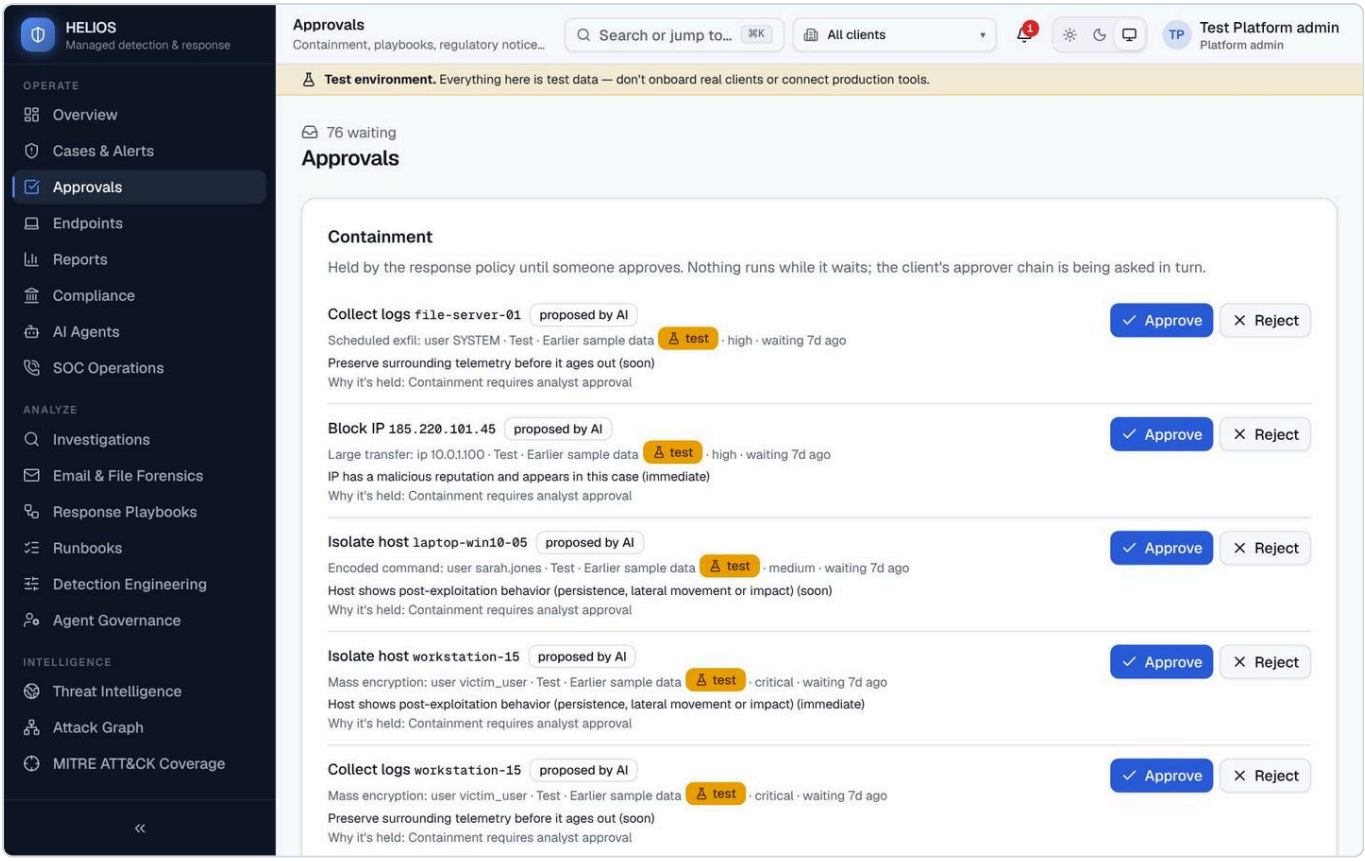


Figure 10. Approvals, as seen by a platform admin (test environment).

How to use it

- Each item says what would happen, against what, why it was held, who proposed it, the client, the severity and how long it has waited. Approve or Reject right there.
- Nothing held runs while it waits. The client's approver chain is asked in turn (in the app, by Slack or Teams, email, SMS or WhatsApp); when the chain runs out, the action holds or becomes a suggestion, as the client's rulebook says.
- A two-person rule in the rulebook needs two different approvers; one approval leaves it waiting for the second.

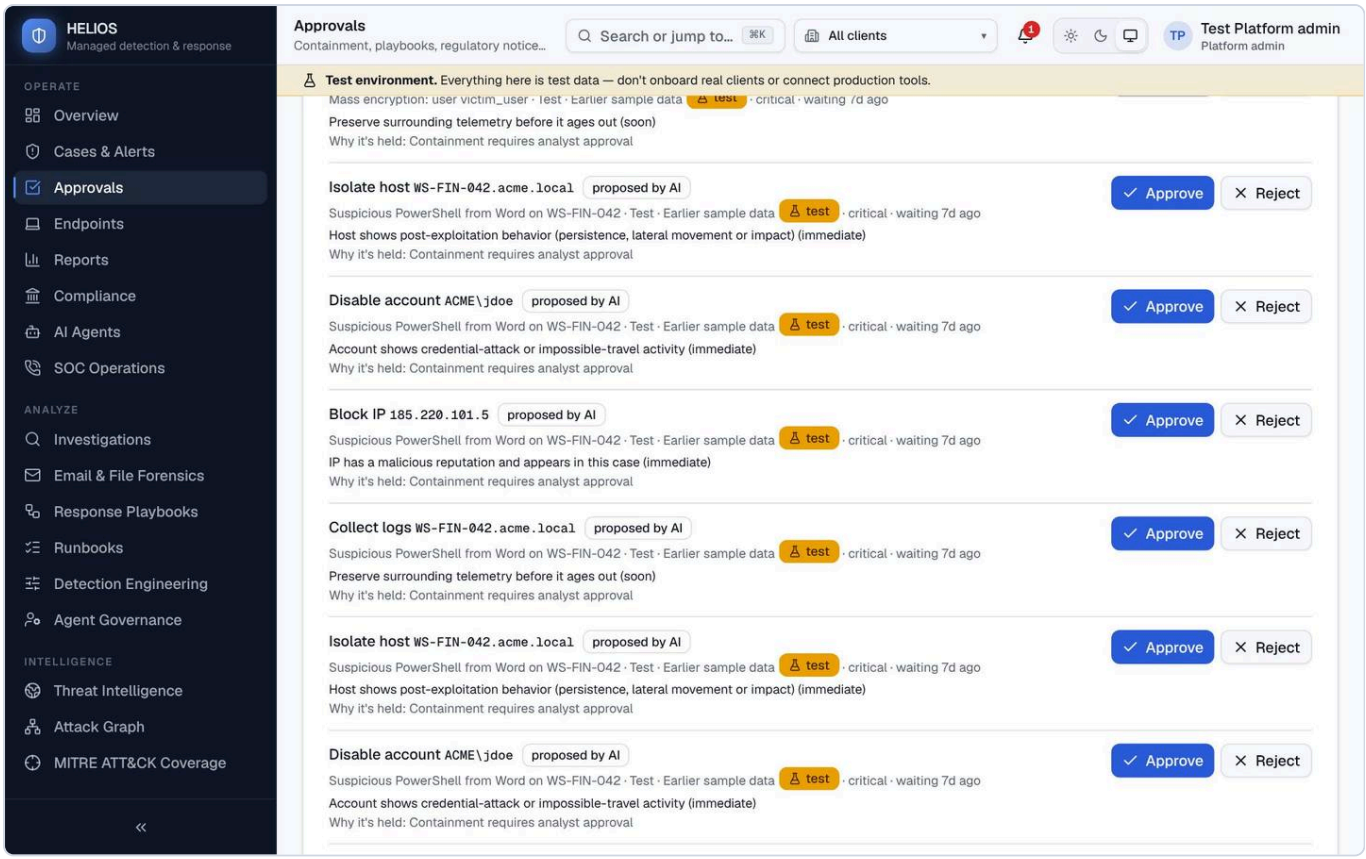


Figure 11. Approvals, further down the page.

Endpoints

The endpoint inventory and the whole attack surface: devices, servers, cloud, network, identity, IoT/OT

and applications, with how well each is protected.

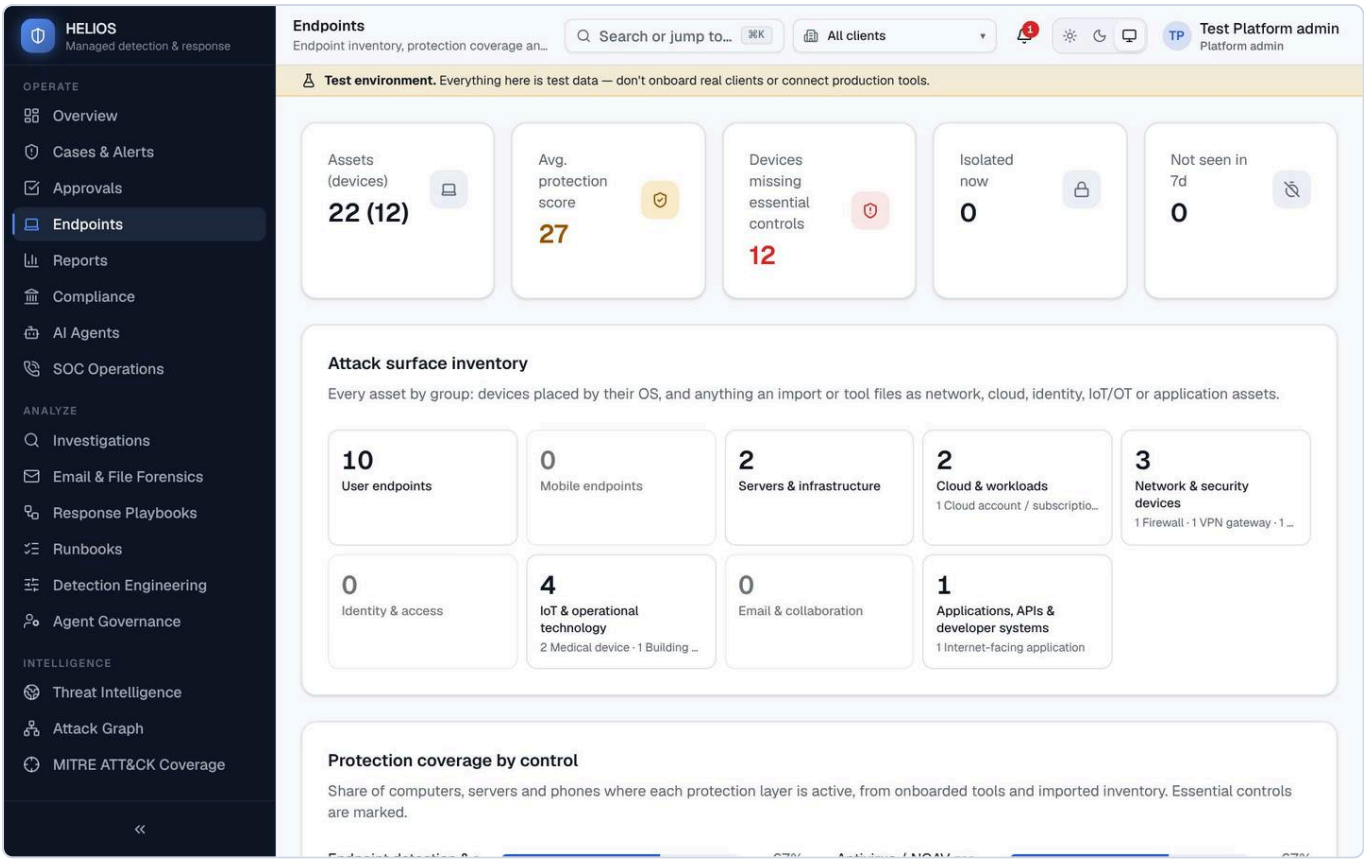


Figure 12. Endpoints, as seen by a platform admin (test environment).

How to use it

- The tiles count assets, the average protection score, devices missing essential controls, devices isolated now and devices not seen for seven days.
- Attack surface inventory groups every asset; Protection coverage by control shows the share of computers, servers and phones where each protection layer (EDR, antivirus, disk encryption, MFA...) is active, with the essential ones marked.
- Isolate or Release a device from its row. With an EDR connector HELIOS asks the vendor and confirms the device's state; without one it records the change and asks you to make it in the EDR console. The client's rulebook still decides: a never-isolate asset can't be isolated, a crown jewel needs approval.

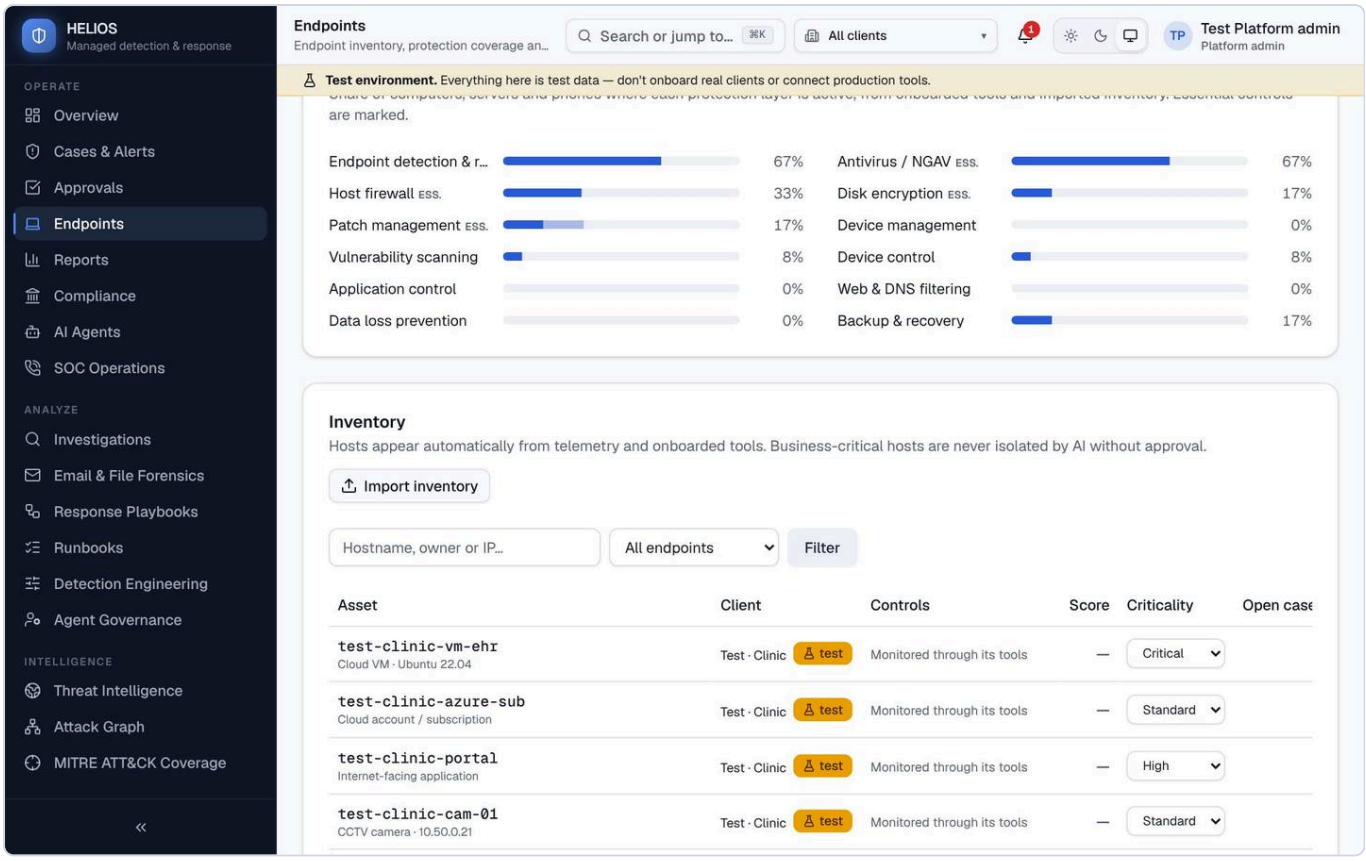


Figure 13. Endpoints, further down the page.

Reports

How the service performed over a period: volume, response times, SLA compliance, AI performance and

analyst workload, each compared with the period before.

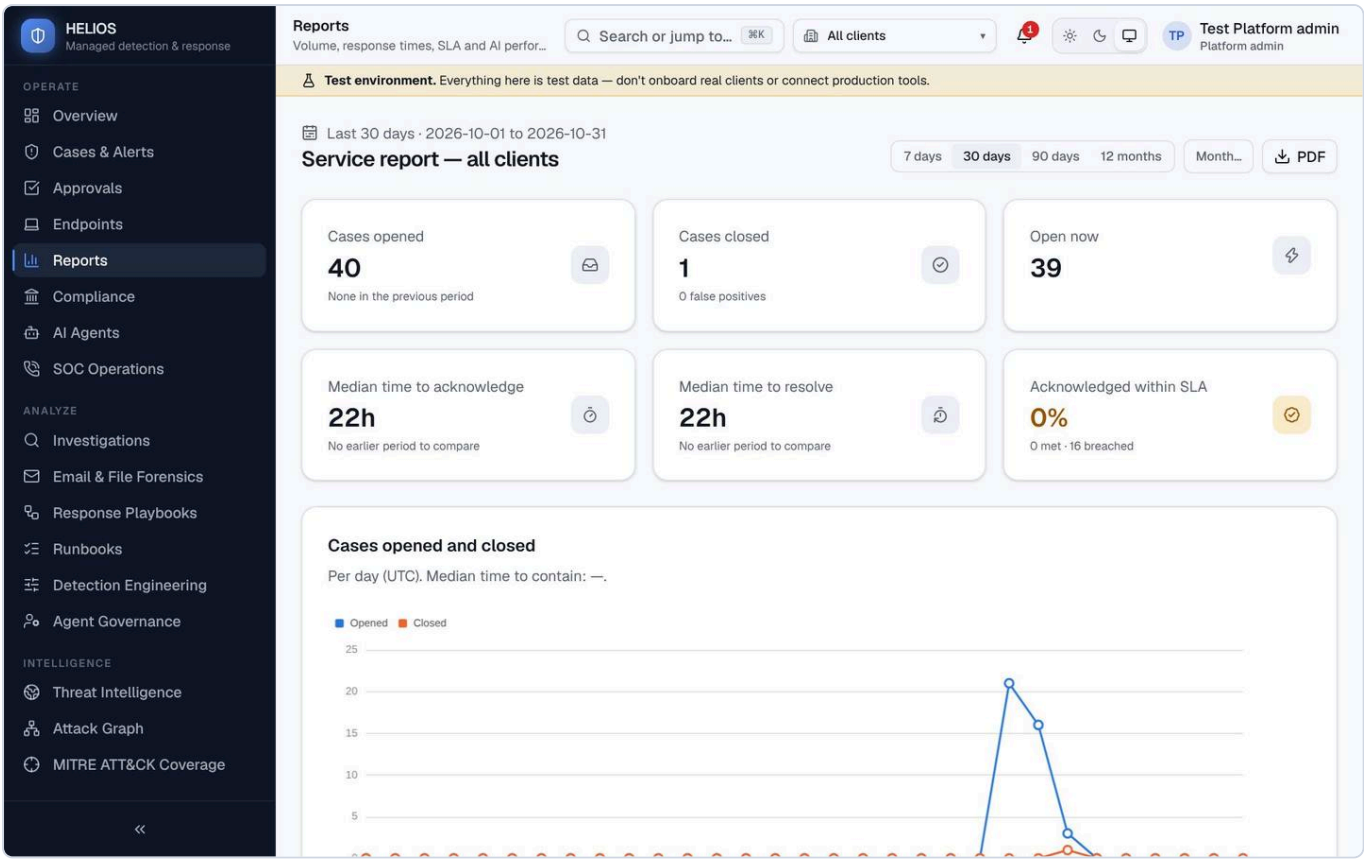


Figure 14. Reports, as seen by a platform admin (test environment).

How to use it

- Pick a period or a month. The figures include median time to acknowledge, contain and resolve, acknowledgement SLA compliance, the mix by severity, source, domain and client, top ATT&CK techniques and indicators, AI runs and analyst agreement, and containment outcomes.
- Download the service report as a PDF for the client.

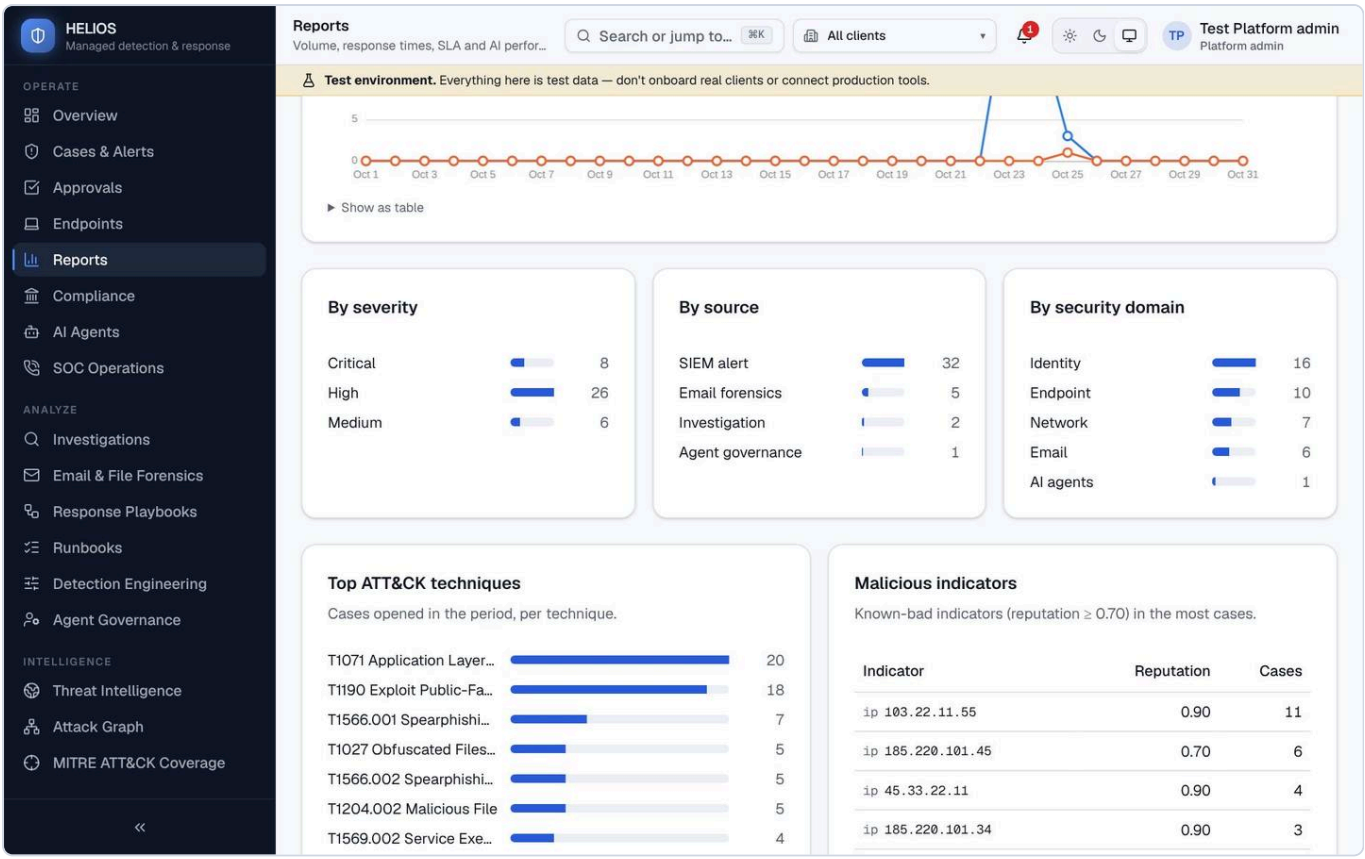


Figure 15. Reports, further down the page.

Compliance

Regulatory clocks and evidence: CERT-In's 6-hour incident reports, the DPDP breach notices, and the

ISO/IEC 27001 evidence map.

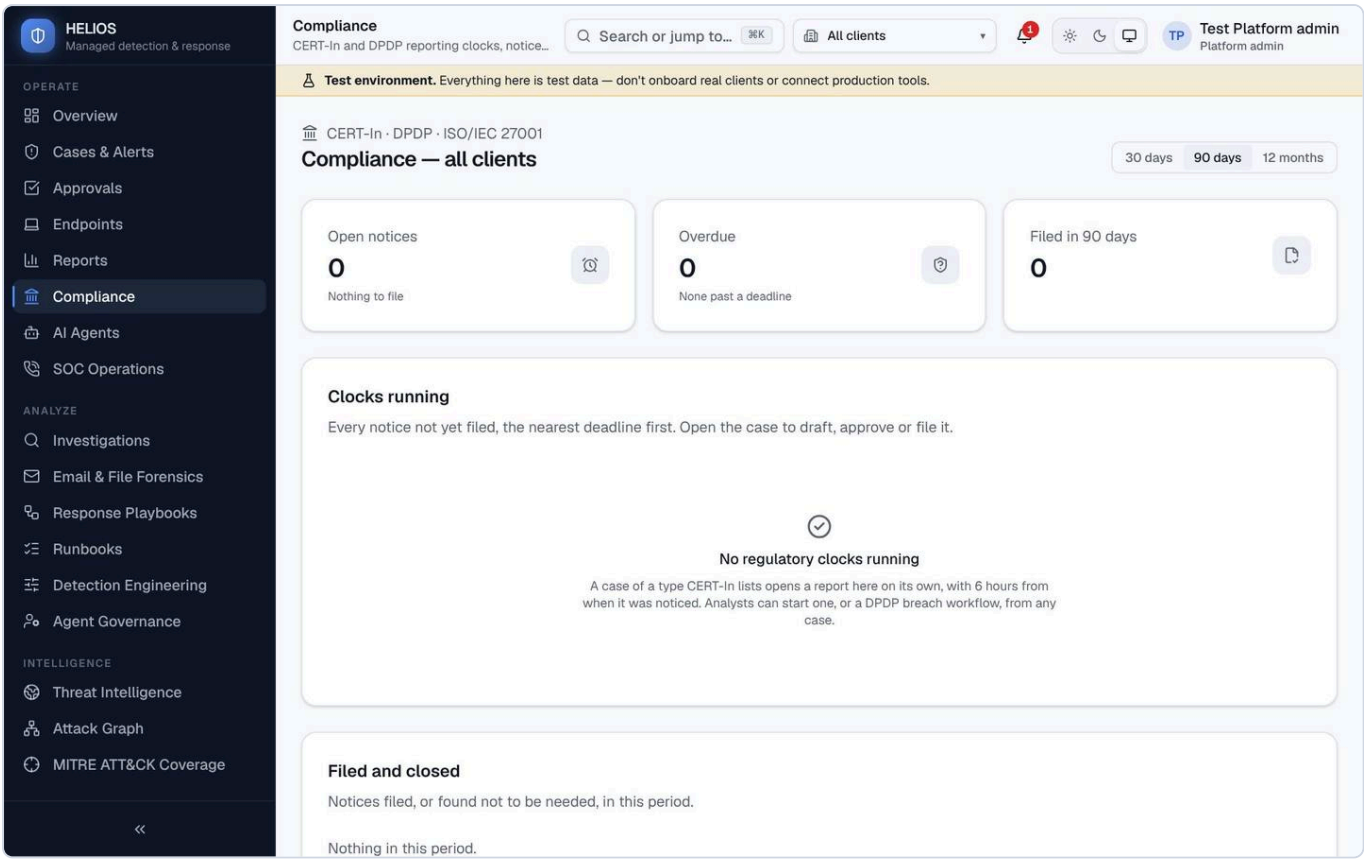


Figure 16. Compliance, as seen by a platform admin (test environment).

How to use it

- Clocks running lists every notice not yet filed, nearest deadline first. CERT-In: incidents of the 20 listed types are reported within 6 hours of being noticed. DPDP: the Board and each affected person are told without delay, and the Board gets a detailed report within 72 hours.
- Open the case to draft, approve or file a notice. Each notice is drafted from the case, approved by the client, then filed (emailed to CERT-In with its PDF, or recorded with the regulator's reference).
- A notice goes to the client for approval before it is filed: approve it here or from Approvals. An edit after approval sends it back.
- The ISO 27001 view maps 15 Annex A controls to what HELIOS's records show over the period, and shows gaps (controls with nothing behind them) plainly.
- Switch the period (30 days, 90 days, 12 months) at the top.

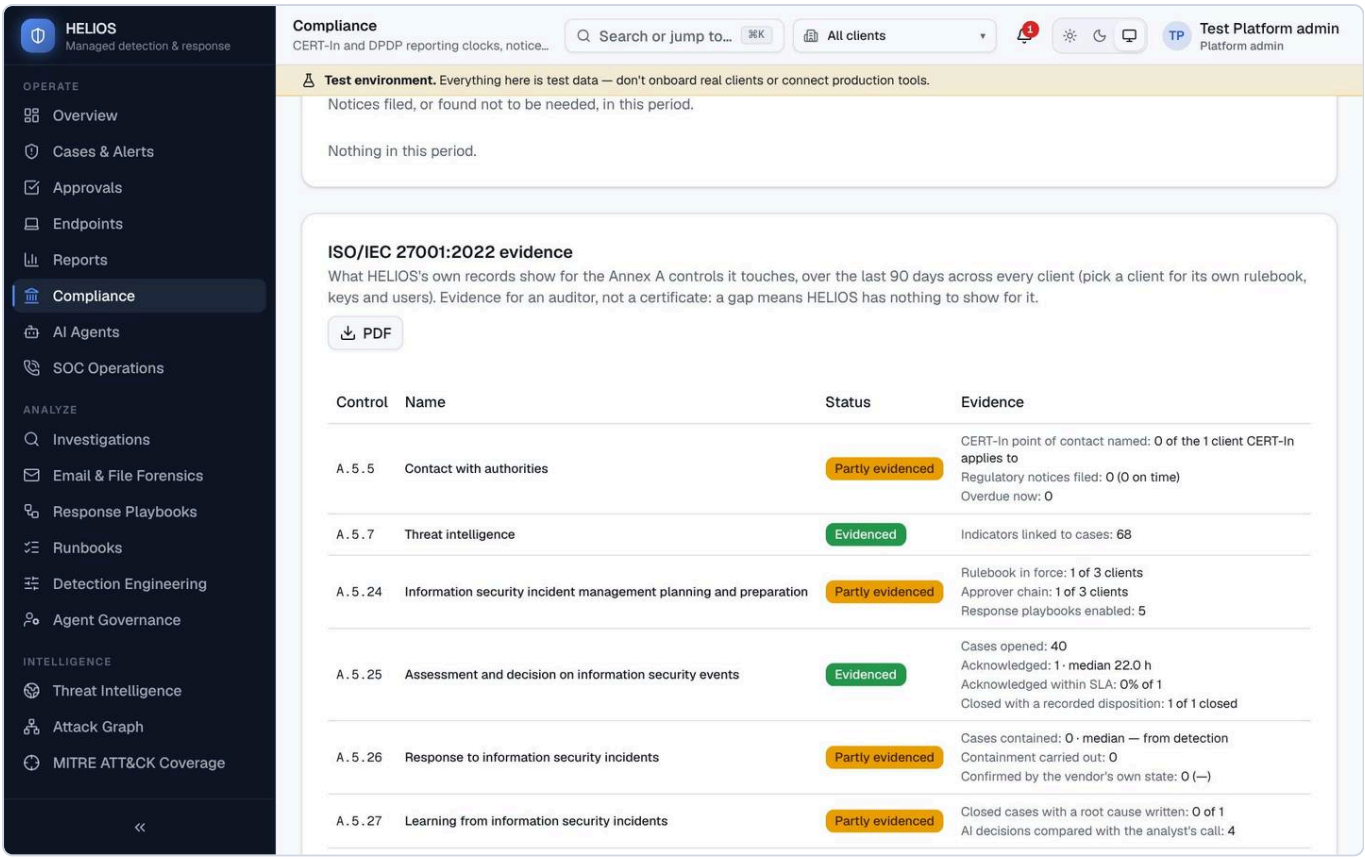


Figure 17. Compliance, further down the page.

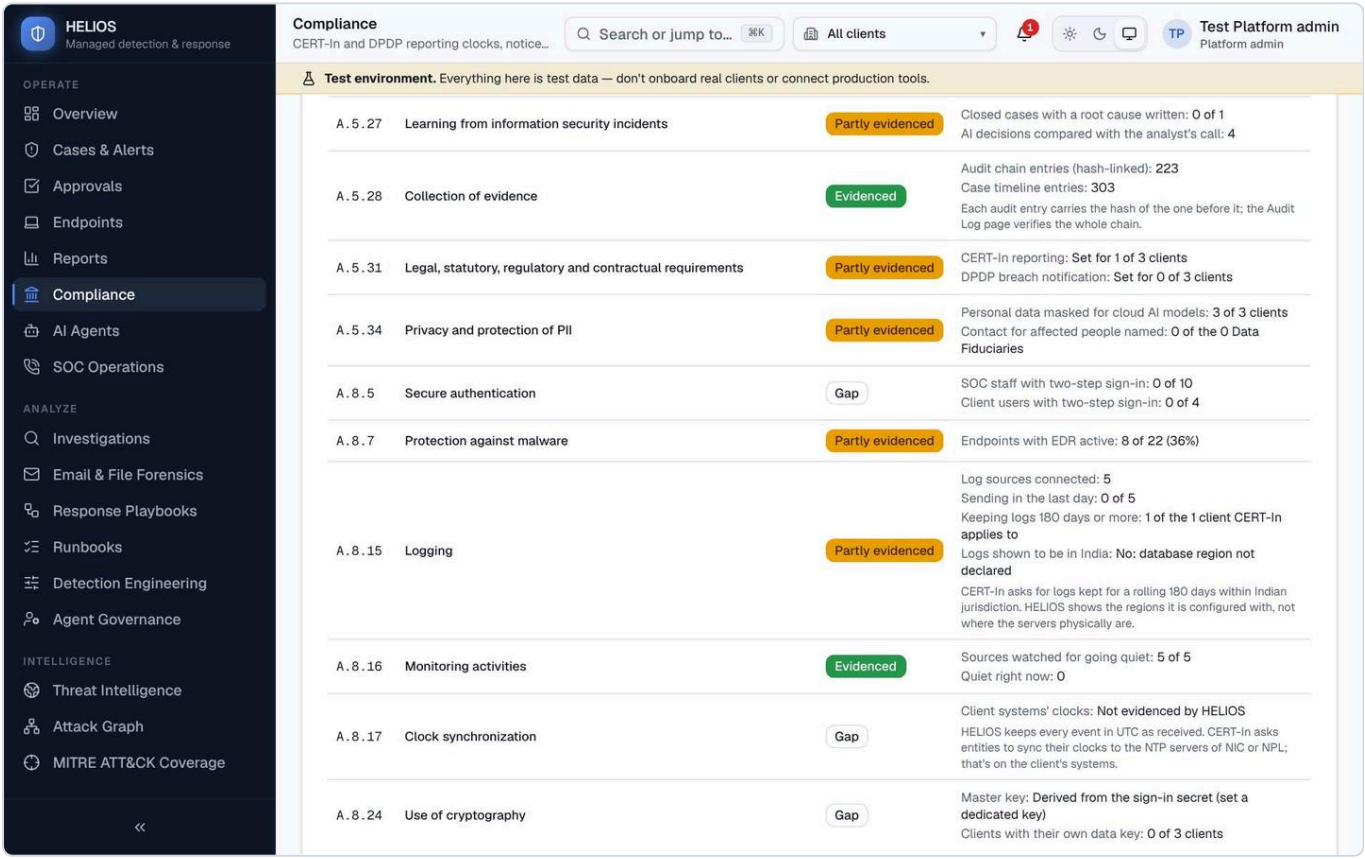


Figure 18. Compliance, further down the page.

AI Agents

The AI agents that work cases: what they did, how far they may act on their own, which models run them, and how those models scored.

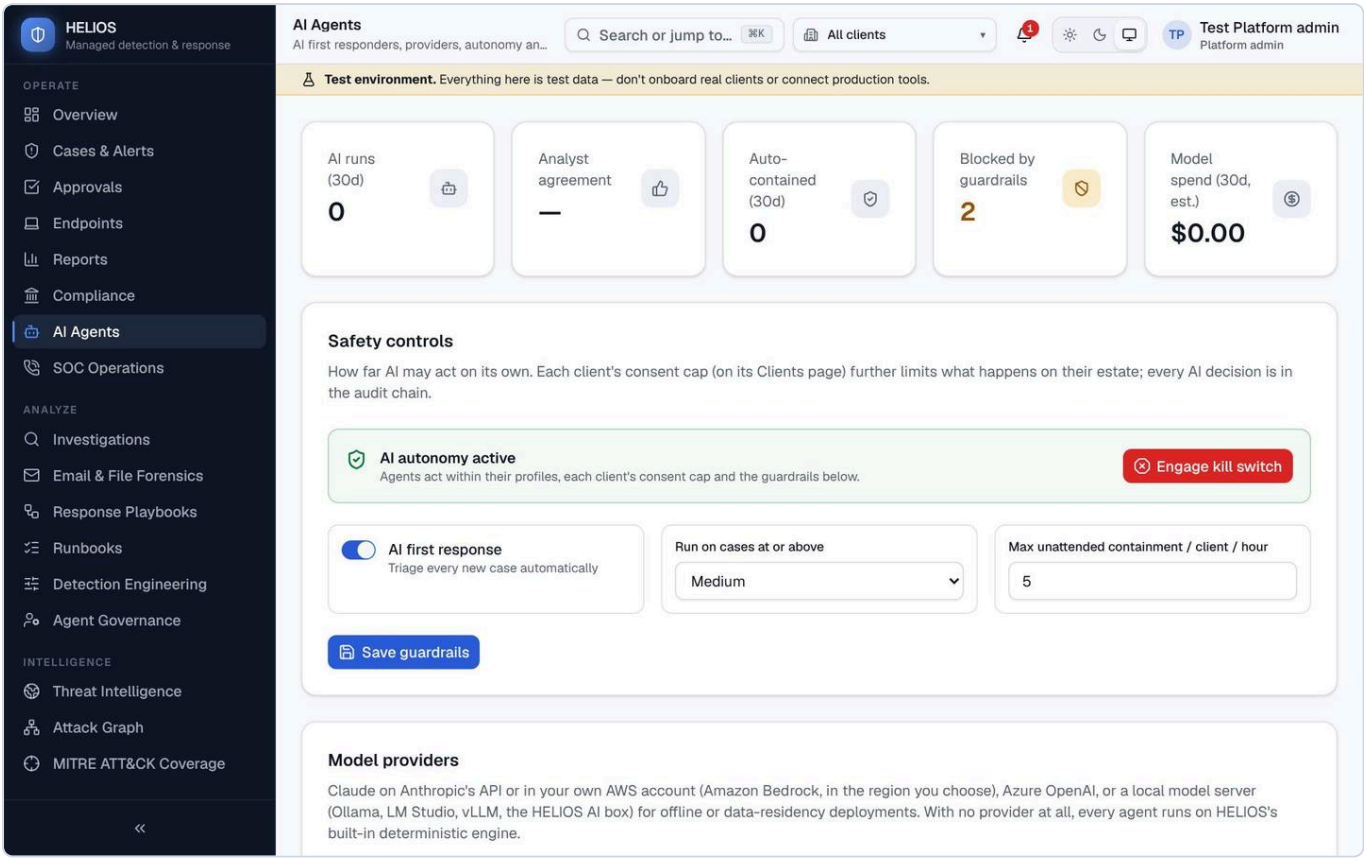


Figure 19. AI Agents, as seen by a platform admin (test environment).

How to use it

- The tiles show AI runs in the last 30 days, analyst agreement, actions contained automatically, actions the guardrails blocked and the estimated model spend (with the share of input read from the prompt cache).
- Safety controls: Engage kill switch stops all AI action at once; AI first response triages every new case at or above a severity; Max unattended containment caps automatic actions per client per hour.
- Model providers: add Claude (Anthropic's API, or in your own AWS account through Amazon Bedrock), Azure OpenAI, or a local model server (Ollama, LM Studio, vLLM, the HELIOS AI box). Test checks the connection; Make default sets the model every agent uses unless told otherwise. Keys are encrypted and never shown again.
- Agent types: for each agent choose the provider, a model override, the reasoning effort, the autonomy (off, suggest only, contain with approval, auto-contain pre-approved actions) and extra instructions.
- Evaluations score a model on a labelled set (attacks, benign activity, look-alikes and planted instructions); the investigation agent is scored on worlds its tools search. A model a client or user picks only suggests until it passes for that agent.
- The run log lists recent runs with their verdict, model, tokens and outcome.

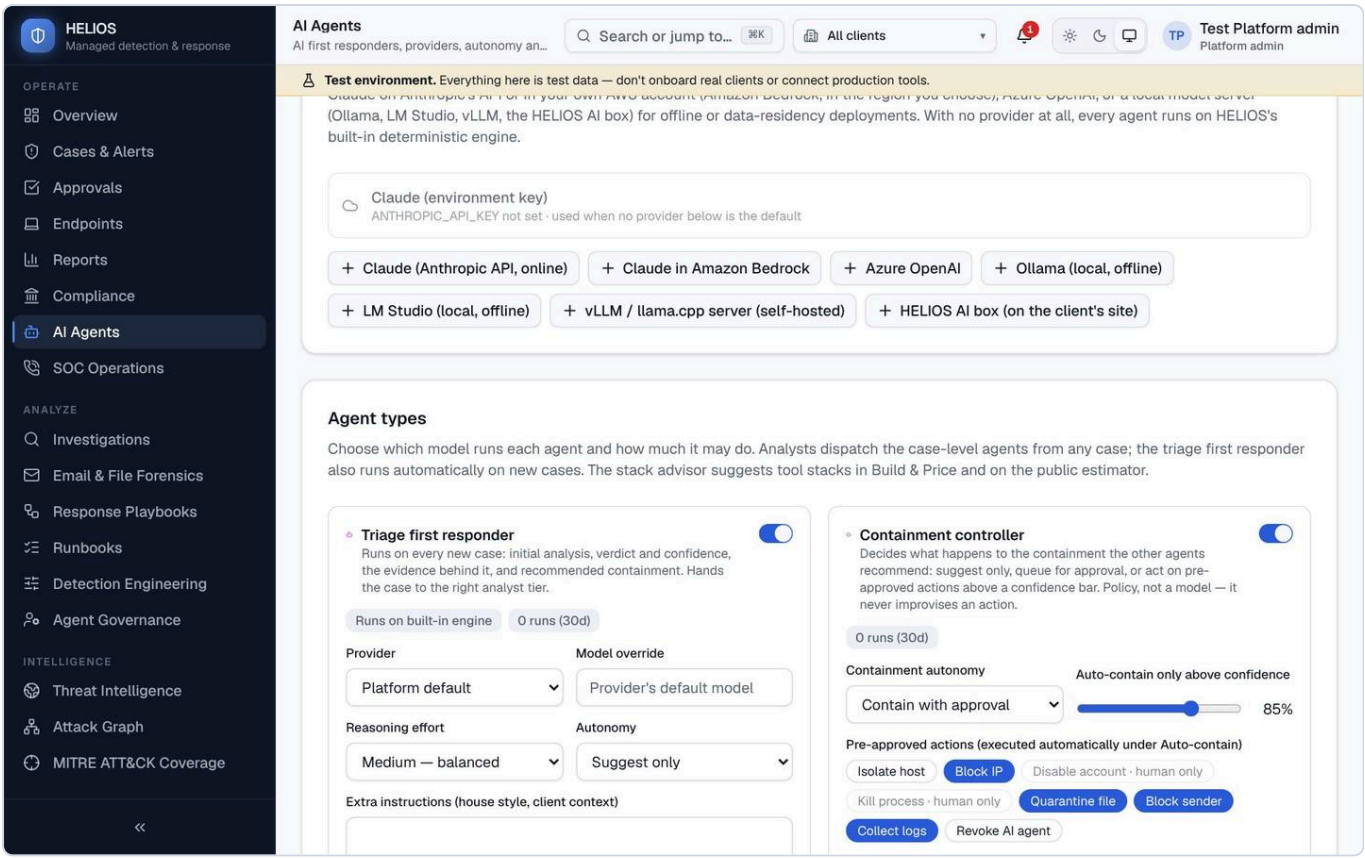


Figure 20. AI Agents, further down the page.

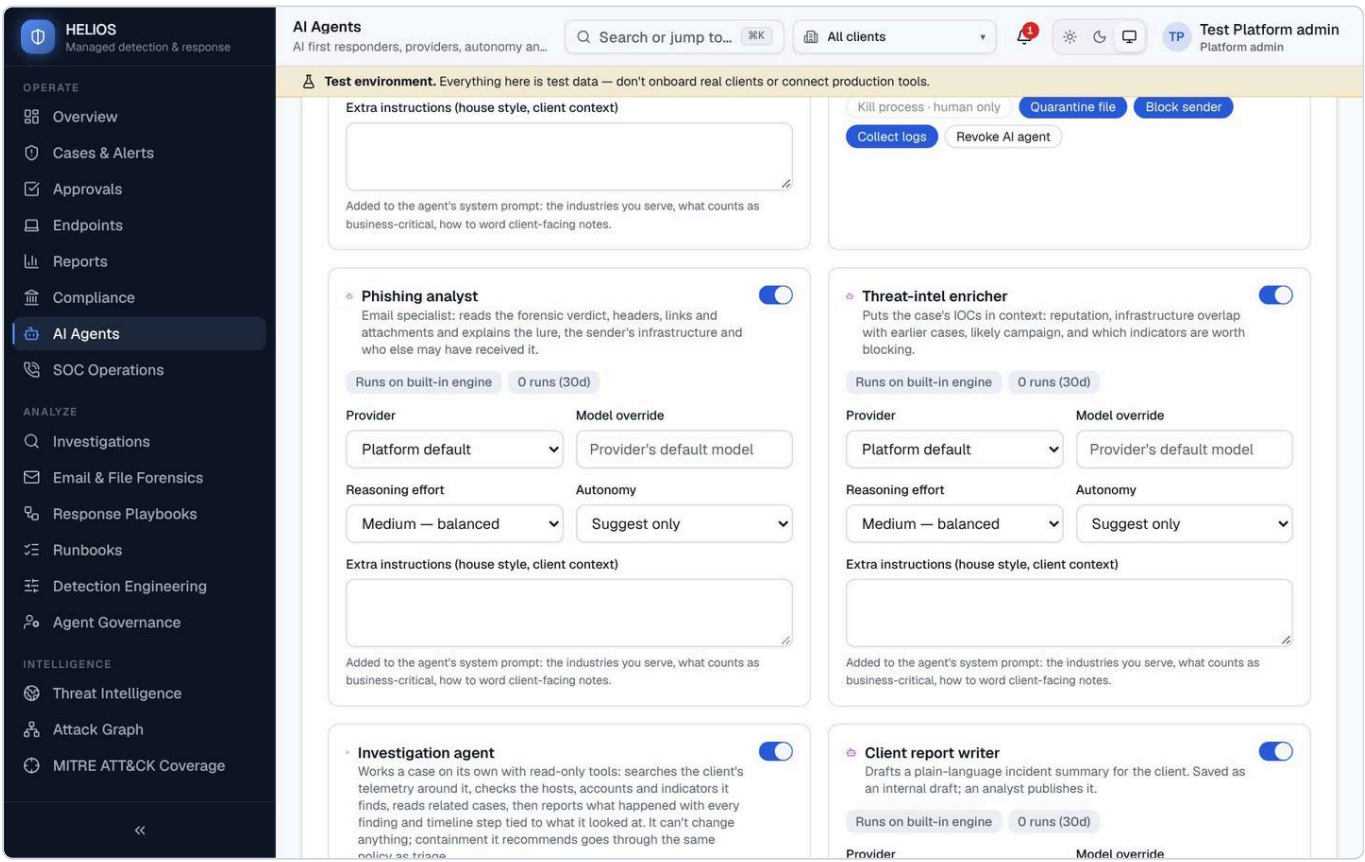


Figure 21. AI Agents, further down the page.

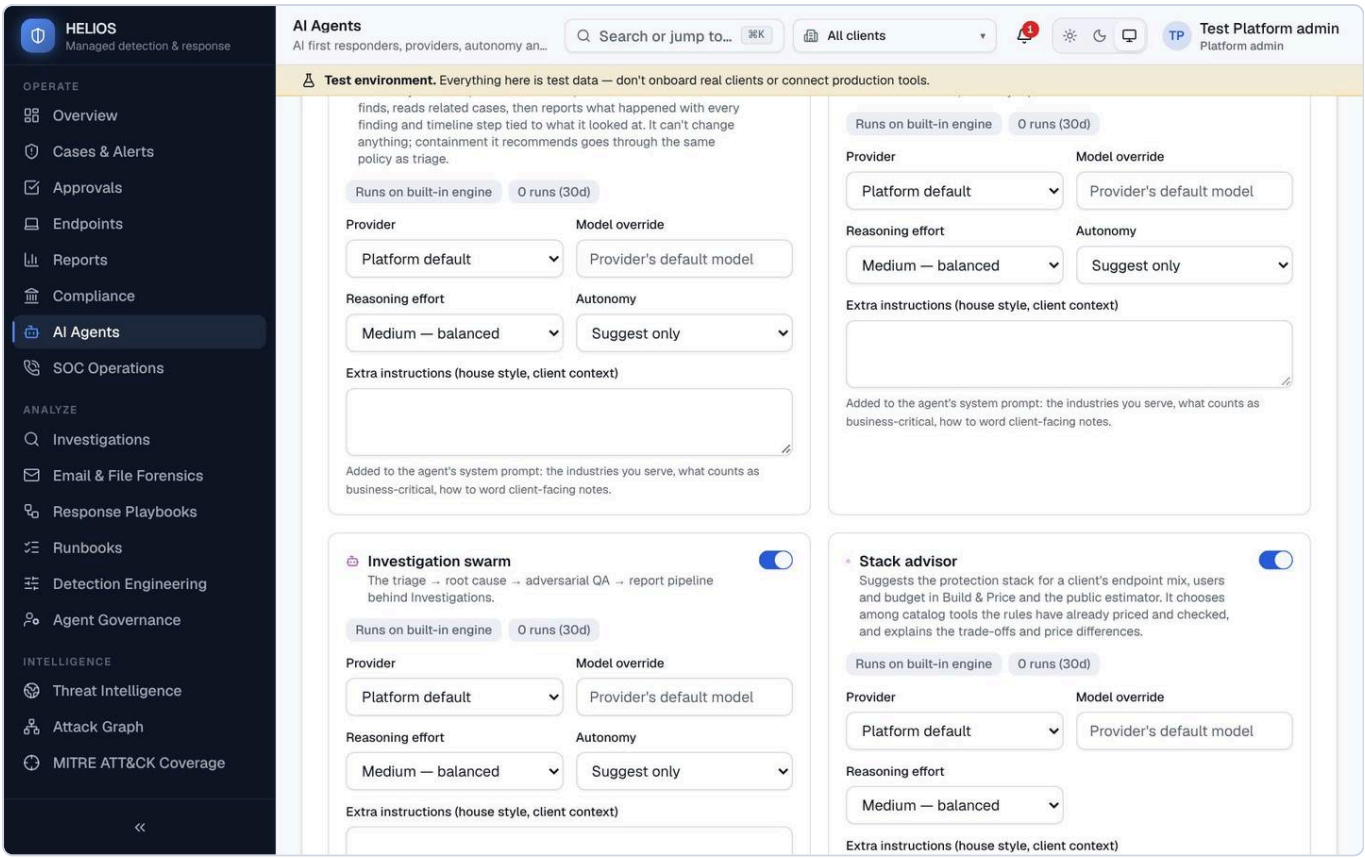


Figure 22. AI Agents, further down the page.

SOC Operations

Who is on call and the pages they get, the weekly QA of what the AI did on its own, and the daily digests.

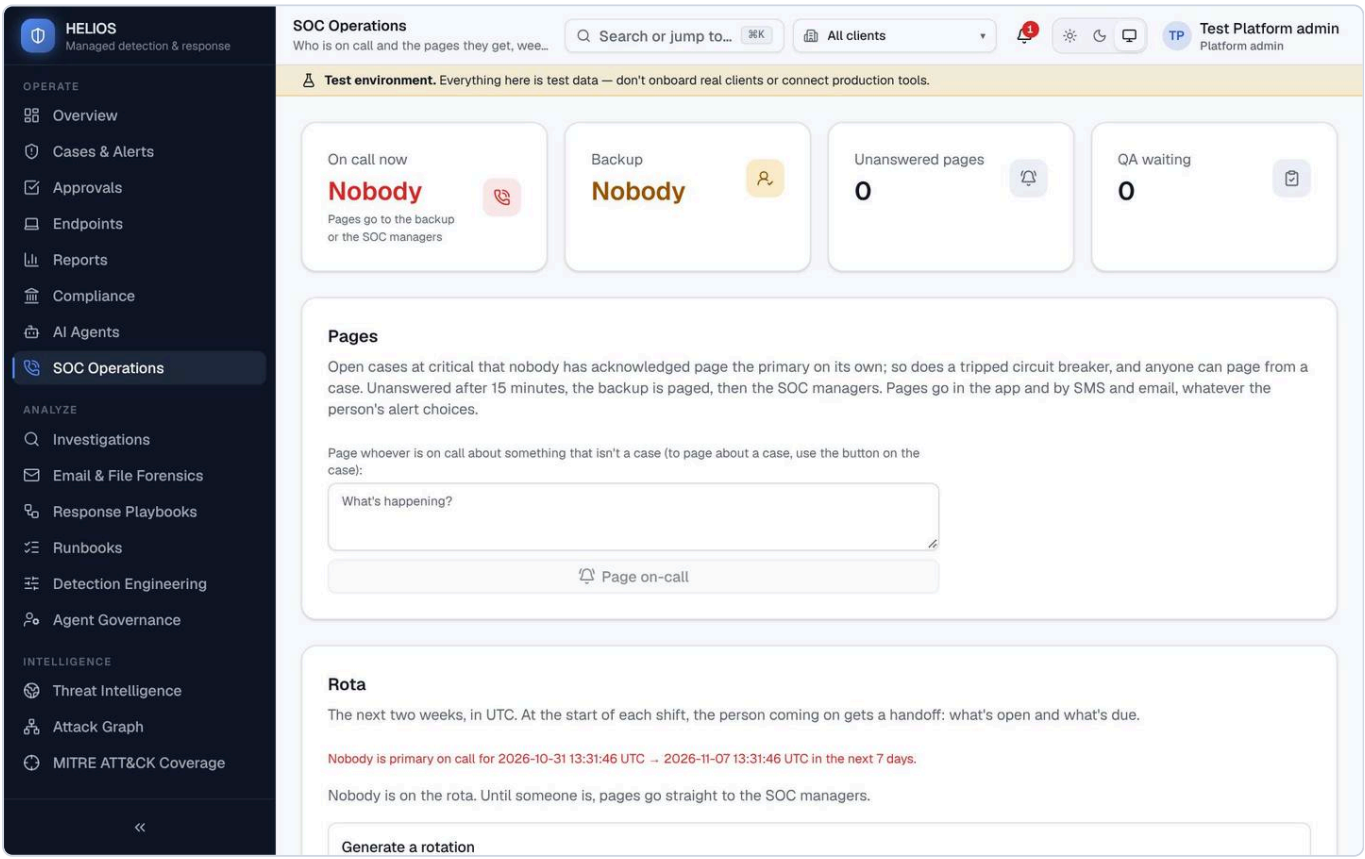


Figure 23. SOC Operations, as seen by a platform admin (test environment).

How to use it

- The tiles show who is primary and backup on call now, unanswered pages and QA reviews waiting.
- Pages: open critical cases nobody acknowledged page the primary on their own; unanswered after the acknowledgement window (15 minutes by default) the backup is paged, then the SOC managers. Pages go in the app and by SMS and email. Page on-call sends a page about something that isn't a case.
- Rota: Generate a rotation from people taking turns (8, 12 or 24-hour shifts, or a week); add cover for a swap or leave. Whoever starts a shift gets a handoff: open high and critical cases, unanswered pages, approvals waiting, notices due, quiet sources.
- QA: each Monday HELIOS samples last week's cases where the AI acted without anyone approving or called a case benign (weighted towards critical cases, new rules and critical assets). Review each one on the case: should the action have run, was the verdict right. A 'shouldn't have run' counts against the AI's autonomy at once.
- Settings: paging severity, acknowledgement window, QA share and floor, and the digest time and time zone.

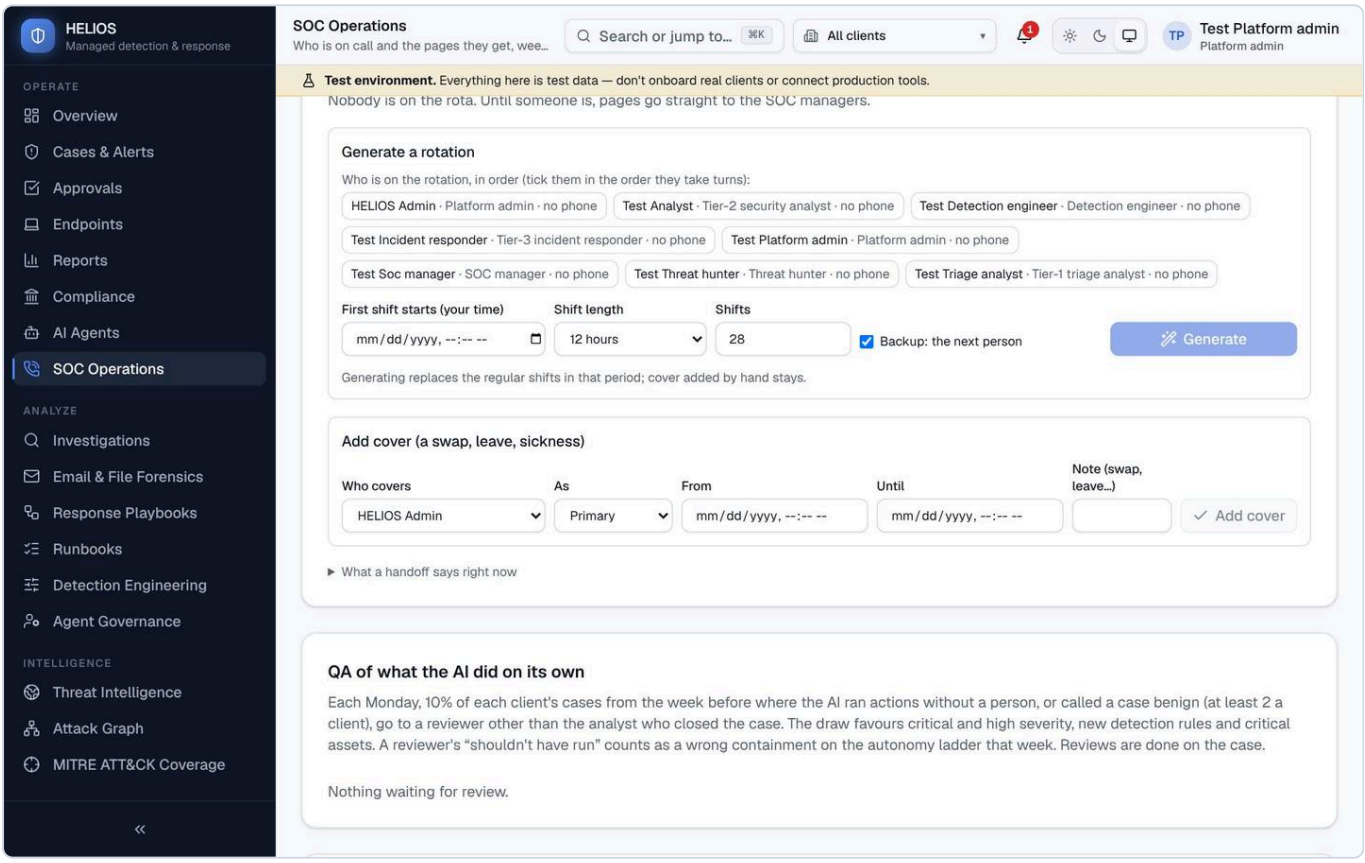


Figure 24. SOC Operations, further down the page.

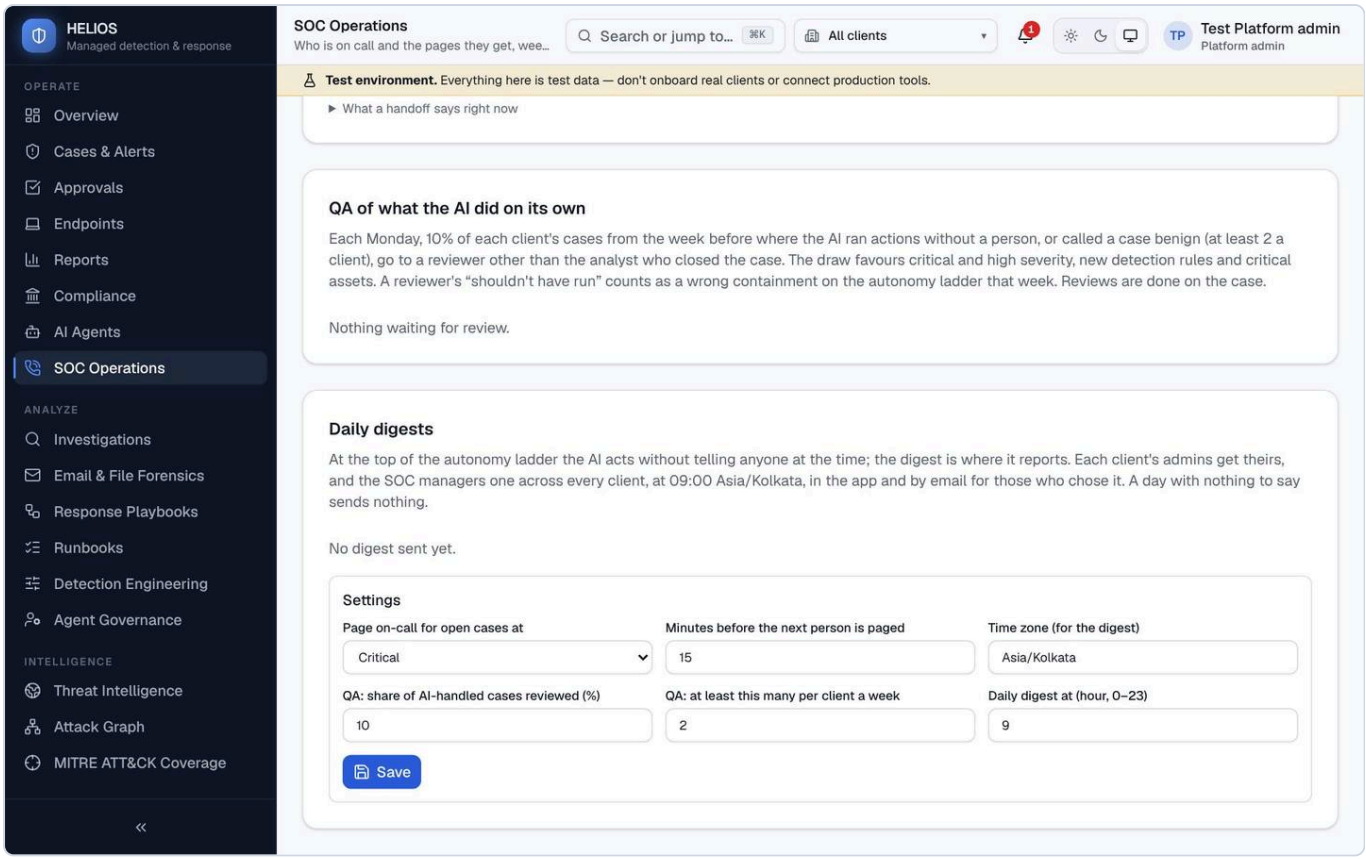


Figure 25. SOC Operations, further down the page.

Investigations

Upload logs from an incident and get a root-cause analysis: an agent pipeline (triage, root cause,

adversarial QA, report) works through them and opens a case with the findings.

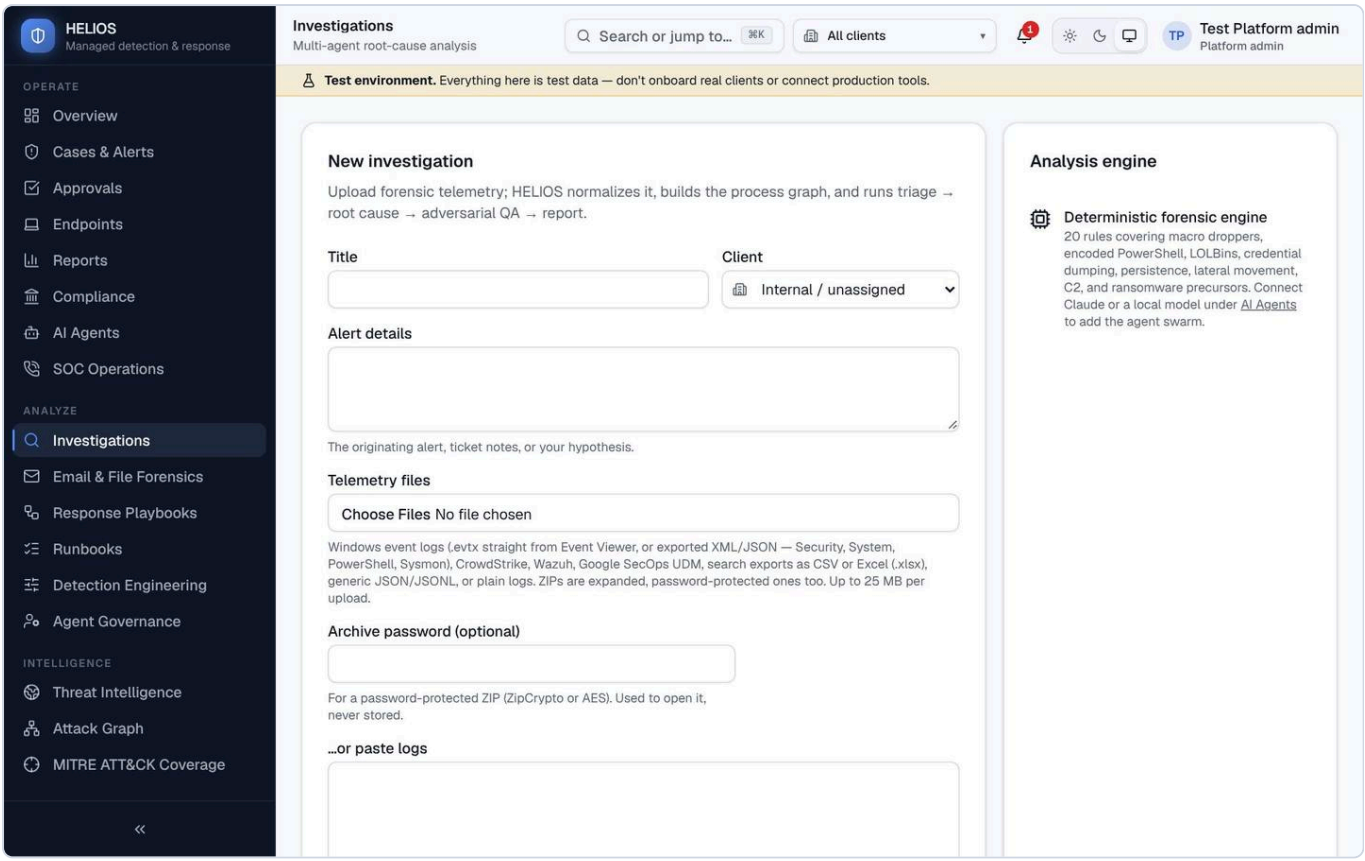


Figure 26. Investigations, as seen by a platform admin (test environment).

How to use it

- New investigation: a title, the client, the alert details (the originating alert, ticket notes or your hypothesis), and the telemetry files: Windows event logs (.evtx, or exported XML/JSON), EDR and SIEM exports, CSV or Excel, JSON/JSONL, plain logs, or a .zip (password-protected ones too, with the archive password, which isn't stored). Up to 25 MB per upload; or paste the logs.
- HELIOS normalizes the telemetry, builds the process graph and runs triage, root cause, adversarial QA and the report. Without a model it uses its deterministic forensic engine (20 rules for droppers, encoded PowerShell, living-off-the-land binaries, credential dumping, persistence, lateral movement, C2 and ransomware precursors).
- Open a past investigation to read its timeline, root cause, the adversarial QA's challenges and the report.

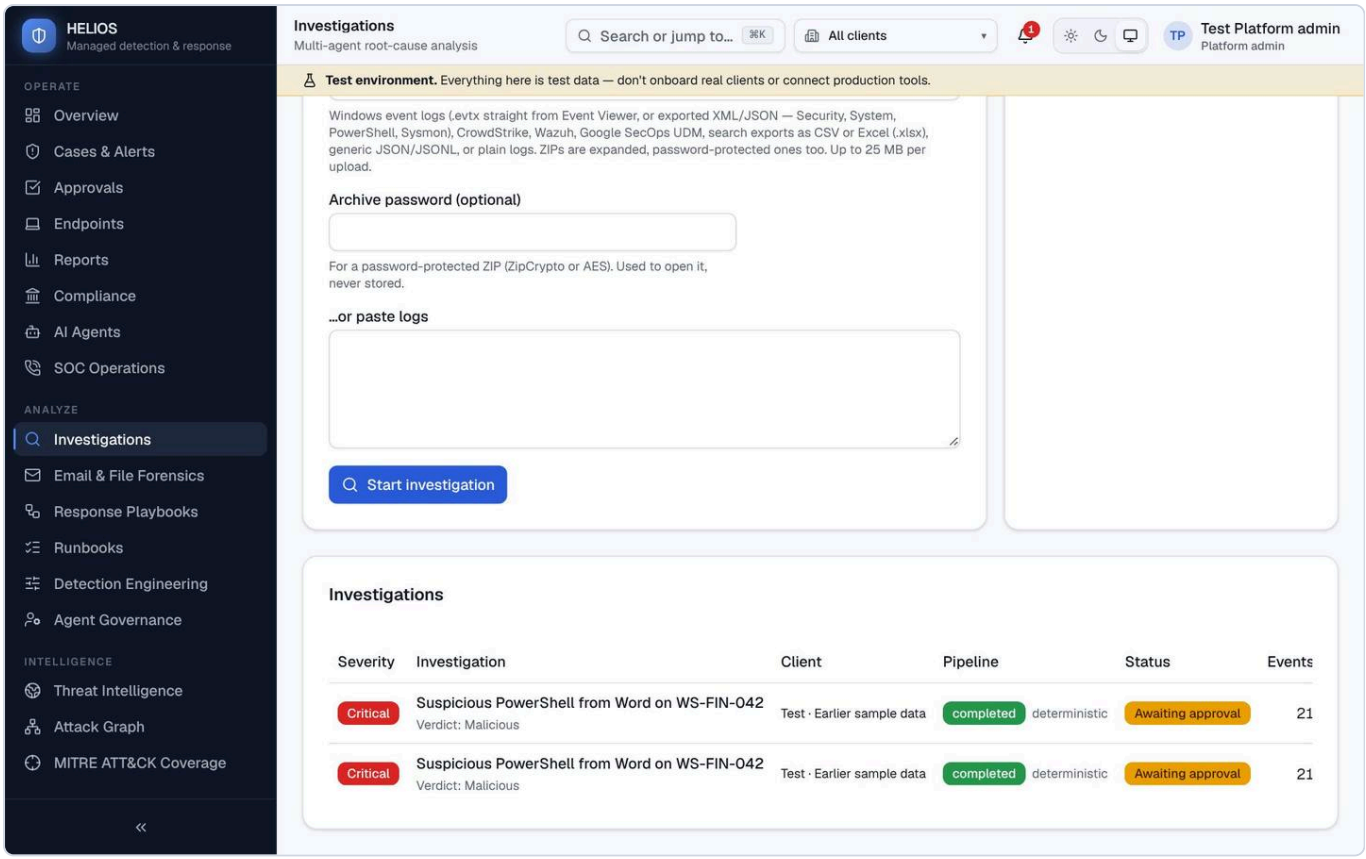


Figure 27. Investigations, further down the page.

Email & File Forensics

Submit a suspicious email or file for analysis: header authentication, links, attachments, macros, archives

and images are examined and scored to a verdict.

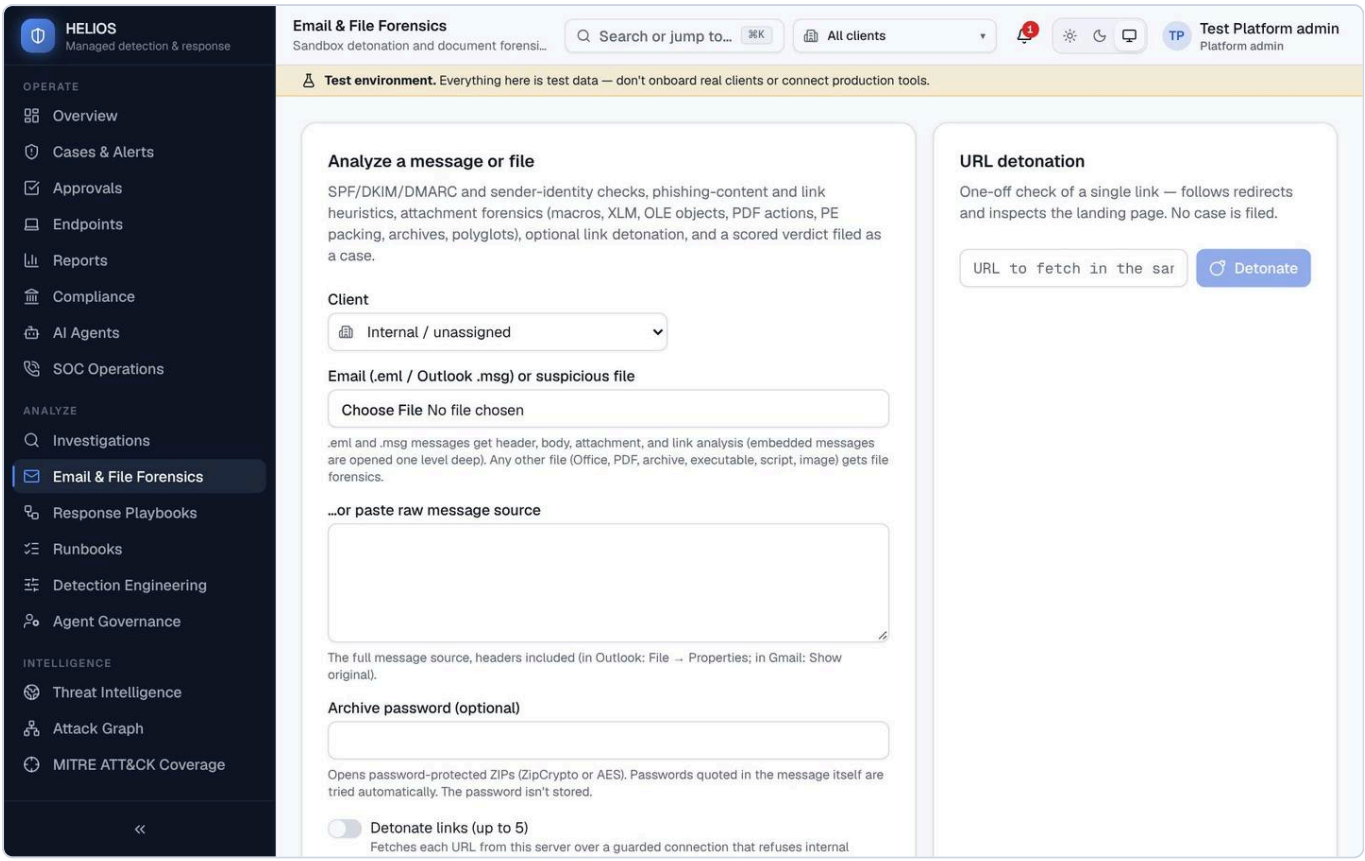


Figure 28. Email & File Forensics, as seen by a platform admin (test environment).

How to use it

- Analyze a message or file: choose the client, then the .eml or Outlook .msg (or any suspicious file), or paste the raw message source. Give an archive password if it's a protected ZIP (it isn't stored), and switch on Detonate links (up to 5) if you want links fetched through the guarded connection.
- HELIOS checks the sender's authentication (SPF, DKIM, DMARC) and identity, phishing content and links, attachments (macros, OLE objects, PDF actions, packing, archives, polyglots) and text and QR codes in images. URL detonation on the right checks a single link without filing a case.
- The verdict comes from the total score across the header, content, attachment, sandbox, intel and sender checks, against the suspicious and malicious thresholds set in Settings. A malicious result opens a case, where the message can be purged from mailboxes and the sender blocked.

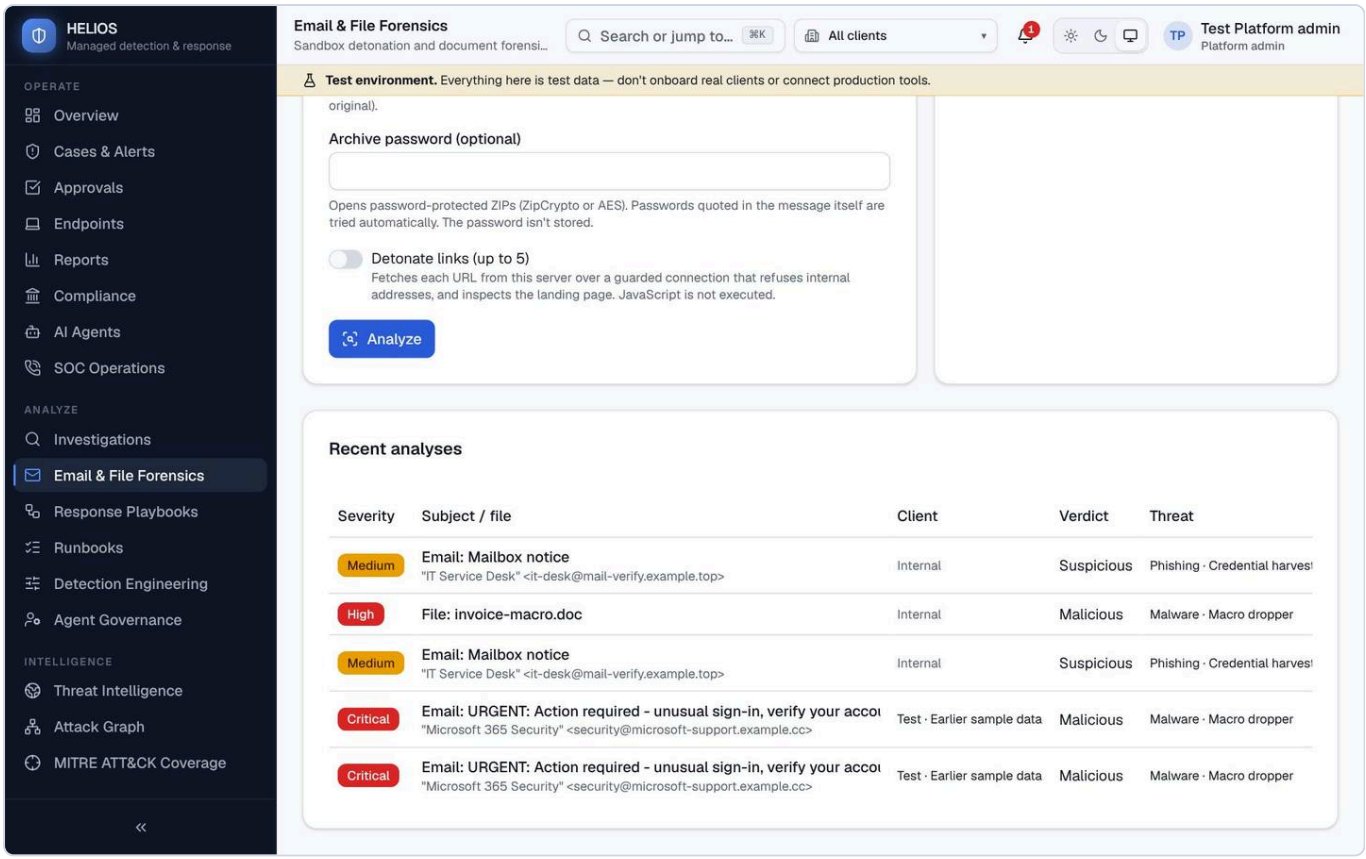


Figure 29. Email & File Forensics, further down the page.

Response Playbooks

Automated response workflows: what triggers them and the steps they run (containment, collecting logs,

notifying, opening tickets).

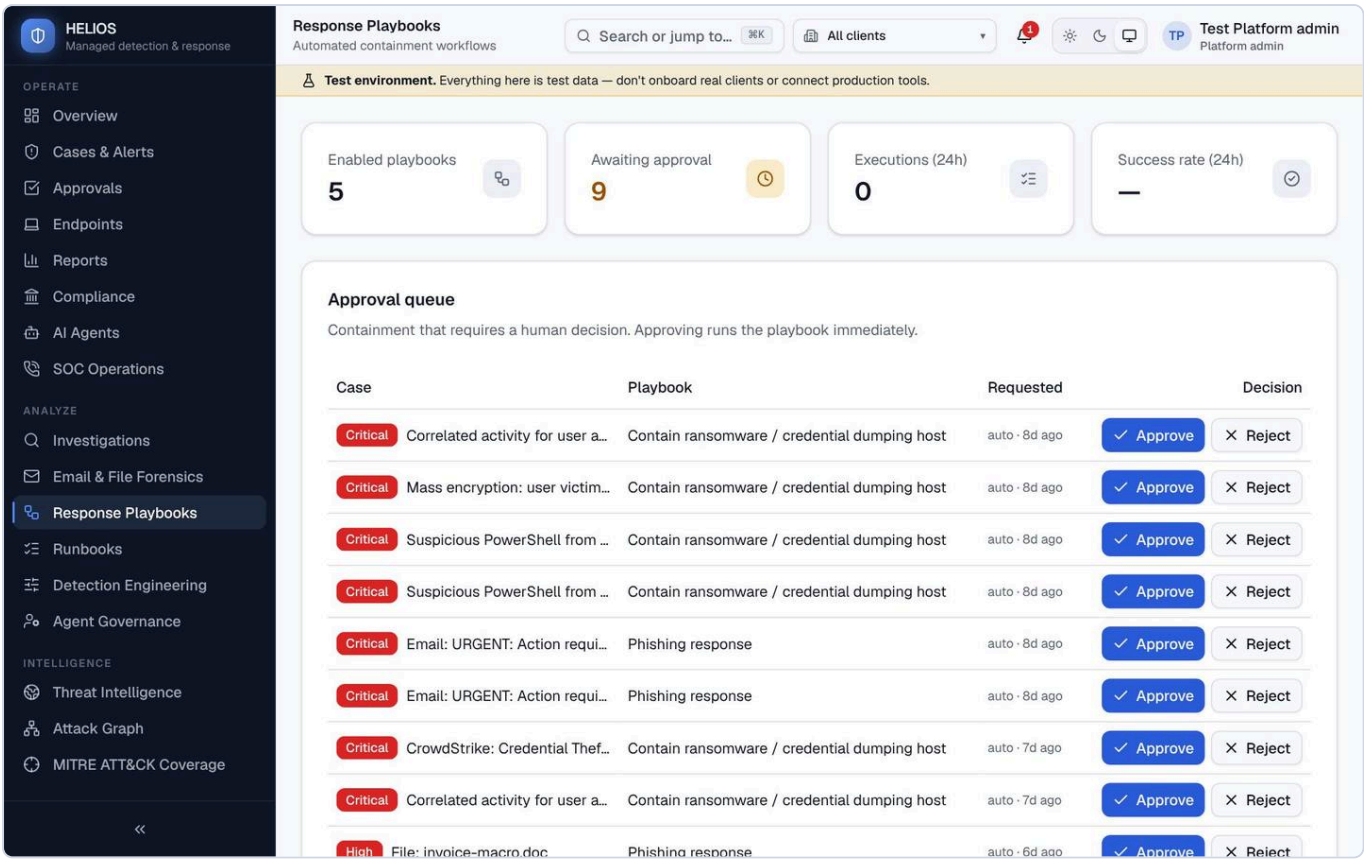


Figure 30. Response Playbooks, as seen by a platform admin (test environment).

How to use it

- Each playbook shows its trigger, its steps and whether it is enabled. Its runs appear on the cases they ran on; steps that need approval wait in Approvals, and steps without a connector are left to do by hand.
- Create or edit a playbook: name, trigger (severity, source, technique), and an ordered list of steps with their parameters. Enable it when it's ready.

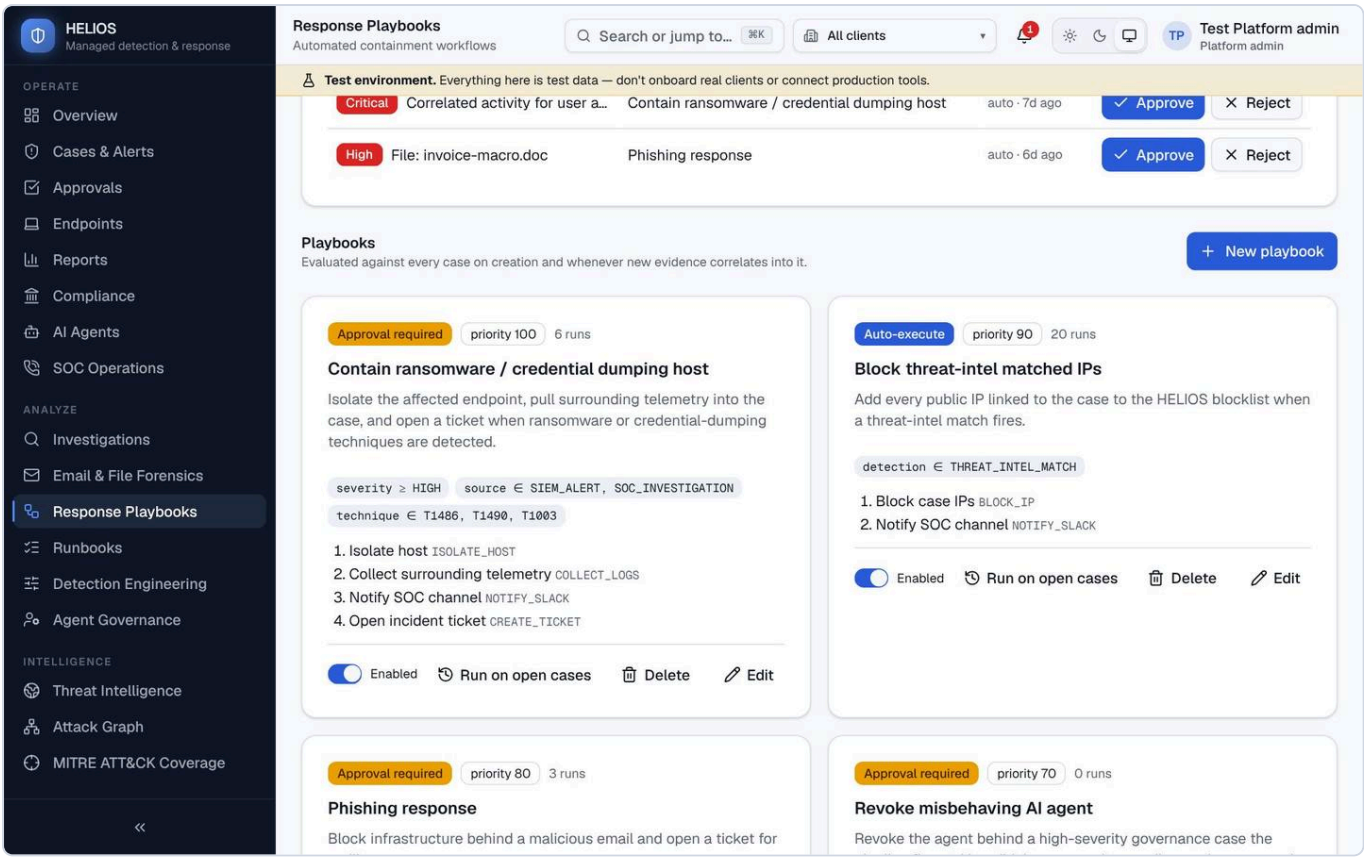


Figure 31. Response Playbooks, further down the page.

Runbooks

Procedures for analysts, attached to the cases they match and worked as checklists on the case.

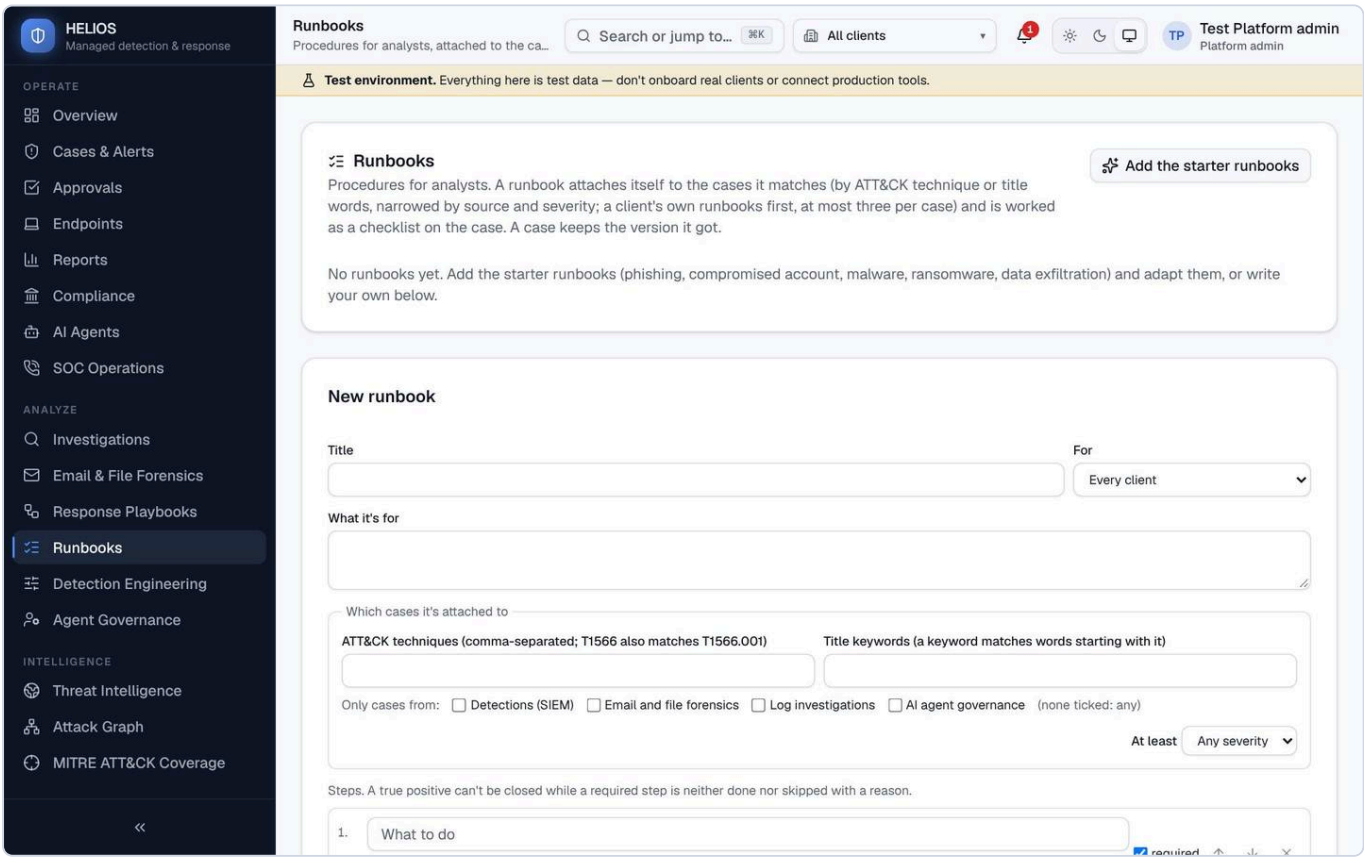


Figure 32. Runbooks, as seen by a platform admin (test environment).

How to use it

- A runbook matches cases by ATT&CK technique or title words, narrowed by source and minimum severity. Up to three matching runbooks attach to a case when it is opened, the client's own first.
- New runbook: a title, which client it's for (or every client), what it's for, the ATT&CK techniques and title keywords it matches, the sources and minimum severity, and its steps (tick required for steps a true positive can't close without). Add the starter runbooks to begin from phishing, compromised account, malware, ransomware and data exfiltration. Editing makes a new version; cases already attached keep the steps they had.

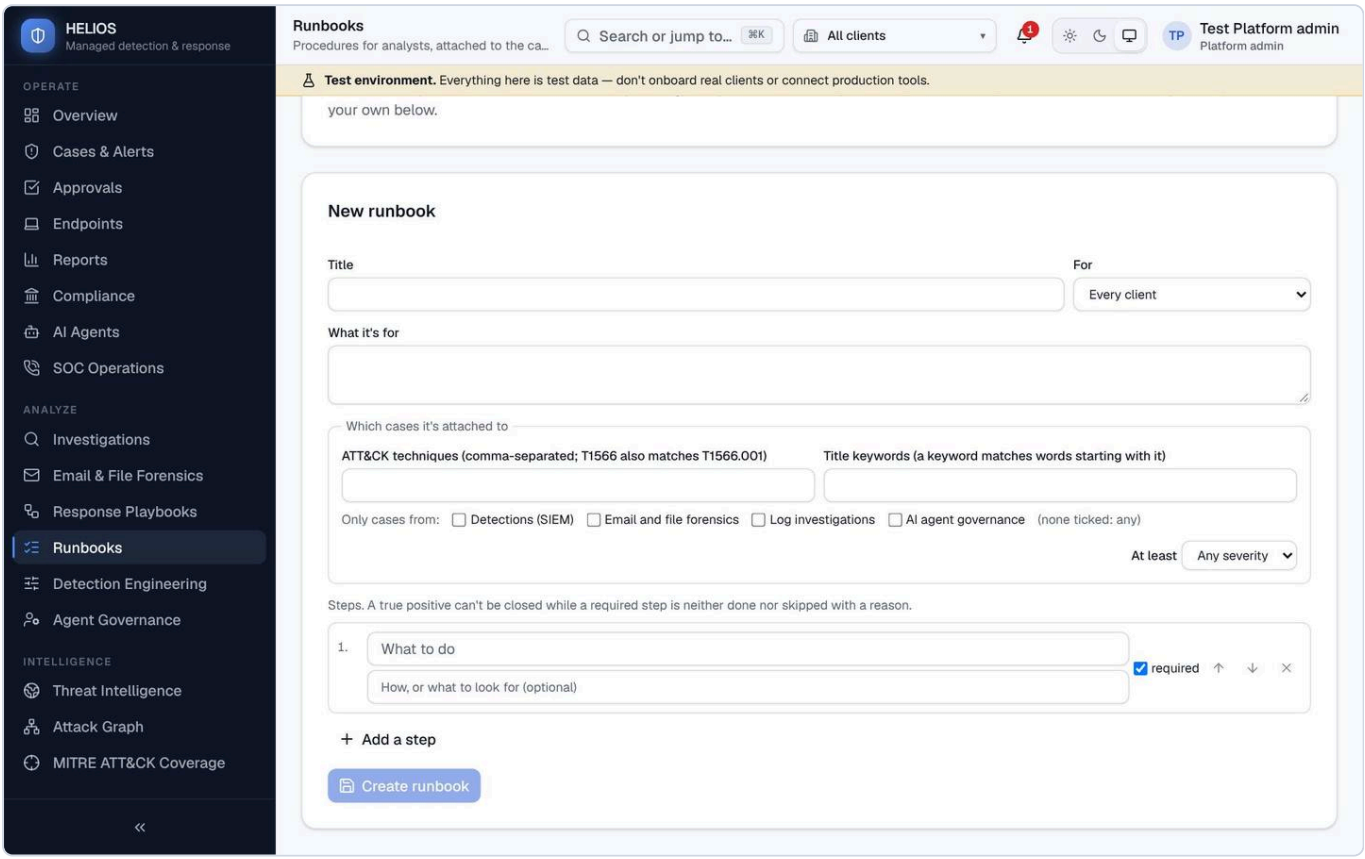


Figure 33. Runbooks, further down the page.

Detection Engineering

The detectors that turn telemetry into cases: the built-in detectors, imported Sigma rules tested before

they go live, and per-client tuning.

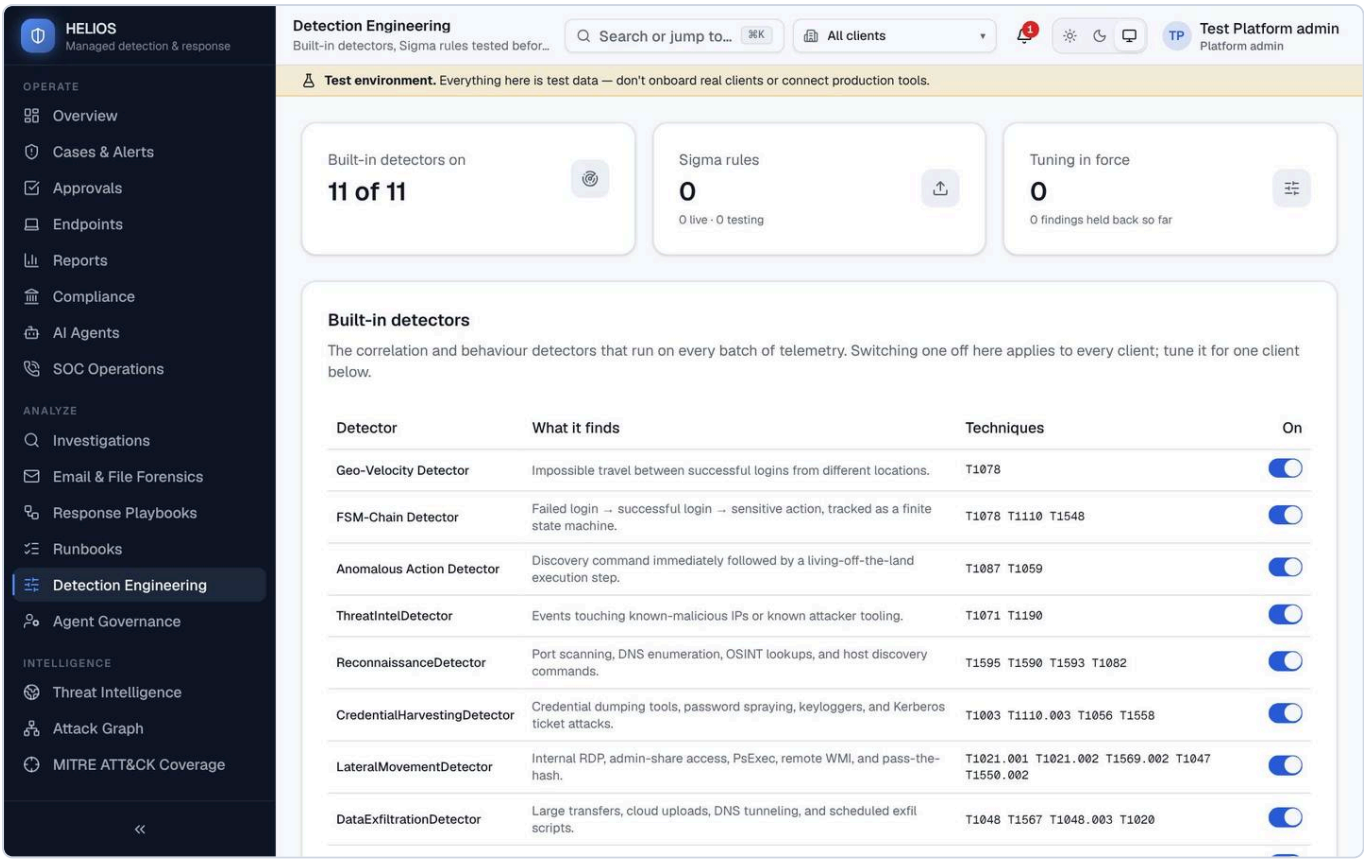


Figure 34. Detection Engineering, as seen by a platform admin (test environment).

How to use it

- Built-in detectors: switch one off for every client here (tune for one client below instead where you can).
- Import a Sigma rule: it is compiled into a matcher (anything it can't evaluate faithfully is refused with the reason), backtested on recent history, run in testing (matches counted, nothing opened) and then made live.
- Tune per client: switch a detector or rule off there, leave out known-good hosts, accounts or processes, or change the severity, each with a reason and an optional expiry. The count shows what tuning held back.

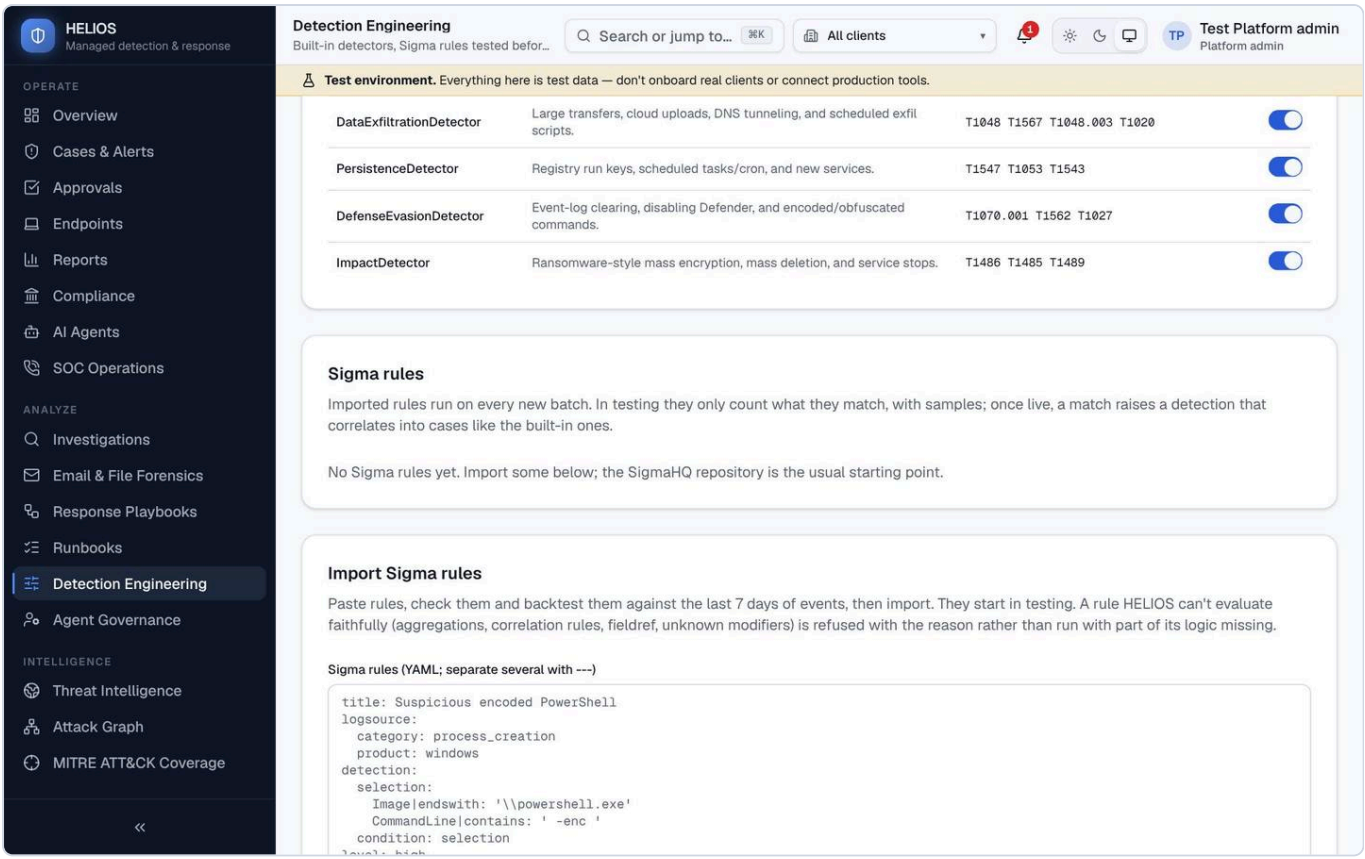


Figure 35. Detection Engineering, further down the page.

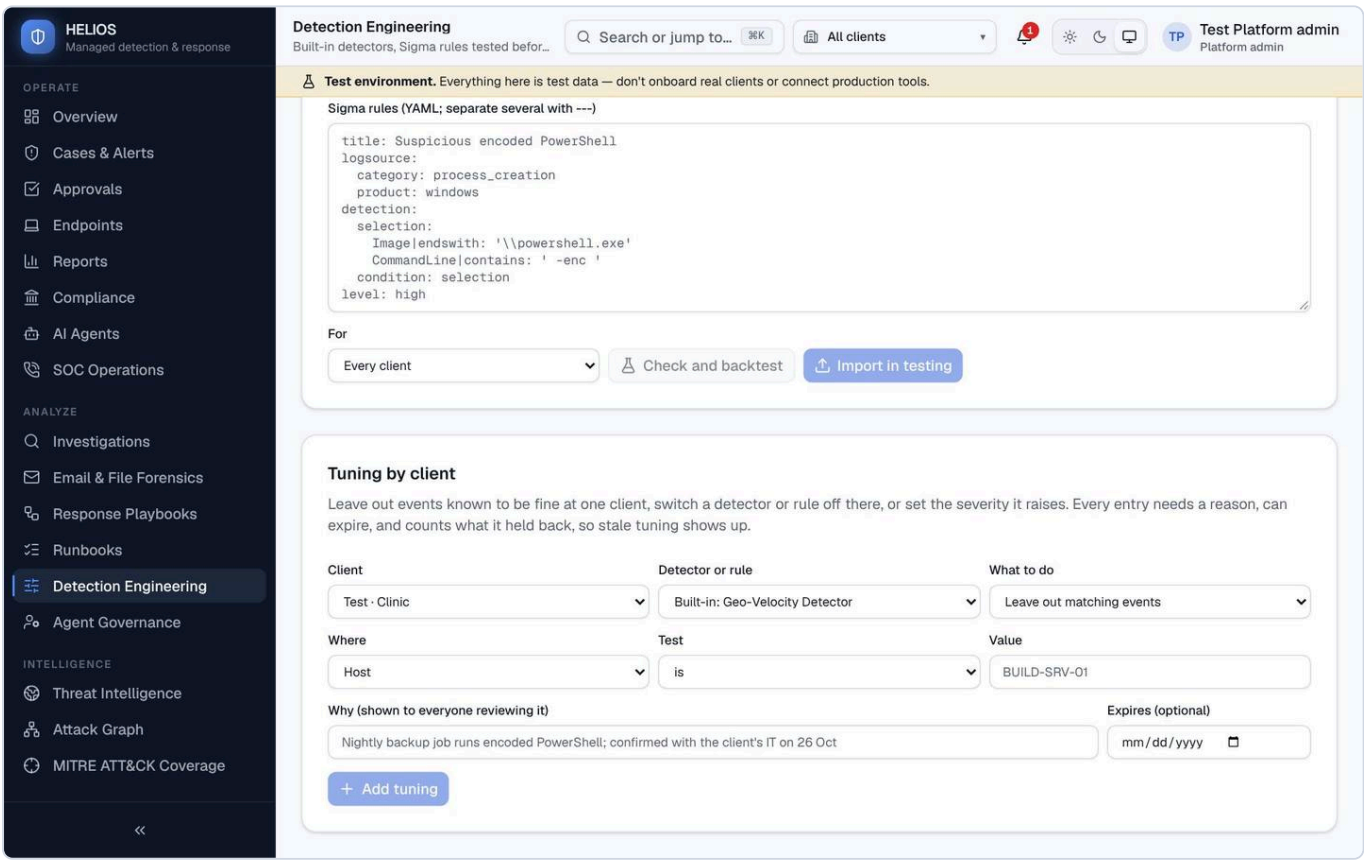


Figure 36. Detection Engineering, further down the page.

Agent Governance

Govern the AI agents your clients run: identities and tokens, the policies they act under, their trust, and revocation.

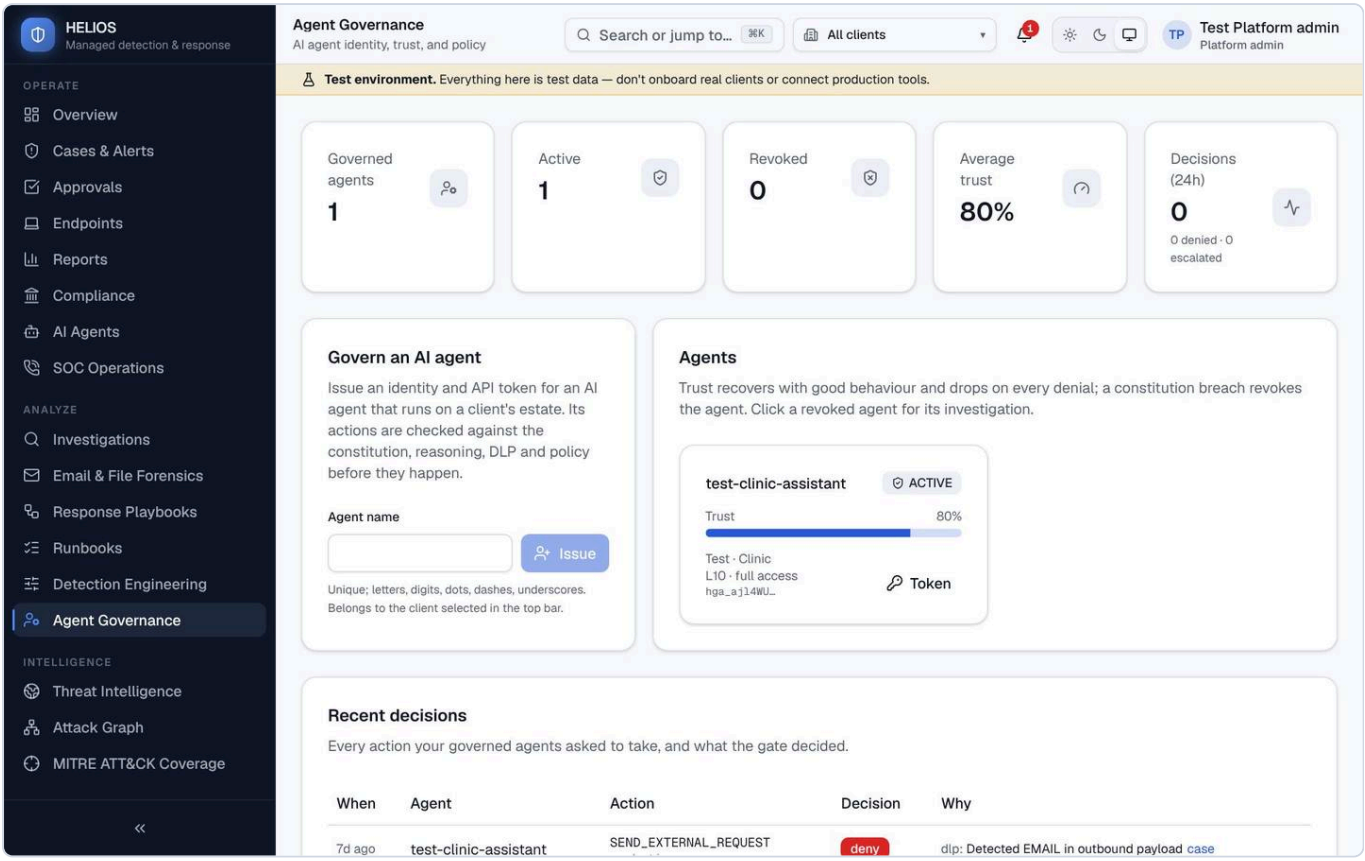


Figure 37. Agent Governance, as seen by a platform admin (test environment).

How to use it

- Govern an AI agent: enter the agent's name and Issue, for the client selected in the top bar. Its token is shown once. The agent asks HELIOS before each action; constitution, reasoning, data-loss and policy checks decide, and Recent decisions shows each request and why it was allowed or denied.
- Trust recovers with good behaviour and drops on every denial.
- A breach revokes the agent and opens a case. Revoke or restore an agent by hand here or from its case.

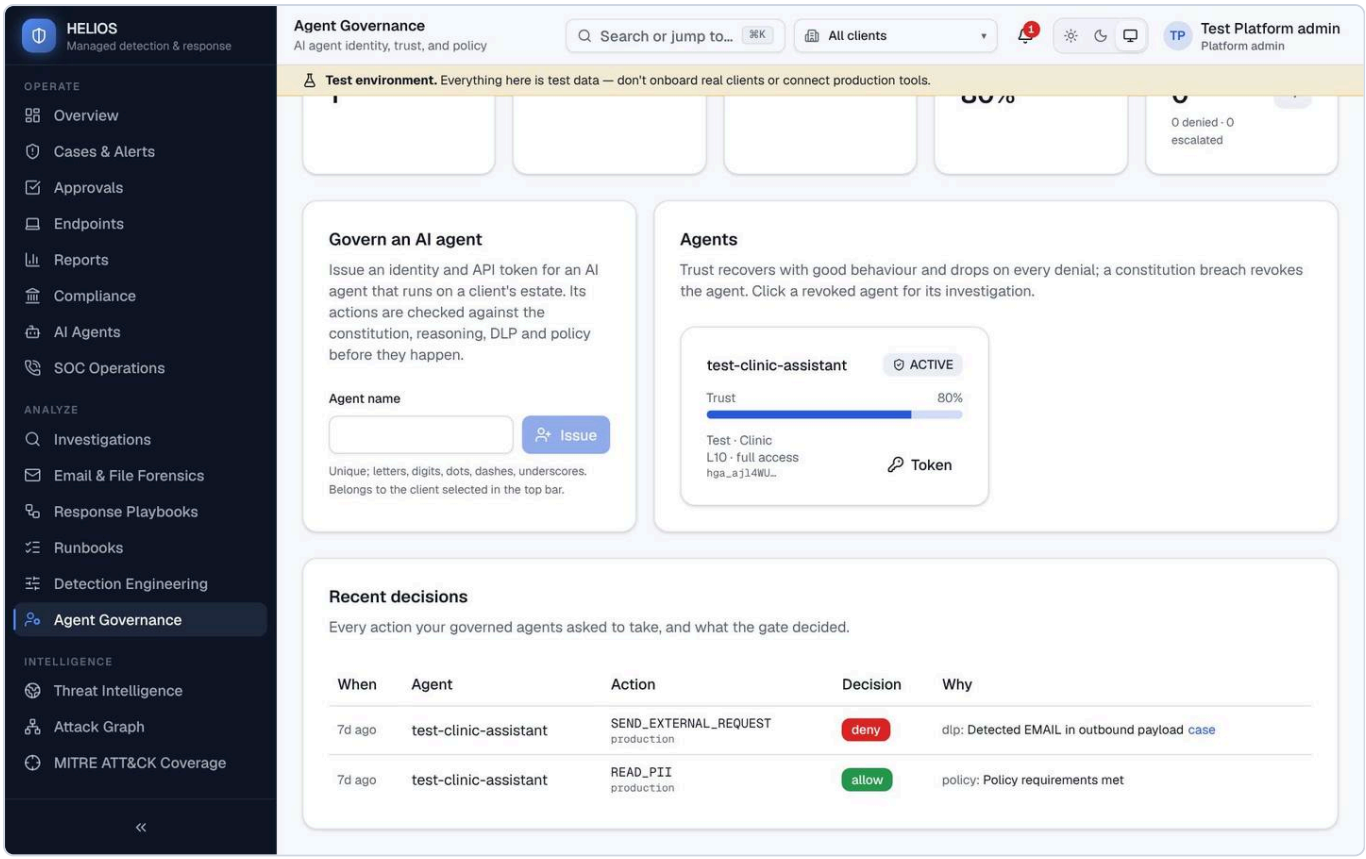


Figure 38. Agent Governance, further down the page.

Threat Intelligence

Look up indicators (IP, domain, URL, file hash, email address) against the SOC's indicator store, the

configured reputation providers and the threat feeds.

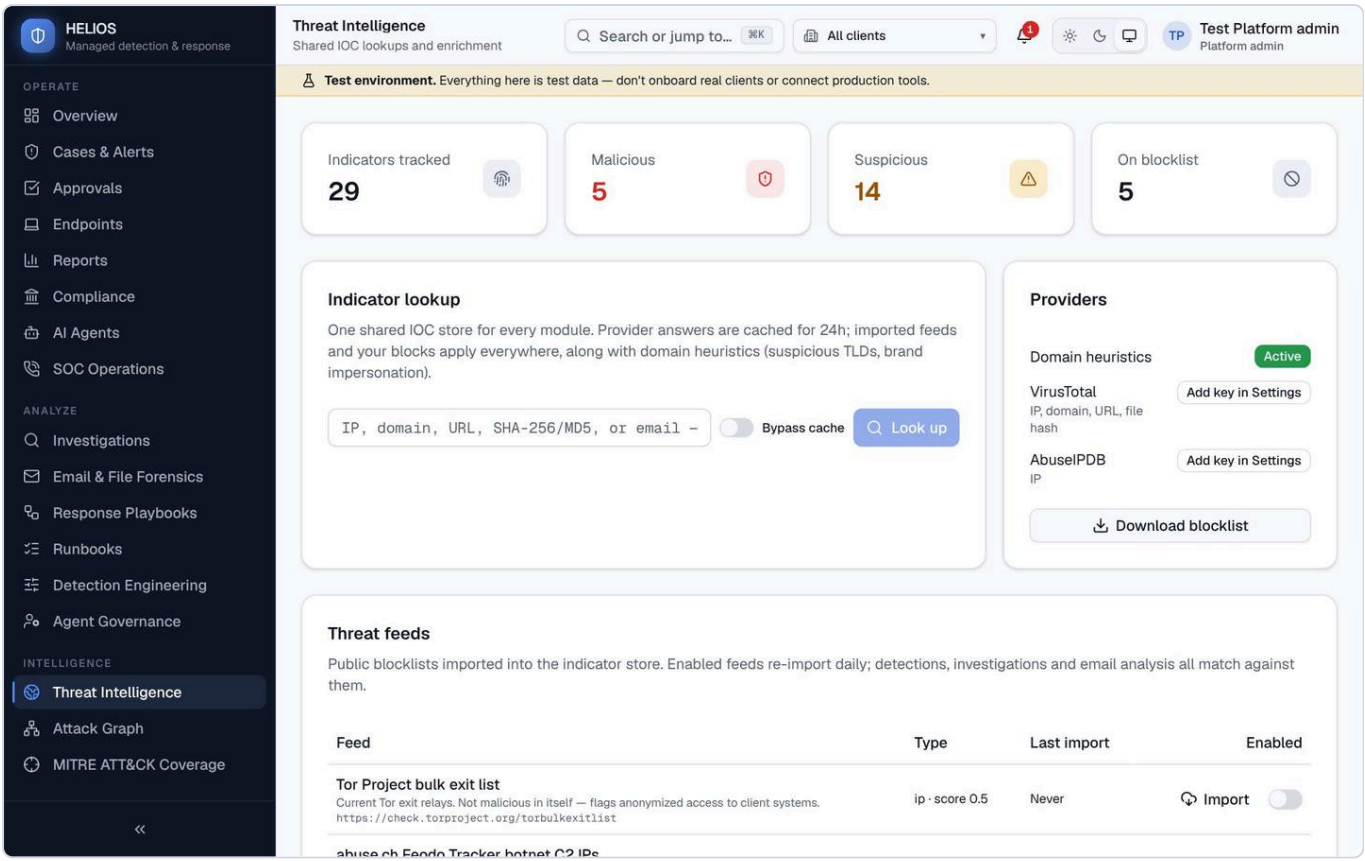


Figure 39. Threat Intelligence, as seen by a platform admin (test environment).

How to use it

- Paste an indicator (defanged forms like hxxp and [.] are fine) and Look up. The result shows the verdict, score, sources, why, and for an IP where it is (country, city and network, from the server's GeoIP database). Results are cached for 24 hours; Bypass cache forces a fresh lookup.
- Block or unblock an indicator: the blocklist feeds detections and firewalls that pull it. Enable and import threat feeds.

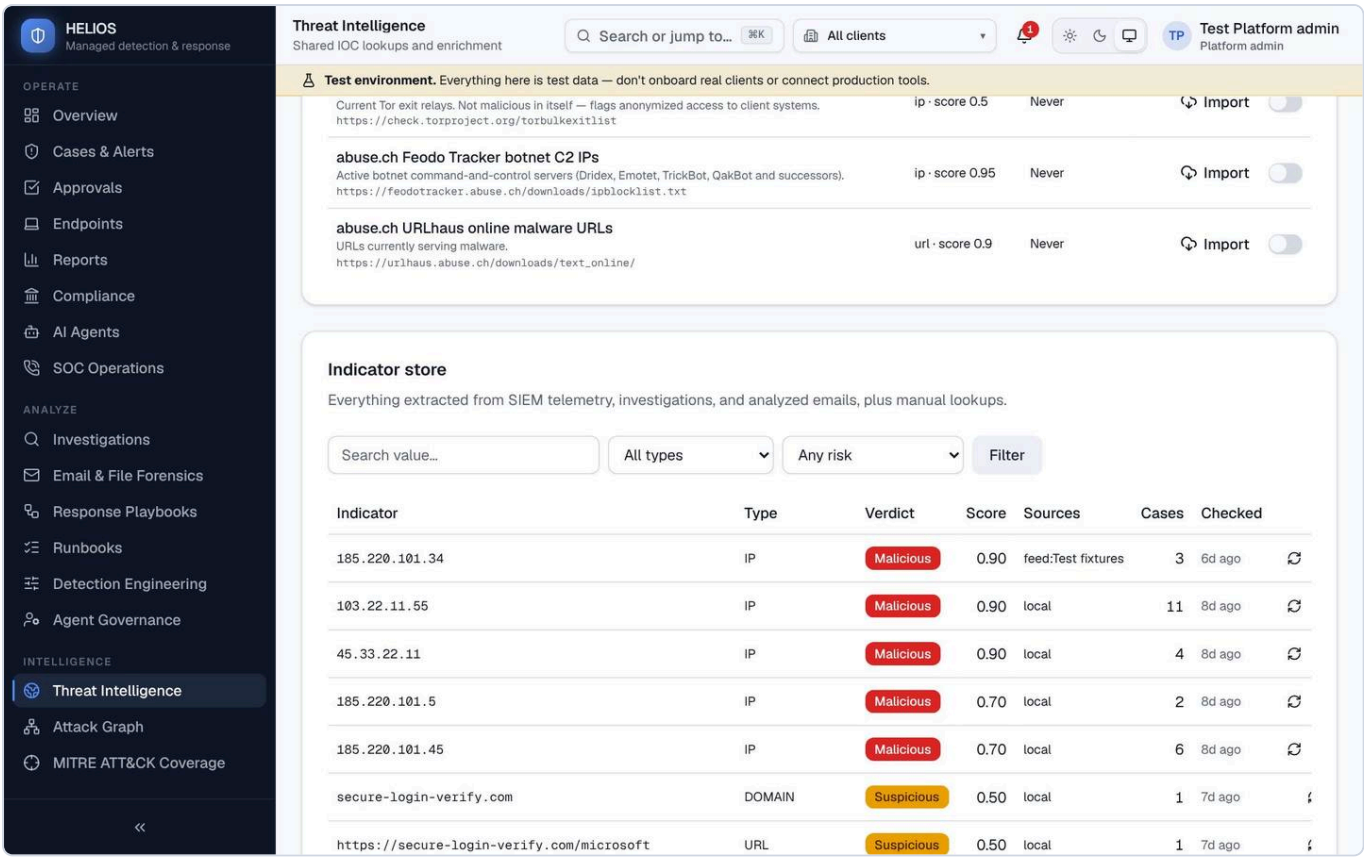


Figure 40. Threat Intelligence, further down the page.

Attack Graph

Cases joined by the hosts, accounts, addresses, indicators and agents they share. An entity is an attack

chain when its cases together reach more ATT&CK tactics than any one alone.

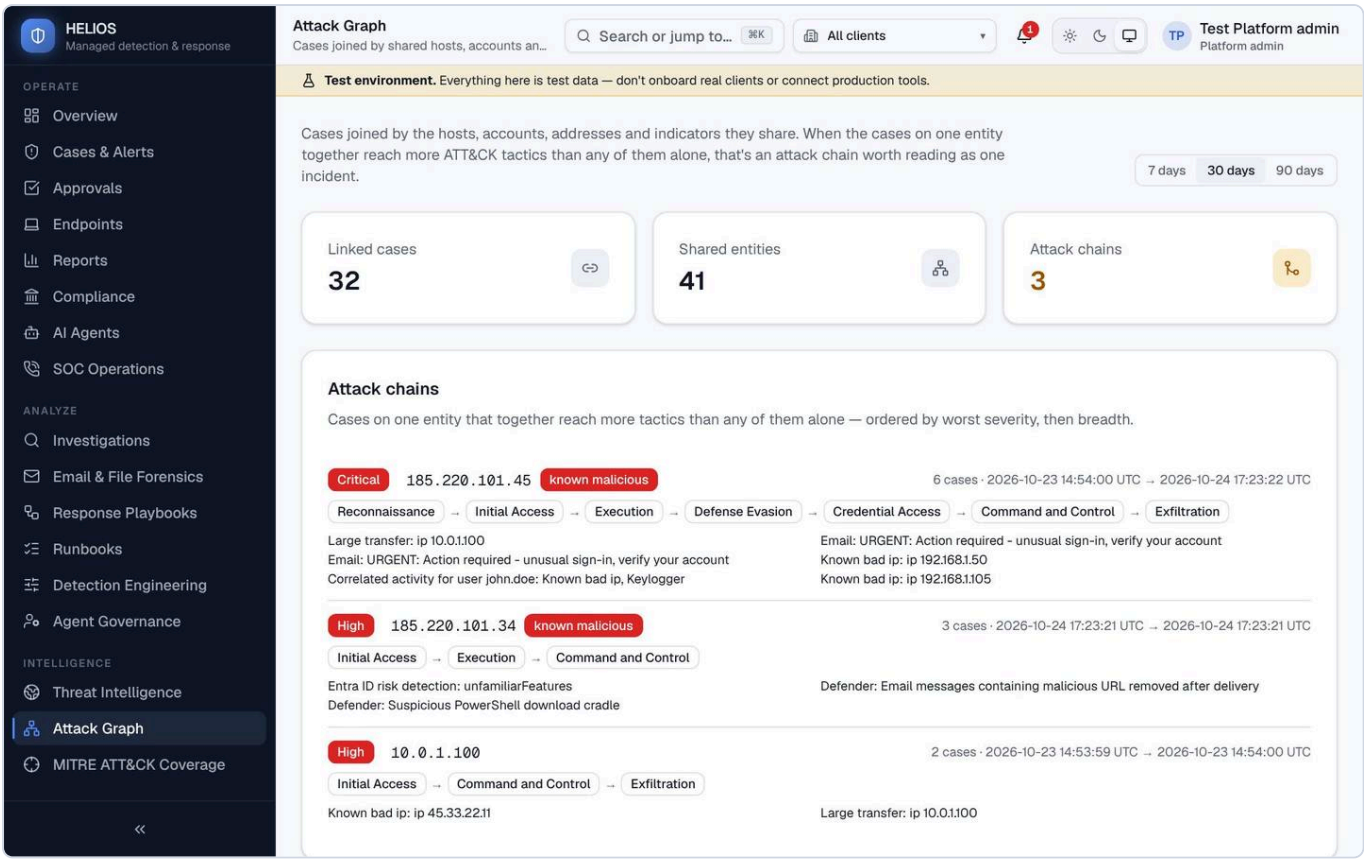


Figure 41. Attack Graph, as seen by a platform admin (test environment).

How to use it

- Hover or click a node to see the cases behind it; follow chains across tactics to see an intrusion that several smaller cases make up.

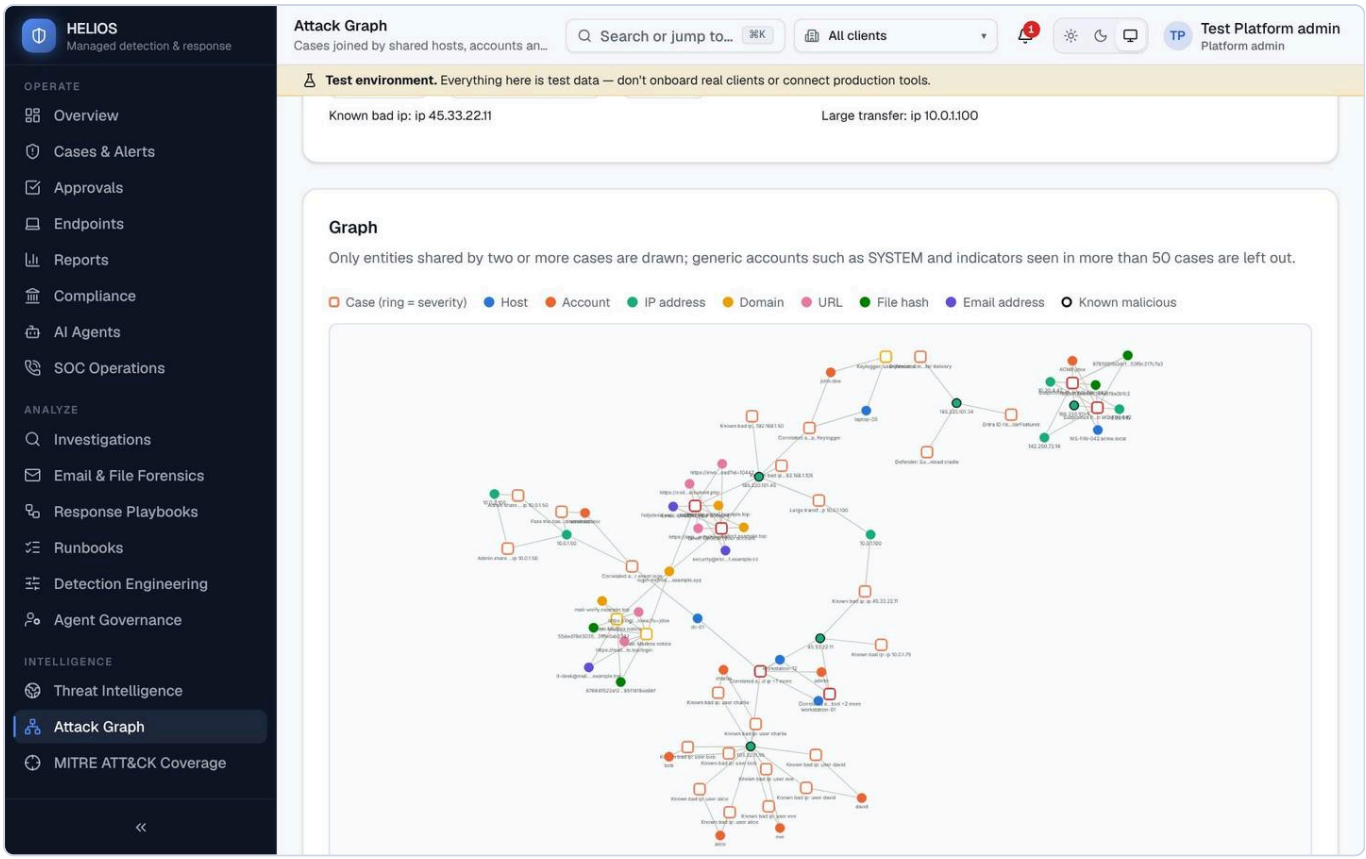


Figure 42. Attack Graph, further down the page.

MITRE ATT&CK Coverage

Which ATT&CK techniques HELIOS can detect, across built-in detectors, Sigma rules and onboarded

tools, and which it has seen.

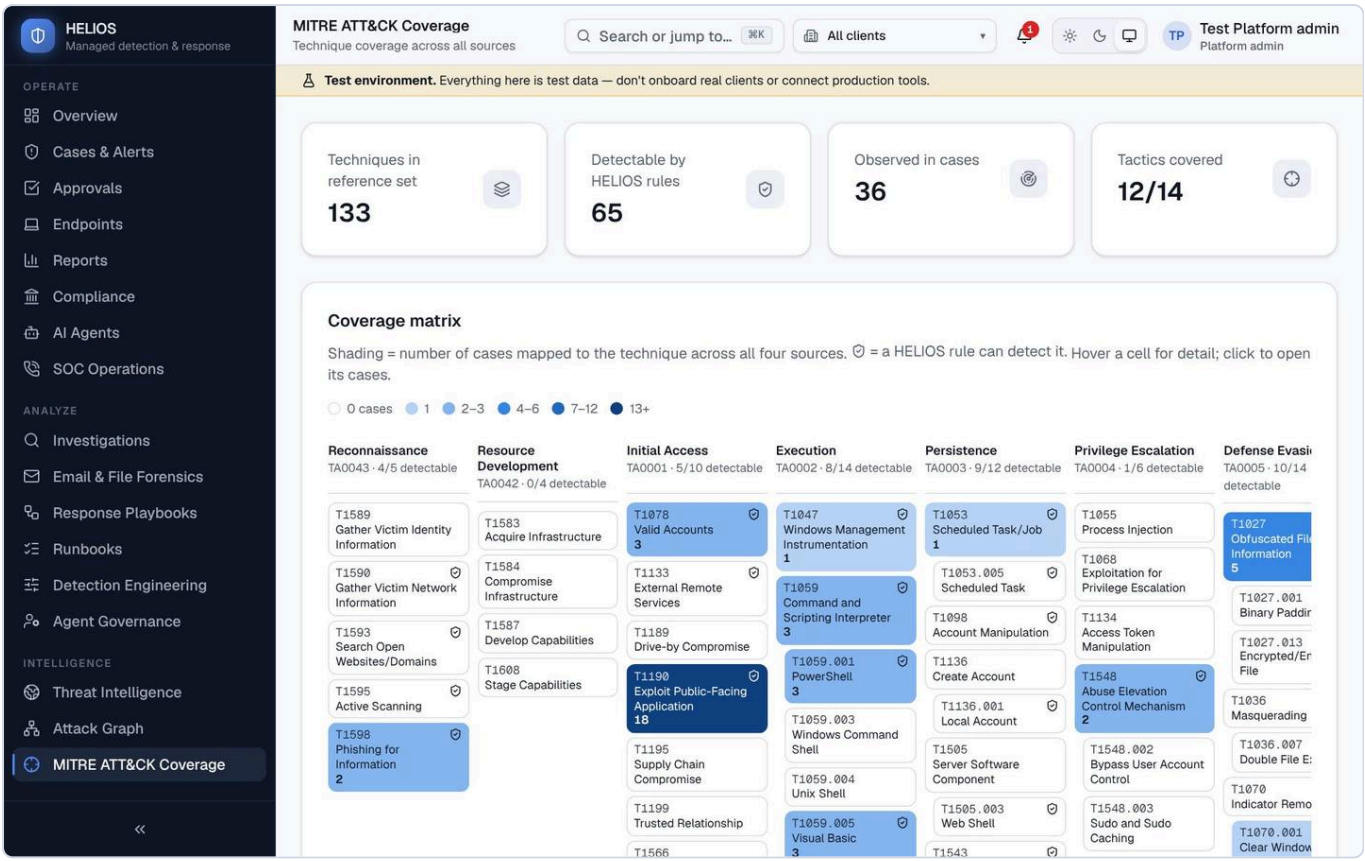


Figure 43. MITRE ATT&CK Coverage, as seen by a platform admin (test environment).

How to use it

- Read coverage by tactic and technique; gaps show where detection is missing.
- Switch built-in detectors on and off from here too.

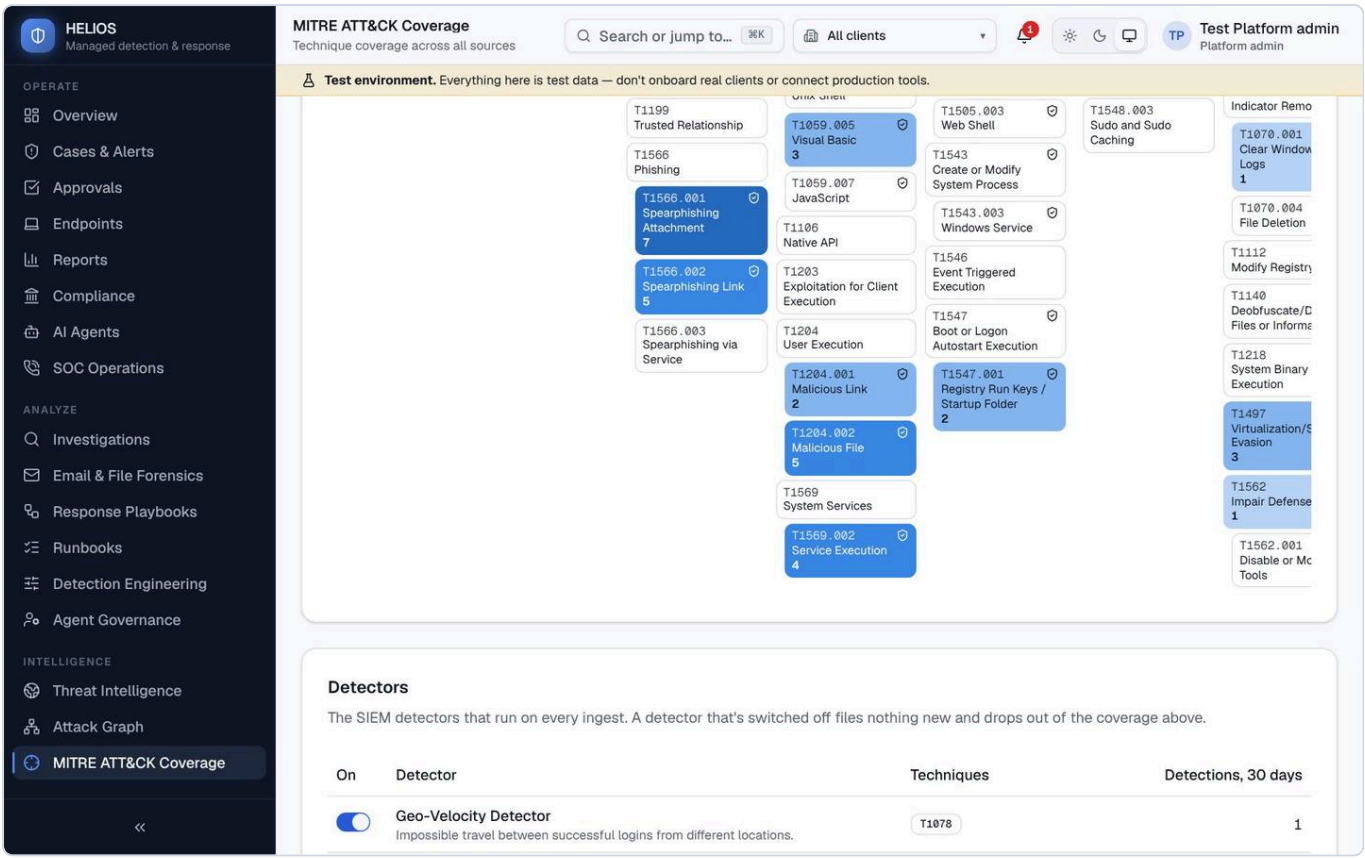


Figure 44. MITRE ATT&CK Coverage, further down the page.

Tools & Integrations

The security tools onboarded for each client: how their logs reach HELIOS and how HELIOS acts through

them.

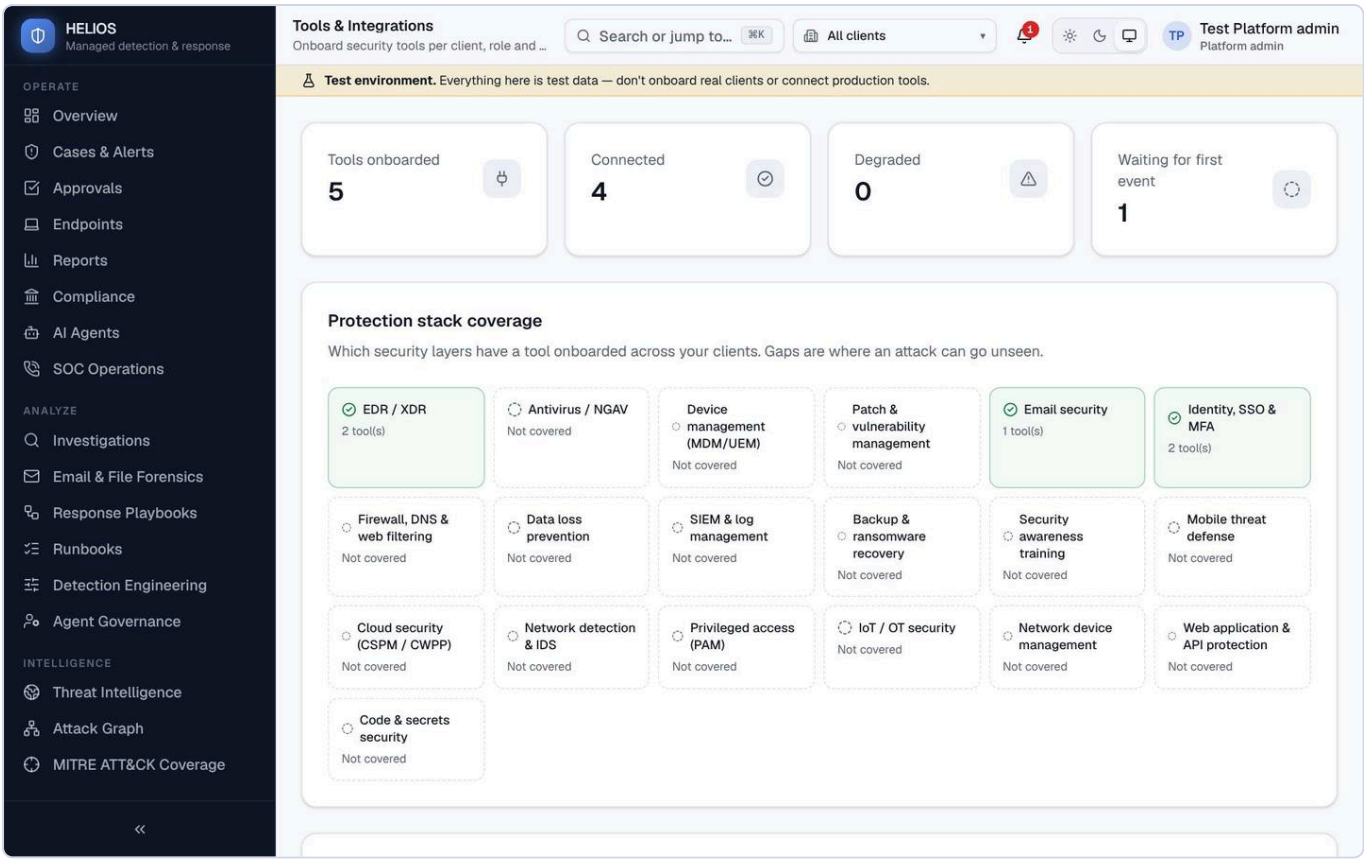


Figure 45. Tools & Integrations, as seen by a platform admin (test environment).

How to use it

- The tiles count tools onboarded, connected, degraded and waiting for their first event. Protection stack coverage shows which security layers have a tool across your clients; gaps are where an attack can go unseen.
- Each tool shows how it sends logs (push to its ingest URL, HELIOS pulling from the vendor's API, or the HELIOS forwarder), its source health, and whether a response connector can act through it.
- Onboard a tool for a client: pick it from the catalog, enter its read-only credentials (encrypted with the client's own key), test the connection, set the collectors running, and turn on the response connector (dry-run first). Generate the forwarder configuration for syslog, Windows event logs and files. Replay records waiting in the dead-letter queue.
- A source silent longer than its limit opens a case: attackers often stop logging first. A client can't go live while a source is dead.

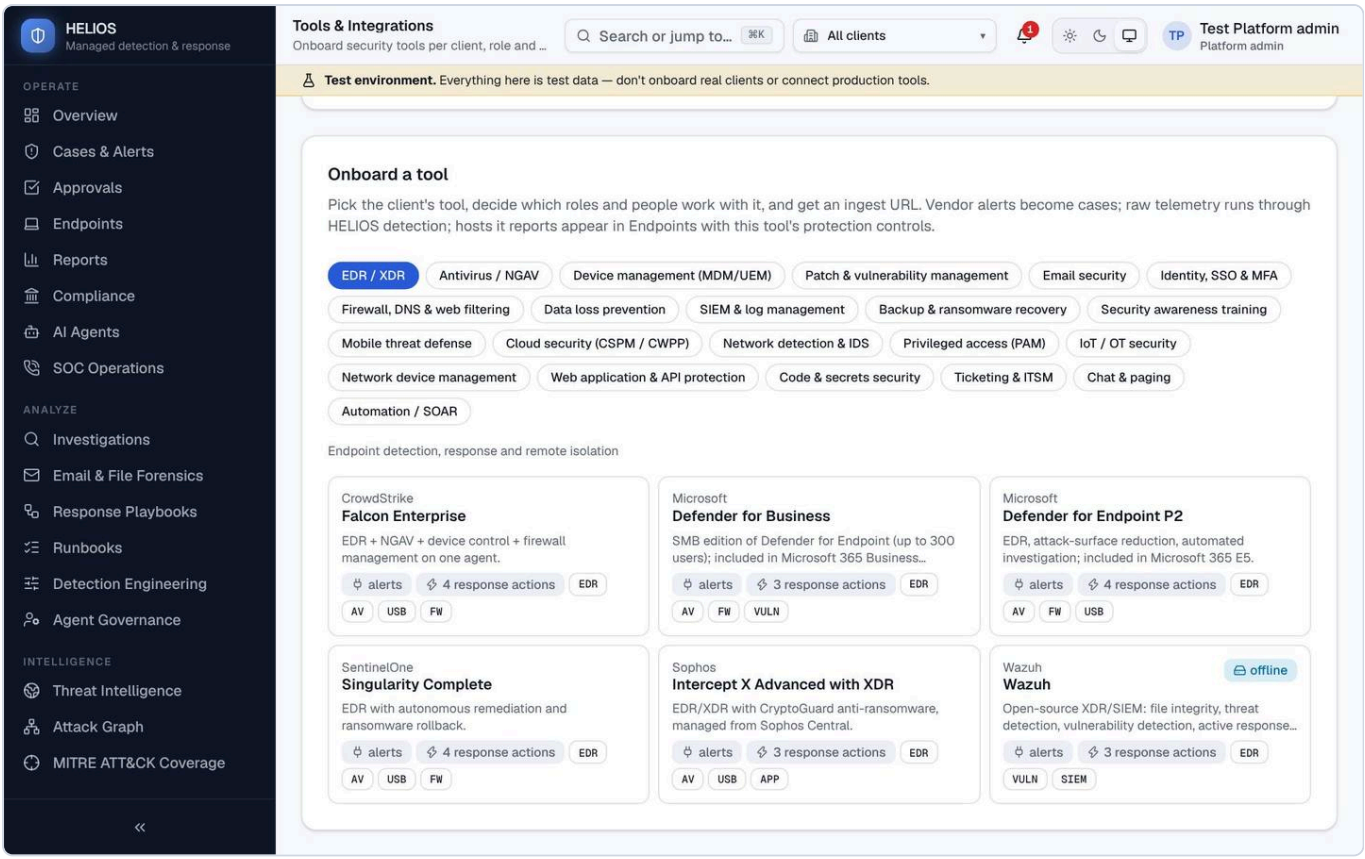


Figure 46. Tools & Integrations, further down the page.

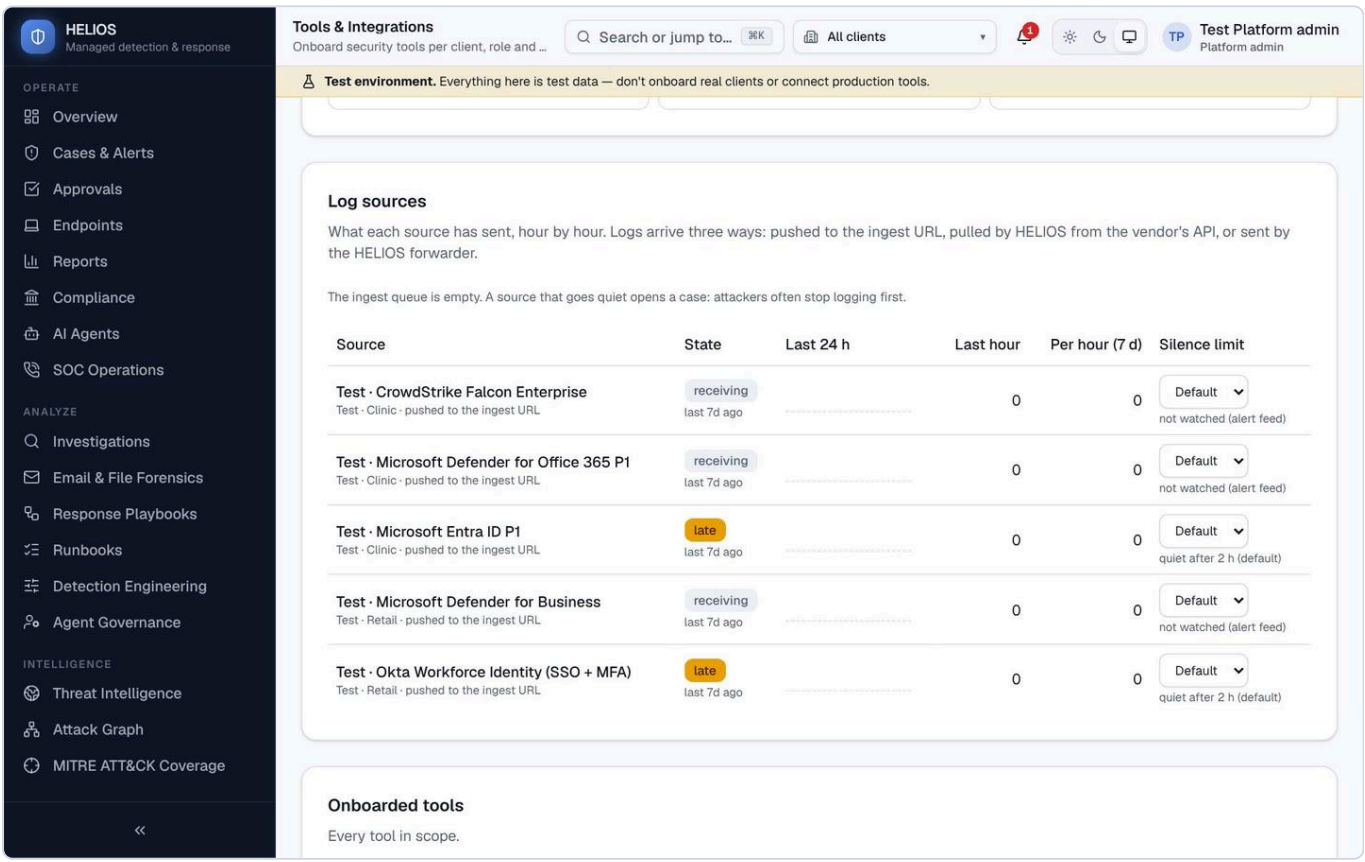


Figure 47. Tools & Integrations, further down the page.

Rulebooks

Each client's signed rules for what HELIOS may do on its estate. The rules only ever restrict: protected assets and accounts, which actions run automatically, need approval or never happen, blast-radius limits, business hours and change freezes, the approver chain, AI data handling and CERT-In contacts.

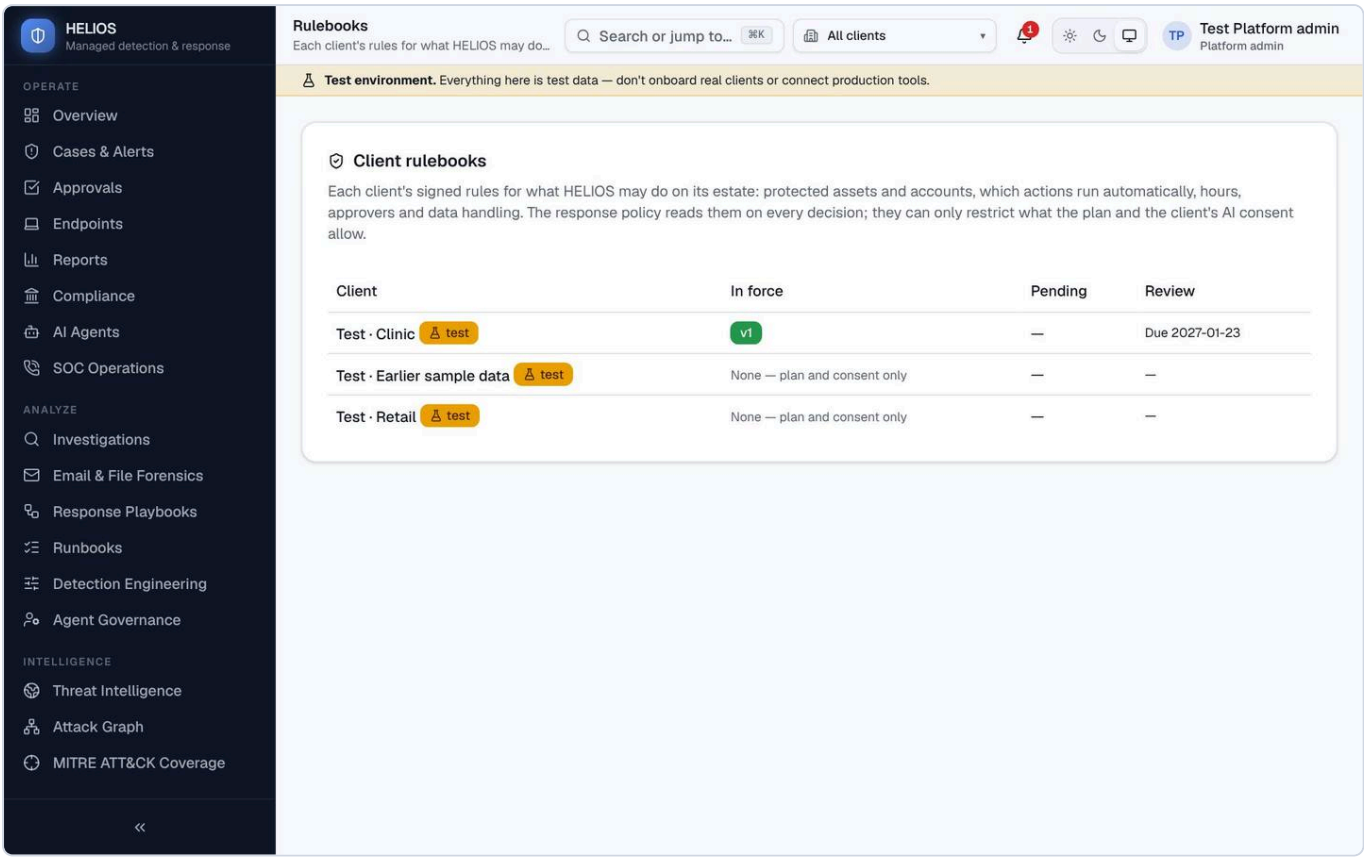


Figure 48. Rulebooks, as seen by a platform admin (test environment).

How to use it

- Pick the client at the top. The page shows the version in force, its hash and when it is due for review.
- The tabs: In force (the rules now applying), Edit, Dry run, Allowlist, Versions, AI models, Autonomy and Policy as code.
- Edit drafts a change, from an industry template or the version in force. Dry run replays the last 30 days of AI proposals under the draft and probes each action. Versions compares any two versions side by side.
- Sign the draft to put it in force. Client admins sign their own; staff activate only on the client's recorded instruction, naming it.
- Allowlist: known-good hosts, IPs and accounts, each with an owner, a reason and an expiry. They open no cases, and AI containment against them is only suggested.
- Policy as code: export or import the rulebook as YAML, or keep it in Git (GitHub, GitLab or an https URL): a changed file becomes a draft that still needs signing here.
- AI models: which models may read the client's evidence (any, region-locked or on-premises only), the client's own providers (bring your own key), a provider per agent, and the monthly AI budget.
- Autonomy: the AI's autonomy ladder per action type for this client, and promotions waiting for approval.

Clients

The client organisations (tenants): their plan and SLA, status, contacts, and how far they consent to AI acting on their estate.

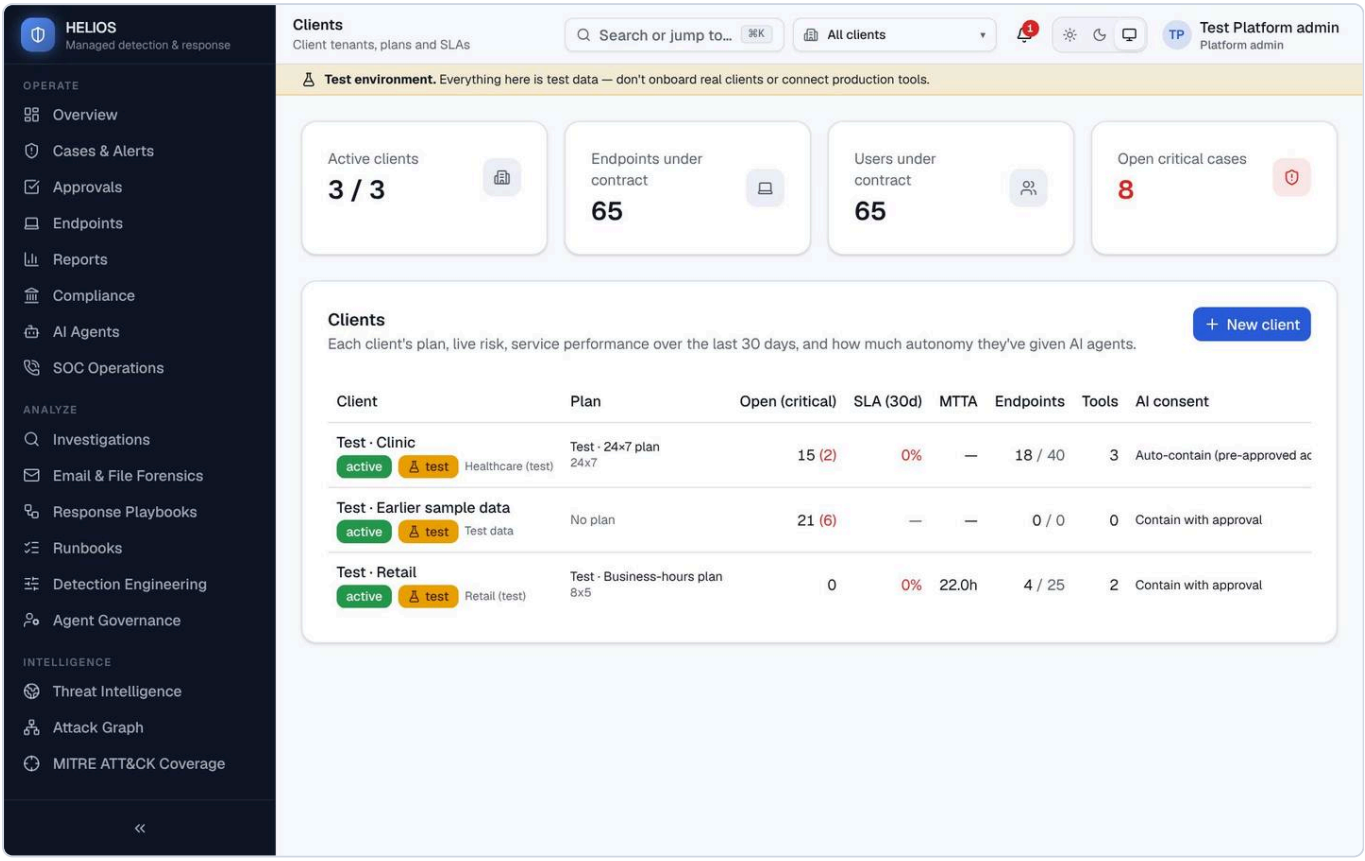


Figure 49. Clients, as seen by a platform admin (test environment).

How to use it

- Each client shows its plan, status (onboarding, active, suspended), licensed endpoints and users, and test clients are badged.
- Add or edit a client: name, industry, contacts, plan, licensed counts, and the AI autonomy cap the client has consented to (it caps every agent on their estate). A client moves from onboarding to active only once every data source has sent something in the last day.

Team & Roles

The people who use HELIOS, their roles, the security domains (queues) analysts work, and what each role

can see and do.

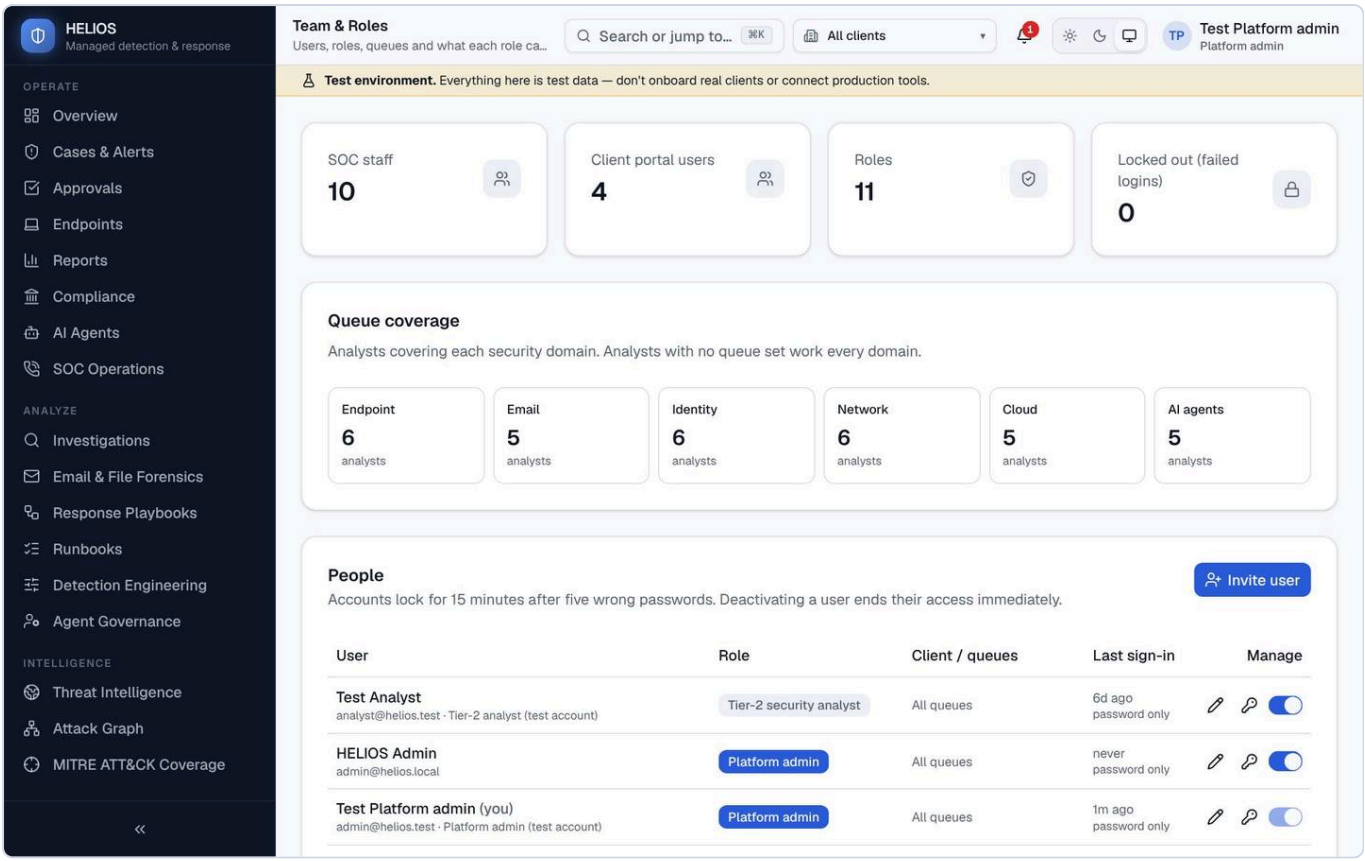


Figure 50. Team & Roles, as seen by a platform admin (test environment).

How to use it

- Invite user: name, email, role, and for client users their organisation. HELIOS shows a temporary password once; pass it to the person securely. Limit an analyst to security domains to give them those queues only.
- On each person's row: the pencil edits the role, organisation and queues; the key resets the password; the switch deactivates the account. Changes apply on the person's next request. Reset two-step sign-in for someone who lost their phone.
- Queue coverage shows how many analysts cover each security domain. Accounts lock for 15 minutes after five wrong passwords; the Locked out tile counts them.
- The role matrix shows every permission each role has.

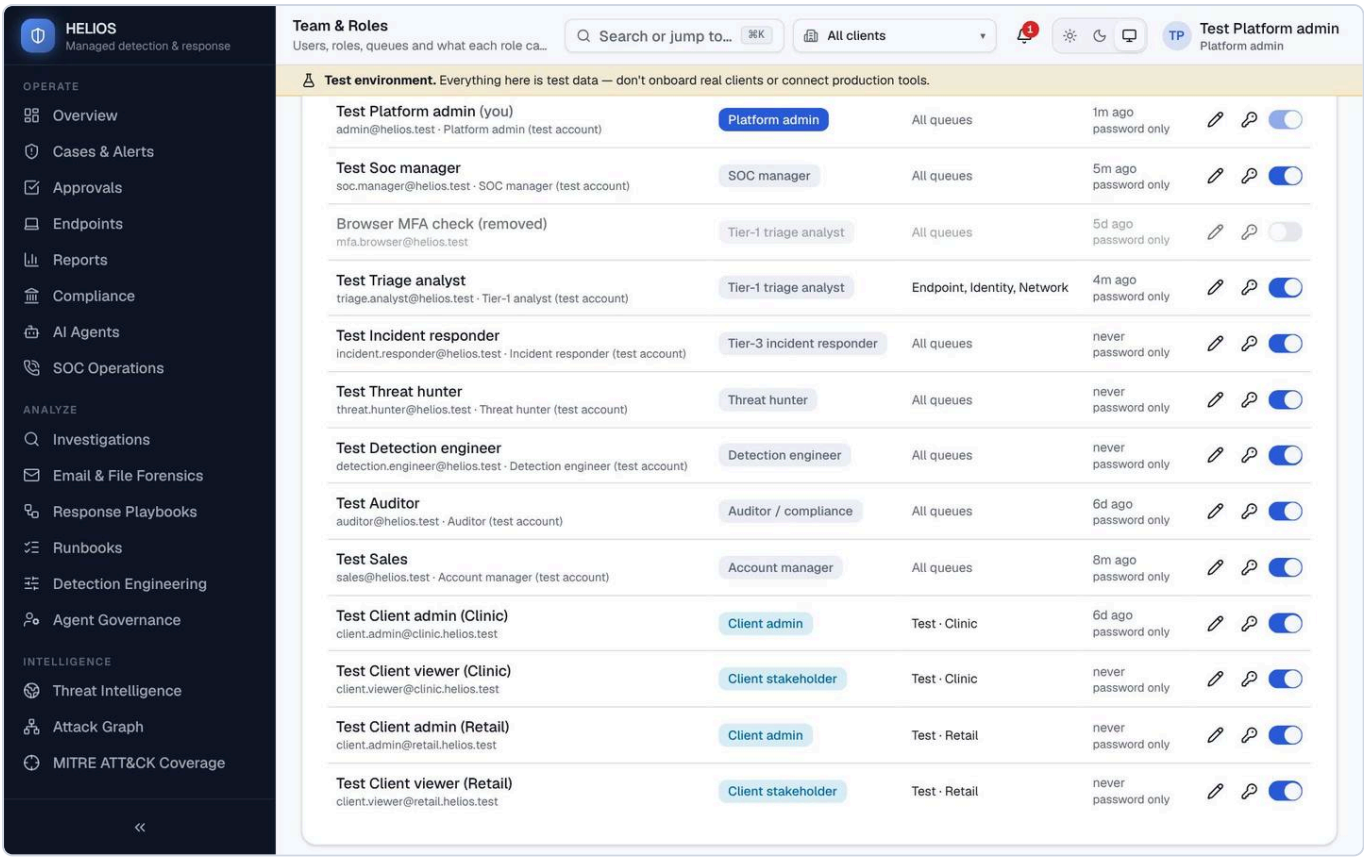


Figure 51. Team & Roles, further down the page.

Data Retention

Where each client's logs are kept and for how long, the encrypted archive, and legal holds.

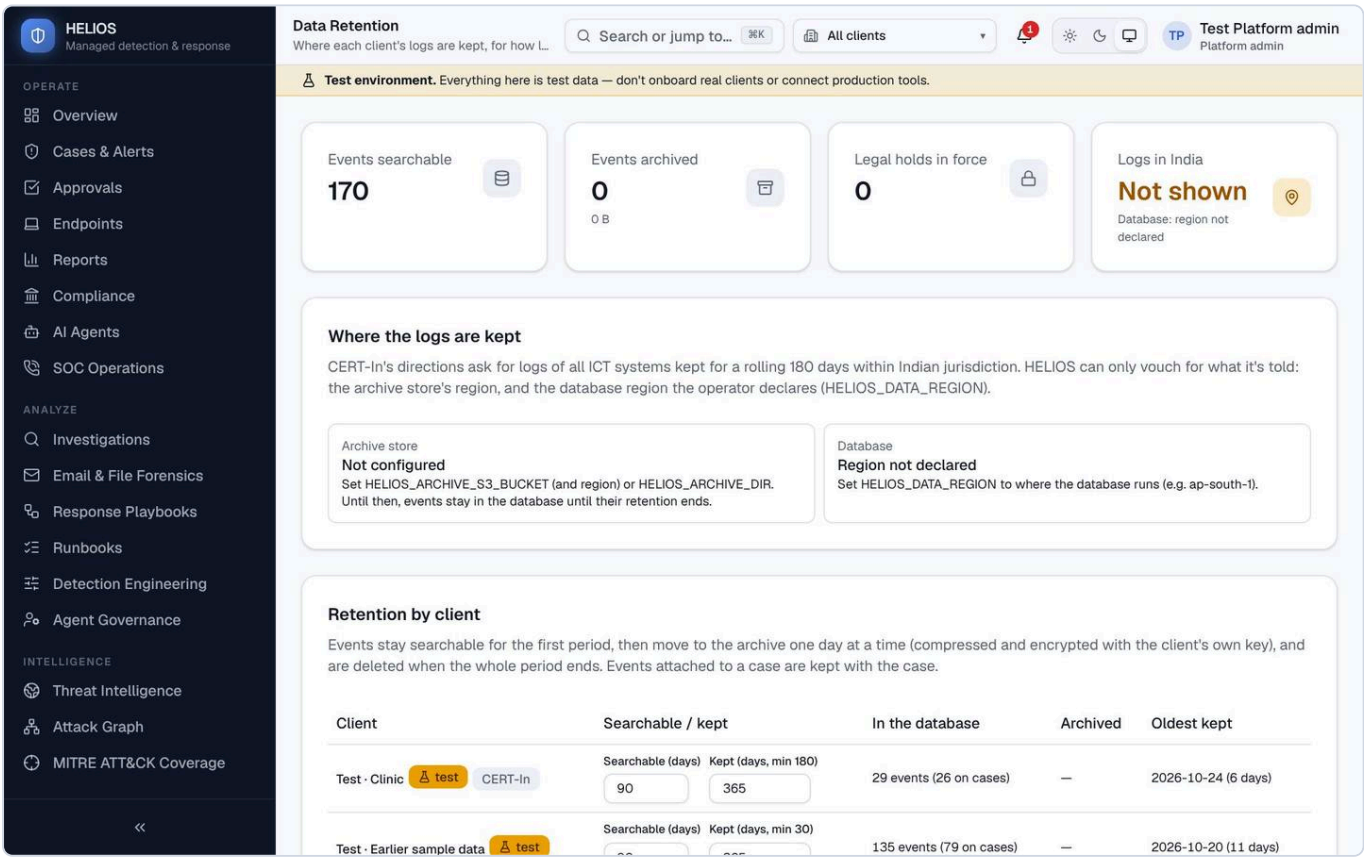


Figure 52. Data Retention, as seen by a platform admin (test environment).

How to use it

- Retention by client: Searchable (days) and Kept (days). Where CERT-In's directions apply, Kept can't go below 180. Events attached to a case are kept with the case.
- The page says whether the logs are shown to be in India (the database region and the archive store's region), from what the server is configured with.
- Legal holds stop deletion over a date range (optionally tied to a case); placing and releasing need a reason. Archives download decrypted, and each download is audited.

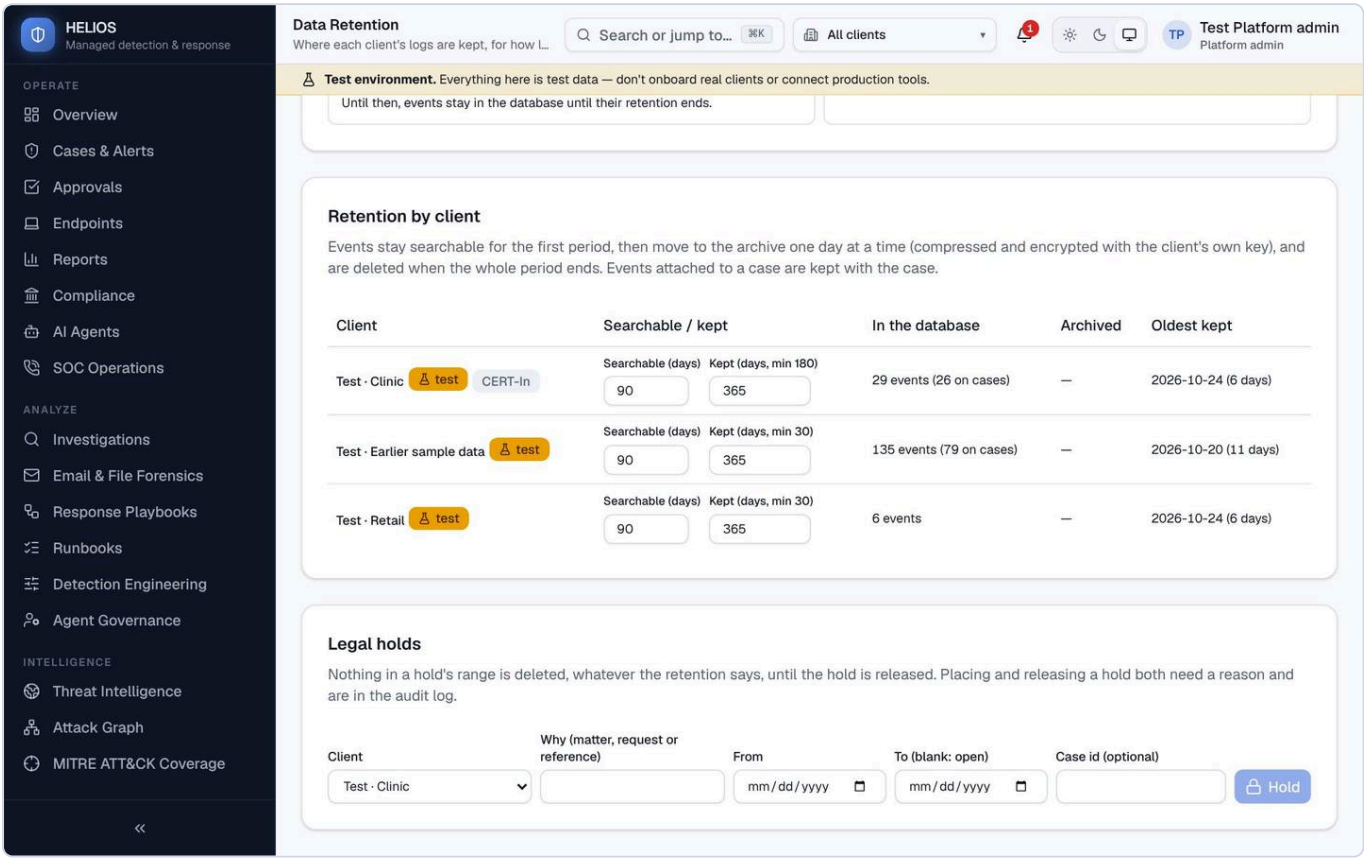


Figure 53. Data Retention, further down the page.

Audit Log

The hash-chained record of every action in HELIOS: who did what, when, to which case or object.

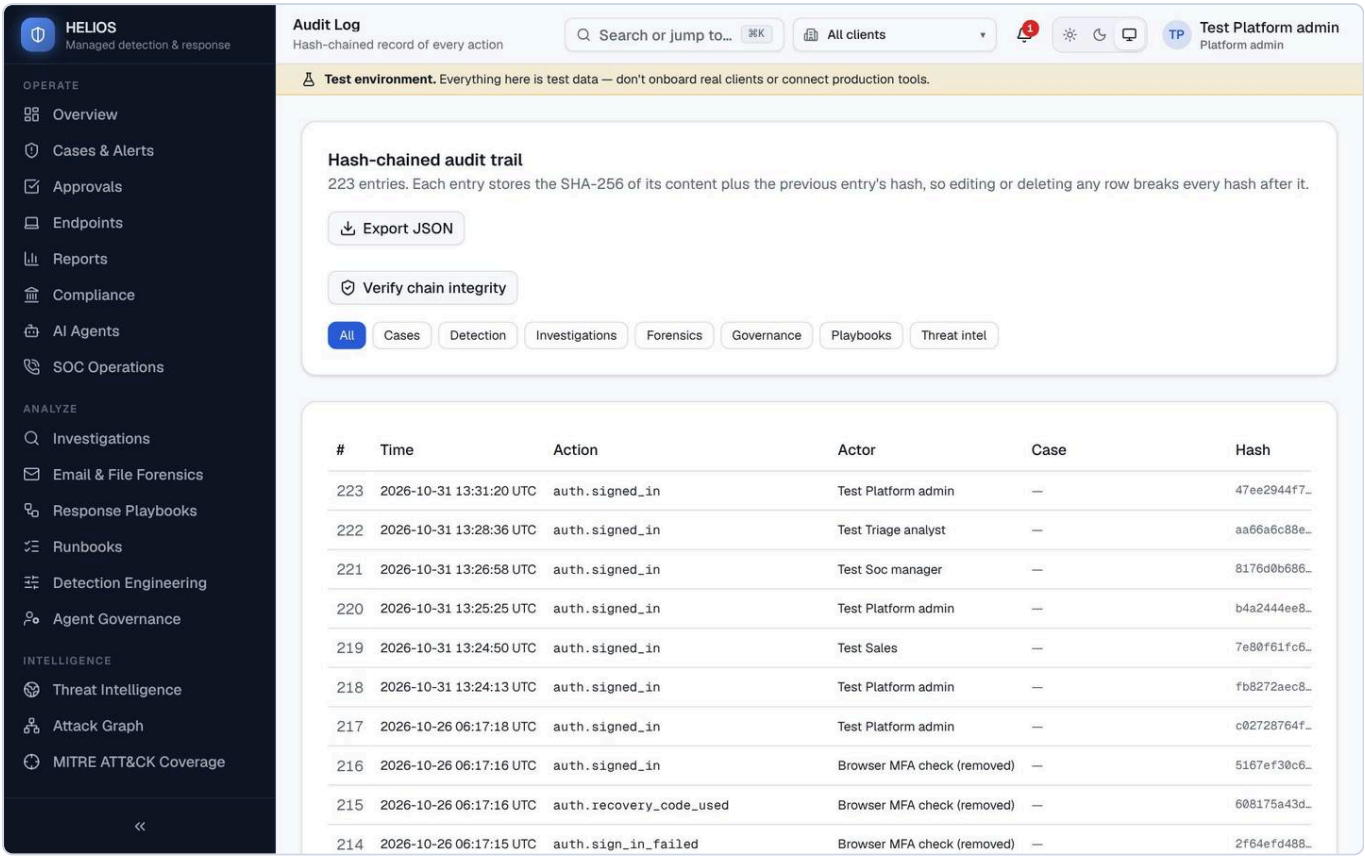


Figure 54. Audit Log, as seen by a platform admin (test environment).

How to use it

- Each entry stores the SHA-256 of its content plus the previous entry's hash, so editing or deleting any row breaks every hash after it. Verify chain integrity recomputes them all and names the first altered entry if there is one.
- Filter by area: cases, detection, investigations, forensics, governance, playbooks, threat intel.
- Export JSON exports the chain for an auditor or a regulator.

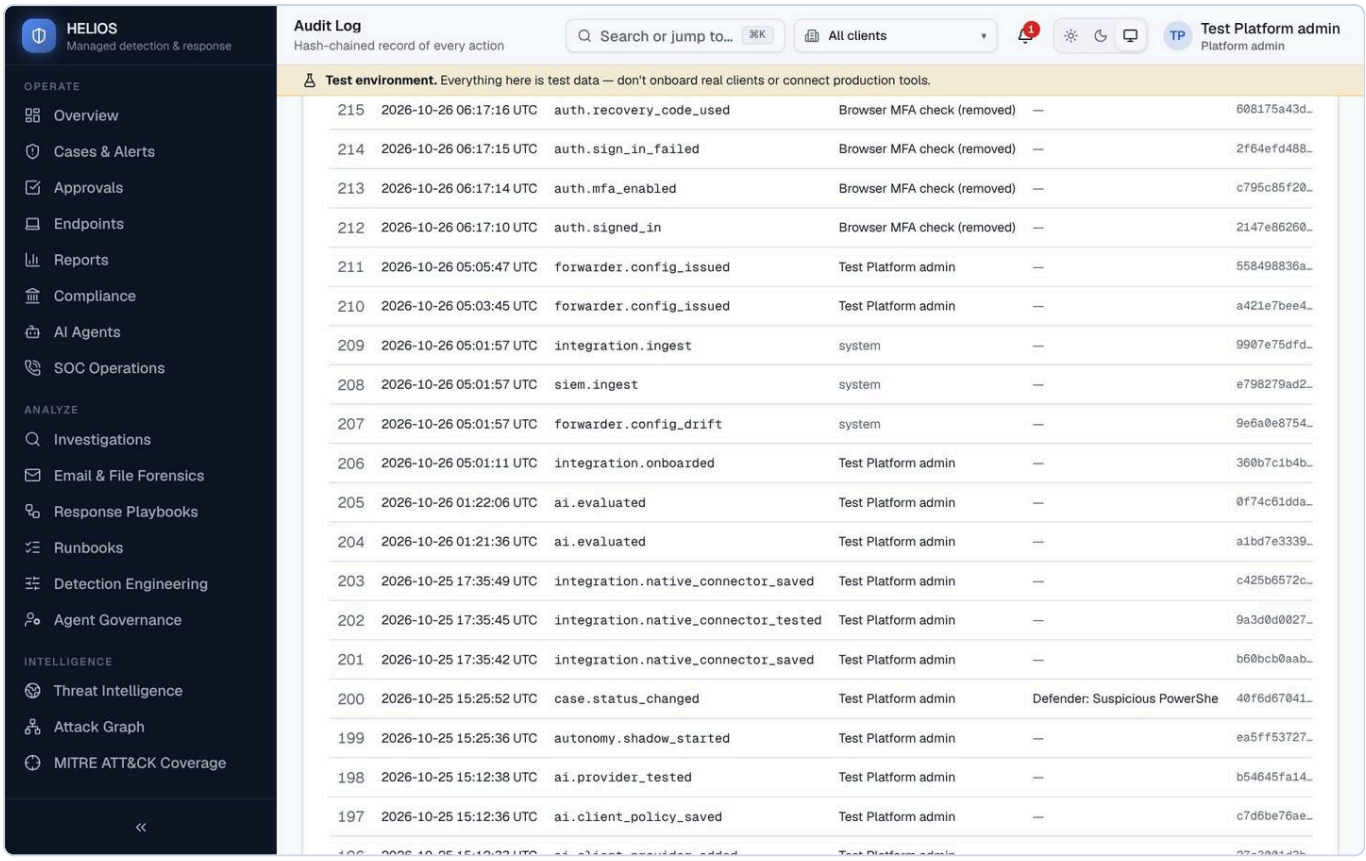


Figure 55. Audit Log, further down the page.

API Tokens

Personal tokens for the REST API, for ticketing systems, SOAR and client portals. A token acts as you:

same role, client and queues.

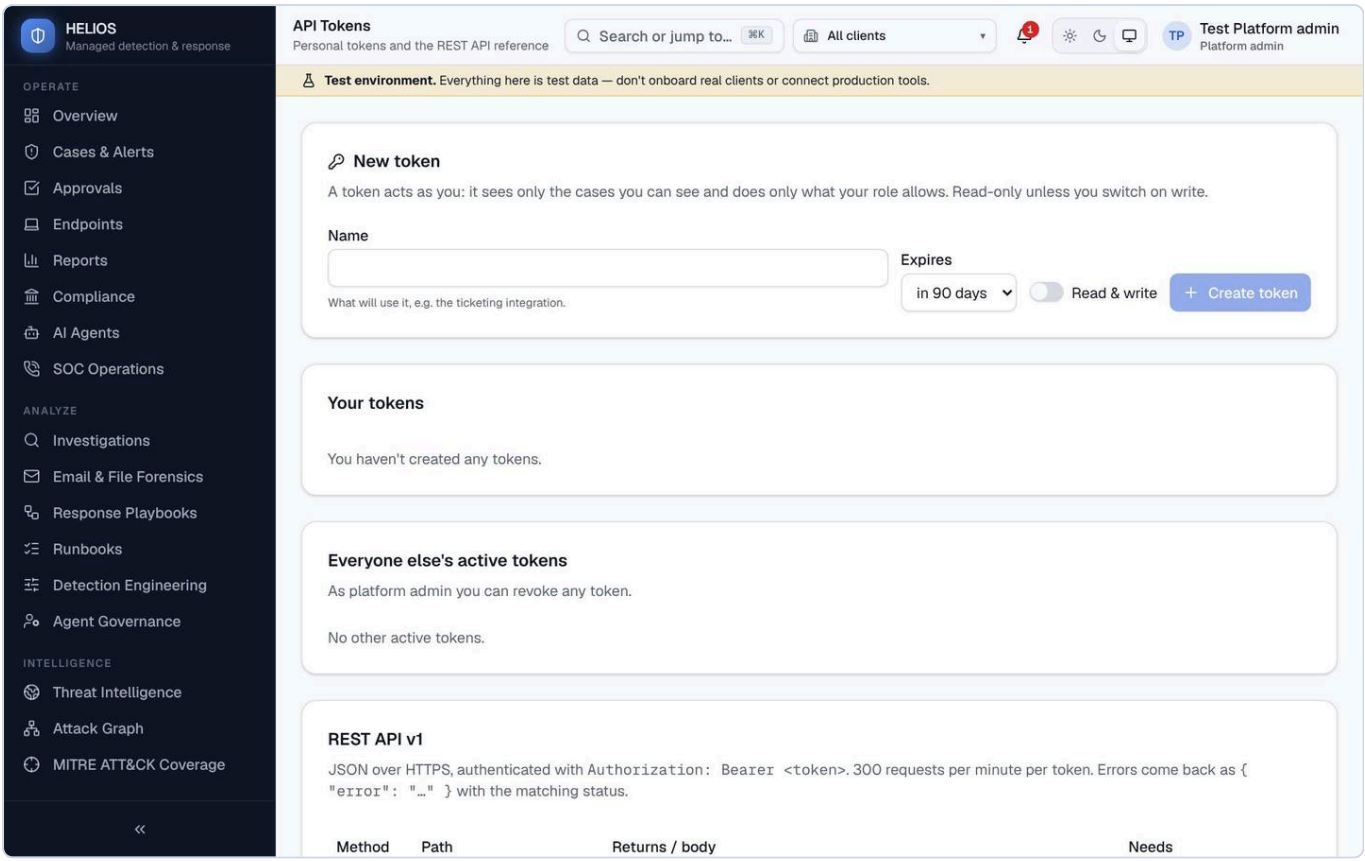


Figure 56. API Tokens, as seen by a platform admin (test environment).

How to use it

- New token: give it a name (what will use it), an expiry (90 days by default), and switch on Read & write only if it must change things. Create token shows it once; copy it then. Your tokens lists yours to revoke.
- As platform admin you also see everyone else's active tokens and can revoke any of them.
- The page includes the API reference: endpoints, parameters and examples. Calls are limited to 300 a minute per token, and every change through the API is audited.

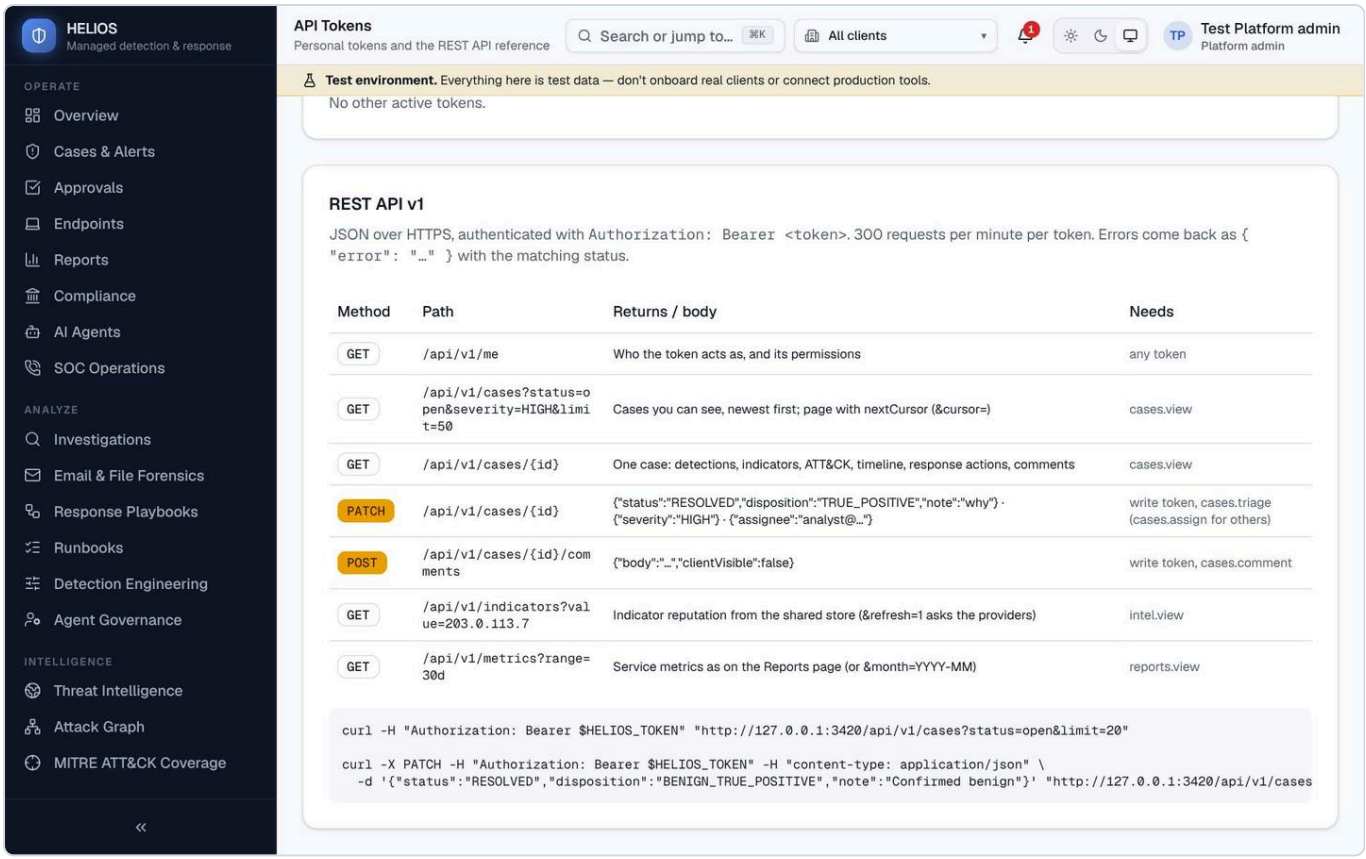


Figure 57. API Tokens, further down the page.

Settings

Platform-wide settings: environment, sign-in and encryption, company profile, public website, alert

channels, provider keys, forensics thresholds and IP geolocation.

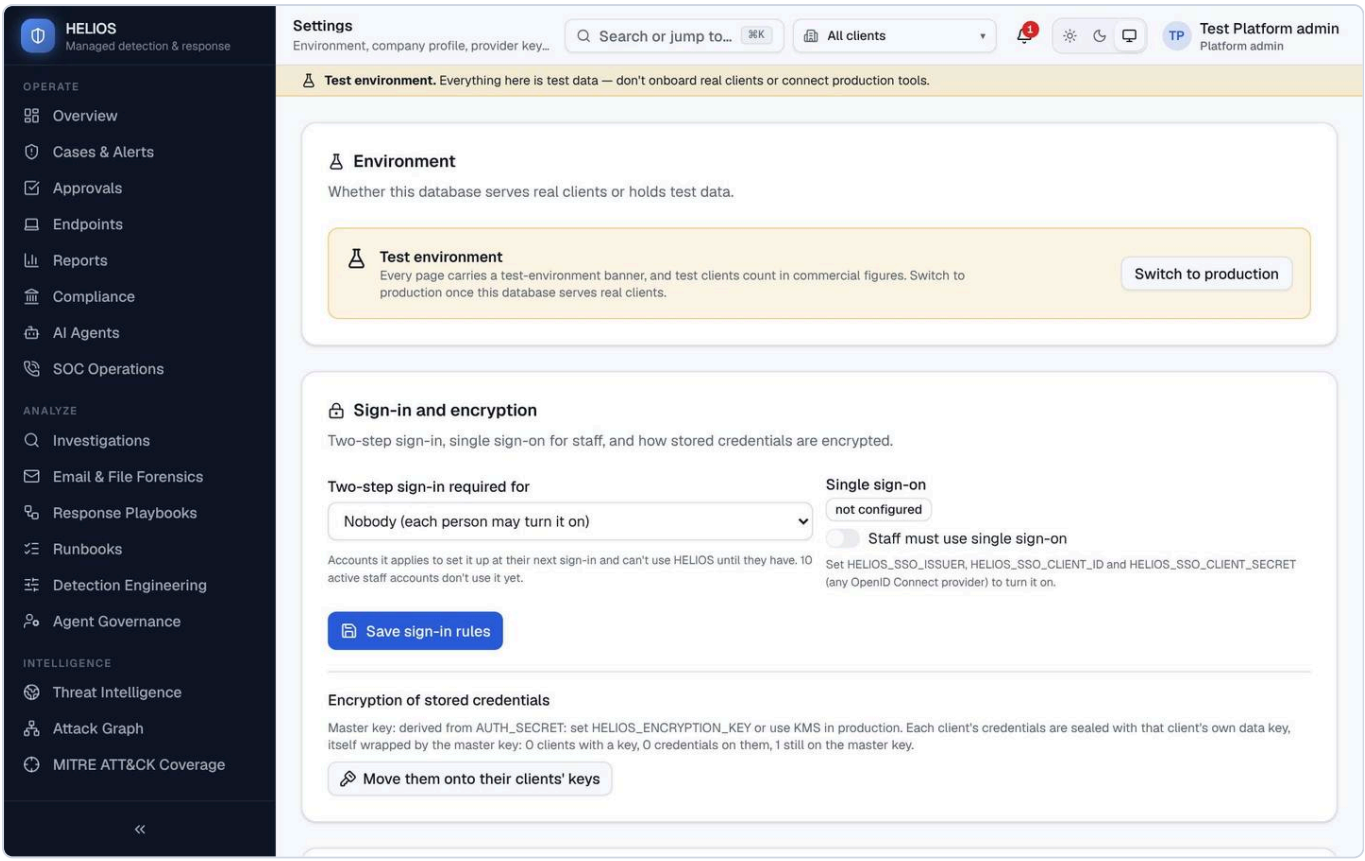


Figure 58. Settings, as seen by a platform admin (test environment).

How to use it

- Environment: production or test (a test environment shows a banner on every page). Switch to production once the database serves real clients.
- Sign-in and encryption: Two-step sign-in required for (nobody, staff or everyone) and Staff must use single sign-on, then Save sign-in rules. Under encryption, Move them onto their clients' keys re-seals older credentials with each client's own key.
- Company profile (printed on quotes, reports and invoices), the public website, alert channels (email, SMS, WhatsApp) and provider keys (VirusTotal, AbuseIPDB, MxToolbox, Slack). Keys are encrypted; only their last four characters are ever shown.
- Forensics thresholds and trusted link domains, and the IP geolocation databases in use.

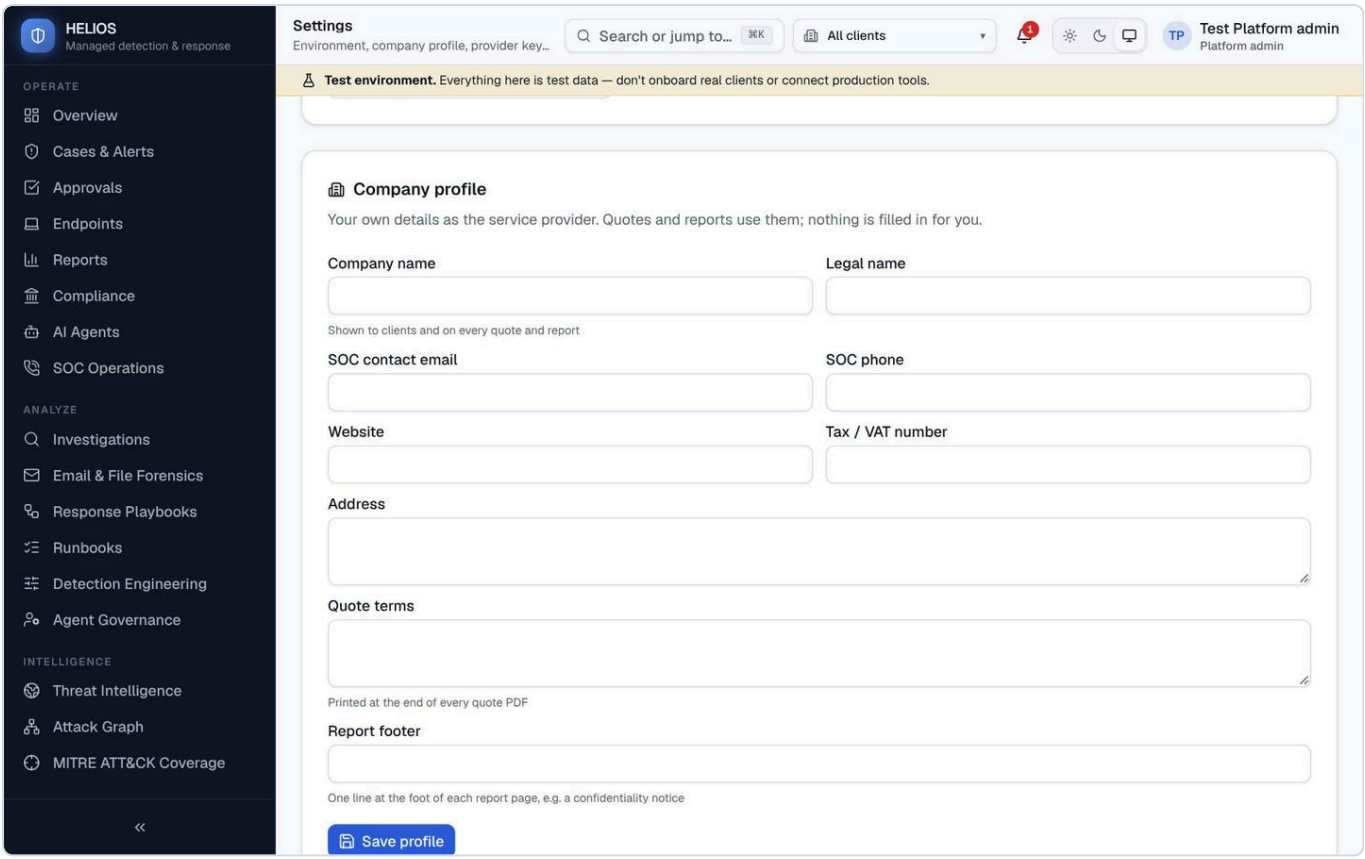


Figure 59. Settings, further down the page.

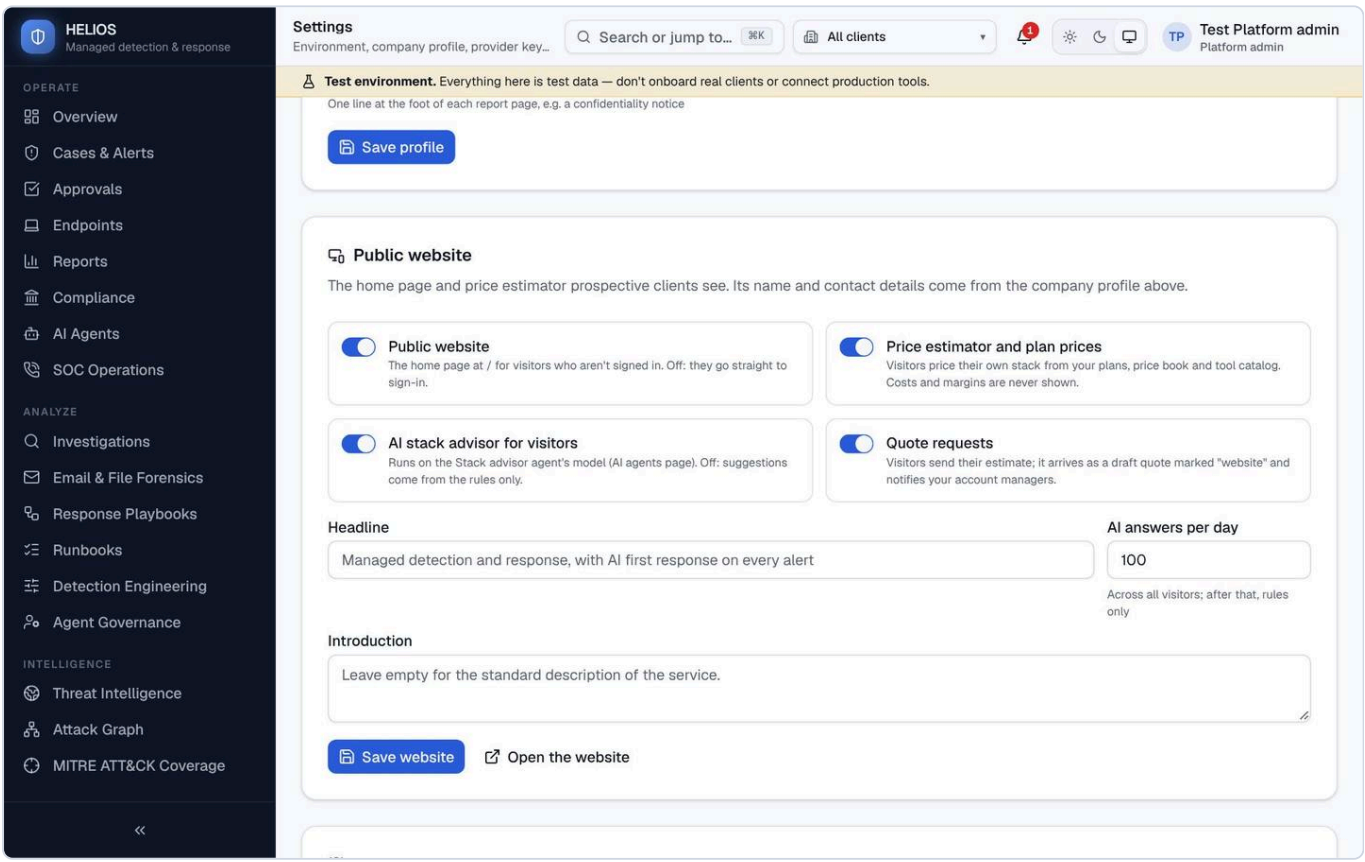


Figure 60. Settings, further down the page.

Build & Price

Configure a client's protection stack, compare vendors and plans, and estimate what it costs: monthly, annual, one-time, maintenance and over three years.

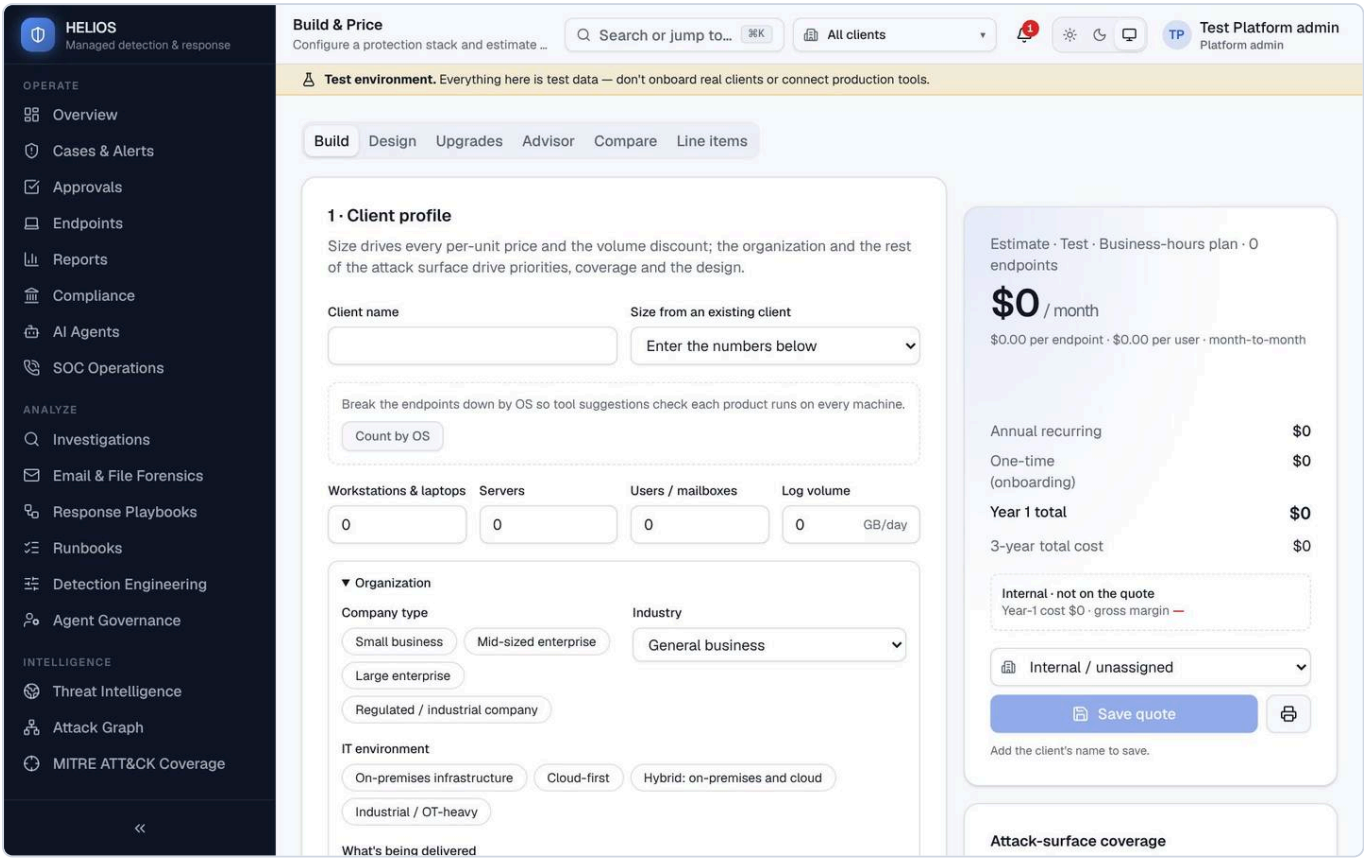


Figure 61. Build & Price, as seen by a platform admin (test environment).

How to use it

- The tabs: Build, Design, Upgrades, Advisor, Compare and Line items. The estimate on the right (monthly, annual recurring, one-time, year 1 and three-year totals) updates as you go.
- Build: describe the client (size from an existing client, or the numbers: workstations, servers, users, log volume, endpoints counted by OS; company type, industry, IT environment) and choose a tool per protection layer. Advisor suggests one per layer for a goal (best value, lowest cost, strongest protection) with its reasons.
- The upgrade review compares each tool with every alternative: what it adds or removes, whether it runs on the client's machines, and the price difference. The solution design lays out priorities, a phased rollout and coverage gaps.
- See internal cost and margin next to the price.
- Save quote (it is versioned), and print or download it as a PDF; a quote never shows cost or margin.
- Edit the price book, plans and exchange rates (Pricing setup).

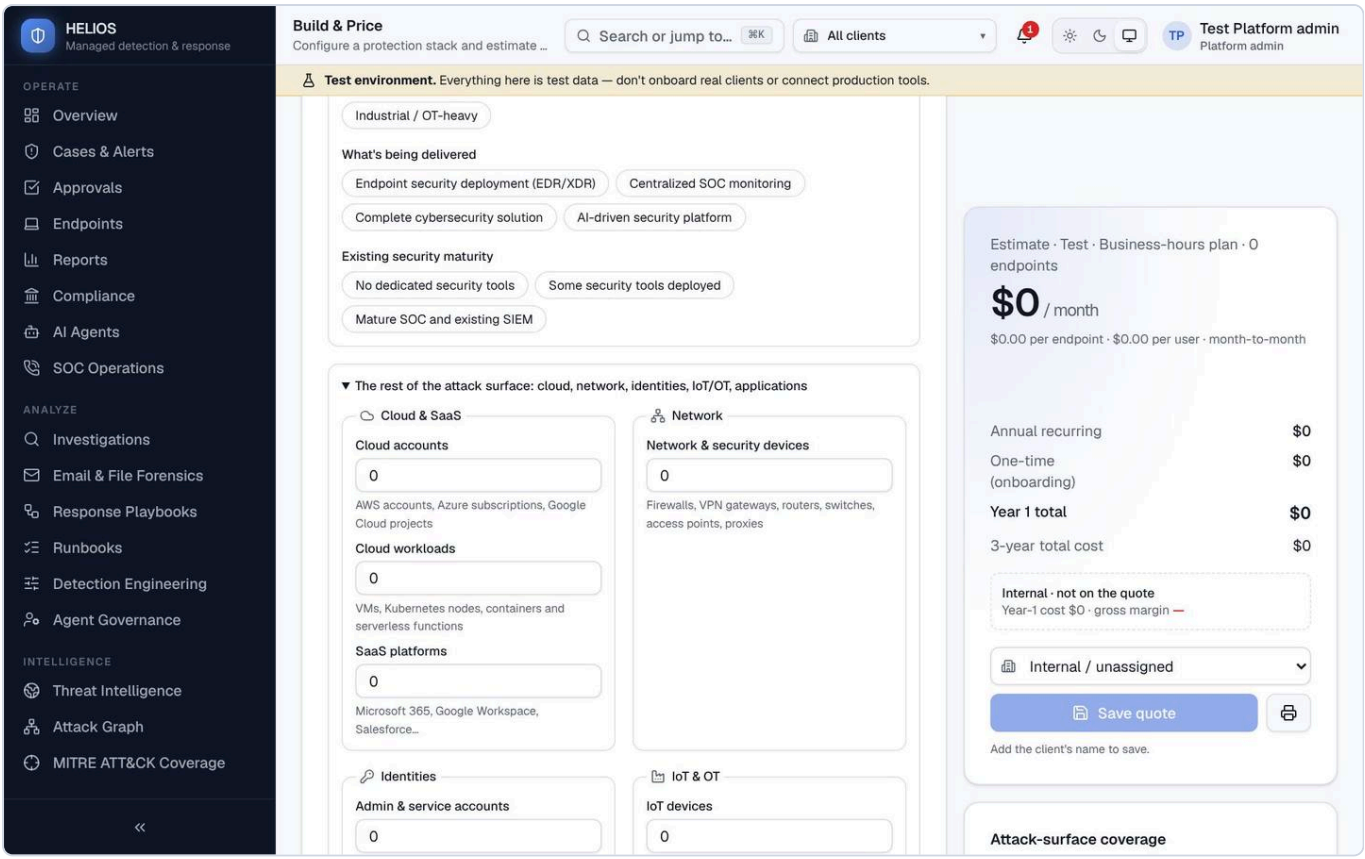


Figure 62. Build & Price, further down the page.

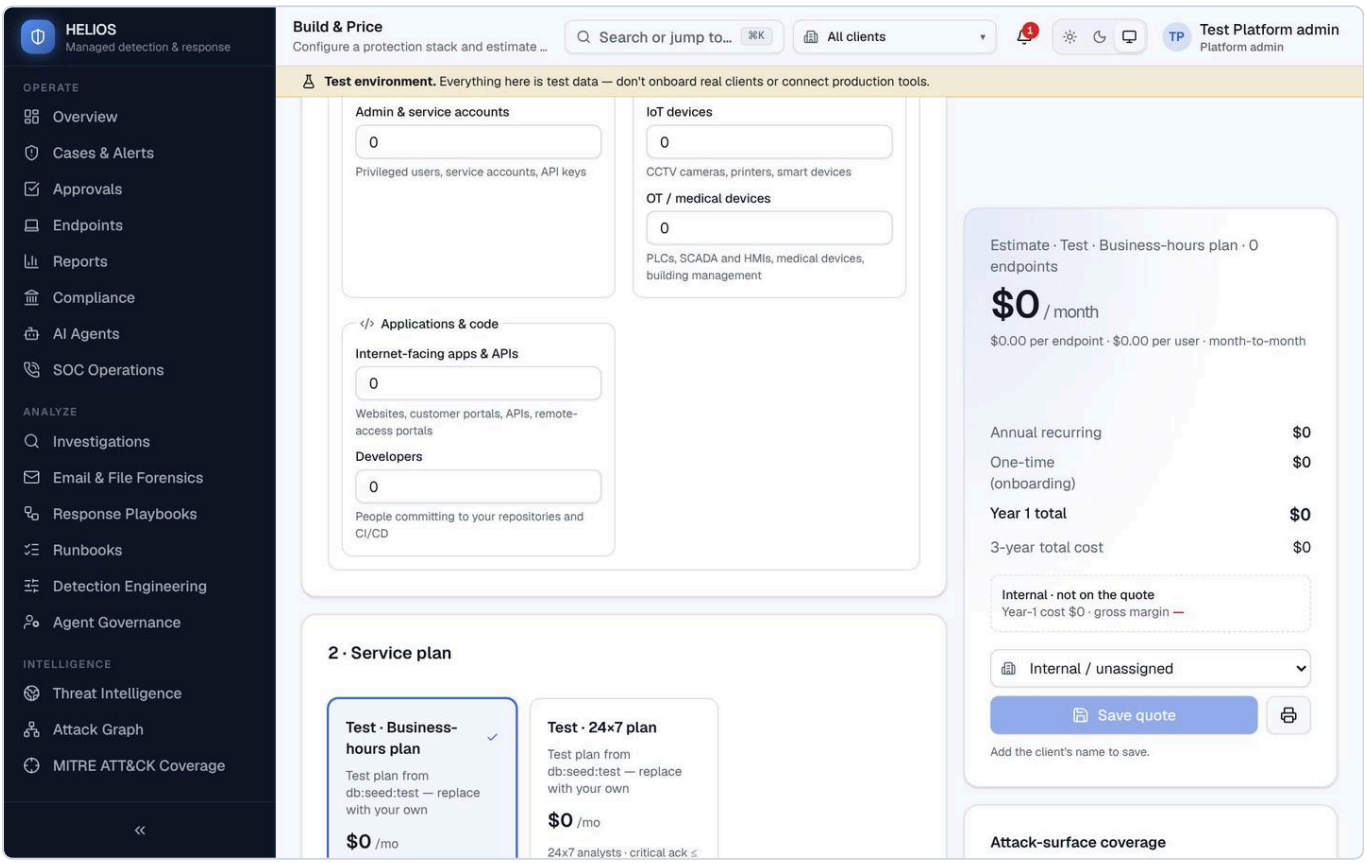


Figure 63. Build & Price, further down the page.

Contracts & Billing

Pilots with measured success criteria, contracts and GST invoices.

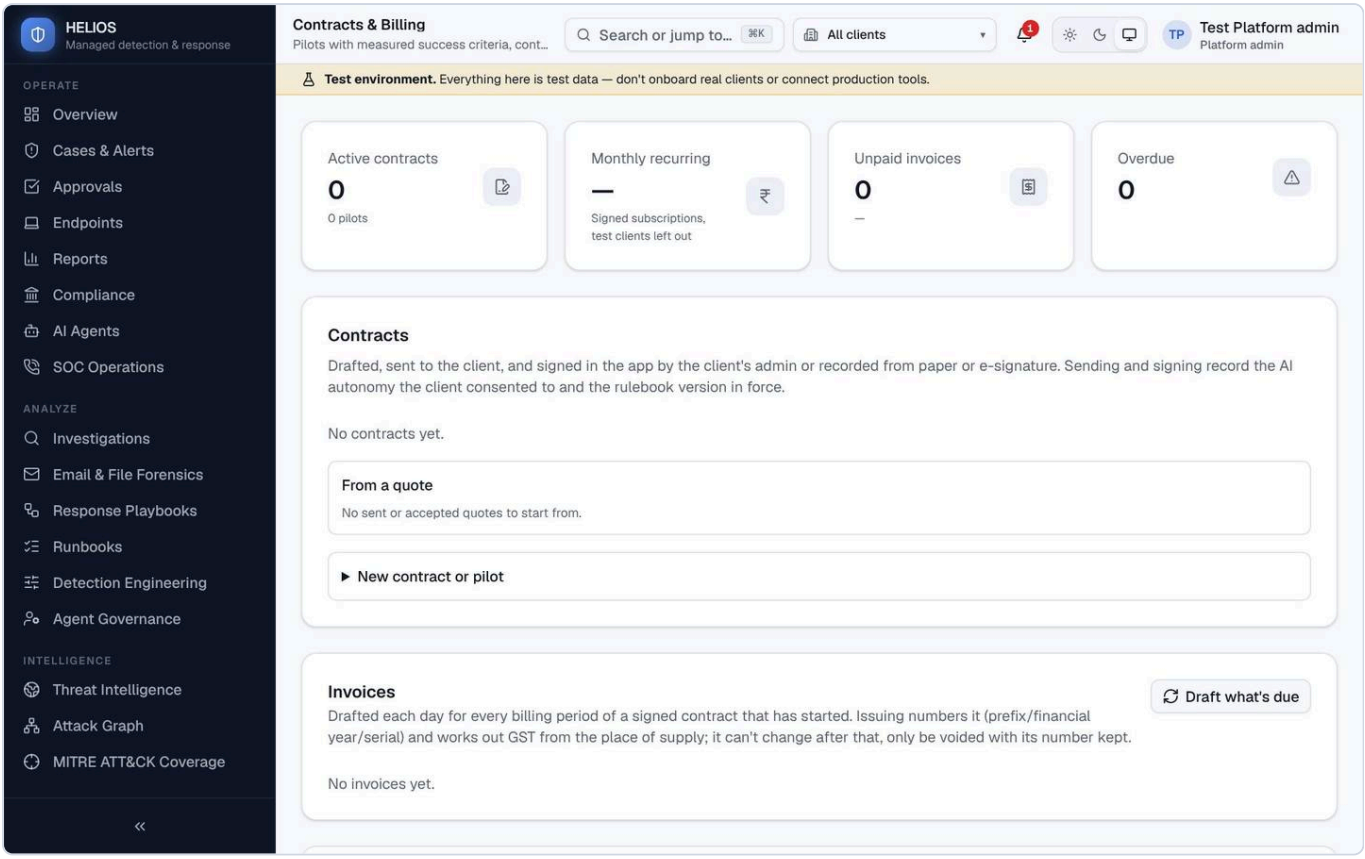


Figure 64. Contracts & Billing, as seen by a platform admin (test environment).

How to use it

- Pilots carry success criteria measured from HELIOS's own records (time to acknowledge and contain, SLA, AI agreement, wrong containments, sources live). A week before the end you get their status; convert the pilot to a subscription or record that it isn't continuing.
- Contracts: start From a quote (sent or accepted) or with New contract or pilot; send it with your terms of service; the client's admin signs it in the app (or you record a signature on paper or an e-signature). Sending and signing record the AI autonomy the client consented to and the rulebook version in force; signing is refused if either changed since sending.
- Invoices are drafted each day for every billing period of a signed contract that has started (Draft what's due does it now). Issuing numbers the invoice (prefix, financial year, serial) and works out GST from the place of supply; it can't change after that, only be voided with its number kept. Mark it paid when the money arrives.

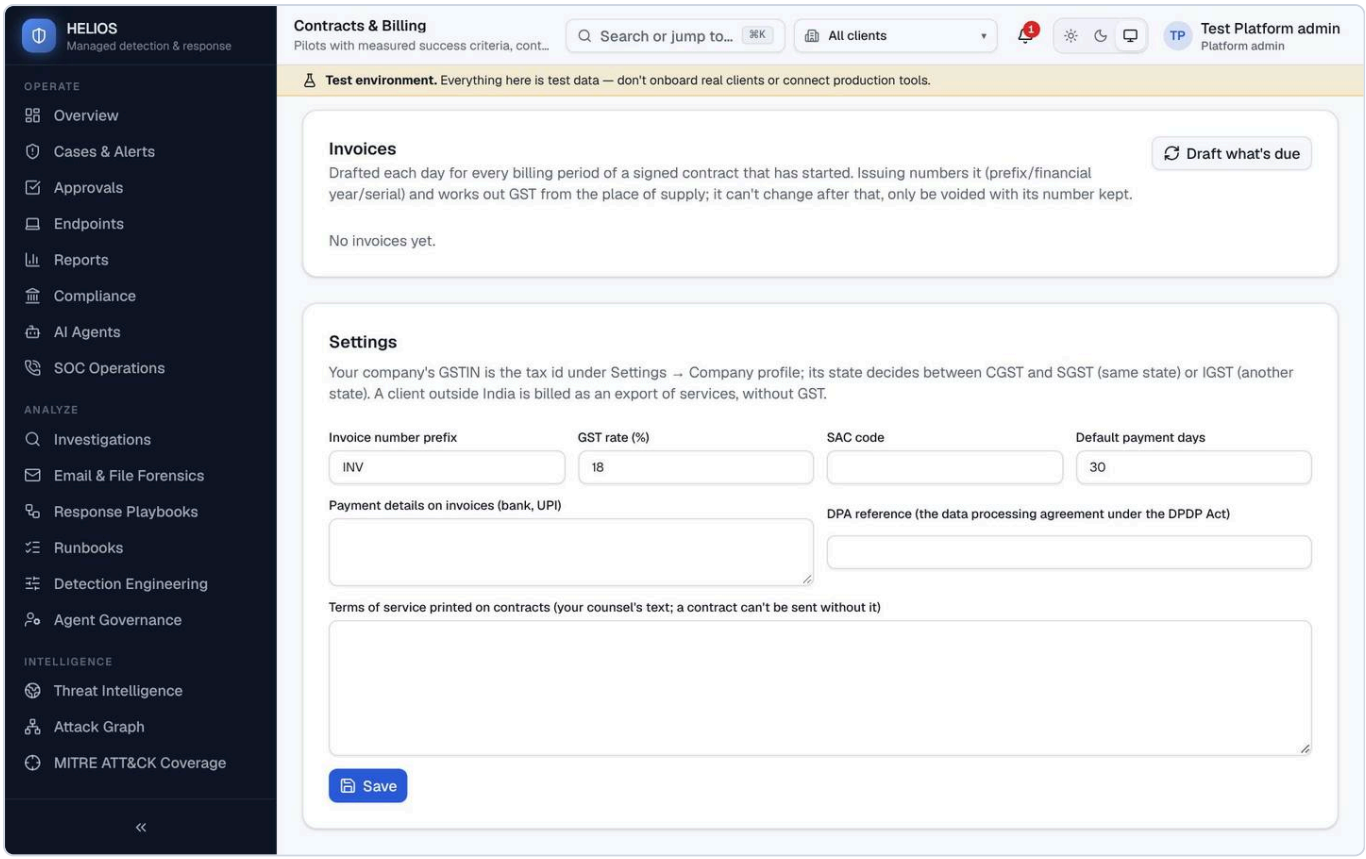


Figure 65. Contracts & Billing, further down the page.

A client's rulebook

The rulebook page for one client (here the test client Clinic).

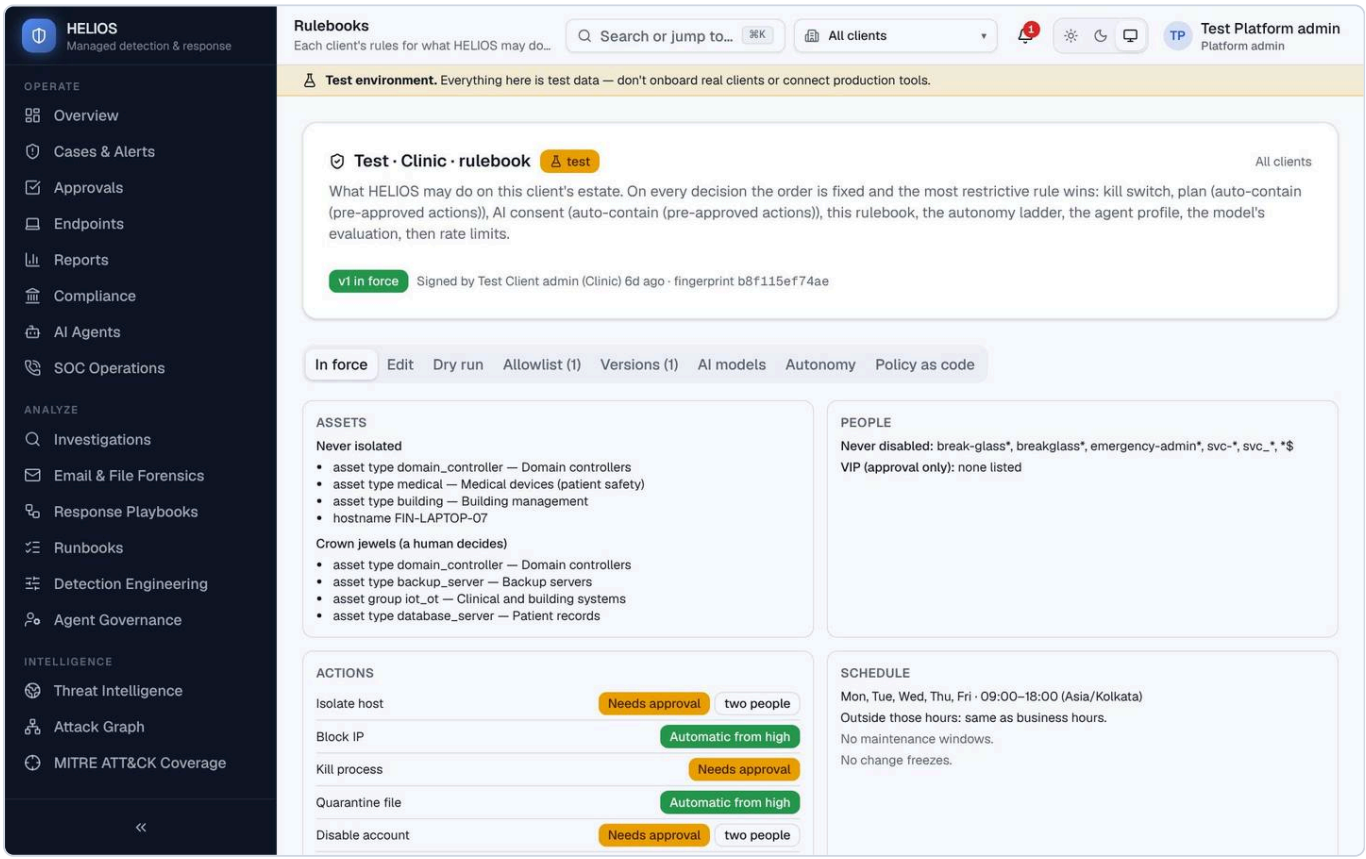


Figure 66. A client's rulebook, as seen by a platform admin (test environment).

How to use it

- In force shows the assets (never isolated, crown jewels), the people (never disabled, VIPs), each action's mode (automatic from a severity, needs approval, two people), the schedule, the approver chain, and data handling and obligations. The other tabs are as described above.

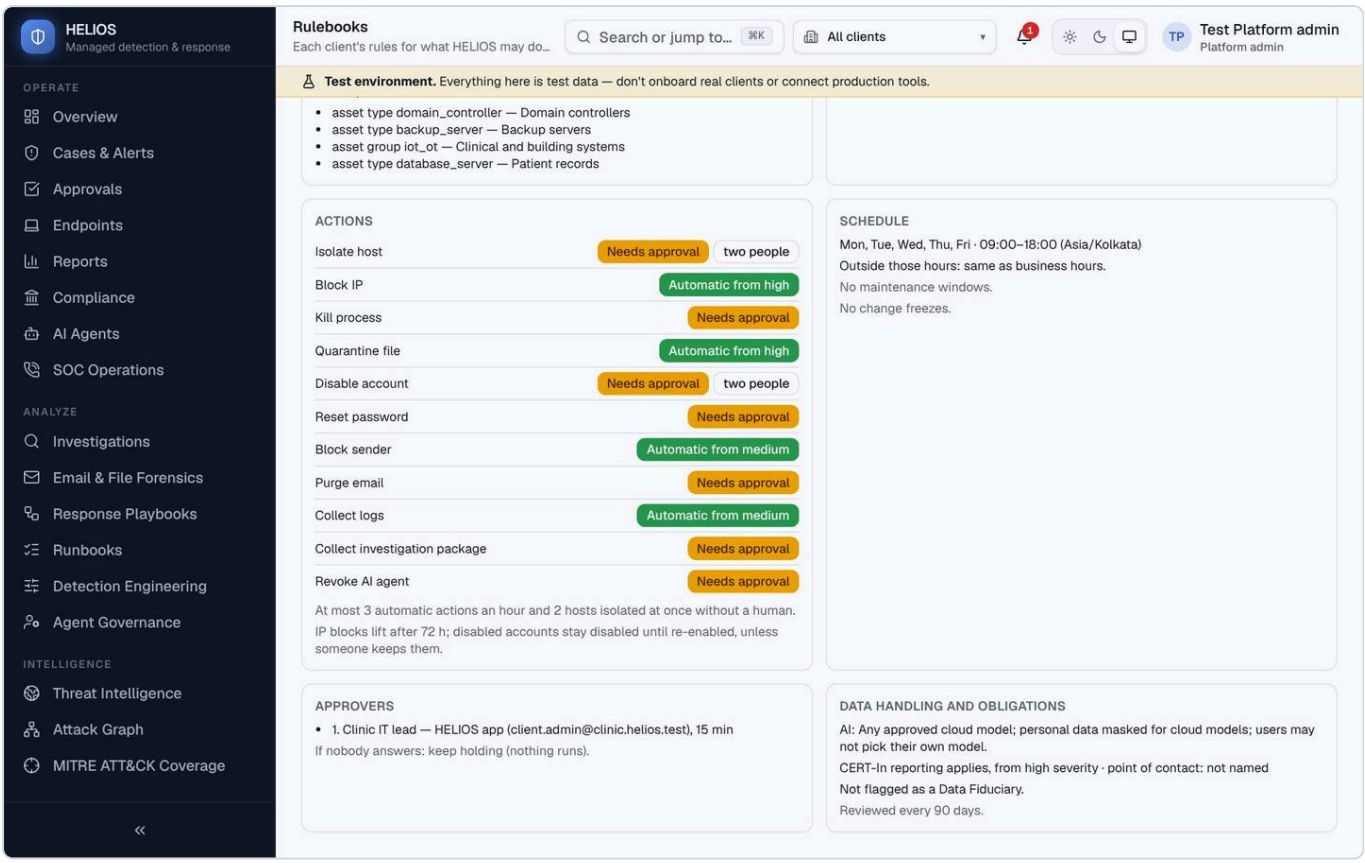


Figure 67. A client's rulebook, further down the page.

A client

One client's page: plan and SLA, contacts, consent, onboarding status, tools, users and data keys.

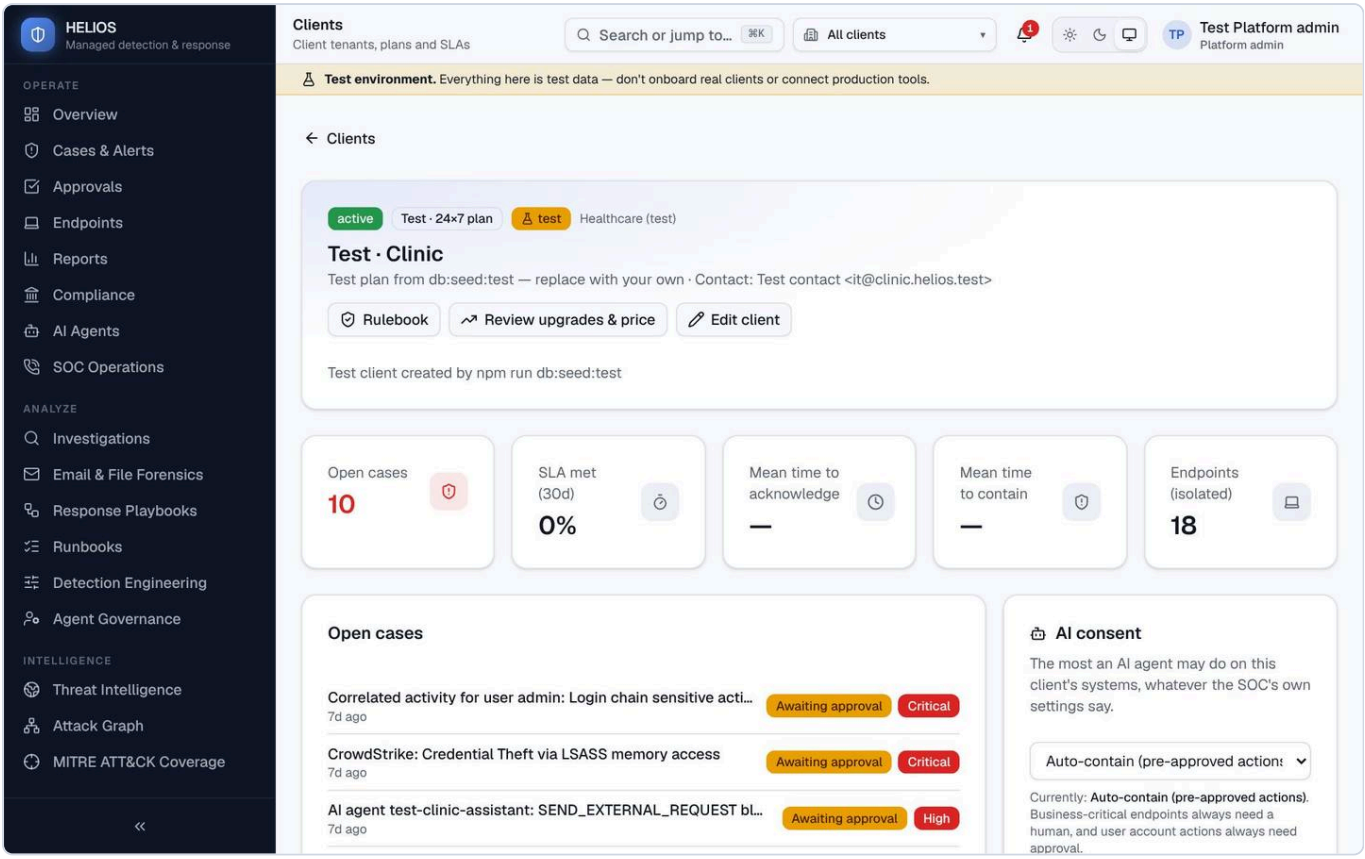


Figure 68. A client, as seen by a platform admin (test environment).

How to use it

- The page shows the client's open cases, SLA met, mean time to acknowledge and contain, endpoints, and buttons for its Rulebook, Review upgrades & price, and Edit client.
- AI consent: the most an AI agent may do on the client's systems, whatever the SOC's own settings say. Set it as the client agreed.
- Encryption key: Rotate key re-seals the client's credentials under a new data key; Destroy keys (for an offboarded, suspended client) makes its stored credentials unreadable for good, and asks you to type the client's name. Only platform admins can.

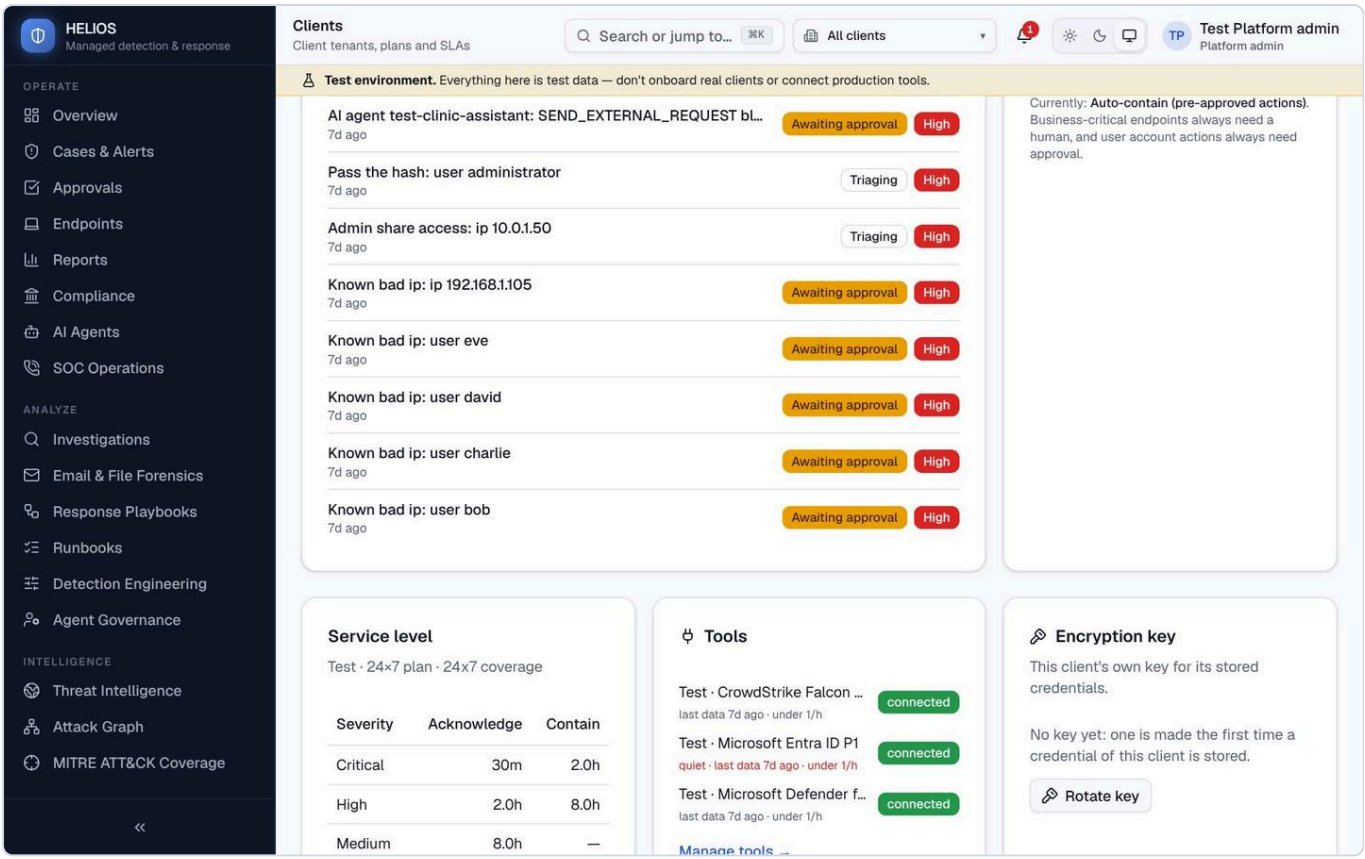


Figure 69. A client, further down the page.

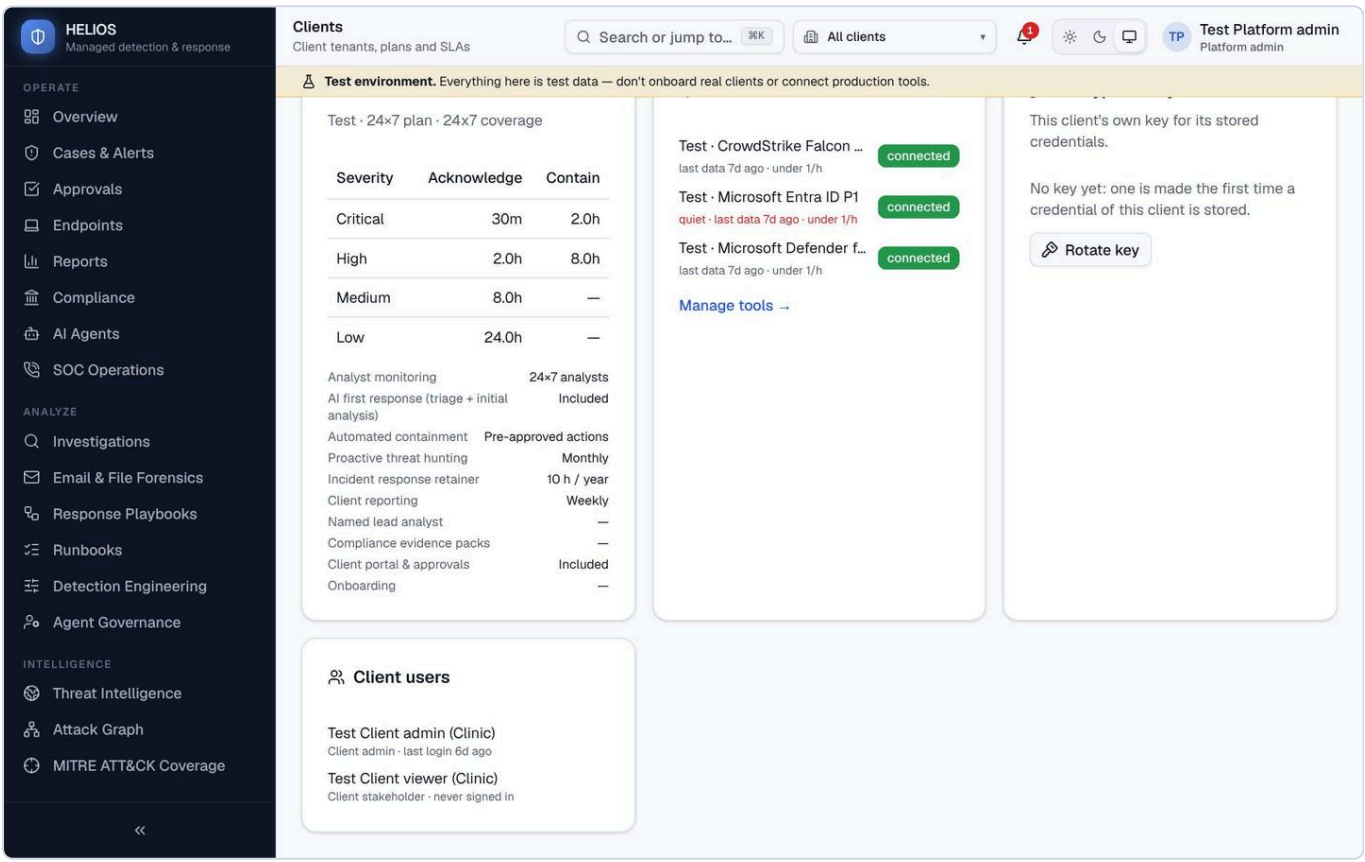


Figure 70. A client, further down the page.

A saved quote

One quote: its configuration, the price breakdown and its versions.

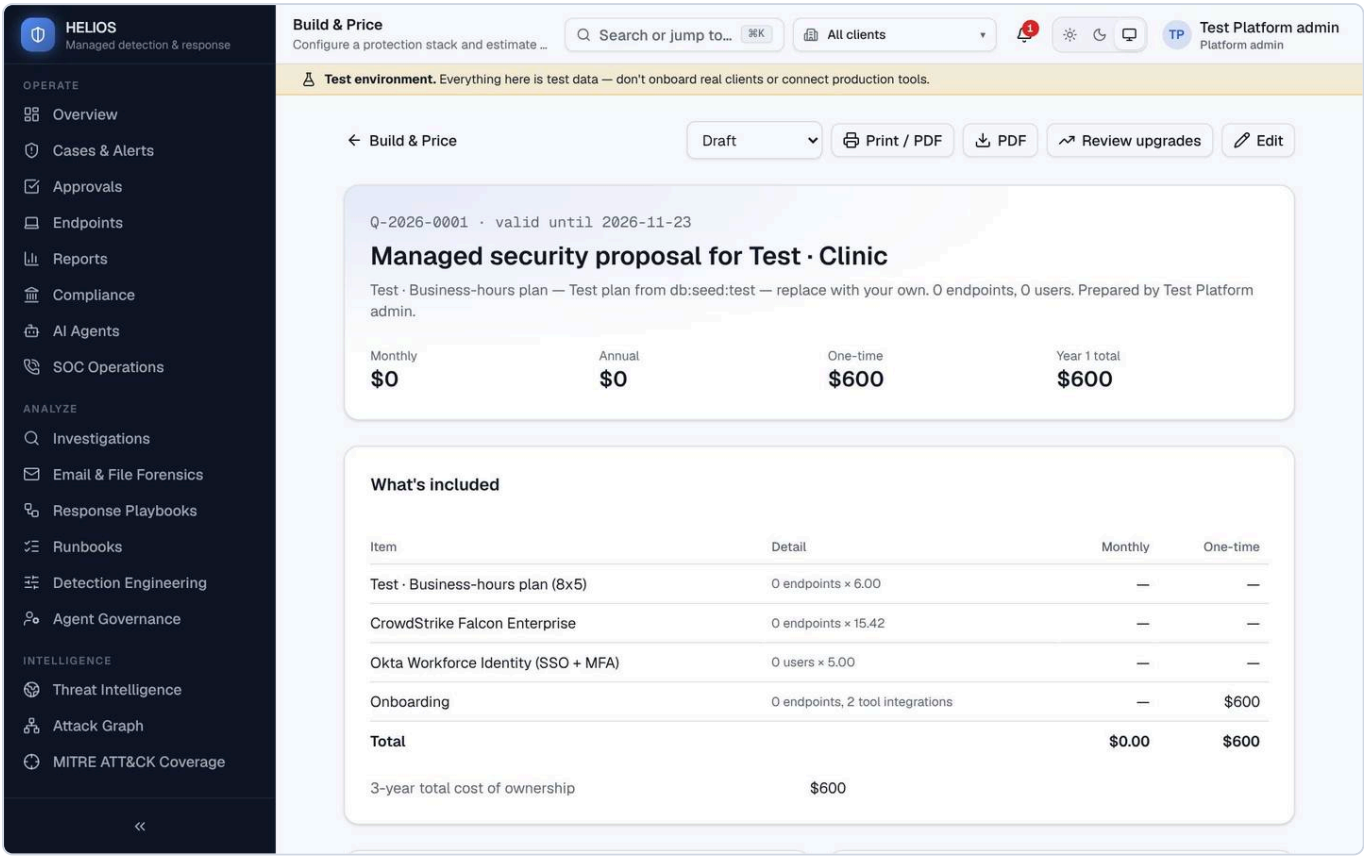


Figure 71. A saved quote, as seen by a platform admin (test environment).

How to use it

- Reopen it in the configurator, download the PDF, or start a contract draft from it on Contracts & Billing.

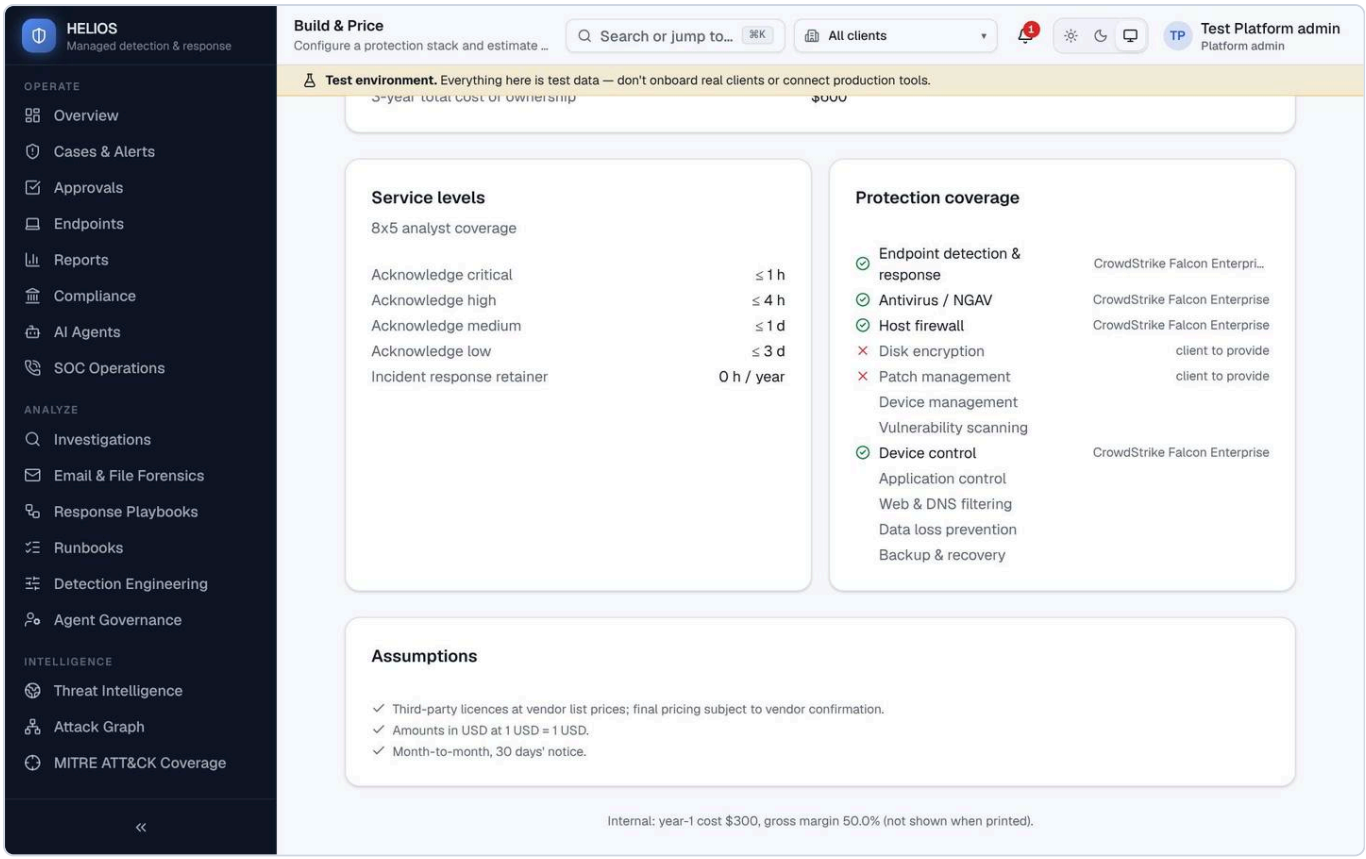


Figure 72. A saved quote, further down the page.

Daily digests

Once a day, what HELIOS did on its own, what was waiting, and the deadlines coming up.

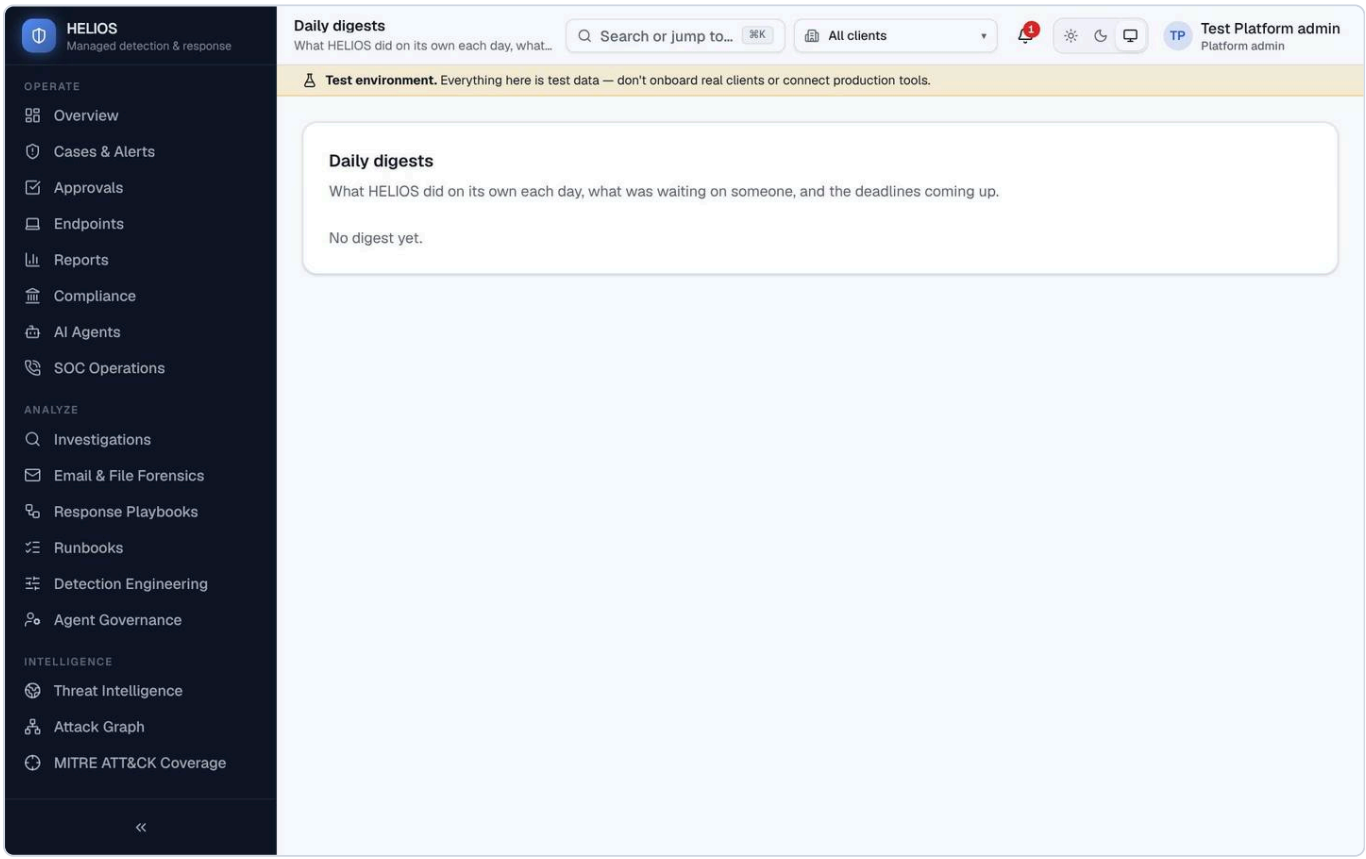


Figure 73. Daily digests, as seen by a platform admin (test environment).

How to use it

- Each digest lists actions the AI took without asking (with its autonomy level, rollbacks and QA calls), cases opened and closed, containment waiting for approval, regulatory deadlines within a day, quiet sources and the week's QA results.
- Choose under My alerts whether it also reaches you by email.

My alerts

Which notifications also reach you outside HELIOS, by email, SMS or WhatsApp.

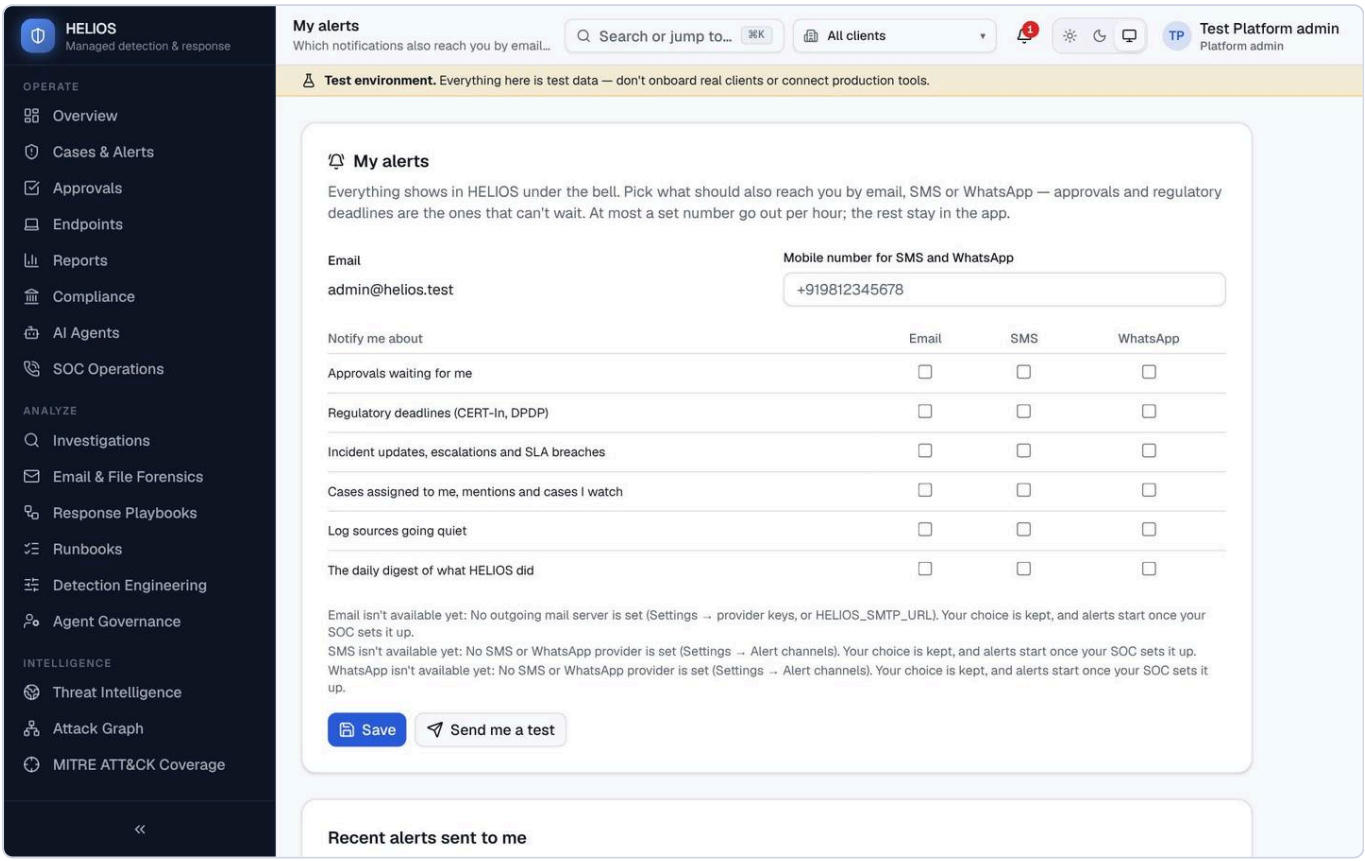


Figure 74. My alerts, as seen by a platform admin (test environment).

How to use it

- Every notification lands in the app (the bell at the top). Here choose, per kind of notification, the channels it also comes by, and give your phone number in international format (+91…).
- Deliveries that couldn't go (a channel not set up, no number, over the hourly cap) are listed with the reason.

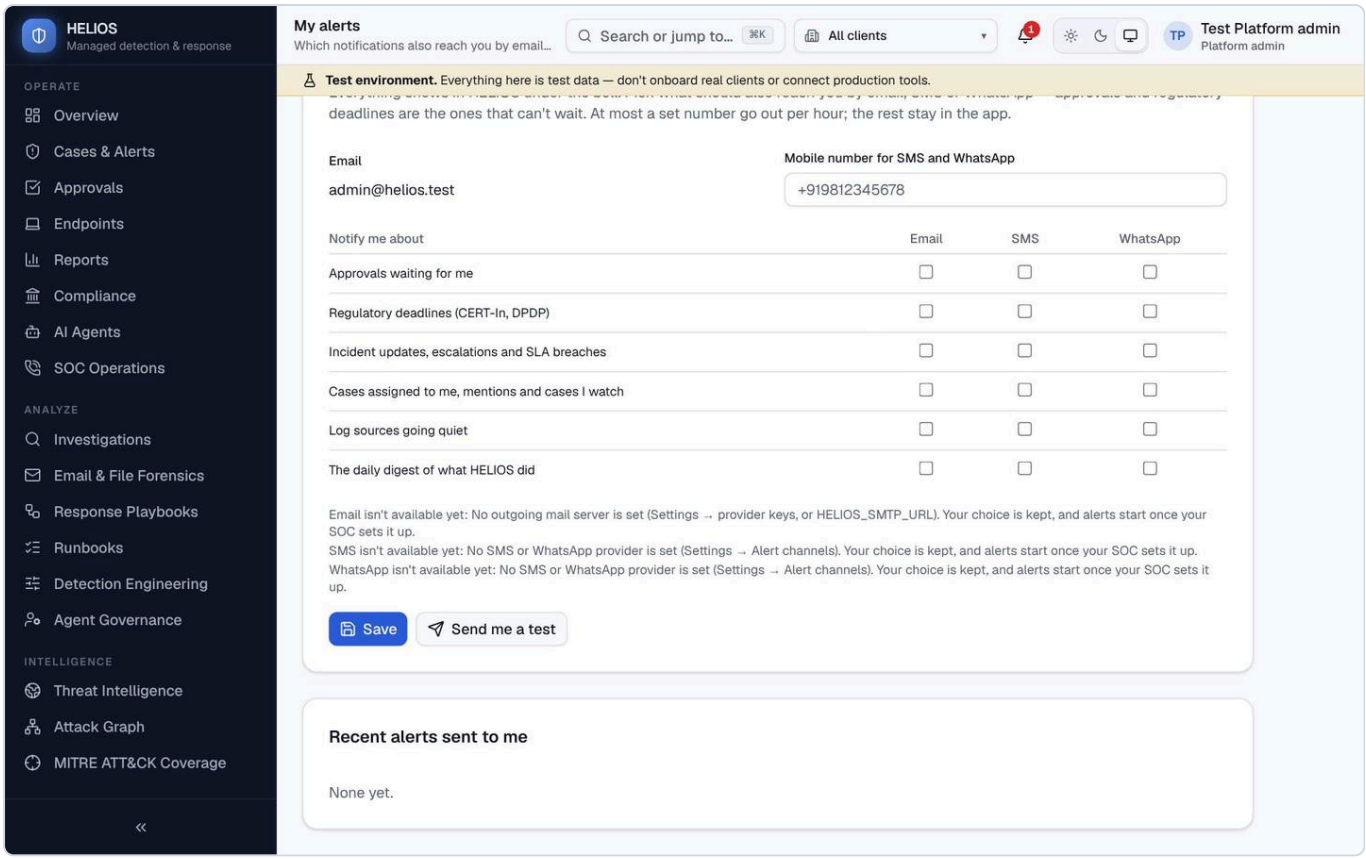


Figure 75. My alerts, further down the page.

5. Procedures

Step by step, for the tasks your role does most.

5.1 First-run set-up

1. Open Overview. Work through Get HELIOS ready from top to bottom.
2. Define your service plans (SLA targets per severity, the most autonomy each plan allows) and fill in the price book under Build & Price → Pricing setup.
3. Price the quote-only tools you resell, or switch off the ones you don't sell.
4. Connect an AI provider on AI Agents (or leave agents on the built-in engine for now).
5. Invite your analysts on Team & Roles.
6. Add your first client on Clients, then onboard its security tools on Tools & Integrations.
7. Turn on the threat feeds you want on Threat Intelligence.

5.2 Invite a person and give them the right access

1. Open Team & Roles and choose Invite user.
2. Enter name and email, and pick the role that matches the job (see the role matrix on the page). Give the least access that does the job.
3. For a client user, choose their organisation; they will only ever see it.
4. For an analyst who works some queues only, tick their security domains (endpoint, identity, email, network, cloud, AI agents).
5. Save. HELIOS shows a temporary password once: pass it to the person by a secure channel. They sign in with it, and set up two-step sign-in if your policy requires it.

5.3 Someone leaves or changes job

1. Open Team & Roles, find the person, and deactivate the account (leaver) or change the role (mover). It applies on their next request; there is no session to wait out.
2. Revoke their personal API tokens if they had any (they stop working with the account in any case).
3. Check the on-call rota on SOC Operations and the approver chains in clients' rulebooks for their name.

5.4 Require two-step sign-in

1. Open Settings → Sign-in and encryption.
2. Choose whether two-step sign-in is required for staff or for everyone. People it applies to are taken to set it up (an authenticator app, plus ten recovery codes) before anything else.
3. Optionally turn on single sign-on only for staff once your identity provider is connected.
4. For someone who lost their phone, open Team & Roles → their account → Reset two-step.

5.5 Configure AI safely

1. On AI Agents → Model providers, add a provider and press Test. Keys are encrypted and never shown again.
2. Run an evaluation for each agent you will use (Evaluations). A model a client picks only suggests until it passes.
3. Under Agent types, set each agent's provider, reasoning effort and autonomy. Start with Suggest only.
4. Under Safety controls, set the minimum severity for AI first response and the cap on unattended containment per client per hour.
5. In an emergency, Engage kill switch stops all AI action at once, for every client.

5.6 Offboard a client

1. Make sure contracts and invoices are settled (Contracts & Billing) and no legal hold is needed (Data Retention).
2. On Clients, set the client to suspended.
3. Once the retention period your contract promises is honoured or the data is returned, open the client's page → Encryption key → Destroy keys and type the client's name to confirm. Their stored credentials become unreadable, backups included.

5.7 Give an auditor or regulator the audit trail

1. Open Audit Log and press Verify chain integrity: it should report the chain intact.
2. Filter to the area asked for, and Export JSON.
3. Record in the case or a comment who received the export and why.

6. Guidelines

Do

- Keep a separate day-to-day account for ordinary SOC work; use the admin account only for administration.
- Require two-step sign-in for every staff account.
- Change roles rather than sharing accounts; every action is attributed to a person in the audit chain.
- Test AI models on Evaluations before letting them act, and start clients at the lowest autonomy.
- Keep the test environment for testing: it shows a banner on every page; never onboard real clients or production tools in it.

Don't

- Don't paste keys, passwords or client secrets into comments or case notes; use the encrypted provider-key and credential fields.
- Don't destroy a client's keys until offboarding is final: it can't be undone.
- Don't edit the database directly: changes outside HELIOS bypass the audit chain and the row-level isolation.

For everyone

- Never share your account or password; every action is recorded under your name.
- Keep client information inside HELIOS and to the people who need it.
- Treat text inside evidence (command lines, emails, file names) as possibly written by an attacker: it may address 'the AI' or 'the analyst' to mislead. HELIOS flags such text; don't follow it.
- Every action in HELIOS is reversible or recorded, and audited; write down why you did it.

> The client's signed rulebook binds everyone: the AI, playbooks and people. If HELIOS refuses an action, the reason names the rule. Don't work around it in a vendor's console; ask the client to change the rule.

7. Reference

Terms you will see

Case statuses

Open	New; nobody has acknowledged it yet.
Triaging	Someone is looking at it.
Investigating	Under analysis.
Awaiting approval	Containment is waiting on someone's yes.
Contained	The threat is stopped; clean-up may continue.
Resolved	Closed with a disposition.
False positive	Closed: it wasn't a threat.

Severities

Critical, High, Medium, Low, Info	Set by the detections and adjustable by analysts. The client's plan sets the acknowledgement SLA for each.

Dispositions (when closing)

True positive	A real threat.
Benign true positive	Real activity, but expected or authorised (an admin's tool, a pen test).
False positive	Not what the detection thought.

Tiers

T1	Triage.
T2	Analysis.
T3	Incident response.

Response action statuses

Suggested	A recommendation only; autonomy doesn't allow acting.
Pending approval	Waiting on an approver.
In progress	The vendor was asked; waiting for it to confirm. Doesn't count as containment yet.
Executed	Done and confirmed by the vendor (or by an analyst for a manual action).
Manual required	Approved, but no connector can do it: someone does it in the tool's console and records it.
Blocked	A guardrail or the client's rulebook refused it; the reason is shown.
Rejected	An approver said no.
Rolled back	Undone.

Failed	The vendor or HELIOS couldn't do it.
Shadow	Shadow mode: what the AI would have done, recorded silently and shown once the case closes.

AI autonomy (per agent)

Off	The agent doesn't run.
Suggest only	It proposes; people decide.
Contain with approval	Its containment waits for approval.
Auto-contain	It may act on pre-approved actions above its confidence bar.

Autonomy ladder (per client and action type)

L0 observe	Shadow mode.
L1 suggest	Suggestions only.
L2 approve	Acts once approved.
L3 act and notify	Acts, and tells people at once.
L4 act	Acts, and reports in the daily digest.

Evidence references

D, E, I, T, C	Detections, events, indicators, timeline entries and context in the case; Q for what the investigation agent's tools returned. Every AI finding cites them.

If something doesn't work

- **A page says you don't have access**: your role doesn't include it. Ask your administrator if your job needs it.
- **An action was refused**: the message names the rule (the client's rulebook, a guardrail, the plan or the kill switch).
- **An action is 'Manual required'**: no connector can do it. Do it in the tool's own console, then record what you did on the case.
- **You don't get alerts by SMS or email**: check My alerts for your channel choices, your phone number, and deliveries skipped with their reason.
- **You can't sign in**: after five failures the account locks for 15 minutes; your administrator can reset your password or two-step sign-in.