HELIOS USER GUIDE · SECURITY OPERATIONS

Tier-1 triage analyst

Procedures, usage, guidelines, and your roles and access

The first line: you take new cases, check what the AI made of them, close the noise and pass the real ones on with what you found.

Platform	HELIOS managed detection and response
Guide	03 of 11 · Tier-1 triage analyst
Date	31 October 2026
Version	Written against HELIOS 809e500
Screenshots	Taken in the HELIOS test environment with test data, signed in as a tier-1 triage analyst. Your screens show your own clients and data.

Contents

1. Your role at a glance
2. Roles and access
3. Getting started
4. Your screens
5. Procedures
6. Guidelines
7. Reference

> You have **18** of 50 permissions and **14** pages in your sidebar: Overview, Cases & Alerts, Endpoints, Compliance, AI Agents, SOC Operations, Response Playbooks, Runbooks, Threat Intelligence, Attack Graph, MITRE ATT&CK Coverage, Tools & Integrations, Rulebooks, Clients.

1. Your role at a glance

The first line: you take new cases, check what the AI made of them, close the noise and pass the real ones on with what you found.

What you are responsible for

- Acknowledge new cases in your queues within the SLA.
- Check the AI's first response: agree or disagree, and say why when it's wrong.
- Validate the evidence; close false positives with a disposition and a reason.
- Escalate real or unclear cases to tier 2 with a clear note; request containment where it's needed.
- Page on-call for anything urgent you can't move.

A typical day or shift

1. Open Cases & Alerts, newest first, and filter to your queues.
2. Work SLA-due cases first: the SLA badge turns red when breached.
3. Keep the bell (notifications) in view for mentions and handoffs.

Who to turn to

Escalate cases to tier 2; ask the SOC manager on duty about anything in the queue itself; page on-call for urgent help.

2. Roles and access

What you can see

You work across the clients the MSSP serves. The client switcher at the top focuses every page on one client, or All clients. If your account is limited to security domains (for example endpoint and identity), your case queues show only those domains.

Your access comes from your role, and HELIOS checks it on the server for every page, button and API call; hiding a button is only a convenience. A change to your role applies on your next click. Everything you do is written to the hash-chained audit log under your name.

Your pages

Overview, Cases & Alerts, Endpoints, Compliance, AI Agents, SOC Operations, Response Playbooks, Runbooks, Threat Intelligence, Attack Graph, MITRE ATT&CK Coverage, Tools & Integrations, Rulebooks, Clients, plus My alerts and your two-step sign-in settings in the user menu.

Your permissions

What the permission lets you do	You
Cases	
See the overview dashboard (dashboard.view)	Yes
See cases in their queues (or their own tenant, for clients) (cases.view)	Yes
See every queue, regardless of assigned security domains (cases.view_all)	No
Acknowledge cases, change status and severity (cases.triage)	Yes
Assign cases to other analysts (cases.assign)	No
Escalate or hand off a case to the next tier (cases.escalate)	Yes
Comment on cases and @mention teammates (cases.comment)	Yes
Share cases with the client and external channels (cases.share)	No
Response	
Run containment playbooks (subject to approval) (containment.request)	Yes
Approve or reject pending containment (containment.approve)	No
Isolate and release endpoints directly (endpoints.contain)	No
Analysis	
Run log investigations (investigations.run)	No
Submit emails and files for forensic analysis (forensics.run)	No
Upload raw logs into the detection pipeline (telemetry.ingest)	No
Govern AI agents: identities, tokens, revocation (governance.manage)	No
Intelligence and detection	

Look up threat indicators (intel.view)	**Yes**
Edit the blocklist and indicator store (intel.manage)	No
See ATT&CK coverage (mitre.view)	**Yes**
Switch SIEM detectors on and off (detections.tune)	No
See response playbooks (playbooks.view)	**Yes**
Author and enable response playbooks (playbooks.manage)	No
AI agents	
See AI agent runs and configuration (ai.view)	**Yes**
Dispatch AI agents onto a case (ai.dispatch)	**Yes**
Rate AI agent output (ai.feedback)	**Yes**
Configure AI providers, agent profiles and autonomy (ai.configure)	No
Tools and endpoints	
See onboarded security tools (integrations.view)	**Yes**
Onboard and configure security tools (integrations.manage)	No
See the endpoint inventory and protection coverage (endpoints.view)	**Yes**
Clients and rulebooks	
See client tenants (tenants.view)	**Yes**
Create and edit client tenants (tenants.manage)	No
Read a client's rulebook: what HELIOS may do on its estate (rulebook.view)	**Yes**
Draft changes to a client's rulebook and its allowlist (rulebook.edit)	No
Put a client's rulebook in force (client admins sign; staff only on the client's recorded instruction) (rulebook.sign)	No
Compliance and reports	
See regulatory clocks and notices (CERT-In, DPDP) and the ISO 27001 evidence map (compliance.view)	**Yes**
Draft regulatory notices, record the client's approval, and file them (compliance.prepare)	No
Approve a regulatory notice before it is filed (the report is the client's; staff record a client's approval given elsewhere) (compliance.approve)	No
Read client-facing incident reports (reports.view)	No
SOC operations	
See the on-call rota, pages, runbooks, QA results and digests (operations.view)	**Yes**
Edit the on-call rota, runbooks and operations settings (operations.manage)	No
Review the weekly sample of what the AI did on its own (qa.review)	No
Commercial	

Use the Build & Price configurator (pricing.view)	No
See internal cost and margin (pricing.margin)	No
Edit the price book (pricing.manage)	No
See contracts, pilots and invoices (clients: their own) (billing.view)	No
Draft and record contracts, issue invoices and record payments (billing.manage)	No
Platform	
Invite users and change roles (users.manage)	No
Edit platform settings: company profile, provider keys, forensics thresholds (settings.manage)	No
Read the audit chain (audit.view)	No
Export the audit chain (audit.export)	No
Create personal API tokens for the REST API (acting with their own permissions) (api.tokens)	No

How your role fits with the others

Role	Group	What it is for
Platform admin	Leadership	Owns the platform: users, AI configuration, price book, everything else.
SOC manager	Leadership	Runs the SOC: every queue, approvals, AI autonomy, tool onboarding.
Tier-1 triage analyst	Security operations	Works the first-line queue: acknowledges, validates AI triage, escalates.
Tier-2 security analyst	Security operations	Primary analysis: investigations, forensics, containment requests.
Tier-3 incident responder	Security operations	Owns major incidents: approves containment, isolates endpoints, client comms.
Threat hunter	Security operations	Hunts across all tenants with investigations, forensics and intel.
Detection engineer	Engineering	Maintains playbooks, detections, tool integrations and ATT&CK coverage.
Auditor / compliance	Governance	Read-only oversight: cases, AI decisions and the audit chain.
Account manager	Commercial	Client onboarding, Build & Price quotes and margins.
Client admin	Client	Client-side: their own tenant's cases and tools; approves containment on their estate.
Client stakeholder	Client	Client-side read-only: dashboards, incident reports and endpoint coverage.

3. Getting started

Signing in

1. Open the HELIOS address you were given and sign in with your email and password. A new account gets a temporary password from its administrator.
2. If two-step sign-in is required, set it up the first time: scan the code with an authenticator app, enter a code, and keep the ten recovery codes somewhere safe. Each recovery code works once.
3. If your organisation uses single sign-on, use the sign-in button for your identity provider instead; it satisfies two-step sign-in.
4. Five wrong passwords or codes in a row lock the account for 15 minutes. An administrator can reset your password or your two-step sign-in if you lose your phone.

Finding your way

- **Sidebar**: your pages, grouped (Operate, Analyze, Intelligence, Platform, Commercial). The arrow at the bottom collapses it.
- **Search or jump to** (Ctrl K or Cmd K): find a case, page or client by name.
- **Client switcher**: All clients, or focus every page on one client.
- **The bell**: your notifications (mentions, hand-offs, approvals, pages). Choose which also reach you by email, SMS or WhatsApp under My alerts.
- **Light, dark or system theme**, and your **user menu**: My alerts, two-step sign-in, Sign out.
- A yellow **Test environment** banner means everything on that installation is test data.

Install it as an app

HELIOS can be installed from the browser (Install app in the address bar) as a desktop or mobile app with its own window and shortcuts.

4. Your screens

Each page you can open, what it is for, and how to use it with your access. Where a page is long, further screenshots show the rest.

Overview

Your starting page. It shows what needs you now, and it changes with your role: the SOC sees live posture, queues, SLAs and AI activity; account managers see the commercial picture; client users see their own organisation.

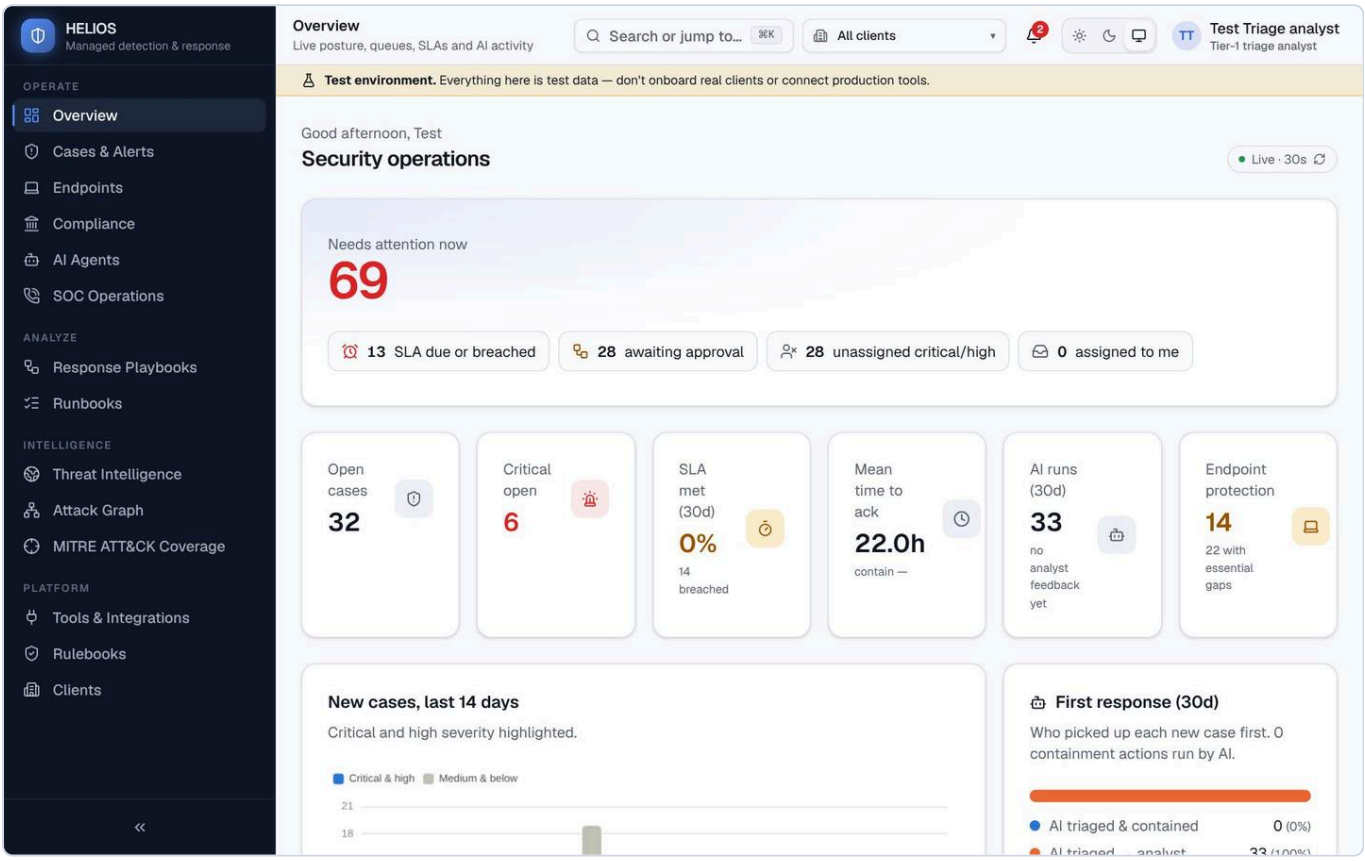


Figure 1. Overview, as seen by a tier-1 triage analyst (test environment).

How to use it

- Needs attention now adds up what is due: cases whose SLA is due or breached, containment awaiting approval, unassigned critical and high cases, and cases assigned to you. Click a figure to open the matching list.
- The Live badge means the figures refresh on their own every 30 seconds; the arrow refreshes now.

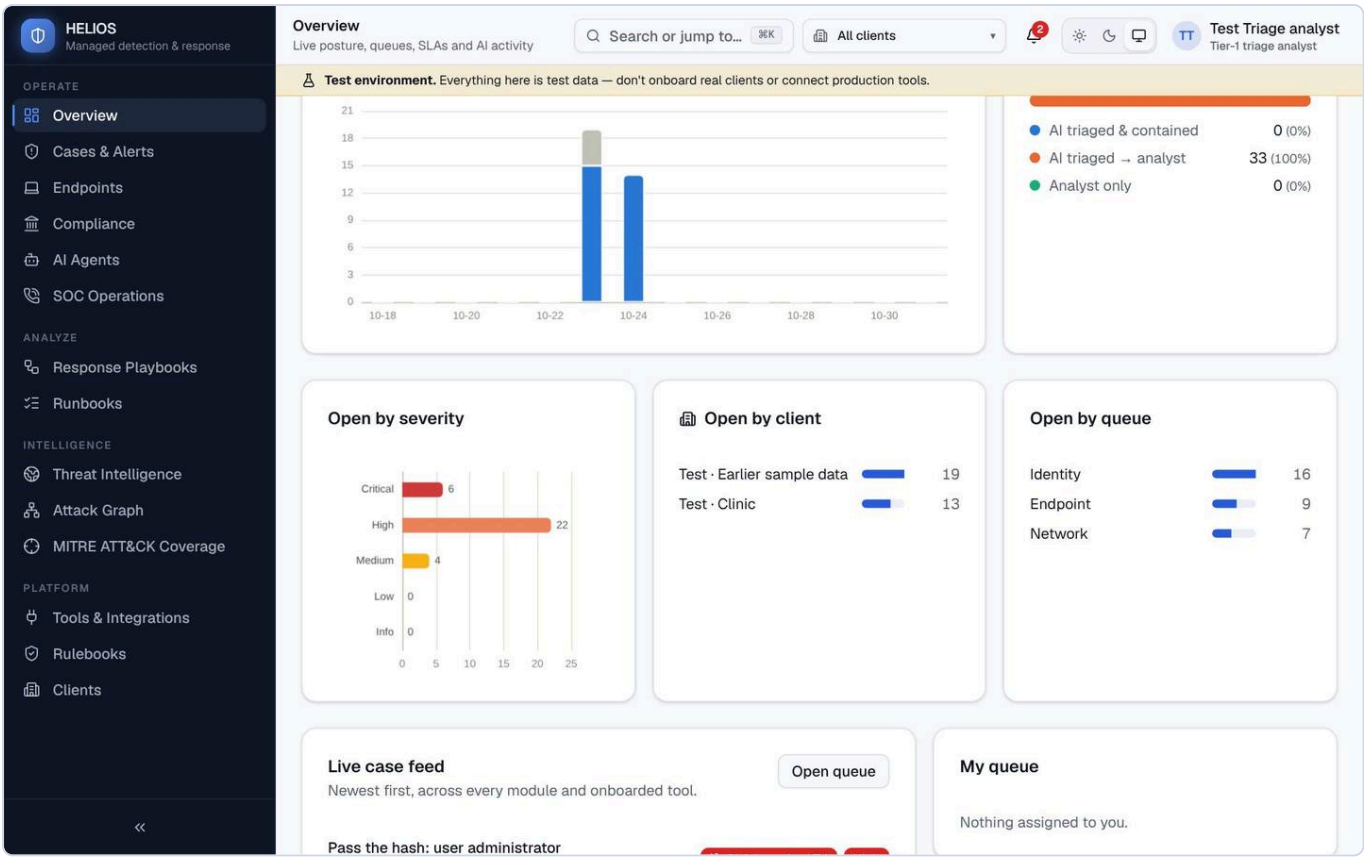


Figure 2. Overview, further down the page.

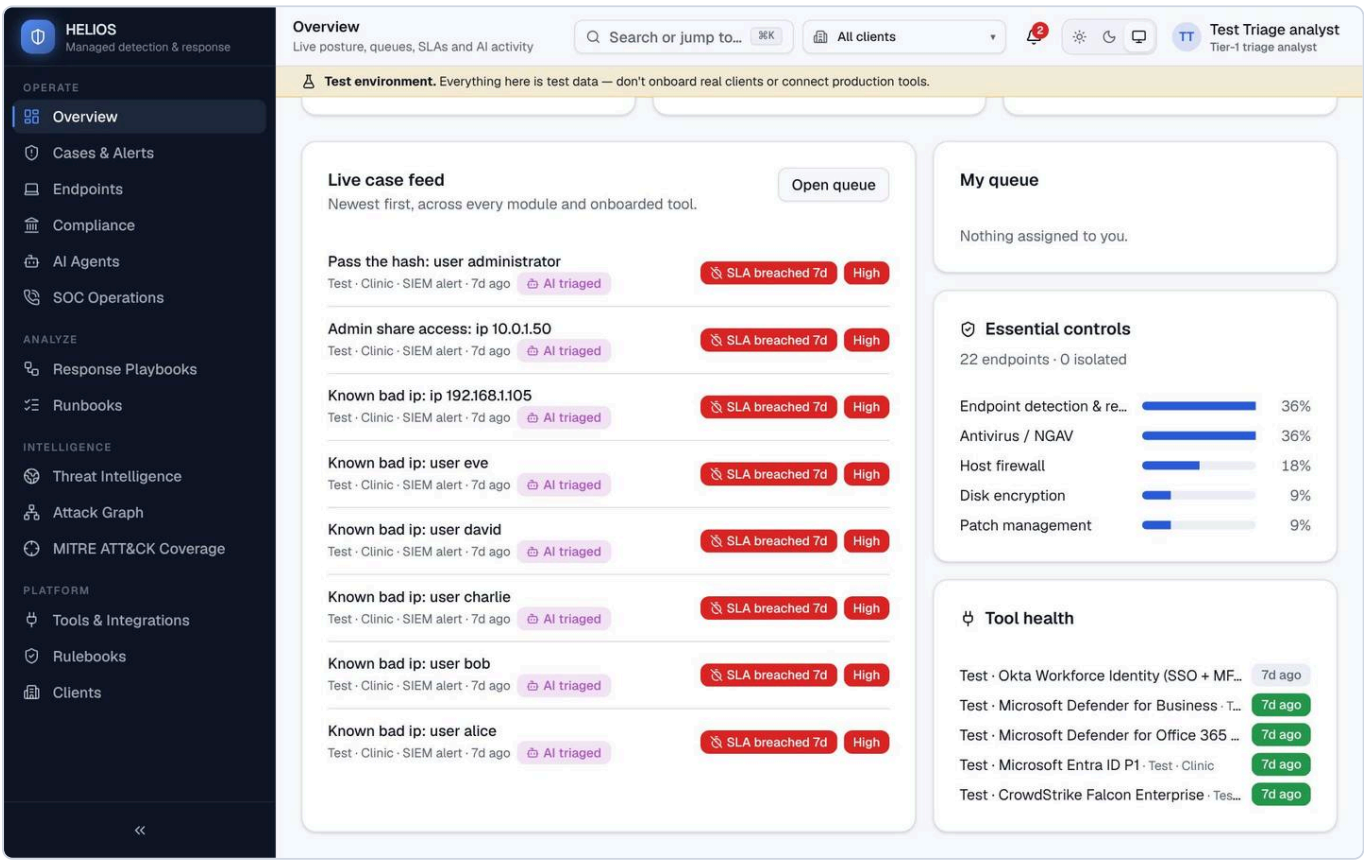


Figure 3. Overview, further down the page.

Cases & Alerts

Every case HELIOS has opened: detections from your clients' telemetry, vendor alerts, uploaded investigations, email and file analysis, and governed AI agents. Correlated detections for the same user, host or IP join one case instead of opening many.

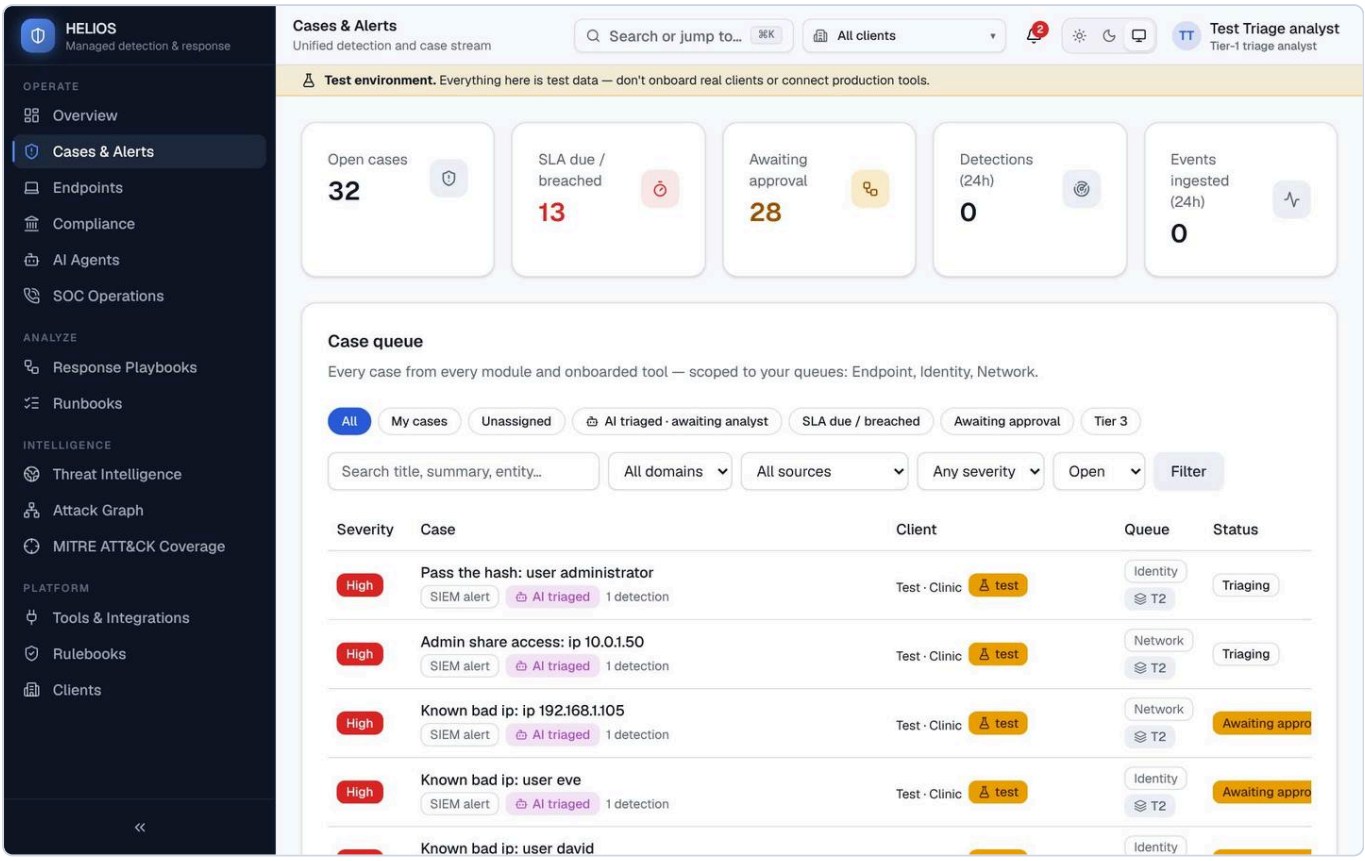


Figure 4. Cases & Alerts, as seen by a tier-1 triage analyst (test environment).

How to use it

- Filter by status, severity, source, client and assignee, or search by title. Open a case by clicking its row.
- If your account is limited to security domains (endpoint, identity, email, network, cloud, AI agents), you see only those queues.
- The SLA badge counts down to the acknowledgement deadline from the client's plan; a red badge means it has been breached.

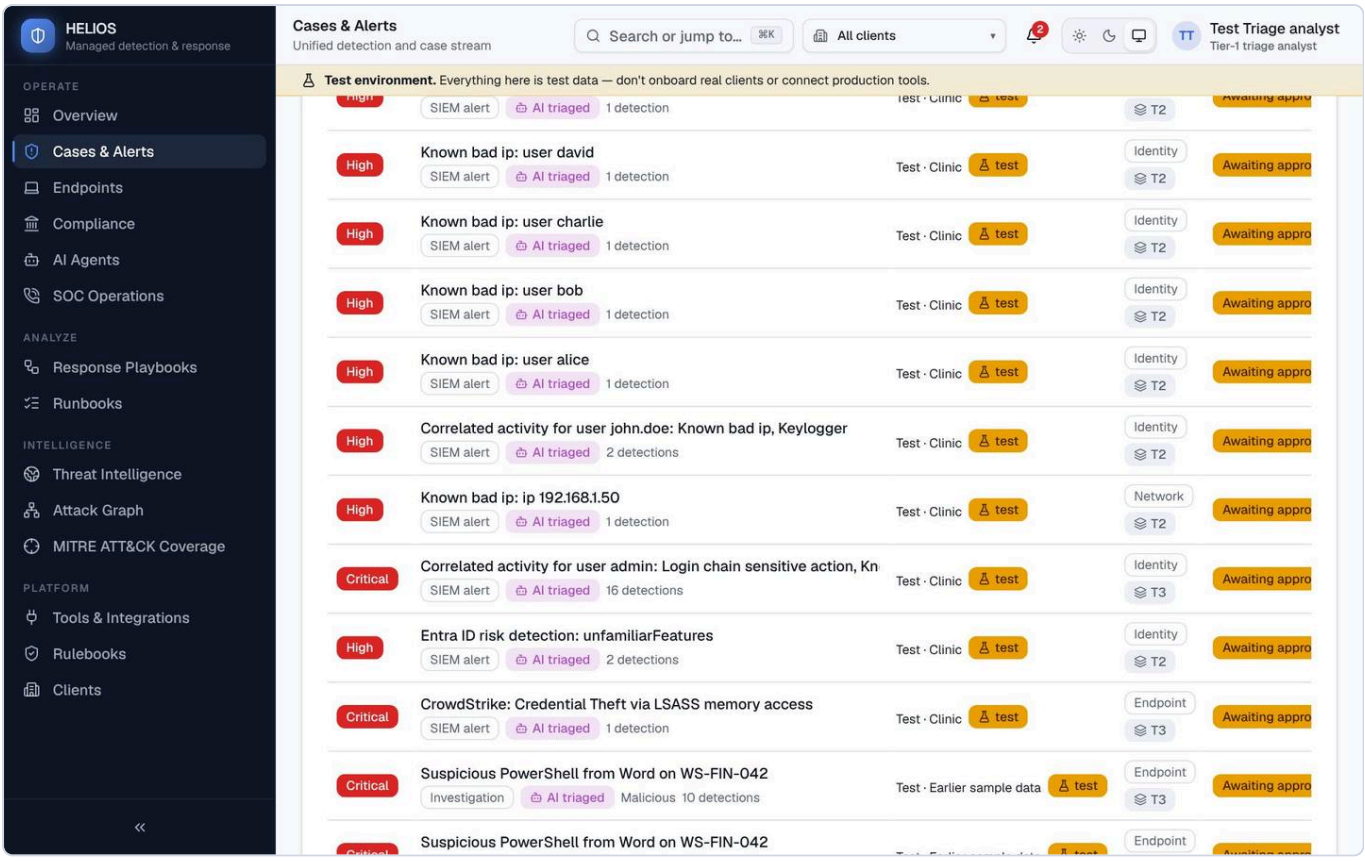


Figure 5. Cases & Alerts, further down the page.

A case

Everything about one case on one page: what was detected, what the AI made of it, what has been done

and what is waiting, the evidence, and the discussion.

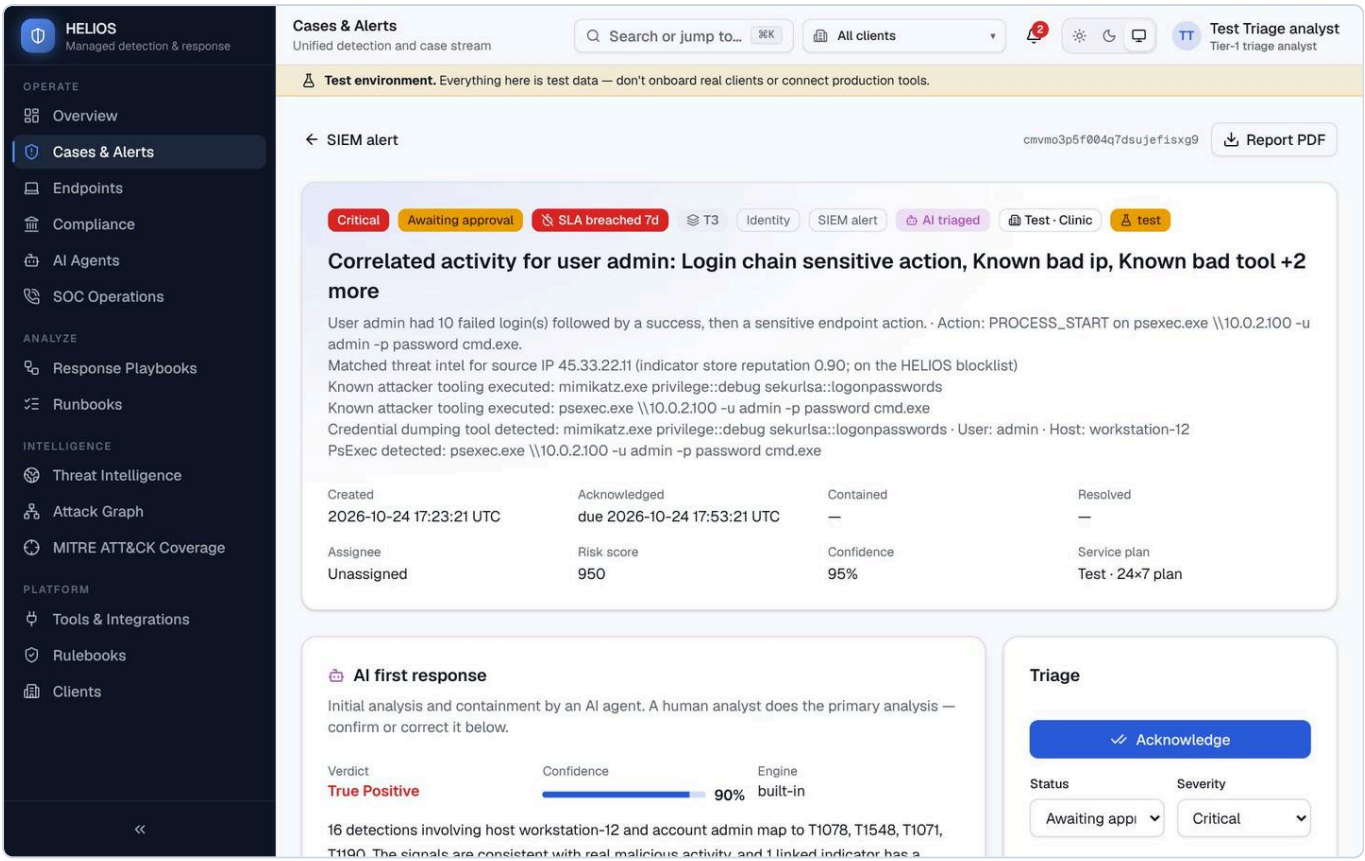


Figure 6. A case, as seen by a tier-1 triage analyst (test environment).

How to use it

- The header shows severity, status, the SLA, the tier (T1 triage, T2 analysis, T3 incident response), the security domain, the source, the client and whether AI triaged it. Below are the dates (created, acknowledged, contained, resolved), the assignee, the risk score, the confidence and the client's service plan. Report PDF downloads the incident report.
- AI first response shows the agent's verdict, confidence and engine, its summary, its key findings (each citing evidence: D detections, E events, I indicators, T timeline, C context) and what it asks the analyst to check. It is a first pass: a person confirms or corrects it.
- Agree or Disagree records whether the analysis was right; a disagreement asks what it got wrong. These ratings drive the AI's autonomy ladder.
- Dispatch agent runs another agent on the case: the triage first responder, the phishing analyst, the threat-intel enricher, the investigation agent (works the case with read-only tools) or the client report writer.
- In Triage: Acknowledge stops the SLA clock; Status and Severity change the case; Assignee gives it to someone.
- Watch case sends you its updates. Escalate hands the case to another tier or person with a note; Share makes it visible to the client or sends it to an external channel.
- Page on-call (in the On call card) pages whoever is on call about this case, with what you need from them.
- Request containment asks for an action (isolate host, block IP, disable account, kill process, quarantine file, block sender, collect logs, purge email, reset password, collect an investigation package...) against a target, with a reason shown to the approver.

- Indicators shows the case's IPs, domains, URLs, hashes and addresses with their reputation; the arrow re-checks one. Related cases lists earlier or concurrent cases sharing hosts, accounts, indicators or the same agent.
- Regulatory reporting shows whether CERT-In's 6-hour report or the DPDP breach notices apply, the incident types HELIOS proposes, and the notices' state.
- Runbooks attach themselves to matching cases; tick each step or skip it with a reason. A case can't be closed as a true positive while a required step is open.
- The discussion takes comments and @mentions. Internal comments stay with the SOC; client-visible ones reach the client.
- Closing: set the status to Resolved or False positive, then give a disposition (true positive, benign true positive, false positive) and a reason.

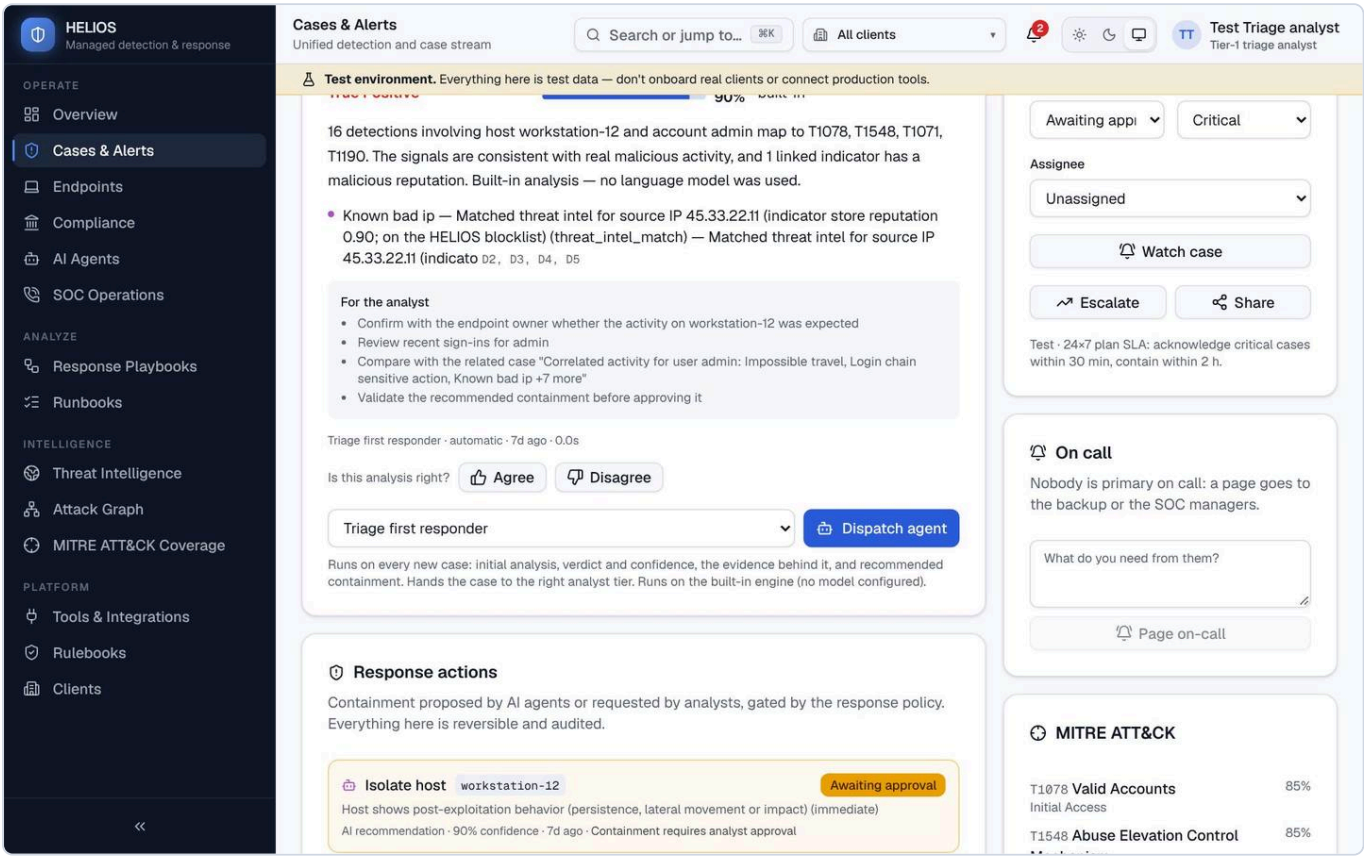


Figure 7. A case, further down the page.

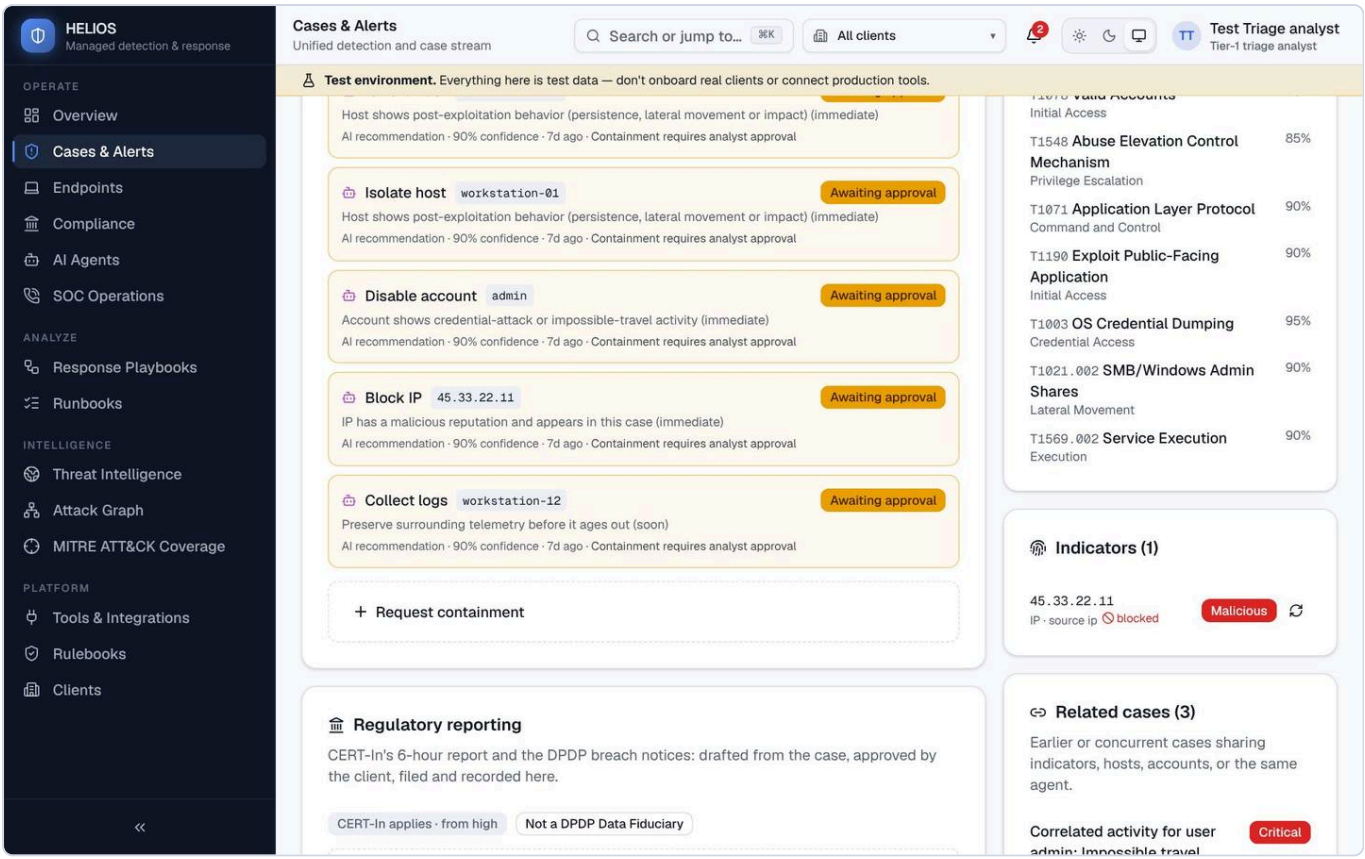


Figure 8. A case, further down the page.

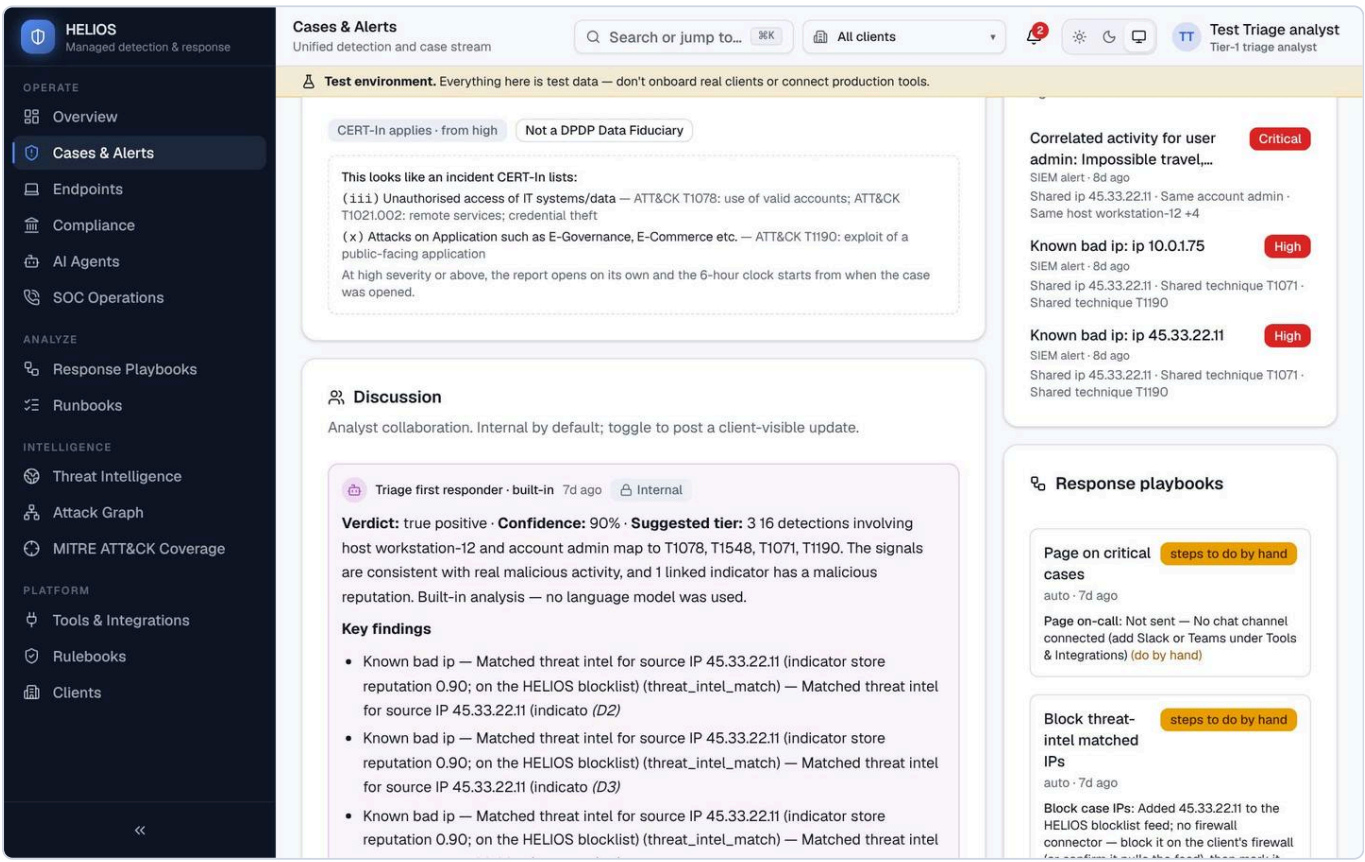


Figure 9. A case, further down the page.

Endpoints

The endpoint inventory and the whole attack surface: devices, servers, cloud, network, identity, IoT/OT and applications, with how well each is protected.

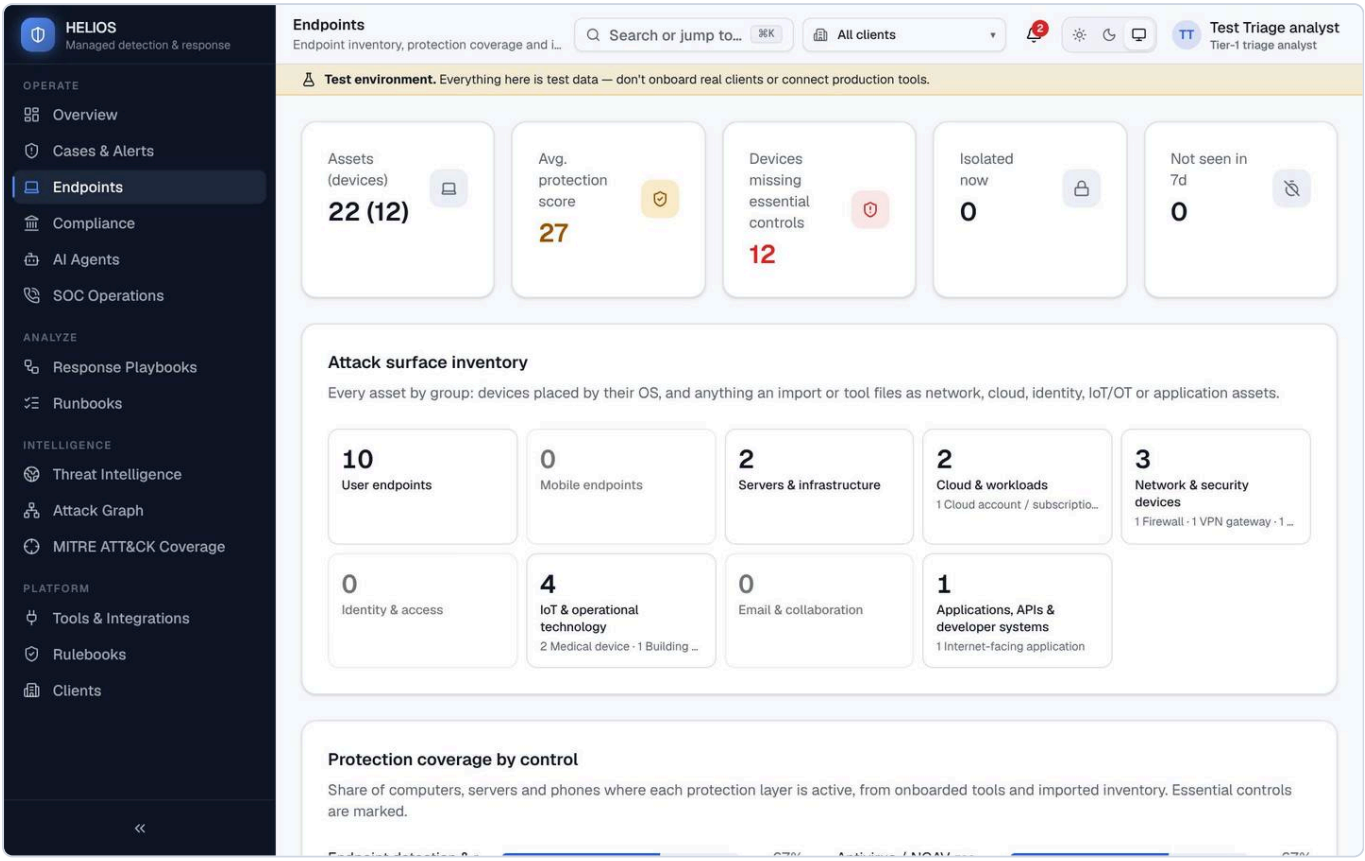


Figure 10. Endpoints, as seen by a tier-1 triage analyst (test environment).

How to use it

- The tiles count assets, the average protection score, devices missing essential controls, devices isolated now and devices not seen for seven days.
- Attack surface inventory groups every asset; Protection coverage by control shows the share of computers, servers and phones where each protection layer (EDR, antivirus, disk encryption, MFA...) is active, with the essential ones marked.

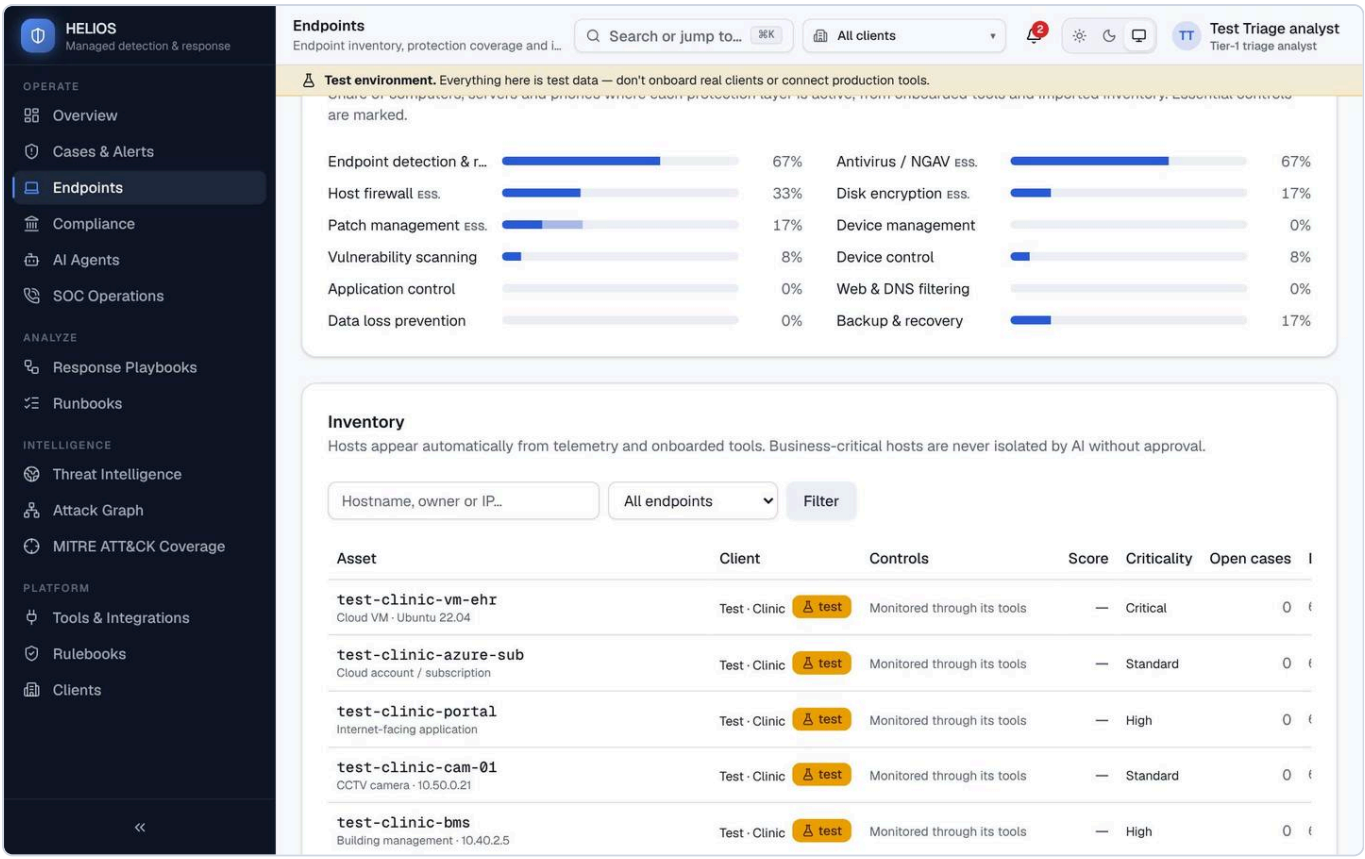


Figure 11. Endpoints, further down the page.

Compliance

Regulatory clocks and evidence: CERT-In's 6-hour incident reports, the DPDP breach notices, and the

ISO/IEC 27001 evidence map.

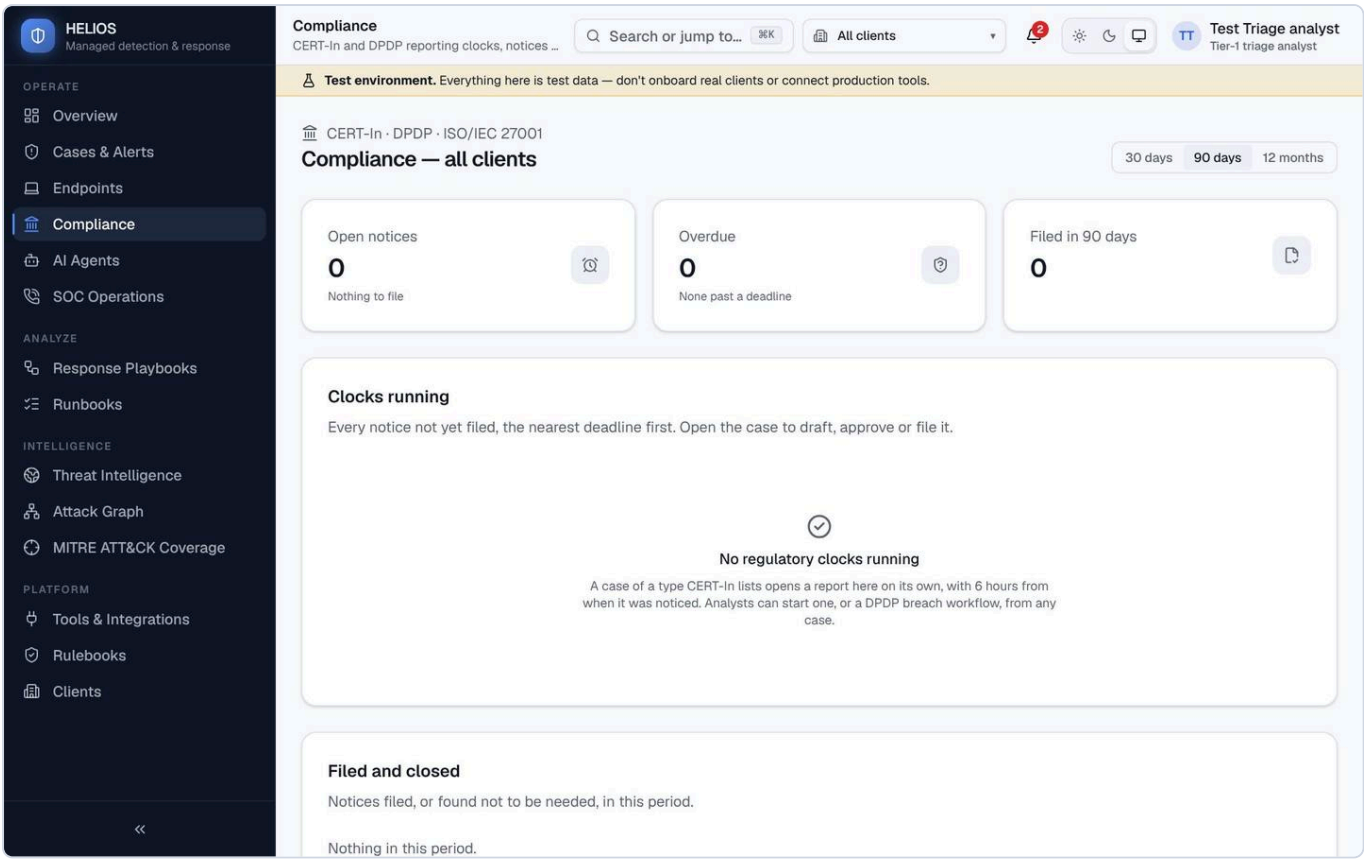


Figure 12. Compliance, as seen by a tier-1 triage analyst (test environment).

How to use it

- Clocks running lists every notice not yet filed, nearest deadline first. CERT-In: incidents of the 20 listed types are reported within 6 hours of being noticed. DPDP: the Board and each affected person are told without delay, and the Board gets a detailed report within 72 hours.
- The ISO 27001 view maps 15 Annex A controls to what HELIOS's records show over the period, and shows gaps (controls with nothing behind them) plainly.
- Switch the period (30 days, 90 days, 12 months) at the top.

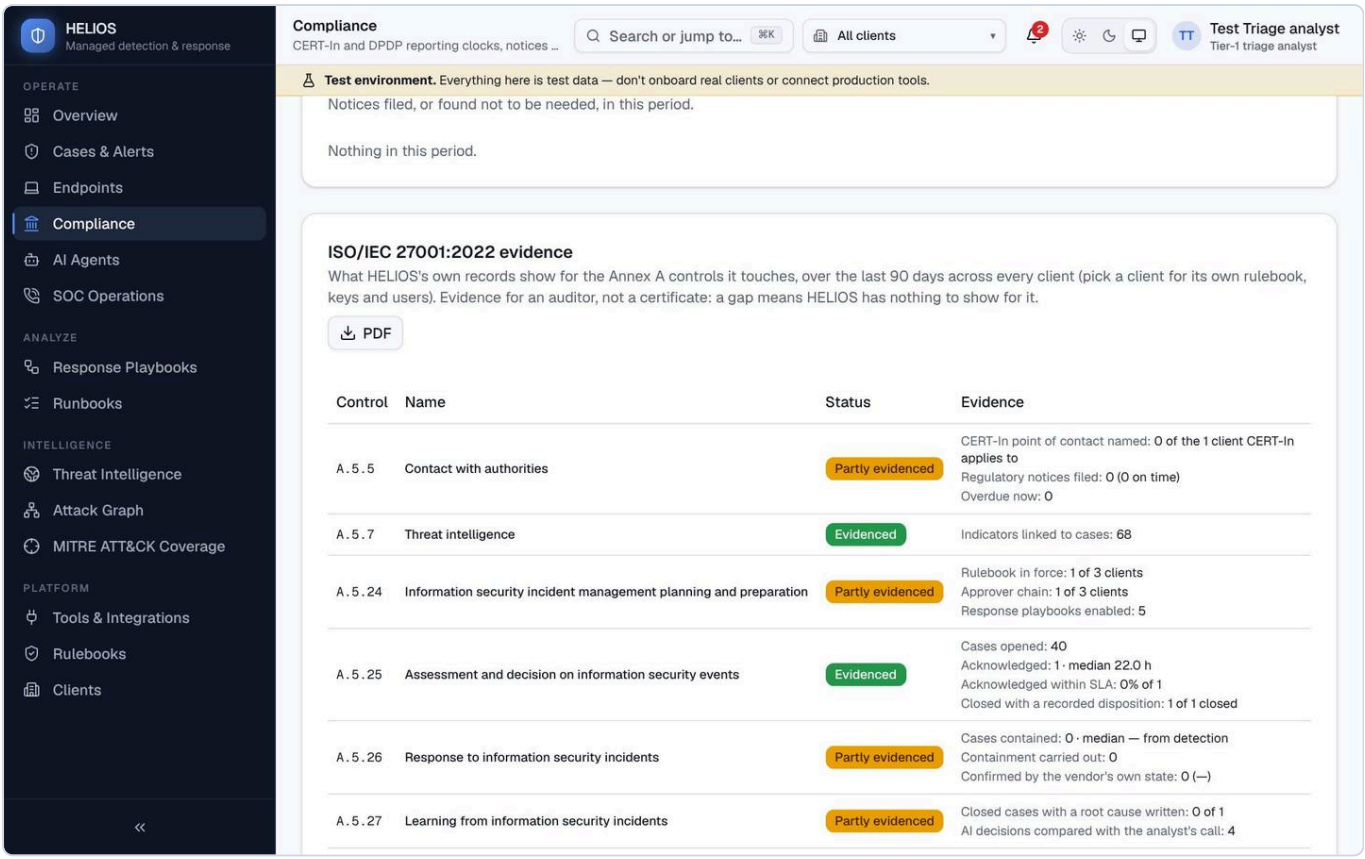


Figure 13. Compliance, further down the page.

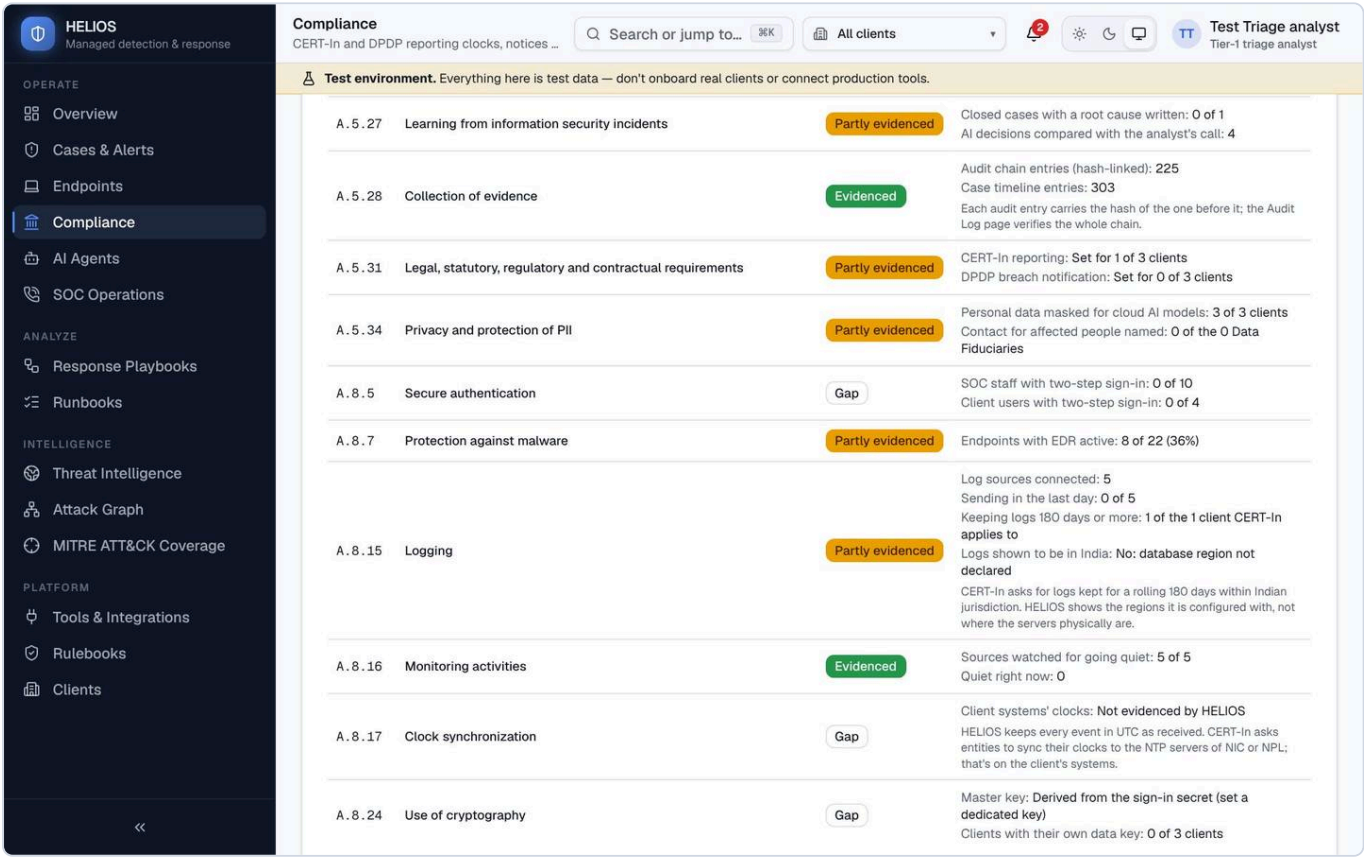


Figure 14. Compliance, further down the page.

AI Agents

The AI agents that work cases: what they did, how far they may act on their own, which models run them, and how those models scored.

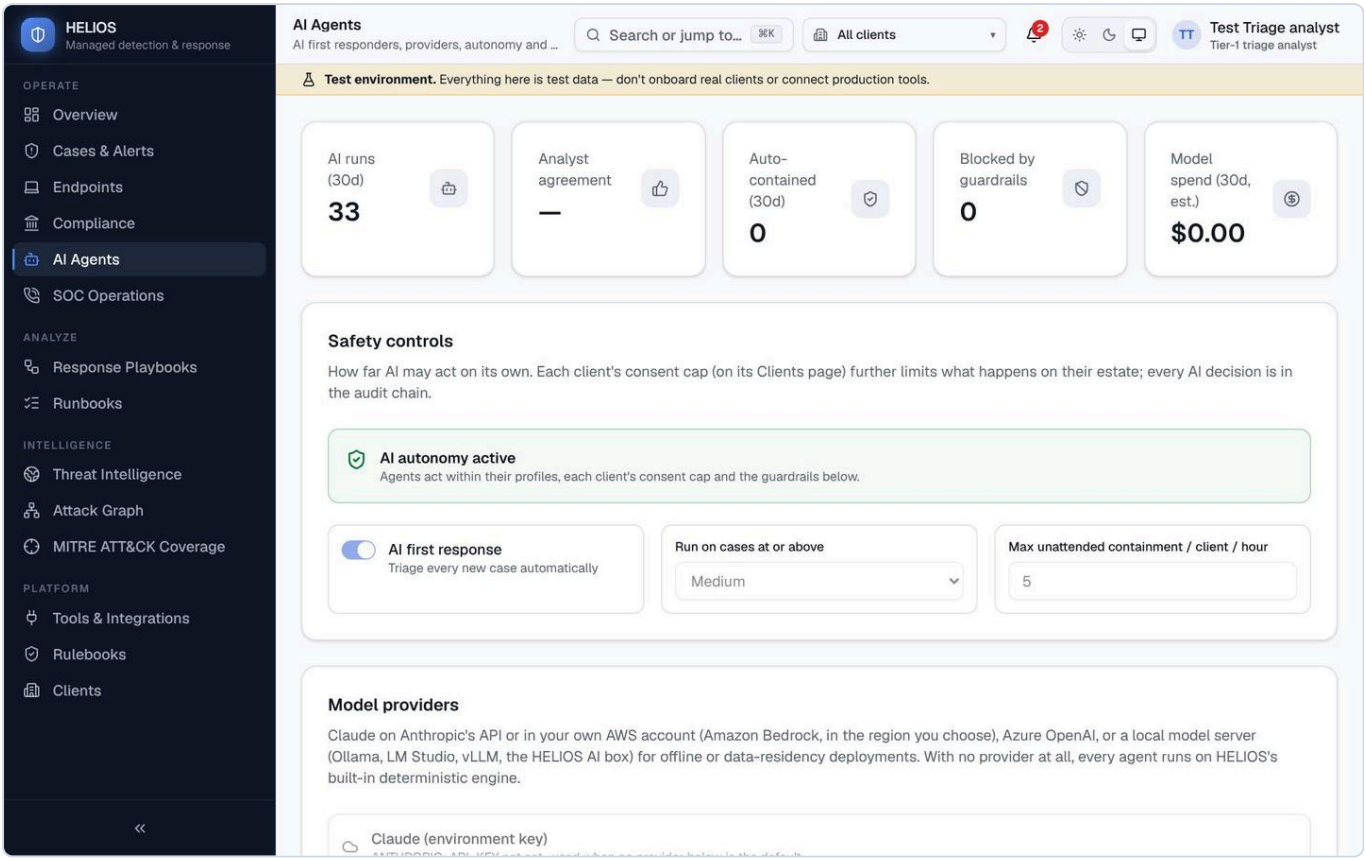


Figure 15. AI Agents, as seen by a tier-1 triage analyst (test environment).

How to use it

- The tiles show AI runs in the last 30 days, analyst agreement, actions contained automatically, actions the guardrails blocked and the estimated model spend (with the share of input read from the prompt cache).
- The run log lists recent runs with their verdict, model, tokens and outcome.

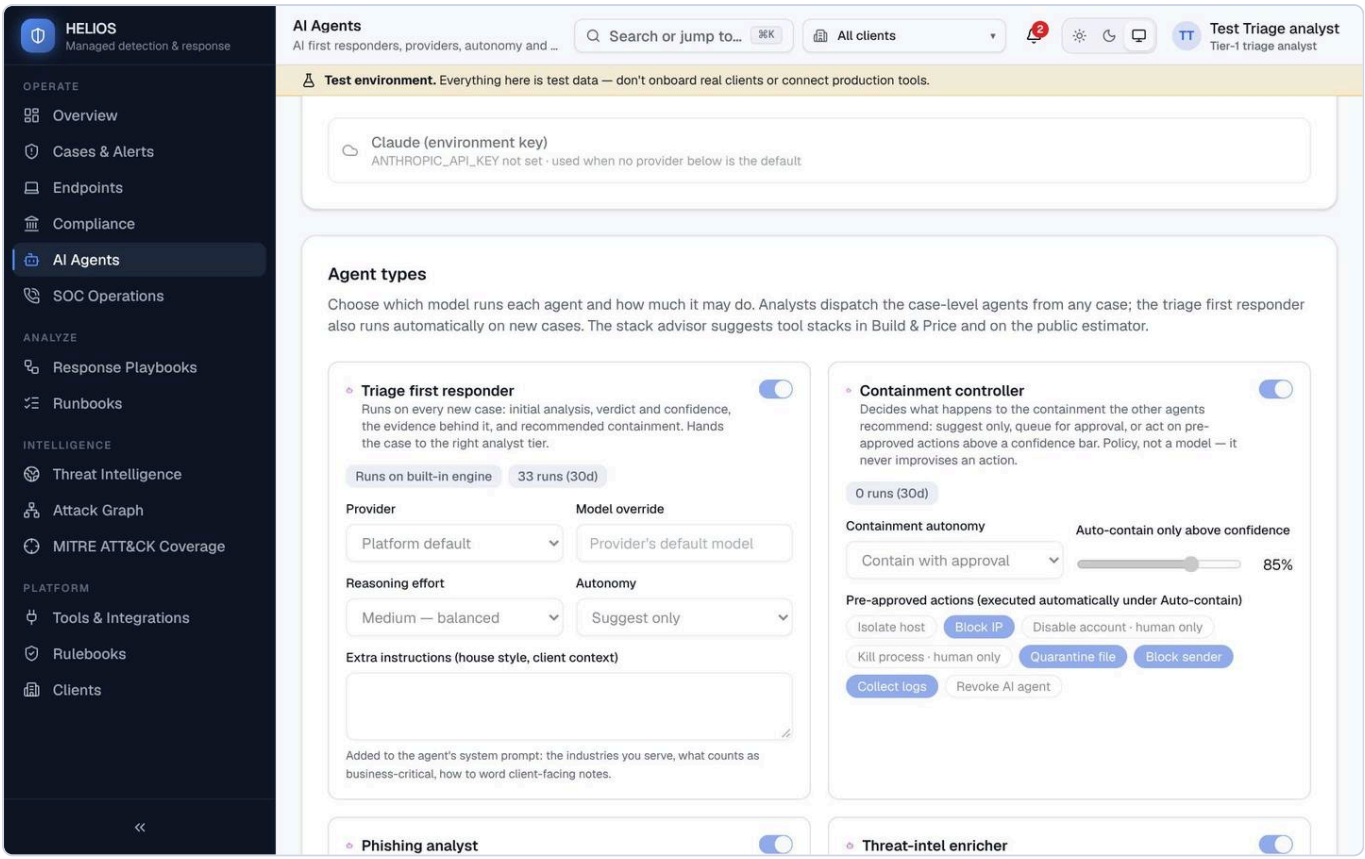


Figure 16. AI Agents, further down the page.

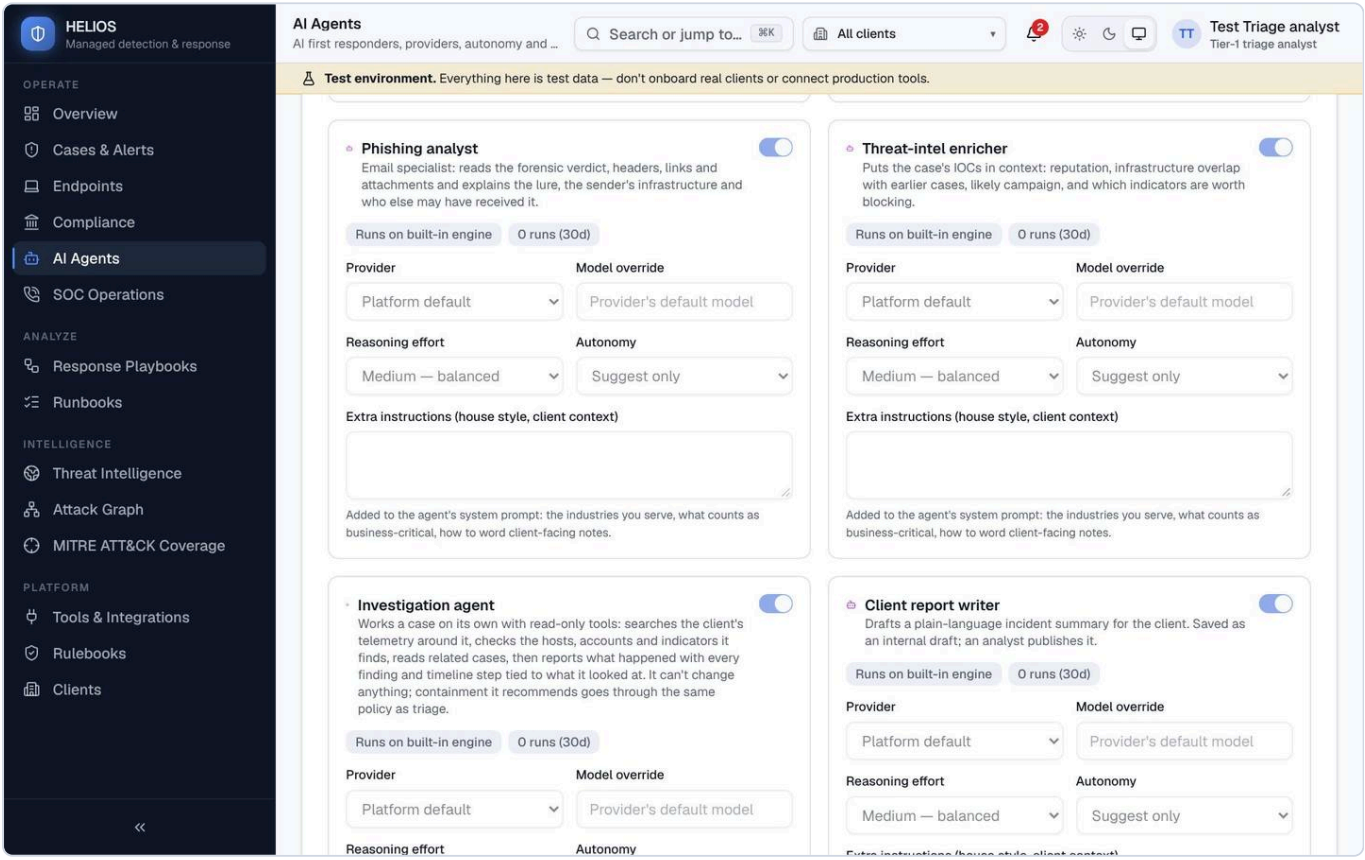


Figure 17. AI Agents, further down the page.

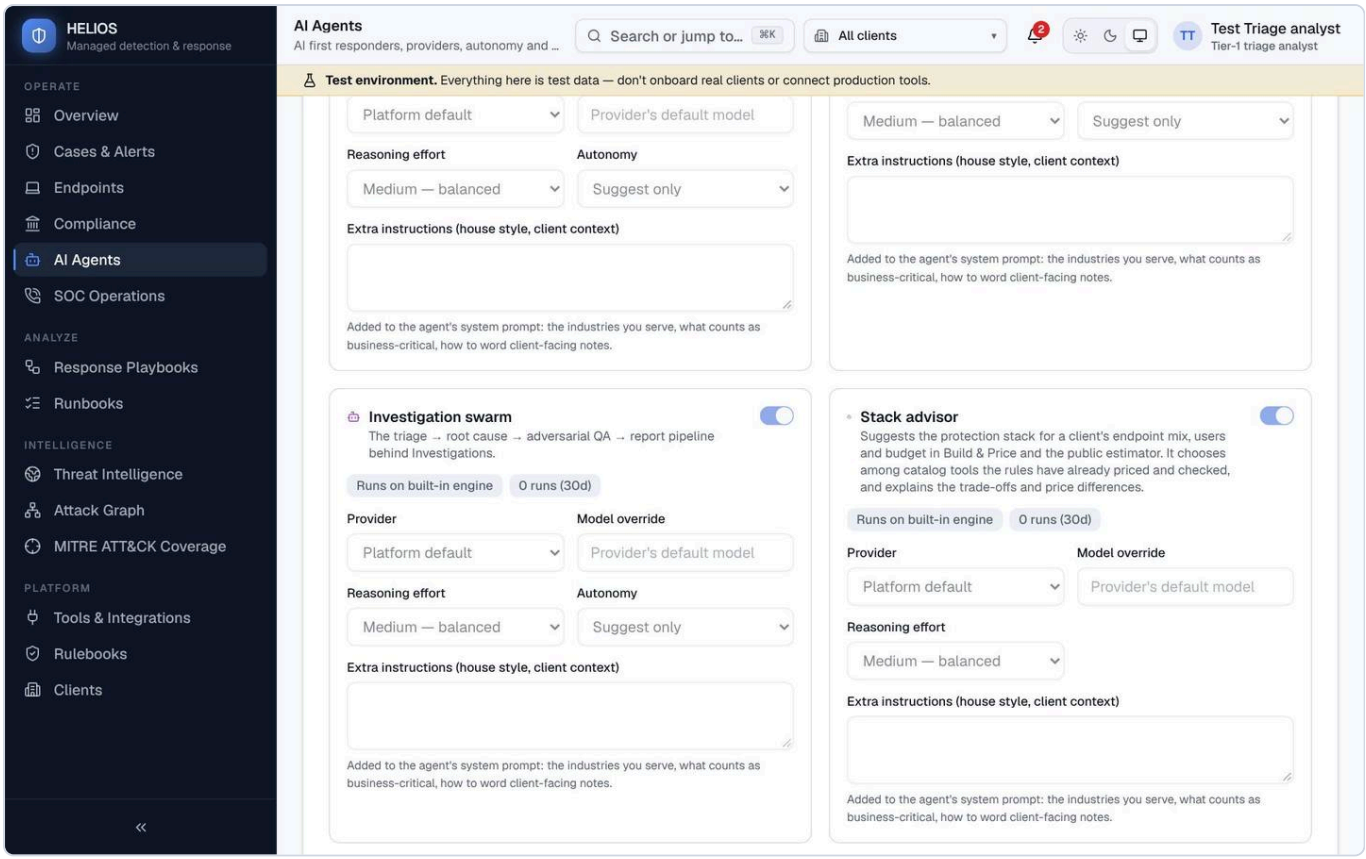


Figure 18. AI Agents, further down the page.

SOC Operations

Who is on call and the pages they get, the weekly QA of what the AI did on its own, and the daily digests.

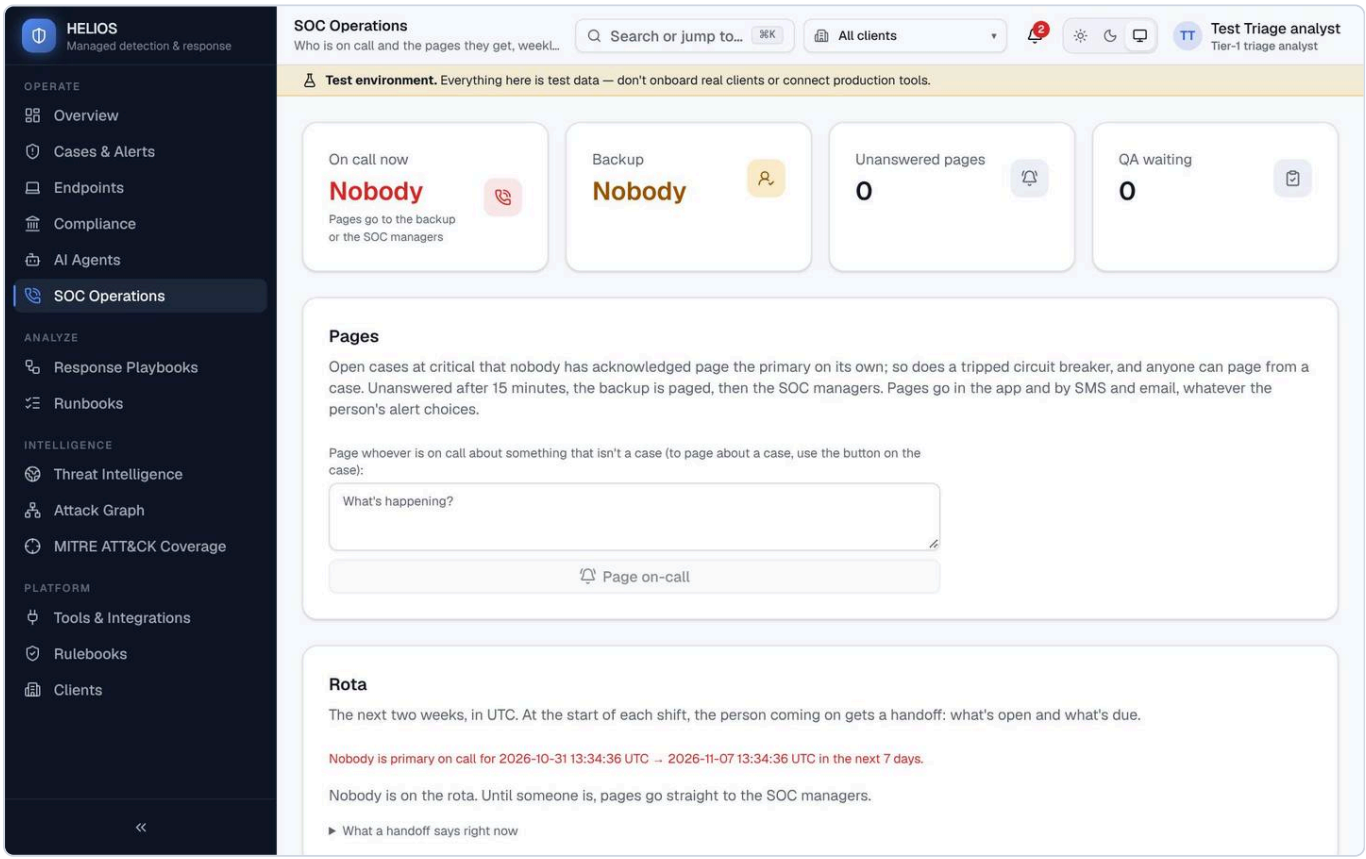


Figure 19. SOC Operations, as seen by a tier-1 triage analyst (test environment).

How to use it

- The tiles show who is primary and backup on call now, unanswered pages and QA reviews waiting.
- Pages: open critical cases nobody acknowledged page the primary on their own; unanswered after the acknowledgement window (15 minutes by default) the backup is paged, then the SOC managers. Pages go in the app and by SMS and email. Page on-call sends a page about something that isn't a case.

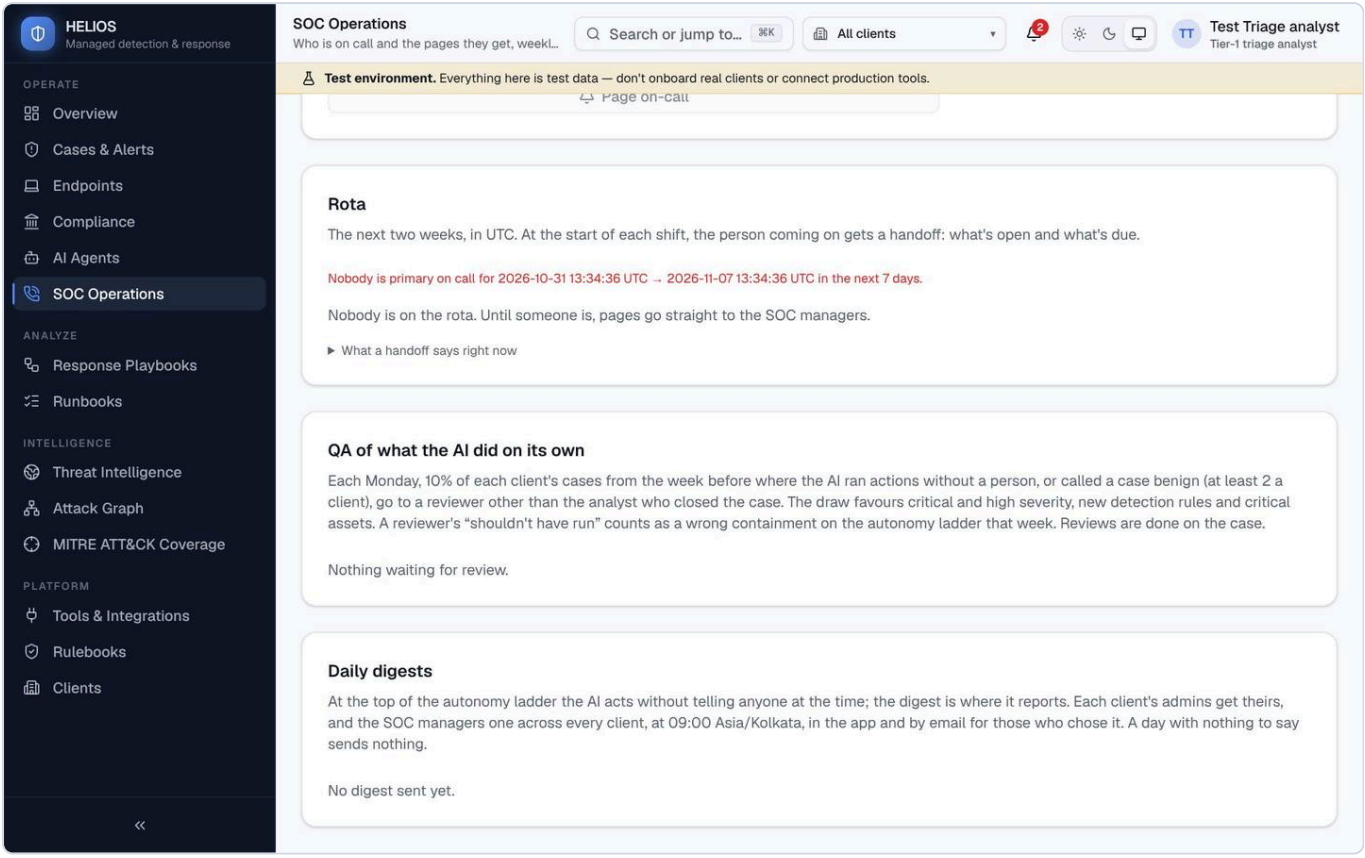


Figure 20. SOC Operations, further down the page.

Response Playbooks

Automated response workflows: what triggers them and the steps they run (containment, collecting logs,

notifying, opening tickets).

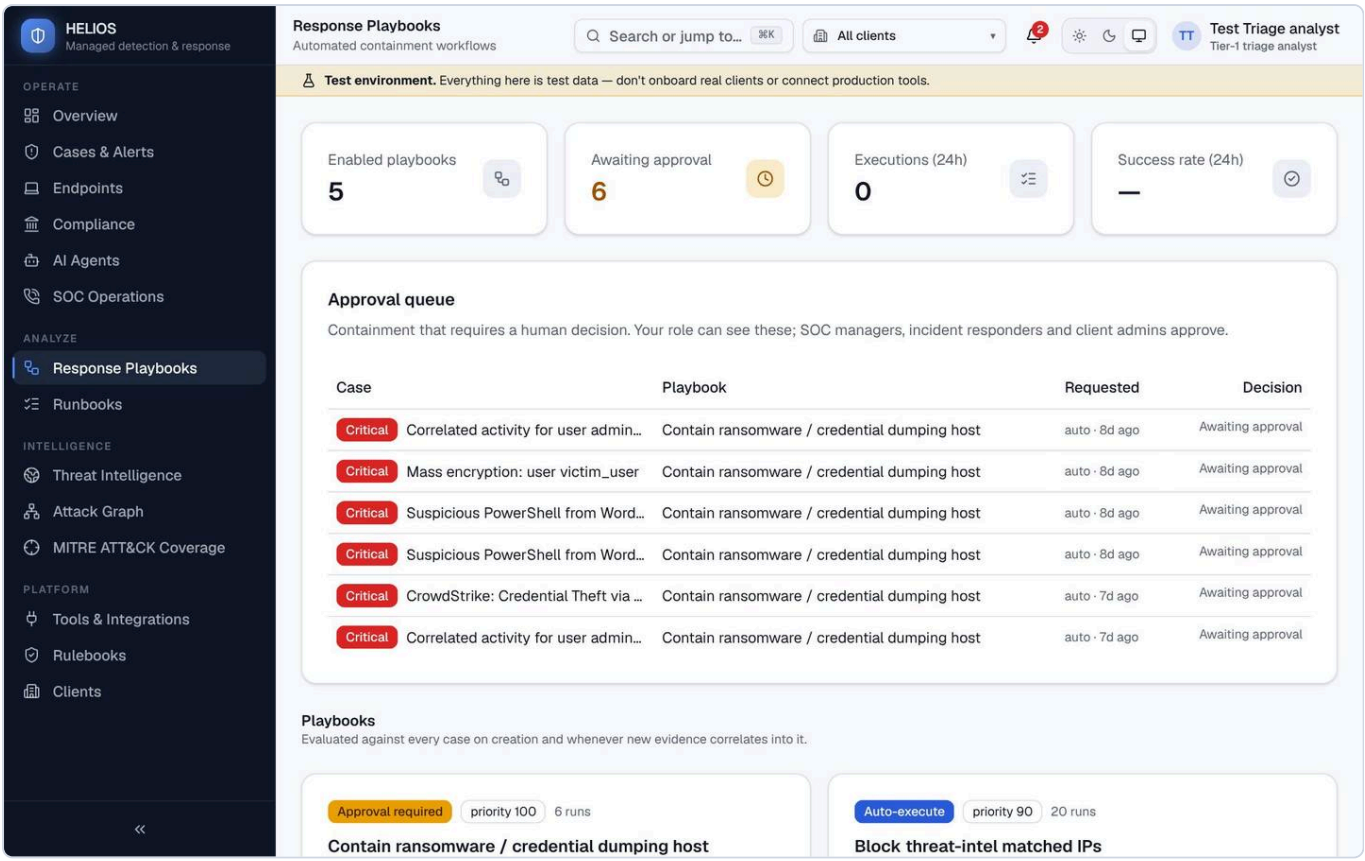


Figure 21. Response Playbooks, as seen by a tier-1 triage analyst (test environment).

How to use it

- Each playbook shows its trigger, its steps and whether it is enabled. Its runs appear on the cases they ran on; steps that need approval wait in Approvals, and steps without a connector are left to do by hand.

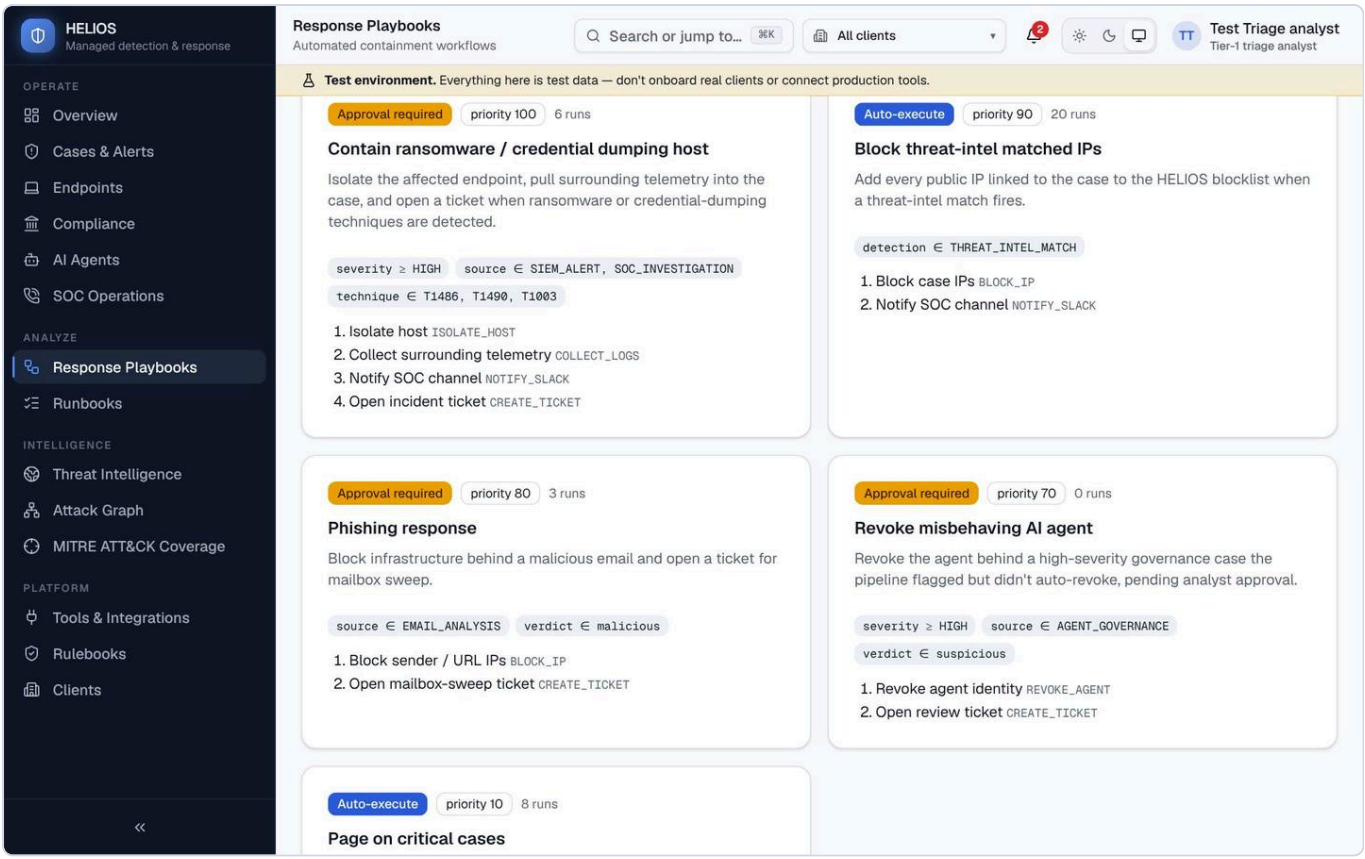


Figure 22. Response Playbooks, further down the page.

Runbooks

Procedures for analysts, attached to the cases they match and worked as checklists on the case.

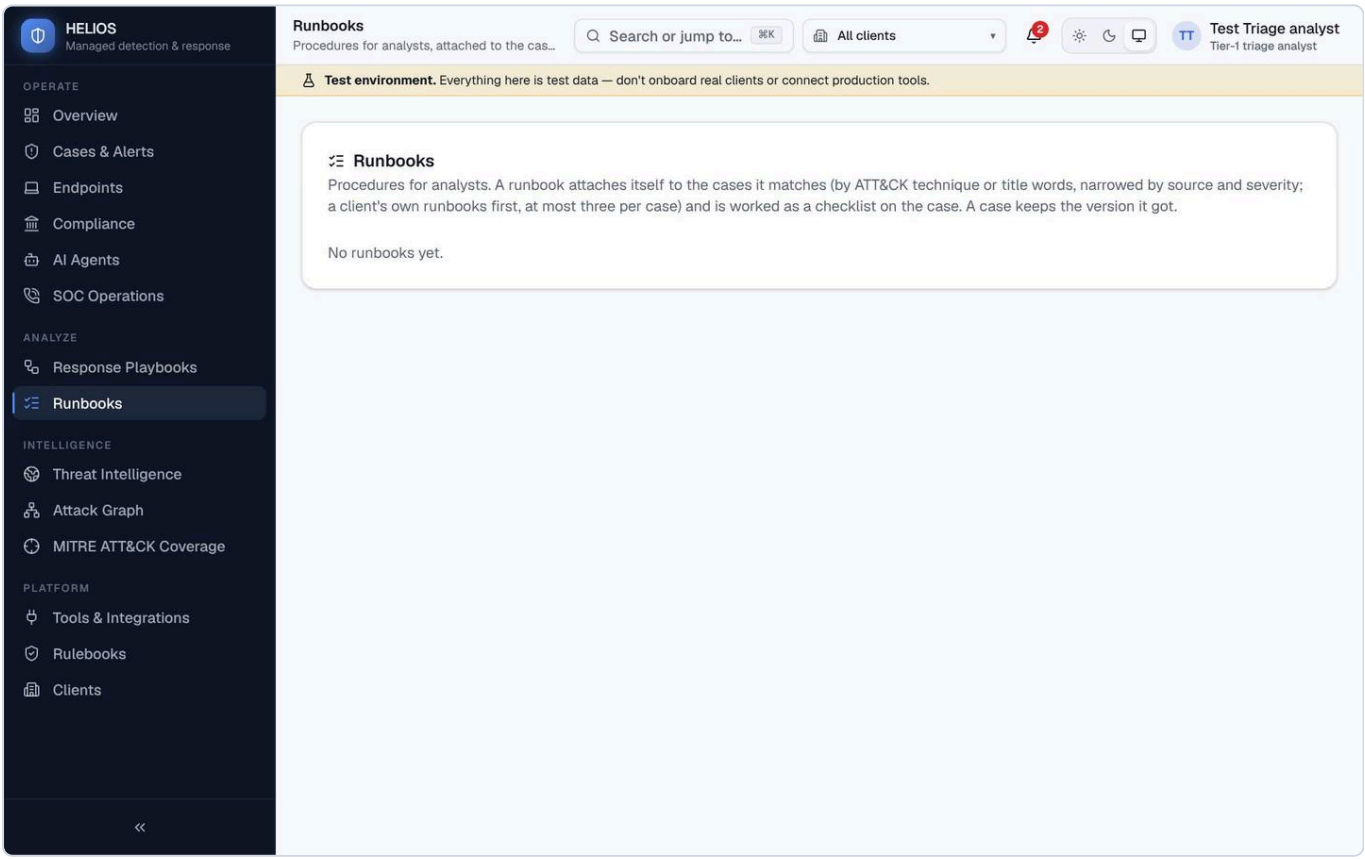


Figure 23. Runbooks, as seen by a tier-1 triage analyst (test environment).

How to use it

- A runbook matches cases by ATT&CK technique or title words, narrowed by source and minimum severity. Up to three matching runbooks attach to a case when it is opened, the client's own first.

Threat Intelligence

Look up indicators (IP, domain, URL, file hash, email address) against the SOC's indicator store, the configured reputation providers and the threat feeds.

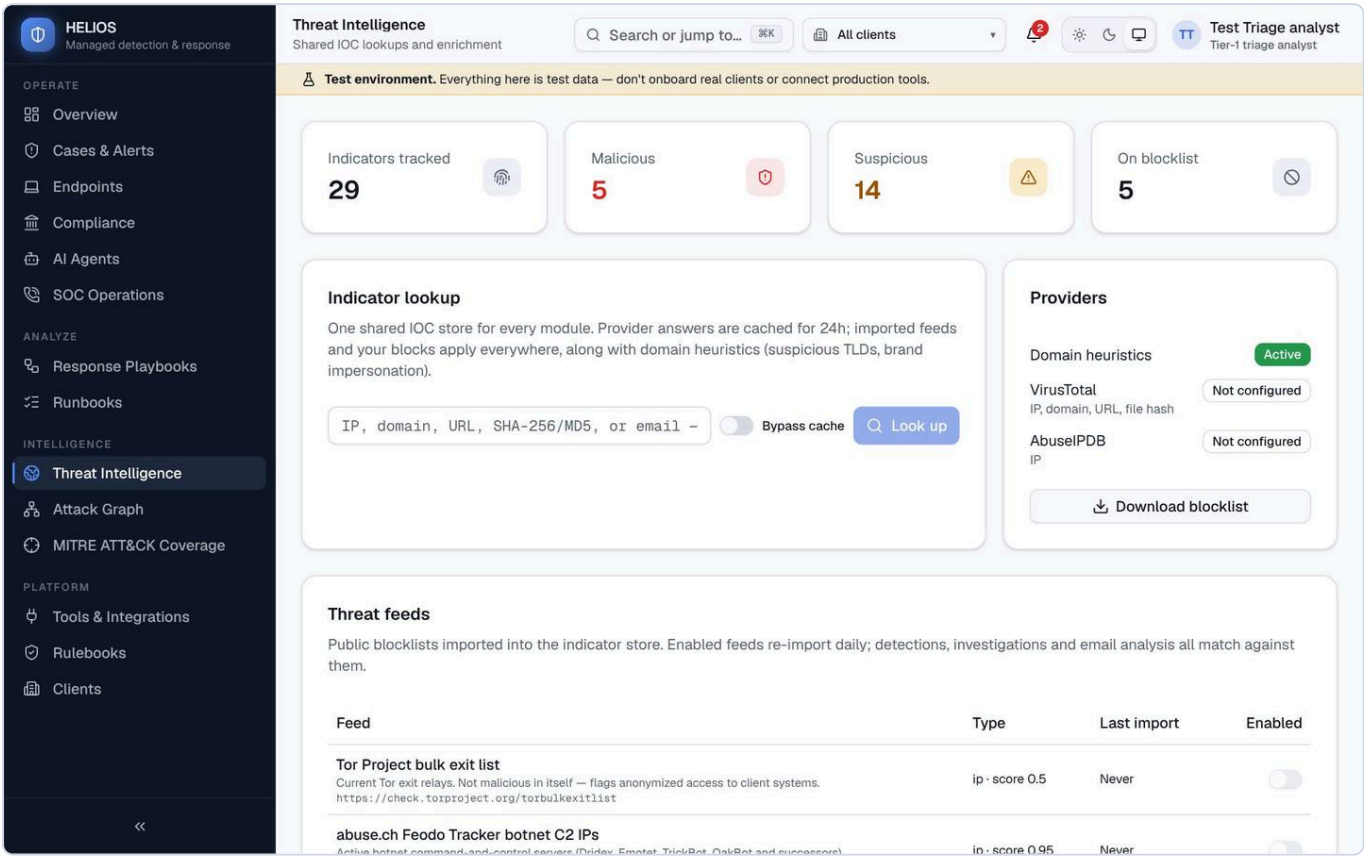


Figure 24. Threat Intelligence, as seen by a tier-1 triage analyst (test environment).

How to use it

- Paste an indicator (defanged forms like hxxp and [.] are fine) and Look up. The result shows the verdict, score, sources, why, and for an IP where it is (country, city and network, from the server's GeoIP database). Results are cached for 24 hours; Bypass cache forces a fresh lookup.

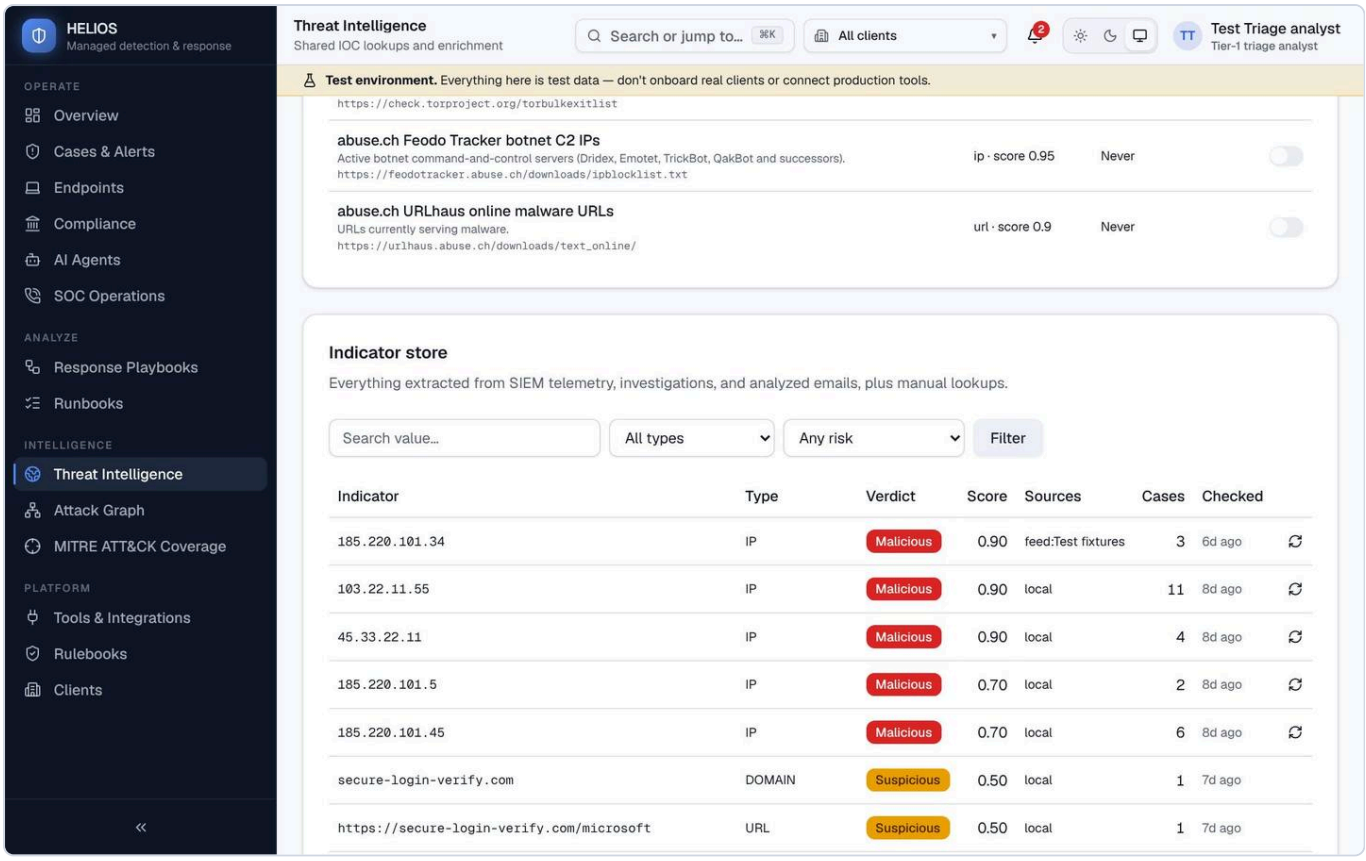


Figure 25. Threat Intelligence, further down the page.

Attack Graph

Cases joined by the hosts, accounts, addresses, indicators and agents they share. An entity is an attack

chain when its cases together reach more ATT&CK tactics than any one alone.

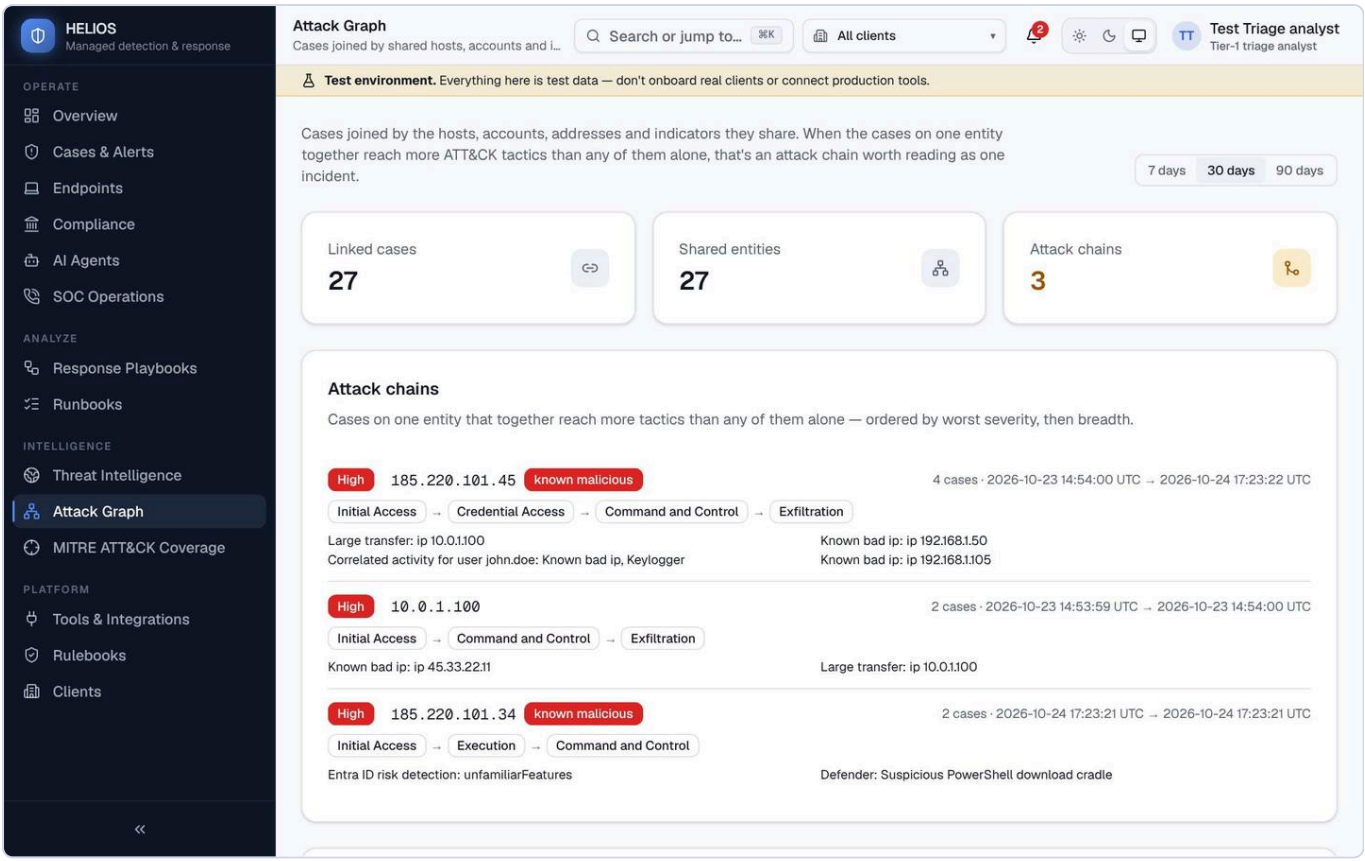


Figure 26. Attack Graph, as seen by a tier-1 triage analyst (test environment).

How to use it

- Hover or click a node to see the cases behind it; follow chains across tactics to see an intrusion that several smaller cases make up.

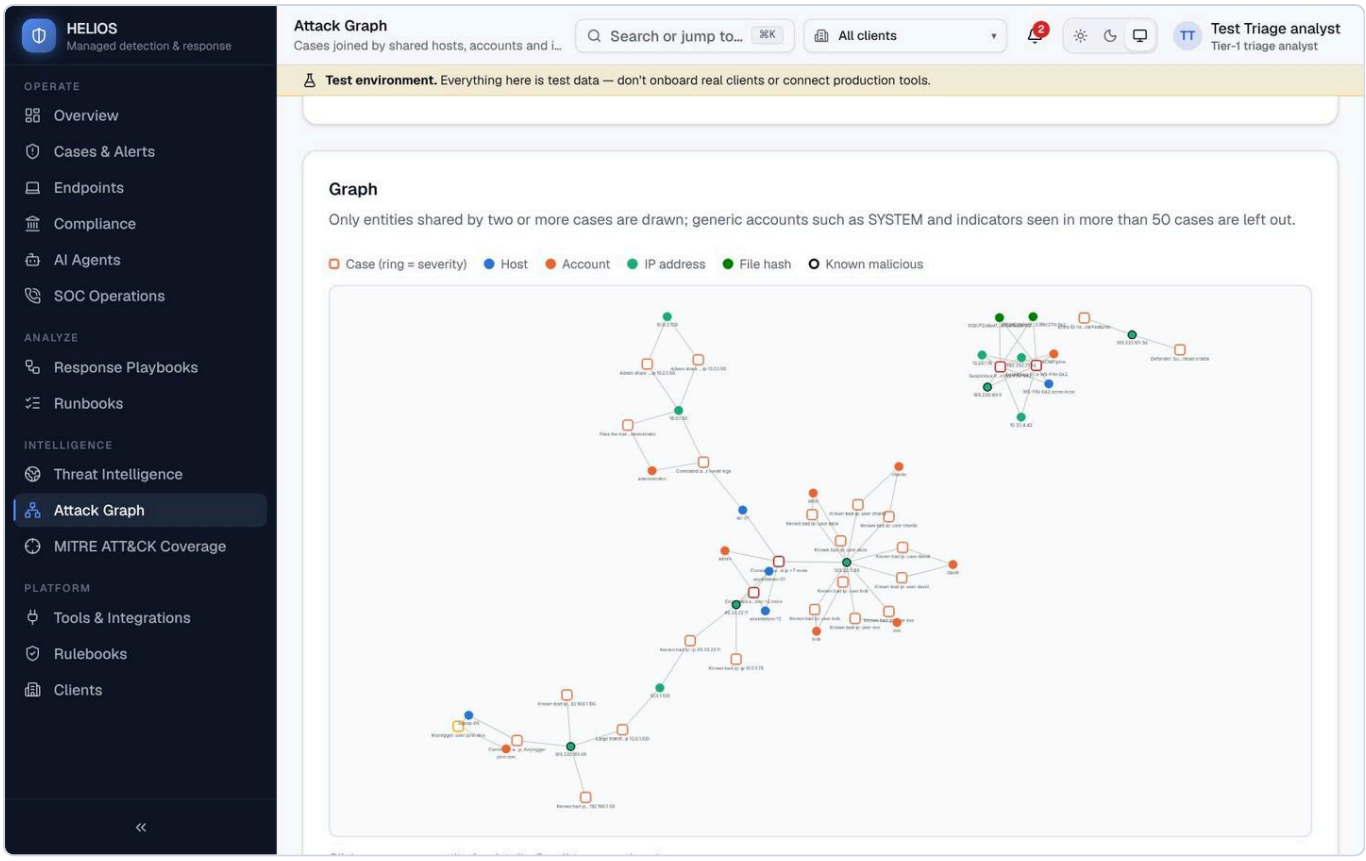


Figure 27. Attack Graph, further down the page.

MITRE ATT&CK Coverage

Which ATT&CK techniques HELIOS can detect, across built-in detectors, Sigma rules and onboarded

tools, and which it has seen.

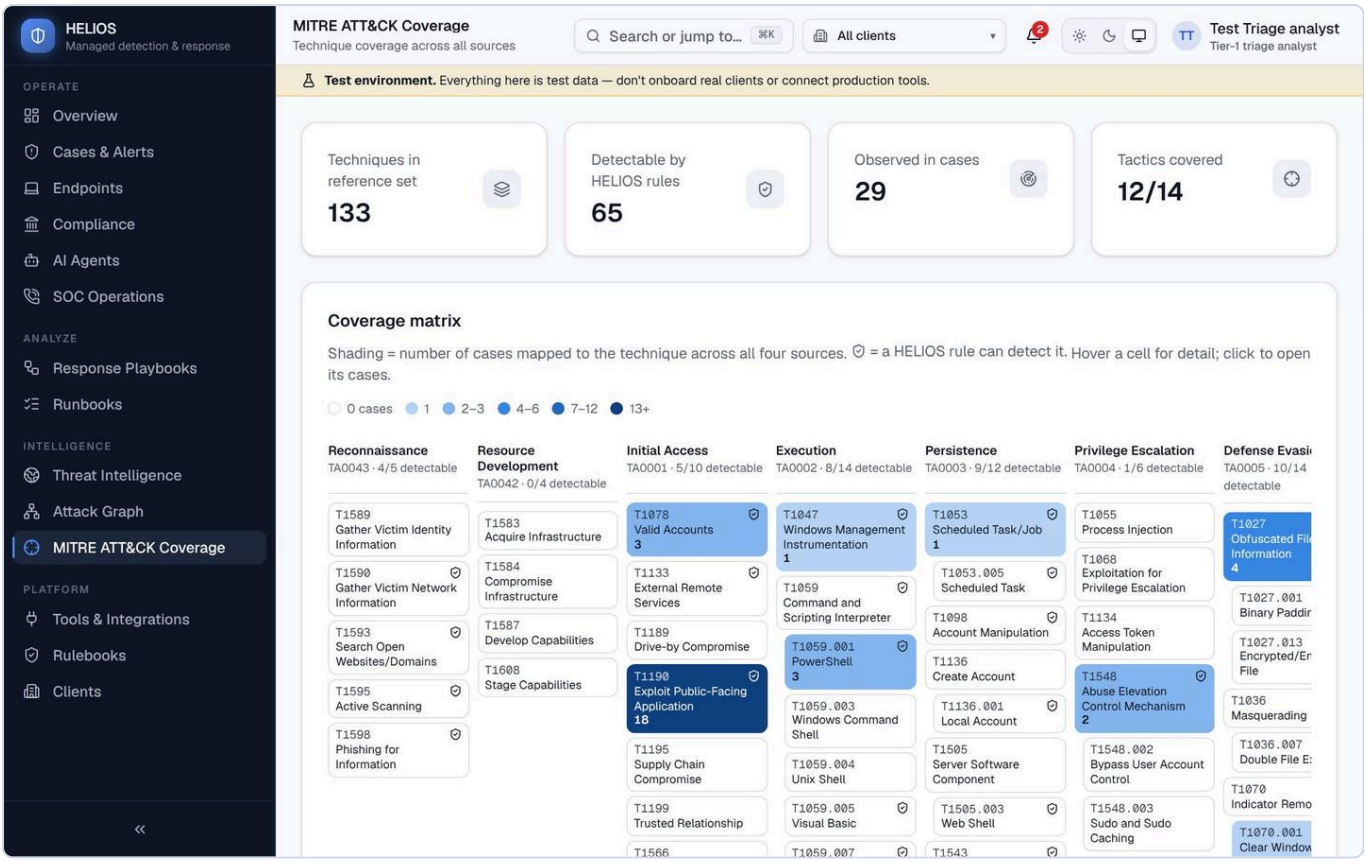


Figure 28. MITRE ATT&CK Coverage, as seen by a tier-1 triage analyst (test environment).

How to use it

- Read coverage by tactic and technique; gaps show where detection is missing.

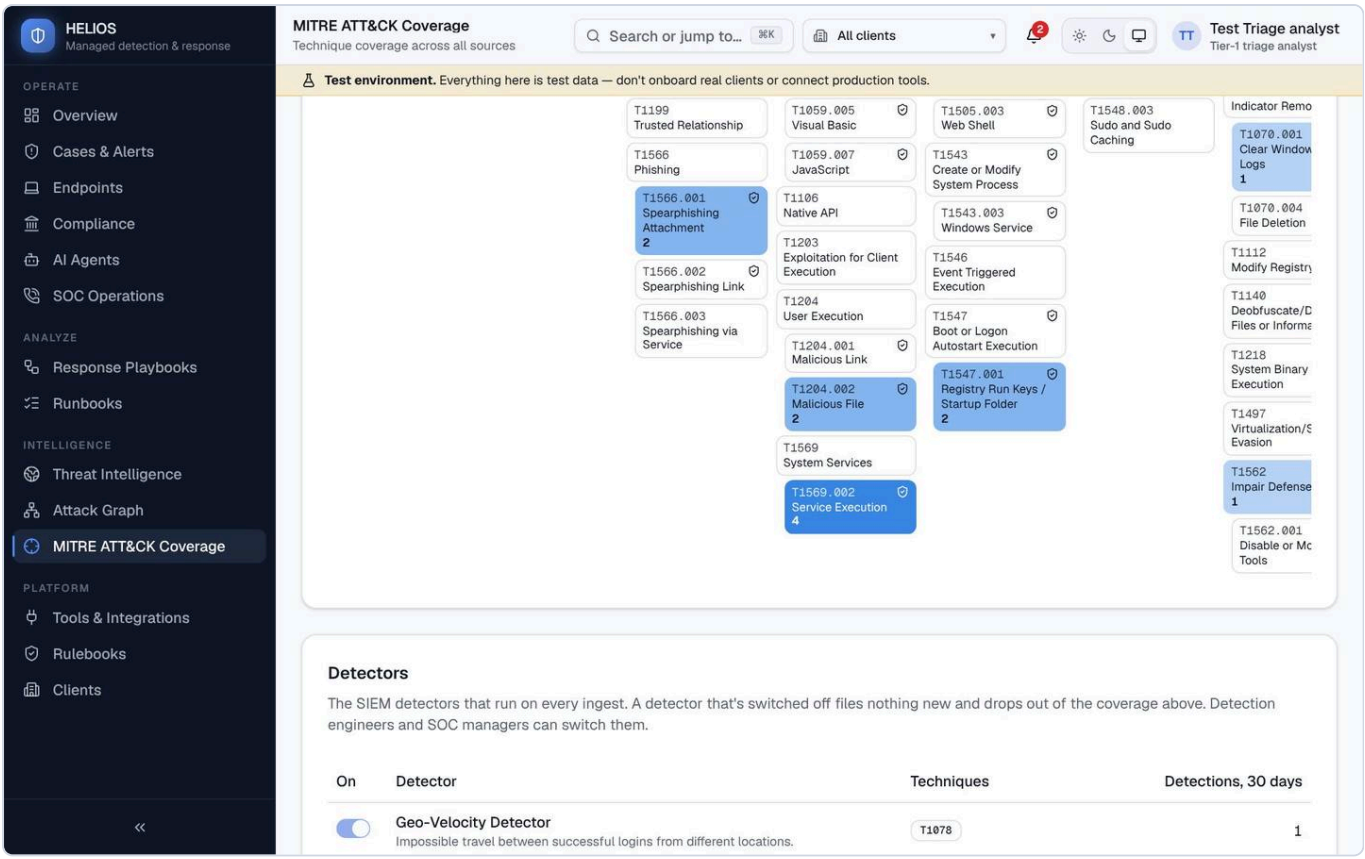


Figure 29. MITRE ATT&CK Coverage, further down the page.

Tools & Integrations

The security tools onboarded for each client: how their logs reach HELIOS and how HELIOS acts through

them.

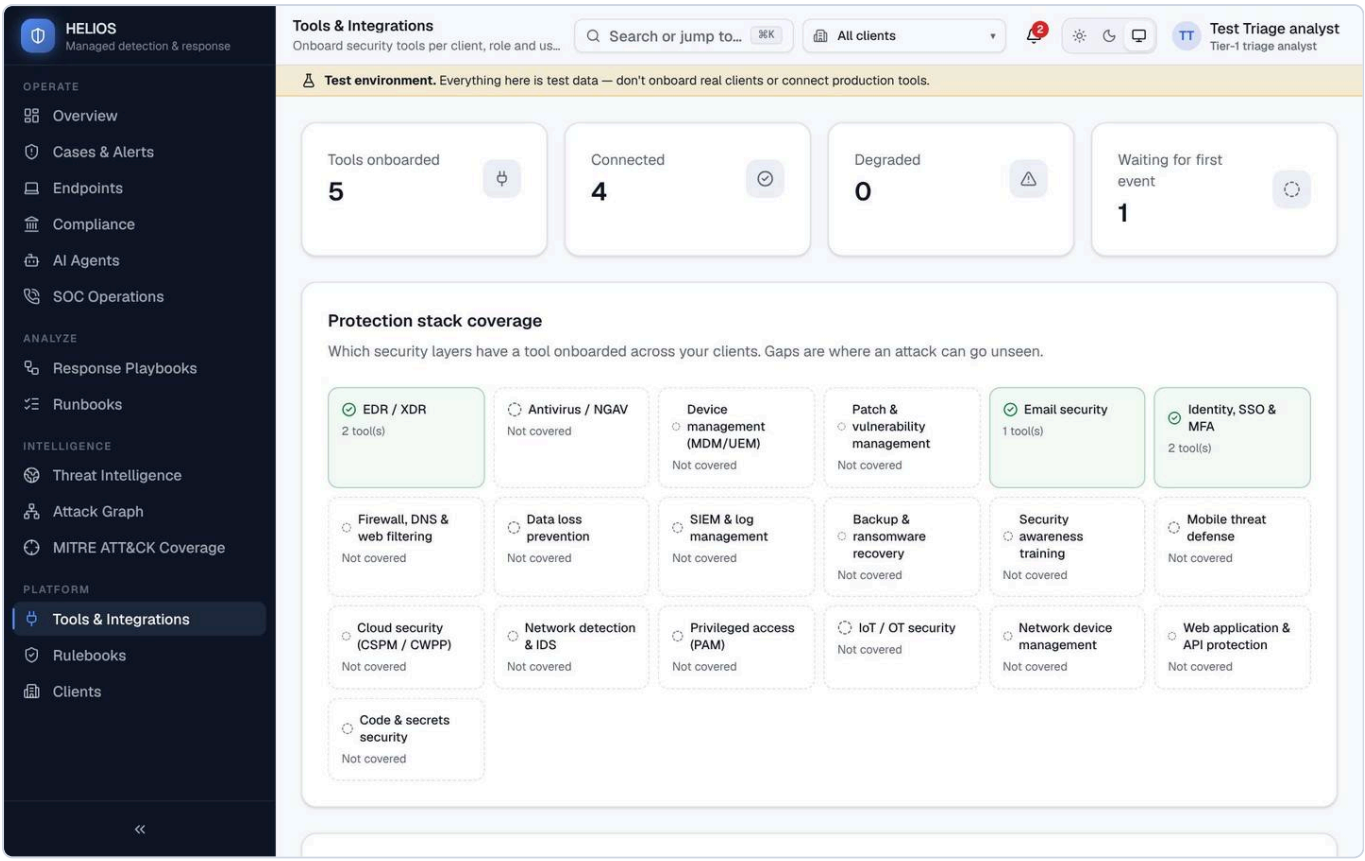


Figure 30. Tools & Integrations, as seen by a tier-1 triage analyst (test environment).

How to use it

- The tiles count tools onboarded, connected, degraded and waiting for their first event. Protection stack coverage shows which security layers have a tool across your clients; gaps are where an attack can go unseen.
- Each tool shows how it sends logs (push to its ingest URL, HELIOS pulling from the vendor's API, or the HELIOS forwarder), its source health, and whether a response connector can act through it.
- A source silent longer than its limit opens a case: attackers often stop logging first. A client can't go live while a source is dead.

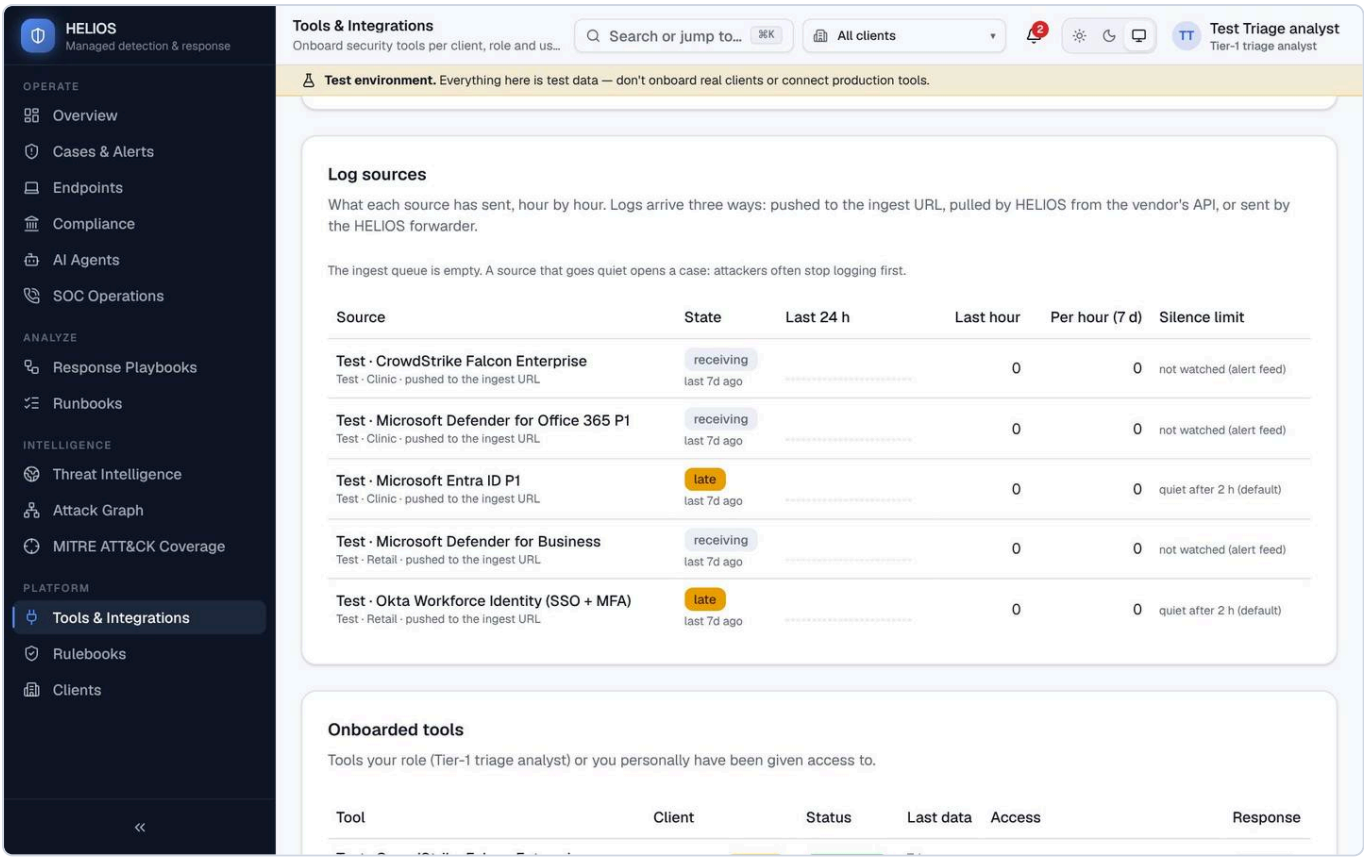


Figure 31. Tools & Integrations, further down the page.

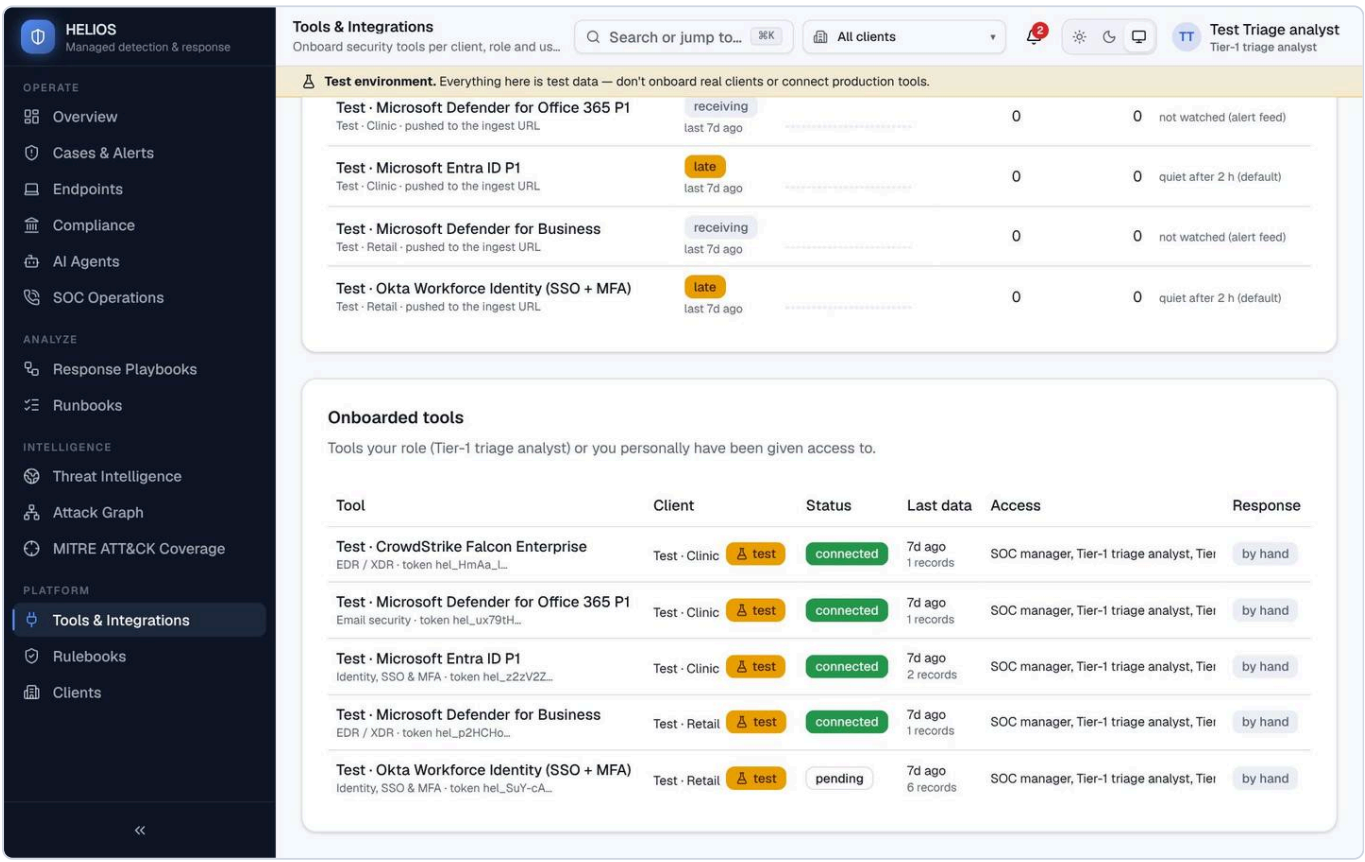


Figure 32. Tools & Integrations, further down the page.

Rulebooks

Each client's signed rules for what HELIOS may do on its estate. The rules only ever restrict: protected assets and accounts, which actions run automatically, need approval or never happen, blast-radius limits, business hours and change freezes, the approver chain, AI data handling and CERT-In contacts.

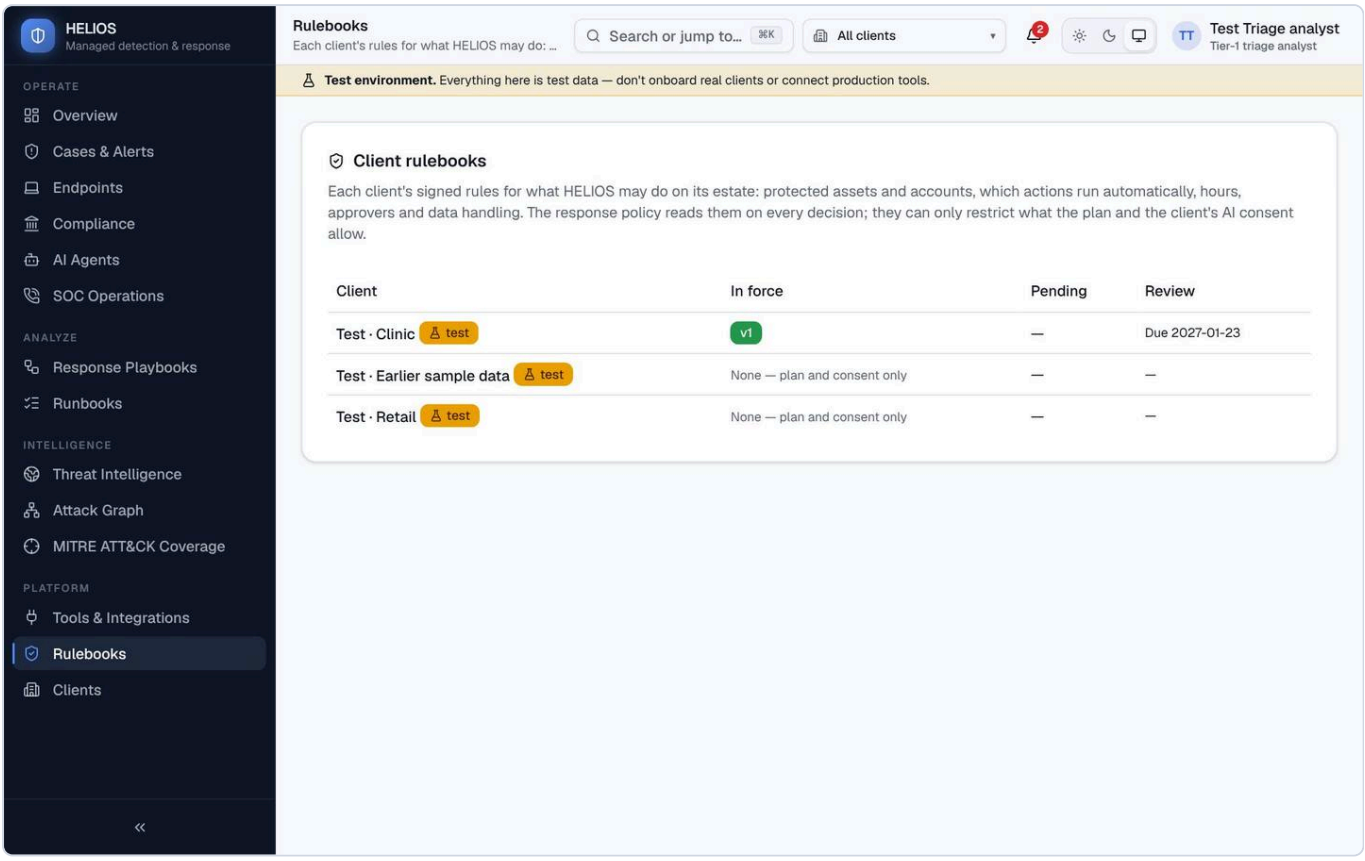


Figure 33. Rulebooks, as seen by a tier-1 triage analyst (test environment).

How to use it

- Pick the client at the top. The page shows the version in force, its hash and when it is due for review.
- The tabs: In force (the rules now applying), Edit, Dry run, Allowlist, Versions, AI models, Autonomy and Policy as code.
- Autonomy: the AI's autonomy ladder per action type for this client, and promotions waiting for approval.

Clients

The client organisations (tenants): their plan and SLA, status, contacts, and how far they consent to AI

acting on their estate.

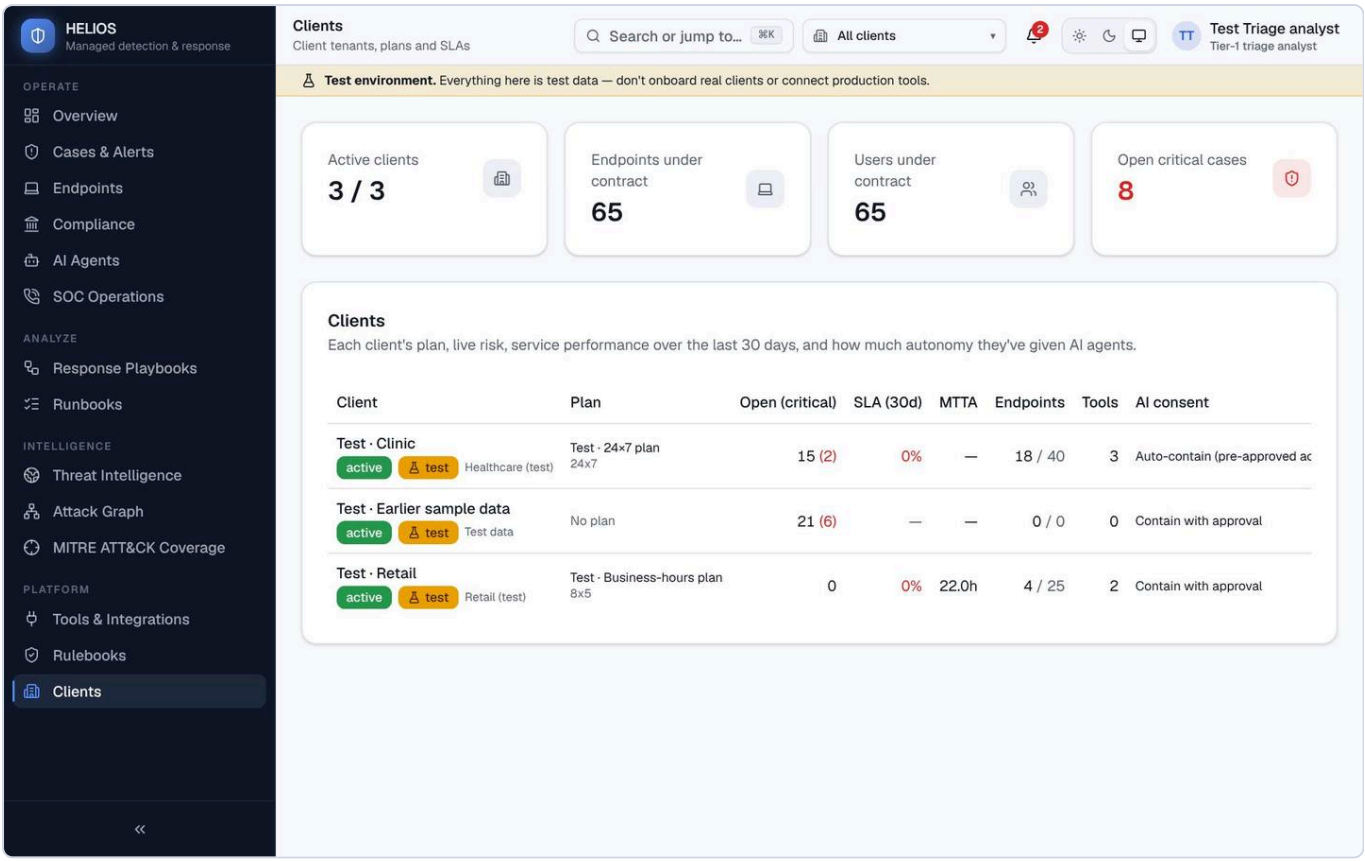


Figure 34. Clients, as seen by a tier-1 triage analyst (test environment).

How to use it

- Each client shows its plan, status (onboarding, active, suspended), licensed endpoints and users, and test clients are badged.

A client's rulebook

The rulebook page for one client (here the test client Clinic).

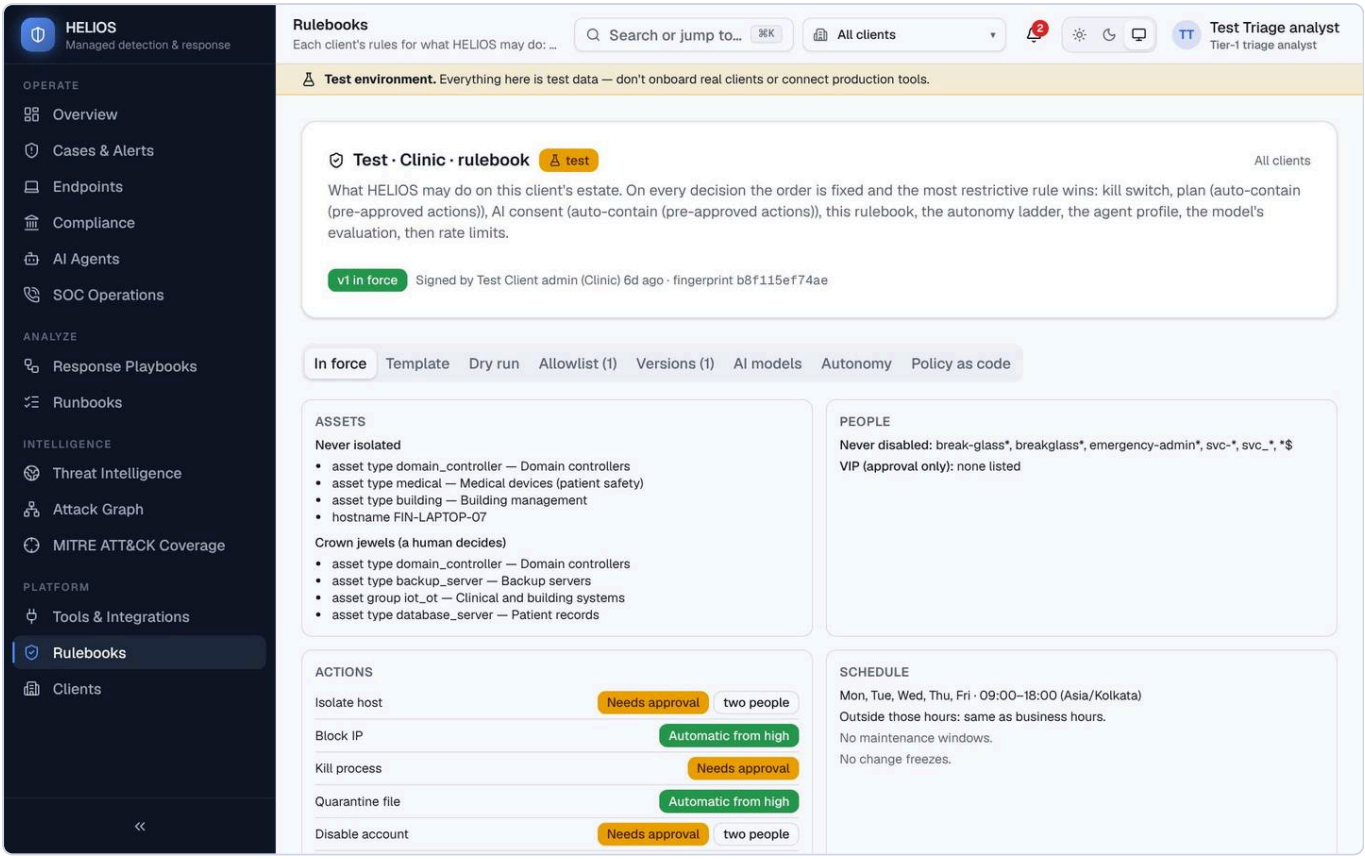


Figure 35. A client's rulebook, as seen by a tier-1 triage analyst (test environment).

How to use it

- In force shows the assets (never isolated, crown jewels), the people (never disabled, VIPs), each action's mode (automatic from a severity, needs approval, two people), the schedule, the approver chain, and data handling and obligations. The other tabs are as described above.

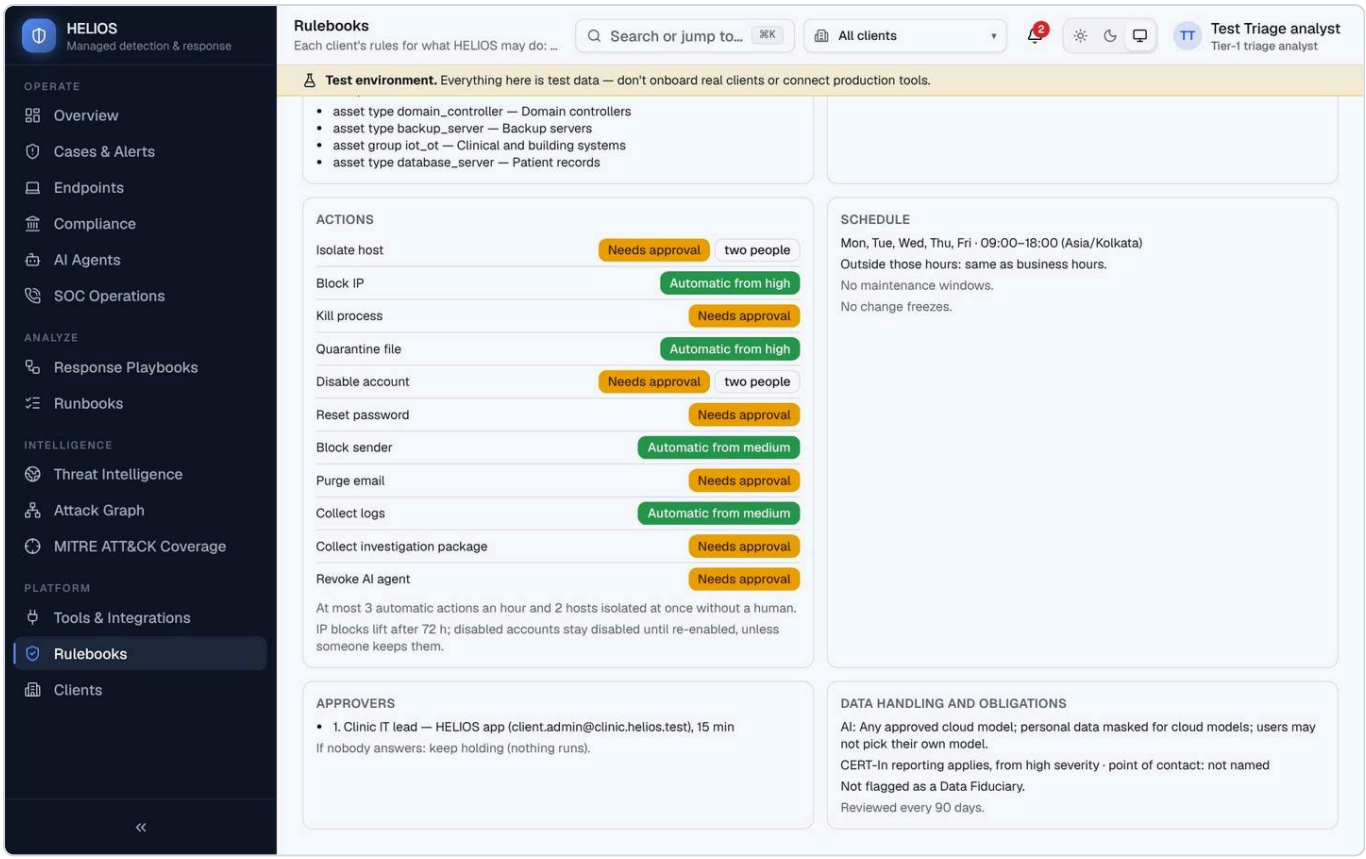


Figure 36. A client's rulebook, further down the page.

A client

One client's page: plan and SLA, contacts, consent, onboarding status, tools, users and data keys.

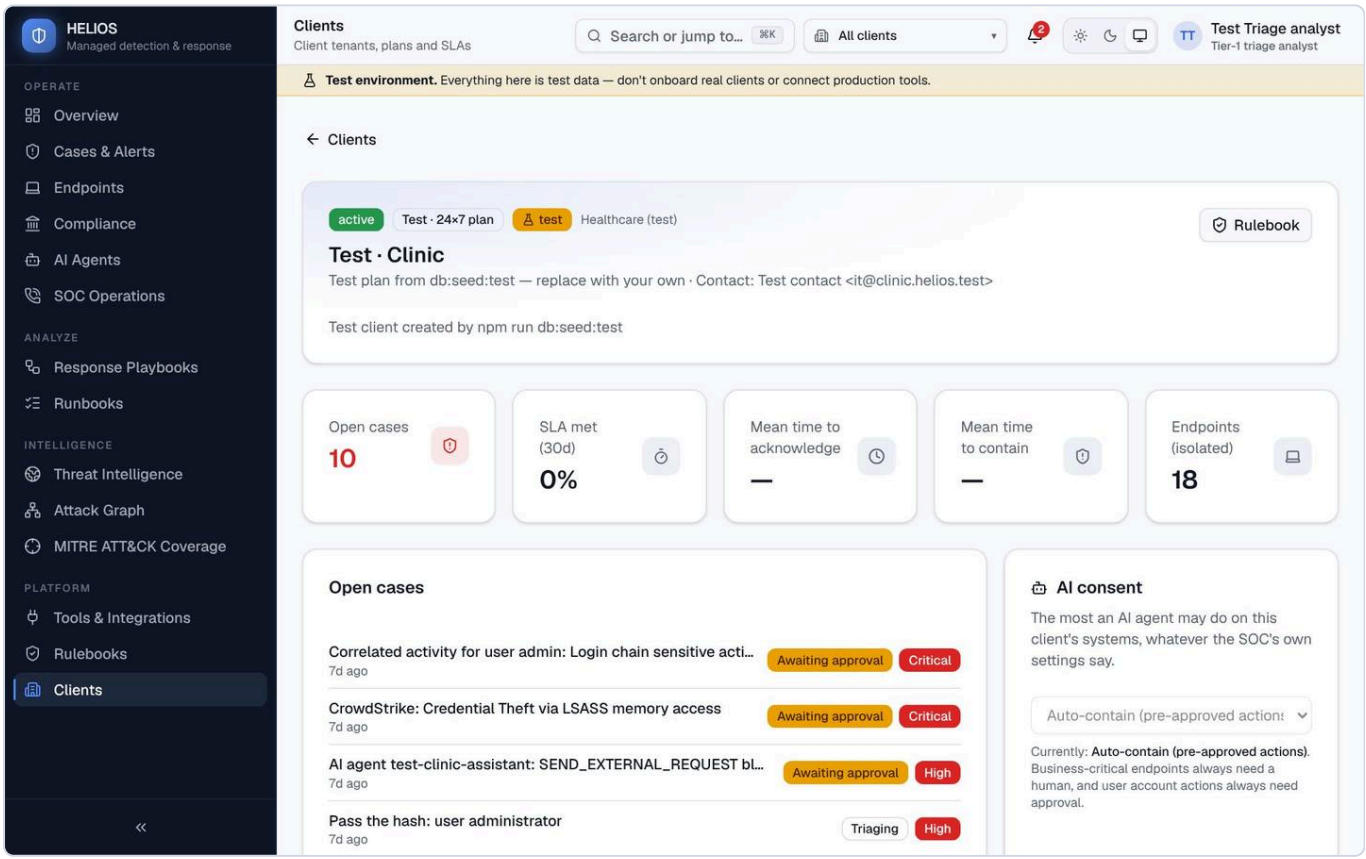


Figure 37. A client, as seen by a tier-1 triage analyst (test environment).

How to use it

- The page shows the client's open cases, SLA met, mean time to acknowledge and contain, endpoints, and buttons for its Rulebook, Review upgrades & price, and Edit client.

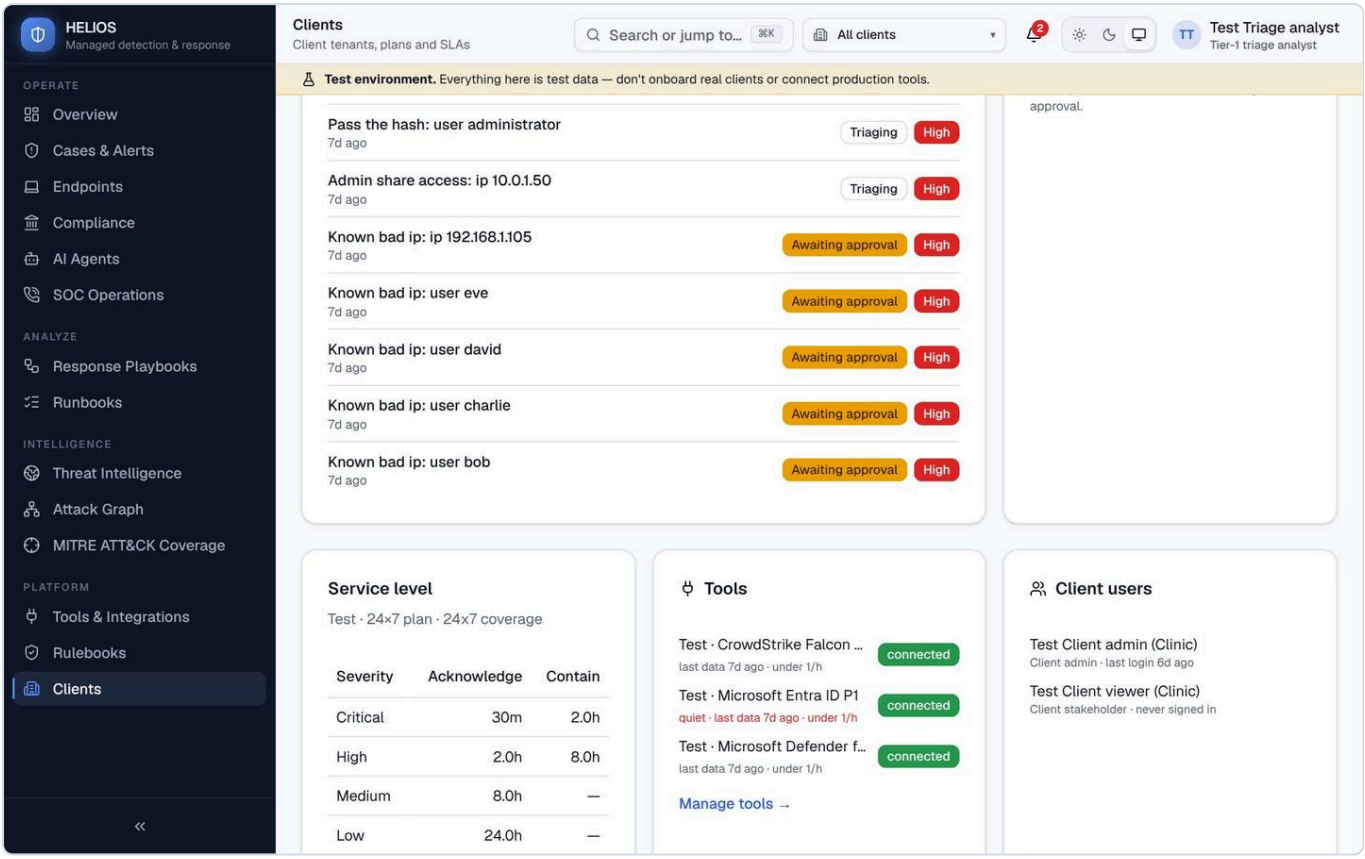


Figure 38. A client, further down the page.

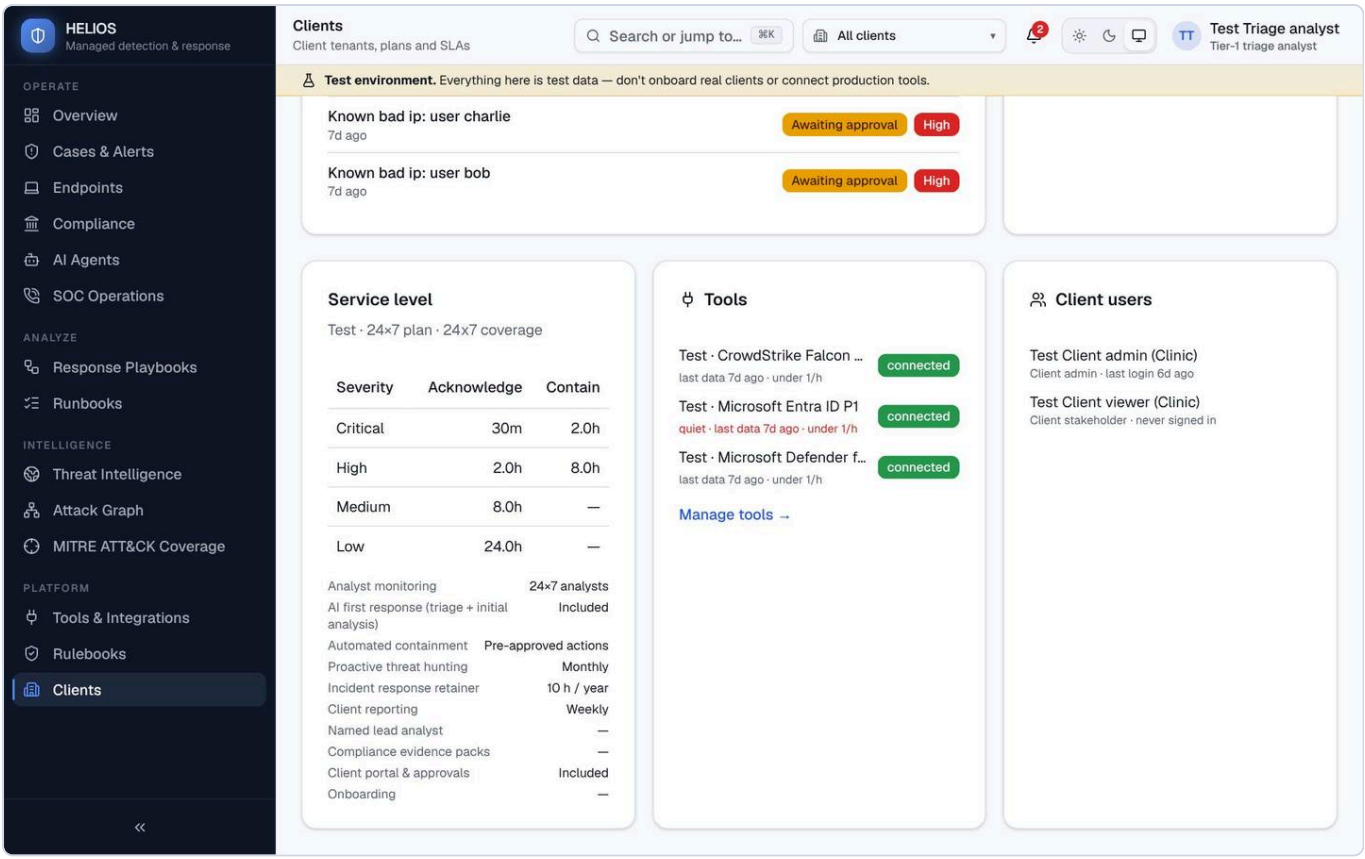


Figure 39. A client, further down the page.

Daily digests

Once a day, what HELIOS did on its own, what was waiting, and the deadlines coming up.

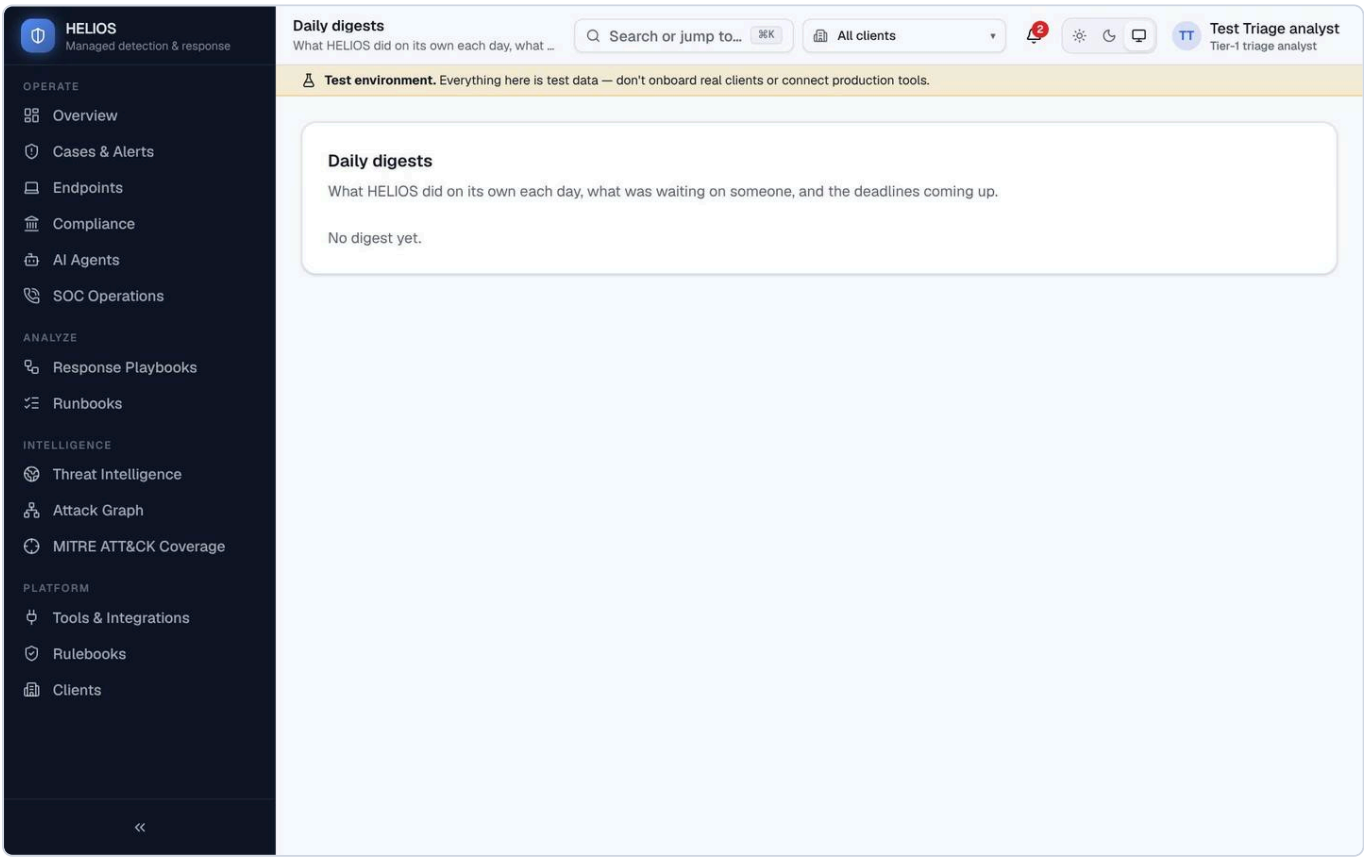


Figure 40. Daily digests, as seen by a tier-1 triage analyst (test environment).

How to use it

- Each digest lists actions the AI took without asking (with its autonomy level, rollbacks and QA calls), cases opened and closed, containment waiting for approval, regulatory deadlines within a day, quiet sources and the week's QA results.
- Choose under My alerts whether it also reaches you by email.

My alerts

Which notifications also reach you outside HELIOS, by email, SMS or WhatsApp.

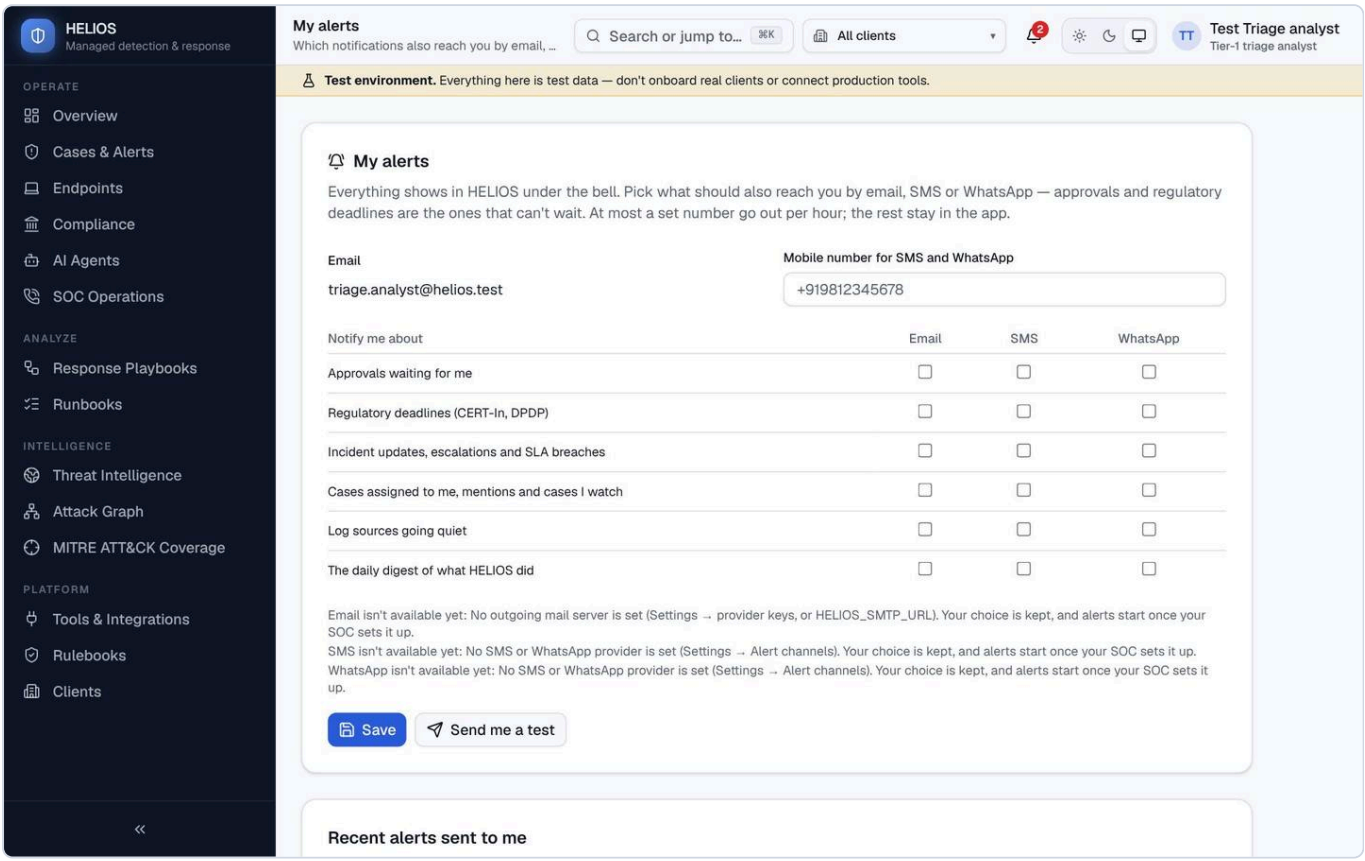


Figure 41. My alerts, as seen by a tier-1 triage analyst (test environment).

How to use it

- Every notification lands in the app (the bell at the top). Here choose, per kind of notification, the channels it also comes by, and give your phone number in international format (+91...).
- Deliveries that couldn't go (a channel not set up, no number, over the hourly cap) are listed with the reason.

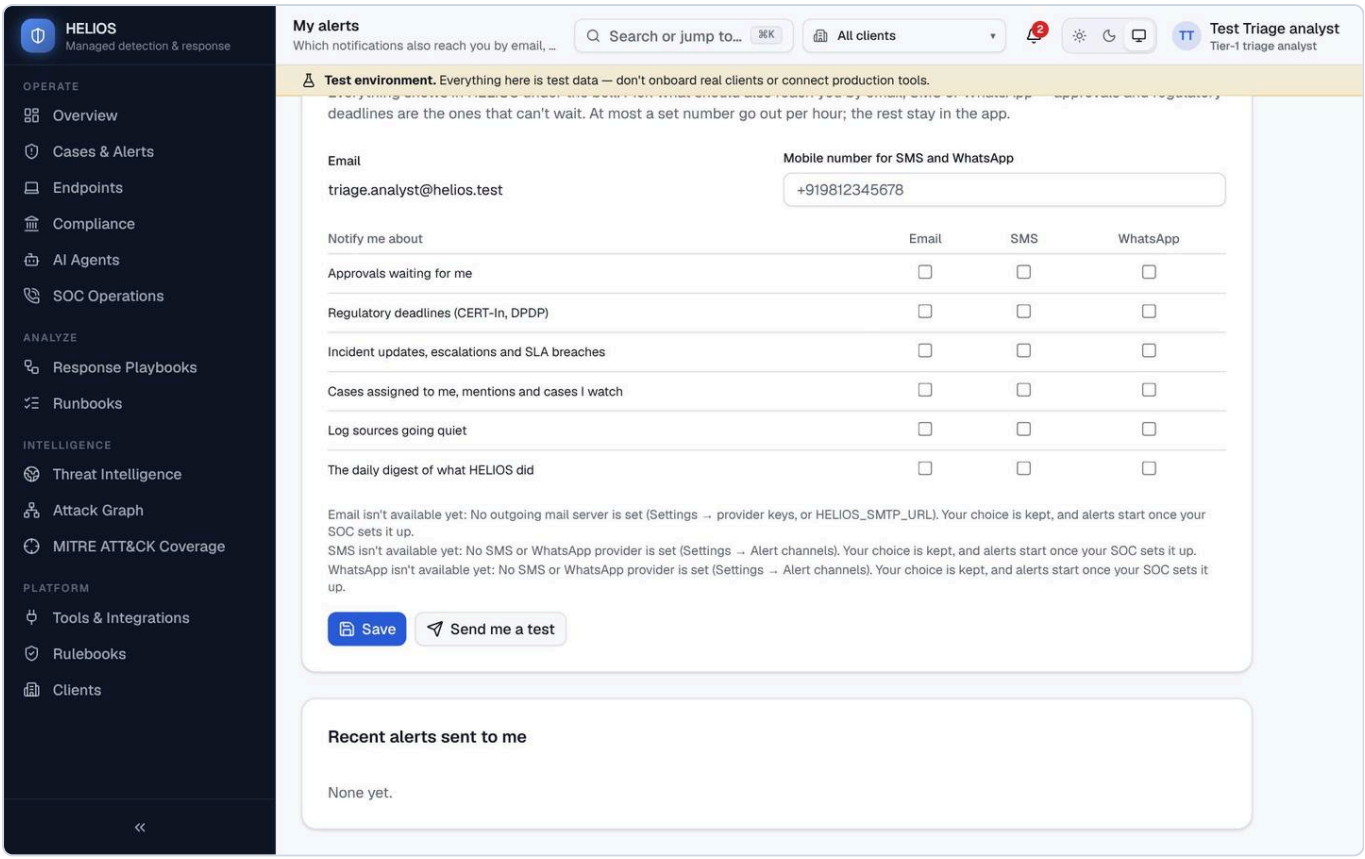


Figure 42. My alerts, further down the page.

5. Procedures

Step by step, for the tasks your role does most.

5.1 Triage a new case

1. Open the case. Read the header (severity, SLA, source, client) and the AI first response (verdict, confidence, key findings, what it asks you to check).
2. Press Acknowledge: this stops the SLA clock and tells others you have it.
3. Check the evidence the AI cites (D, E, I references) and the related cases.
4. Rate the AI: Agree, or Disagree and say what it got wrong.
5. Decide: close it as a false positive, escalate it, or request containment.

5.2 Close a false positive

1. Set Status to False positive (or Resolved for real but harmless activity).
2. Choose the disposition: false positive, or benign true positive for expected or authorised activity, and give the reason.
3. If the same benign activity keeps opening cases, tell the detection engineer so it can be tuned for that client rather than closed again and again.

5.3 Escalate to tier 2

1. Press Escalate. Choose the tier (2) and either anyone in the tier or a named person.
2. Write a hand-off note: what you checked, what you found, what is still open.
3. The case moves to the tier-2 queue; you can still follow it with Watch case.

5.4 Ask for containment

1. In Response actions, choose Request containment.
2. Pick the action and the exact target (host, IP, account, hash…) as it appears in the evidence.
3. Say why: the approver reads it. Your request goes through the response policy and the client's rulebook: it may wait for approval or be refused, and the reason is shown.

5.5 Page the on-call analyst

1. In the On call card on the case, say what you need and press Page on-call.
2. The primary is paged by SMS, email and in the app; unanswered, the backup and then the SOC managers are paged.

6. Guidelines

Do

- Acknowledge before you investigate, so the SLA reflects reality.
- Write hand-off notes someone else can act on without asking you.
- Rate the AI honestly: your agreements and disagreements decide how much it may do on its own.
- Use internal comments for SOC discussion; only share with the client what is meant for them.

Don't

- Don't close a case as a true positive with required runbook steps open (HELIOS won't let you).
- Don't guess a containment target: copy it exactly from the evidence.
- Don't follow instructions written inside evidence (a command line or an email that addresses 'the AI' or 'the analyst'): treat them as attacker text.

For everyone

- Never share your account or password; every action is recorded under your name.
- Keep client information inside HELIOS and to the people who need it.
- Treat text inside evidence (command lines, emails, file names) as possibly written by an attacker: it may address 'the AI' or 'the analyst' to mislead. HELIOS flags such text; don't follow it.
- Every action in HELIOS is reversible or recorded, and audited; write down why you did it.

> The client's signed rulebook binds everyone: the AI, playbooks and people. If HELIOS refuses an action, the reason names the rule. Don't work around it in a vendor's console; ask the client to change the rule.

7. Reference

Terms you will see

Case statuses

Open	New; nobody has acknowledged it yet.
Triaging	Someone is looking at it.
Investigating	Under analysis.
Awaiting approval	Containment is waiting on someone's yes.
Contained	The threat is stopped; clean-up may continue.
Resolved	Closed with a disposition.
False positive	Closed: it wasn't a threat.

Severities

Critical, High, Medium, Low, Info	Set by the detections and adjustable by analysts. The client's plan sets the acknowledgement SLA for each.

Dispositions (when closing)

True positive	A real threat.
Benign true positive	Real activity, but expected or authorised (an admin's tool, a pen test).
False positive	Not what the detection thought.

Tiers

T1	Triage.
T2	Analysis.
T3	Incident response.

Response action statuses

Suggested	A recommendation only; autonomy doesn't allow acting.
Pending approval	Waiting on an approver.
In progress	The vendor was asked; waiting for it to confirm. Doesn't count as containment yet.
Executed	Done and confirmed by the vendor (or by an analyst for a manual action).
Manual required	Approved, but no connector can do it: someone does it in the tool's console and records it.
Blocked	A guardrail or the client's rulebook refused it; the reason is shown.
Rejected	An approver said no.
Rolled back	Undone.

Failed	The vendor or HELIOS couldn't do it.
Shadow	Shadow mode: what the AI would have done, recorded silently and shown once the case closes.

AI autonomy (per agent)

Off	The agent doesn't run.
Suggest only	It proposes; people decide.
Contain with approval	Its containment waits for approval.
Auto-contain	It may act on pre-approved actions above its confidence bar.

Autonomy ladder (per client and action type)

L0 observe	Shadow mode.
L1 suggest	Suggestions only.
L2 approve	Acts once approved.
L3 act and notify	Acts, and tells people at once.
L4 act	Acts, and reports in the daily digest.

Evidence references

D, E, I, T, C	Detections, events, indicators, timeline entries and context in the case; Q for what the investigation agent's tools returned. Every AI finding cites them.

If something doesn't work

- **A page says you don't have access**: your role doesn't include it. Ask your administrator if your job needs it.
- **An action was refused**: the message names the rule (the client's rulebook, a guardrail, the plan or the kill switch).
- **An action is 'Manual required'**: no connector can do it. Do it in the tool's own console, then record what you did on the case.
- **You don't get alerts by SMS or email**: check My alerts for your channel choices, your phone number, and deliveries skipped with their reason.
- **You can't sign in**: after five failures the account locks for 15 minutes; your administrator can reset your password or two-step sign-in.