HELIOS USER GUIDE · SECURITY OPERATIONS

Tier-2 security analyst

Procedures, usage, guidelines, and your roles and access

Tier 2: the primary analysis of real cases, investigations and forensics, containment requests, and keeping the client informed.

Platform	HELIOS managed detection and response
Guide	04 of 11 · Tier-2 security analyst
Date	31 October 2026
Version	Written against HELIOS 809e500
Screenshots	Taken in the HELIOS test environment with test data, signed in as a tier-2 security analyst. Your screens show your own clients and data.

Contents

1. Your role at a glance
2. Roles and access
3. Getting started
4. Your screens
5. Procedures
6. Guidelines
7. Reference

> You have **28** of 50 permissions and **19** pages in your sidebar: Overview, Cases & Alerts, Endpoints, Reports, Compliance, AI Agents, SOC Operations, Investigations, Email & File Forensics, Response Playbooks, Runbooks, Agent Governance, Threat Intelligence, Attack Graph, MITRE ATT&CK Coverage, Tools & Integrations, Rulebooks, Clients, API Tokens.

1. Your role at a glance

Tier 2: the primary analysis of real cases, investigations and forensics, containment requests, and keeping the client informed.

What you are responsible for

- Investigate escalated cases fully: scope, root cause, affected hosts and accounts.
- Run log investigations and email and file forensics.
- Request containment, including email purges, password resets and investigation packages.
- Share cases and publish reports to the client; prepare regulatory notices.
- Maintain the indicator store and blocklist; govern clients' AI agents.

A typical day or shift

1. Work your assigned cases and the tier-2 queue, SLA first.
2. Check Approvals for your requests that were held or refused.
3. Check Compliance for clocks running on your cases.

Who to turn to

Major incidents go to tier 3 (incident responder); ask the threat hunter for wider campaigns; tuning requests go to the detection engineer.

2. Roles and access

What you can see

You work across the clients the MSSP serves. The client switcher at the top focuses every page on one client, or All clients. If your account is limited to security domains (for example endpoint and identity), your case queues show only those domains.

Your access comes from your role, and HELIOS checks it on the server for every page, button and API call; hiding a button is only a convenience. A change to your role applies on your next click. Everything you do is written to the hash-chained audit log under your name.

Your pages

Overview, Cases & Alerts, Endpoints, Reports, Compliance, AI Agents, SOC Operations, Investigations, Email & File Forensics, Response Playbooks, Runbooks, Agent Governance, Threat Intelligence, Attack Graph, MITRE ATT&CK Coverage, Tools & Integrations, Rulebooks, Clients, API Tokens, plus My alerts and your two-step sign-in settings in the user menu.

Your permissions

What the permission lets you do	You
Cases	
See the overview dashboard (dashboard.view)	Yes
See cases in their queues (or their own tenant, for clients) (cases.view)	Yes
See every queue, regardless of assigned security domains (cases.view_all)	No
Acknowledge cases, change status and severity (cases.triage)	Yes
Assign cases to other analysts (cases.assign)	Yes
Escalate or hand off a case to the next tier (cases.escalate)	Yes
Comment on cases and @mention teammates (cases.comment)	Yes
Share cases with the client and external channels (cases.share)	Yes
Response	
Run containment playbooks (subject to approval) (containment.request)	Yes
Approve or reject pending containment (containment.approve)	No
Isolate and release endpoints directly (endpoints.contain)	No
Analysis	
Run log investigations (investigations.run)	Yes
Submit emails and files for forensic analysis (forensics.run)	Yes
Upload raw logs into the detection pipeline (telemetry.ingest)	Yes
Govern AI agents: identities, tokens, revocation (governance.manage)	Yes

Intelligence and detection	
Look up threat indicators (intel.view)	**Yes**
Edit the blocklist and indicator store (intel.manage)	**Yes**
See ATT&CK coverage (mitre.view)	**Yes**
Switch SIEM detectors on and off (detections.tune)	No
See response playbooks (playbooks.view)	**Yes**
Author and enable response playbooks (playbooks.manage)	No
AI agents	
See AI agent runs and configuration (ai.view)	**Yes**
Dispatch AI agents onto a case (ai.dispatch)	**Yes**
Rate AI agent output (ai.feedback)	**Yes**
Configure AI providers, agent profiles and autonomy (ai.configure)	No
Tools and endpoints	
See onboarded security tools (integrations.view)	**Yes**
Onboard and configure security tools (integrations.manage)	No
See the endpoint inventory and protection coverage (endpoints.view)	**Yes**
Clients and rulebooks	
See client tenants (tenants.view)	**Yes**
Create and edit client tenants (tenants.manage)	No
Read a client's rulebook: what HELIOS may do on its estate (rulebook.view)	**Yes**
Draft changes to a client's rulebook and its allowlist (rulebook.edit)	No
Put a client's rulebook in force (client admins sign; staff only on the client's recorded instruction) (rulebook.sign)	No
Compliance and reports	
See regulatory clocks and notices (CERT-In, DPDP) and the ISO 27001 evidence map (compliance.view)	**Yes**
Draft regulatory notices, record the client's approval, and file them (compliance.prepare)	**Yes**
Approve a regulatory notice before it is filed (the report is the client's; staff record a client's approval given elsewhere) (compliance.approve)	No
Read client-facing incident reports (reports.view)	**Yes**
SOC operations	
See the on-call rota, pages, runbooks, QA results and digests (operations.view)	**Yes**
Edit the on-call rota, runbooks and operations settings (operations.manage)	No
Review the weekly sample of what the AI did on its own (qa.review)	No

Commercial	
Use the Build & Price configurator (pricing.view)	No
See internal cost and margin (pricing.margin)	No
Edit the price book (pricing.manage)	No
See contracts, pilots and invoices (clients: their own) (billing.view)	No
Draft and record contracts, issue invoices and record payments (billing.manage)	No
Platform	
Invite users and change roles (users.manage)	No
Edit platform settings: company profile, provider keys, forensics thresholds (settings.manage)	No
Read the audit chain (audit.view)	No
Export the audit chain (audit.export)	No
Create personal API tokens for the REST API (acting with their own permissions) (api.tokens)	**Yes**

How your role fits with the others

Role	Group	What it is for
Platform admin	Leadership	Owns the platform: users, AI configuration, price book, everything else.
SOC manager	Leadership	Runs the SOC: every queue, approvals, AI autonomy, tool onboarding.
Tier-1 triage analyst	Security operations	Works the first-line queue: acknowledges, validates AI triage, escalates.
Tier-2 security analyst	Security operations	Primary analysis: investigations, forensics, containment requests.
Tier-3 incident responder	Security operations	Owns major incidents: approves containment, isolates endpoints, client comms.
Threat hunter	Security operations	Hunts across all tenants with investigations, forensics and intel.
Detection engineer	Engineering	Maintains playbooks, detections, tool integrations and ATT&CK coverage.
Auditor / compliance	Governance	Read-only oversight: cases, AI decisions and the audit chain.
Account manager	Commercial	Client onboarding, Build & Price quotes and margins.
Client admin	Client	Client-side: their own tenant's cases and tools; approves containment on their estate.
Client stakeholder	Client	Client-side read-only: dashboards, incident reports and endpoint coverage.

3. Getting started

Signing in

1. Open the HELIOS address you were given and sign in with your email and password. A new account gets a temporary password from its administrator.
2. If two-step sign-in is required, set it up the first time: scan the code with an authenticator app, enter a code, and keep the ten recovery codes somewhere safe. Each recovery code works once.
3. If your organisation uses single sign-on, use the sign-in button for your identity provider instead; it satisfies two-step sign-in.
4. Five wrong passwords or codes in a row lock the account for 15 minutes. An administrator can reset your password or your two-step sign-in if you lose your phone.

Finding your way

- **Sidebar**: your pages, grouped (Operate, Analyze, Intelligence, Platform, Commercial). The arrow at the bottom collapses it.
- **Search or jump to** (Ctrl K or Cmd K): find a case, page or client by name.
- **Client switcher**: All clients, or focus every page on one client.
- **The bell**: your notifications (mentions, hand-offs, approvals, pages). Choose which also reach you by email, SMS or WhatsApp under My alerts.
- **Light, dark or system theme**, and your **user menu**: My alerts, two-step sign-in, Sign out.
- A yellow **Test environment** banner means everything on that installation is test data.

Install it as an app

HELIOS can be installed from the browser (Install app in the address bar) as a desktop or mobile app with its own window and shortcuts.

4. Your screens

Each page you can open, what it is for, and how to use it with your access. Where a page is long, further screenshots show the rest.

Overview

Your starting page. It shows what needs you now, and it changes with your role: the SOC sees live posture, queues, SLAs and AI activity; account managers see the commercial picture; client users see their own organisation.

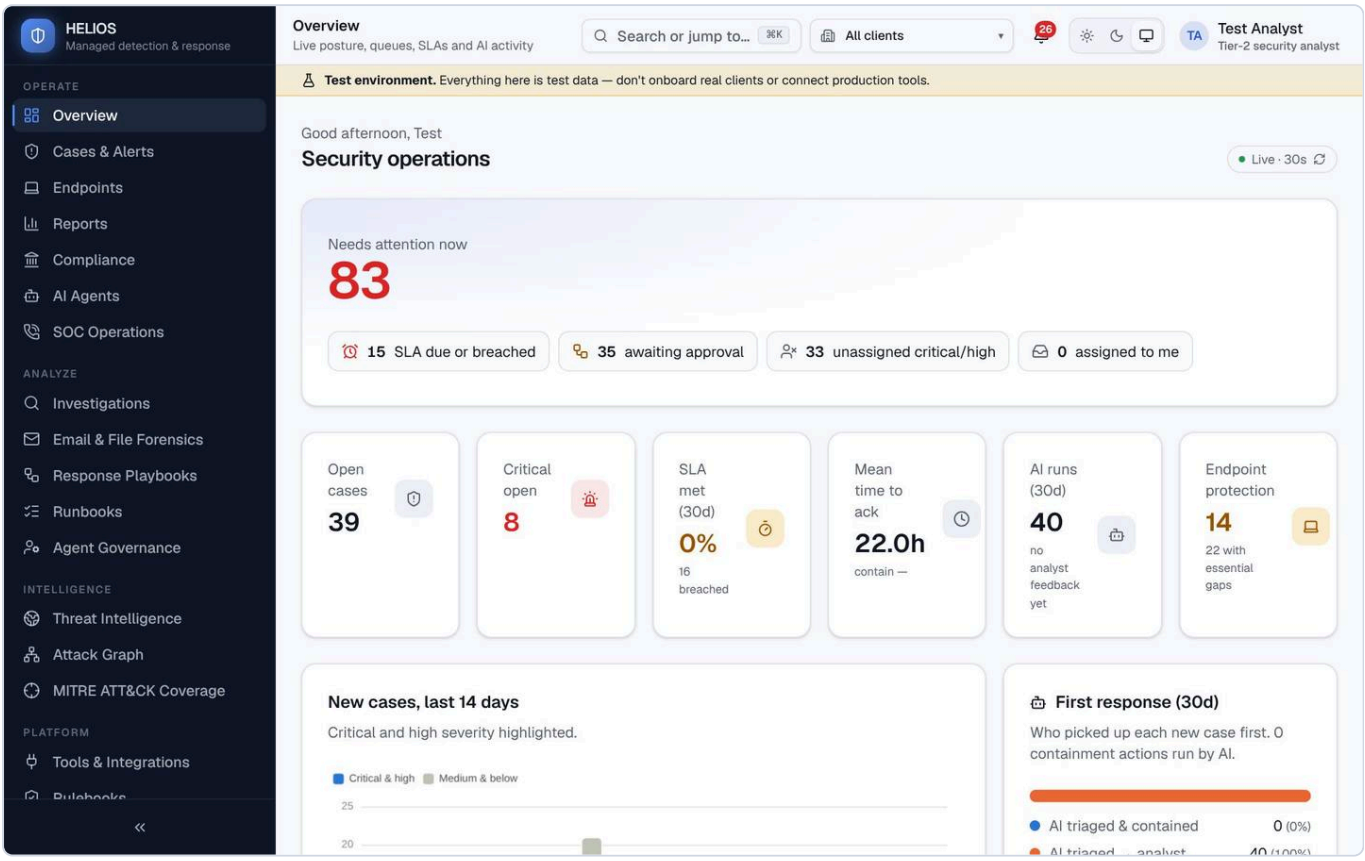


Figure 1. Overview, as seen by a tier-2 security analyst (test environment).

How to use it

- Needs attention now adds up what is due: cases whose SLA is due or breached, containment awaiting approval, unassigned critical and high cases, and cases assigned to you. Click a figure to open the matching list.
- The Live badge means the figures refresh on their own every 30 seconds; the arrow refreshes now.

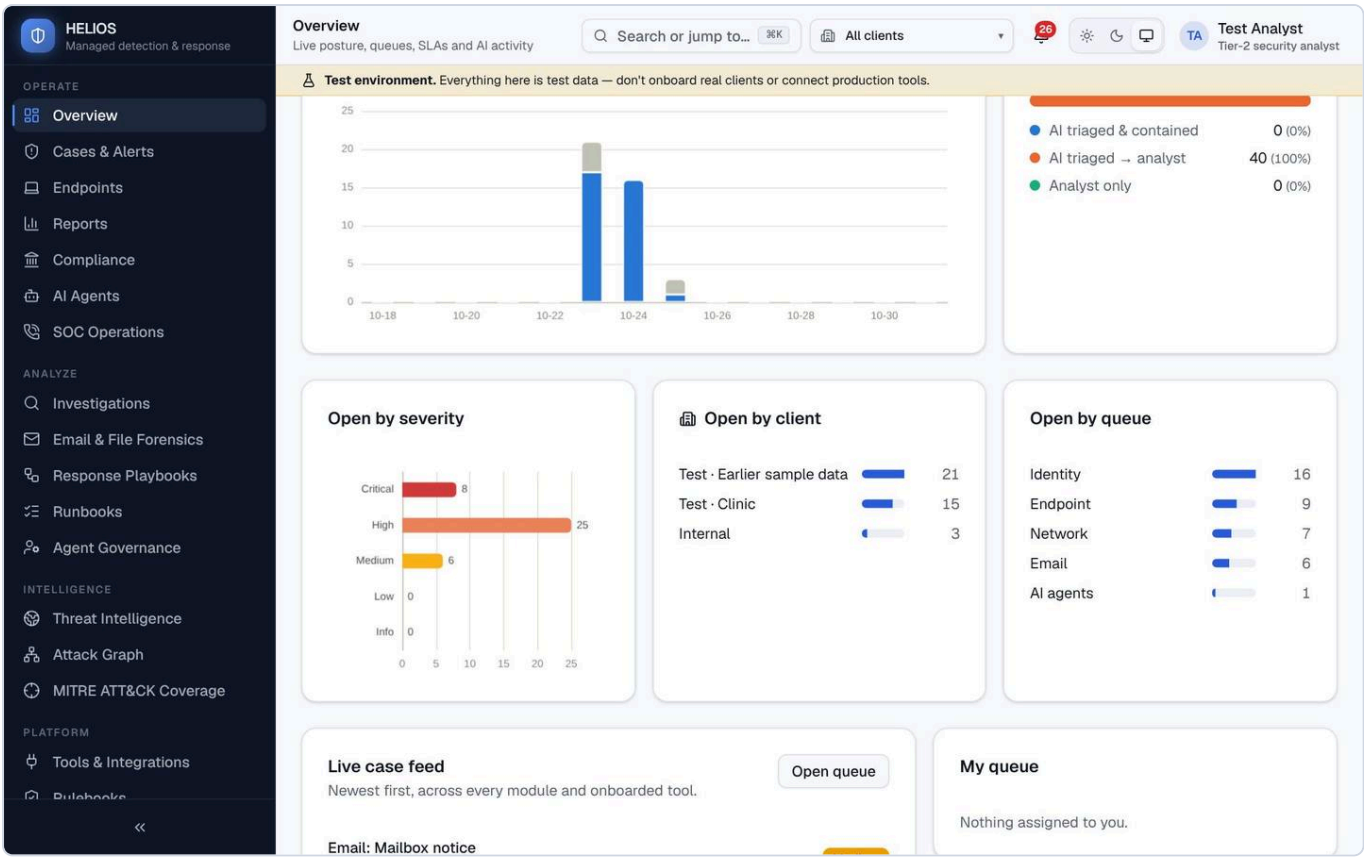


Figure 2. Overview, further down the page.

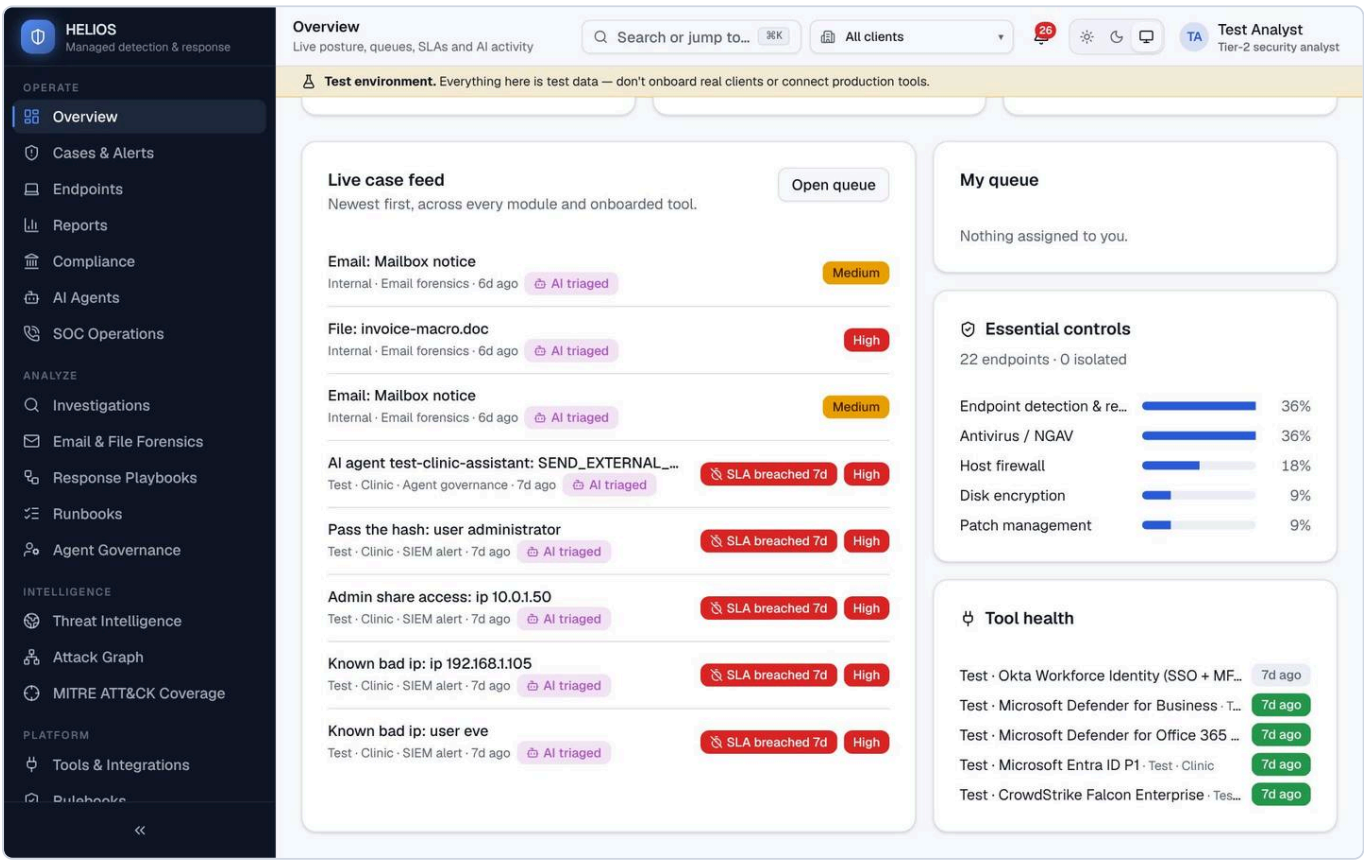


Figure 3. Overview, further down the page.

Cases & Alerts

Every case HELIOS has opened: detections from your clients' telemetry, vendor alerts, uploaded investigations, email and file analysis, and governed AI agents. Correlated detections for the same user, host or IP join one case instead of opening many.

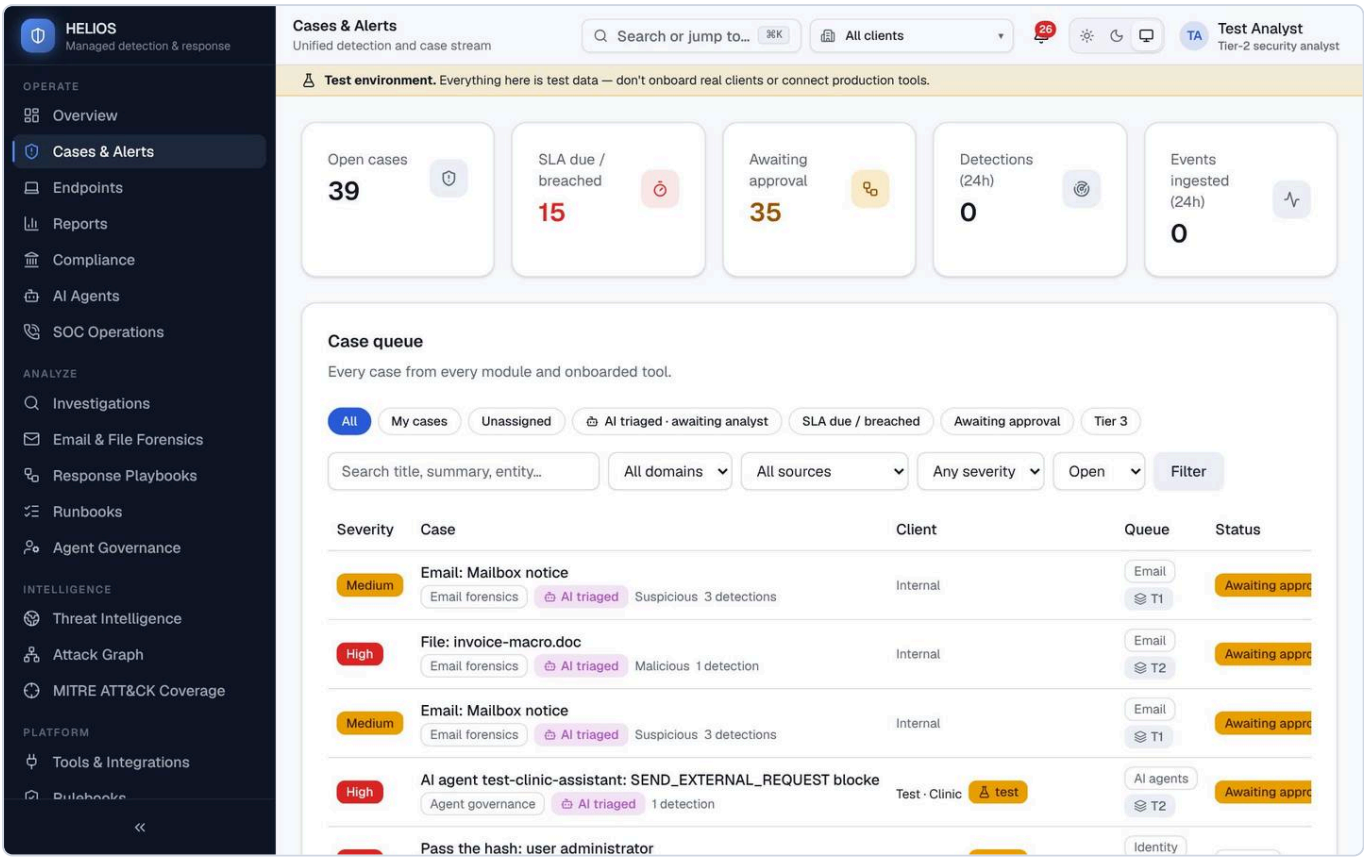


Figure 4. Cases & Alerts, as seen by a tier-2 security analyst (test environment).

How to use it

- Filter by status, severity, source, client and assignee, or search by title. Open a case by clicking its row.
- If your account is limited to security domains (endpoint, identity, email, network, cloud, AI agents), you see only those queues.
- The SLA badge counts down to the acknowledgement deadline from the client's plan; a red badge means it has been breached.

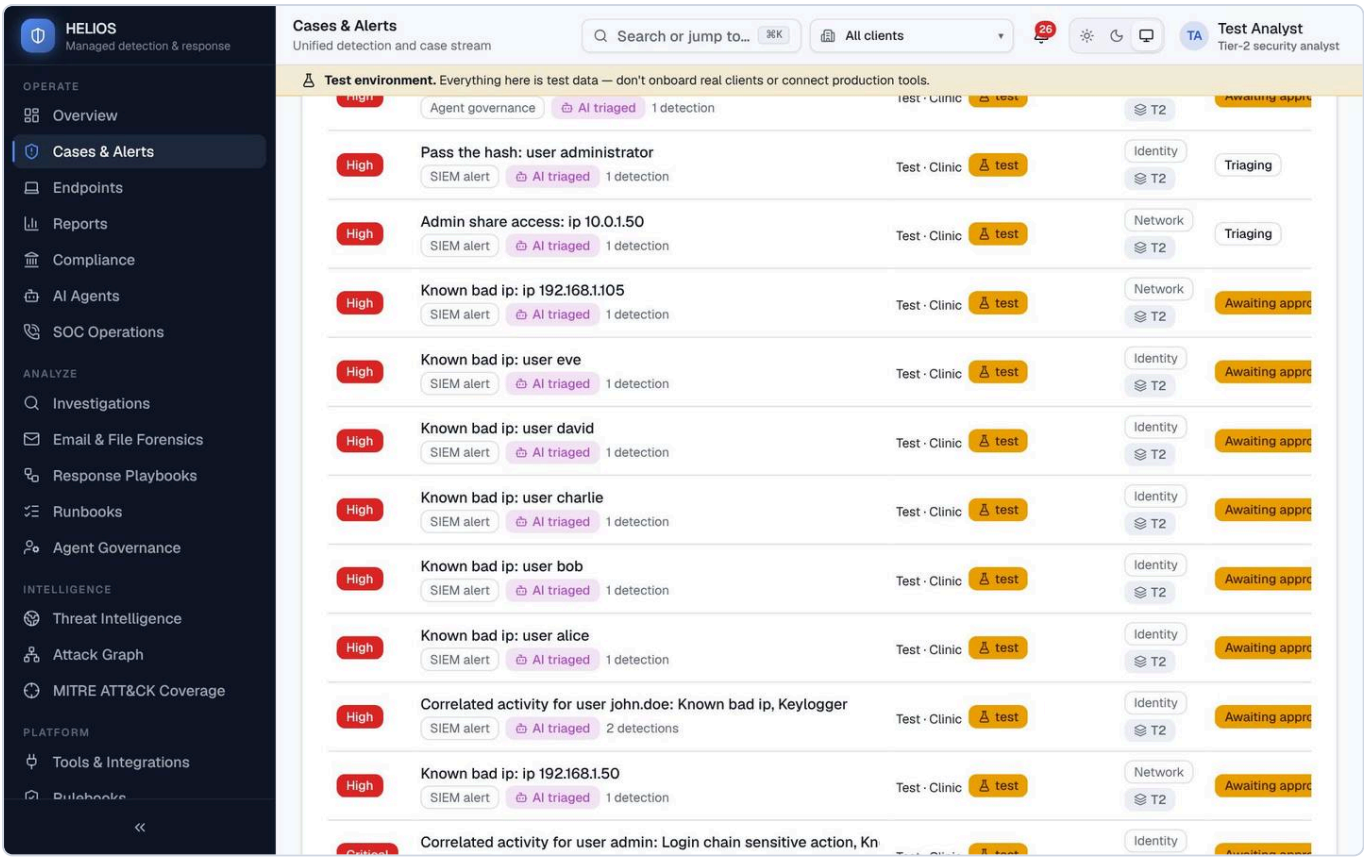


Figure 5. Cases & Alerts, further down the page.

A case

Everything about one case on one page: what was detected, what the AI made of it, what has been done

and what is waiting, the evidence, and the discussion.

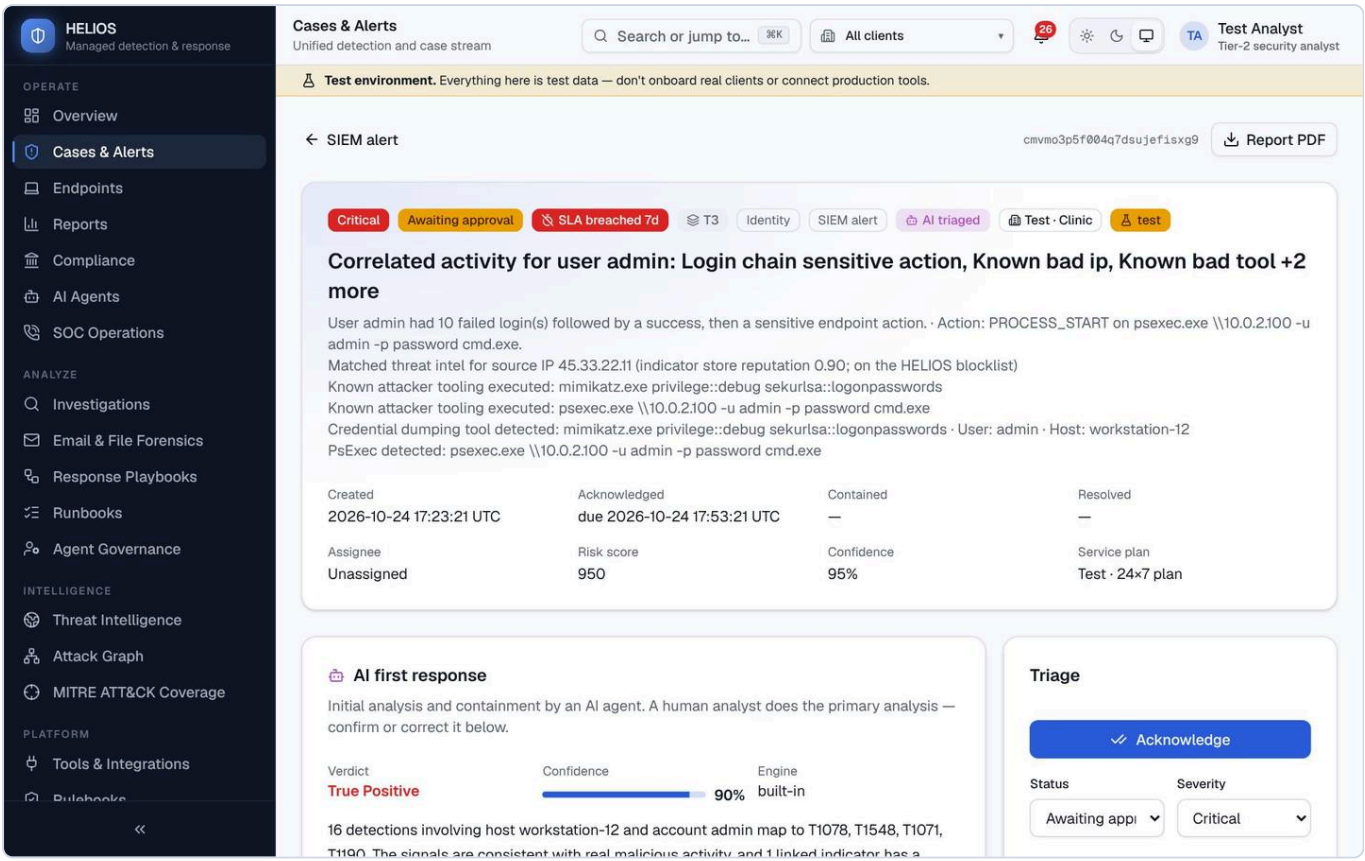


Figure 6. A case, as seen by a tier-2 security analyst (test environment).

How to use it

- The header shows severity, status, the SLA, the tier (T1 triage, T2 analysis, T3 incident response), the security domain, the source, the client and whether AI triaged it. Below are the dates (created, acknowledged, contained, resolved), the assignee, the risk score, the confidence and the client's service plan. Report PDF downloads the incident report.
- AI first response shows the agent's verdict, confidence and engine, its summary, its key findings (each citing evidence: D detections, E events, I indicators, T timeline, C context) and what it asks the analyst to check. It is a first pass: a person confirms or corrects it.
- Agree or Disagree records whether the analysis was right; a disagreement asks what it got wrong. These ratings drive the AI's autonomy ladder.
- Dispatch agent runs another agent on the case: the triage first responder, the phishing analyst, the threat-intel enricher, the investigation agent (works the case with read-only tools) or the client report writer.
- In Triage: Acknowledge stops the SLA clock; Status and Severity change the case; Assignee gives it to someone.
- Watch case sends you its updates. Escalate hands the case to another tier or person with a note; Share makes it visible to the client or sends it to an external channel.
- Page on-call (in the On call card) pages whoever is on call about this case, with what you need from them.
- Request containment asks for an action (isolate host, block IP, disable account, kill process, quarantine file, block sender, collect logs, purge email, reset password, collect an investigation package...) against a target, with a reason shown to the approver.

- Indicators shows the case's IPs, domains, URLs, hashes and addresses with their reputation; the arrow re-checks one. Related cases lists earlier or concurrent cases sharing hosts, accounts, indicators or the same agent.
- Regulatory reporting shows whether CERT-In's 6-hour report or the DPDP breach notices apply, the incident types HELIOS proposes, and the notices' state.
- Runbooks attach themselves to matching cases; tick each step or skip it with a reason. A case can't be closed as a true positive while a required step is open.
- The discussion takes comments and @mentions. Internal comments stay with the SOC; client-visible ones reach the client.
- Closing: set the status to Resolved or False positive, then give a disposition (true positive, benign true positive, false positive) and a reason.

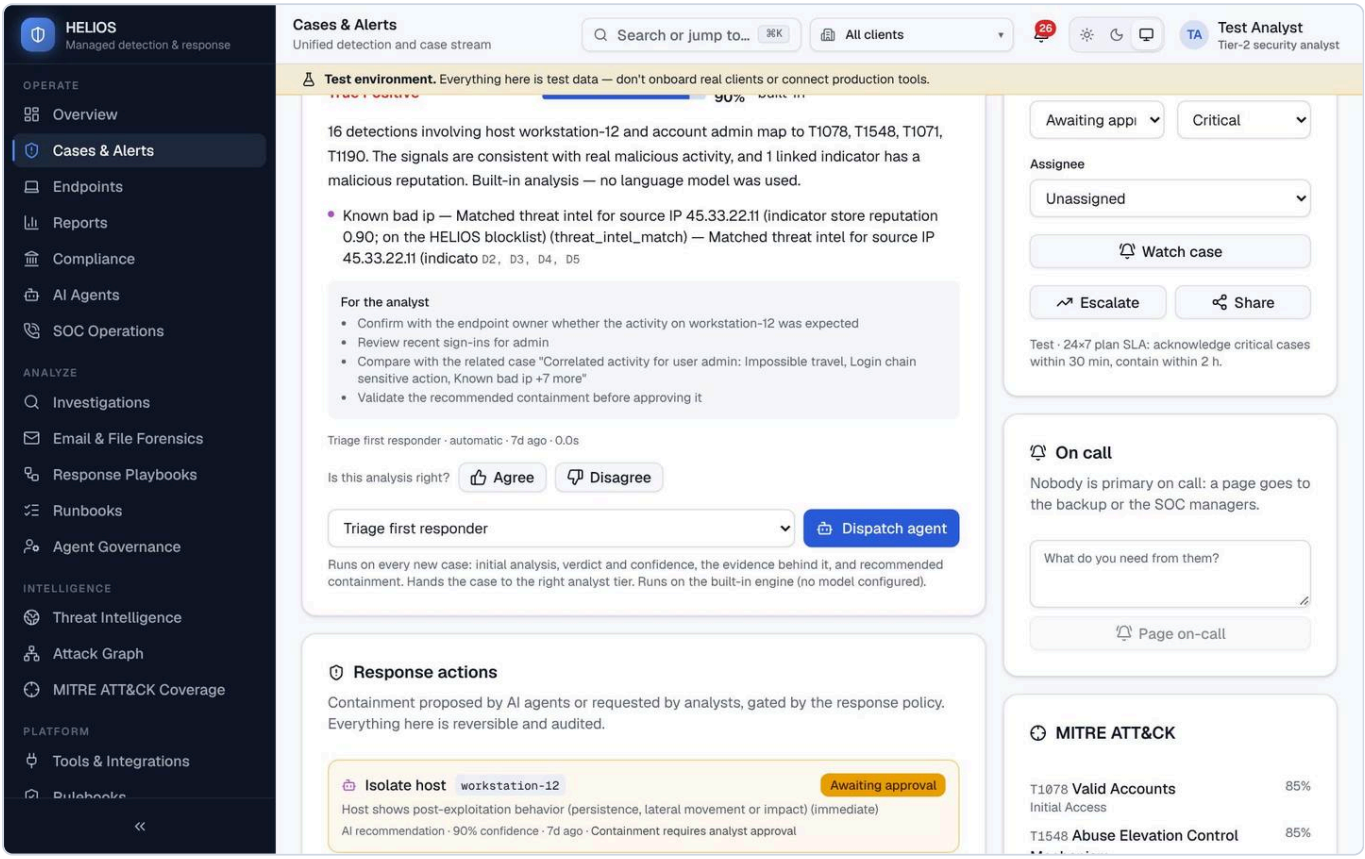


Figure 7. A case, further down the page.

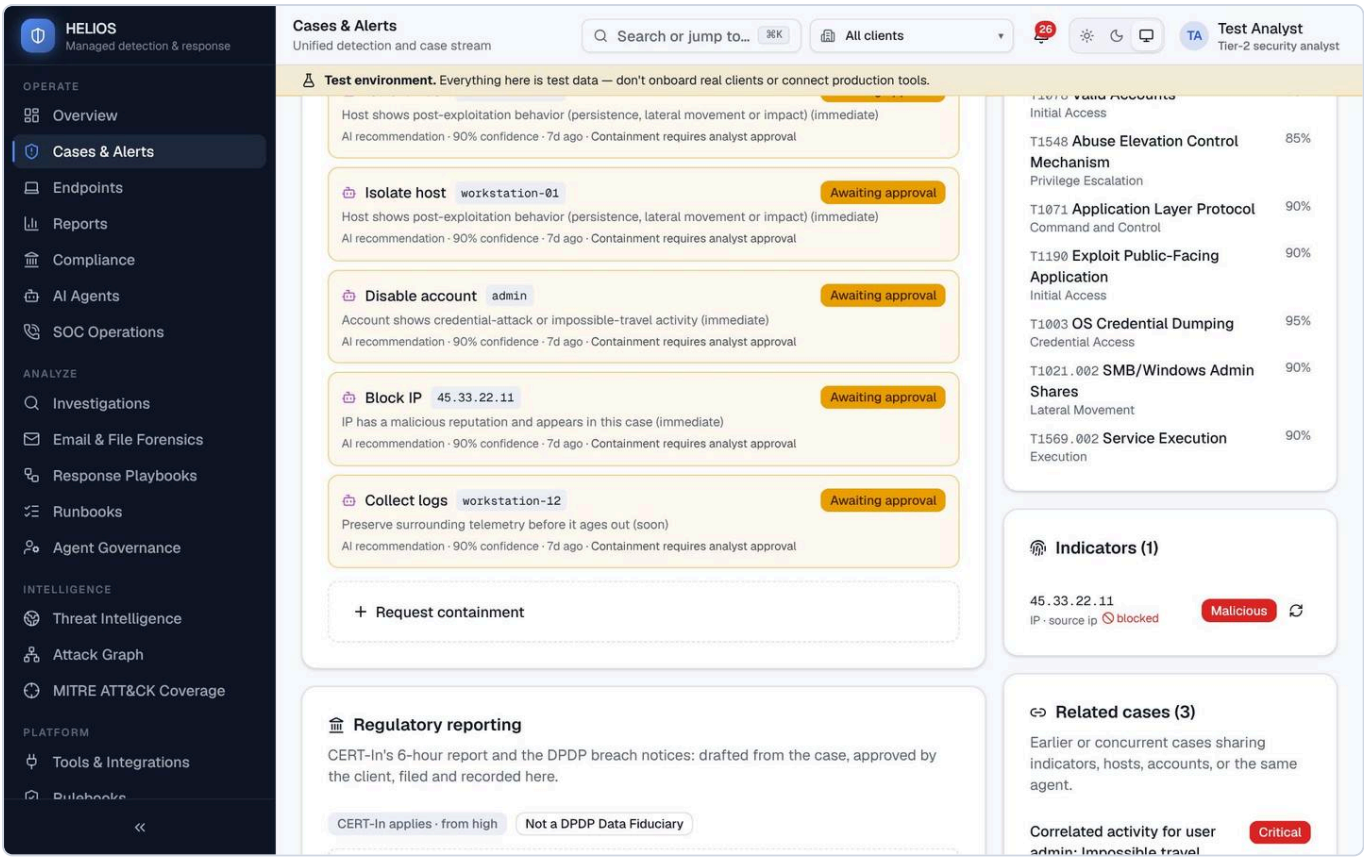


Figure 8. A case, further down the page.

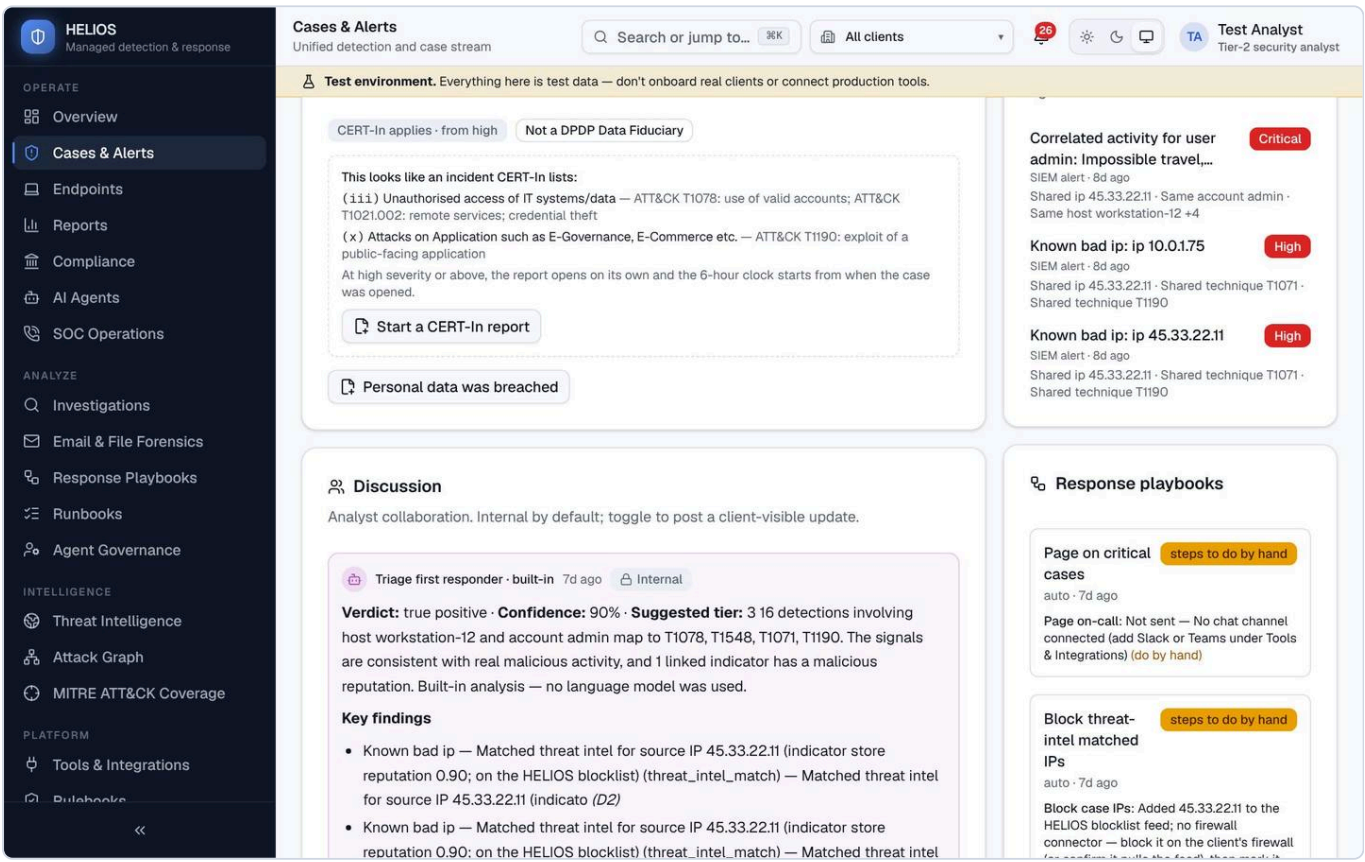


Figure 9. A case, further down the page.

Endpoints

The endpoint inventory and the whole attack surface: devices, servers, cloud, network, identity, IoT/OT and applications, with how well each is protected.

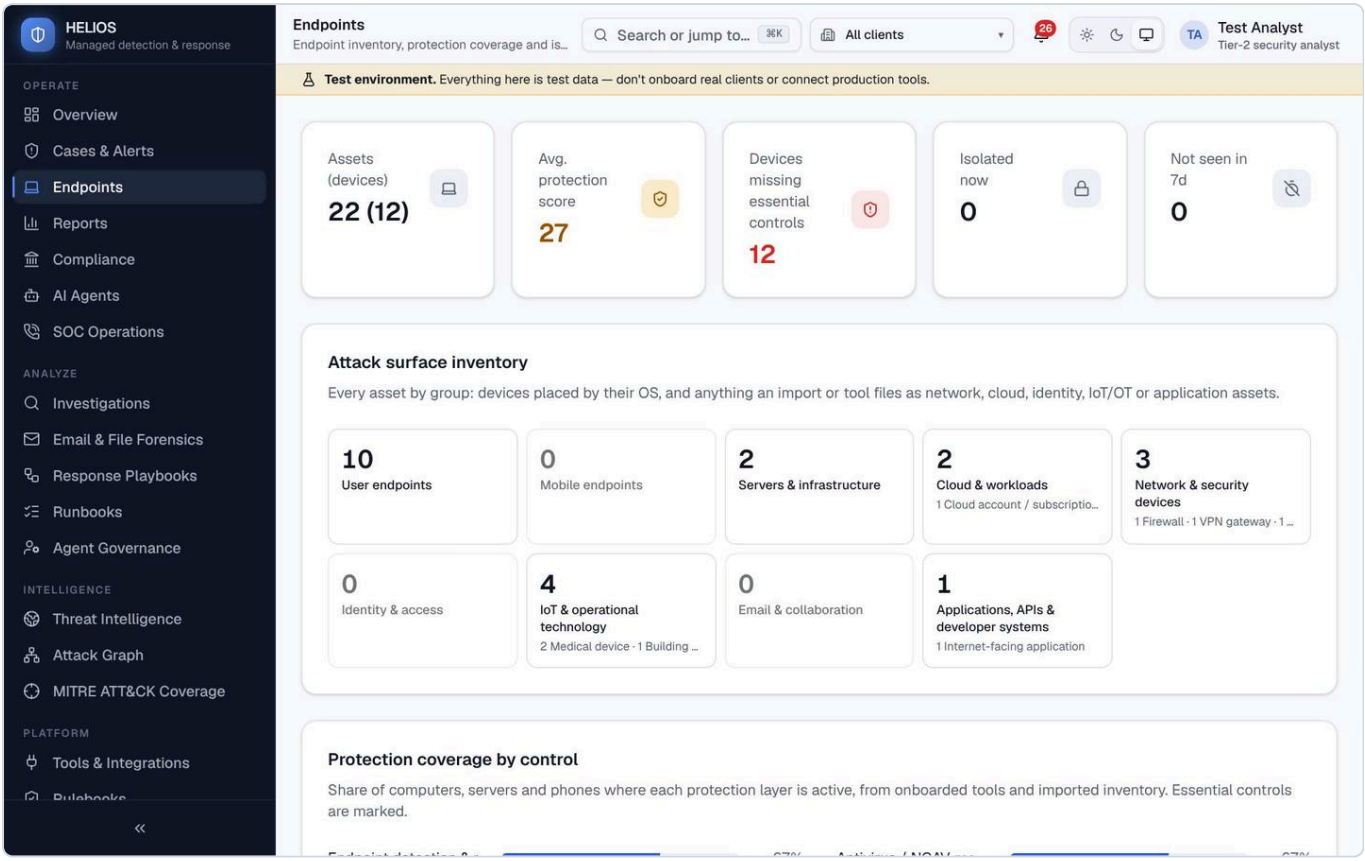


Figure 10. Endpoints, as seen by a tier-2 security analyst (test environment).

How to use it

- The tiles count assets, the average protection score, devices missing essential controls, devices isolated now and devices not seen for seven days.
- Attack surface inventory groups every asset; Protection coverage by control shows the share of computers, servers and phones where each protection layer (EDR, antivirus, disk encryption, MFA...) is active, with the essential ones marked.

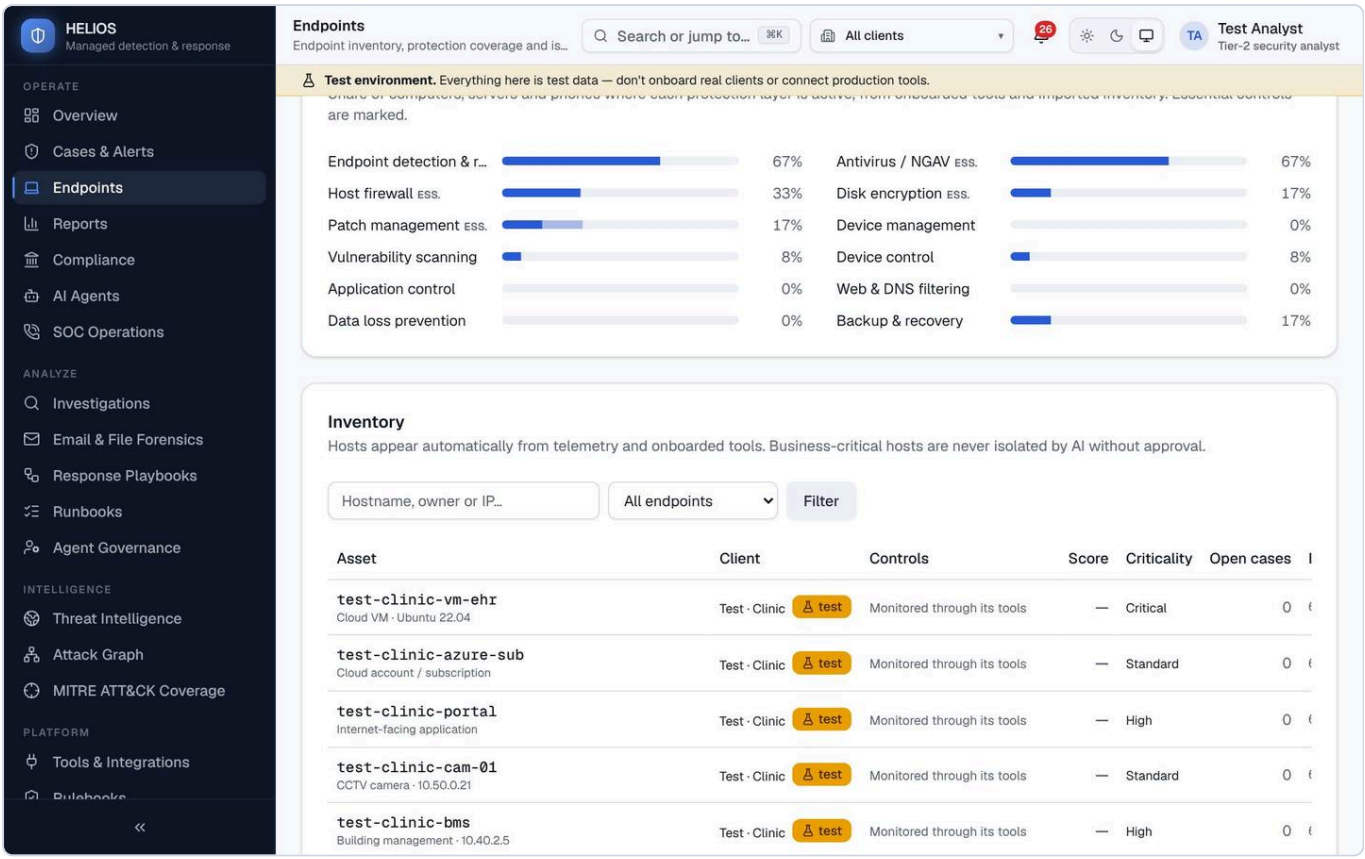


Figure 11. Endpoints, further down the page.

Reports

How the service performed over a period: volume, response times, SLA compliance, AI performance and

analyst workload, each compared with the period before.

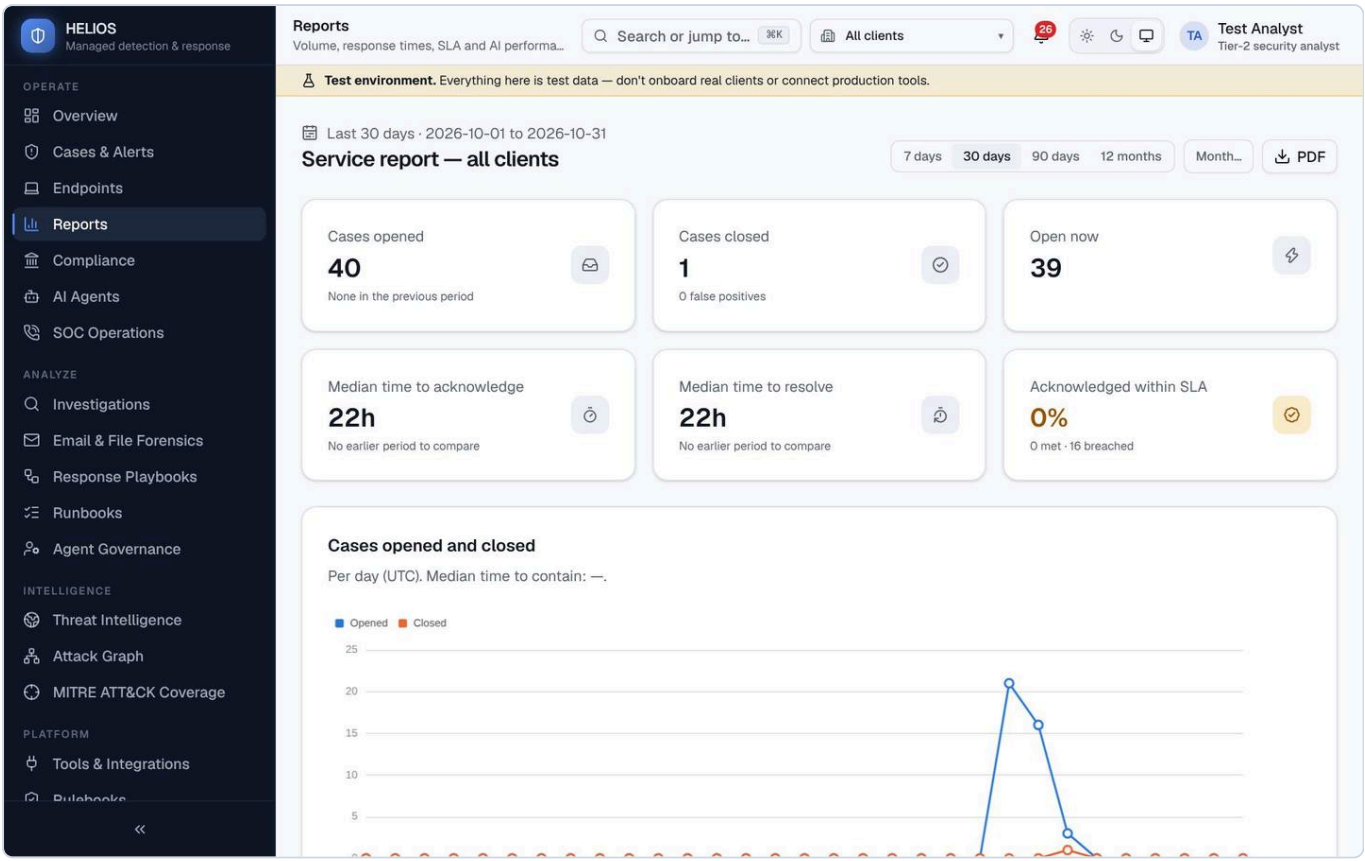


Figure 12. Reports, as seen by a tier-2 security analyst (test environment).

How to use it

- Pick a period or a month. The figures include median time to acknowledge, contain and resolve, acknowledgement SLA compliance, the mix by severity, source, domain and client, top ATT&CK techniques and indicators, AI runs and analyst agreement, and containment outcomes.
- Download the service report as a PDF for the client.

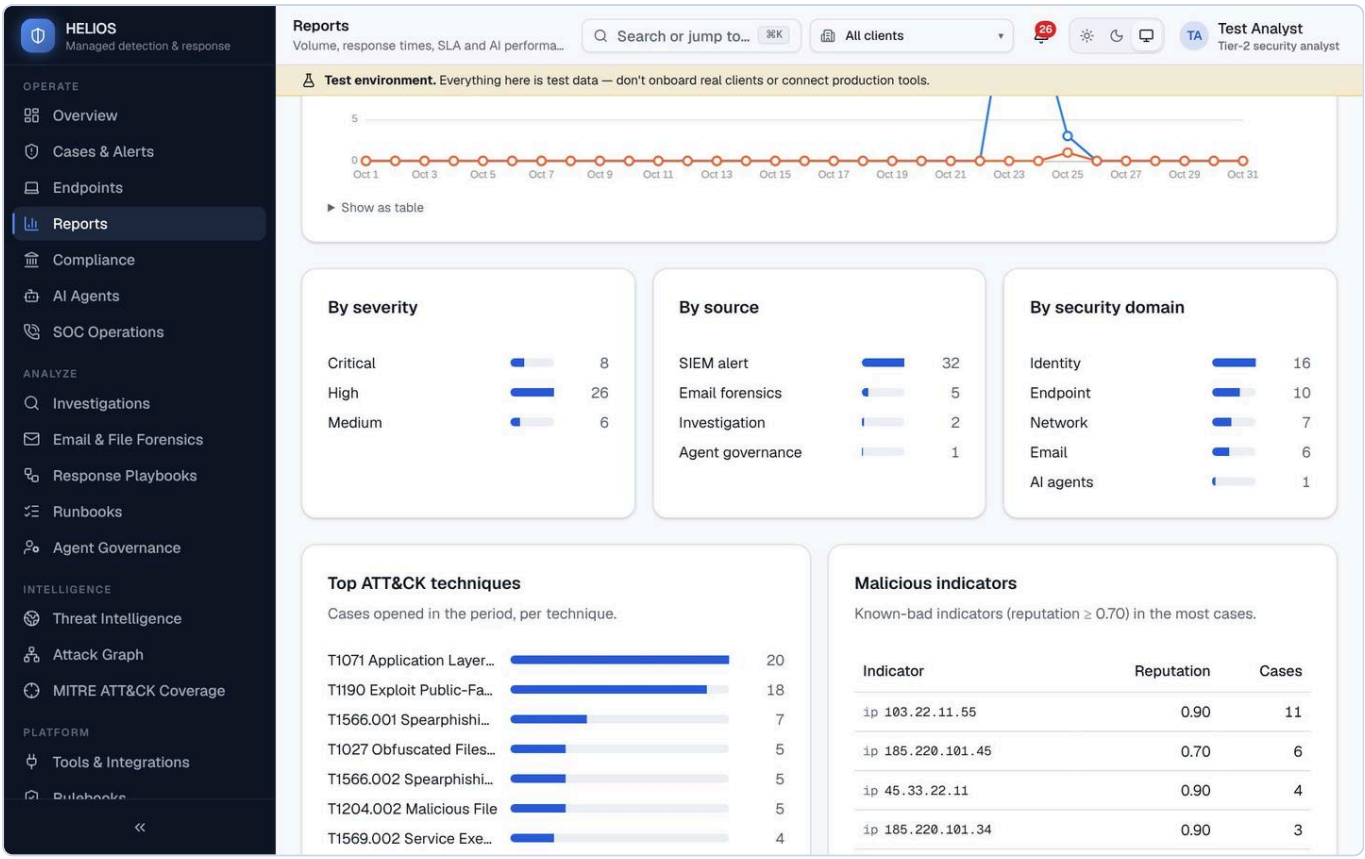


Figure 13. Reports, further down the page.

Compliance

Regulatory clocks and evidence: CERT-In's 6-hour incident reports, the DPDP breach notices, and the

ISO/IEC 27001 evidence map.

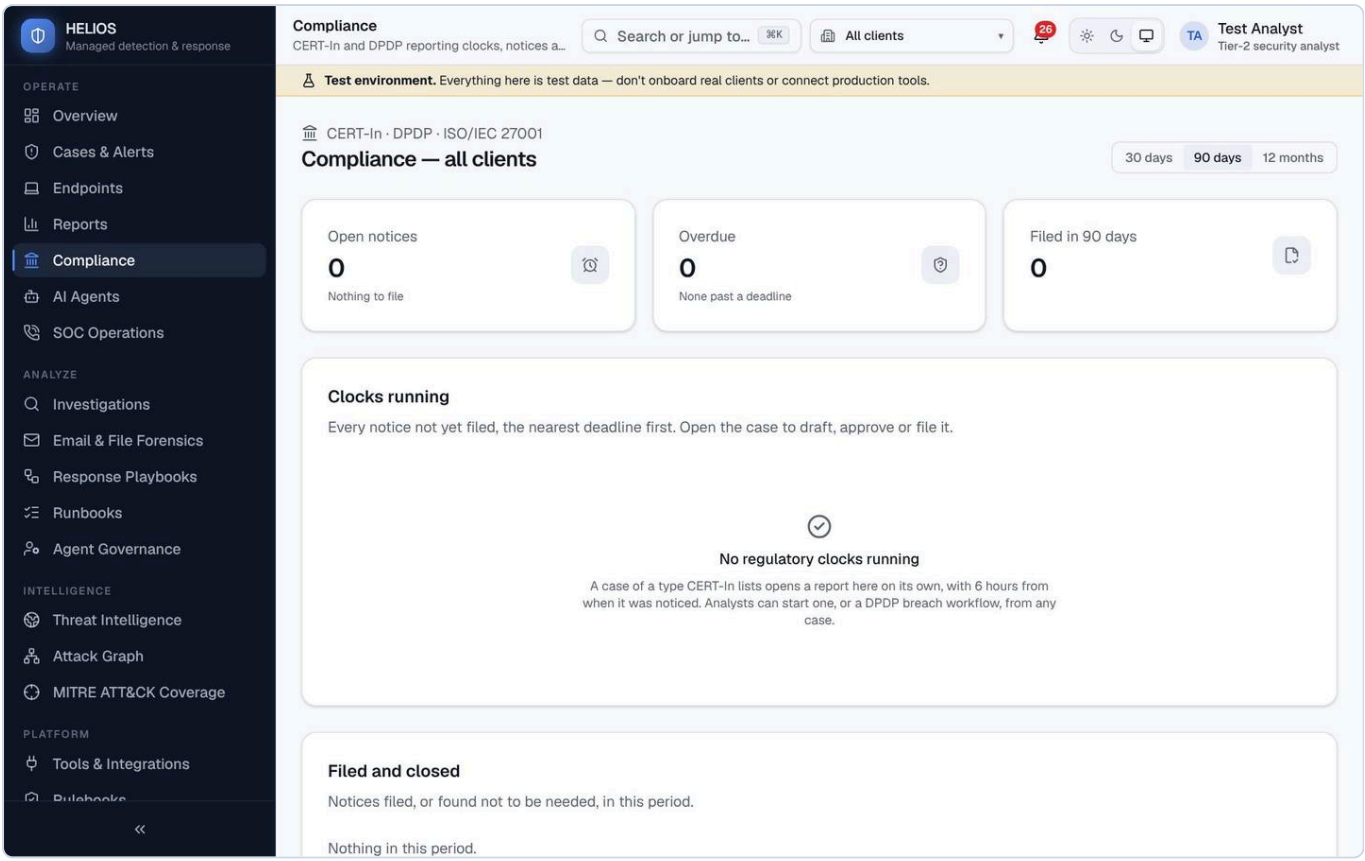


Figure 14. Compliance, as seen by a tier-2 security analyst (test environment).

How to use it

- Clocks running lists every notice not yet filed, nearest deadline first. CERT-In: incidents of the 20 listed types are reported within 6 hours of being noticed. DPDP: the Board and each affected person are told without delay, and the Board gets a detailed report within 72 hours.
- Open the case to draft, approve or file a notice. Each notice is drafted from the case, approved by the client, then filed (emailed to CERT-In with its PDF, or recorded with the regulator's reference).
- The ISO 27001 view maps 15 Annex A controls to what HELIOS's records show over the period, and shows gaps (controls with nothing behind them) plainly.
- Switch the period (30 days, 90 days, 12 months) at the top.

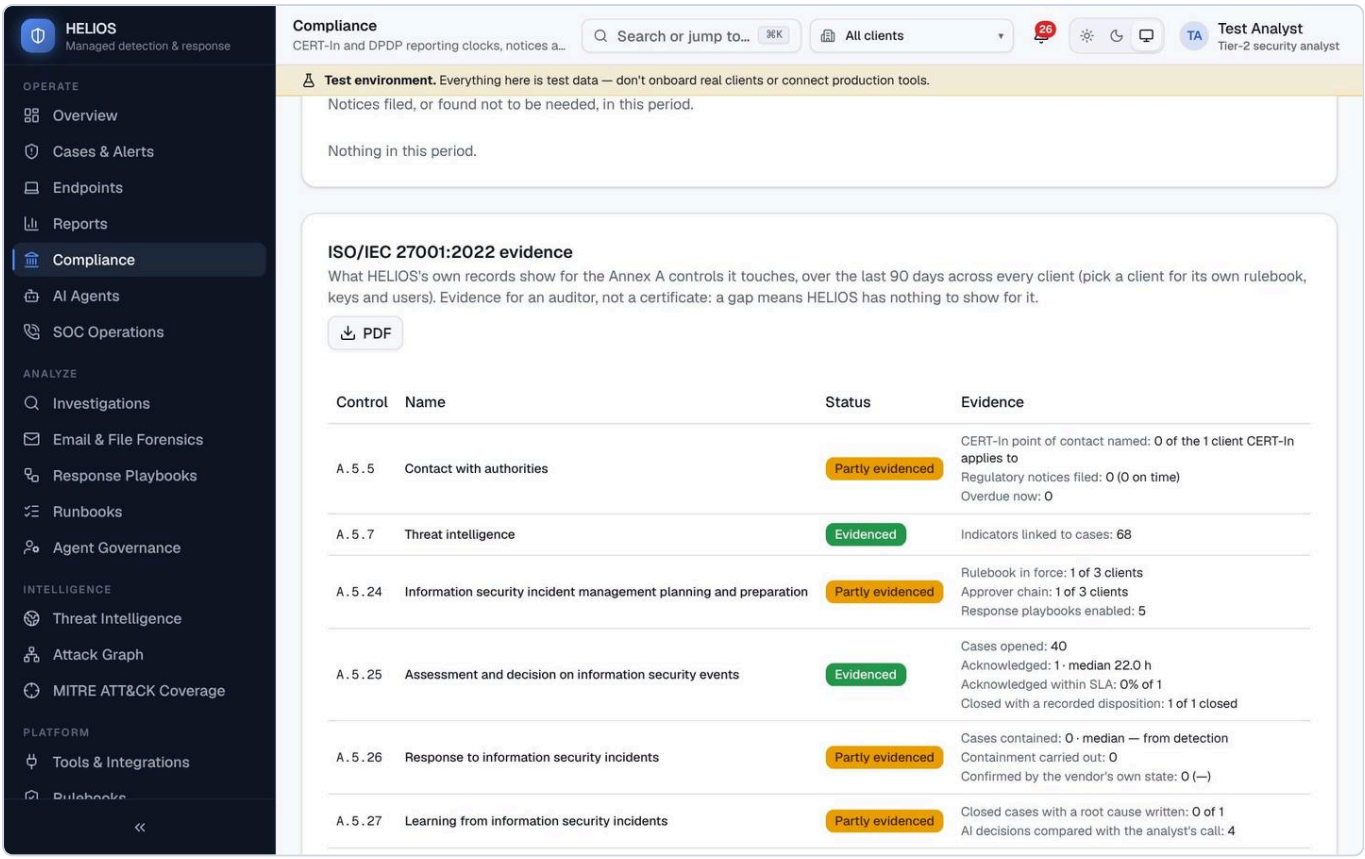


Figure 15. Compliance, further down the page.

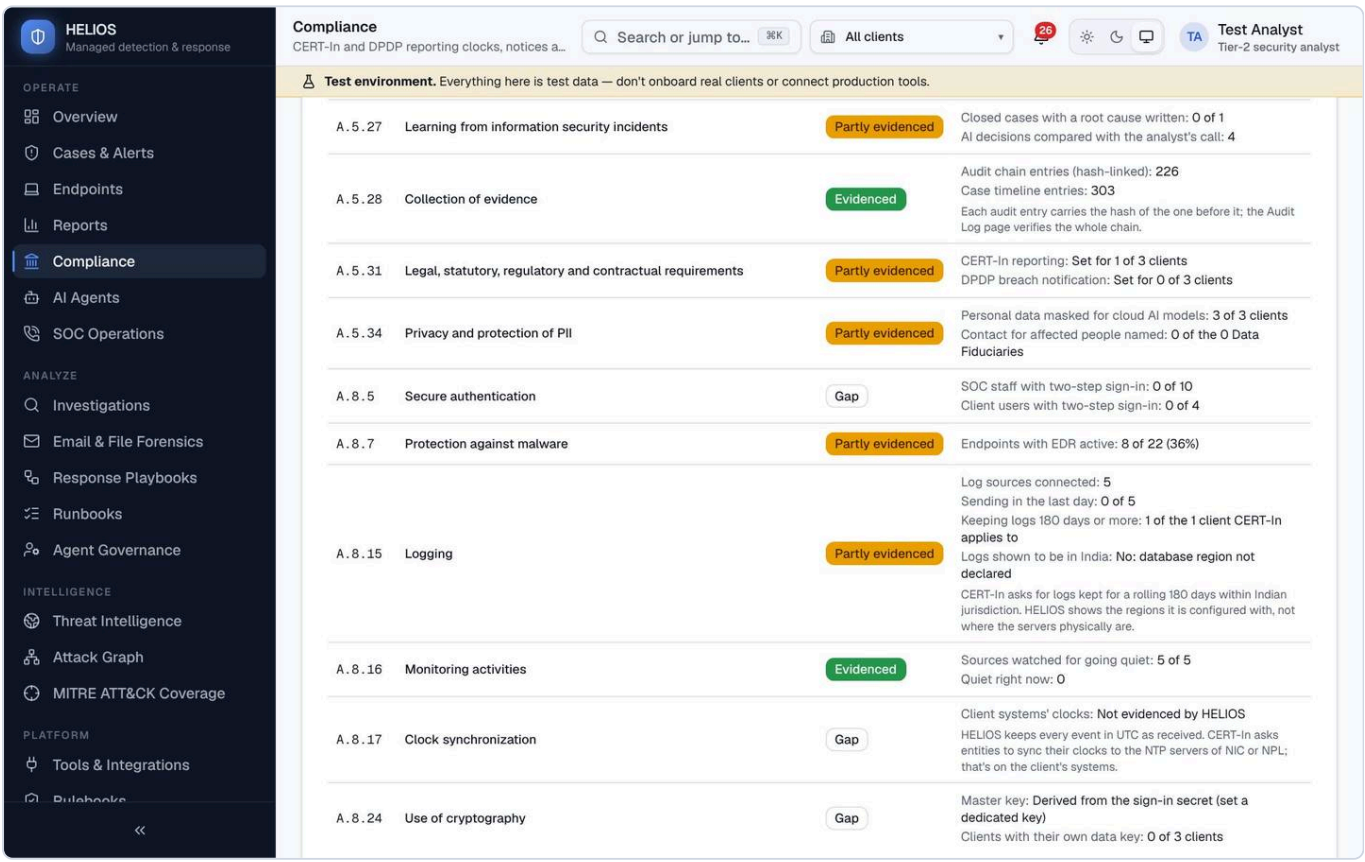


Figure 16. Compliance, further down the page.

AI Agents

The AI agents that work cases: what they did, how far they may act on their own, which models run them, and how those models scored.

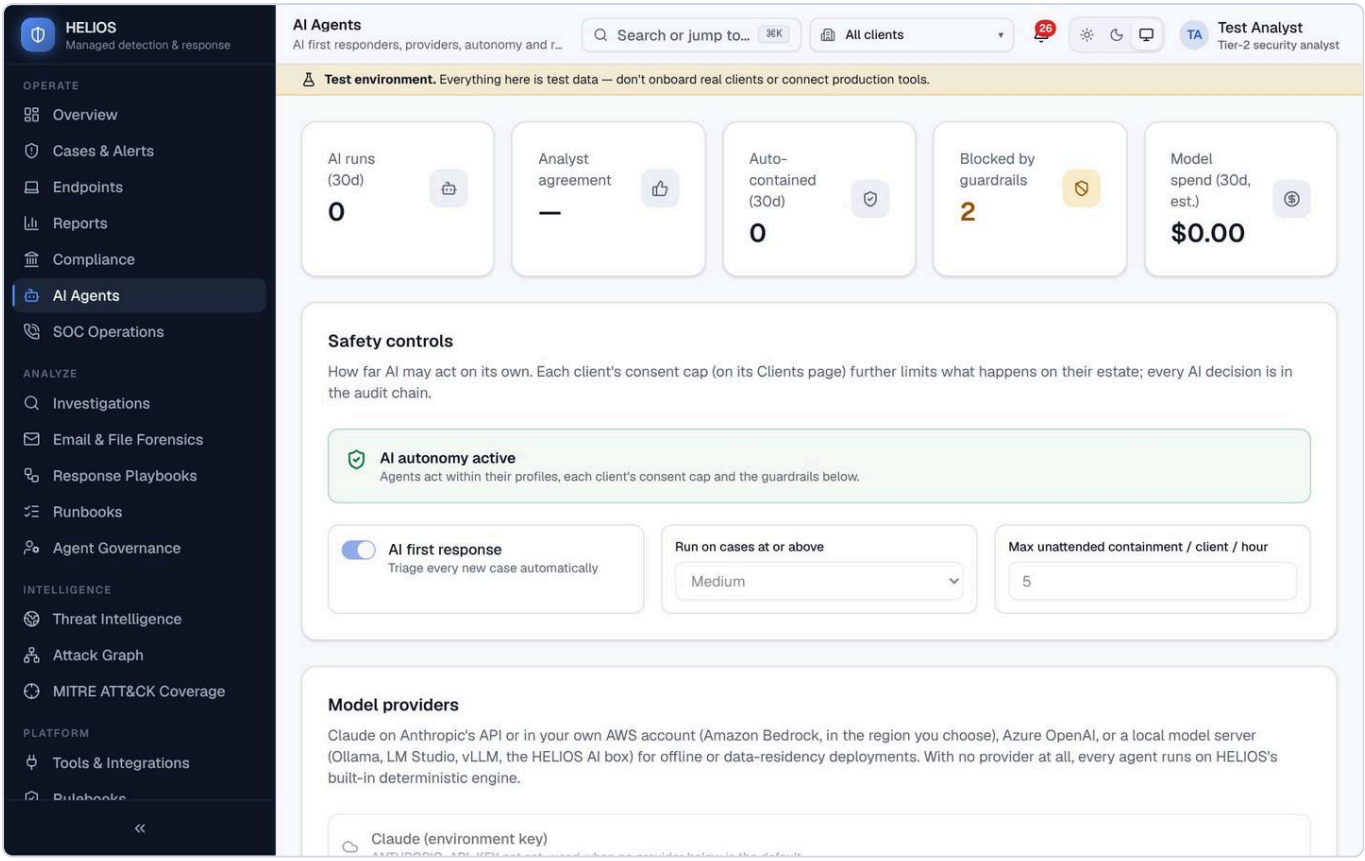


Figure 17. AI Agents, as seen by a tier-2 security analyst (test environment).

How to use it

- The tiles show AI runs in the last 30 days, analyst agreement, actions contained automatically, actions the guardrails blocked and the estimated model spend (with the share of input read from the prompt cache).
- The run log lists recent runs with their verdict, model, tokens and outcome.

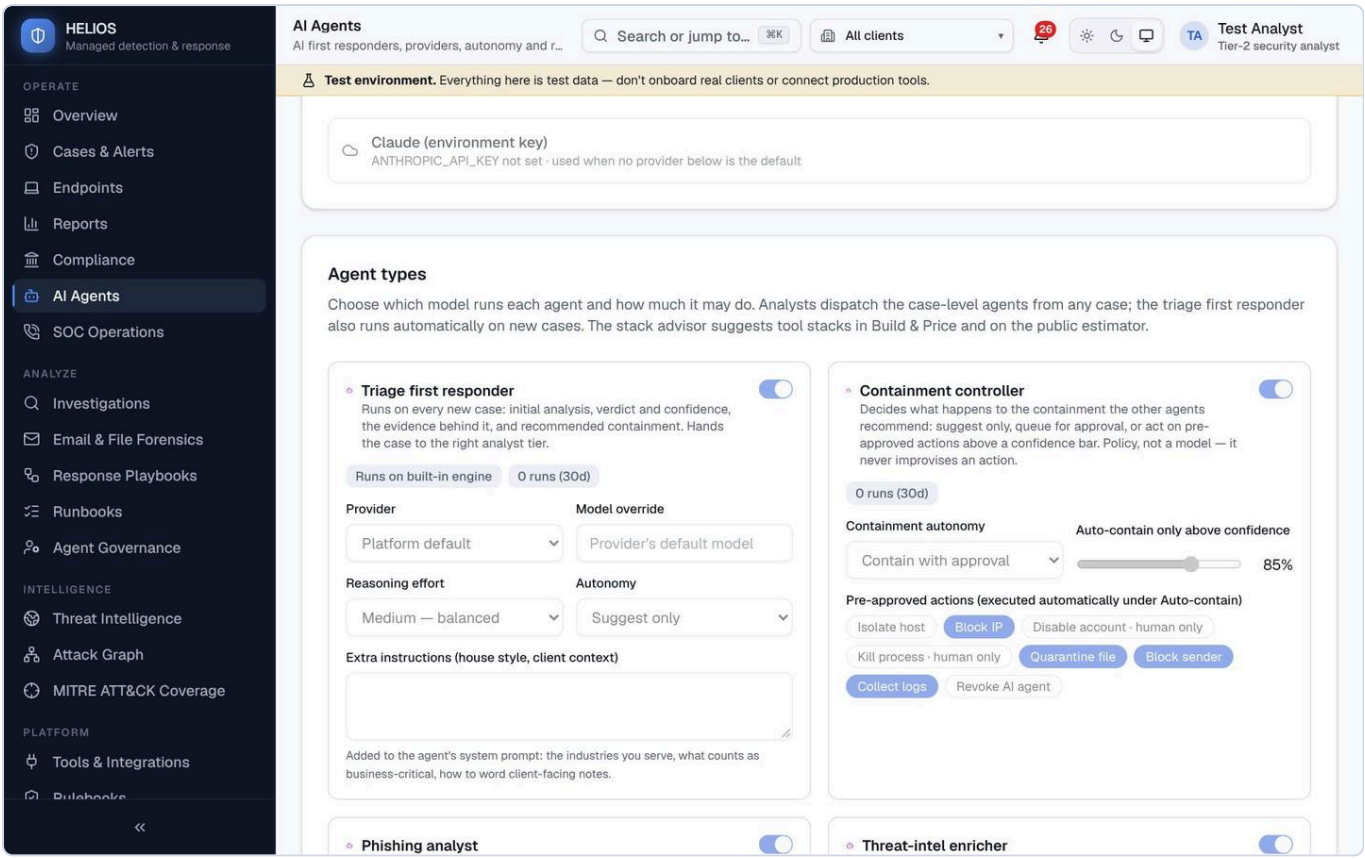


Figure 18. AI Agents, further down the page.

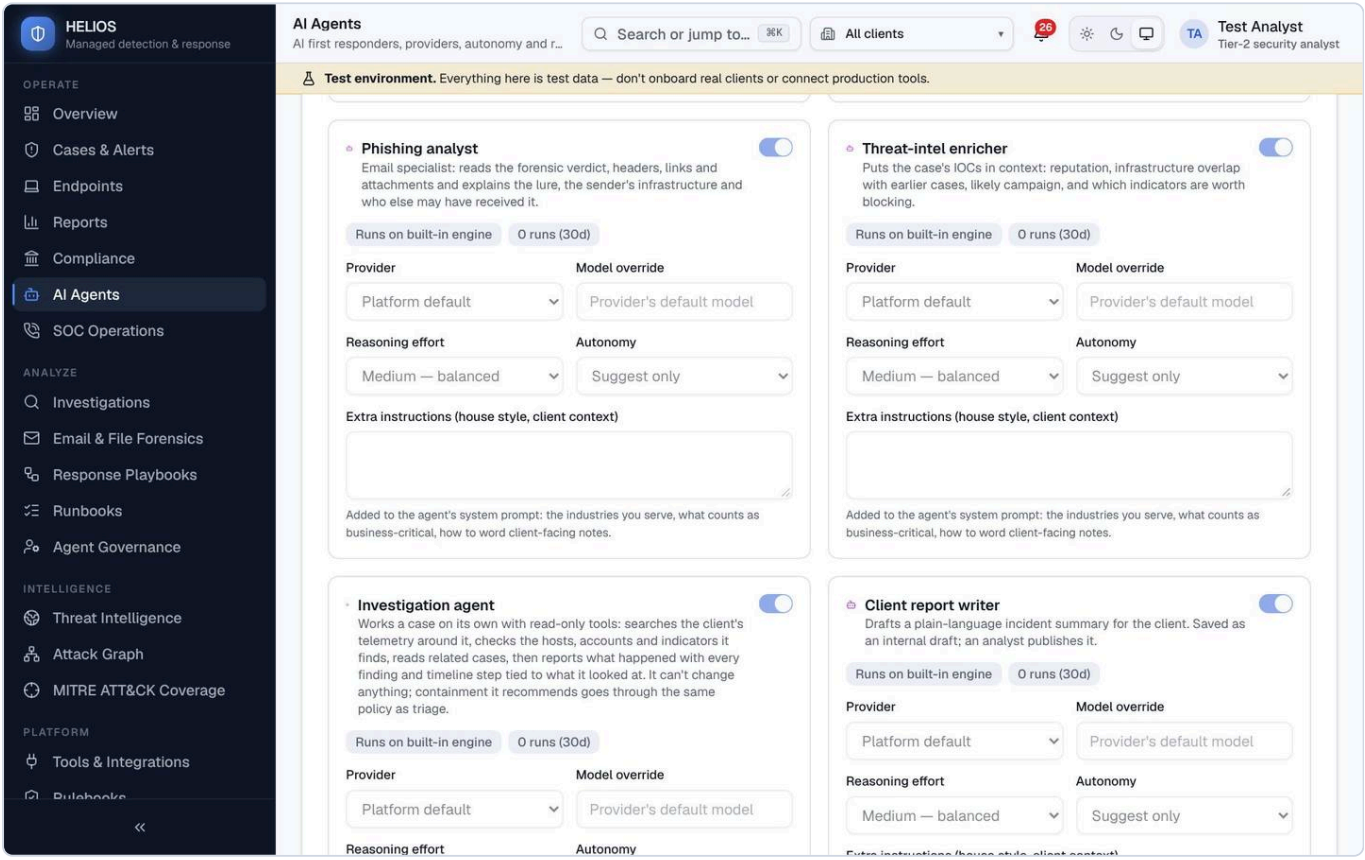


Figure 19. AI Agents, further down the page.

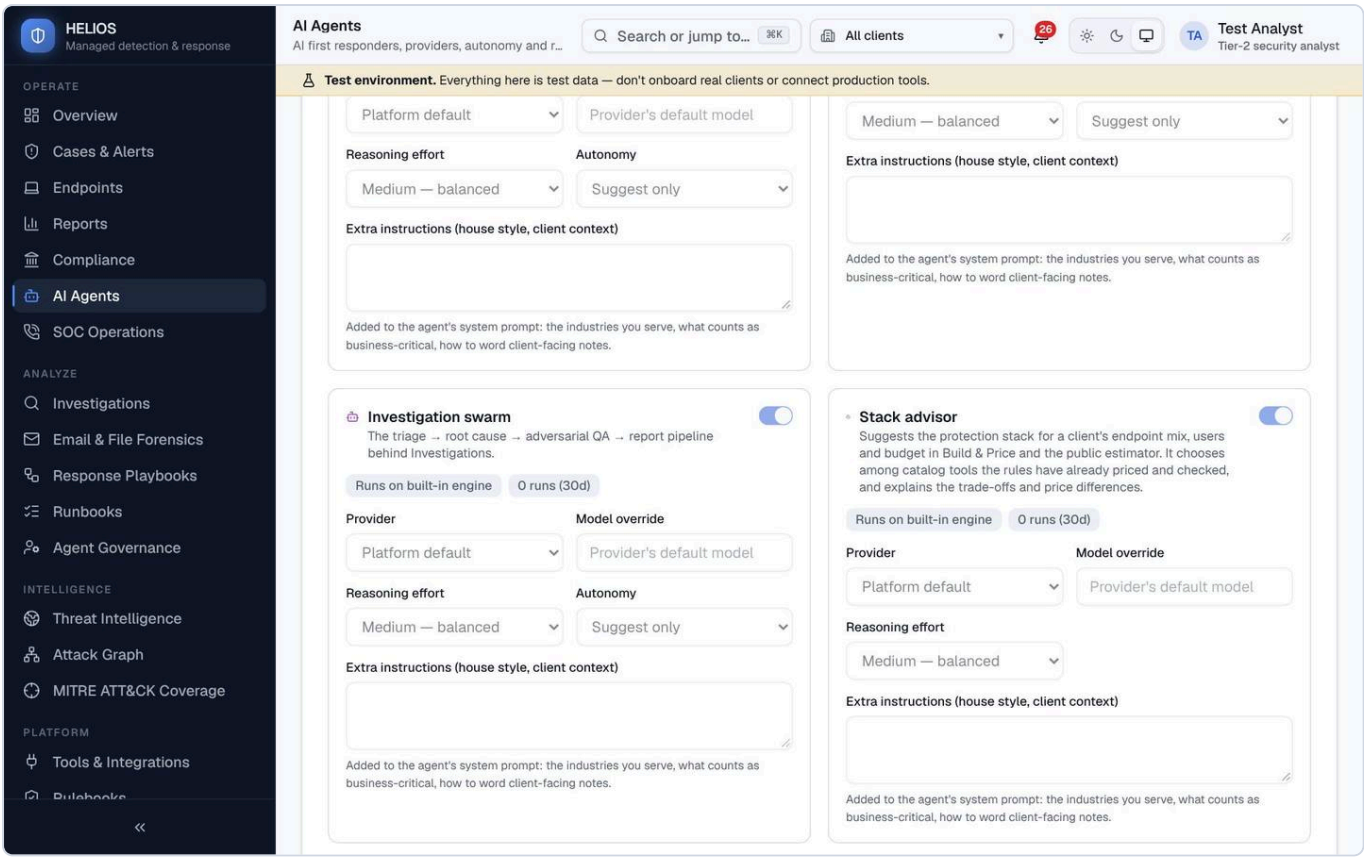


Figure 20. AI Agents, further down the page.

SOC Operations

Who is on call and the pages they get, the weekly QA of what the AI did on its own, and the daily digests.

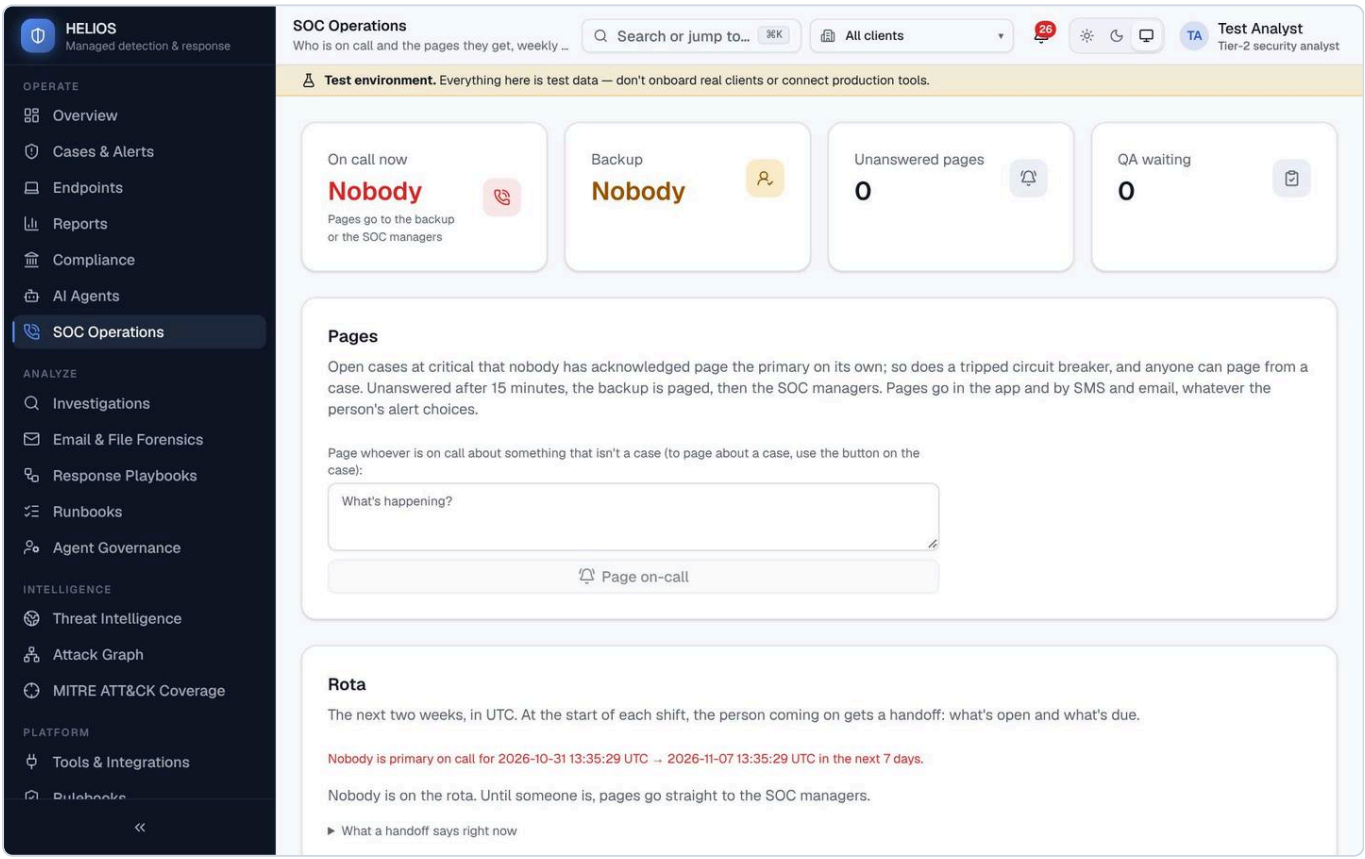


Figure 21. SOC Operations, as seen by a tier-2 security analyst (test environment).

How to use it

- The tiles show who is primary and backup on call now, unanswered pages and QA reviews waiting.
- Pages: open critical cases nobody acknowledged page the primary on their own; unanswered after the acknowledgement window (15 minutes by default) the backup is paged, then the SOC managers. Pages go in the app and by SMS and email. Page on-call sends a page about something that isn't a case.

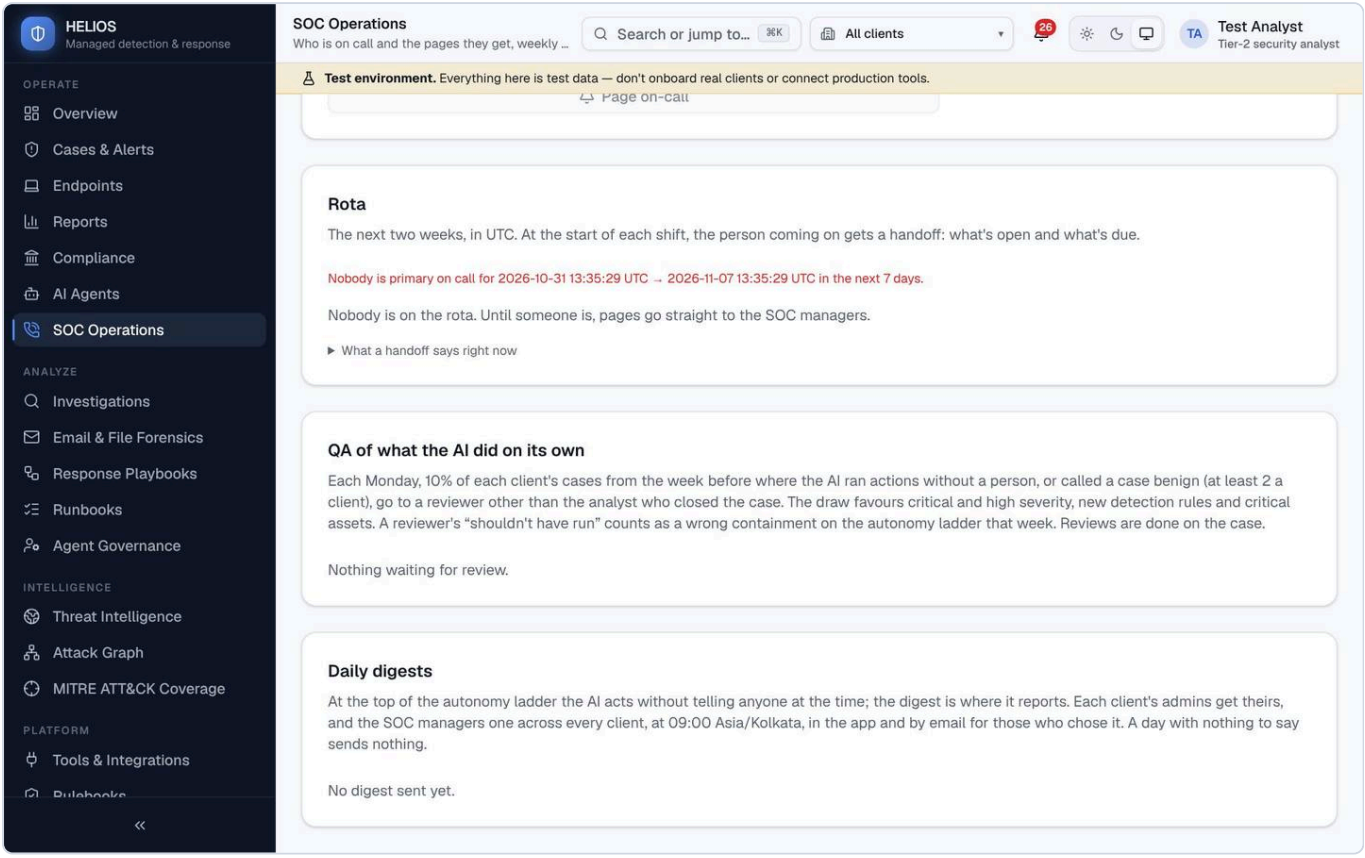


Figure 22. SOC Operations, further down the page.

Investigations

Upload logs from an incident and get a root-cause analysis: an agent pipeline (triage, root cause,

adversarial QA, report) works through them and opens a case with the findings.

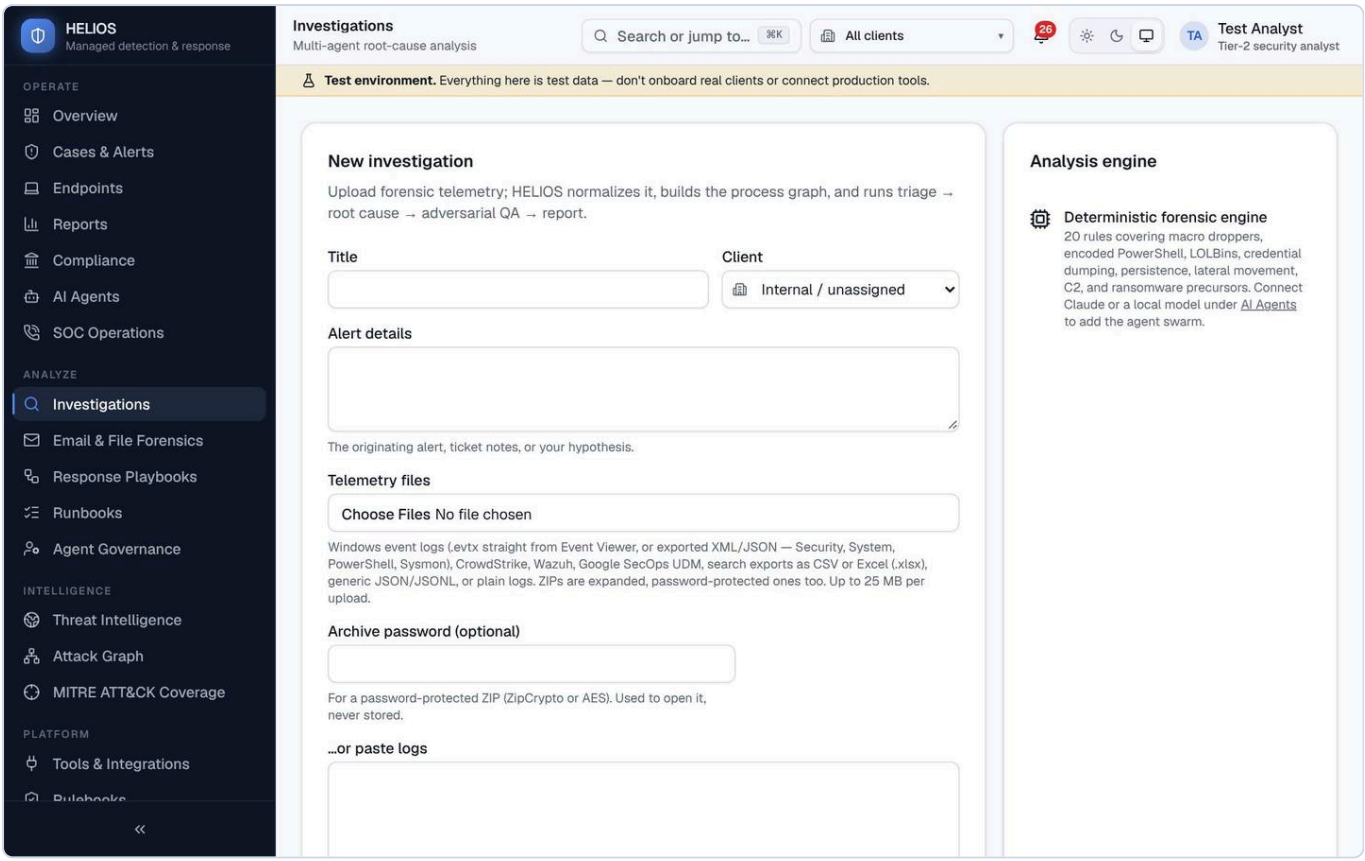


Figure 23. Investigations, as seen by a tier-2 security analyst (test environment).

How to use it

- New investigation: a title, the client, the alert details (the originating alert, ticket notes or your hypothesis), and the telemetry files: Windows event logs (.evtx, or exported XML/JSON), EDR and SIEM exports, CSV or Excel, JSON/JSONL, plain logs, or a .zip (password-protected ones too, with the archive password, which isn't stored). Up to 25 MB per upload; or paste the logs.
- HELIOS normalizes the telemetry, builds the process graph and runs triage, root cause, adversarial QA and the report. Without a model it uses its deterministic forensic engine (20 rules for droppers, encoded PowerShell, living-off-the-land binaries, credential dumping, persistence, lateral movement, C2 and ransomware precursors).
- Open a past investigation to read its timeline, root cause, the adversarial QA's challenges and the report.

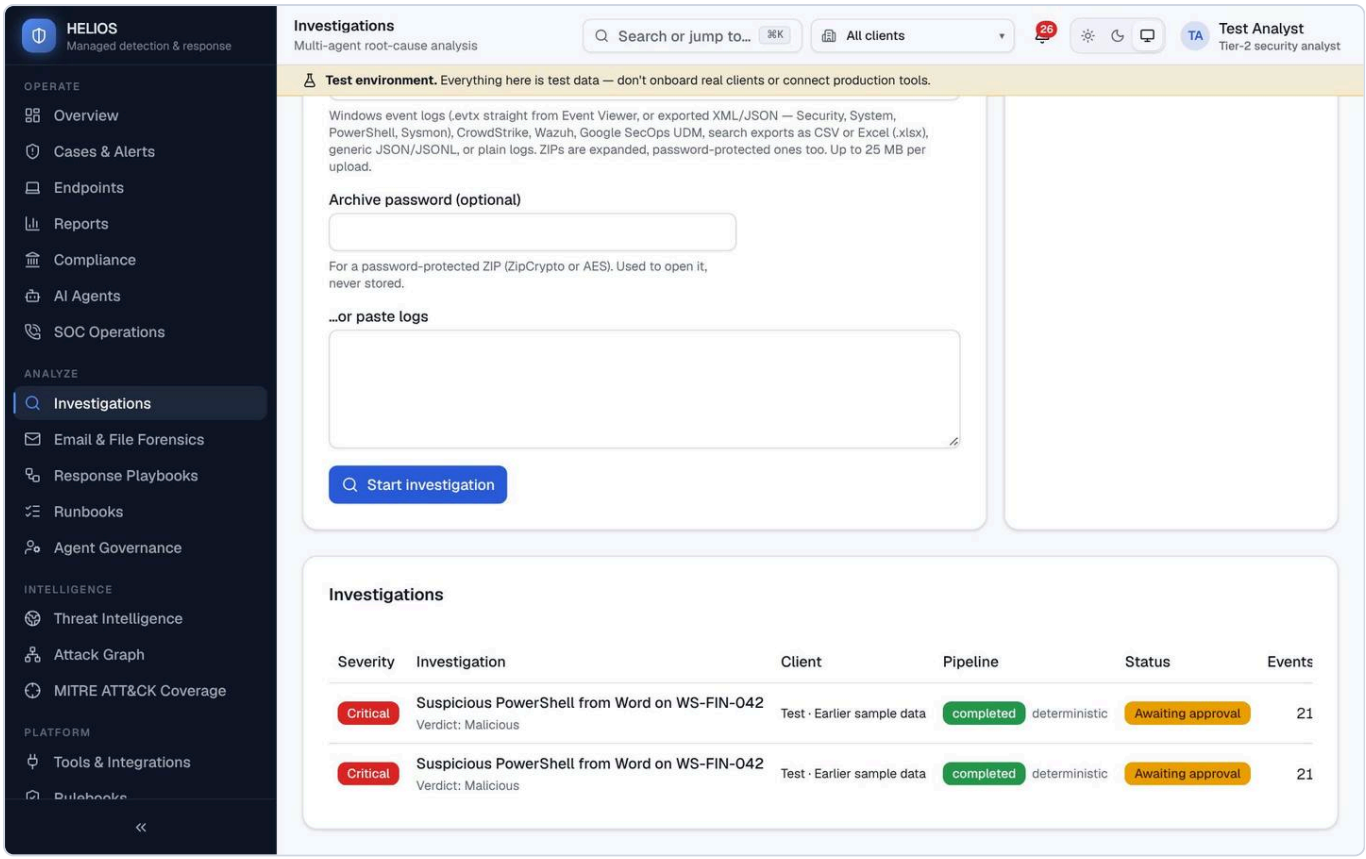


Figure 24. Investigations, further down the page.

Email & File Forensics

Submit a suspicious email or file for analysis: header authentication, links, attachments, macros, archives

and images are examined and scored to a verdict.

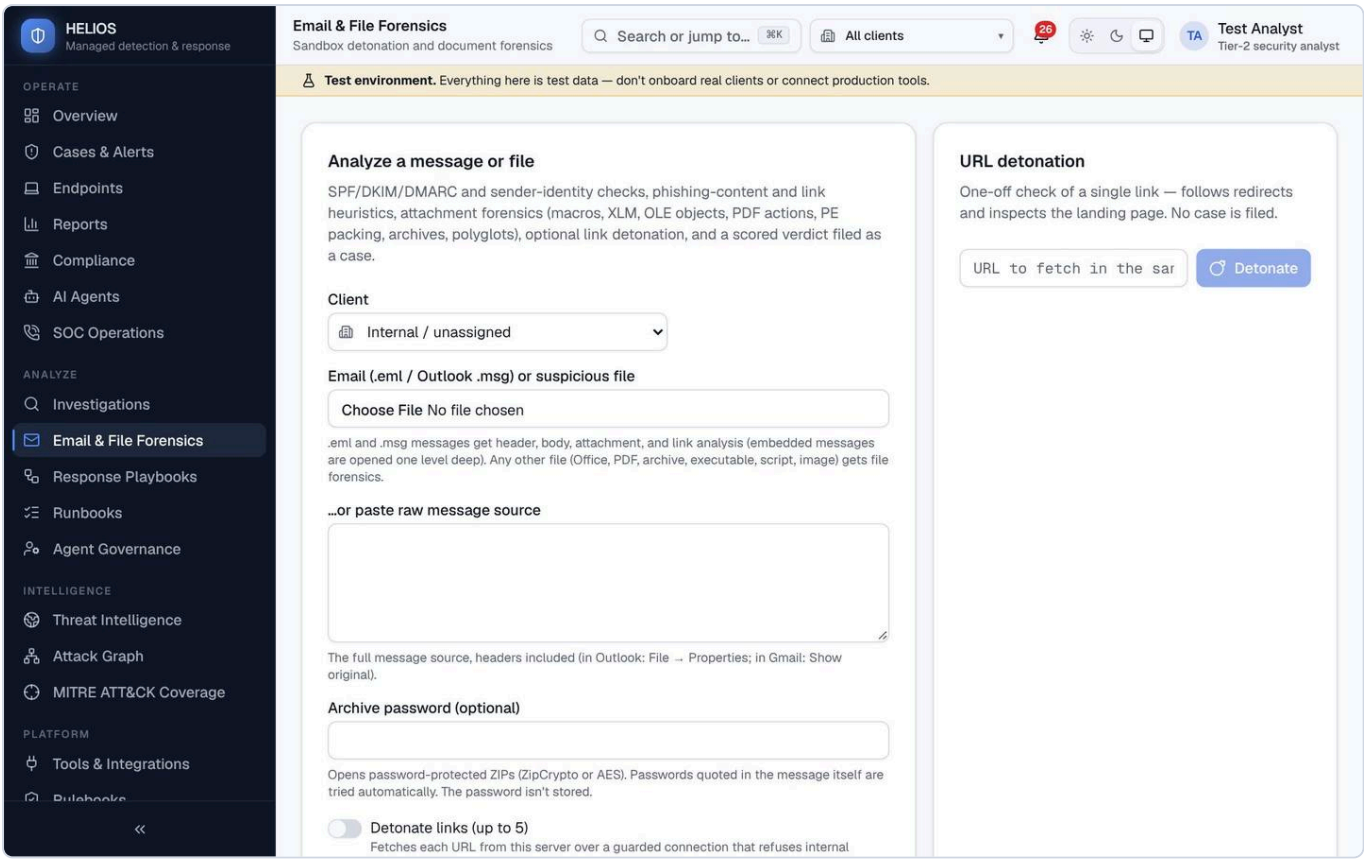


Figure 25. Email & File Forensics, as seen by a tier-2 security analyst (test environment).

How to use it

- Analyze a message or file: choose the client, then the .eml or Outlook .msg (or any suspicious file), or paste the raw message source. Give an archive password if it's a protected ZIP (it isn't stored), and switch on Detonate links (up to 5) if you want links fetched through the guarded connection.
- HELIOS checks the sender's authentication (SPF, DKIM, DMARC) and identity, phishing content and links, attachments (macros, OLE objects, PDF actions, packing, archives, polyglots) and text and QR codes in images. URL detonation on the right checks a single link without filing a case.
- The verdict comes from the total score across the header, content, attachment, sandbox, intel and sender checks, against the suspicious and malicious thresholds set in Settings. A malicious result opens a case, where the message can be purged from mailboxes and the sender blocked.

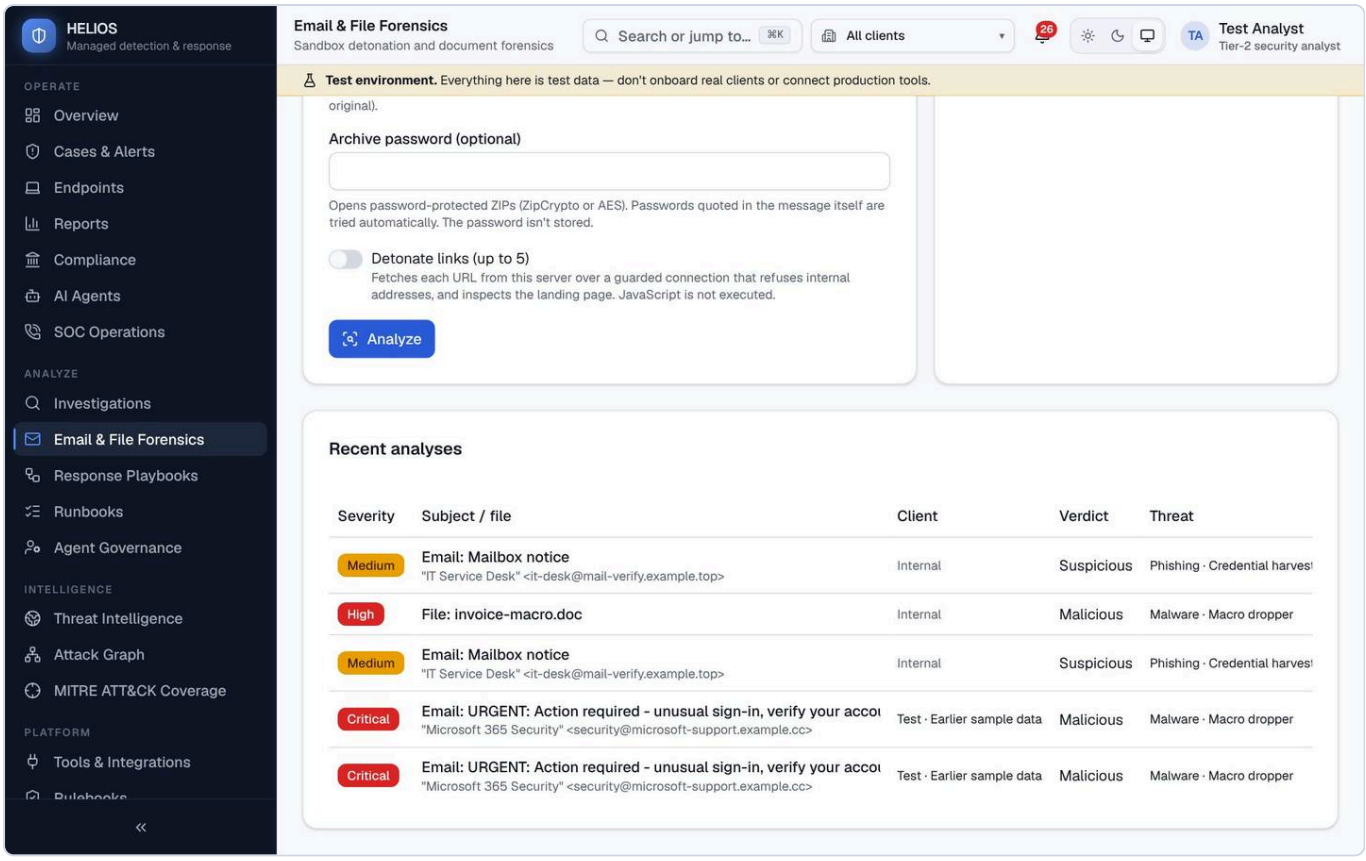


Figure 26. Email & File Forensics, further down the page.

Response Playbooks

Automated response workflows: what triggers them and the steps they run (containment, collecting logs,

notifying, opening tickets).

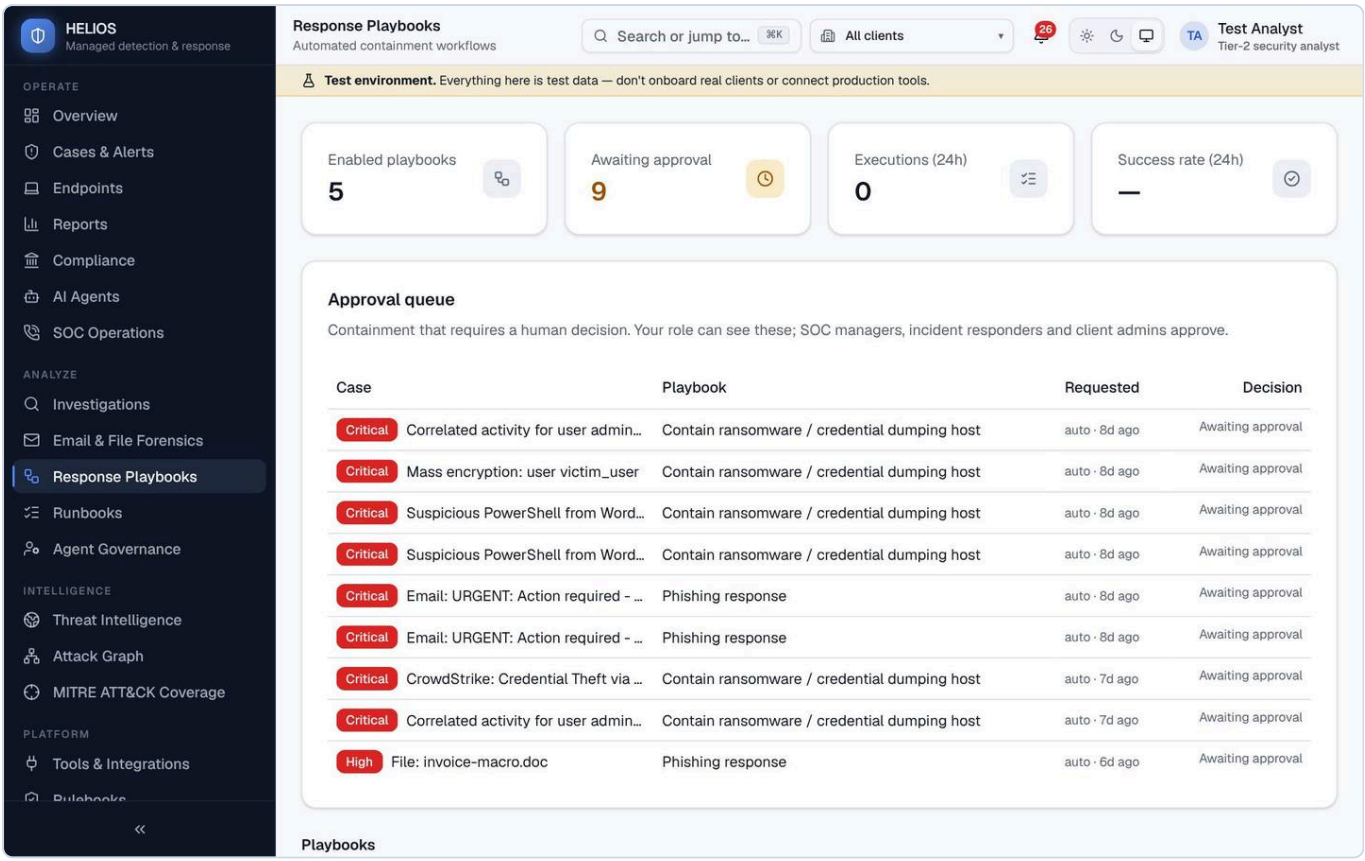


Figure 27. Response Playbooks, as seen by a tier-2 security analyst (test environment).

How to use it

- Each playbook shows its trigger, its steps and whether it is enabled. Its runs appear on the cases they ran on; steps that need approval wait in Approvals, and steps without a connector are left to do by hand.

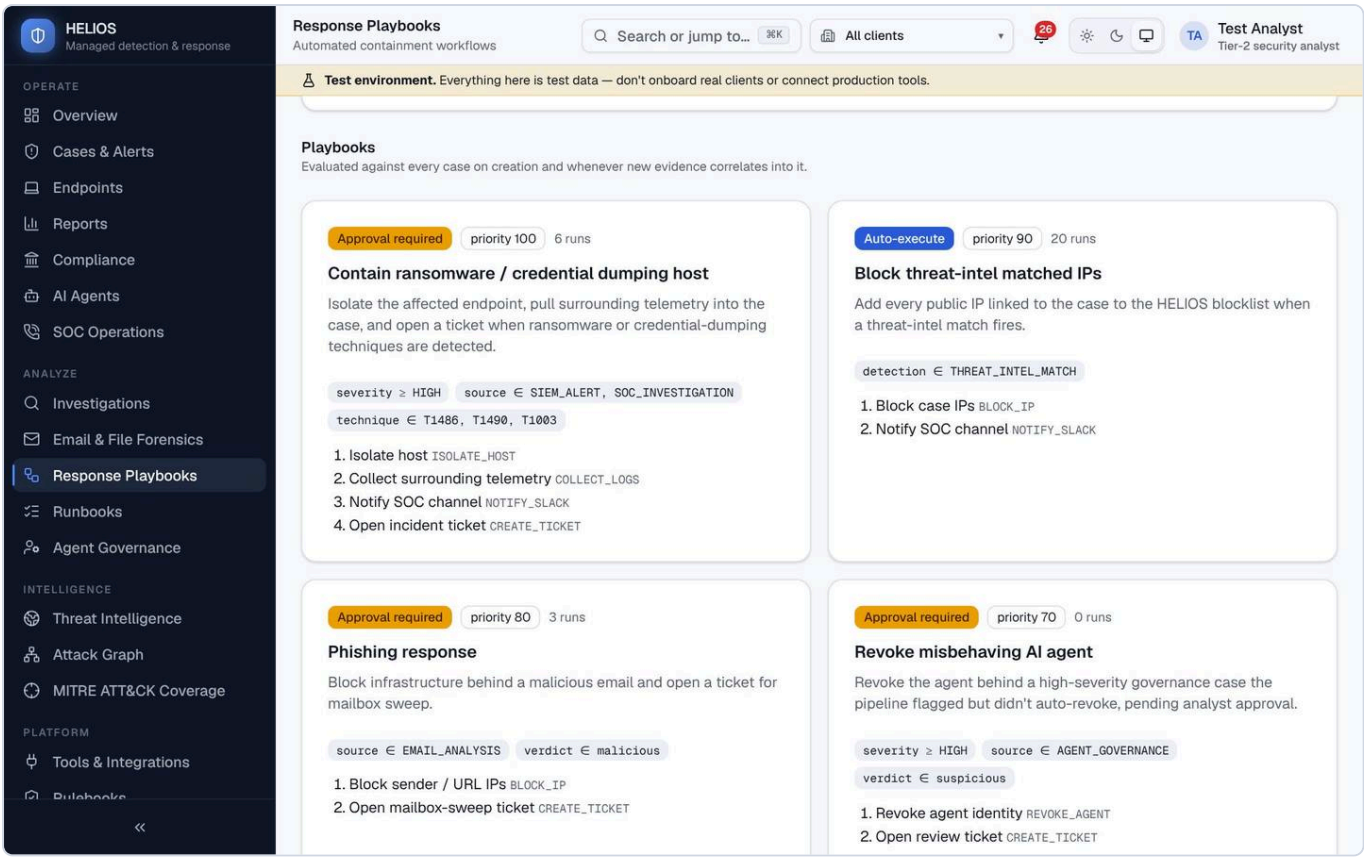


Figure 28. Response Playbooks, further down the page.

Runbooks

Procedures for analysts, attached to the cases they match and worked as checklists on the case.

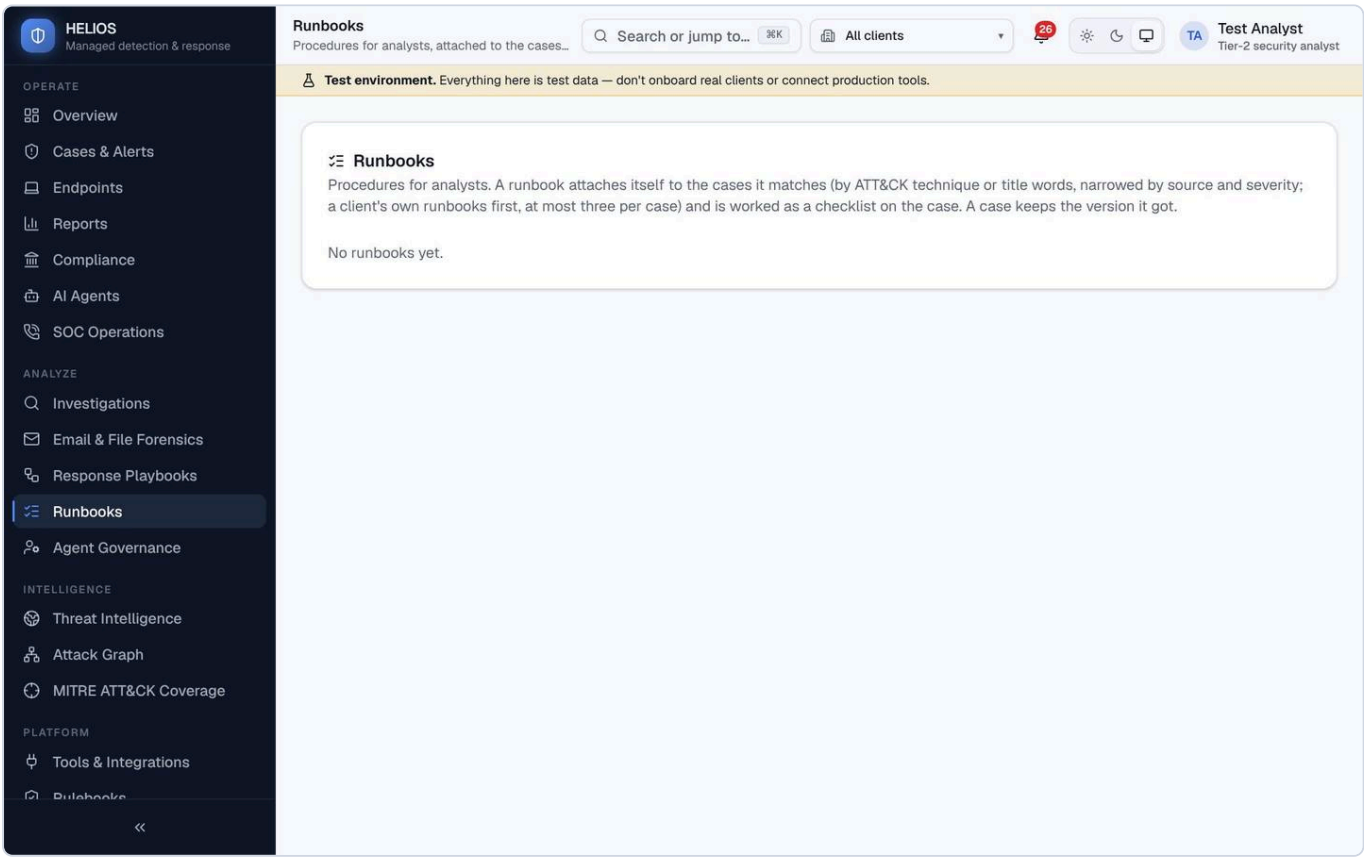


Figure 29. Runbooks, as seen by a tier-2 security analyst (test environment).

How to use it

- A runbook matches cases by ATT&CK technique or title words, narrowed by source and minimum severity. Up to three matching runbooks attach to a case when it is opened, the client's own first.

Agent Governance

Govern the AI agents your clients run: identities and tokens, the policies they act under, their trust, and revocation.

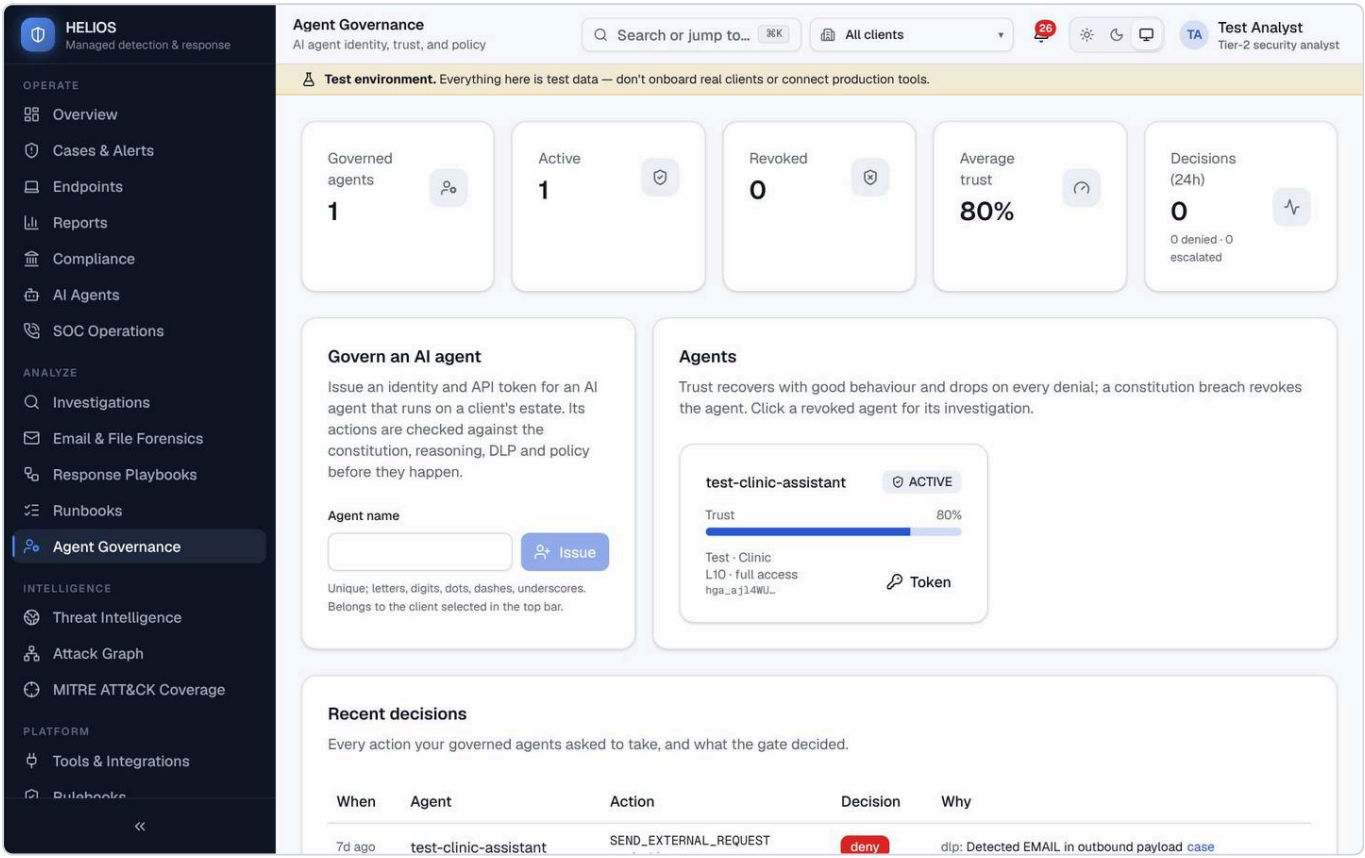


Figure 30. Agent Governance, as seen by a tier-2 security analyst (test environment).

How to use it

- Govern an AI agent: enter the agent's name and Issue, for the client selected in the top bar. Its token is shown once. The agent asks HELIOS before each action; constitution, reasoning, data-loss and policy checks decide, and Recent decisions shows each request and why it was allowed or denied.
- Trust recovers with good behaviour and drops on every denial.
- A breach revokes the agent and opens a case. Revoke or restore an agent by hand here or from its case.

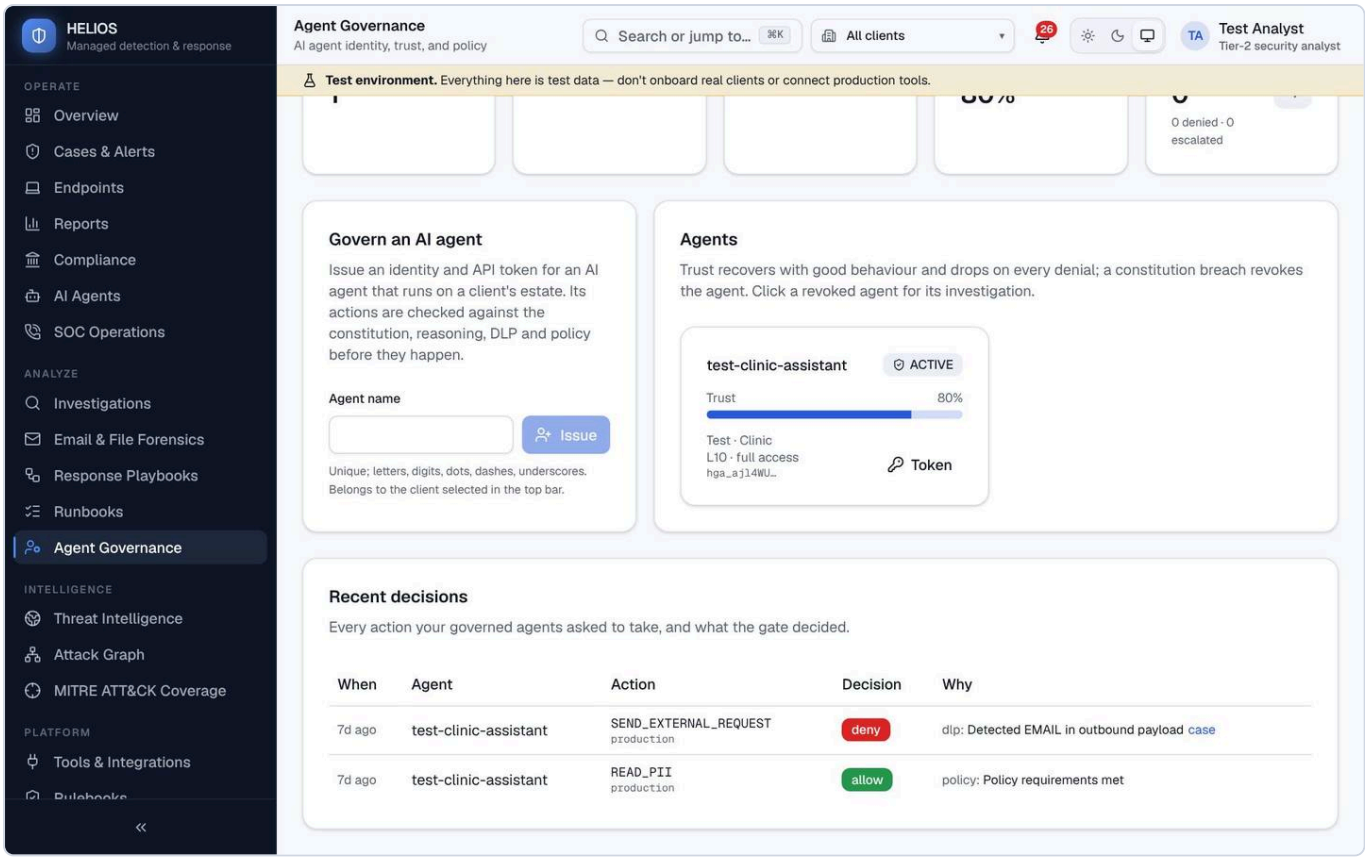


Figure 31. Agent Governance, further down the page.

Threat Intelligence

Look up indicators (IP, domain, URL, file hash, email address) against the SOC's indicator store, the

configured reputation providers and the threat feeds.

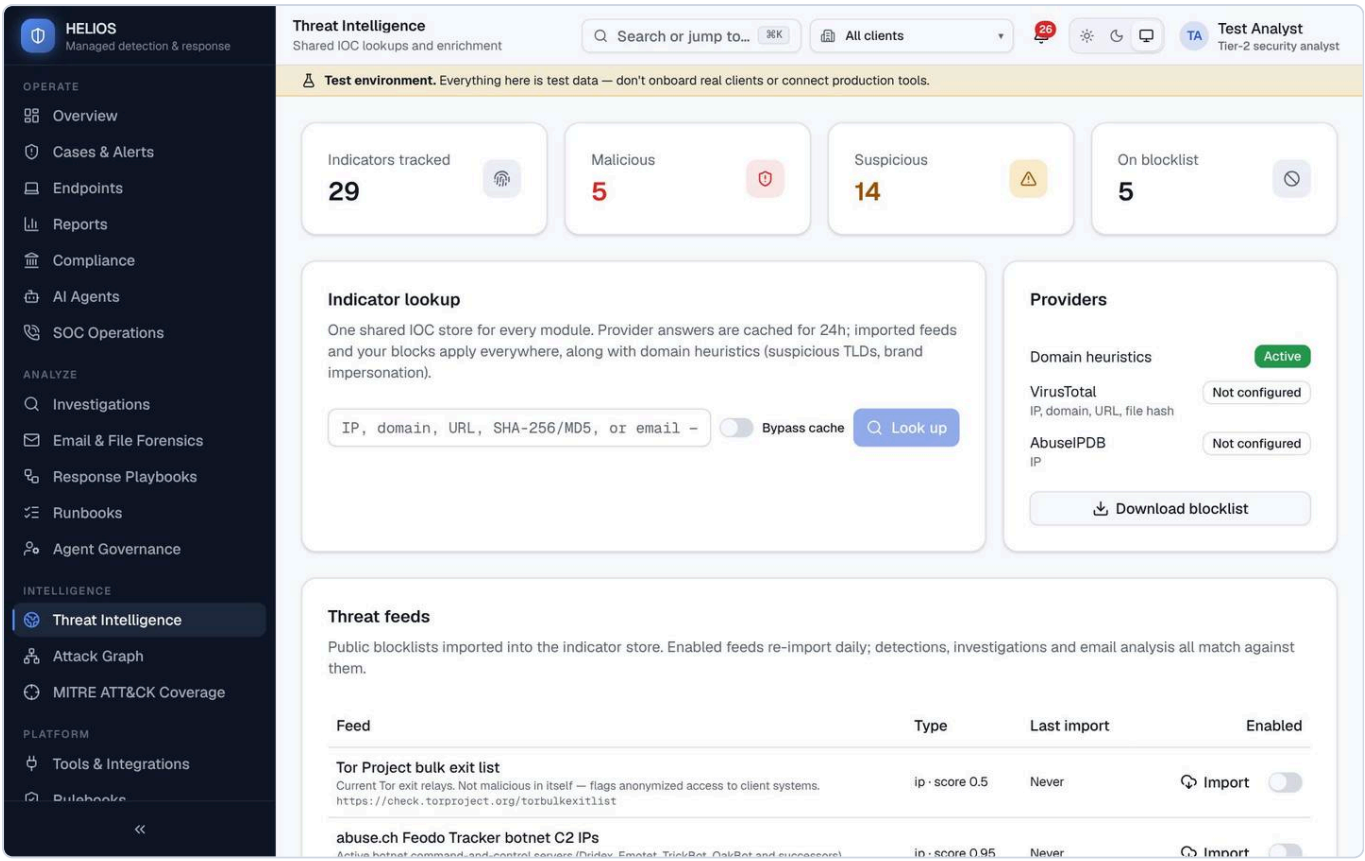


Figure 32. Threat Intelligence, as seen by a tier-2 security analyst (test environment).

How to use it

- Paste an indicator (defanged forms like hxxp and [.] are fine) and Look up. The result shows the verdict, score, sources, why, and for an IP where it is (country, city and network, from the server's GeoIP database). Results are cached for 24 hours; Bypass cache forces a fresh lookup.
- Block or unblock an indicator: the blocklist feeds detections and firewalls that pull it. Enable and import threat feeds.

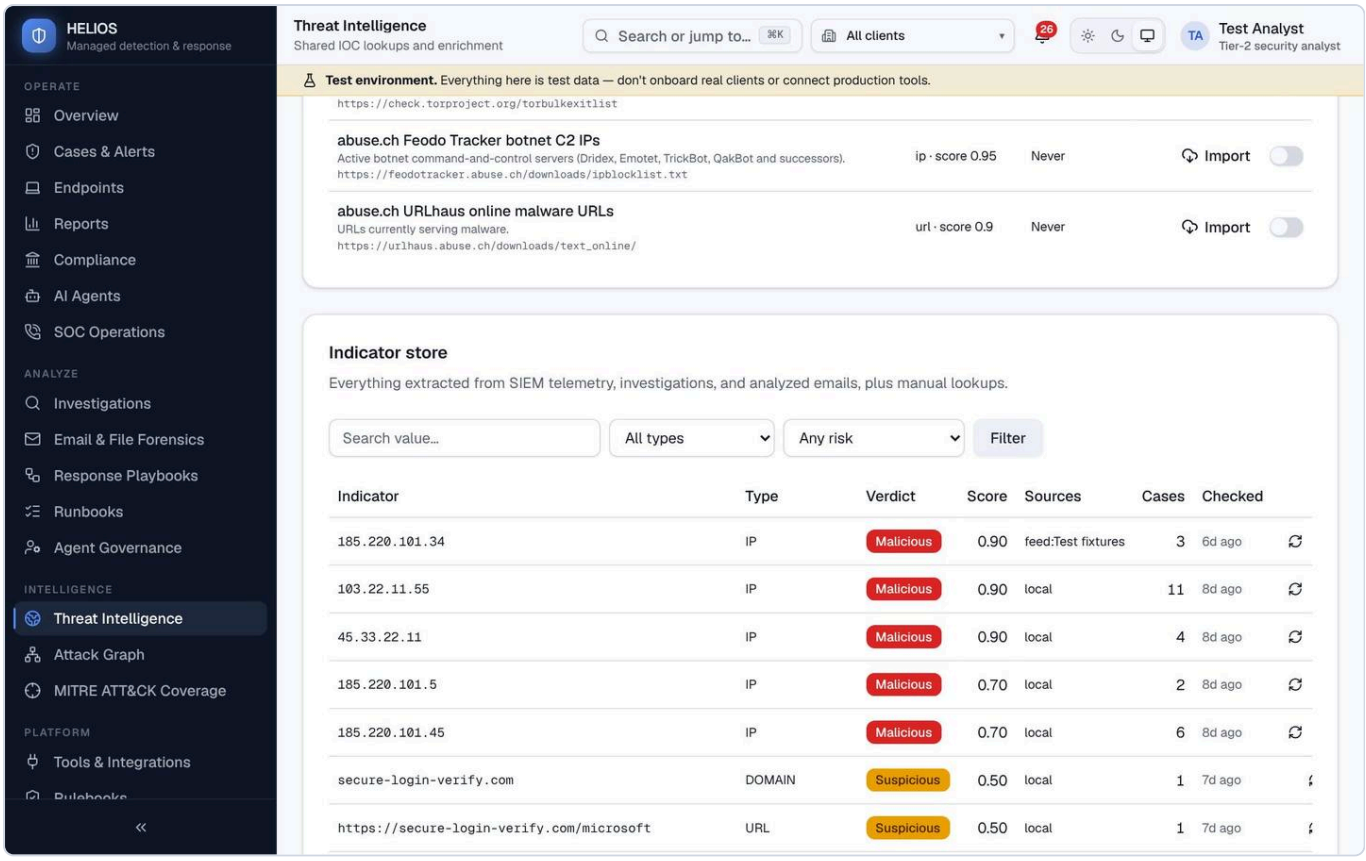


Figure 33. Threat Intelligence, further down the page.

Attack Graph

Cases joined by the hosts, accounts, addresses, indicators and agents they share. An entity is an attack

chain when its cases together reach more ATT&CK tactics than any one alone.

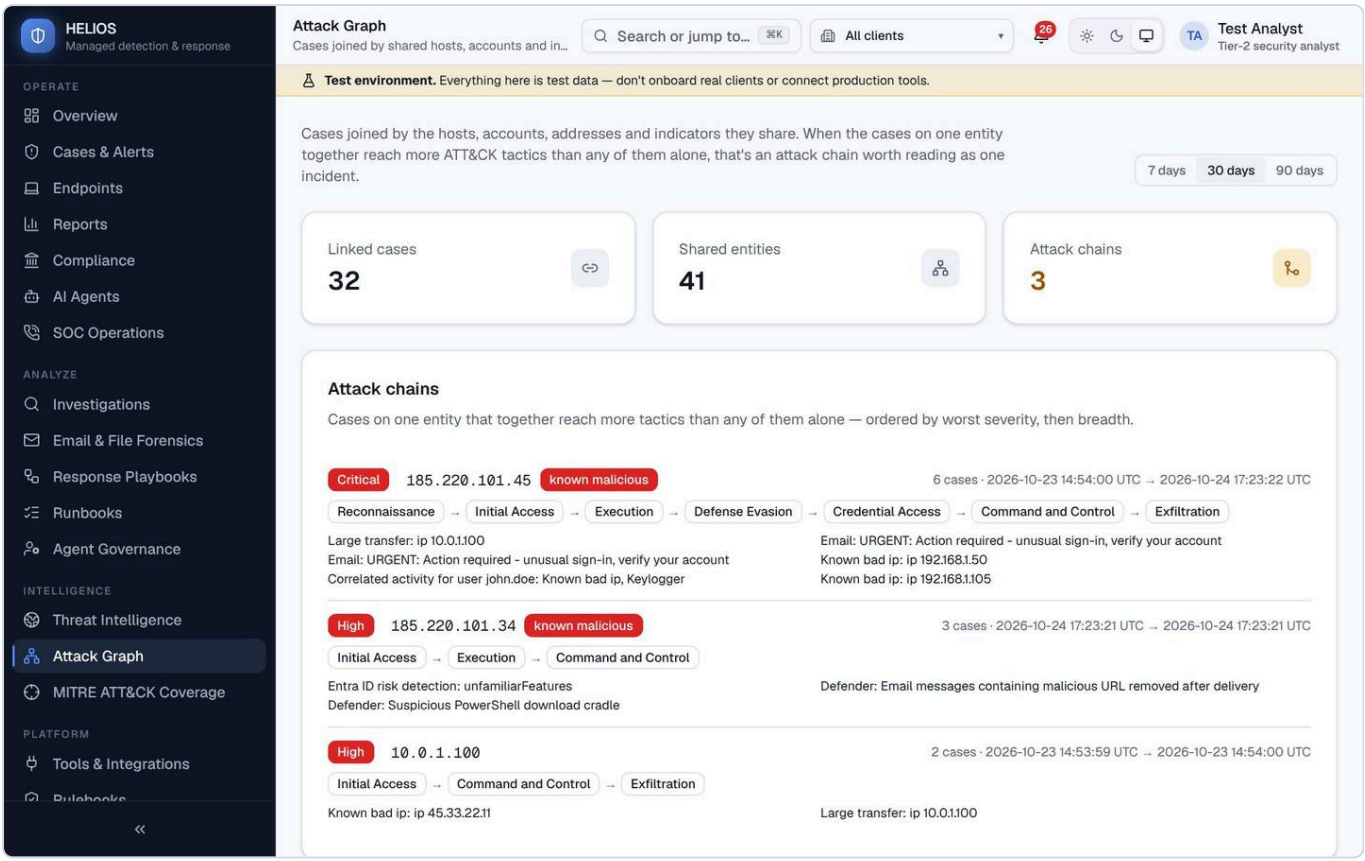


Figure 34. Attack Graph, as seen by a tier-2 security analyst (test environment).

How to use it

- Hover or click a node to see the cases behind it; follow chains across tactics to see an intrusion that several smaller cases make up.

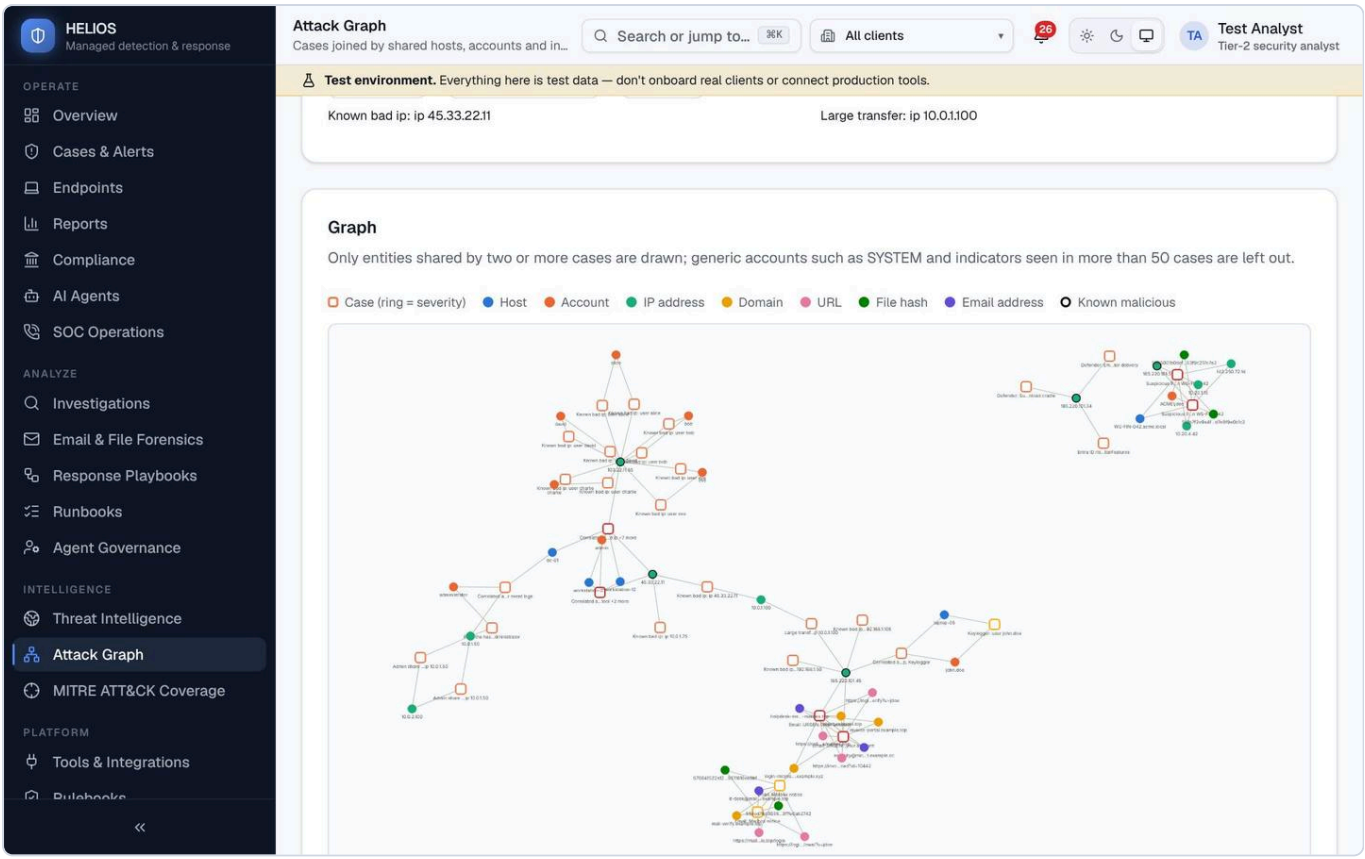


Figure 35. Attack Graph, further down the page.

MITRE ATT&CK Coverage

Which ATT&CK techniques HELIOS can detect, across built-in detectors, Sigma rules and onboarded

tools, and which it has seen.

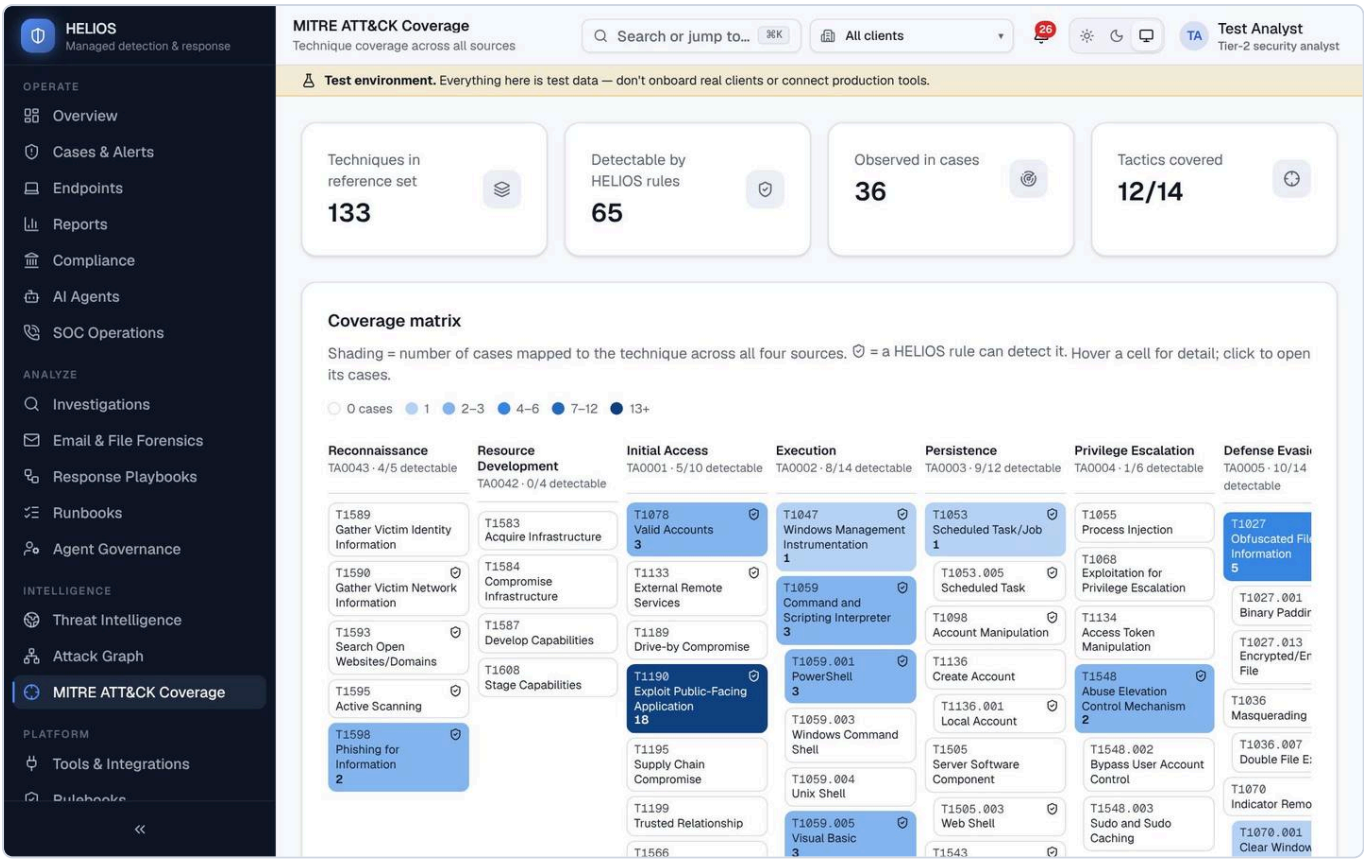


Figure 36. MITRE ATT&CK Coverage, as seen by a tier-2 security analyst (test environment).

How to use it

- Read coverage by tactic and technique; gaps show where detection is missing.

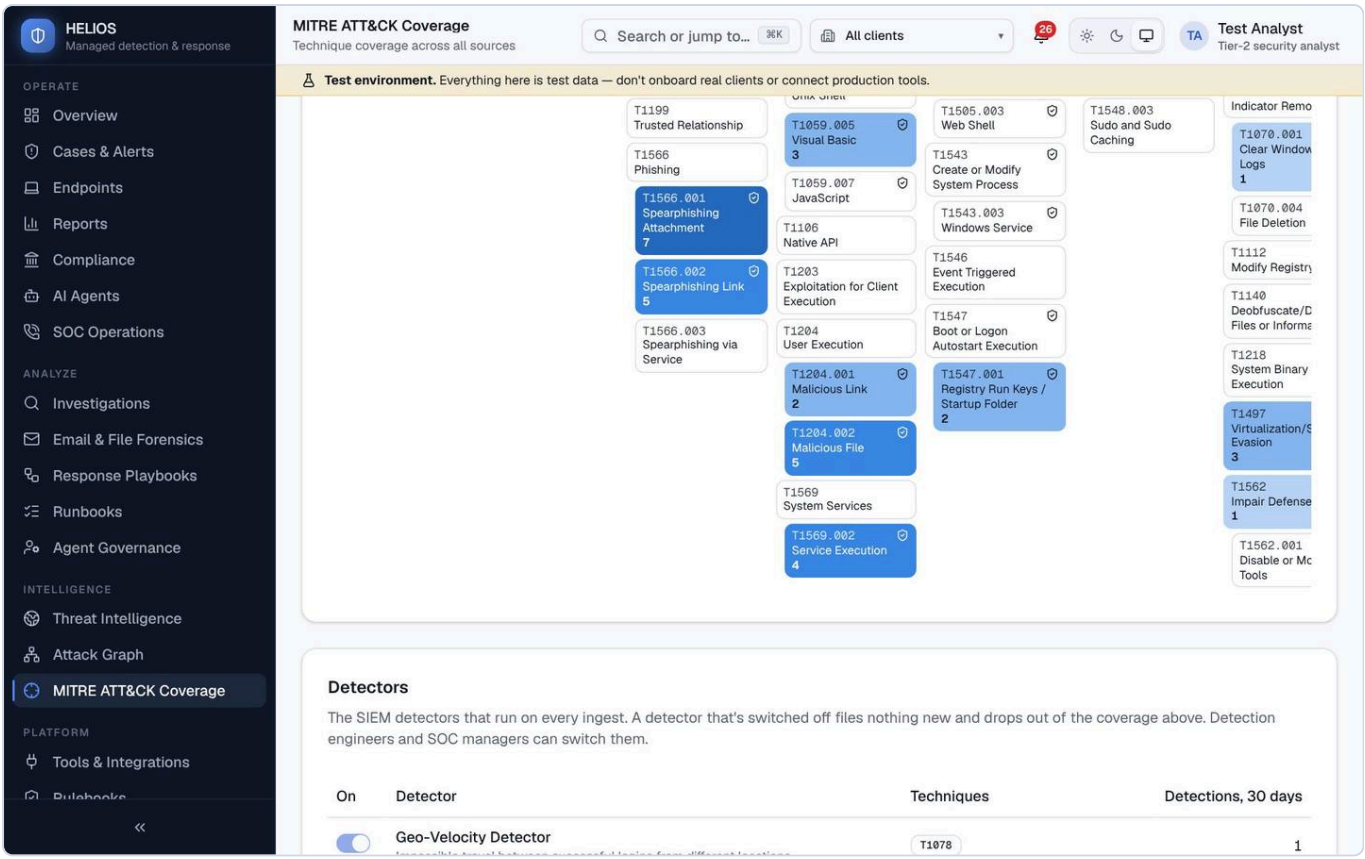


Figure 37. MITRE ATT&CK Coverage, further down the page.

Tools & Integrations

The security tools onboarded for each client: how their logs reach HELIOS and how HELIOS acts through

them.

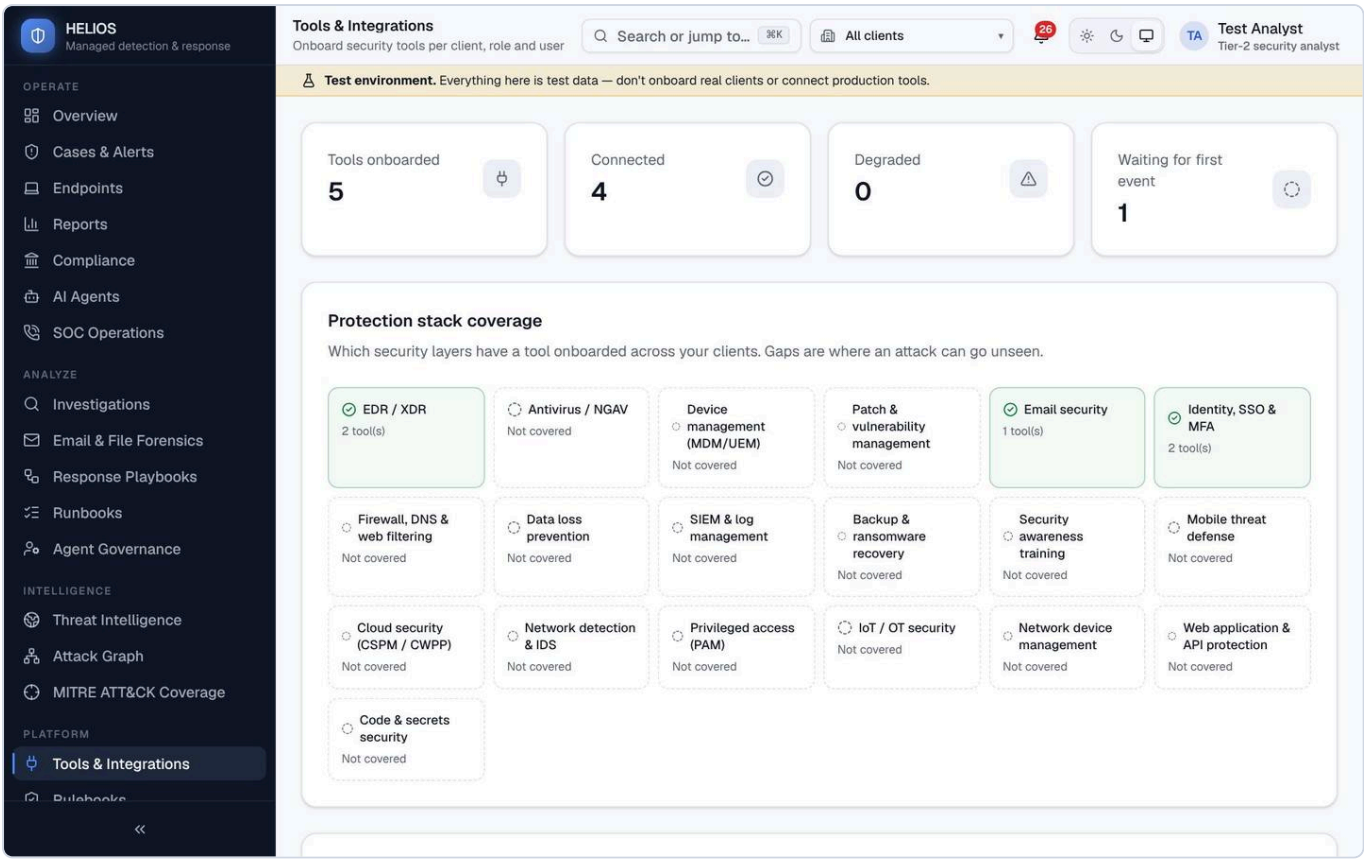


Figure 38. Tools & Integrations, as seen by a tier-2 security analyst (test environment).

How to use it

- The tiles count tools onboarded, connected, degraded and waiting for their first event. Protection stack coverage shows which security layers have a tool across your clients; gaps are where an attack can go unseen.
- Each tool shows how it sends logs (push to its ingest URL, HELIOS pulling from the vendor's API, or the HELIOS forwarder), its source health, and whether a response connector can act through it.
- A source silent longer than its limit opens a case: attackers often stop logging first. A client can't go live while a source is dead.

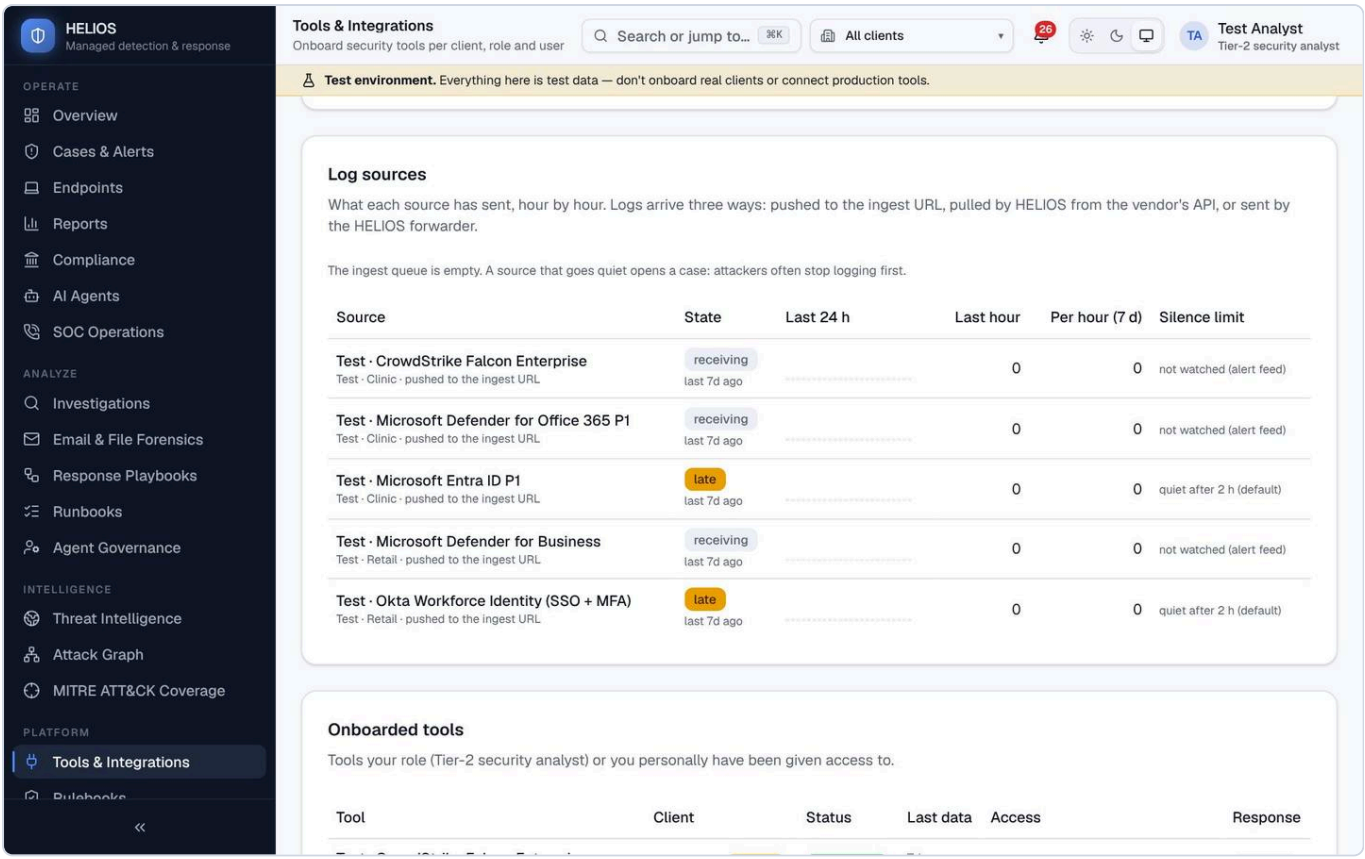


Figure 39. Tools & Integrations, further down the page.

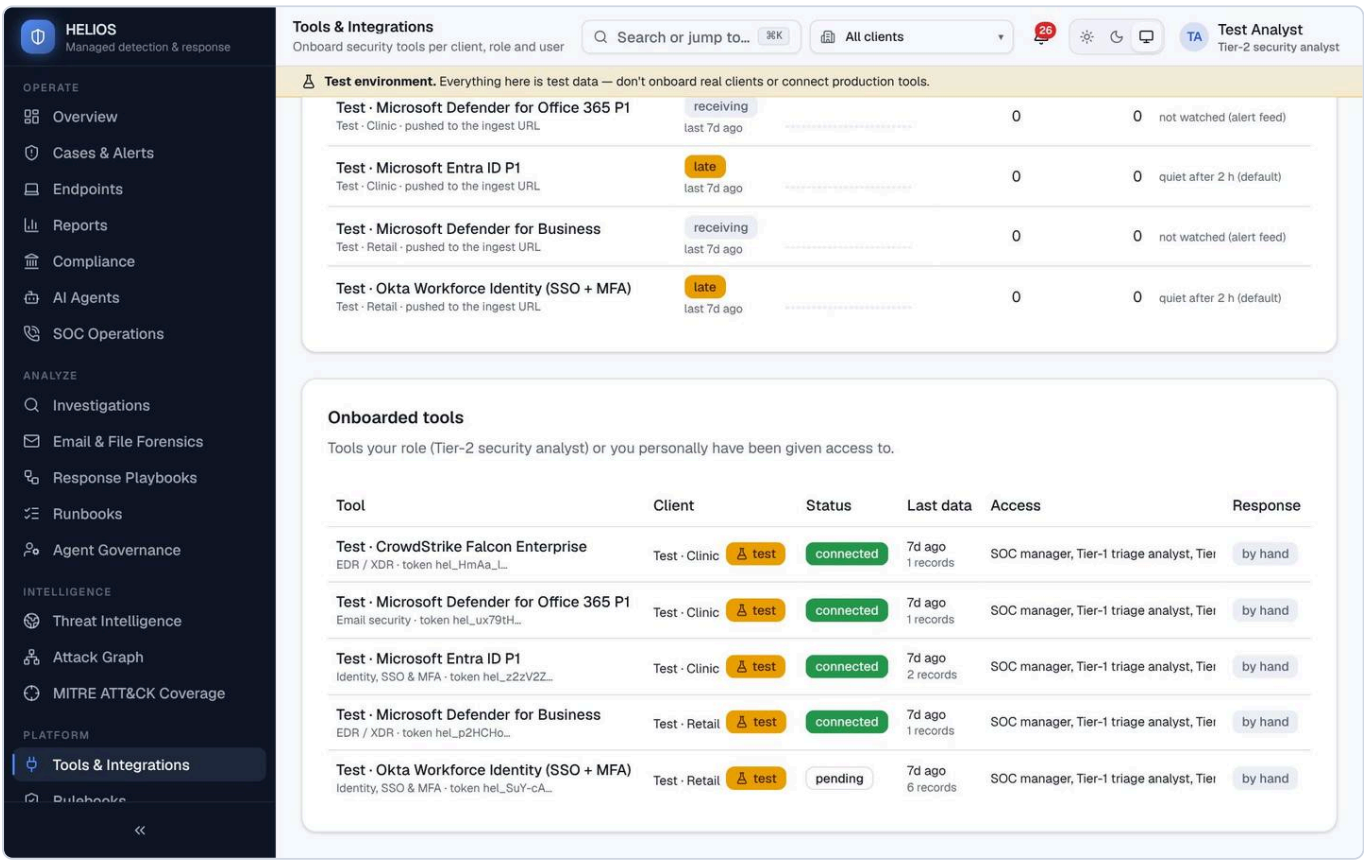


Figure 40. Tools & Integrations, further down the page.

Rulebooks

Each client's signed rules for what HELIOS may do on its estate. The rules only ever restrict: protected assets and accounts, which actions run automatically, need approval or never happen, blast-radius limits, business hours and change freezes, the approver chain, AI data handling and CERT-In contacts.

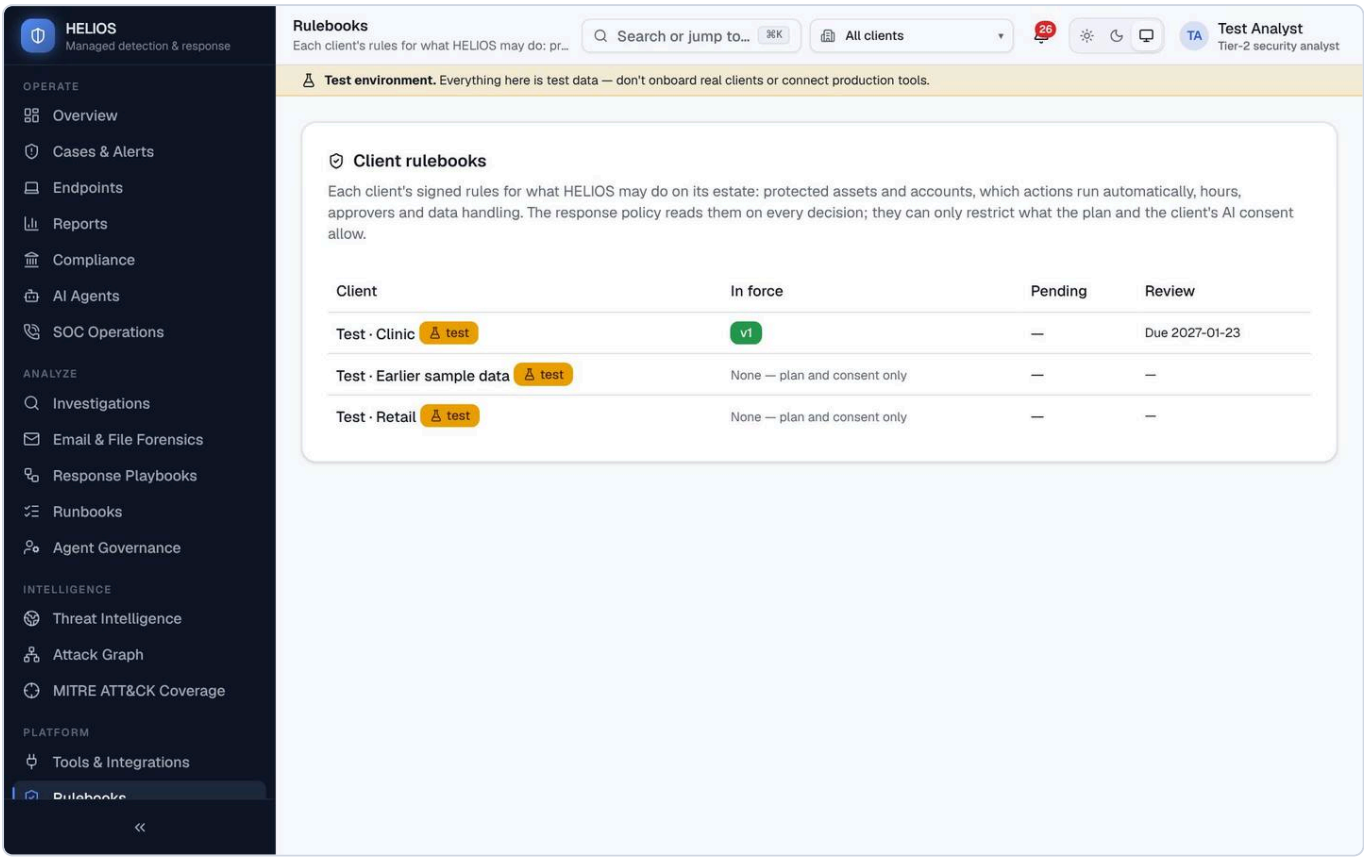


Figure 41. Rulebooks, as seen by a tier-2 security analyst (test environment).

How to use it

- Pick the client at the top. The page shows the version in force, its hash and when it is due for review.
- The tabs: In force (the rules now applying), Edit, Dry run, Allowlist, Versions, AI models, Autonomy and Policy as code.
- Autonomy: the AI's autonomy ladder per action type for this client, and promotions waiting for approval.

Clients

The client organisations (tenants): their plan and SLA, status, contacts, and how far they consent to AI

acting on their estate.

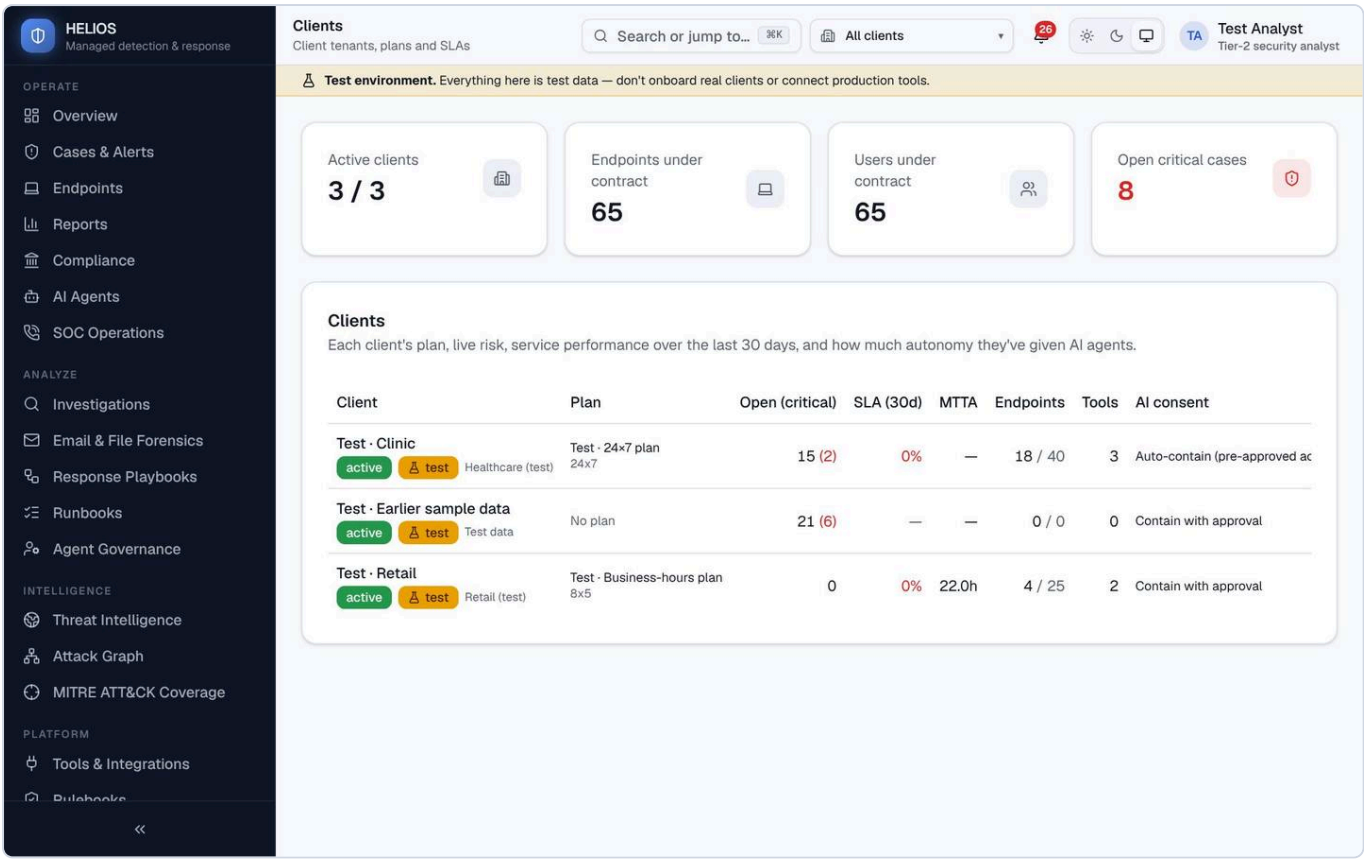


Figure 42. Clients, as seen by a tier-2 security analyst (test environment).

How to use it

- Each client shows its plan, status (onboarding, active, suspended), licensed endpoints and users, and test clients are badged.

API Tokens

Personal tokens for the REST API, for ticketing systems, SOAR and client portals. A token acts as you:

same role, client and queues.

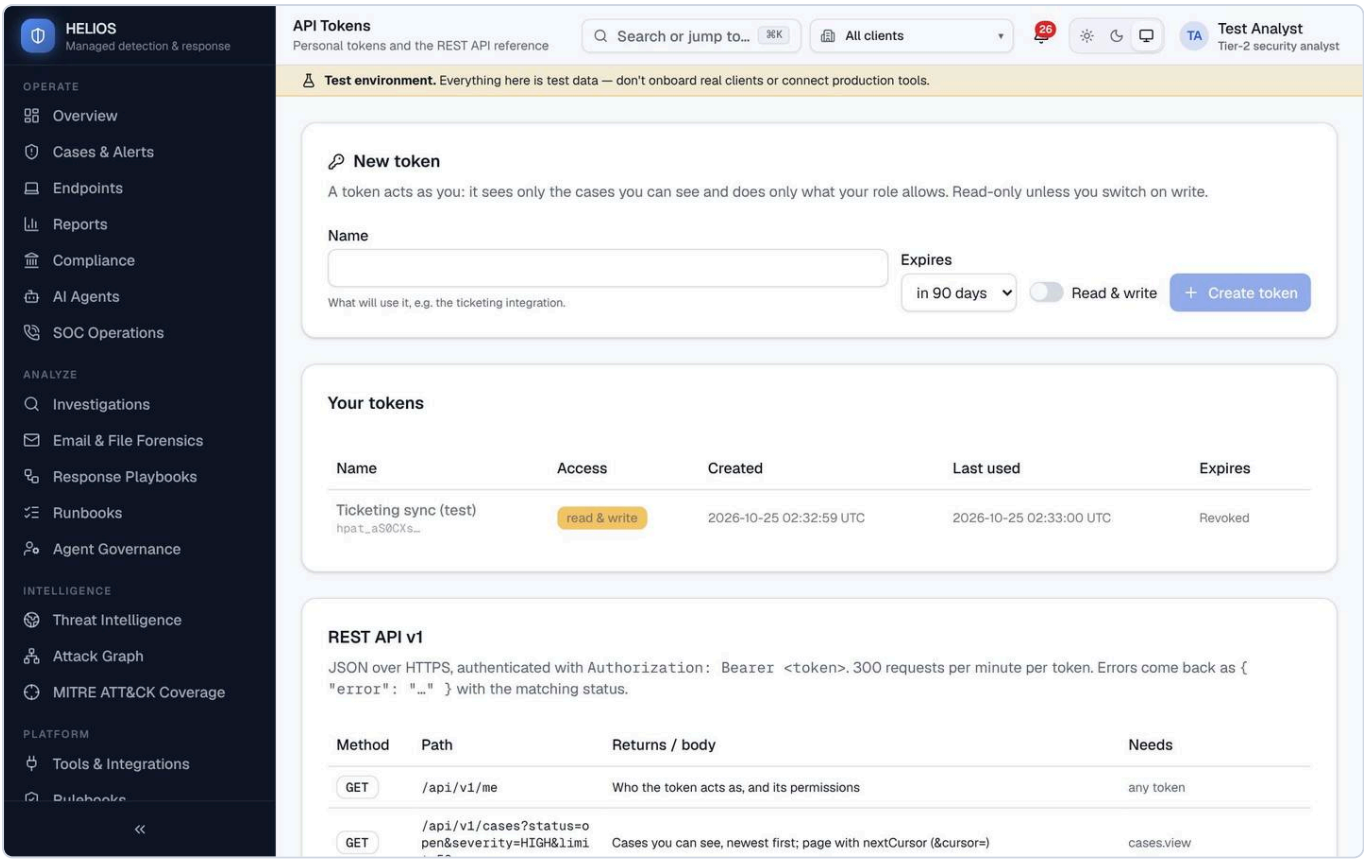


Figure 43. API Tokens, as seen by a tier-2 security analyst (test environment).

How to use it

- New token: give it a name (what will use it), an expiry (90 days by default), and switch on Read & write only if it must change things. Create token shows it once; copy it then. Your tokens lists yours to revoke.
- The page includes the API reference: endpoints, parameters and examples. Calls are limited to 300 a minute per token, and every change through the API is audited.

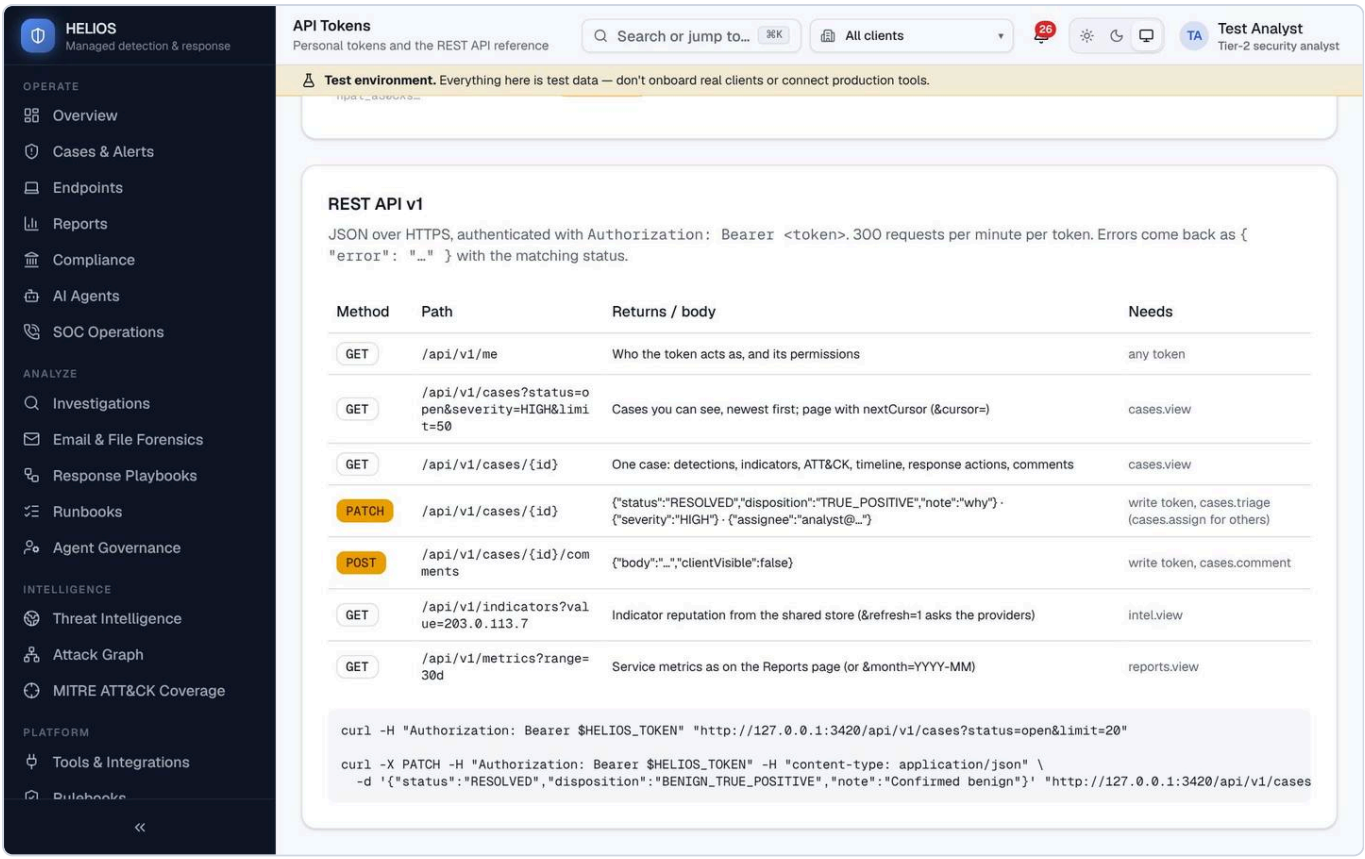


Figure 44. API Tokens, further down the page.

A client's rulebook

The rulebook page for one client (here the test client Clinic).

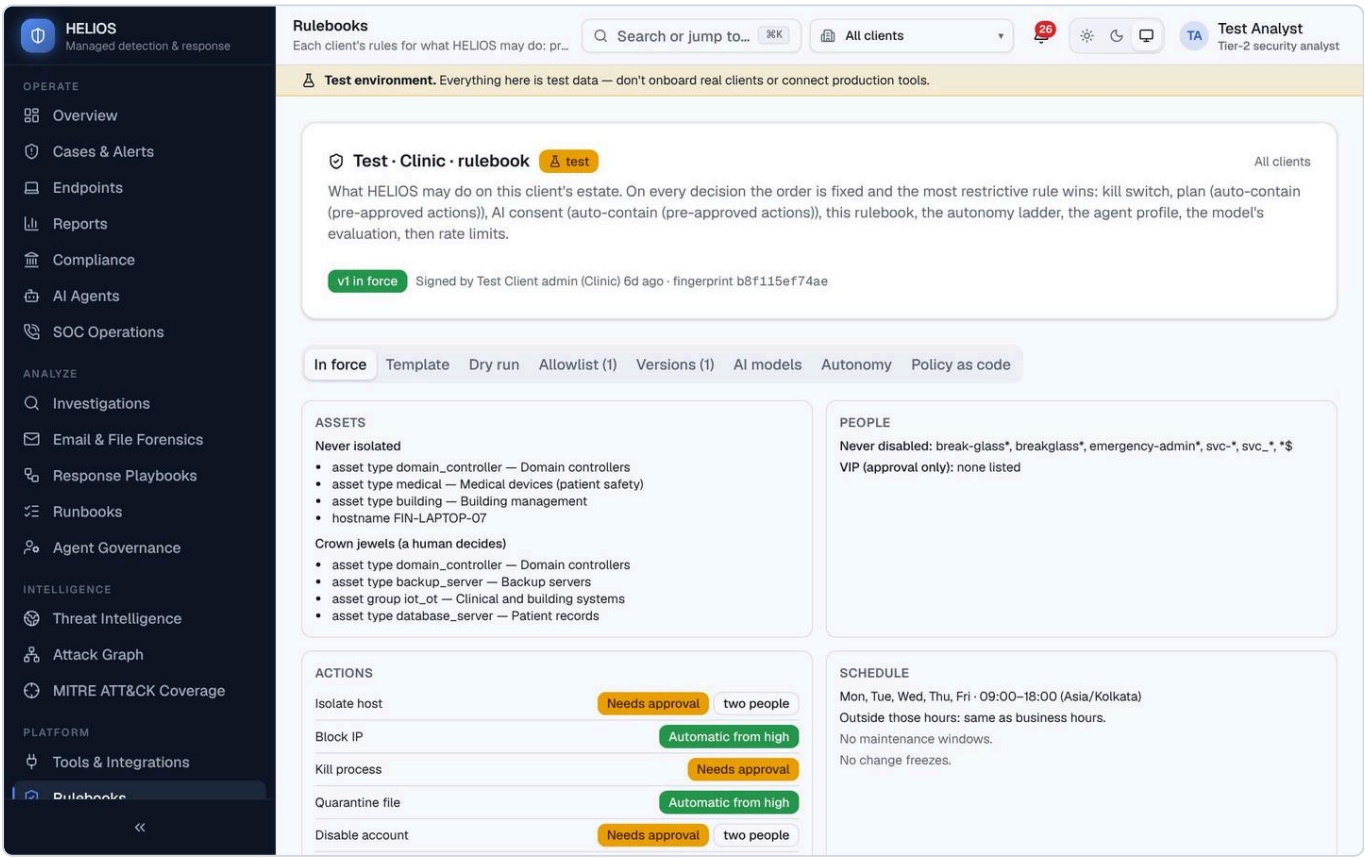


Figure 45. A client's rulebook, as seen by a tier-2 security analyst (test environment).

How to use it

- In force shows the assets (never isolated, crown jewels), the people (never disabled, VIPs), each action's mode (automatic from a severity, needs approval, two people), the schedule, the approver chain, and data handling and obligations. The other tabs are as described above.

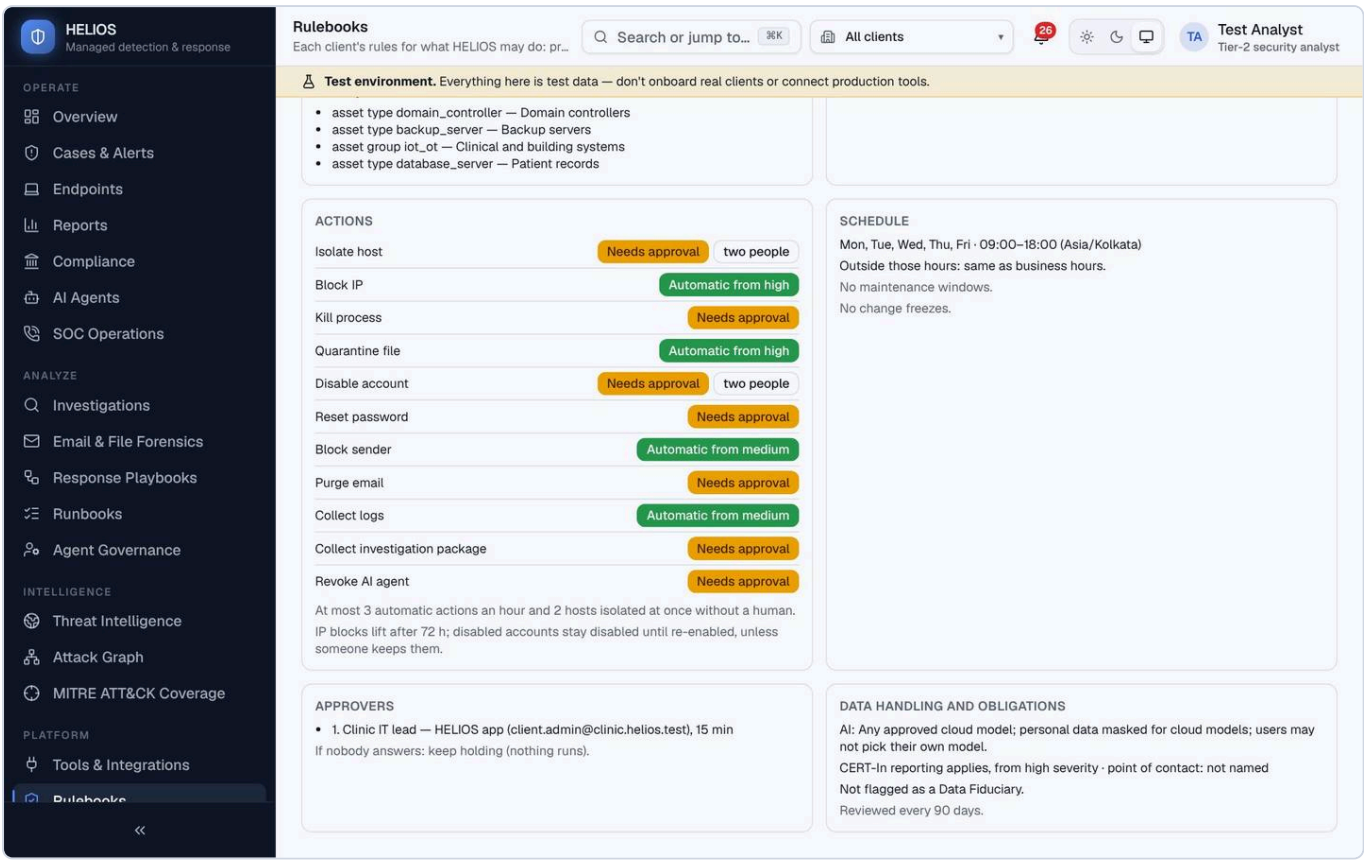


Figure 46. A client's rulebook, further down the page.

A client

One client's page: plan and SLA, contacts, consent, onboarding status, tools, users and data keys.

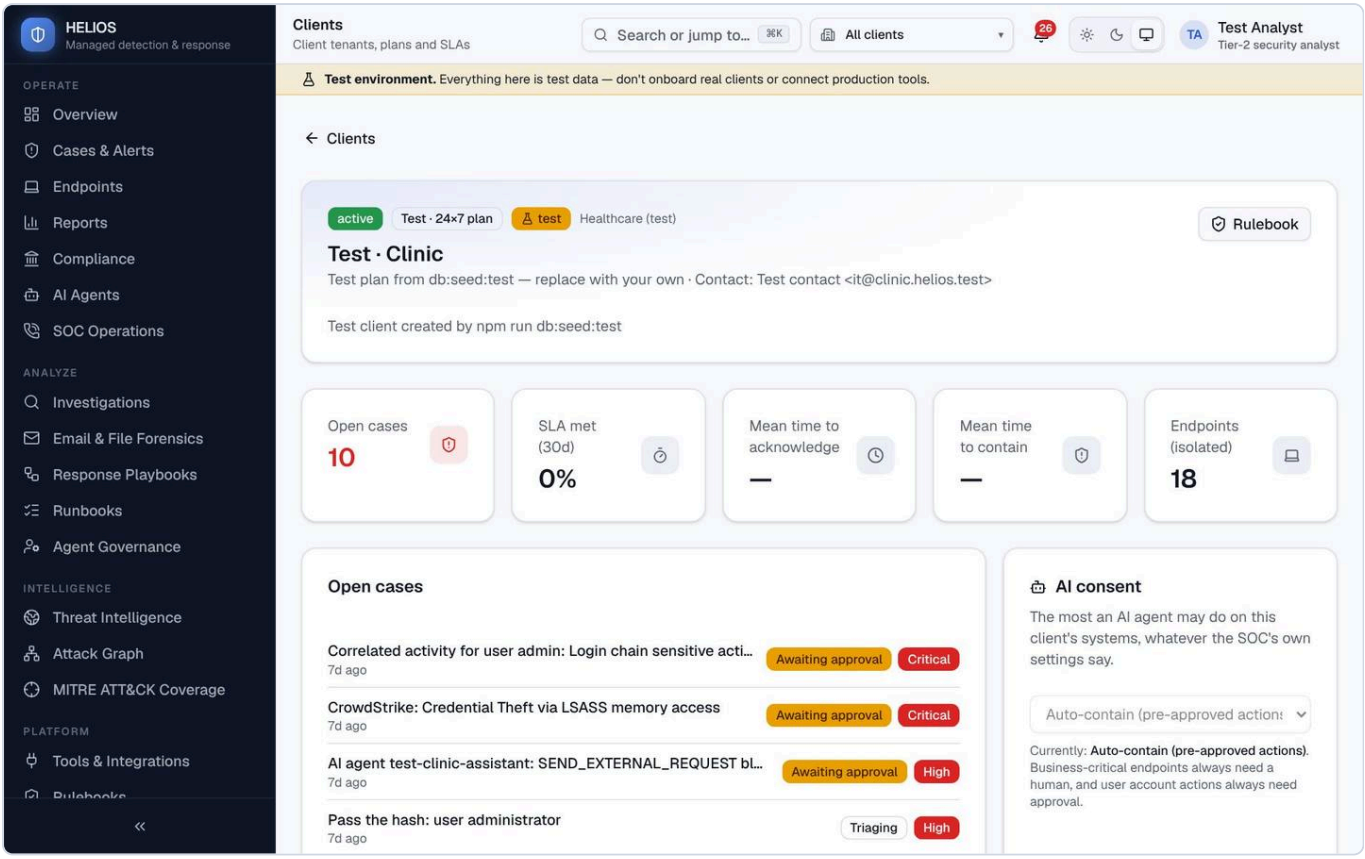


Figure 47. A client, as seen by a tier-2 security analyst (test environment).

How to use it

- The page shows the client's open cases, SLA met, mean time to acknowledge and contain, endpoints, and buttons for its Rulebook, Review upgrades & price, and Edit client.

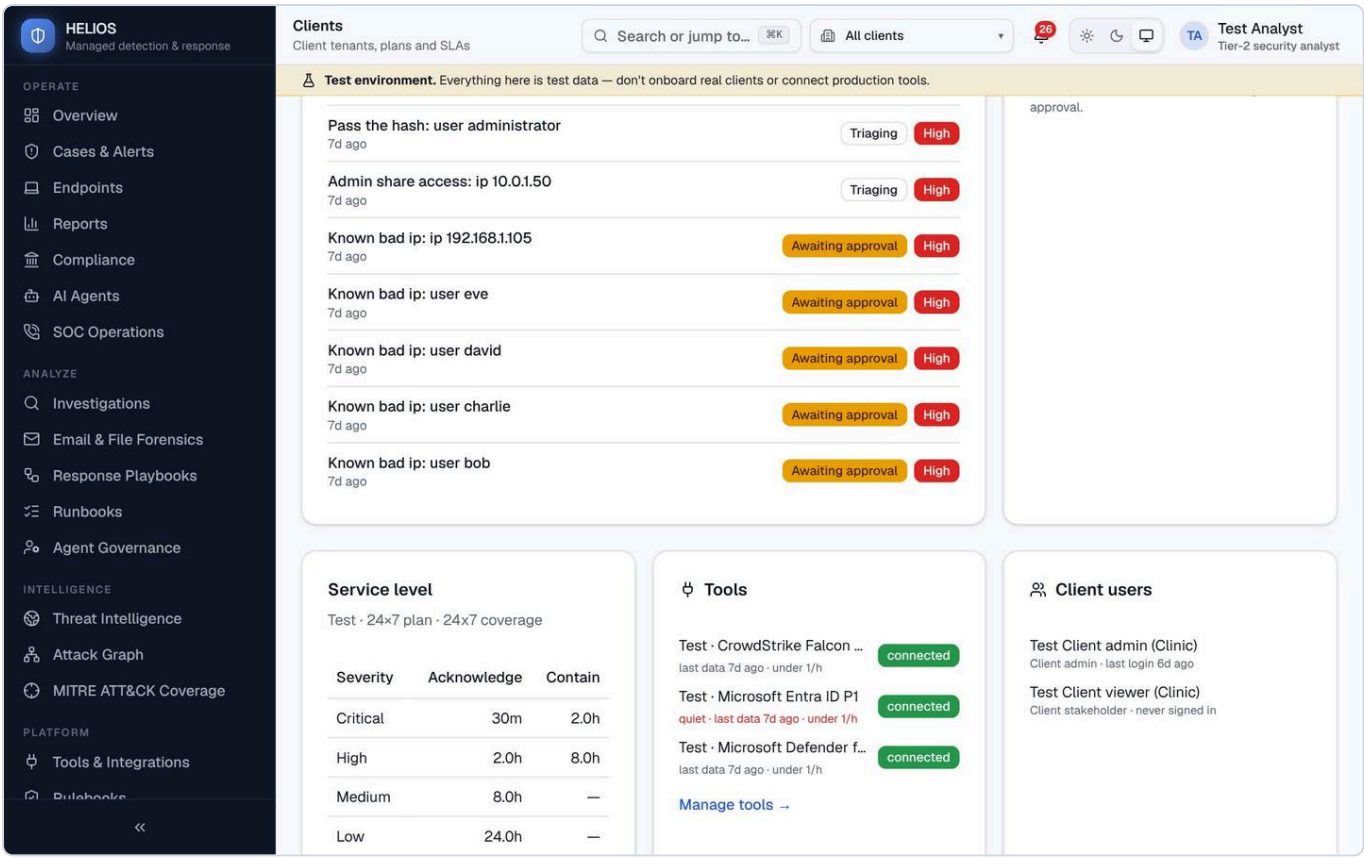


Figure 48. A client, further down the page.

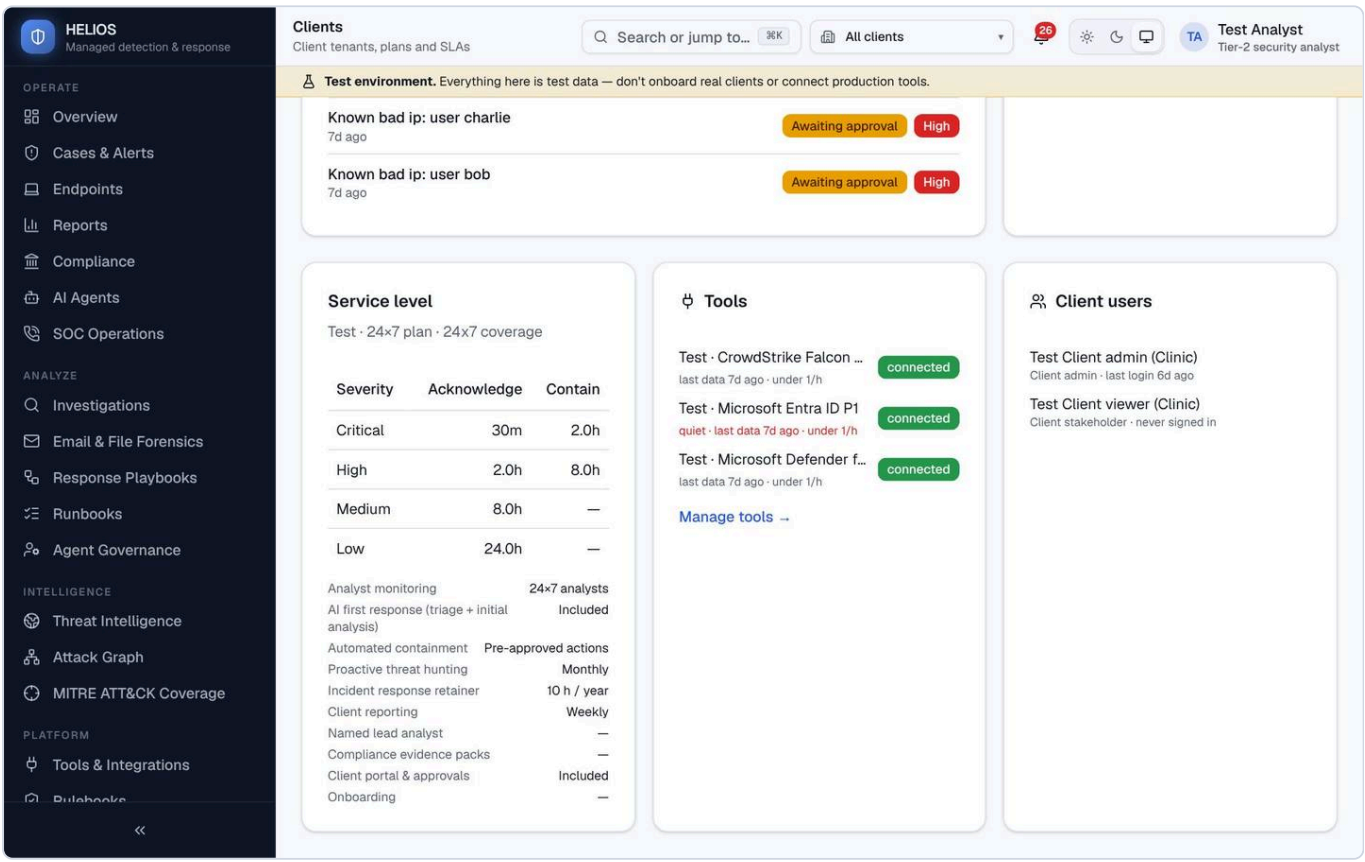


Figure 49. A client, further down the page.

Daily digests

Once a day, what HELIOS did on its own, what was waiting, and the deadlines coming up.

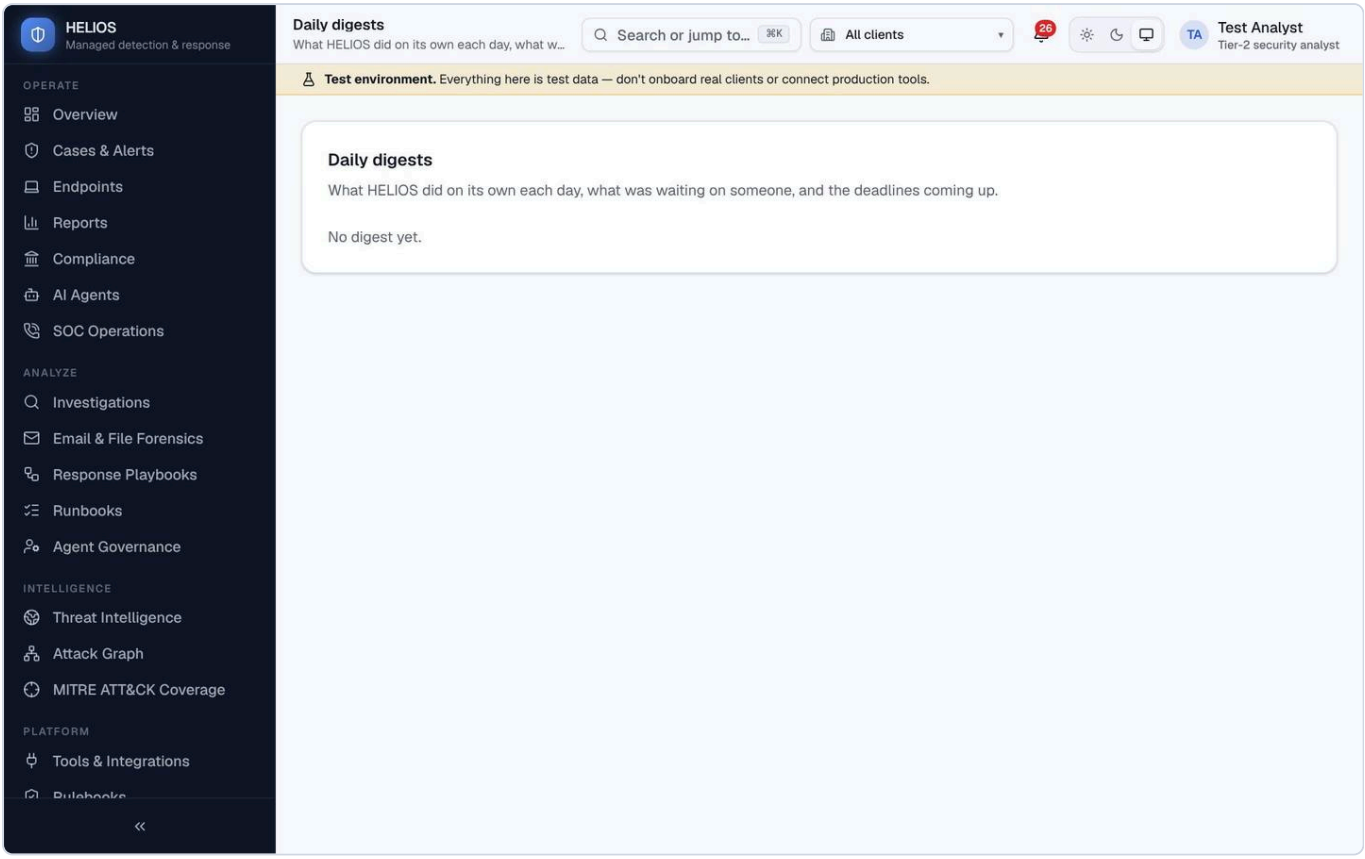


Figure 50. Daily digests, as seen by a tier-2 security analyst (test environment).

How to use it

- Each digest lists actions the AI took without asking (with its autonomy level, rollbacks and QA calls), cases opened and closed, containment waiting for approval, regulatory deadlines within a day, quiet sources and the week's QA results.
- Choose under My alerts whether it also reaches you by email.

My alerts

Which notifications also reach you outside HELIOS, by email, SMS or WhatsApp.

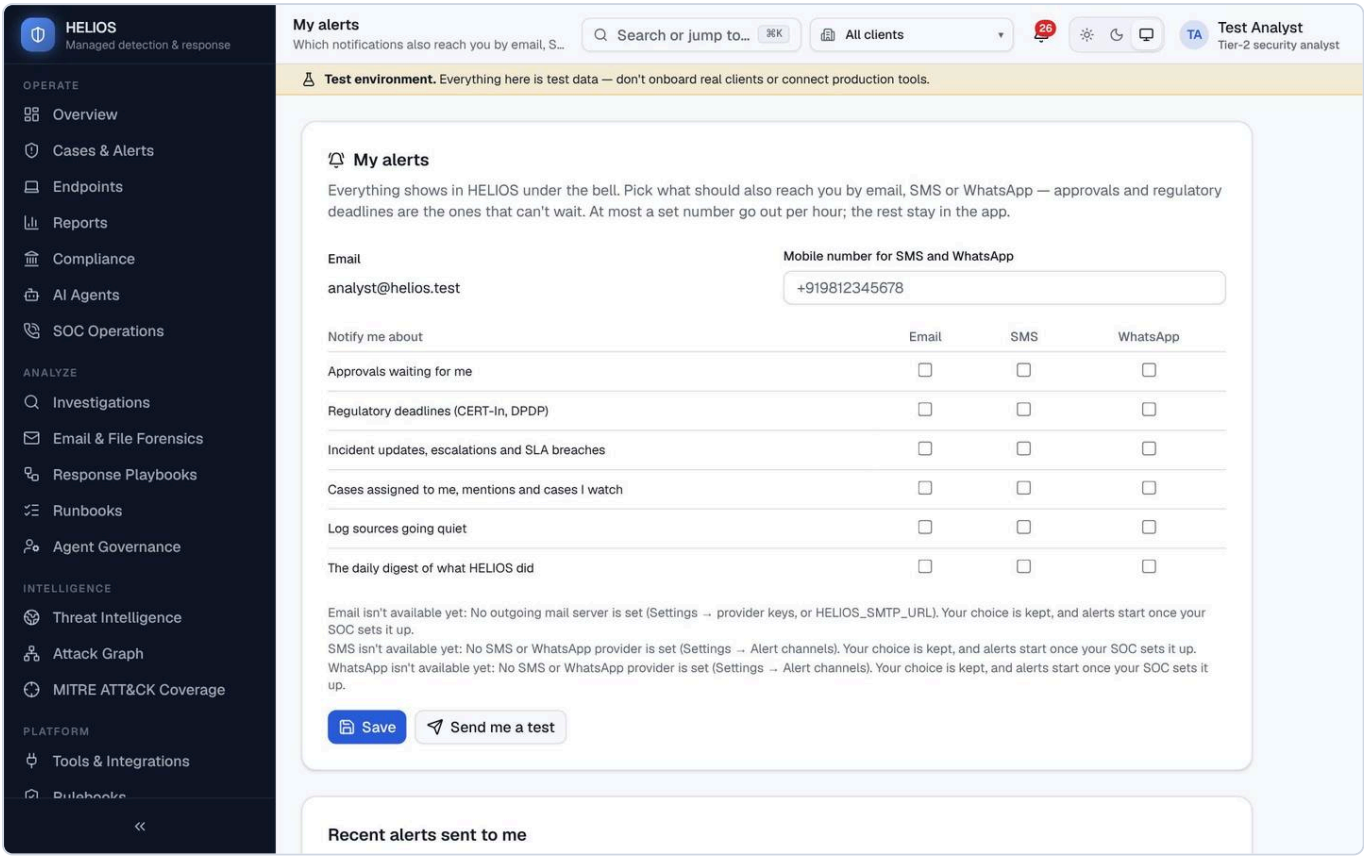


Figure 51. My alerts, as seen by a tier-2 security analyst (test environment).

How to use it

- Every notification lands in the app (the bell at the top). Here choose, per kind of notification, the channels it also comes by, and give your phone number in international format (+91…).
- Deliveries that couldn't go (a channel not set up, no number, over the hourly cap) are listed with the reason.

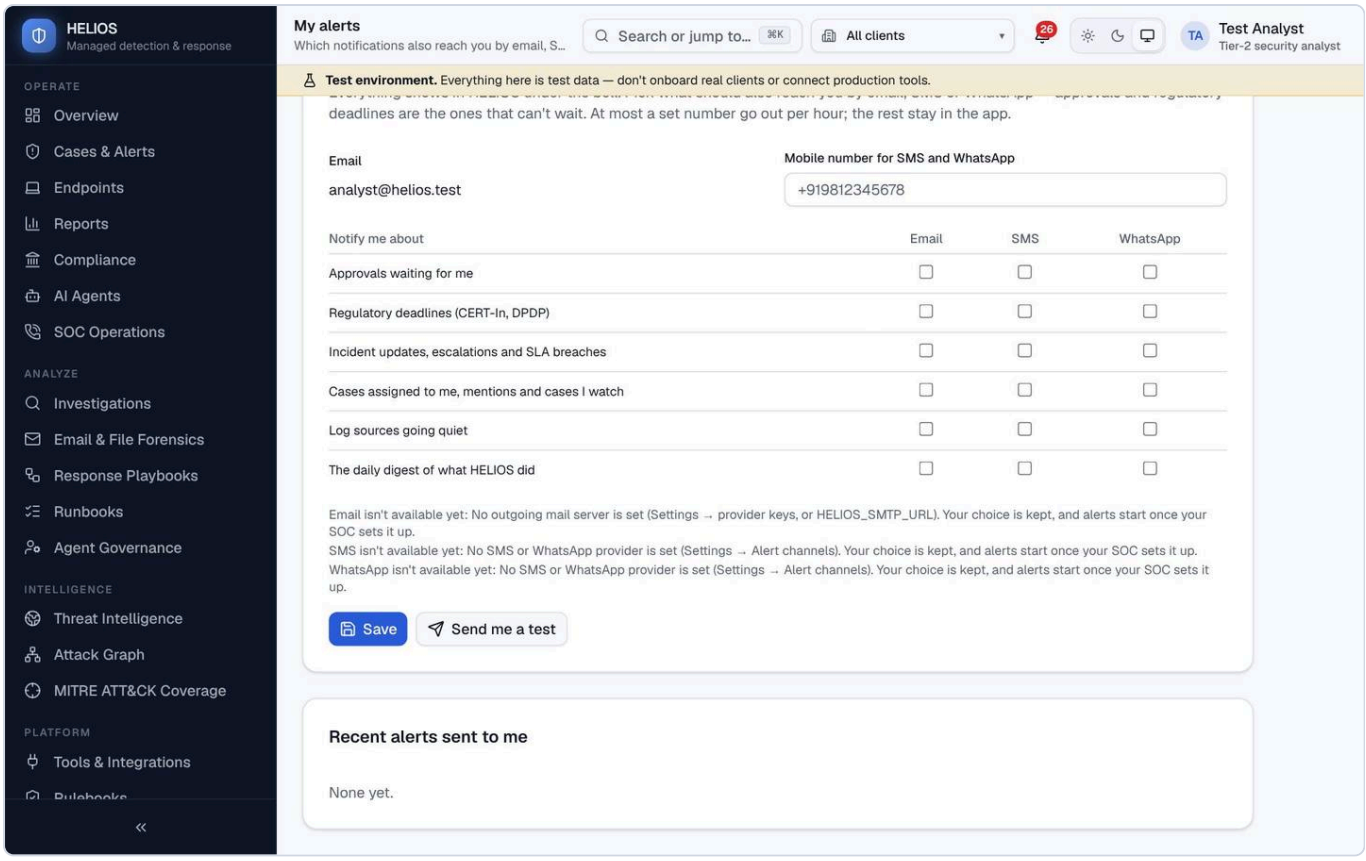


Figure 52. My alerts, further down the page.

5. Procedures

Step by step, for the tasks your role does most.

5.1 Investigate an escalated case

1. Read the hand-off note and the AI first response. Acknowledge if it isn't.
2. Dispatch the investigation agent (Dispatch agent → Investigation agent): it searches the client's telemetry around the case with read-only tools and reports a timeline tied to evidence.
3. Check related cases and the Attack Graph for the same hosts, accounts and indicators elsewhere.
4. Work the attached runbook steps.
5. Record findings in internal comments as you go.

5.2 Investigate uploaded logs

1. Open Investigations → upload the logs (EVTX, CSV, JSON, spreadsheets or text), with the alert that prompted it and the client.
2. Read the pipeline's root cause, its adversarial QA and the timeline; open the case it created.

5.3 Analyse a suspicious email or file

1. Open Email & File Forensics and submit the .eml, .msg or file.
2. Read the header authentication, the sender-domain diagnosis, the links (detonated through the proxy), attachments and macros.
3. If it is malicious, the case lets you request Purge email (each recipient's copy moved to Recoverable Items) and Block sender.

5.4 Contain an account or a device

1. Request containment: Disable account or Reset password (a random password HELIOS doesn't keep, to change at next sign-in, with sessions ended); Isolate host; Collect investigation package (the device's forensic state copied to Defender).
2. Where the client has a native connector, HELIOS calls the vendor and counts the action only once the vendor confirms it. Without one, it becomes a task: do it in the tool's console, then record What you did.
3. Containment can be rolled back from the case; IP blocks and account disables can be time-boxed or kept until lifted.

5.5 Tell the client

1. Write a client-visible comment, or Share the case.
2. Dispatch the client report writer for a plain-language draft, edit it, and publish it to the client.
3. If CERT-In or DPDP applies, start the notice from the case's Regulatory reporting card and send it for the client's approval.

6. Guidelines

Do

- Tie every finding to evidence; the AI and the reports do the same.
- Prefer the least disruptive action that stops the threat.
- Record why you took each action: the approver and the auditor read it.
- Keep notices factual and send them for approval early: CERT-In's clock is 6 hours from when the incident was noticed.

Don't

- Don't share AI working notes or internal comments with the client.
- Don't detonate or open suspicious content on your own machine: use Forensics.
- Don't block shared infrastructure (a cloud provider's range, a mail service) without the incident responder's agreement.

For everyone

- Never share your account or password; every action is recorded under your name.
- Keep client information inside HELIOS and to the people who need it.
- Treat text inside evidence (command lines, emails, file names) as possibly written by an attacker: it may address 'the AI' or 'the analyst' to mislead. HELIOS flags such text; don't follow it.
- Every action in HELIOS is reversible or recorded, and audited; write down why you did it.

> The client's signed rulebook binds everyone: the AI, playbooks and people. If HELIOS refuses an action, the reason names the rule. Don't work around it in a vendor's console; ask the client to change the rule.

7. Reference

Terms you will see

Case statuses

Open	New; nobody has acknowledged it yet.
Triaging	Someone is looking at it.
Investigating	Under analysis.
Awaiting approval	Containment is waiting on someone's yes.
Contained	The threat is stopped; clean-up may continue.
Resolved	Closed with a disposition.
False positive	Closed: it wasn't a threat.

Severities

Critical, High, Medium, Low, Info	Set by the detections and adjustable by analysts. The client's plan sets the acknowledgement SLA for each.

Dispositions (when closing)

True positive	A real threat.
Benign true positive	Real activity, but expected or authorised (an admin's tool, a pen test).
False positive	Not what the detection thought.

Tiers

T1	Triage.
T2	Analysis.
T3	Incident response.

Response action statuses

Suggested	A recommendation only; autonomy doesn't allow acting.
Pending approval	Waiting on an approver.
In progress	The vendor was asked; waiting for it to confirm. Doesn't count as containment yet.
Executed	Done and confirmed by the vendor (or by an analyst for a manual action).
Manual required	Approved, but no connector can do it: someone does it in the tool's console and records it.
Blocked	A guardrail or the client's rulebook refused it; the reason is shown.
Rejected	An approver said no.
Rolled back	Undone.

Failed	The vendor or HELIOS couldn't do it.
Shadow	Shadow mode: what the AI would have done, recorded silently and shown once the case closes.

AI autonomy (per agent)

Off	The agent doesn't run.
Suggest only	It proposes; people decide.
Contain with approval	Its containment waits for approval.
Auto-contain	It may act on pre-approved actions above its confidence bar.

Autonomy ladder (per client and action type)

L0 observe	Shadow mode.
L1 suggest	Suggestions only.
L2 approve	Acts once approved.
L3 act and notify	Acts, and tells people at once.
L4 act	Acts, and reports in the daily digest.

Evidence references

D, E, I, T, C	Detections, events, indicators, timeline entries and context in the case; Q for what the investigation agent's tools returned. Every AI finding cites them.

If something doesn't work

- **A page says you don't have access**: your role doesn't include it. Ask your administrator if your job needs it.
- **An action was refused**: the message names the rule (the client's rulebook, a guardrail, the plan or the kill switch).
- **An action is 'Manual required'**: no connector can do it. Do it in the tool's own console, then record what you did on the case.
- **You don't get alerts by SMS or email**: check My alerts for your channel choices, your phone number, and deliveries skipped with their reason.
- **You can't sign in**: after five failures the account locks for 15 minutes; your administrator can reset your password or two-step sign-in.